Microsoft Security Copilot

Use Security Copilot to bring the full power of OpenAI architecture to defend your organization at machine speed and scale.

About Security Copilot

OVERVIEW

What is Microsoft Security Copilot?

Where is Security Copilot experienced?

Get started

GET STARTED

Get started with Security Copilot

Using prompts in Security Copilot

Security Copilot in your workflows

Try promptbooks

Security Copilot in your workflows

CONCEPT

Explore capabilities in Security Copilot

Discover agents

Microsoft Security Copilot agents

Partner-built agents

Security store

Build custom agents

Plugins

Connectors

Promptbooks

Understand Security Copilot concepts

CONCEPT

Prompting tips

Navigation tips

Use cases

TUTORIAL

Triage alerts with enriched threat intelligence

Investigate an incident's malicious script

How-to guides

HOW-TO GUIDE

Navigate Microsoft Security Copilot

Configure owner settings

Manage plugins

Build your own promptbook

Access the Security Copilot audit log

Use the Admin Activity Export API

Security Copilot experiences

REFERENCE

Security Copilot in Microsoft Defender XDR

Security Copilot in Microsoft Entra

Security Copilot in Microsoft Intune

Security Copilot in Microsoft Purview

Security Copilot integration in Azure Firewall

Security Copilot integration in Azure Web Application Firewall

Security Copilot in Microsoft Sentinel

Security Copilot integration in Defender External Attack Surface Management

Security Copilot in Microsoft Defender for Cloud and Cloud Security Posture Management

Security Copilot for Threat Intelligence

Reference

REFERENCE

Data security and privacy

Frequently Asked Questions

Application card for Security Copilot

Application card for Security Copilot agents

Learn live - Get started with Microsoft Security Copilot

Microsoft Security Copilot Adoption hub

Microsoft Security Copilot video hub

Developer content

REFERENCE

Custom agent overview

Build an agent using natural language

Build an agent using the agent builder form

Build an agent by uploading a YAML

Build an agent using a manifest

Build an agent using MCP

Extend your agent with tools

Build promptbooks

What is Microsoft Security Copilot?

Microsoft Security Copilot (Security Copilot) is a generative AI-powered security solution that helps increase the efficiency and capabilities of defenders to improve security outcomes at machine speed and scale.

Security Copilot provides a natural language, assistive copilot experience. Security Copilot helps support security professionals in various end-to-end scenarios such as incident response, threat hunting, intelligence gathering, posture management, and more. For more information, see Security Copilot primary use cases.

Designed with integration in mind, Security Copilot offers a standalone experience and also seamlessly integrates with products in the Microsoft Security portfolio. Security Copilot integrates with products such as Microsoft Defender XDR, Microsoft Sentinel, Microsoft Intune, Microsoft Entra, and other third-party services such as ServiceNow and Jamf. For more information, see Security Copilot experiences.

> **Tip**
>
> Visit the **Microsoft Security Copilot Adoption hub** to access useful links to training, videos, GitHub repository for sample plugins, and other technical readiness information.

The solution leverages the full power of OpenAI architecture to generate a response to a user prompt by using security-specific plugins, including organization-specific information, authoritative sources, and global threat intelligence. By using plugins as data point sources, security professionals have wider visibility into threats and gain more context. Incident responders also have the opportunity to extend the solution's functionalities. For more information about plugins, read Manage plugins.

> **Note**
>
> **Disclaimer:** This documentation is only intended for customers using commercial clouds. Currently, Security Copilot is not designed for use by customers using US government clouds, including but not limited to GCC, GCC High, DoD, and Microsoft Azure Government. For more information, consult with your Microsoft representative.

Security Copilot primary use cases

Security Copilot focuses on making the following use cases easy to accomplish.

- **Investigate and remediate security threats**: Gain context for incidents to quickly triage complex security alerts into actionable summaries and remediate quicker with step-by-step response guidance.
- **Build KQL queries or analyze suspicious scripts**: Eliminate the need to manually write query-language scripts or reverse engineer malware scripts with natural language translation to enable every team member to execute technical tasks.
- **Understand risks and manage security posture of the organization**: Get a broad picture of your environment with prioritized risks to uncover opportunities to improve posture more easily.
- **Troubleshoot IT issues faster**: Synthesize relevant information rapidly and receive actionable insights to identify and resolve IT issues quickly.
- **Define and manage security policies**: Define a new policy, cross-reference it with others for conflicts, and summarize existing policies to manage complex organizational context quickly and easily.
- **Configure secure lifecycle workflows**: Build groups and set access parameters with step-by-step guidance to ensure a seamless configuration to prevent security vulnerabilities.
- **Develop reports for stakeholders**: Get a clear and concise report that summarizes the context and environment, open issues, and protective measures prepared for the tone and language of the report's audience.
- **Build and add agents**: Developers can build and add agents to Security Copilot ecosystem.

Read Use cases for Security Copilot to delve deeper into the different security team roles like CISOs, threat intelligence analysts, IT admins, and the like, who can benefit from each highlighted use case.

How does Security Copilot work?

Security Copilot capabilities can be accessed through an immersive standalone experience and through intuitive embedded experiences available in other Microsoft security products. The foundation language model and proprietary Microsoft technologies work together in an underlying system that helps increase the efficiency and capabilities of defenders.

- **Microsoft security solutions** such as Microsoft Defender XDR, Microsoft Sentinel, and Microsoft Intune integrate seamlessly with Security Copilot. There are some embedded

experiences available in Microsoft security solutions that give access to Security Copilot and prompt capabilities in the context of their work within those solutions.

- **Plugins from Microsoft and third-party security products** are a means to extend and integrate services with Security Copilot. Plugins bring more context from event logs, alerts, incidents, and policies from both Microsoft security products and supported third-party solutions such as ServiceNow.

- **Security Copilot also has access to threat intelligence and authoritative content** through plugins. These plugins can search across Microsoft Defender Threat Intelligence articles and intel profiles, Microsoft Defender XDR threat analytics reports, and vulnerability disclosure publications, among others.

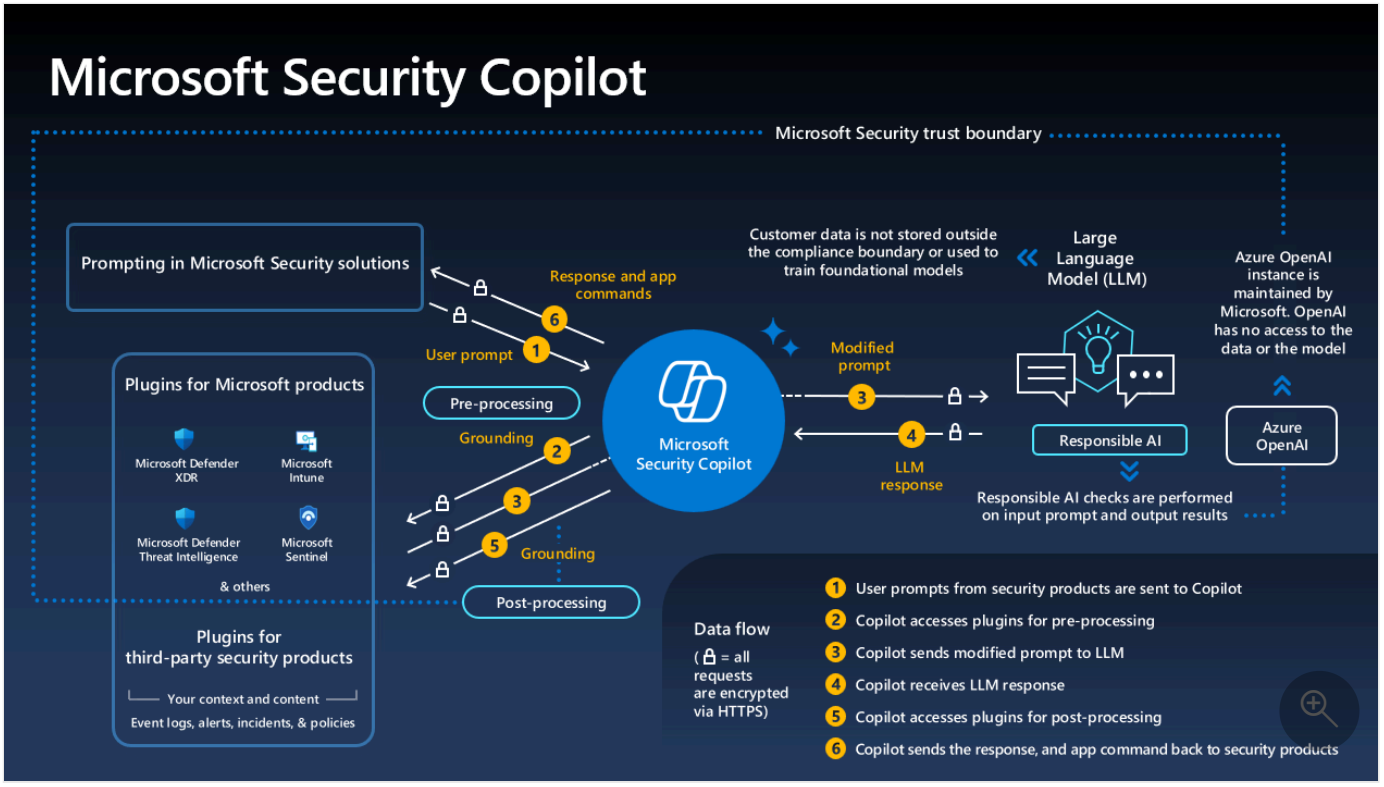

Here's an explanation of how Security Copilot works:

- User prompts from security products are sent to Security Copilot.

- Security Copilot then preprocesses the input prompt through an approach called grounding, which improves the specificity of the prompt to help you get answers that are relevant and actionable to your prompt. Security Copilot accesses plugins for preprocessing, then sends the modified prompt to the language model.

- Security Copilot takes the response from the language model and post-processes it. This post-processing includes accessing plugins to gain contextualized information.

- Security Copilot returns the response, where the user can review and assess the response.

Security Copilot iteratively processes and orchestrates these sophisticated services to help produce results that are relevant to your organization because they're contextually based on your organizational data.

Next steps

- Get started with onboarding to Microsoft Security Copilot

Security Copilot training

Go to Enhance security operations by using Microsoft Security Copilot to learn about Microsoft Security Copilot, an AI-powered security analysis tool that enables analysts to process security signals and respond to threats at a machine speed, and the AI concepts upon which it's built.

Security Copilot Customer Connection Program

Join the Security Copilot Customer Connection Program (CCP) to stay up to date with Security Copilot. CCP community members have access to:

- The latest technical product information and access to private previews.
- Free weekly technical trainings and product skilling webinars.
- A Teams Community to discuss with Security Copilot product experts and engineers.

Opt in to join the community .

Last updated on 04/20/2026

Microsoft Security Copilot experiences

Security Copilot offers an immersive standalone experience and intuitive embedded experiences that are available in other Microsoft security products.

In addition to the standalone and embedded experiences, Security Copilot offers developer scenarios that enable you to build, test, publish, and integrate agents and plugins to meet unique security needs. These scenarios empower developers to streamline workflows, automate tasks, and extend Security Copilot's capabilities across Microsoft security products.

This topic helps you:

- Identify the differences between standalone and embedded Security Copilot experiences
- Explore available embedded experiences across Microsoft security solutions
- Find guidance for hiding or showing the Copilot sidecar panel
- Understand developer scenarios such as building custom agents and automating workflows with Security Copilot

Watch the following video to learn more about the Security Copilot experiences:

https://learn-video.azurefd.net/vod/player?id=257c1a95-04fb-422e-add2-51e91388ae5f&locale=en-us&embedUrl=%2Fcopilot%2Fsecurity%2Fexperiences-security-copilot

Standalone and embedded experiences

Standalone experience
Security Copilot, accessed through https://securitycopilot.microsoft.com, is considered the **standalone** experience.

Embedded experience
Accessing Security Copilot embedded experiences in other Microsoft security products is considered an **embedded** experience.

> **Important**
>
> Guidance on specific embedded experiences can be found in the documentation library of the corresponding service. For example, if you access a Microsoft Defender XDR embedded experience, then the corresponding documentation for that Security Copilot experience can be found in the Microsoft Defender XDR documentation. This ensures that you receive service-specific guidance wherever you access Security Copilot.

A feature common to all embedded experiences is the ability to hide or show the Copilot sidecar. The panel behavior persists per Copilot embedded experience, remembering that user's option. The embedded copilot functionality remains the same whether it's hidden or not, so closing the panel doesn't affect SCU consumption.

For example, to hide the Copilot in Microsoft Defender XDR experience, select the **X**. The incident summary will continue to generate in the background.

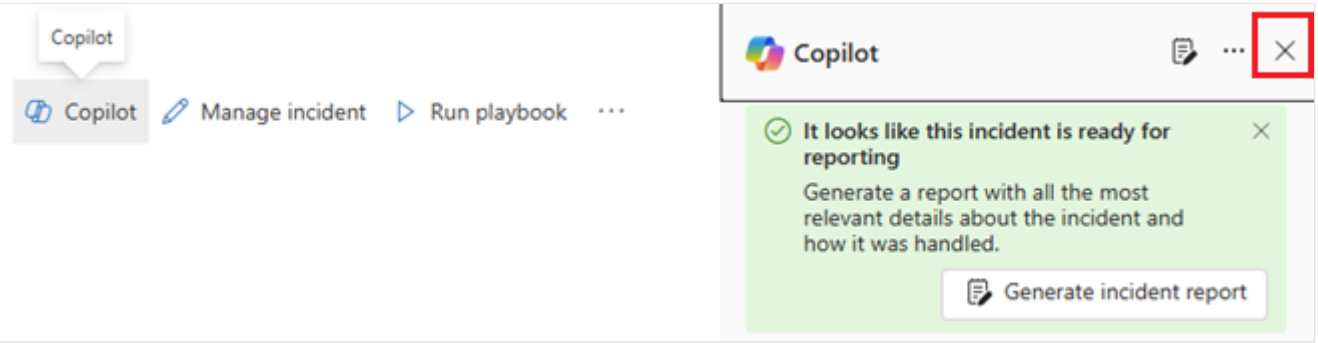

To make the panel visible again, select the **Copilot** button.

Embedded experiences

The following table lists the available embedded experiences.

Expand table

Product	Embedded experience
Azure Firewall	- Enrich the threat profile of an IDPS signature beyond log information - Generate recommendations to secure your environment using Azure Firewall's IDPS feature - Look for a given IDPS signature across your tenant, subscription, or resource group - Retrieve the top IDPS signature hits for an Azure Firewall
Microsoft Defender for Cloud	- Analyze recommendations - Delegate recommendations - Remediate code - Remediate recommendations

Product	Embedded experience
	- Summarize recommendations
Microsoft Defender Threat Intelligence	- Using Security Copilot for threat intelligence
Microsoft Defender XDR	- Analyze files - Analyze scripts and codes - Create incident reports - Generate KQL queries for hunting - Summarize device information - Summarize incidents - Summarize identities - Use guided response
Microsoft Entra	- Investigate app risk - Investigate incidents - Investigate risky users - Manage lifecycle workflows
Microsoft Intune	- Device query - Policy and setting management - Use Microsoft Copilot in Intune to troubleshoot devices
Microsoft Purview	- Investigate a Microsoft Purview Data Loss Prevention alert - Investigate insider risk management activities - Summarize Communication Compliance messages by using Security Copilot - Summarize an eDiscovery message by using Security Copilot

Developer scenarios

While Security Copilot offers immersive standalone and embedded experiences for end users, it also supports developer scenarios for organizations seeking to extend, automate, or integrate security workflows.

Developer scenarios empower you to:

- Build custom agents and plugins to automate security tasks
- Integrate Security Copilot with other Microsoft security products
- Build, publish, test and manage agents tailored to your organization's needs

Explore the Security Copilot Developer content.

Last updated on 11/13/2025

What's new in Microsoft Security Copilot?

Security Copilot receives improvements on an ongoing basis. To stay up to date with the most recent developments, this article provides you with information about:

- The latest releases
- Known issues
- Plans for changes

Each item describes where the enhancement can be experienced. For more information, see Security Copilot experiences.

This page updates monthly, so revisit it regularly.

April 2026

Security Analyst Agent - Public preview

Type: New feature
Experience: Standalone / Embedded

The Security Analyst Agent in Security Copilot helps security analysts quickly identify, assess, and prioritize risks. The agent performs deep, multi-step investigations across Microsoft Defender and Microsoft Sentinel telemetry to surface high-impact risks and deliver prioritized insights. Each finding includes clear reasoning and supporting evidence, enabling analysts to quickly understand and act on the results.

Key capabilities include:

- **Flexible analysis**: Perform ready-to-use or custom analyses on security data. Get actionable and prioritized insights, recommendations, and reports to uncover top vulnerabilities and risks.
- **Data integration**: Analyze data from Microsoft Defender XDR, Sentinel Log Analytics, or Sentinel Data Lake, based on your instructions.
- **Interactive exploration**: Visualize data to spot anomalies and risks faster.
- **Conversation assistance**: Chat with the agent, ask follow-up questions, and perform related analyses to deepen your understanding.

The agent can perform single or multi-step analysis on large volumes of data, iteratively reasoning to uncover hidden risks and prioritizing findings with a detailed evidence trail, without the need to write any code or queries.

For more information, see Security Analyst Agent.

November 2025

Microsoft Security Copilot for Microsoft 365 E5 customers

Type: New capability
Experience: Standalone / Embedded

Microsoft Security Copilot is included for all Microsoft 365 E5 customers. Expanded agent portfolio: 40 new Microsoft and partner-built agents added. 12 new Microsoft-built agents across Microsoft Defender, Entra, Intune, and Purview are available today in preview. Additionally, more than 30 new partner-built agents extend protection end-to-end. These agents automate large-scale tasks, which allows security teams to dedicate more time to strategic initiatives.

Existing Security Copilot customers with Microsoft 365 E5 subscriptions can get started with the agents:

- Security operations in Microsoft Defender
- Data security in Microsoft Purview
- Identity and access in Microsoft Entra
- Endpoint management in Microsoft Intune

To learn more on what's included for E5 customers, see Security Copilot inclusion.

September 2025

Microsoft Sentinel and Security Copilot integration - General availability

Type: New feature
Experience: Embedded

Microsoft Sentinel data lake is now generally available. New capabilities such as Microsoft Sentinel graph and the Model Context Protocol (MCP) server are in public preview. This

integration with Security Copilot lets agents access connected data across Sentinel, combining graph, structured, and semantic context to reason and act with greater precision. This foundation enhances AI-driven detection and response, helping your teams resolve incidents faster and uncover deeper insights across your environments.

For more information, see Microsoft Sentinel platform.

Build your own Security Copilot agent - Public preview

Type: New feature
Experience: Standalone

You can create custom Security Copilot agents. Use the standalone agent builder (no-code experience) or use developer tools to design, test, and deploy agents that automate the workflows you need most. Your team controls how the agents work and what they do.

For more information on building custom agents using the agent builder, see Agents in standalone.

For information on building agents using MCP, see MCP overview.

Microsoft and partner agents - Public preview

Type: New capability
Experience: Standalone / Embedded

These new agents help teams address common security and IT challenges faster and smarter:

- **Access Review Agent in Microsoft Entra**: Streamline access reviews, flag unusual patterns, and reduce fatigue for security and compliance teams. It helps maintain governance and compliance by automatically analyzing ongoing access reviews and highlighting potential risks.

  For more information, see Access Review Agent.

- **Phishing Triage Agent in Microsoft Defender**: In this new customer spotlight, St. Luke is seeing the impact of integrating Security Copilot agents into their daily workflows saving nearly 200 hours a month. With routine triage automated, security teams can shift from reactive response to proactive threat hunting, freeing up time for higher-value work and faster threat mitigation.

  For more information, see Customer Success.

About 30 new partner-built agents are now available in the Microsoft Security Store. These agents deliver solutions such as:

- **Forensic Agent by glueckkanja AG**: Delivers deep-dive analysis of Defender Extended Detection and Response (XDR) incidents to accelerate investigations and uncover root causes faster.
- **Privileged Admin Watchdog Agent by glueckkanja AG**: Helps enforce zero standing privilege principles by removing persistent admin identities, reducing risk, and strengthening administrative security.
- **Ransomware Kill Chain Investigator Agent by adaQuest**: Automates ransomware triage to quickly detect and respond to threats, enabling security teams to focus on high-priority incidents.
- **Entity Guard Investigator Agent by adaQuest**: Investigates Defender incidents and provides actionable insights to accelerate incident resolution and strengthen security posture.
- **Admin Guard Insight Agent by adaQuest**: Analyzes administrative activity, detects anomalies, evaluates risk exposure and compliance, and delivers actionable guidance to improve administrative security.
- **Identity Workload ID Agent by Invoke**: Empowers identity administrators and security teams to manage and secure Workload Identities in Microsoft Entra, reducing risk, strengthening compliance, and controlling identity sprawl.

Learn about these agents and more in the Microsoft Security Store.

Microsoft Security Store - Public preview

Type: New feature
Experience: Standalone

The Microsoft Security Store makes it simple to discover, deploy, and buy Security Copilot agents and partner solutions. You can use any of the 30 new agents or 50 SaaS solutions to power your SOC, IT, privacy, and compliance workflows.

For more information, see Security Store.

July 2025

Copilot in Microsoft Intune – General availability

Type: Changed
Experience: Embedded

IT admins can now use Security Copilot in Intune, which includes a dedicated data exploration experience, allowing them to ask questions, extract insights, and take action—all from within the Intune admin center. Whether it's identifying noncompliant devices, managing updates, or automating remediation, Copilot simplifies complex workflows and brings data and actions together in one place.

For more information, see Microsoft Copilot in Intune

Copilot in Microsoft Entra – General availability

Type: Changed
Experience: Embedded

Security Copilot in Microsoft Entra now brings AI-assisted investigation and identity management directly into the Microsoft Entra admin center. Admins can ask natural language questions to troubleshoot sign-ins, review access, monitor tenant health, and analyze role assignments—without writing queries or switching tools. With expanded coverage and improved performance, Copilot helps teams move faster, close gaps, and stay ahead of threats.

For more information, see Microsoft Copilot in Microsoft Entra.

Conditional Access Optimization Agent – General availability

Type: Changed
Experience: Embedded

The Conditional Access Optimization Agent in Microsoft Entra brings AI-powered automation to identity workflows. The agent runs autonomously to detect gaps, overlaps, and outdated policy assignments—then recommends precise, one-click remediations to close them fast. Key benefits include:

- Autonomous protection: Automatically identifies users and apps not covered by policies
- Explainable decisions: Plain-language summaries and visual activity maps
- Custom adaptability: Learns from natural-language feedback and supports business rules
- Full auditability: All actions logged for compliance and transparency

For more information, see Microsoft Entra Conditional Access optimization agent with Microsoft Security Copilot.

Phishing Triage Agent in Microsoft Defender - Public preview

Type: New capability
Experience: Embedded

The Phishing Triage Agent in Microsoft Defender is now in public preview, bringing autonomous, AI-powered threat detection to your SOC workflows. Powered by large language models, the agent performs deep semantic analysis of emails, URLs, and files to determine whether a submission is a phishing threat or a false alarm—without relying on static rules.

It learns from analyst feedback, adapts to your organization's patterns, and provides clear, natural language explanations for every verdict. A visual decision map shows exactly how the agent reached its conclusion, making the process fully transparent and reviewable.

For more information, see Phishing Triage Agent.

Threat Intelligence Briefing Agent is now - Public Preview

Type: New capability
Experience: Standalone

The Threat Intelligence Briefing agent has entered public preview in the Security Copilot standalone experience, transforming how security teams stay ahead of emerging threats. With this powerful agent, creating highly relevant, organization-specific threat intelligence briefings now takes minutes rather than hours or days, empowering teams to act with speed and confidence. Through real-time dynamic reasoning, the agent surfaces the most relevant threat intelligence based on attributes such as the organization's industry, geographic location, and unique attack surface to deliver critical context and invaluable situational awareness.

For more information, see Threat Intelligence Briefing Agent.

Security Copilot Capacity Calculator

The Security Copilot Capacity Calculator is now available in the standalone experience (Azure account required), helping teams estimate how many SCUs they may need. Security Copilot supports:

- Provisioned SCUs for predictable workloads
- Overage SCUs to scale with variable workloads

Teams can estimate initial capacity using the capacity calculator, monitor usage in the in-product usage dashboard, and adjust their SCU allocation as needed. Learn more about Security Copilot pricing here.

Learn more: New tools for Security Copilot management and capacity planning.

Audit agent changes with Purview Unified Audit Log (UAL) integration - General availability

Type: New capability
Experience: Embedded

Agent administration auditing is now generally available in Microsoft Purview Unified Audit Log, allowing teams to trace agent creation, updates, and deletions with detailed metadata for improved visibility and compliance.

For more information, see Access the Security Copilot audit log.

GPT 4.1 support - General availability

Type: Changed
Experience: Standalone / Embedded

GPT 4.1 Support is now generally available for all experiences at evaluation level in Security Copilot. With GPT-4.1 support, Security Copilot can now offer significantly larger context windows with better interactions and have observed accuracy increased.

Large output support - General availability

Type: Changed
Experience: Standalone / Embedded

Large output support is now generally available for all experiences. With this capability, Security Copilot can now support over 2MB of data for usage in LLM with restrictions only on OpenAI token limits.

Agent administration auditing - General availability

Type: Changed
Experience: Standalone / Embedded

Agent administration auditing is now generally available for Microsoft Purview Unified Audit Log. Agent administration actions whether creating, updating, or deleting agents and triggers are entered into the unified audit log as unique events. Each event type includes metadata pertaining to the agent, tenant, and user of the activation event providing admins the ability to trace changes across complex agent environments.

Dynamic suggested prompts for Microsoft Entra skills

Type: Changed
Experience: Standalone / Embedded

Dynamic suggested prompts are now generally available for Microsoft Entra skills, offering faster and more deterministic follow-up suggestions using direct skill invocation—bypassing the orchestrator for improved performance.

FedRAMP High authorization for Security Copilot

Security Copilot is now included within the Federal Risk and Authorization Management Program (FedRAMP) High Authorization for Azure Commercial. This Provisional Authorization to Operate (P-ATO) within the existing FedRAMP High Azure Commercial environment was approved by the FedRAMP Joint Authorization Board (JAB). This milestone marks a significant step forward in our mission to bring Microsoft Security Copilot's cutting-edge AI-powered security capabilities to our Government Community Cloud (GCC) customers. Stay tuned for updates on when Security Copilot will be fully available for GCC customers.

June 2025

Workspace level plugin management - General availability

Type: Changed
Experience: Standalone

Security Copilot now supports workspaces, giving organizations a flexible way to segment environments by team, region, or business unit. With workspaces now in public preview, admins can align access, data boundaries, and SCU capacity with operational and compliance needs. Each workspace supports role-based access control, localized prompt history, and independent capacity planning – making it easier to manage complex, distributed security and IT operations.

As part of this model, workspace-level plugin management is now generally available, allowing admins to configure plugin settings at the workspace or organization level. This eliminates the need for per-user setup and improves efficiency across large environments.

For more information, see Manage workspaces.

Microsoft Entra embedded experience powered by the NL2API skill - General availability

Type: New capability
Experience: Embedded

Enables Security Copilot to perform advanced reasoning over APIs (MS Graph APIs) and answer complex, multi-stage questions across Microsoft Entra resources.

Korean language

Type: Changed
Experience: Standalone

Korean language is now supported in security copilot standalone and embedded experiences.

For more information, see Supported languages.

Swiss region data residency

Type: Changed
Experience: Standalone

Swiss region data residency is now available enabling storage of Security Copilot data within Switzerland boundary.

For more information, see Availability and recovery.

May 2025

Censys - Public preview

Type: New capability
Experience: Standalone

This plugin allows Security Copilot users to enrich investigations using threat intelligence from the Censys Platform.

For more information, see Censys

HP Workforce Experience Platform - Public preview

Type: New capability
Experience: Standalone

Allows users to query data about their fleet and get responses back quickly and efficiently. This plugin can answer questions about PC warranty, Windows 11 readiness, resource consumption. It can also answer device age/stability/health, security, blue screens, and app performance including hangs and crashes.

For more information, see HP Workforce Experience Platform

Copilot Studio connector - General Availability

Type: New capability
Experience: Standalone

This connector helps users submit natural language prompts to create a new Security Copilot investigation through Copilot Studio. After completion, the evaluation result is returned to the workflow.

For more information, see Copilot Studio connector.

April 2025

Overage SCUs - General Availability

Type: New capability
Experience: Standalone

Customers can set an overage amount of SCUs to ensure that additional SCUs are available when initially provisioned units are depleted during unexpected workload spikes. Overage units are billed on-demand and can be set as unlimited or a maximum amount.

For more information, see Manage usage.

Splunk - Public preview

Type: New feature
Experience: Standalone

The Splunk plugin enables customers to perform searches in Splunk, retrieve alerts, and others.

For more information, see Splunk.

Microsoft Threat Intelligence recommended actions - General Availability

Type: Changed
Experience: Standalone

When investigating a vulnerability or attack utilizing an exploit, it's critical to understand the recommended actions for mitigation. This information is now available through the Microsoft Threat Intelligence plugin.

Microsoft Threat Intelligence - Malware encyclopedia integration - Public preview

Type: New capability
Experience: Standalone

When investigating an incident involving a malware it's critical to quickly get the information about this malware, relevant Defender for Endpoint version and guidance for next steps. This capability allows customers to query this information directly in Security Copilot.

Microsoft Threat Intelligence - Reason for inclusion - General Availability

Type: New capability
Experience: Standalone

When the analyst receives a response with multiple results – such as multiple actors or attacks, multiple articles, and others, it's sometimes difficult to understand which result is the most relevant for the next level of drill-down. With this enhancement it's easier to understand the relevance of each result.

Microsoft Threat Intelligence - Suggested prompts - General Availability

Type: New capability
Experience: Standalone

During a session in Security Copilot customers often need help selecting the next prompt for their investigation. This is a critical part of a guided experience.

Microsoft Threat Intelligence - Indicator skills - General Availability

Type: New capability
Experience: Standalone

Microsoft Threat Intelligence indicator skills allow direct access to all suspicious IOC and the mass datasets of web scan data including DNS, WHOIS, host pairs, certificates, detonation data, and others.

March 2025

Global Copilot experience - Public preview

Type: New feature
Experience: Embedded

The Global Copilot experience is a new global entry point in the Suite Header on the Microsoft Purview compliance portal home page and all solution pages. This capability allows for easy invocation of the Copilot experience. For each page, Copilot offers readily available prompts for you to get started, learn more about Purview, and troubleshoot issues.

For more information, see Security Copilot in Microsoft Purview.

Microsoft Surface Management integration - Public preview

Type: New feature
Experience: Embedded

With Copilot in Intune, IT Admins can quickly search for and resolve specific device issues, summarize warranty information, and access support tickets and service orders related to their organization's Surface devices. This reduces the time and effort needed for routine maintenance tasks, creating more time to focus on other initiatives. In addition, Copilot pulls contextually relevant data from your Intune-enrolled Surface devices along with public

information into a single view, streamlining the management process and enhancing overall efficiency.

For more information, see Security Copilot in Microsoft Surface Management Portal .

CheckPhish - General Availability

Type: New feature
Experience: Standalone

This plugin allows users to utilize the CheckPhish AI to analyze URLs for potential phishing threats, tech support scams, cryptojacking, and other security risks.

For more information, see CheckPhish.

Plugin management - General Availability

Type: New feature
Experience: Standalone

You can personalize the way Security Copilot interacts with you by managing plugins.

For more information, see Manage plugins.

KQL response explanation feature in Query assistant - General Availability

Type: New feature
Experience: Standalone / Embedded

The Security Copilot Query Assistant in Advanced hunting generates KQL queries from user requests in natural language, allowing junior and first-tier analysts that aren't familiar with KQL to easily hunt for threats. With this new feature, customers can see the logic behind the KQL queries generated by Copilot, including a breakdown of the query. This enhancement helps customers validate that the query aligns with their intent and needs, even if they don't have a deep understanding of KQL. Additionally, this feature helps customers improve their knowledge and proficiency in KQL, and will also increase their trust in the Query Assistant.

For more information, see Microsoft Security Copilot in advanced hunting.

Set prompts to continue on failure - General Availability

Type: New feature
Experience: Standalone

The **Continue on failure** feature is now available in Security Copilot promptbooks. This feature allows users to set a promptbook to continue running the next prompt even after a prompt fails, instead of halting the entire promptbook.

For more information, see Set prompts to continue on failure.

Call system capabilities directly - Public preview

Type: New feature
Experience: Standalone

You can now directly call a Security Copilot system capability during promptbook creation. System capabilities are the various actions that Security Copilot can perform for you. This feature allows you to generate responses more efficiently and might help lower SCU consumption.

For more information, see Call capabilities directly in the prompt.

Last updated on 04/17/2026

Apply principles of Zero Trust to Microsoft Security Copilot

08/22/2025

Summary: To apply Zero Trust principles to your environment for Microsoft Security Copilot, you need to apply five layers of protection:

1. Protect admin and SecOps staff user accounts with identity and access policies.
2. Apply least privilege access to admin and SecOps staff user accounts, including assigning the minimum user account roles.
3. Manage and protect admin and SecOps staff devices.
4. Deploy or validate your threat protection.
5. Secure access to third-party security products that you integrate with Security Copilot.

Introduction

As a part of introducing Microsoft Security Copilot into your environment, Microsoft recommends that you build a strong foundation of security for your admin and SecOps staff user accounts and devices. Microsoft also recommends ensuring you have configured threat protection tools. If you're integrating third-party security products with Security Copilot, also ensure you've protected access to these products and related data.

Fortunately, guidance for a strong security foundation exists in the form of Zero Trust. The Zero Trust security strategy treats each connection and resource request as though it originated from an uncontrolled network and a bad actor. Regardless of where the request originates or what resource it accesses, Zero Trust teaches us to "never trust, always verify."

From within security portals, Security Copilot provides a natural language, assistive copilot experience that helps support:

- Security professionals in end-to-end scenarios such as incident response, threat hunting, intelligence gathering, and posture management.
- IT professionals in policy evaluation and configuration, device and user access troubleshooting, and performance monitoring.

Security Copilot uses data from event logs, alerts, incidents, and policies for your Microsoft and third-party subscriptions and security products. If an attacker compromises an admin or security staff user account that has been assigned a Security Copilot role, they can use Security Copilot and its results to understand how your SecOps team is addressing attacks in progress. An attacker can then use this information to thwart attempts to respond to an incident, possibly one that they initiated.

Consequently, it's critical to ensure that you've applied appropriate mitigations within your environment.

Logical architecture

The first line of defense when introducing Security Copilot is to apply the principles of Zero Trust to the accounts and devices of admin and SecOps staff. It's also important to ensure your organization applies the principle of least privilege. In addition to Copilot-specific roles, the assigned roles for admin and SecOps staff within your security tools determine what data they have access to while using Security Copilot.

It's easy to understand why these mitigations are important by looking at the logical architecture of Security Copilot shown here.

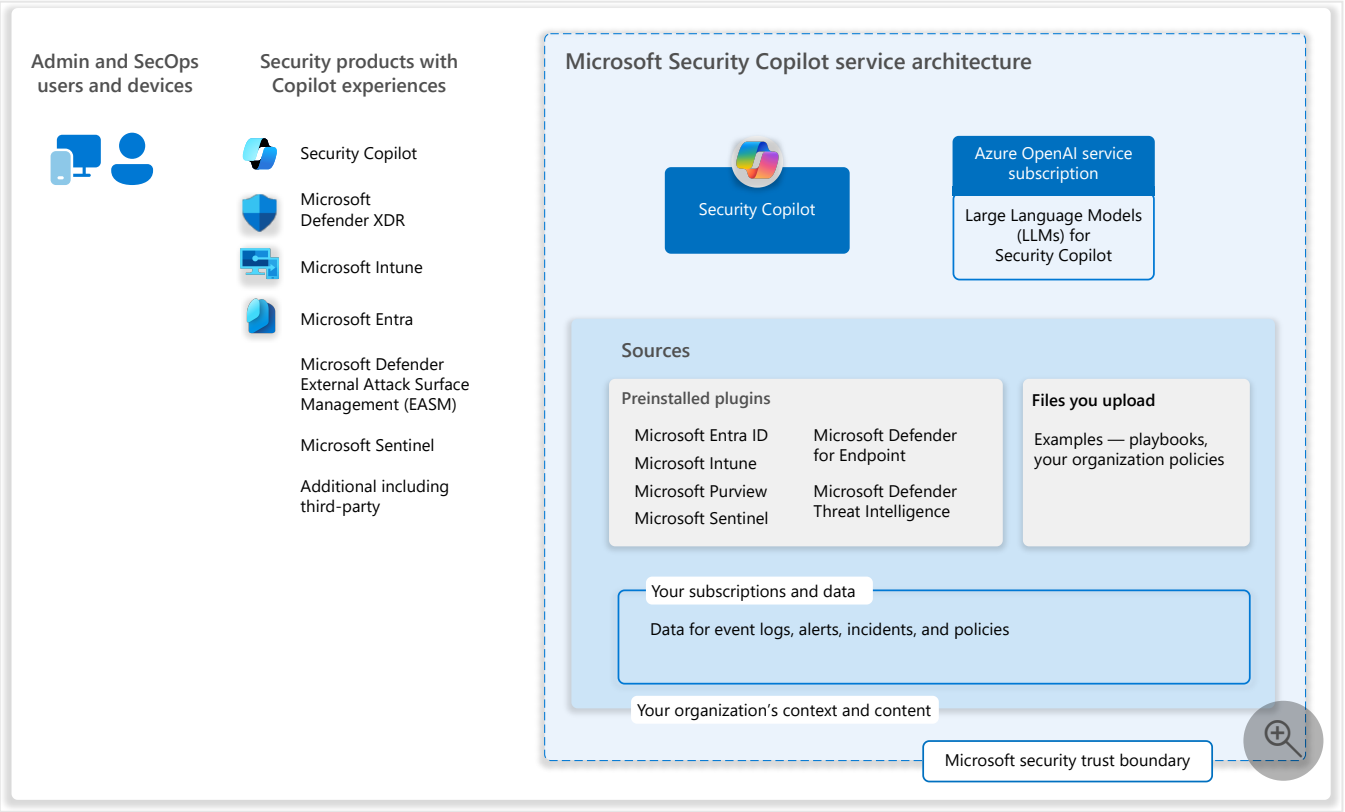

In the diagram:

- SecOps team members can prompt using a copilot experience, such as those offered by Security Copilot, Microsoft Defender XDR, and Microsoft Intune.
- Security Copilot components include:
  - The Security Copilot service, which orchestrates responses to user and skill-based prompts.
  - A set of Large Language Models (LLMs) for Security Copilot.
  - Plugins for specific products. Preinstalled plugins for Microsoft products are provided. These plugins preprocess and post-process prompts.
  - Your subscription data. The SecOps data for event logs, alerts, incidents, and policies stored in the subscriptions. For more information, see this Microsoft Sentinel article for the most common data sources for security products.
  - Files you upload. You can upload specific files to Security Copilot and include these in the scope of prompts.

Each Microsoft security product with a copilot experience only provides access to the data set associated with that product, such as event logs, alerts, incidents, and policies. Security Copilot provides access to all the data sets to which the user has access.

For more information, see Get started with Microsoft Security Copilot.

How does on-behalf-of authentication work with Security Copilot?

Security Copilot uses on-behalf-of (OBO) authentication provided by OAuth 2.0. This is an authentication flow provided by delegation in OAuth. When a SecOps user issues a prompt, Security Copilot passes the user's identity and permissions through the request chain. This prevents the user gaining permission to resources for which they shouldn't have access.

For more information about OBO authentication, see Microsoft identity platform and OAuth2.0 On-Behalf-Of flow.

Prompting within a Microsoft security product: Embedded example for Microsoft Intune

When you use one of the embedded experiences of Security Copilot, the scope of the data is determined by the context of the product you're using. For example, if you issue a prompt within Microsoft Intune, the results are produced from data and context provided by Microsoft Intune only.

Here's the logical architecture when issuing prompts from within the Microsoft Intune embedded experience.

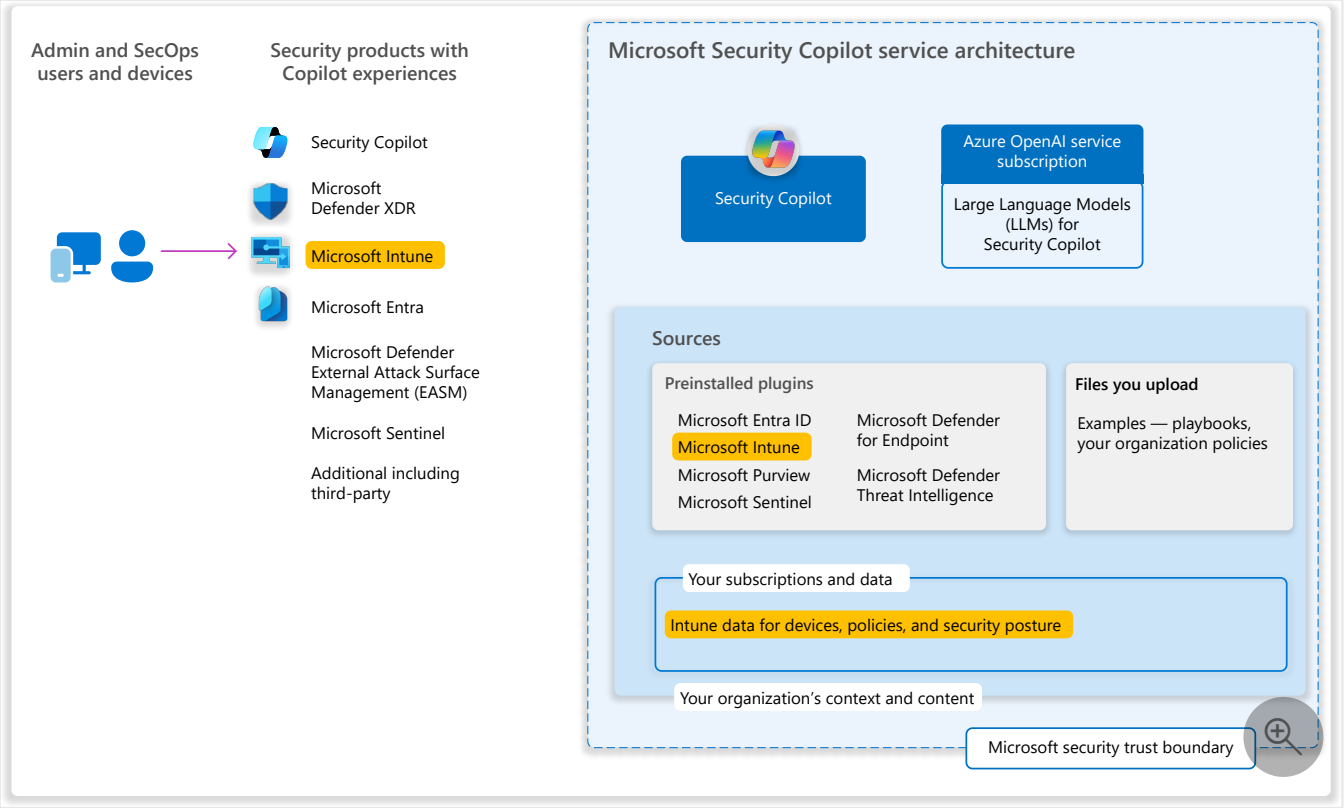

In the diagram:

- Intune admins use the Microsoft Copilot in Intune experience to submit prompts.

- The Security Copilot component orchestrates responses to the prompts using:
  - The LLMs for Security Copilot.
  - The Microsoft Intune preinstalled plugin.
  - The Intune data for devices, policies, and security posture stored in your Microsoft 365 subscription.

Integrating with third-party security products

Security Copilot provides the ability to host plugins for third-party products. These third-party plugins provide access to their associated data. These plugins and their associated data live outside the Microsoft security trust boundary. Consequently, it's important to ensure you've secured access to these applications and their associated data.

Here's the logical architecture of Security Copilot with third-party security products.

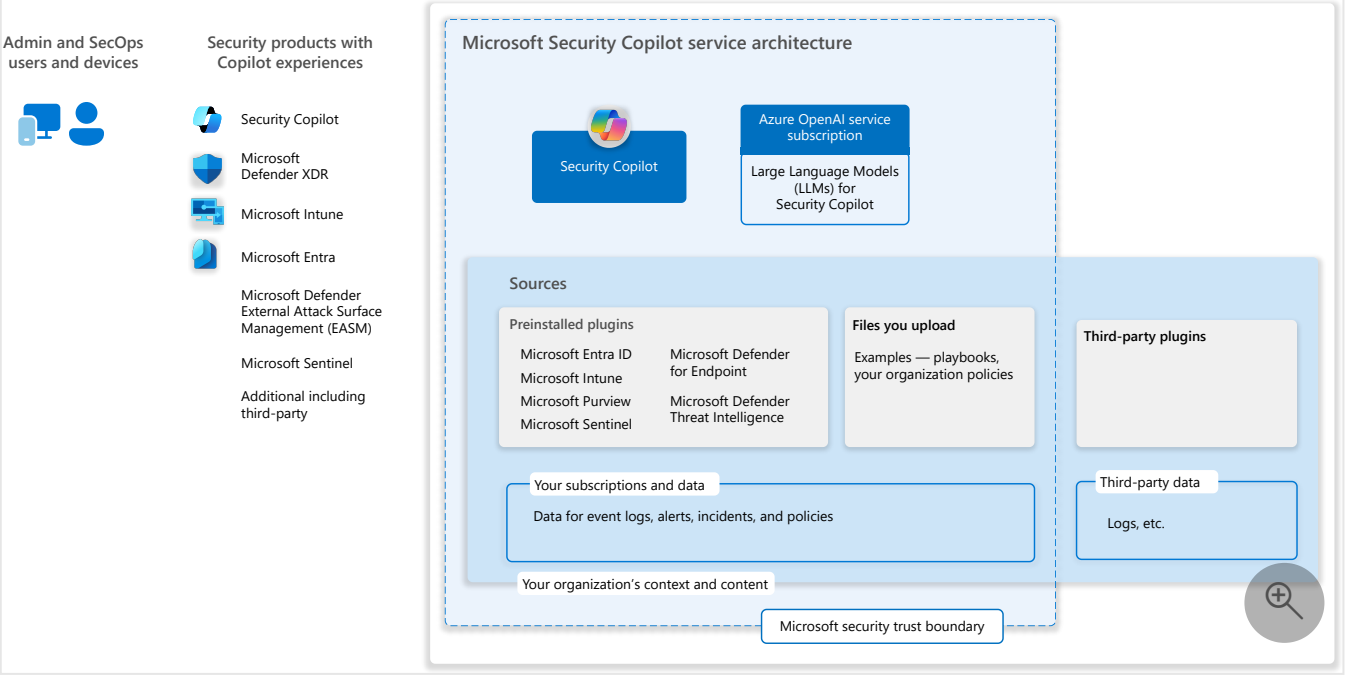

In the diagram:

- Security Copilot integrates with third-party security products through plugins.
- These plugins provide access to the data associated with the product, such as logs and alerts.
- These third-party components reside outside the Microsoft security trust boundary.

Applying security mitigations to your environment for Security Copilot

The rest of this article walks you through the steps to apply the principles of Zero Trust to prepare your environment for Security Copilot.

Expand table

Step	Task	Zero Trust principles applied
1	Deploy or validate identity and access policies for admin and SecOps staff.	Verify explicitly
2	Apply least privilege to admin and SecOps user accounts.	Use least privileged access
3	Secure devices for privileged access.	Verify explicitly
4	Deploy or validate your threat protection services.	Assume breach
5	Secure access to third-party security products and data.	Verify explicitly Use least privileged access Assume breach

There are several approaches you can take to onboarding admin and SecOps staff to Security Copilot while you configure protections for your environment.

Per-user onboarding to Security Copilot

At a minimum, walk through a checklist for your admin and SecOps staff before you assign a role for Security Copilot. This works well for small teams and organizations that want to start with a test or pilot group.

Phased deployment of Security Copilot

For large environments, a more standard phased deployment works well. In this model, you address groups of users at the same time to configure protection and assign roles.

Here is an example model.

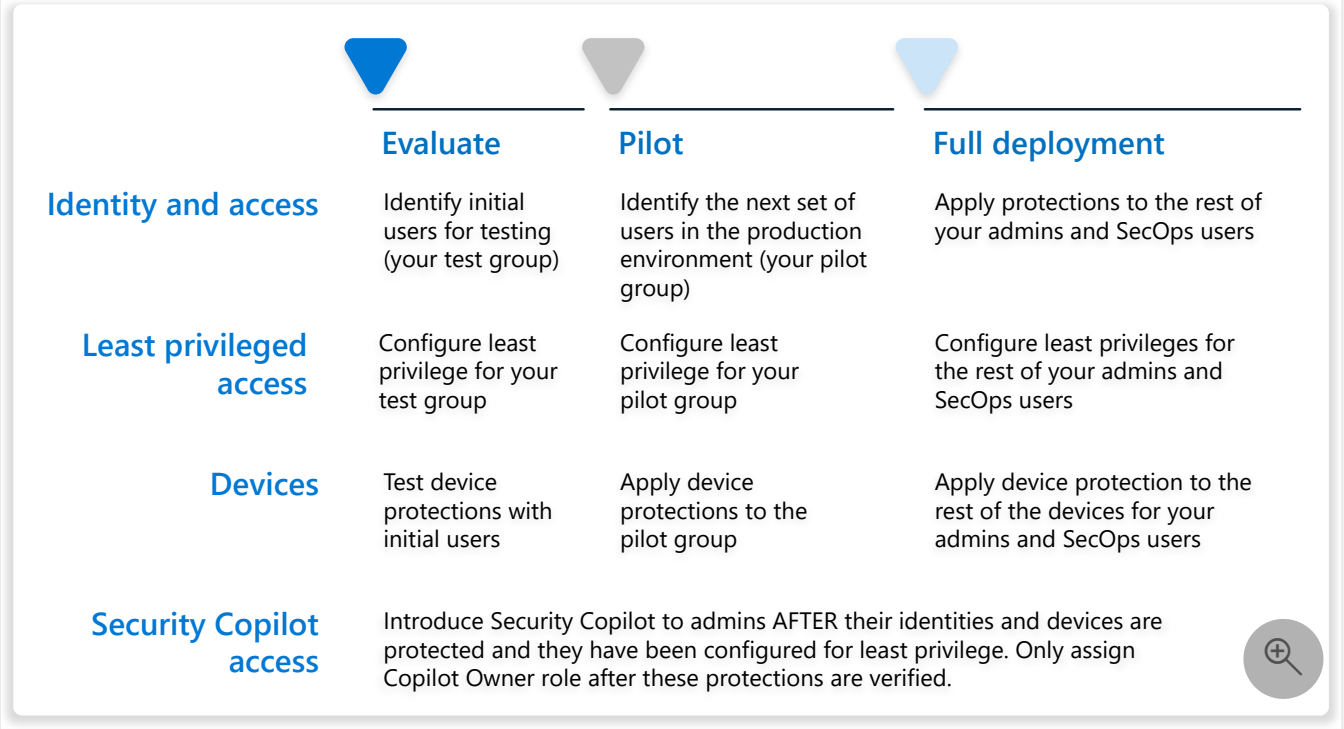

In the illustration:

- In the **Evaluate** phase, you pick a small set of admin and SecOps users that you want to have access to Security Copilot and apply identity and access and device protections.
- In the **Pilot** phase, you pick the next set of admin and SecOps users and apply identity and access and device protections.
- In the **Full deployment** phase, you apply identity and access and device protections for the rest of your admin and SecOps users.
- At the end of each phase, you assign the appropriate role in Security Copilot to the user accounts.

Because different organizations can be at various stages of deploying Zero Trust protections for their environment, in each of these steps:

- If you're NOT using any of the protections described in the step, take the time to pilot and deploy them to your admin and SecOps staff prior to assigning roles that include Security Copilot.
- If you're already using some of the protections described in the step, use the information in the step as a checklist and verify that each protection stated has been piloted and deployed prior to assigning roles that include Security Copilot.

Step 1. Deploy or validate identity and access policies for admin and SecOps staff

To prevent bad actors from using Security Copilot to quickly get information on cyberattacks, the first step is to prevent them from gaining access. You must ensure that admin and SecOps staff:

- User accounts are required to use multifactor authentication (MFA) (so their access can't be compromised by guessing user passwords alone) and they're required to change their passwords when high-risk activity is detected.
- Devices must comply with Intune management and device compliance policies.

For identity and access policy recommendations, see the identity and access step in Zero Trust for Microsoft 365 Copilot. Based on the recommendations in this article, ensure that your resulting configuration applies the following policies for all SecOps staff user accounts and their devices:

- Always use MFA for sign-ins
- Block clients that don’t support modern authentication
- Require compliant PCs and mobile devices
- High risk users must change password (for Microsoft 365 E5 only)
- Require adherence to Intune device compliance policies

These recommendations align with the **Specialized security** protection level in Microsoft's Zero Trust identity and device access policies. The following diagram illustrates the recommended three levels of protection: Starting point, Enterprise, and Specialized. The Enterprise protection level is recommended **as a minimum** for your privileged accounts.

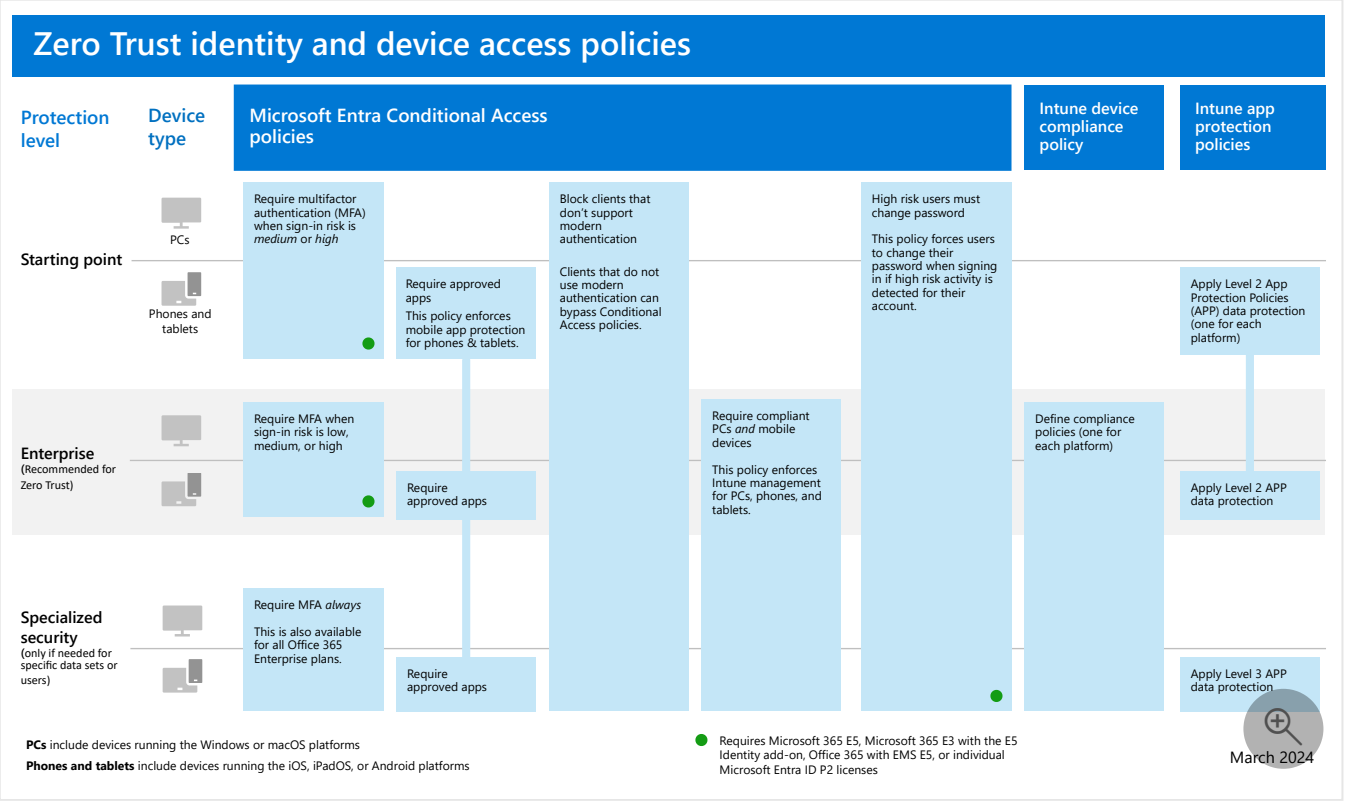

In the diagram, the recommended policies for Microsoft Entra Conditional Access, Intune device compliance, and Intune app protection are illustrated for each of the three levels:

- **Starting point**, which doesn't require device management.
- **Enterprise** is recommended for Zero Trust and as a minimum for access to Security Copilot and your third-party security products and related data.
- **Specialized security** is recommended for access to Security Copilot and your third-party security products and related data.

Each of these policies is described in greater detail in Common Zero Trust identity and device access policies for Microsoft 365 organizations.

Configure a separate set of policies for privileged users

When configuring these policies for your admin and SecOps staff, create a separate set of policies for these privileged users. For example, don't add your admins to the same set of policies that govern access of unprivileged users to apps such as Microsoft 365 and Salesforce. Use a dedicated set of policies with protections that are appropriate for privileged accounts.

Include security tools in the scope of Conditional Access policies

For now, there's not an easy way to configure Conditional Access for Security Copilot. However, because on-behalf-of authentication is used to access data within security tools, be sure you have configured Conditional Access for these tools, which can include Microsoft Entra ID and Microsoft Intune.

Step 2. Apply least privilege to admin and SecOps user accounts

This step includes configuring the appropriate roles within Security Copilot. It also includes reviewing your admin and SecOps user accounts to ensure they are assigned the least amount of privileges for the work they're intended to do.

Assign user accounts to Security Copilot roles

The permissions model for Security Copilot includes roles in both Microsoft Entra ID and Security Copilot.

Expand table

Product	Roles	Description
Microsoft Entra ID / Microsoft Purview	Microsoft Entra ID and Mirosoft Purview supported roles	These Microsoft Entra and Microsoft Purview roles inherit the **Copilot owner** role in Security Copilot. Only use these privileged roles to onboard Security Copilot to your organization.
Security Copilot	Copilot owner Copilot contributor	These two roles include access to use Security Copilot. Most of your admin and SecOps staff can use the **Copilot contributor** role. The **Copilot owner** role includes the ability to publish custom plugins and to manage settings that affect all of Security Copilot.

It's important to know that, by default, *all users in the tenant are given Copilot contributor access*. With this configuration, access to your security tool data is governed by the permissions you configured for each of the security tools. An advantage of this configuration is that the embedded experiences of Security Copilot are immediately available to your admin and SecOps staff within the products they use daily. This works well if you've already adopted a strong practice of least privileged access within your organization.

If you'd like to take a staged approach to introducing Security Copilot to your admin and SecOps staff while you tune up least privileged access in your organization, remove **All users** from the **Copilot contributor** role and add security groups as you're ready.

For more information, see these Microsoft Security Copilot resources:

- Get started
- Understand authentication

Configuring or reviewing least privilege access for admin and SecOps user accounts

Introducing Security Copilot is a great time to review the access of your admin and SecOps staff user accounts to be sure you're following through with the principle of least privilege for their access to specific products. This includes the following tasks:

- Review the privileges granted for the specific products your admin and SecOps staff work with. For example, for Microsoft Entra, see Least privileged roles by task.
- Use Microsoft Entra Privileged Identity Management (PIM) to gain greater control over access to Security Copilot.
- Use Microsoft Purview Privileged Access Management to configure granular access control over privileged admin tasks in Office 365.

Using Microsoft Entra Privileged Identity Management together with Security Copilot

Microsoft Entra Privileged Identity Management (PIM) allows you to manage, control, and monitor the roles required to access Security Copilot. With PIM, you can:

- Provide role activation that is time-based.
- Require approval to activate privileged roles.
- Enforce MFA to activate any role.
- Get notifications when privileged roles are activated.
- Conduct access reviews to ensure admin and SecOps staff user accounts still need their assigned roles.
- Perform audits on access and role changes for admin and SecOps staff.

Using privileged access management together with Security Copilot

Microsoft Purview Privileged Access Management helps protect your organization from breaches and helps to meet compliance best practices by limiting standing access to sensitive data or access to critical configuration settings. Instead of administrators having constant

access, just-in-time access rules are implemented for tasks that need elevated permissions. Instead of administrators having constant access, just-in-time access rules are implemented for tasks that need elevated permissions. For more information, see Privileged access management.

Step 3. Secure devices for privileged access

In Step 1, you configured Conditional Access policies that required managed and compliant devices for your admin and SecOps staff. For additional security, you can deploy privileged access devices for your staff to use when accessing security tools and data, including Security Copilot. A privileged access device is a hardened workstation that has clear application control and application guard. The workstation uses credential guard, device guard, app guard, and exploit guard to protect the host from attackers.

For more information on how to configure a device for privileged access, see Securing devices as part of the privileged access story.

To require these devices, be sure to update your Intune device compliance policy. If you're transitioning admin and SecOps staff to hardened devices, transition your security groups from the original device compliance policy to the new policy. The Conditional Access rule can remain the same.

Step 4. Deploy or validate your threat protection services

To detect the activities of bad actors and keep them from gaining access to Security Copilot, ensure that you can detect and respond to security incidents with a comprehensive suite of threat protection services, which include Microsoft Defender XDR with Microsoft 365, Microsoft Sentinel, and other security services and products.

Use the following resources.

Expand table

Scope	Description and resources
Microsoft 365 and SaaS apps integrated with Microsoft Entra	See the Zero Trust for Microsoft 365 Copilot article for guidance on how to ramp up threat protection beginning with Microsoft 365 E3 plans and progressing with Microsoft E5 plans. For Microsoft 365 E5 plans, also see Evaluate and pilot Microsoft Defender XDR security.

Scope	Description and resources
Your Azure cloud resources Your resources in other cloud providers, such as Amazon Web Services (AWS)	Use the following resources to get started with Defender for Cloud: - Microsoft Defender for Cloud - Apply Zero Trust principles to IaaS applications in AWS
Your digital estate with all Microsoft XDR tools and Microsoft Sentinel	The Implement Microsoft Sentinel and Microsoft Defender XDR for Zero Trust solution guide walks through the process of setting up Microsoft eXtended detection and response (XDR) tools together with Microsoft Sentinel to accelerate your organization's ability to respond to and remediate cybersecurity attacks.

Step 5. Secure access to third-party security products and data

If you're integrating third-party security products with Security Copilot, be sure you have secured access to these products and related data. Microsoft Zero Trust guidance includes recommendations for securing access to SaaS apps. These recommendations can be used for your third-party security products.

For protection with identity and device access policies, changes to common policies for SaaS apps are outlined in red in the following diagram. These are the policies to which you can add your third-party security products.

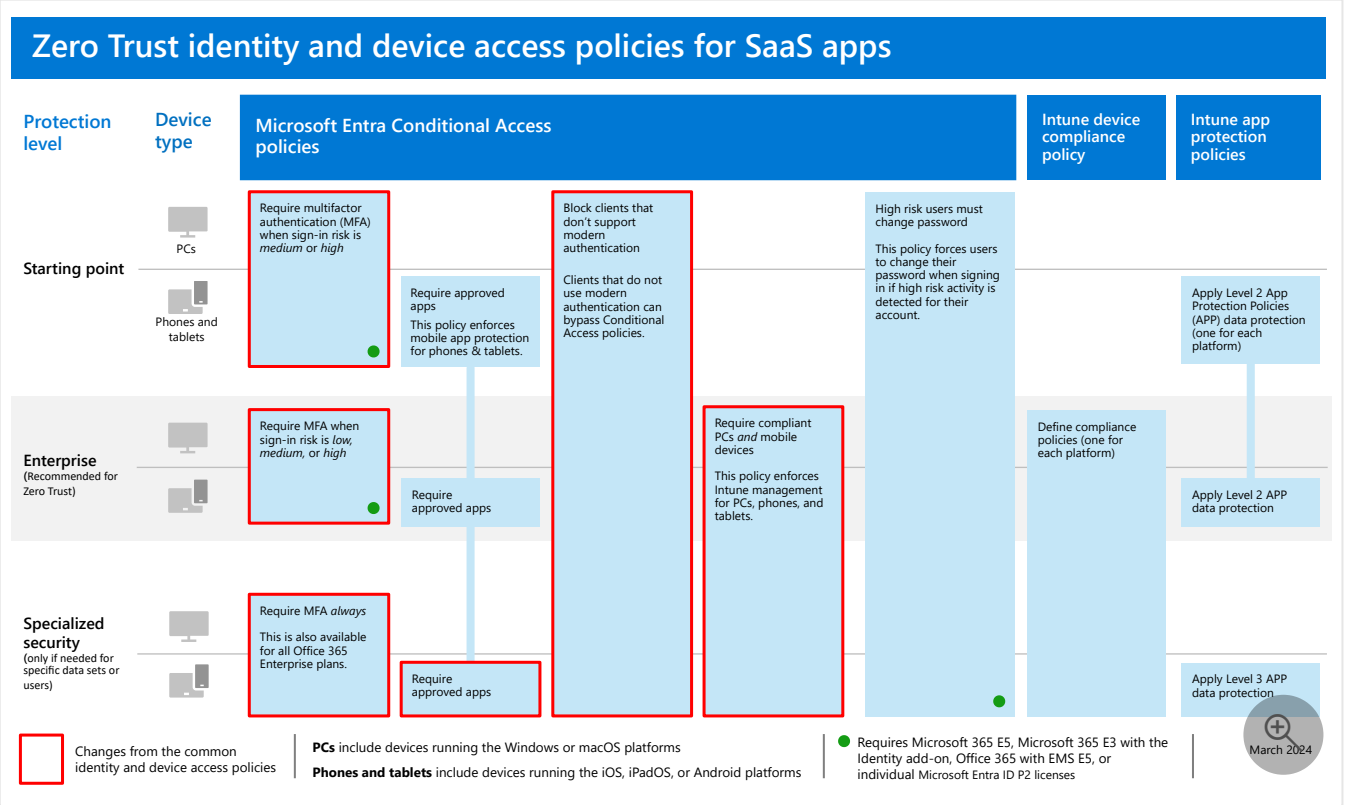

For your third-party security products and apps, consider creating a dedicated set of policies for these. This allows you to treat your security products with greater requirements compared to productivity apps, like Dropbox and Salesforce. For example, add Tanium and all other third-party security products to the same set of Conditional Access policies. If you want to enforce stricter requirements for devices for your admin and SecOps staff, also configure unique policies for Intune device compliance and Intune app protection and assign these policies to your admin and SecOps staff.

For more information about adding your security products to Microsoft Entra ID and to the scope of your Conditional Access and related policies (or configuring a new set of policies), see Add SaaS apps to Microsoft Entra ID and to the scope of policies.

Depending on the security product, it might be appropriate to use Microsoft Defender for Cloud Apps to monitor the use of these apps and apply session controls. Additionally, if these security apps include storage of data in any of the file types supported by Microsoft Purview, you can use Defender for Cloud to monitor and protect this data using sensitivity labels and data loss prevention (DLP) policies. For more information, see Integrate SaaS apps for Zero Trust with Microsoft 365.

Example for Tanium SSO

Tanium is a provider of endpoint management tools and offers a custom Tanium Skills plugin for Security Copilot. This plugin helps ground prompts and responses that leverage Tanium-gathered information and insights.

Here's the logical architecture of Security Copilot with the Tanium Skills plugin.

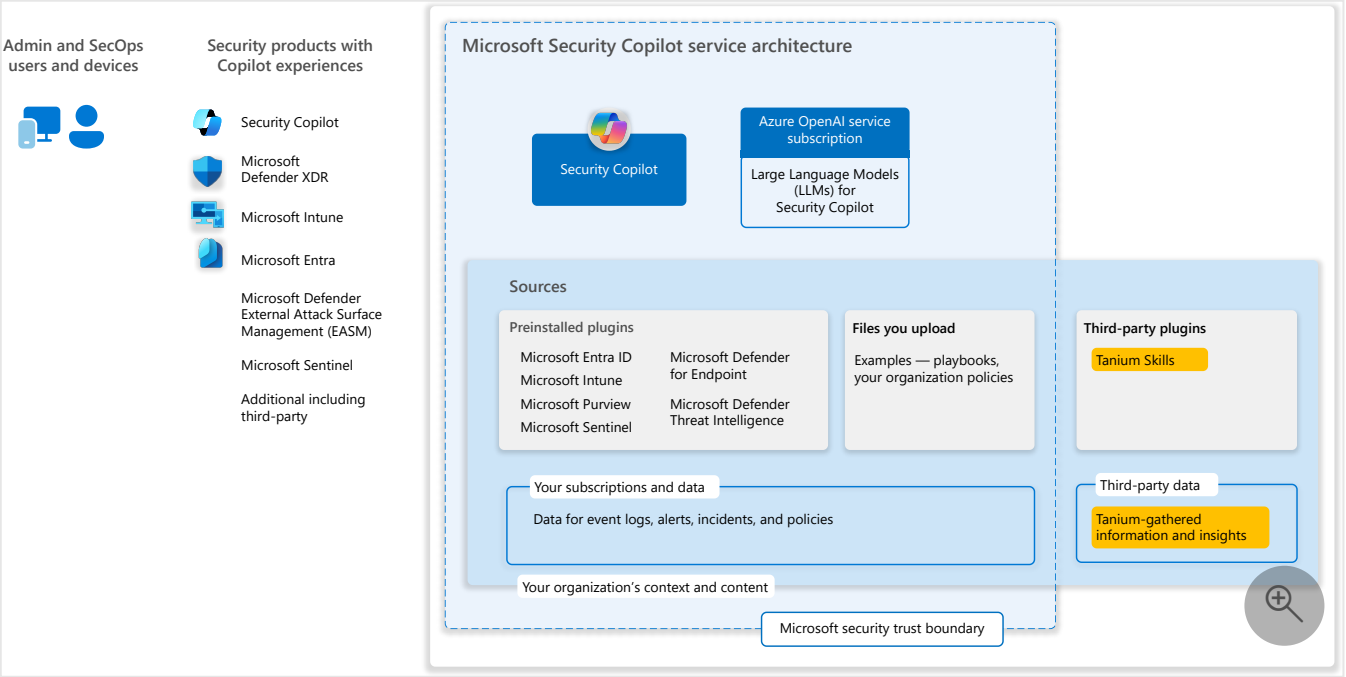

In the diagram:

- Tanium Skills is a custom plugin for Microsoft Security Copilot.
- Tanium Skills provides access to and helps ground both prompts and responses that use Tanium-gathered information and insights.

To secure access to Tanium products and related data:

1. Use the Microsoft Entra ID Application Gallery to find and add Tanium SSO to your tenant. See Add an enterprise application. For a Tanium-specific example, see Microsoft Entra SSO integration with Tanium SSO.
2. Add Tanium SSO to the scope of your Zero Trust identity and access policies.

Next steps

Watch the Discover Microsoft Security Copilot video.

See these additional articles for Zero Trust and Microsoft's Copilots:

- Overview
- Microsoft Copilot
- Microsoft 365 Copilot

Also see the Microsoft Security Copilot documentation.

References

Refer to these links to learn about the various services and technologies mentioned in this article.

- Get started with Microsoft Security Copilot
- Onboarding to Security Copilot
- Data sources for Microsoft Sentinel
- Microsoft Entra Privileged Identity Management (PIM)
- Privileged access management
- Privileged access devices

Get started with onboarding to Microsoft Security Copilot

Microsoft Security Copilot is a generative AI security product that empowers security and IT professionals to respond to cyber threats, process signals, and assess risk exposure at the speed and scale of AI. For an overview, see What is Microsoft Security Copilot?.

> **Note**
>
> **Disclaimer:** This documentation is only intended for customers using commercial clouds. Currently, Security Copilot isn't designed for use by customers using US government clouds, including but not limited to GCC, GCC High, DoD, and Microsoft Azure Government. For more information, consult with your Microsoft representative.

> **Note**
>
> When you first sign in to Security Copilot, your experience might differ depending on your onboarding and rollout stage.
>
> - **New Microsoft 365 E5 or E7 customers** might see an **agents-first landing experience**, where agents are the default homepage and primary way to get started.
> - **Existing customers** might continue to see the current experience with chat and prompts on the homepage.
>
> In the agents-first experience, you access chat from **All history** instead of the homepage.

There are two categories of Security Copilot customers:

- **Microsoft 365 E5 and E7 Security Copilot included customers**

  Microsoft Security Copilot is included with your Microsoft 365 E5 or E7 license. However, eligibility alone does not enable access, Security Copilot must be rolled out in your tenant for access. Once provisioning is complete, Security Copilot is ready for use. See, How do I know if my tenant is enabled for Security Copilot?.

What you see when you first sign in

In some environments, Security Copilot opens with an **agents-first homepage** and a role-based welcome experience.
    - **Security analysts** are guided to set up and try agents.
    - **Administrators** are guided to assign roles, review usage and capacity, and configure agents.

    To start a chat session in this experience, go to **All history** and select **New session**.

- **Non-Microsoft 365 E5 or E7 Security Copilot customers**

    If you're not a Microsoft 365 E5 and E7 customer, you must follow the manual onboarding steps to provision Security Compute Units (SCUs) so that you can use Security Copilot.

    After onboarding is complete, your experience might vary depending on rollout (for example, an agents-first landing page).

To get started with Security Copilot, use the following guidance:

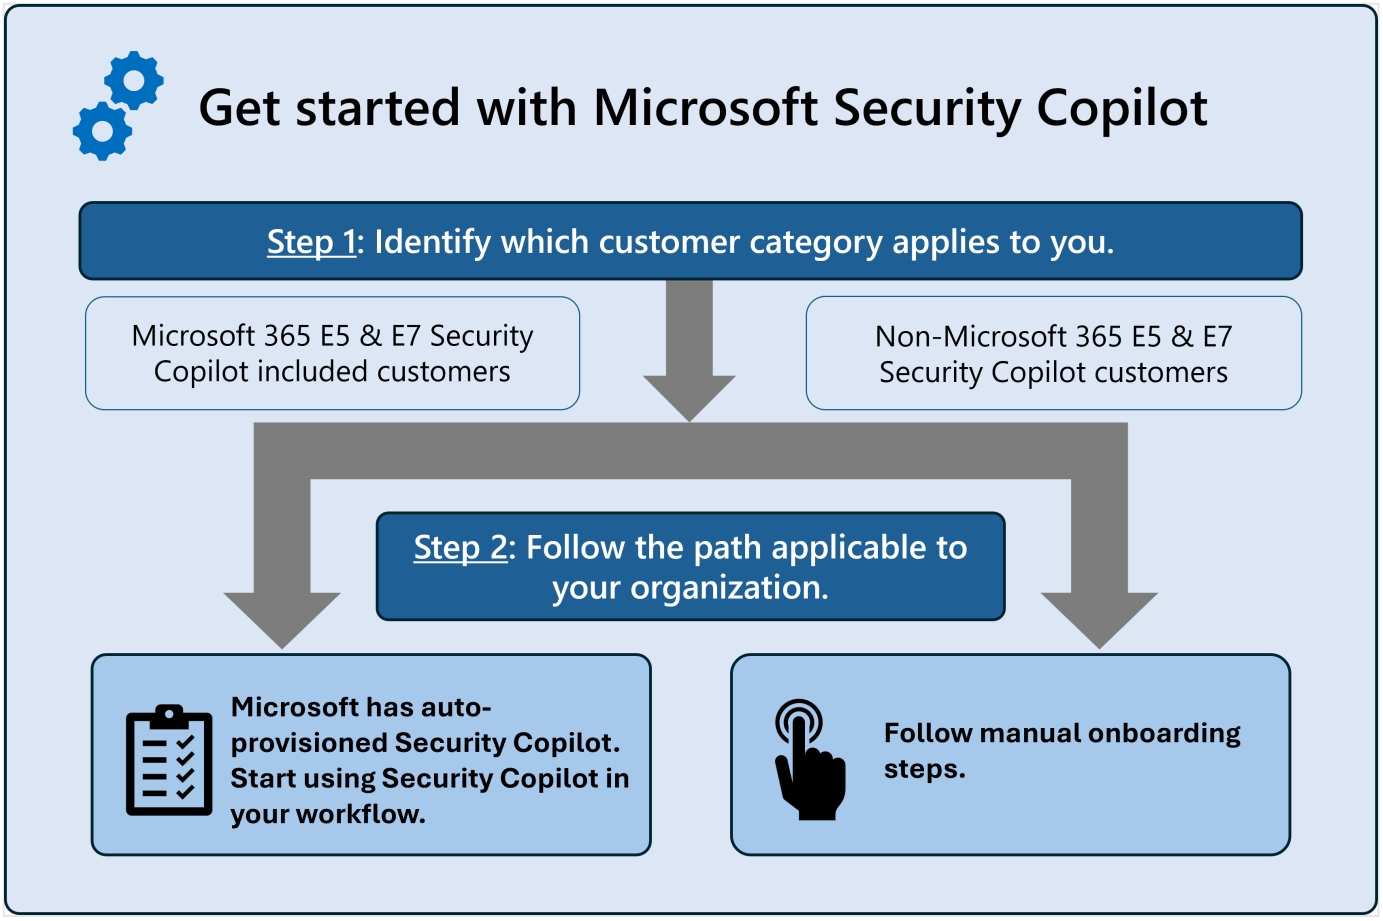

This illustration describes the steps to get started with Security Copilot:

Step 1: Identify which customer category applies to you

Before you begin onboarding to Microsoft Security Copilot, you must first determine what type of Security Copilot customer you are.

Your onboarding experience depends on this license status:

- **Microsoft 365 E5 and E7 customers** might already have Security Copilot included and automatically provisioned.

  > **Important**
  >
  > To determine if your organization is eligible, see Eligibility and access requirements.

  > **Tip**
  >
  > Your license determines onboarding steps, but your in-product experience might differ during rollout (for example, agents-first vs chat-first landing experience).

  To learn what's included with Microsoft 365 E5 and E7, see Learn about Security Copilot in Microsoft 365 E5 and E7.

- **Non–Microsoft 365 E5 and E7 customers** must provision Security Compute Units (SCUs) to enable Security Copilot.

  > **Important**
  >
  > Don't begin onboarding until you've confirmed whether your organization has a Microsoft 365 E5 or E7 license.
  > The steps you follow and whether you need to purchase capacity depend on this determination.

Step 2: Follow the path applicable to your organization

After you have confirmed your Microsoft 365 E5 or E7 license status, use the following guidance to continue.

Expand table

Microsoft 365 E5 and E7 eligible and included Security Copilot customers	Non–Microsoft 365 E5 and E7 Security Copilot customers
If you are eligible for Microsoft 365 E5 or E7 and Microsoft has onboarded you to this benefit, then Security Copilot is auto provisioned. You can start using Security Copilot in your workflows. For more information, see Understand how Security Copilot is auto provisioned for Microsoft 365 E5 and E7 customers.	Follow the manual onboarding steps. For more information, see Onboarding for non-Microsoft 365 E5 and E7 customers.

Start using Security Copilot

After onboarding is complete, you can begin using Security Copilot:

- **Agents (recommended starting point)** – In newer experiences, agents are the primary entry point and homepage.
- **Chat and prompts** – To start a chat session, go to **All history**, then select **New session**.

Role-based onboarding guidance

In the agents-first experience, the initial guidance can vary by role:

- **Administrators**:
  - Assign access and roles
  - Review usage and capacity
  - Discover and set up agents
- **Security analysts**:
  - Set up an agent
  - Browse the agent store
  - Try an agent in Microsoft Defender

First thing to do

When you first open Security Copilot, start by choosing how you want to work.

Start with an agent (recommended)

Agents are the primary entry point in newer experiences. They help you perform structured and repeatable tasks.

1. Select **Home**.
2. Select an agent or choose **Discover agents**.

3. Run the agent to begin your workflow.

Use agents when you want to:

- Automate investigation steps
- Follow a guided workflow
- Perform repeatable tasks

Start a chat session

If you want to explore or investigate using prompts, start a chat session.

1. Select **History**
2. Select **New session**
3. Enter your prompt

Use prompts when you want to:

- Ask questions
- Investigate specific issues
- Explore data interactively

What to try next

- Run an agent to complete a common task
- Ask a question using a prompt
- Save your prompts as a promptbook for reuse

Next steps

Once you've completed onboarding, you're ready to start using Security Copilot. Depending on your experience, you might begin with agents or start a chat session from **All history**.

- Explore agents, plugins, connectors, and promptbooks in your workflows
- Navigate Security Copilot
- Start with prompts (use All history to begin a chat session)

Related content

Use the following Security Copilot articles for guidance:

- Understand use cases

Last updated on 06/17/2026

Onboarding to Security Copilot for non-Microsoft 365 E5 customers

> **Note**
>
> The following steps are only applicable for customers who are **NOT** Microsoft 365 E5 customers.
> If you're a Microsoft 365 E5 included customer and have been notified that you're eligible to use Security Copilot these steps don't apply to you.
> Microsoft will automatically provision and onboard you, so that you can immediately get started with Security Copilot. For more information, see Understand how Security Copilot is auto provisioned for Microsoft 365 E5 customers.

If Security Copilot isn't included in your license, you need to provision SCUs.

You can choose from the following options to provision capacity:

- Option 1 (Recommended): Provision capacity within Security Copilot
- Option 2: Provision capacity through Azure

> **Note**
>
> Regardless of the method you choose, you'll need to purchase a minimum of 1 provisioned SCU. You can purchase up to a maximum of 100 SCUs. The recommended number of units to conduct an introductory exploration of Security Copilot is 3 units with overage set to unlimited overage. You can use the usage monitoring dashboard to keep track of usage and adjust capacity based on workload. For more information, see Manage usage.

Minimum requirements

- **Azure subscription**: Required to purchase SCUs. Create your free Azure account.
- **Supported roles**: Anyone with the supported Microsoft Entra or Microsoft Purview roles can complete onboarding steps. For more information, see Security Copilot roles.

Microsoft Entra and Microsoft Purview roles

The following Microsoft Entra and Microsoft Purview roles automatically inherit **Copilot owner** access, ensuring Security Copilot always has at least two owners.

Microsoft Entra roles:

- Billing Administrator
- Entra Compliance Administrator
- Global Administrator
- Intune Administrator
- Security Administrator

For more information on continuity in Microsoft Entra, see Manage emergency access accounts.

Microsoft Purview roles:

- Purview Compliance Administrator
- Purview Data Governance Administrator
- Purview Organization Management

As part of the Security Copilot inclusion in the Microsoft 365 E5 license, the Compliance pack is automatically added.

For capacity provisioning

- Azure Contributor or Owner role on subscription/resource group
- Security Administrator or higher in the tenant you're onboarding Security Copilot.

These roles are needed to attach the provision of SCU capacities to the Security Copilot workspace.

> **Note**
>
> Whoever sets up the default environment (including any other owner roles that that person holds) will be designated as the default owner. For example, if the person setting up Security Copilot has an Microsoft Entra Compliance Administrator role and Purview Organization Management role, then both of those roles are added as owners by default.

Option 1 (Recommended): Provision capacity through Security Copilot

When you first open Security Copilot (https://securitycopilot.microsoft.com), you're guided through the steps in setting up capacity for your organization.

1. Sign in to Security Copilot (https://securitycopilot.microsoft.com).

2. Select **Get started**.

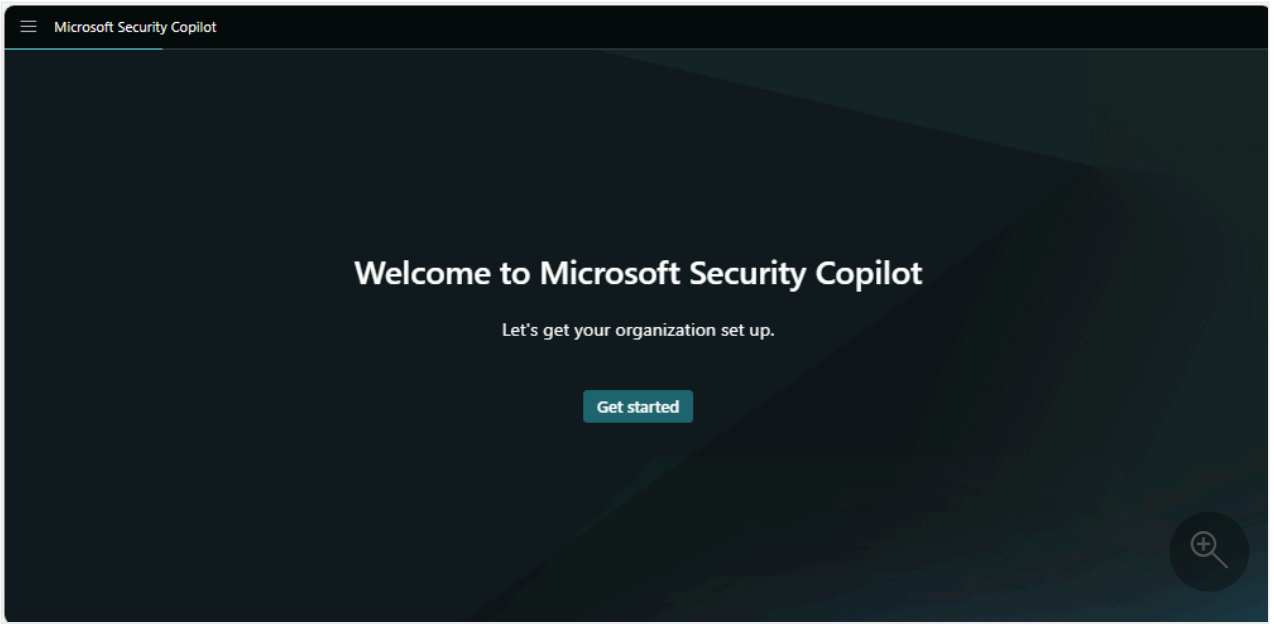

3. Set up a workspace. A workspace is a scoped, tenant-bound environment where users, automations, and agents operate. Enter a workspace name following the workspace naming convention, then select **Continue**. For more information, see Naming rules and restrictions for Azure resources.

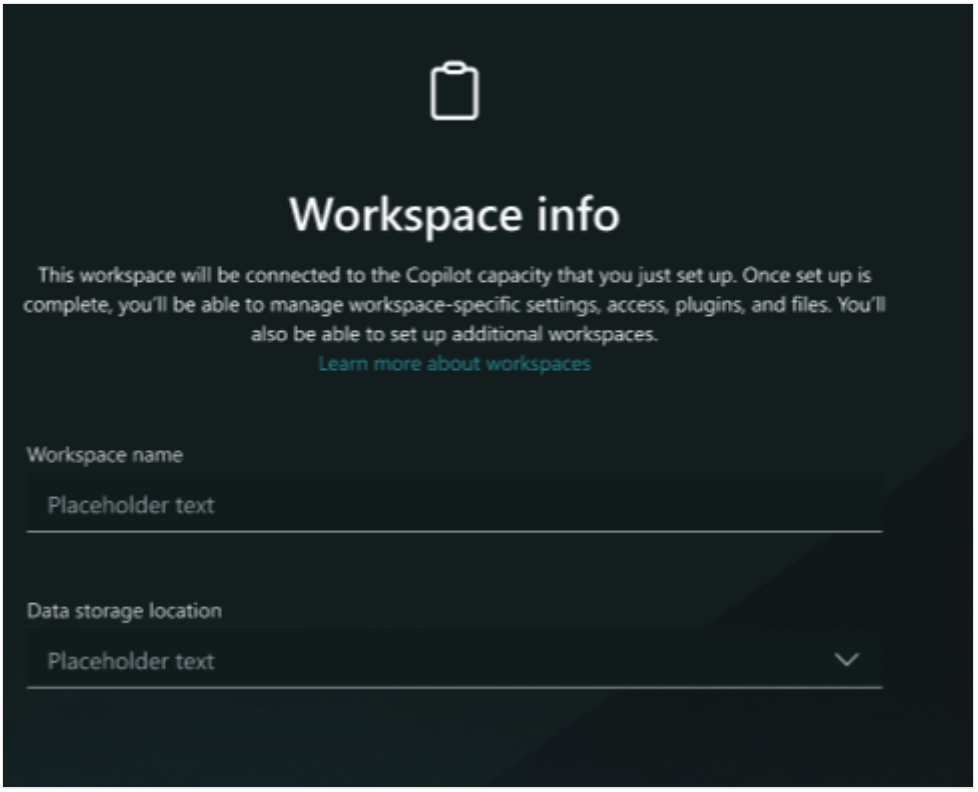

4. Set up your security capacity:
   Select the Azure subscription, associate capacity to a resource group, add a name to the capacity, select the prompt evaluation location, and specify the number of Security Compute Units (SCUs). You can also specify the number of overage units. These units

offer the flexibility to have additional security compute units available when the initially provisioned SCUs are depleted during usage spikes.

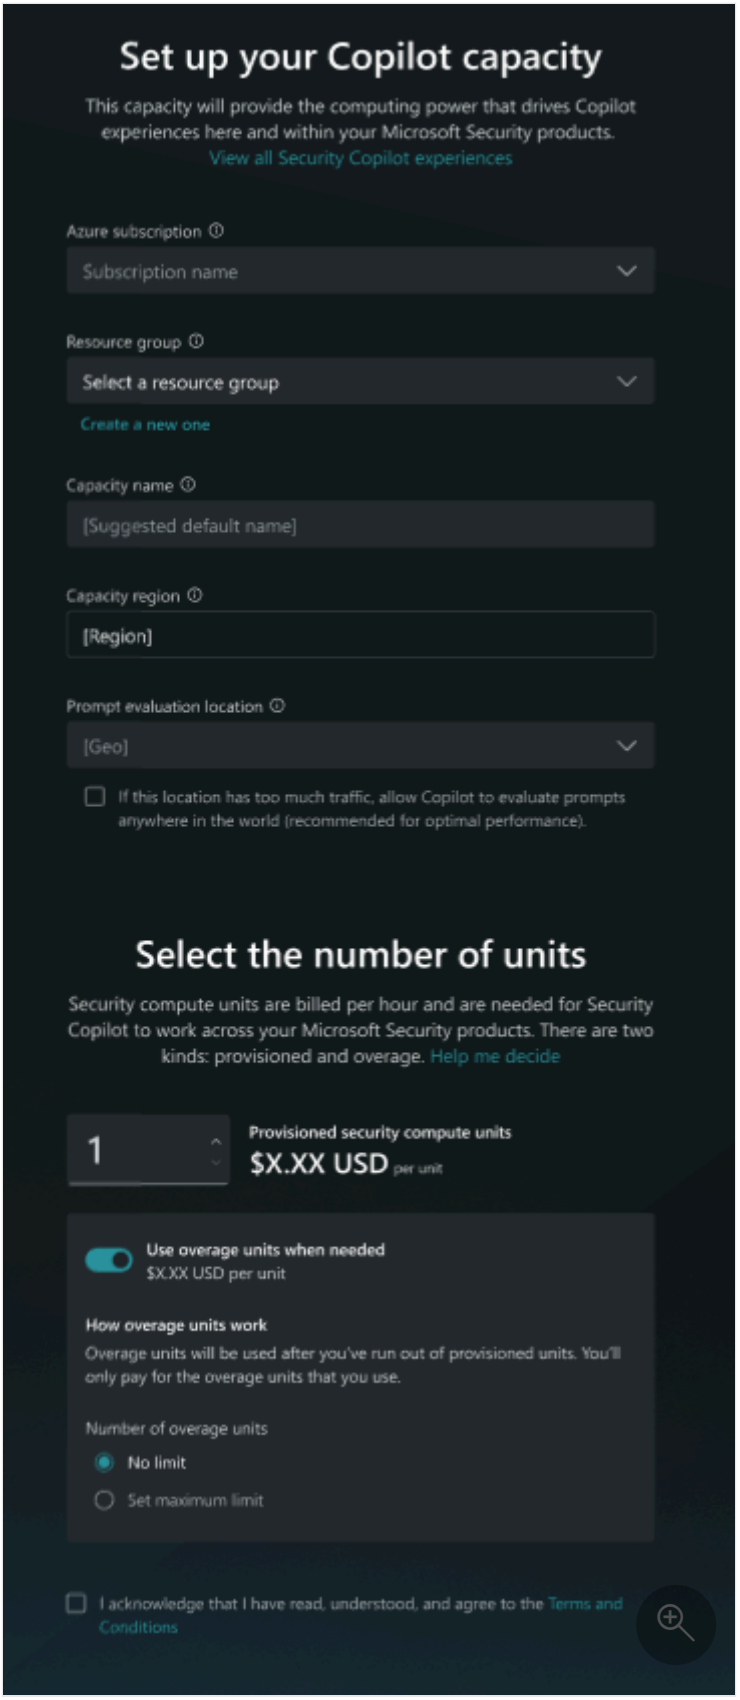

> **Note**
>
> The number of SCUs is provisioned on an hourly basis, and the estimated monthly cost is displayed.

If your selected geo location is too busy, you can also evaluate the prompts anywhere in the world. This can be done by selecting the appropriate option in the capacity creation screen.

5. Select among the data sharing options. Select **Continue**. For more information on data sharing, see Privacy and data security.

6. Confirm that you acknowledge and agree to the terms and conditions, then select **Continue**.

   After you've created the capacity, it will take a few minutes to deploy the Azure resource on the backend.

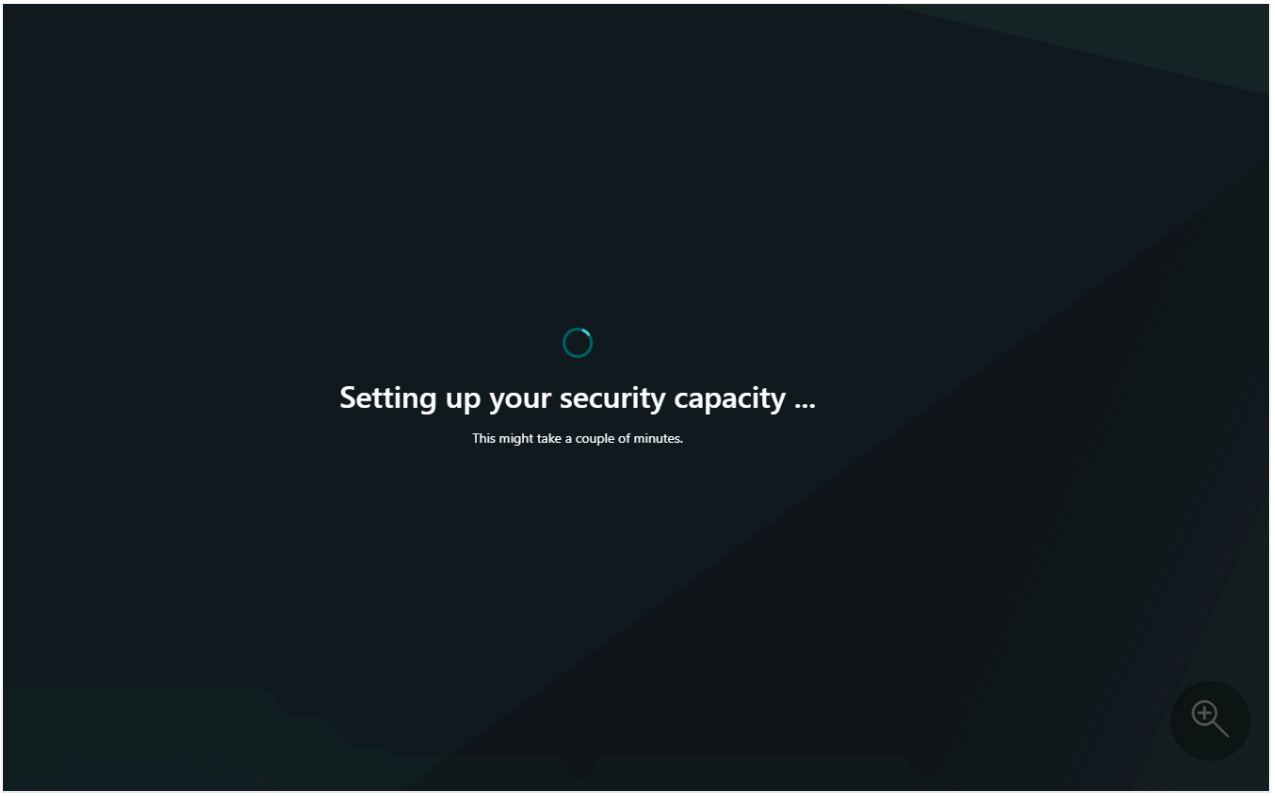

7. Select if you want Security Copilot to capture and store administrator actions, user actions, and system responses. Select **Continue**. For more information, see Accessing data from Microsoft 365 services.

   > ⓘ **Note**
   >
   > Microsoft Purview will store your Customer Data in the region where your Microsoft 365 data is stored. For more information, see Privacy and data security.

8. Review the roles that can access Security Copilot. You can choose to select other recommended owner roles or add contributors. Select **Continue**.

Copilot usage

Owners

Owner access to Copilot includes additional functionality like owner settings, plugin management, usage monitoring, role assignments, and more.

You and your roles are added as owners by default

- C — Copilot Compliance Admin User
- EC — Entra Compliance Administrator
  Role • Manage in Microsoft Entra ID
- PC — Purview Compliance Admin
  Role • Manage in Microsoft Purview
- PO — Purview Organization Management
  Role • Manage in Microsoft Purview

Add other recommended owner roles

- ☐ BA — Billing Administrator
  Role • Manage in Microsoft Entra ID
- ☐ GA — Global Administrator
  Role • Manage in Microsoft Entra ID
- ☐ IA — Intune Administrator
  Role • Manage in Microsoft Entra ID
- ☐ PDG — Purview Data Governance Administrator
  Role • Manage in Microsoft Purview
- ☐ SA — Security Administrator
  Role • Manage in Microsoft Entra ID

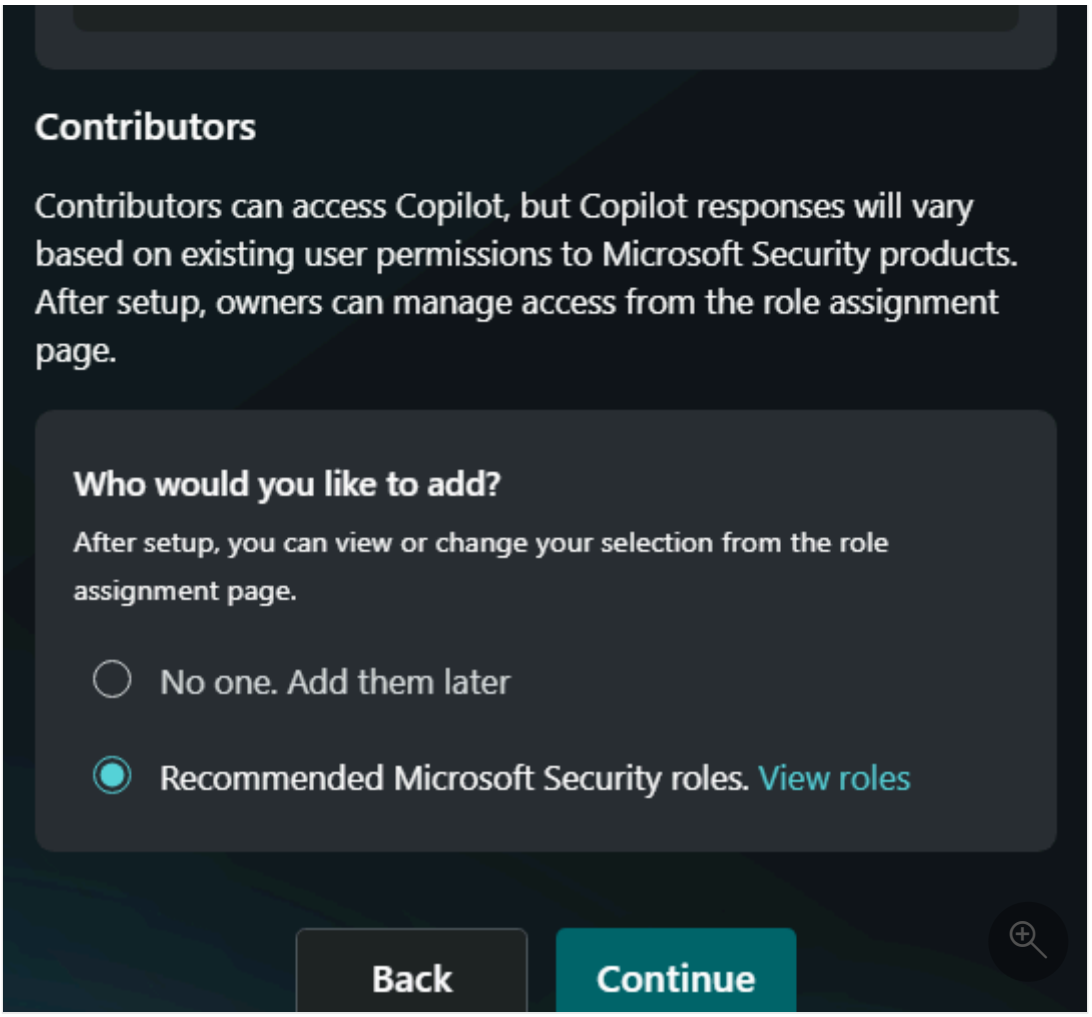

9. A confirmation page is displayed. Select **Finish**.

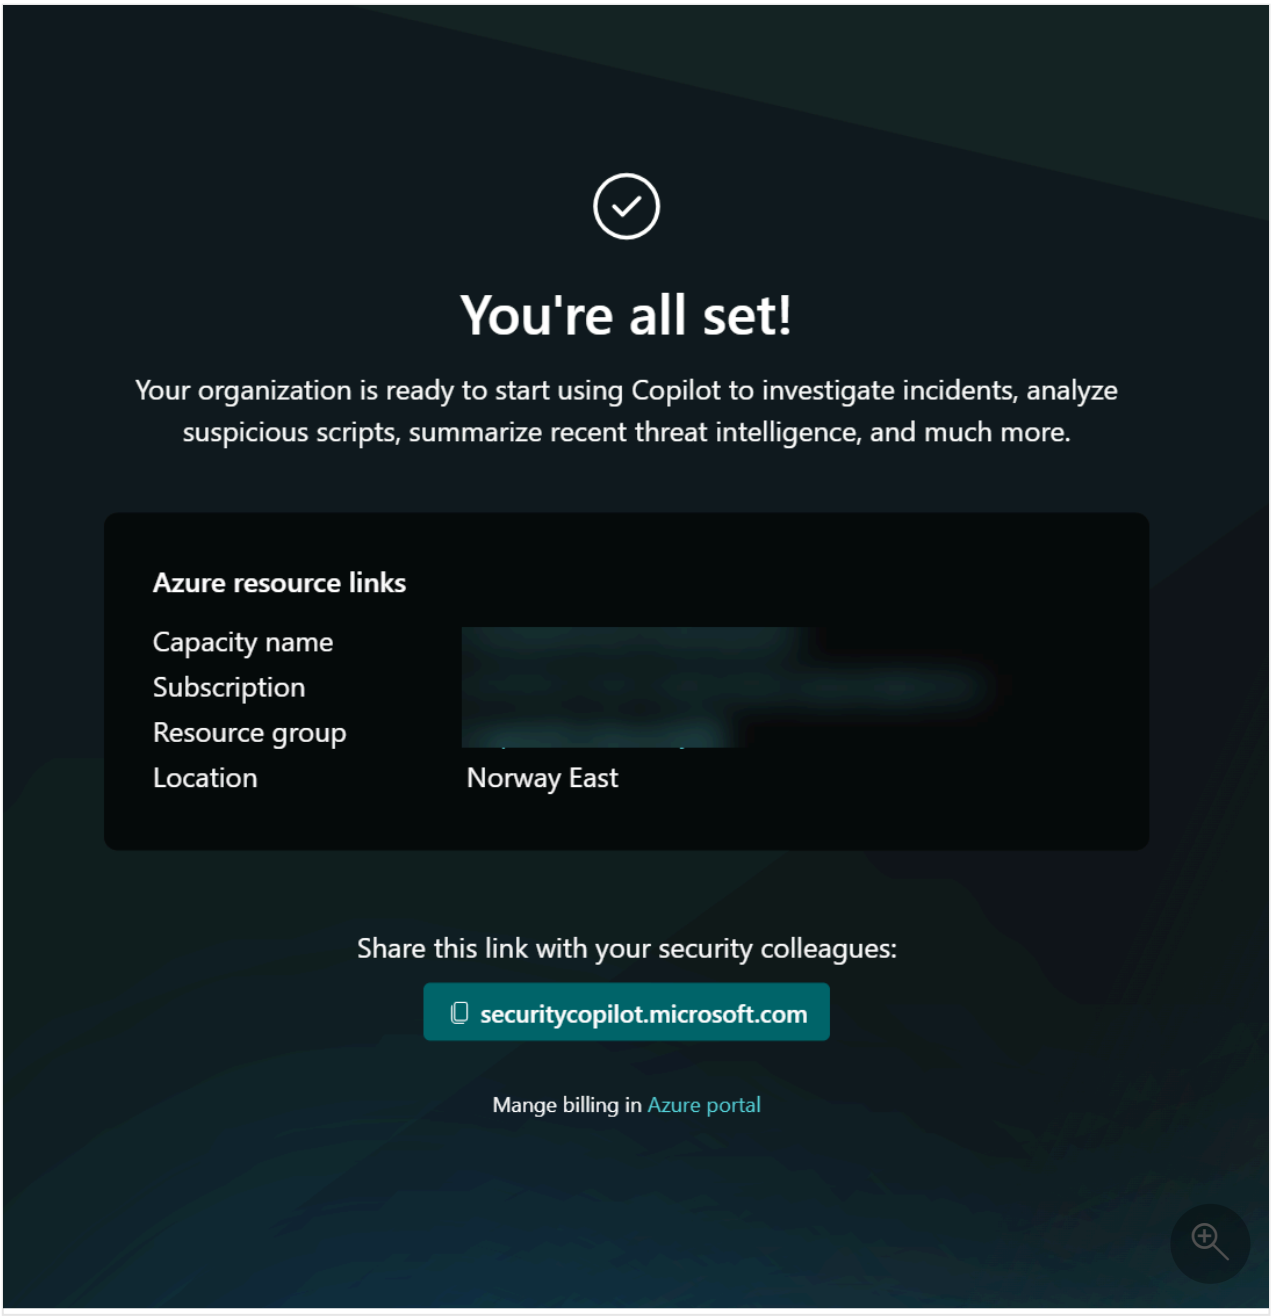

Option 2: Provision capacity in Azure

The initial setup in this method starts in the Azure portal. Then, you need to complete the setup in the Security Copilot portal.

> **Note**
>
> Billing begins as soon as capacity is created, regardless of whether the SCU is attached to an environment.

1. Sign in to the Azure portal.
2. Search for Security compute in the list of services, then select **Microsoft Security compute capacities**.

3. Select **Resource groups**.
4. Under **Plan**, select **Microsoft Security Copilot**. Then select **Create**.
5. Select a subscription and resource group, add a name to the capacity, select the prompt evaluation location and select the number of Security Compute Units (SCUs). You can also specify the number of overage units. These units offer the flexibility to have additional security compute units available when the initially provisioned SCUs are depleted during usage spikes. Data is always stored in your home tenant geo.

   > **Note**
   >
   > The number of SCUs is provisioned on an hourly basis, and the estimated monthly cost is displayed.
   >
   > If your selected geo location is too busy, you can also evaluate the prompts anywhere in the world. This can be done by selecting the appropriate option in the capacity creation screen.

6. Confirm that you acknowledge and have read the terms and conditions, then select **Review + create**.
7. Verify that all the information is correct, then select **Create**. A confirmation page is displayed.
8. Select **Finish setup in the Security Copilot portal**.

Recommended next steps

- Assign roles to users
- Try out standalone and embedded experiences
- Learn about integrations
- Check out primary use cases
- Join the Microsoft Security Copilot Customer Connection Program (CCP)

See also

- What is Security Copilot?
- Understand Security Copilot in Microsoft 365 E5
- Understand how Security Copilot is auto provisioned for Microsoft 365 E5 customers
- Learn about Security Copilot for Microsoft 365 E5 included customers

Last updated on 03/25/2026

Learn about Security Copilot for Microsoft 365 E5 and E7 included customers

At Ignite 2025, Microsoft announced that Security Copilot agents will be directly built into the flow of work for security teams using Microsoft Defender, Microsoft Entra, Microsoft Intune, and Microsoft Purview.

We introduced dozen new agents across these products to bring agentic defense across workflows, enabling autonomous and proactive protection.

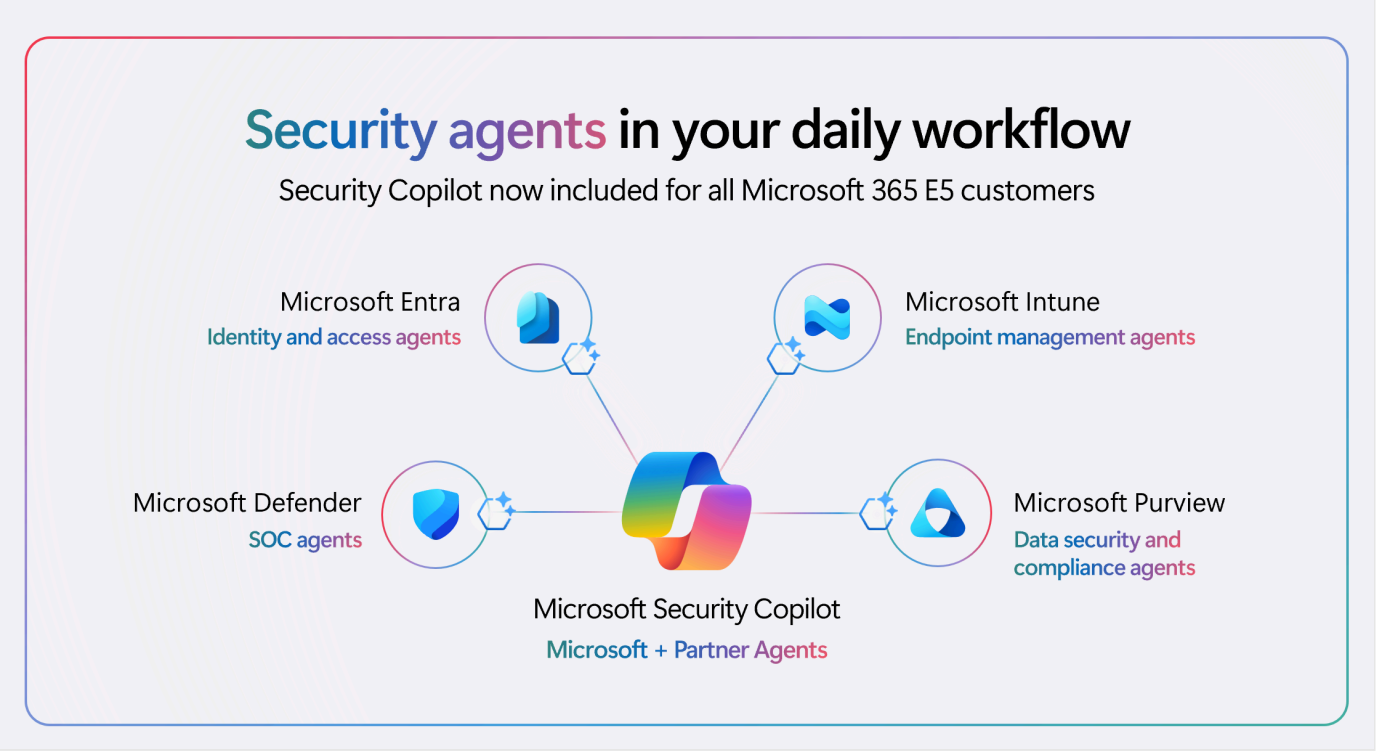

To make the agents easily accessible and help security teams get started more quickly, we're excited to announce that Security Copilot will be available to all eligible Microsoft 365 E5 and E7 customers.

Rollout started November 18, 2025, for existing Security Copilot customers with Microsoft 365 E5 as of November 18, 2025, and will continue in the upcoming months for all eligible Microsoft 365 E5 and E7 customers. For more information on what's included, see What capacity is included?.

Microsoft 365 E5 and E7 customers who already have Security Copilot can get started today with:

- Security operations in Microsoft Defender: https://security.microsoft.com
- Data security in Microsoft Purview: https://purview.microsoft.com
- Identity and access in Microsoft Entra: https://entra.microsoft.com

- Endpoint management in Microsoft Intune: https://intune.microsoft.com

All other Microsoft 365 E5 and E7 customers can get access to Security Copilot as part of a phased rollout in the upcoming months. Read the FAQ for more information.

FAQ

When will Security Copilot inclusion in Microsoft 365 E5 and E7 be available?

Security Copilot will be included for all eligible Microsoft 365 E5 and E7 customers in the upcoming months - bringing agentic AI in the daily workflow.

Rollout started November 18, 2025 for existing Security Copilot Microsoft 365 E5 customers and will continue in the upcoming months for all eligible Microsoft 365 E5 and E7 customers.

Is Security Copilot available as part of Microsoft 365 E7 license?

Yes, Security Copilot is included for all Microsoft 365 E5 and E7 customers. Rollout started November 18, 2025 for existing Security Copilot Microsoft 365 E5 customers and will continue in the upcoming months for all Microsoft 365 E5 and E7 customers.

Do customers need to take any steps to configure Security Copilot when they're enabled through Microsoft 365 E5 and E7?

Eligible Microsoft 365 E5 and E7 customers who were notified of their eligibility will automatically receive Security Copilot through zero click activation. For eligible customers, Security Copilot is provisioned automatically as part of their subscription. This means no Azure setup is needed or capacity provisioning required. Eligible customers can start using Security Copilot right away.

For more information, see Understand how Security Copilot is auto provisioned for Microsoft 365 E5 and E7.

What configuration settings are preselected during auto provisioning?

- The geography where Customer Data (such as customer prompts and responses) is stored
- The GPU processing location
- Customer Data sharing preferences (turned off by default)
- Security Copilot access to data from Microsoft 365 services (turned on by default)
- Default roles that grant Security Copilot owner and contributor access

For more information on auto provisioning, see Understand how Security Copilot is auto provisioned for Microsoft 365 E5 and E7.

How does Security Copilot use Microsoft 365 data?

After enablement, Security Copilot can process Microsoft 365 Customer Data from supported products in order to generate security insights and responses to user prompts.

As part of this experience:

- Relevant Customer Data may be **copied**, **processed**, and **stored** by Security Copilot to provide responses.
- Security Copilot uses this data only to deliver the service and in accordance with Microsoft's product terms and data protection commitments.
- Access is scoped to the products and permissions your organization has configured.

For more information on admin control and choices, see Understand how Security Copilot is auto provisioned for Microsoft 365 E5 and E7.

What roles can access Security Copilot?

If Security Copilot is auto-provisioned for your organization, these default roles inherit the Security Copilot owner and contributor access. For more information, see default roles.

Licensing and capacity

What capacity is included?

Customers with Microsoft 365 E5 and E7 will have 400 Security Compute Units (SCU) each month for every 1,000 paid user license, up to 10,000 SCUs each month at no additional cost.

This amount scales by user license count, including for customers with fewer than 1,000 user licenses. This included capacity is expected to support typical scenarios as mentioned.

- Example 1: An organization with 400 user licenses gets 160 SCUs/month.
- Example 2: An organization with 4,000 user licenses gets 1,600 SCUs/month.

> **ⓘ Important**
>
> Accessing Security Copilot through a Microsoft 365 E5 and E7 license uses the inclusion capacity model. For more information on the capacity models, see Understand Security

Copilot Compute Units and capacity.

Your tenant is assigned a Default Security Copilot Capacity.

Default Security Copilot Capacity is the automatically created inclusion capacity that appears in your tenant when Security Copilot inclusion is enabled through Microsoft 365 E5 and E7.

- It's created automatically with a default workspace
- It represents tenant-wide Security Copilot compute
- It cannot be modified
- It's shared across all users and experiences in the tenant
- It isn't billed hourly (SCUs come from the monthly Microsoft 365 E5 and E7 inclusion bucket)
- UI cost values shown for this capacity are informational only, not charges

What is a Security Compute Unit?

A Security Compute Unit is a unit measure of the compute power to run Security Copilot workloads for AI capabilities within the standalone and embedded experiences.

For more information, see Understand Security Copilot Compute Units and capacity.

Is there an example of SCU consumption under current provisioned vs inclusion models?

Let's assume an enterprise company has provisioned five SCUs for each hour under the current billing model. Within a two-hour period, the customer runs the following scenarios:

- Prompt Execution: three SCUs over 40 seconds in the first hour
- Incident Summarization: 0.5 SCU over 10 seconds in the second hour

Total usage: 3.5 SCUs across two hours. However, because billing is based on provisioned capacity, the customer is charged for 5 SCUs every hour, resulting in 10 SCUs billed for two hours.

With the inclusion model, which provides 400 SCUs per 1,000 user licenses, there's no hourly provisioned capacity. Instead, the same two scenarios would deduct 3.5 SCUs from the monthly SCU allocation.

Will Security Compute Unit (SCU) allocation reset monthly or roll over if unused?

SCU allocations reset monthly and don't roll over if unused. Customers can't carry forward unused SCUs to future months.

Eligibility and access requirements

Who is eligible for Security Copilot inclusion?

All Microsoft 365 E5 and E7 customers are eligible.

How do I know if my tenant is enabled for the Security Copilot inclusion?

You can confirm that your tenant is enabled for the Microsoft 365 E5 and E7 inclusion in several ways:

- You receive notifications in the Microsoft 365 admin center portal. The roles that receive notifications relevant to their workload are as follows:
  - Global Administrator
  - Message Center Reader
  - Security Admin (Microsoft Defender and Microsoft Entra)
  - Purview Compliance Admin (Microsoft Purview)
  - Intune Admin (Microsoft Intune)
- You see in-product banners in Microsoft Defender, Purview, Entra, Intune, or the standalone Security Copilot portal.

What do the notifications look like?

The following image is an example of the message center post:

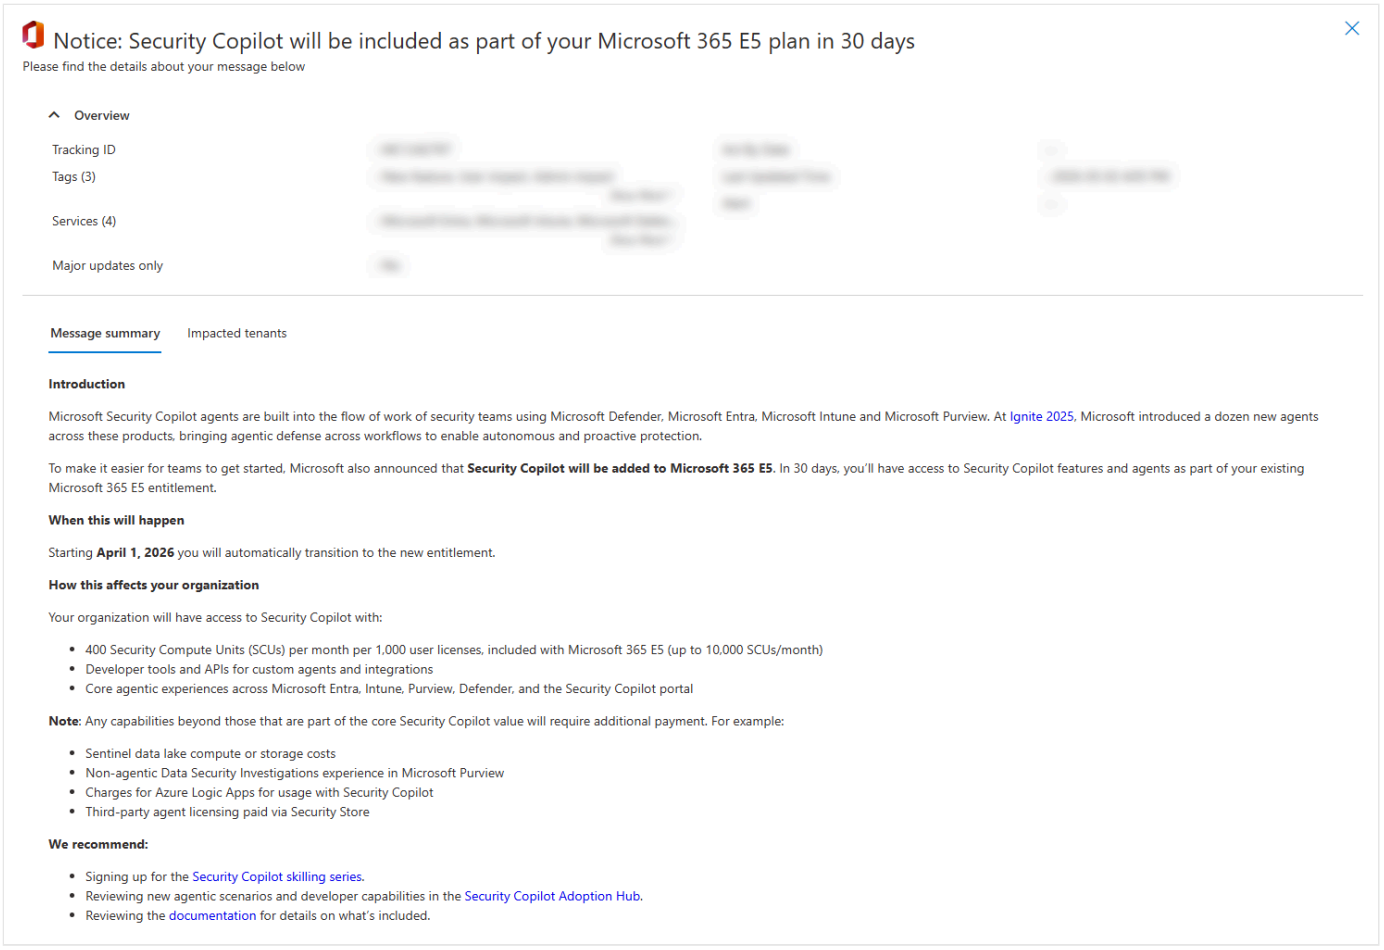

Notice: Security Copilot will be included as part of your Microsoft 365 E5 plan in 30 days

Please find the details about your message below

Overview

Tracking ID

Tags (3)

Services (4)

Major updates only

Message summary | Impacted tenants

Introduction

Microsoft Security Copilot agents are built into the flow of work of security teams using Microsoft Defender, Microsoft Entra, Microsoft Intune and Microsoft Purview. At Ignite 2025, Microsoft introduced a dozen new agents across these products, bringing agentic defense across workflows to enable autonomous and proactive protection.

To make it easier for teams to get started, Microsoft also announced that **Security Copilot will be added to Microsoft 365 E5**. In 30 days, you'll have access to Security Copilot features and agents as part of your existing Microsoft 365 E5 entitlement.

When this will happen

Starting **April 1, 2026** you will automatically transition to the new entitlement.

How this affects your organization

Your organization will have access to Security Copilot with:

- 400 Security Compute Units (SCUs) per month per 1,000 user licenses, included with Microsoft 365 E5 (up to 10,000 SCUs/month)
- Developer tools and APIs for custom agents and integrations
- Core agentic experiences across Microsoft Entra, Intune, Purview, Defender, and the Security Copilot portal

Note: Any capabilities beyond those that are part of the core Security Copilot value will require additional payment. For example:

- Sentinel data lake compute or storage costs
- Non-agentic Data Security Investigations experience in Microsoft Purview
- Charges for Azure Logic Apps for usage with Security Copilot
- Third-party agent licensing paid via Security Store

We recommend:

- Signing up for the Security Copilot skilling series.
- Reviewing new agentic scenarios and developer capabilities in the Security Copilot Adoption Hub.
- Reviewing the documentation for details on what's included.

As an existing Security Copilot customer with already provisioned capacity, should I delete my current capacity now that I qualify for Security Copilot inclusion with Microsoft 365 E5 and E7?

No. If you're an existing customer, we recommend that you *don't delete* the Security Copilot capacity you previously provisioned. Retain the capacity to ensure that you're able to use the product.

How do customers prepare their organization and tenants to configure Security Copilot?

- Security Copilot is automatically provisioned for all eligible Microsoft 365 E5 and E7 tenants. No Azure setup or consent flows are required.
- Customers see in-product banners and guided onboarding to help them get started. No manual provisioning of SCUs is needed.

Is there a minimum Microsoft 365 E5 and E7 user license count consumption commitment to qualify for the Security Copilot inclusion?

No. All customers with Microsoft 365 E5 and E7 licenses qualify for the offer, regardless of user license count.

Is Security Copilot still available to customers without qualifying Microsoft 365 E5 and E7 licenses?

Yes. Customers without Microsoft 365 E5 and E7 licenses can continue to access Security Copilot using the existing pricing model. Learn more about Security Copilot pricing.

Are Microsoft Sentinel customers without Microsoft 365 E5 and E7 licenses excluded from eligibility?

Yes. Microsoft Sentinel customers without Microsoft 365 E5 and E7 licenses aren't eligible.

Microsoft 365 E5 and E7 eligible customers get Security Copilot at no additional cost as part of their license.

Usage and billing

What happens if a customer exceeds the included SCUs?

Usage beyond the allocated SCUs will be throttled at a future date. Customers have an option to scale beyond the allocated amount at $6 USD for each SCU on a pay-as-you go basis at that time. Customers will get a 30-day advanced notification when this option is available.

Do customers need to provision SCUs separately if they're eligible for this offer?

No. Eligible customers don't need to manually provision SCUs.

Do customers need to manage other licenses or setup?

No. Security Copilot is ready to use inside the tools teams already rely on, with no action needed and guided onboarding for immediate value.

Features and capabilities

What capabilities are included?

- Core Experiences: All chat, promptbook, and agentic scenarios across Microsoft Entra, Microsoft Intune, Microsoft Purview, Microsoft Defender, and the standalone Security Copilot portal. Additionally, customers with Microsoft 365 E5 and E7 who also use Microsoft

Sentinel can apply their included SCU allocation to run Security Copilot scenarios in Microsoft Sentinel.

- Developer Experiences: Tools such as Agent Builder and APIs for creating custom agents, promptbooks, and integrations via MCP and Graph APIs.
- Partner-built agents: SCU costs included until further notice.

What isn't covered as part of this inclusion?

Any capabilities beyond those that are part of the core Security Copilot value require additional payment.

For example:

- Microsoft Sentinel data lake compute or storage costs.
- Non-agentic Data Security Investigations experience in Microsoft Purview.
- Charges for Azure Logic Apps for usage with Security Copilot.
- Partner-built agent licensing paid through Security Store.
- Some agents may require prerequisites that involve products outside of Microsoft 365 E5 and E7, which aren't covered.

Are Partner-built agents included?

Partner-built agents may require a separate license from the partner (purchased via Security Store); SCU costs for third-party agents are included, but this might change in the future as we roll out Security Copilot inclusion more broadly.

What can customers do with Security Copilot?

Security Copilot supports various use cases, such as:

- SOC teams can simplify phishing triage and incident response with agents that classify alerts, summarize threats, and guide remediation.
- Identity teams can strengthen access controls and reduce risk with agents that optimize Conditional Access policies and automate access reviews.
- Data security teams can strengthen data protection and compliance with agents built to help triage alerts and discover sensitive data.

- IT admins can better maintain secure, compliant endpoints more easily, with agents that help with tasks like assessing changes before they impact productivity.
- For all other use cases, you can build custom Security Copilot agents, tailored to your organization's unique workflows.

Read more about Security Copilot use cases in the Adoption Hub.

Product experience

Where can customers track their usage?

Customers are able to track their usage in the in-product usage dashboard in the Security Copilot standalone portal. For more information, see Monitor security compute units use.

What is the in-product experience for customers?

Customers are presented with a recommendation card that suggests high-value use cases for using Security Copilot in a particular product experience. With the provided capacity, customers can use Security Copilot agents to automate tasks like user-submitted phish triage, alert triage, access reviews, and vulnerability remediation.

Will Microsoft Security Copilot agents be automatically enabled?

No, a customer's organization is required to set up and deploy agents in the relevant standalone or embedded experiences. For more information, see Manage agents.

Are there other Security Copilot Integration Experiences and Configurations available?

Security Copilot integrates seamlessly with multiple Microsoft 365 and Microsoft security products that your organization has licensed. This integration allows users to query information directly from those products in the Security Copilot standalone and embedded experiences. Copilot copies, processes, and stores this data.

Customers can choose not to share their Microsoft 365 service data with Copilot in the **Owner settings** page. However, the organization won't be able to use Copilot with Microsoft 365 products such as like Microsoft Purview.

Special scenarios

What does this mean for Managed Service Providers (MSPs)?

If an MSP after November 18, 2025 deploys Security Copilot for eligible customers that are either Microsoft 365 E5 and E7 customers, then that customer also is no longer separately billed for Security Copilot and can continue to use all Security Copilot capabilities and experiences.

Last updated on 06/19/2026

Understand how Security Copilot is auto provisioned for eligible Microsoft 365 E5 and E7 customers

If Security Copilot is included in your Microsoft 365 E5 or E7 license and your tenant has been enabled, Microsoft will automatically provision and onboard, so that you can immediately get started with Security Copilot in the flow of your work. You can start using Security Copilot in the following experiences:

- Data security in Microsoft Purview: https://purview.microsoft.com
- Identity and access in Microsoft Entra: https://entra.microsoft.com
- Endpoint management in Microsoft Intune: https://intune.microsoft.com
- Access Security Copilot portal: https://securitycopilot.microsoft.com

For more information on Security Copilot capabilities included in your license, see Learn about Security Copilot inclusion in Microsoft 365 E5 and E7.

What gets autoprovisioned?

As part of the autoprovisioning process, the following steps are done for you so that you can easily get started:

- Creation of a default Security Copilot virtual capacity and a default workspace
- Assign default workspace settings
  - Customer Data storage location
  - Customer Data sharing preferences
- Assign additional settings
  - Prompt evaluation location
  - Accessing data from Microsoft 365 services
- Assign default roles for the default workspace

Workspace and capacity

A workspace is a logical container that defines three things:

- Where your data is stored (Customer Data storage location)
- Which capacity powers your experience, and
- Who has access (Security Copilot owners and contributors are scoped per workspace)

When your customer tenant is provisioned, Security Copilot automatically creates a default workspace and links it to your Security Copilot capacity. If your tenant already has workspaces, the earliest-created workspace becomes the default.

Security Copilot experiences across Microsoft Defender, Entra, Purview, and Intune use the default workspace.

Customer Data storage location

Customer Data storage location is the geography where customer's prompts, responses, and any Customer Data associated with a workspace are stored. For the default workspace that Microsoft creates for you, the Customer Data storage location that is preselected is your Microsoft Entra geography.

If your customer tenant has Microsoft 365 geo override (Preferred Data location, Go Local move programs, Advanced Data Residency (ADR add-on) etc.), then we use this Microsoft 365 geo override and map to one of the supported geos by Security Copilot.

You can view the Customer Data storage location that is preselected for your Security Copilot default workspace within the **Owner settings** page in the Security Copilot portal. For more information, see Configure owner settings.

Customer Data sharing preferences

By default, the following settings are set to **Off**:

- **Allow Microsoft to capture data from Security Copilot to validate product performance using human review**
- **Allow Microsoft to capture and human review data from Security Copilot to build and validate Microsoft's security AI model**

You can change the data sharing settings within the **Owner settings** page in the Security Copilot portal.

For more information, see Privacy and data security.

Prompt evaluation location

Prompt evaluation location determines where your prompts are processed (evaluated) using GPU resources. When a user enters a prompt, Security Copilot evaluates it on GPU clusters in Azure datacenters. This configuration setting is preselected during autoprovisioning. By default,

- If your tenant's Customer Data storage location is in the EU, then prompts are processed in the EU
- If your tenant's Customer Data storage location is not in the EU, then prompts are processed globally (US, UK, EU, ANZ) depending on locality and GPU availability.

Accessing data from Microsoft 365 services

This setting allows customers to configure whether Security Copilot can query information directly from Microsoft 365 services such as Microsoft Purview.

By default, this setting is set to ON.

If you want to disable Security Copilot's access to Microsoft 365 service data, administrators can adjust this in Owner settings.

> **Note**
>
> Choosing to disable Security Copilot or limit it's access to Microsoft 365 service data will limit Security Copilot's ability to work with Microsoft 365 products.

What admins should do once enabled:

- Review Owner settings to confirm your organization's preferred configuration.
- Inform users that Security Copilot is available as part of your Microsoft 365 E5 and E7 license and how it can be used within your organization's policies.

For more information, see Privacy and data security.

Default roles

The following Microsoft Entra, Intune, Defender, and Purview roles automatically inherit **Security Copilot owner** access:

- Global Administrator
- Security Administrator
- Conditional Access Administrator

- Intune Administrator
- ComplianceBundle
  - Microsoft Entra Compliance Administrator
  - Purview Compliance Admin
  - Purview Organization Management
- Purview Data Governance Administrator

The following Microsoft Defender and Purview roles automatically inherit **Security Copilot contributor** access:

- Microsoft Defender roles
  - *Any custom Microsoft Defender XDR roles that include Security Copilot permission*. For more information, see Create custom roles.
- Microsoft Purview roles. For more information, see Purview roles in Security Copilot.
- Microsoft Intune roles
  - Any Microsoft Intune roles that include Security Copilot permission. For more information, see Intune RBAC roles in Security Copilot.

These roles that were added by default, can be modified within the **Role assignment** page of Security Copilot portal.

Last updated on 06/18/2026

Microsoft Security Copilot Security Compute Units and capacity

Security Compute Units (SCUs) are the compute resources that power Microsoft Security Copilot experiences. This article explains how SCU capacity works, compares provisioned, overage, and Microsoft 365 E5 and E7 inclusion models, and covers billing and monitoring so you can choose the right setup for your organization.

Security compute units

Security Compute Units (SCUs) represent the compute capacity required to run Security Copilot workloads. SCUs are the units of compute resources needed to deliver dependable and consistent performance across Security Copilot experiences.

You consume SCUs in the following scenarios:

- Using the Security Copilot standalone portal
- Running embedded Security Copilot experiences
- Invoking Security Copilot through Microsoft-developed agents
- Invoking Security Copilot through partner-developed agents
- Running other Security Copilot features and capabilities

Depending on your scenario and licenses, you might need to provision or purchase SCUs to use Security Copilot.

For eligible and onboarded Microsoft 365 E5 customers, Security Copilot is auto provisioned with a default capacity. Non-Microsoft 365 E5 customers need to purchase SCUs by provisioning them.

Capacity

Capacity in Security Copilot refers to the computational resources provisioned to run the service, measured in SCUs. You can scale capacity up or down and monitor usage over time to ensure you have the right level of SCUs for your organization's needs. For more information, see Managing usage.

Non-Microsoft 365 E5 and E7 customers

Provisioned capacity is the baseline allocation of SCUs that you set up to support your organization's ongoing use of Security Copilot. You must provision this capacity before you can use Security Copilot.

Overage capacity provides additional SCUs when your usage exceeds your provisioned capacity, such as during peak investigation periods. You can add overage capacity on top of your provisioned capacity to maintain uninterrupted access when demand increases.

The following section provides information about provisioned and overage capacity.

Provisioned capacity

Key characteristics

- SCUs are provisioned in advance (minimum of one SCU)
- Capacity refreshes every full billing hour (for example, 9:00–10:00)
- Unused SCUs expire at the end of the hour and don't roll over
  - Example: If you set up a capacity plan with three SCUs, you get three SCUs each full hour. Refresh happens on fixed clock-hour blocks (9:00–10:00, 10:00–11:00), not rolling blocks (11:05–12:05).
- If provisioned SCUs run out, Security Copilot requests stop unless another capacity type is available

Provisioned capacity is billed per hour and is best for steady, predictable usage.

For more information, see How provisioned and overage SCUs are billed.

Overage capacity

Key characteristics

- SCUs are consumed as used
- Capacity is consumed as Security Copilot activity occurs
- Can be configured with a maximum limit or set to unlimited

Overage capacity helps absorb unexpected usage spikes without requiring extra hourly provisioning.

For more information, see How provisioned and overage SCUs are billed.

Microsoft 365 E5 and E7 customers

Default Security Copilot Capacity is the automatically created inclusion capacity that appears in your tenant when Security Copilot inclusion is enabled through Microsoft 365 E5 and E7. For more information, see Understand how Security Copilot is auto provisioned for Microsoft 365 E5 and E7 customers.

How provisioned and overage SCUs are billed

Security Copilot operates on a provisioned and overage capacity. Provisioned capacity is billed by the hour while the overage capacity is billed on usage.

You can flexibly provision Security Compute Units (SCUs) to accommodate regular workloads and adjust them anytime without long-term commitments.

To manage unexpected demand spikes, you can set an overage amount to ensure that extra SCUs are available when initially provisioned units are depleted during unexpected workload spikes. Overage units are billed on-demand and can be set as unlimited or a maximum amount. This approach enables predictable billing while providing the flexibility to handle both regular and unexpected usage.

Provisioned capacity billing

Billing is calculated on hourly blocks based on provisioned capacity rather than by 60-minute increments and has a minimum of one hour. Any usage consumed within the same hour is billed as a full SCU for provisioned capacity, regardless of start or end times within that hour.

For more information, see Microsoft Security Copilot pricing and the Pricing calculator page.

Overage capacity billing

For overage units, SCUs are billed up to one decimal increments for the exact consumed units. Consumed units aren't rounded up to whole numbers. This means that you're charged precisely based on your usage (to one decimal place).

For instance, if you provision an SCU at 9:05 a.m., then deprovision it at 9:35 am, and then provision another SCU at 9:45 am, you are charged for two units within the 9:00 a.m. to 10:00 a.m. hour. To maximize usage, make SCU provisioning changes at the beginning of the hour. For more information, see Manage usage.

For more information, see Microsoft Security Copilot pricing and the Pricing calculator page.

Example billing scenarios for overages

This section provides some scenarios to illustrate how overages are billed.

An enterprise company provisioned four SCUs and set an overage limit of six SCUs to stay within the monthly budget.

- **Scenario 1**:
  - A user runs a prompt consuming 3 SCUs and uses the incident summarization in Defender consuming 0.5 SCU.
  - The total consumption is calculated as 3.5 SCUs. However, the charge for that hour will be based on four provisioned SCUs.

Expand table

Activity	SCU consumed
Runs a prompt	3.0 SCUs
Uses incident feature	0.5 SCU
Total Consumption	3.5 SCUs

- **Scenario 2**:
  - Building on Scenario 1, a user also runs a promptbook consuming an additional 3.7 SCUs, bringing the total to 7.2 SCUs for the hour.
  - The charge for that hour will now be based on four provisioned SCUs, and 3.2 overage SCUs.

Expand table

Activity	SCU consumed
Runs a prompt	3.0 SCUs
Uses incident feature	0.5 SCU
Runs a promptbook	3.7 SCUs
Total Consumption	7.2 SCUs

Managing and monitoring capacity

You can manage Security Copilot capacity by increasing or decreasing provisioned SCUs through the Azure portal or the Security Copilot portal. Security Copilot also provides a usage monitoring dashboard for tracking consumption over time.

For more information, see Managing usage.

Last updated on 04/29/2026

Navigate Security Copilot

The Security Copilot platform includes navigation features that help you work efficiently. This article covers two landing experiences, home menu options, how to start chat sessions, and prompt features.

Understand the landing experience

When you open Security Copilot, the landing page can vary by onboarding and rollout stage.

- **Agents-first experience (new)** – Agents are the default homepage, with guided actions to help you get started.
- **Chat-first experience (existing)** – Chat and prompts are available directly on the homepage.

Agents-first experience

In the agents-first experience, the **Agents** area is the primary entry point when you open Security Copilot. From the **Agents** area, you can browse available agents, view active agents in your workspace, and access setup options.

> **Note**
>
> The agents-first homepage doesn't include the general chat prompt bar. To start a general chat session, go to **All history** and select **New session**.

Chat-first experience

In the chat-first experience, the homepage opens directly to the chat prompt bar where you can start typing prompts immediately. The **Agents** area is accessible from the home menu but is not the default view.

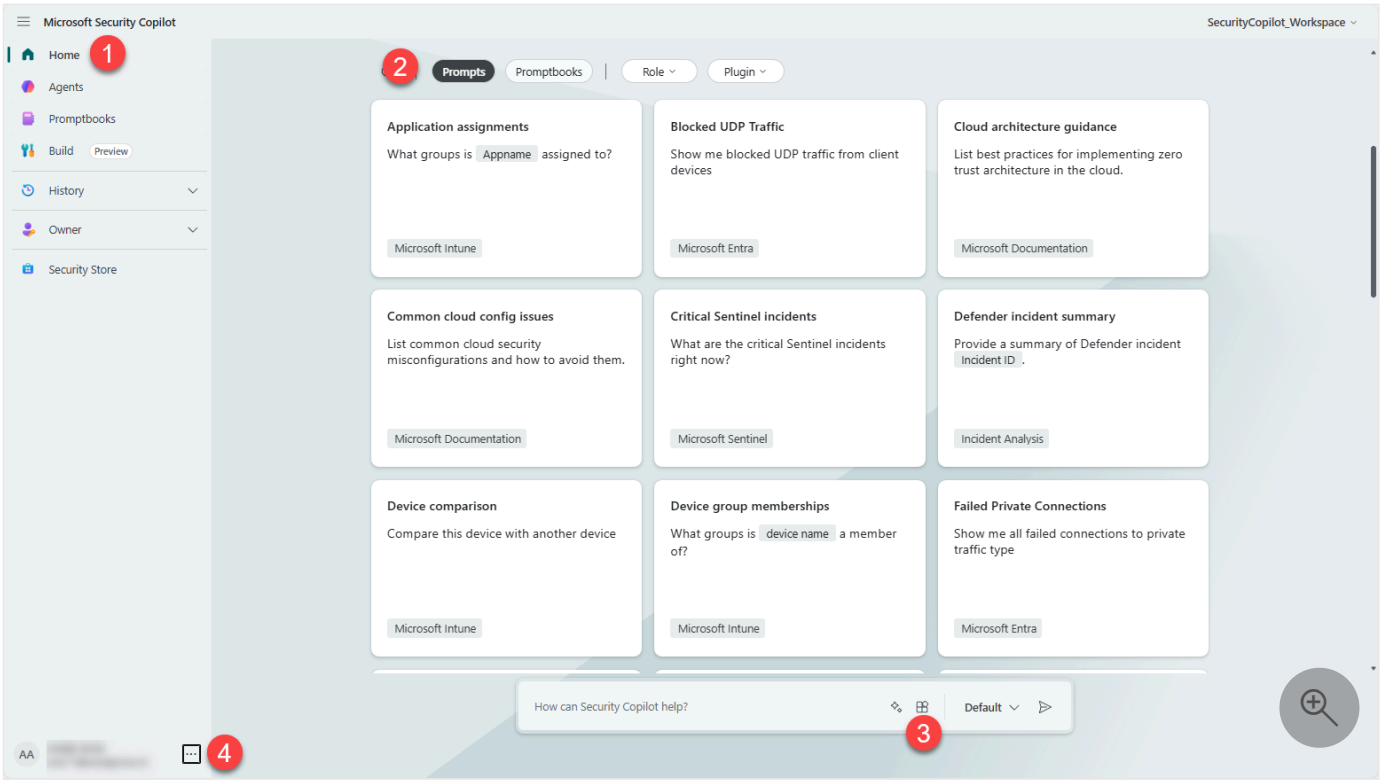

Explore interface landmarks

Homepage landmarks vary by experience: agents-first or chat-first.

Agents-first landmarks

In the agents-first experience, navigation centers on the **Agents** area and the home menu.

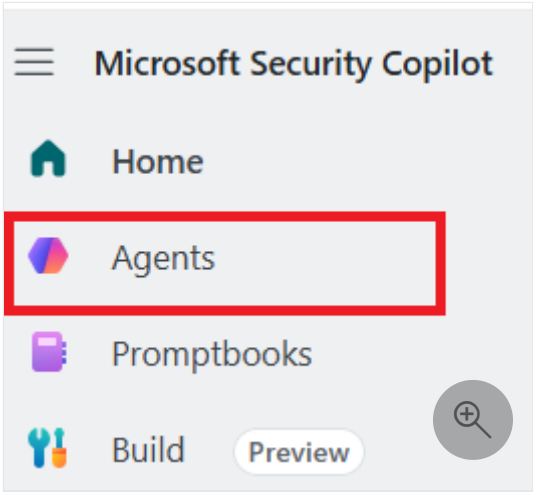

For more information on finding and using agents, see Discover agents.

Chat-first landmarks

In the chat-first experience, the homepage includes key landmarks.

In this context, plugins are data connectors and capabilities that Security Copilot can use while responding to prompts.

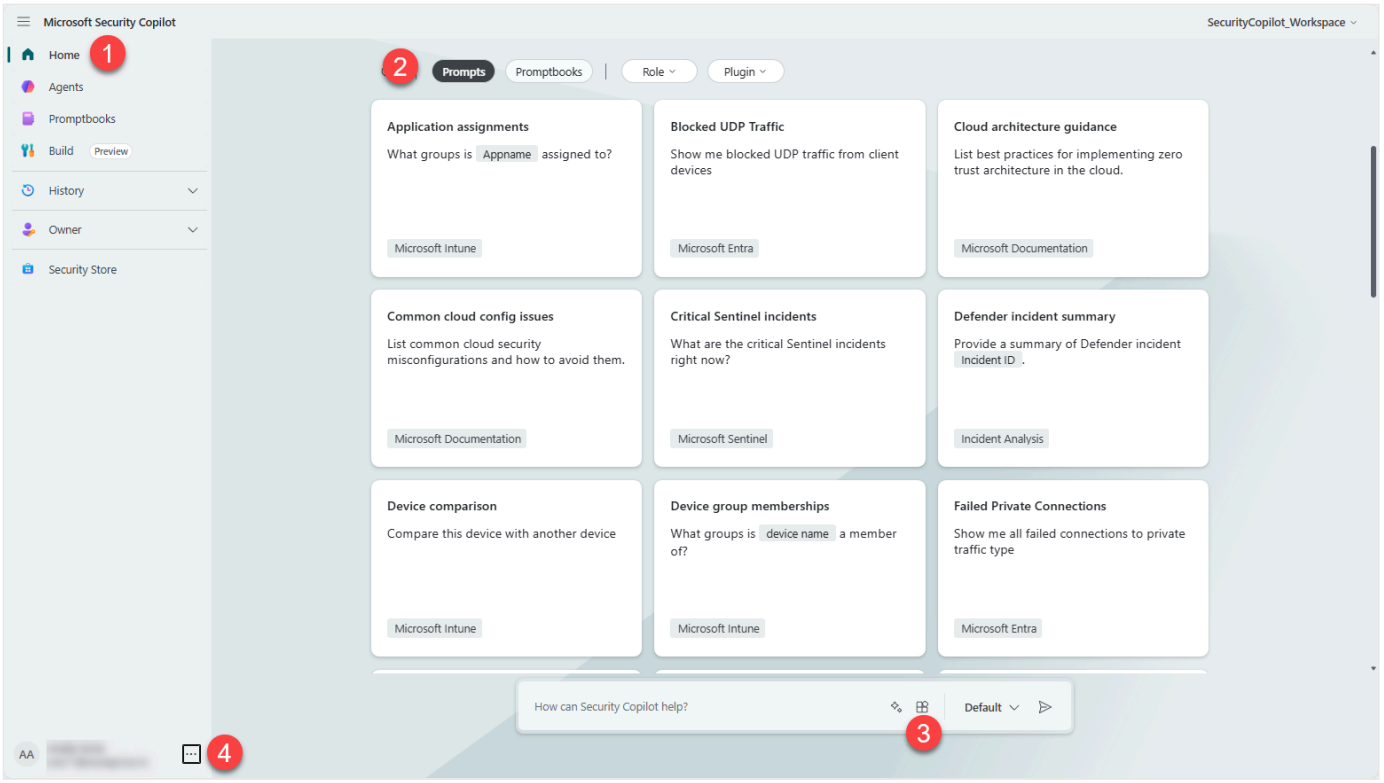

1. Use the home menu
2. Use prompts
3. Sources
4. Help

After you submit your first prompt, additional features become available.

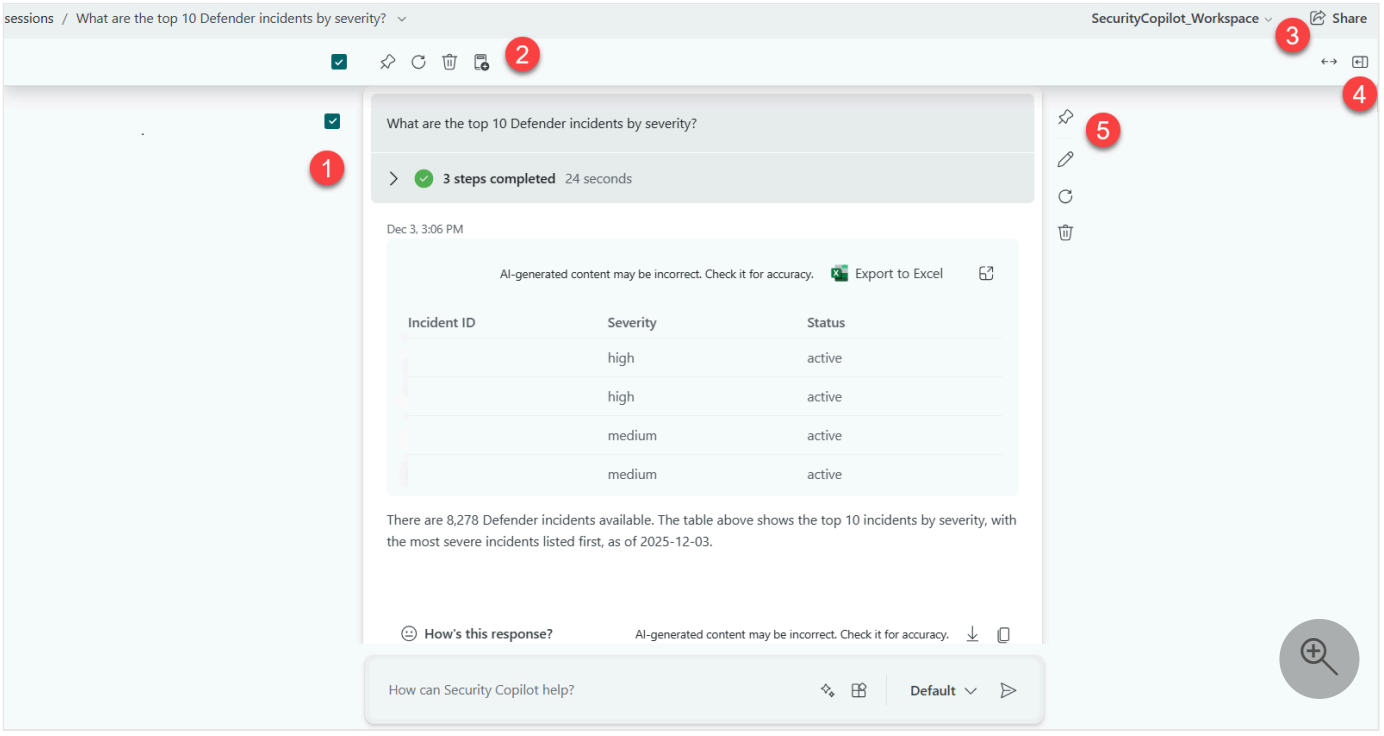

1. View the process log to see what Security Copilot is up to.
2. Create a promptbook based on the selected prompts.
3. Share your session with your team.
4. Side panel where the Pin board displays.
5. Pin a prompt to the pin board.

Use the home menu

The home menu provides access to key areas of Security Copilot. The same options are available in both experiences, but the default homepage view differs.

- **Home** gets you back to the initial page.
- **Agents** – In the agents-first experience, this is the default homepage. In the chat-first experience, select **Agents** from the home menu to access the agent library. Browse the agent library, view agents active in your workspace, and, depending on your role, set up and manage agents. For more information, see Discover agents and Setup and manage agents.
- **Promptbooks** - For more information, see Using promptbooks.
- **Build** - For more information, see Build custom agents.
- **All history** - Provides access to your session history and allows you to create a new session.
- **Owner** - This menu is available only if you have the Security Copilot Owner role.
  - **Owner settings** - For more information, see Configure owner settings.
  - **Plugin settings** - For more information, see Manage plugin restrictions.
  - **Role assignment** - For more information, see Assign roles.
  - **Usage monitoring** - For more information, see Monitor security compute units use.
- **Settings** - Includes access to set the theme, time zone, and language preferences.
- **Tenant switching** - Allows you to switch between tenants. Requires specific authentication.

When a user signs in to Security Copilot in their home tenant, but the security data is in a different tenant, they must explicitly switch to the tenant that contains the security data.

With the tenant switch feature, an external member or guest account that has the appropriate role in the target tenant is authorized to conduct sessions by using data in that tenant.

The Security Copilot portal reloads and initially shows the tenant ID in the URL, for example:

```
https://securitycopilot.microsoft.com/?tenantId=xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx
```

> **Tip**
>
> Manually modifying the tenant ID portion of the Security Copilot URL will also switch the tenant.

For more information, see the following articles:

- Understand authentication
- Grant MSSP access

Access chat and sessions

> **Note**
>
> If you have the **agents-first experience** (the **Agents** page opens first), the homepage doesn't include the general chat prompt bar. Use the steps in this section to start a general chat session. If you have the **chat-first experience** (the prompt bar opens first), the prompt bar is available directly on the homepage. You can also use **All history** to resume a previous session or start a new one.

To start a chat session:

1. Select **All history**.
2. Review your previous sessions.
3. Select **New session**.

You can continue previous sessions or start a new one from the **All history** view.

Use prompts

Use these resources for promptbooks, starter prompts, and prompt-writing guidance.

> **Note**
>
> In some experiences, prompts are accessed through a chat session instead of directly from the homepage.

- Try promptbooks
- Prompting in Security Copilot
- Create effective prompts

Manage sources

Manage plugin enablement and file sources.

To manage plugins and file inputs, see:

- Manage plugins
- File uploads

Get help

Access the support link and documentation articles from the **Get help** menu.

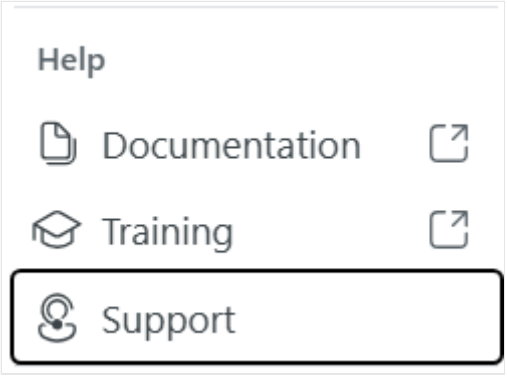

For documentation, training videos, and support options, see:

- Security Copilot documentation
- Security Copilot training videos
- Security Copilot support options

Prompt features

The following features are available after you submit a prompt.

View the process log

The process log appears directly under your prompt after you send it. As the response forms, notice the selected plugin.

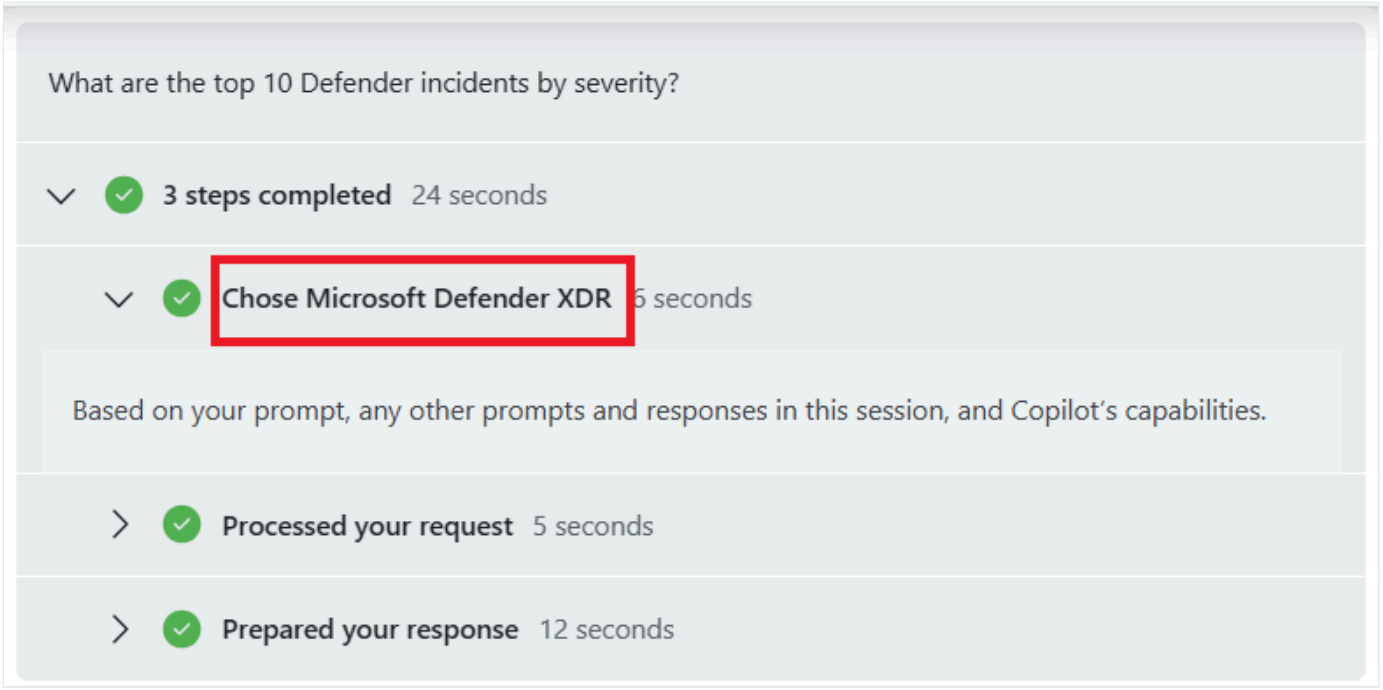

Each plugin and its capabilities directly correspond to prompts found when selecting **See all system capabilities**.

For example, in the process log, the Microsoft Defender XDR plugin is selected. The prompt icon precedes the prompt `Generate guided response`, which indicates that this prompt was a specific capability of the plugin.

When you select the **Prompts** button, use the search to find capabilities. Select **See all system capabilities** to see all the prompts listed for each plugin.

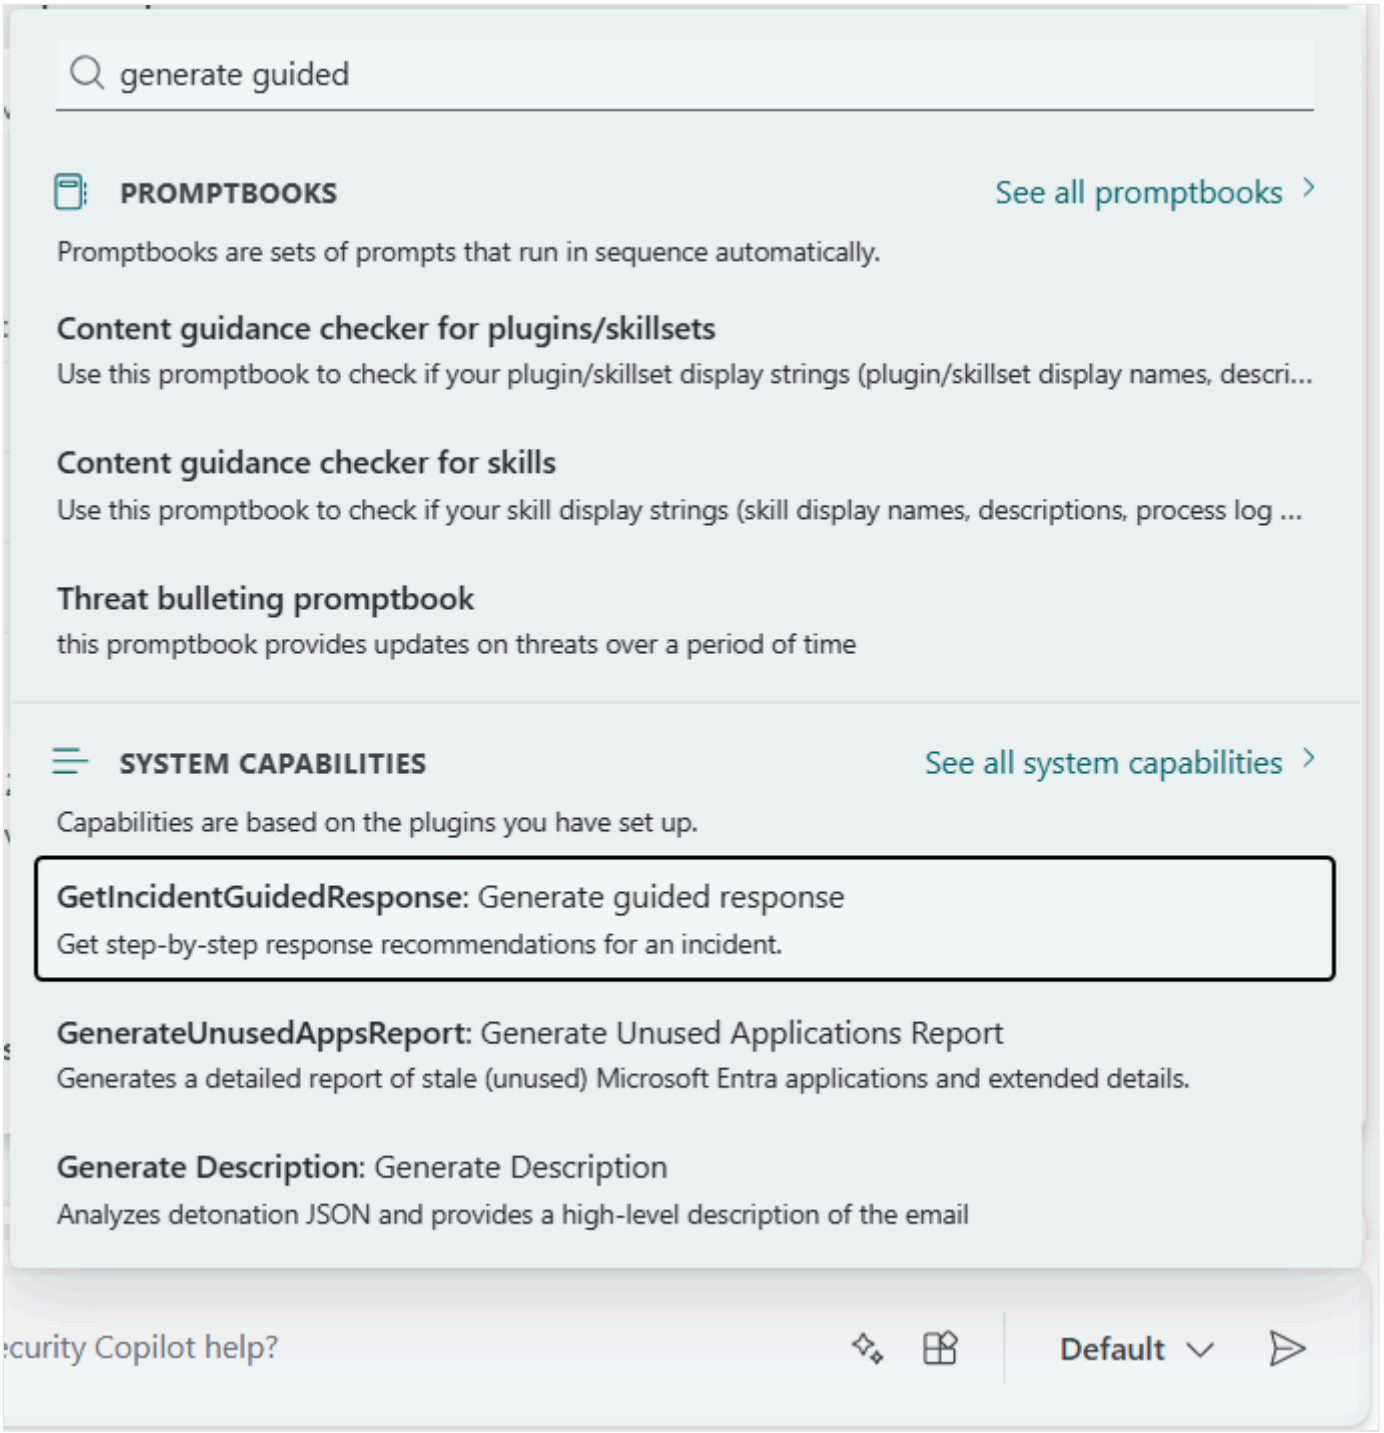

> **Note**
>
> The `Generate guided response` session is just one of the possible sessions created by the Defender XDR embedded experience. When you interact with an embedded Copilot experience to produce results, the session will appear in your session list.

Many responses include direct links to data provided by the Microsoft service plugin used to generate that response. Use the direct links and the process log to build trust in Security Copilot while you verify it's delivering accurate, appropriate responses.

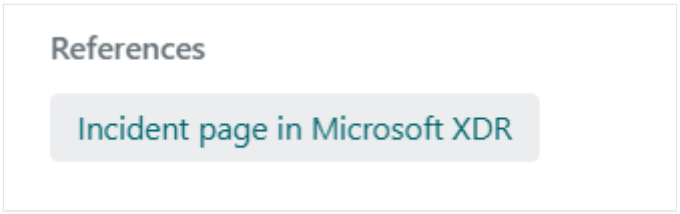

Pin important items to the session

Use the pin board to keep track of important responses in your session.

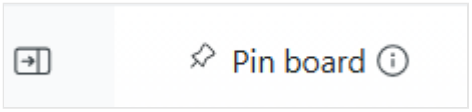

When a response is pinned for the first time, Security Copilot generates a summary of the session for the pin board.

Pin multiple prompt-response pairs to expand what is included in the summary. Pinned items show a condensed response that is expandable. The last prompt pinned is ordered last in the pin board. If you pin a prompt from the middle of the session last, that's where it's pinned on the pin board.

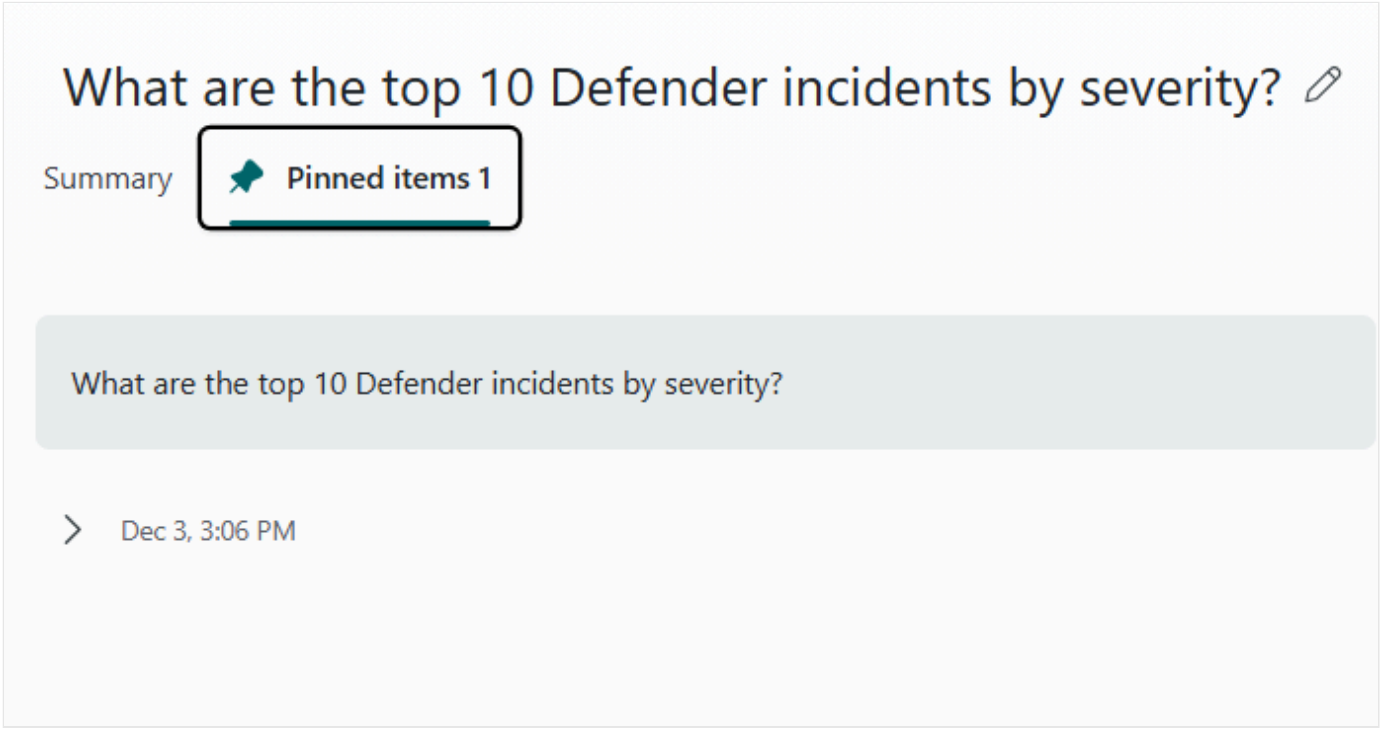

Along with the summary, the pin board features tags under the session name to help provide context to the session.

Share a session

At any time during a Security Copilot session, share your work with your team or decision makers. Select the **Share** icon to create a link.

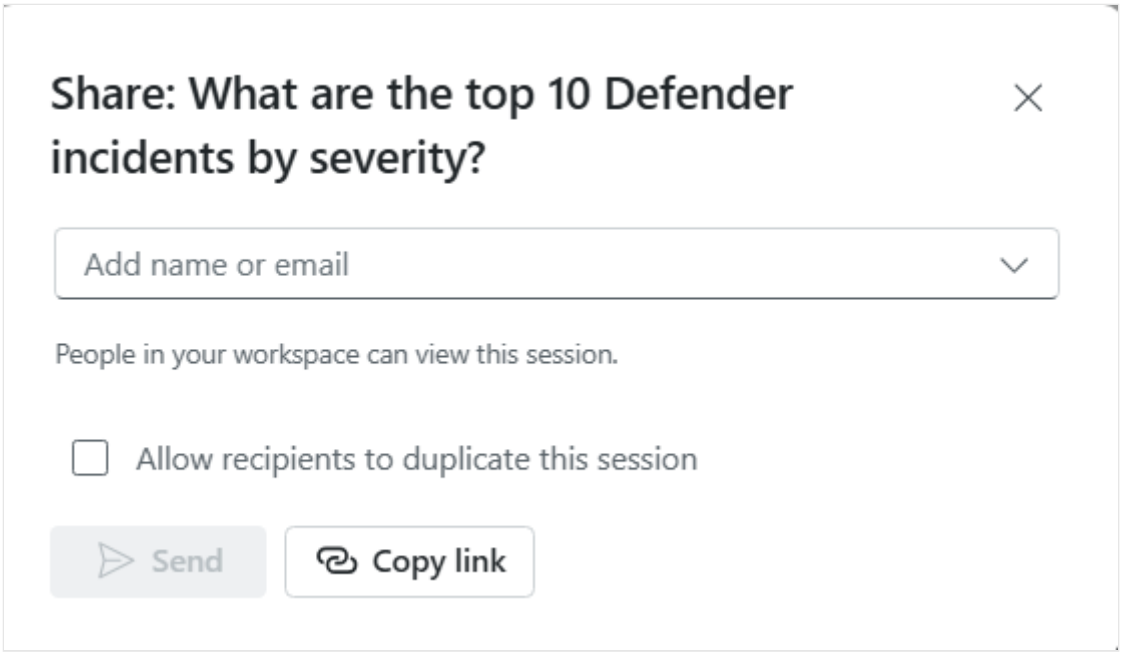

The send via email option isn't available in these regions. Use the **Copy link** and send manually for the following regions:

- `SouthAfricaNorth`
- `UAENorth`

> **Tip**
>
> Keep in mind that the shared session link references the entire session, not just the portion completed before you selected the **Share** feature.

The **Share** feature is also available while reviewing **All history**.

> **⚠ Warning**
>
> When the link is shared by the owner to anyone with Security Copilot access in the same tenant, they are able to view the shared session. Shared sessions show a static display of the entire session, whether they normally have access to the underlying security data or not.

While accessing a shared session, there's an option to **Export** the results.

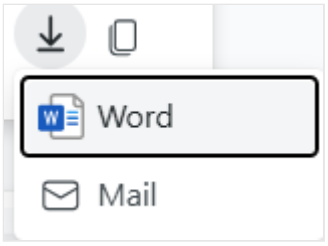

View the shared session summary and pinned items to quickly understand the shared content. For more information, see Shared session access.

Create a promptbook

For more information, see Build your own promptbook

Last updated on 06/04/2026

Understand authentication in Microsoft Security Copilot

Copilot uses on-behalf-of authentication to access security-related data through active Microsoft plugins. Specific Security Copilot roles must be assigned in order for a group or individual to access the Security Copilot platform. After you're authenticated to the platform, your Microsoft Entra and Azure Role Based Access Control (RBAC) determines what plugins are available in prompts. Your Security Copilot role controls what other activities you have access to on the platform, such as configuring settings, assigning permissions, and performing tasks.

Security Copilot RBAC roles are not Microsoft Entra roles. Security Copilot roles are defined and managed within Copilot and only grant access to Security Copilot features.

Microsoft Entra RBAC grants access across the Microsoft portfolio of products including services that contain security data. These roles are managed through the Microsoft Entra admin center. For more information, see Assign Microsoft Entra roles to users.

Azure RBAC controls access to Azure resources like Security Capacity Units (SCU) in a resource group, or Microsoft Sentinel enabled workspaces. For more information, see Assign Azure roles.

Access Security Copilot platform

After Security Copilot is onboarded for your organization, configure Security Copilot RBAC to determine users' access to the Security Copilot platform. Then layer your security coverage with conditional access policies. For more information about securing access to the Security Copilot platform beyond RBAC, see Protect AI with Conditional Access policy.

Security Copilot roles

Security Copilot introduces two roles that function like access groups but aren't Microsoft Entra ID roles. Instead, they only control access to the capabilities of the Security Copilot platform and provide no access by themselves to security data.

- Security Copilot owner
- Security Copilot contributor

> **Note**

To prevent accidental removal of owners, Security Copilot enforces retention of two owners at all times. These two owners cannot be removed.

Microsoft Entra, Microsoft Intune, and Microsoft Purview roles

The following Microsoft Entra, Microsoft Intune, and Microsoft Purview roles automatically inherit **Copilot owner** access, ensuring Security Copilot always has at least two owners.

Microsoft Entra roles:

- Billing Administrator
- Entra Compliance Administrator
- Global Administrator
- Intune Administrator
- Security Administrator

For more information on continuity in Microsoft Entra, see Manage emergency access accounts.

Microsoft Intune role:

- The Microsoft Entra ID Intune Administrator

Microsoft Purview roles:

- Purview Compliance Administrator
- Purview Data Governance Administrator
- Purview Organization Management

As part of the Security Copilot inclusion in the Microsoft 365 E5 and E7 license, the Compliance pack is automatically added.

For capacity provisioning

- Azure Contributor or Owner role on subscription/resource group
- Security Administrator or higher in the tenant you're onboarding Security Copilot.

These roles are needed to attach the provision of SCU capacities to the Security Copilot workspace.

Recommended Security Copilot owner roles

You can update the set of Security Copilot at any time. The list of recommended roles includes Microsoft Entra ID roles and some service specific roles, like Microsoft Purview roles and Microsoft Defender roles.

> **Note**
>
> Only the selected owner roles will inherit the owner role on that workspace. For more information on the supported roles see, Microsoft Entra or Microsoft Purview role.

If a Microsoft plugin requires a different role to access its security data, you need to ensure that the appropriate service role is assigned as well. For example, the bundle works well for an analyst assigned the **Compliance Administrator** role. This role is one of the Microsoft Purview recommended roles and gives them access to Microsoft Purview plugin data. That same analyst needs additional role assignments to access security data like Microsoft Sentinel though. For more examples, see access the capabilities of Microsoft plugins.

Recommended Security Copilot contributor roles

Provision access to Security Copilot with the **Recommended Microsoft Security roles** group, which uses a balanced approach to security and administrative efficiency. Users who already have security permission through Microsoft Entra roles are given contributor access to the Security Copilot platform with this bundle. If the **Everyone** group is assigned, you must remove it before you add the recommended security roles.

Expand table

Role assignment	Advantage	Weakness
Recommended Microsoft Security roles (Default)	Quick, secure way to give users in your organization who already have access to security data the access they need to the platform. More roles, users and groups can still be added.	The group is all or nothing. If there's a role in the recommended group you don't want to have access, the entire group must be removed.
Everyone	This group was initially given contributor access by default to simplify provisioning.	If a user in the everyone group doesn't have access to any security data, the experience in Security Copilot is confusing.
Custom	Complete control over users and groups with access to the platform.	Requires more administrative complexity.

> **Note**
>
> The assignment of the recommended Microsoft security roles to Copilot is the default for new Security Copilot instances. Existing customers who have the everyone group assigned may continue to use the assignment. If the everyone group is removed though, it can't be assigned again. Security Copilot only supports assigning permissions to Role-assignable groups. For more information about these groups, see Creating role-assignable groups in Entra ID.

Access the capabilities of Microsoft plugins

Security Copilot doesn't go beyond the access you have, aligning with Microsoft's Security and privacy RAI principle. Each Microsoft plugin has its own role requirements that remain in effect to access the plugin's service and its data. Verify that you have the proper roles and licenses assigned to use the capabilities of the Microsoft plugins that are activated.

Consider these examples:

1. **Copilot contributor**

   As an analyst, you're assigned the **Copilot contributor** role for access to the Copilot platform and the ability to create sessions. This assignment by itself doesn't give you access to your organization's security data. Following the least privilege model, you don't have any sensitive Microsoft Entra roles like **Security Administrator**. In order to use Copilot to access security data like the Microsoft Sentinel plugin, you still need an appropriate Azure RBAC role like **Microsoft Sentinel Reader**. This role gives you access to incidents in your Microsoft Sentinel workspace from Copilot. For your Copilot session to access the devices, policies, and postures available through the Intune plugin, you need an Intune role like the **Endpoint Security Manager**. The same pattern is true for access to Microsoft Defender XDR data through your Copilot session or the embedded Security Copilot experiences.

   For more information on service specific RBAC, see the following articles:

   - Microsoft Sentinel roles and permissions
   - Microsoft Intune RBAC
   - Microsoft Defender XDR Unified RBAC

2. **Microsoft Entra security group**

Although the **Security Administrator** role inherits access to Copilot and certain plugin capabilities, this role includes PRIVILEGED permissions. Don't assign users this role purely for Copilot access. Instead, create a security group and add that group to the appropriate Copilot role (Owner or Contributor).

For more information, see Best practices for Microsoft Entra roles.

3. **Recommended Microsoft Security roles**

   An analyst with access to Microsoft's unified security operations platform has a custom Defender XDR role assigned that gives them access to multiple workloads in the Defender portal. If the custom role security operations include Security Copilot, then the **Recommende Microsoft Security roles** gives them access to the Security Copilot platform as well as Copilot in Microsoft Defender, simplifying security management.

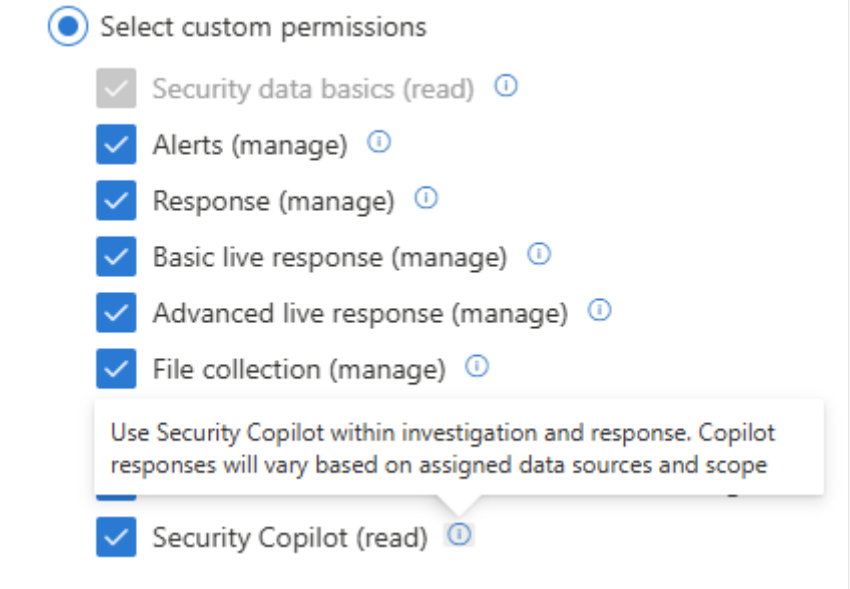

Access embedded experiences

In addition to the Copilot contributor role, verify the requirements for each Security Copilot embedded experience to understand what extra roles and licenses are required.

For more information, see Security Copilot experiences.

Assign roles

The following table illustrates the default access granted to starting roles.

> ⓘ **Note**

Some organizations may still have the **Everyone** group assigned to **Copilot contributor** access. Consider replacing this broad access with **Recommended Microsoft Security roles** group.

Expand table

Capability	Copilot owner	Copilot contributor
Create sessions	Yes	Yes
Manage personal custom plugins	Yes	Default No
Allow contributors to manage personal custom plugins	Yes	No
Allow contributors to publish custom plugins for the tenant	Yes	No
Change availability of preinstalled plugins for the tenant	Yes	No
Upload files	Yes	Default Yes
Manage upload file usage	Yes	No
Run promptbooks	Yes	Yes
Manage personal promptbooks	Yes	Yes
Share promptbooks with tenant	Yes	Yes
Update data sharing and feedback options	Yes	No
Capacity management	Yes*	No
View usage dashboard	Yes	No
Select language	Yes	Yes

Assign Security Copilot access

Assign Copilot roles within Security Copilot settings.

1. Select the home menu.
2. Select **Role assignment** > **Add members**.
3. Start typing the name of the person or group in the **Add members** dialog box.
4. Select the person or group.
5. Select the Security Copilot role to assign (Copilot owner or Copilot contributor).
6. Select **Add**.

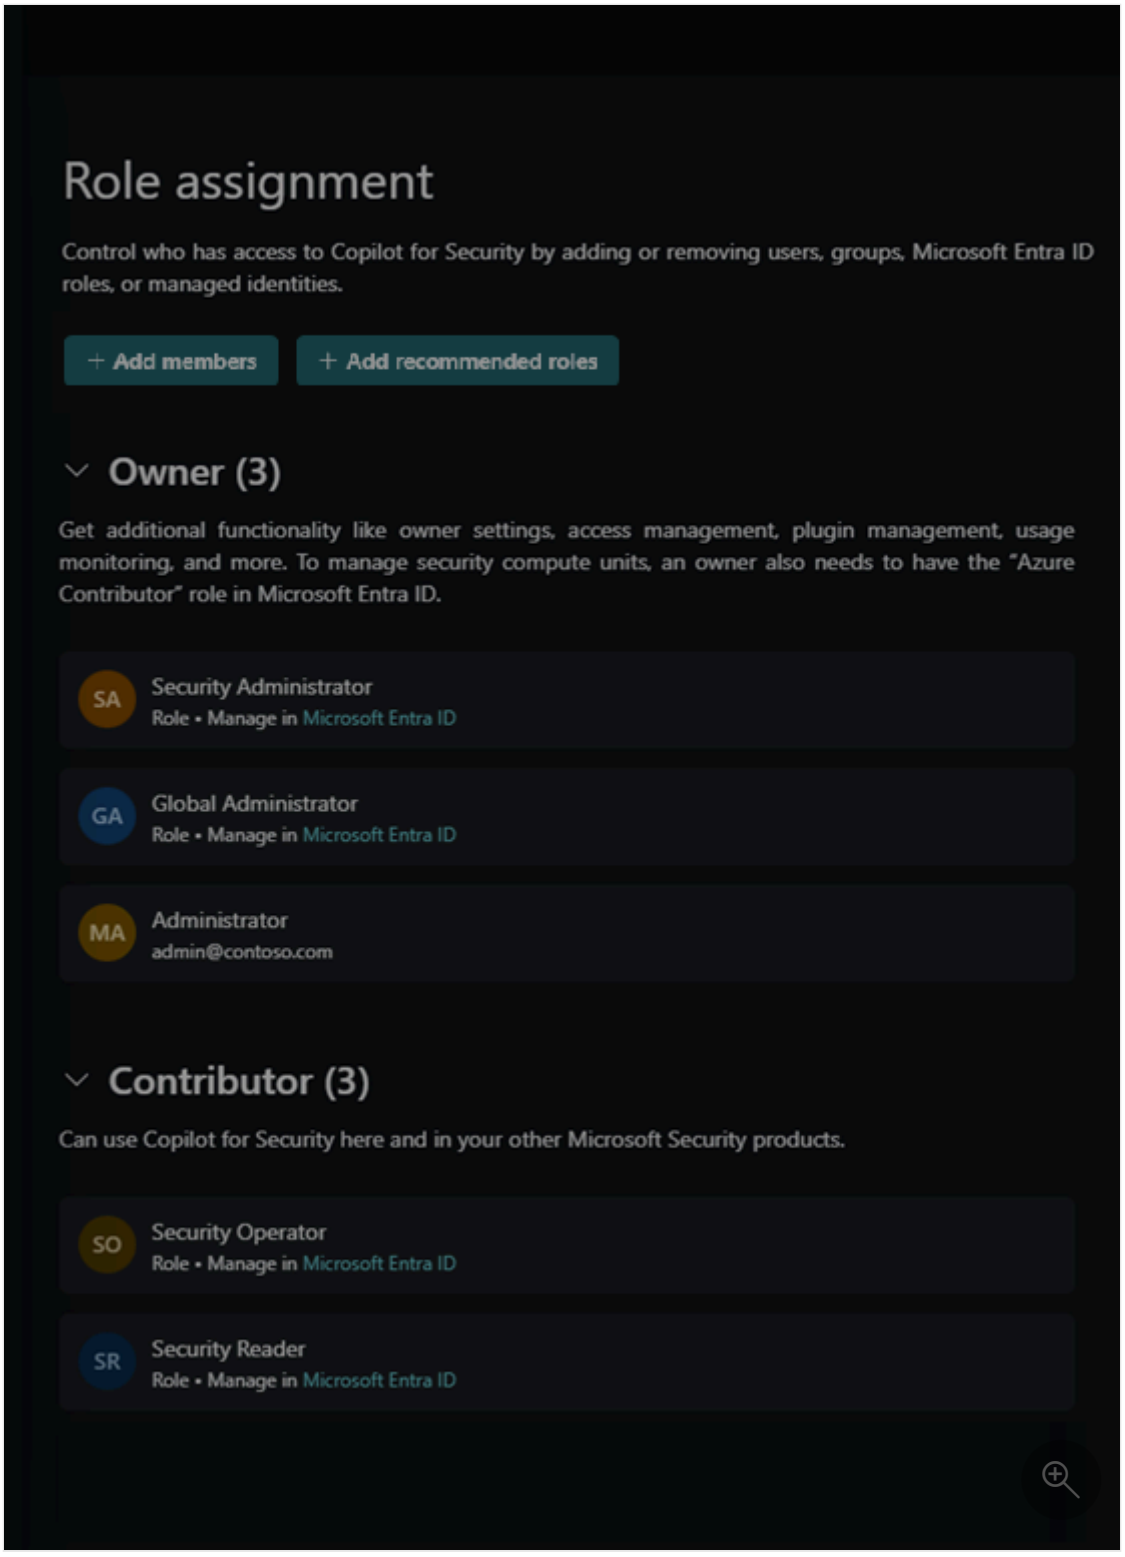

Tip

We recommend using security groups to assign Security Copilot roles instead of individual users. This reduces administrative complexity.

The **Everyone** group is removable from Contributor access. Consider adding the **recommended roles** after you remove the everyone group.

Microsoft Entra role membership is only manageable from the Microsoft Entra admin center. For more information, see Manage Microsoft Entra user roles.

Shared sessions

Copilot contributor role is the only requirement for sharing a session link or viewing it from that tenant.

When you share a session link, consider these access implications:

- Security Copilot needs to access a plugin's service and data to generate a response, but that same access isn't evaluated when viewing the shared session. For example, if you have access to devices and policies in Intune, and the Intune plugin is utilized to generate a response that you share, the recipient of the shared session link doesn't need Intune access to view the full results of the session.
- A shared session contains all the prompts and responses included in the session, whether it was shared after the first prompt or the last.
- Only the user that creates a session controls which Copilot users can access that session. If you receive a link for a shared session from the session creator, you have access. If you forward that link to someone else, it doesn't grant them access.
- Shared sessions are read-only.
- Sessions can only be shared with users in the same tenant that have access to Copilot.
- Some regions don't support session sharing via email.
  - `SouthAfricaNorth`
  - `UAENorth`

For more information on shared sessions, see Navigating Security Copilot.

Multitenant

If your organization has multiple tenants or you leverage a managed service provider partner (MSSP), Security Copilot can accommodate authentication across them to access security data where Security Copilot is provisioned. The tenant that is provisioned for Security Copilot doesn’t need to be the tenant your security analysts logs in from. We currently support the authentication models below through different access approaches.

- Azure B2B & Guest Accounts are able to leverage “Tenant Switching” which is natively built into the Security Copilot standalone portal. For more information, see Navigating Security Copilot tenant switching.

- Granular Delegated Admin Privileges (GDAP) through the Partner Center. For more information, see Grant MSSPs access to Microsoft Security Copilot.
- Azure Lighthouse through the Azure Portal. For more information, see Grant MSSPs access to Microsoft Security Copilot.

Tenant-switching is available through the Security Copilot standalone portal but only supports B2B collaboration accounts or guest account. For more information, see Navigating Security Copilot tenant switching.

Cross tenant sign-in example

Contoso recently merged with Fabrikam. Both tenants have security analysts, but only Contoso purchased and provisioned Security Copilot. Angus MacGregor, an analyst from Fabrikam wants to use their Fabrikam credential to use Security Copilot. Here are the steps to accomplish this access:

1. Ensure Angus MacGregor's Fabrikam account has an external member account in the Contoso tenant.
2. Assign the external member account the necessary roles to access Security Copilot and the desired Microsoft plugins.
3. Sign in to the Security Copilot portal with the Fabrikam account.
4. Switch tenants to Contoso.

For more information, see Grant MSSP access.

Manage promptbooks

Promptbook creation is available to all roles, including the ability to publish a custom promptbook for the tenant. Choose whether to publish a promptbook for yourself or the tenant at the time of creation.

For more information, see Build your own promptbook.

Preinstalled plugin authentication

Preinstalled plugins, like **Microsoft Sentinel** and **Azure AI Search**, require more setup. The plugin provider determines the type of authentication. Any plugin with the gear or **Set up** button is configured per user. Whether a preinstalled plugin is restricted, all users who have access to the plugin configure that plugin's setup for themselves.

> **Note**
>
> Website plugins use anonymous authentication to access content.

For more information, see Manage preinstalled plugins.

Last updated on 06/04/2026

Prompting in Microsoft Security Copilot

Once you're all set up in Security Copilot, you can use prompts to interact with the system.

Depending on your experience, you might start with either **agents** or **chat**:

- In newer experiences, agents are the primary entry point.
- To begin prompting, open **All history** and select **New session** to start a chat session.

Prompts help you ask questions, investigate issues, and explore data using natural language. Promptbooks are a series of prompts that are grouped together to accomplish specific security-related tasks.

Start a prompt session

To begin prompting:

1. Select **History**.
2. Review your existing sessions (if any).
3. Select **New session**.
4. Enter your prompt.

This opens a chat session where you can interact with Security Copilot.

Use the prompts and promptbooks library

The prompts and promptbooks library in Security Copilot helps you quickly harness the power of the platform in a way that aligns with your role. This set of prompts and promptbooks are designed to guide you through features and capabilities most relevant to your work.

The curated prompts and promptbooks provide easy access to role-based examples to help you get started quickly and streamline your workflow.

> ⓘ **Note**
>
> Prompts and promptbooks are available within a chat session and through the promptbooks library.

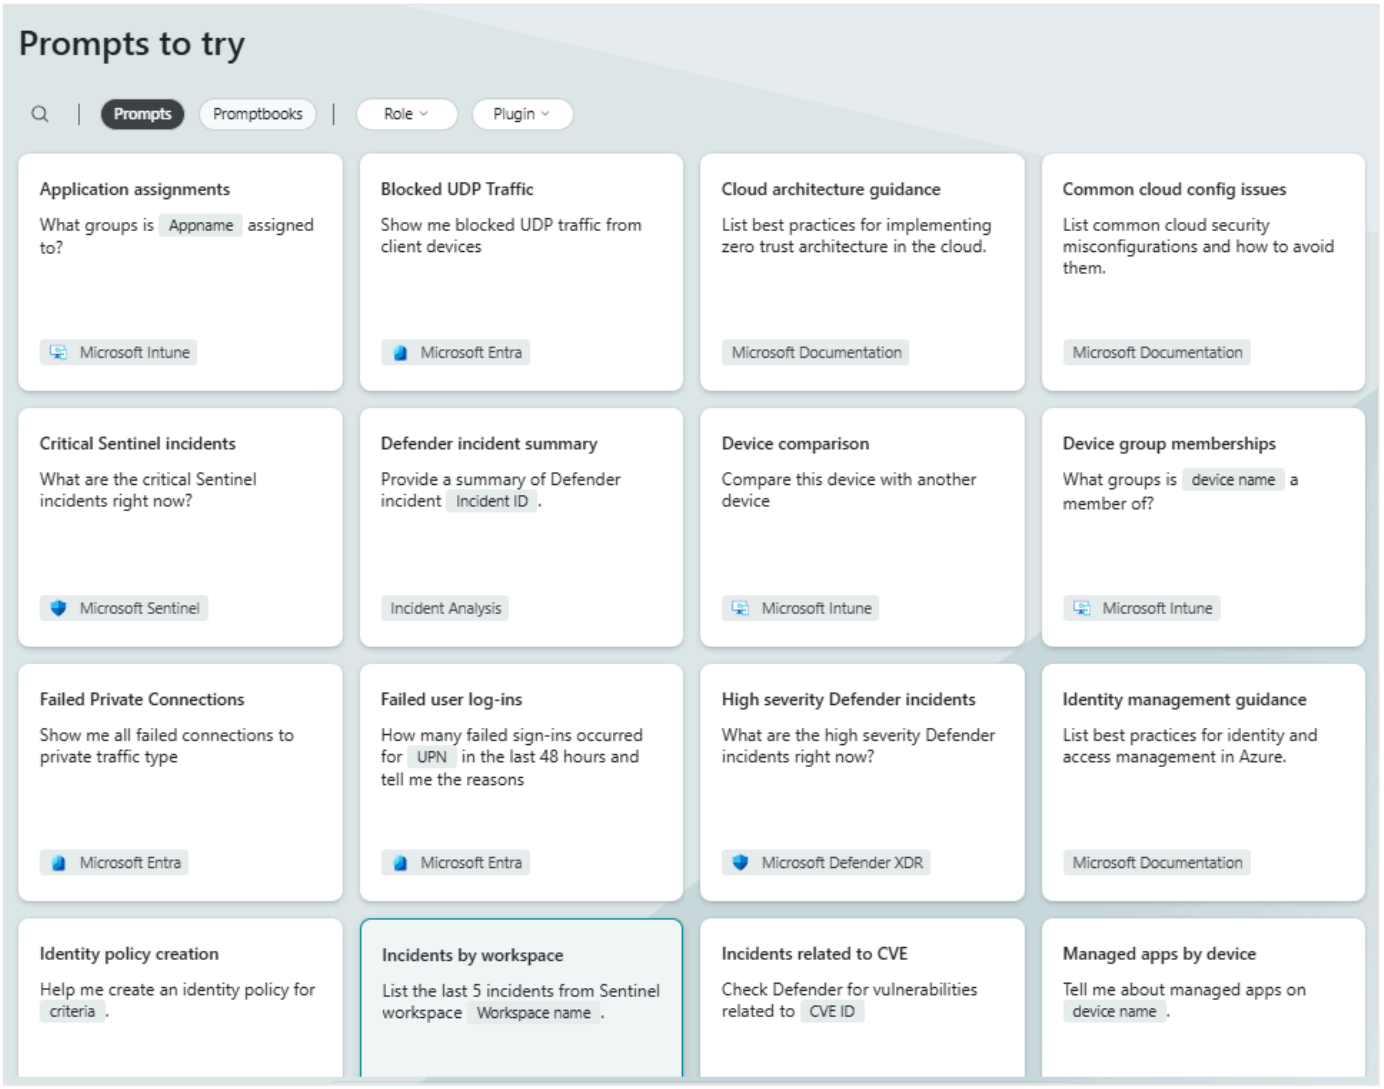

Apply filters

You can apply filters to find prompts or promptbooks that are most relevant to you. Narrow down search results by applying multiple filters to better suit your role or scenario.

Customize your search by selecting filters based on:

- **Roles** - Examples include: CISO, SOC analyst, threat intel analyst, and IT administrator.
- **Plugins** - Examples include: Microsoft Defender XDR, Microsoft Threat Intelligence, and Natural language to KQL for advanced hunting.

You can also use the search function to look up prompts or promptbooks by title.

Using the prompts library

Follow these steps to use ready-made prompts:

1. Select **Prompts** in an active session and apply filters as needed.
2. Use the search function to find the title of a prompt or select a prompt from the library.

3. If the prompt requires input (for example, an app name or incident ID), provide the necessary details.
4. Select **Send** or press Enter.

Using the promptbooks library

There are two ways you can access the promptbooks library:

- Through an active session
- Through the Security Copilot menu

Access promptbooks from a session

There are several actions available when working with promptbooks:

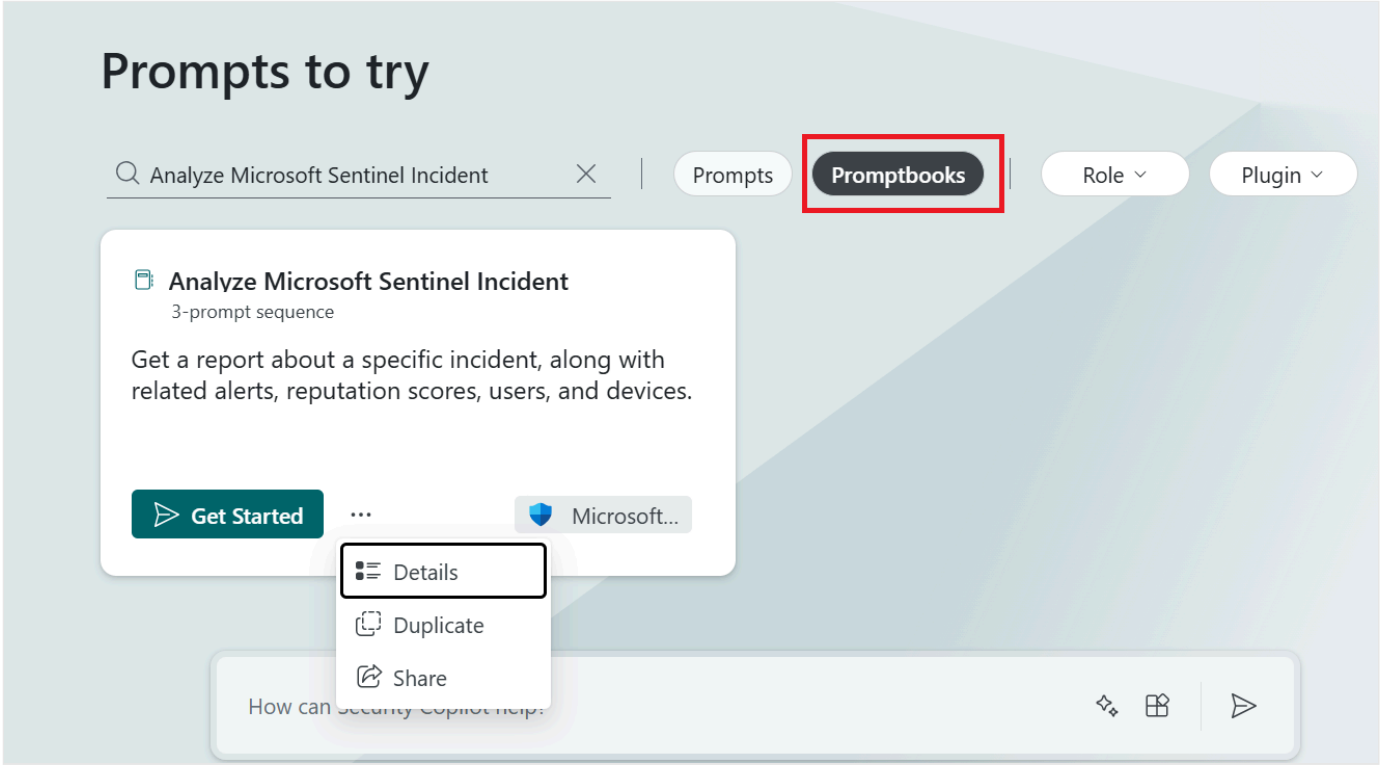

You can choose any of the following actions:

- Filter by **Promptbooks**
- Select **Get started** to run a promptbook
- View details of a promptbook
- Duplicate a promptbook
- Share a promptbook

Access promptbooks from the menu

Both prebuilt and user-built promptbooks across your organization appear in the promptbook library. View promptbooks by going to the Security Copilot menu and selecting **Promptbooks**.

To view prebuilt promptbooks, see Try promptbooks.

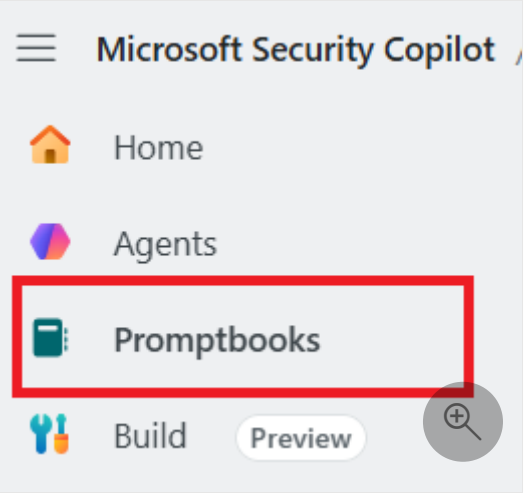

The promptbook library displays all available promptbooks. They are listed by name, owner, description, number of prompts, required plugins, and visibility.

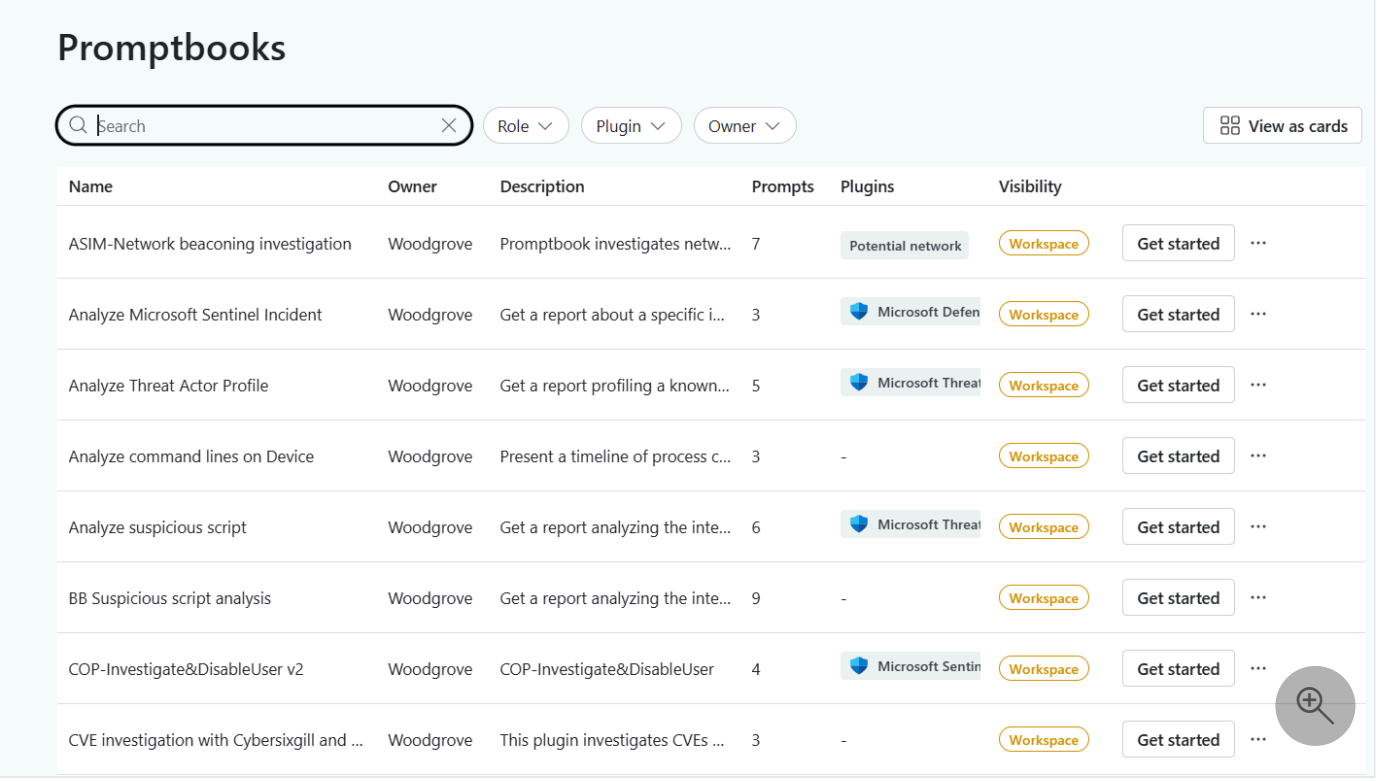

Use the search bar in the **Promptbook library** to locate promptbooks by name or apply filters to refine your results.

Choose individual prompts and promptbooks

In an active session, select the prompts icon to view available prompts and promptbooks.

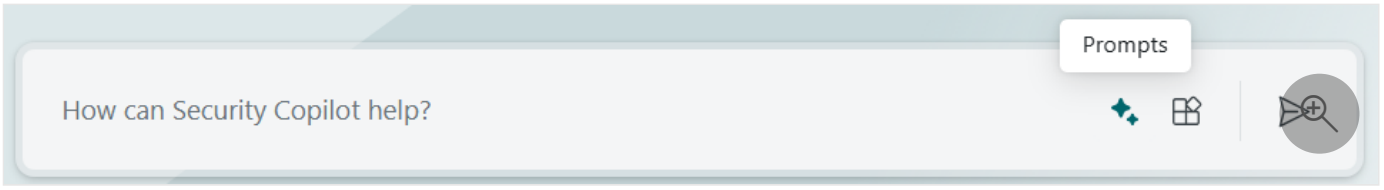

Promptbooks are collections of prompts that run in sequence to complete a task, with each prompt building on the previous one.

For more information, see Using promptbooks.

When to use prompts and agents

You can use both prompts and agents depending on your task:

- Use **agents** for structured workflows and repeatable tasks
- Use **prompts** for ad-hoc questions and investigations

You can move between agents and prompts as needed to complete your work.

Create your own prompts

Using the prompt bar, you can enter requests in your own words for Security Copilot to answer.

1. Enter your prompt in an active session.
2. Select **Send** or press Enter.
3. Wait for Security Copilot to generate a response.

   As the response is being formed, Security Copilot shows steps in the process log. The process log provides visibility into the actions taken, sources used, and processing time.

   You can cancel, edit, or delete your prompt during response generation.
4. Review the response and verify that the results meet your needs.

You can continue the session with additional prompts to refine or expand your results. If needed, rephrase your prompt, provide more details, or specify the desired format.

For guidance on writing effective prompts, see Elements of an effective prompt.

Watch the following video to learn more about prompting:

https://learn-video.azurefd.net/vod/player?id=d98ee5da-a374-43ac-b0d1-977de22cd32e&locale=en-us&embedUrl=%2Fcopilot%2Fsecurity%2Fprompting-security-copilot

Providing feedback

Whether you use promptbooks or your own prompt, providing feedback helps improve Security Copilot. Feedback options appear below each response.

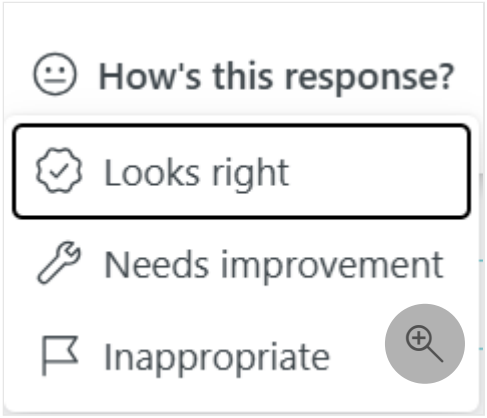

Remember, Security Copilot doesn't always get everything right, and AI-generated content can contain inaccuracies.

See also

- Prompting best practices
- Using promptbooks

Last updated on 06/01/2026

Microsoft Security Copilot in your workflows

This article provides an overview of the key personas and how you can discover, set up and use agents, and the related components in the workflow in Microsoft Security Copilot. Advanced users and developers can also build and publish custom agents tailored to their organization's needs.

Start with agents

In newer onboarding experiences, agents are the primary way to begin working in Security Copilot.

Agents help you:

- Automate investigation workflows
- Perform repeatable tasks
- Guide analysis using structured steps

You can still use prompts and promptbooks alongside agents as needed.

Personas

Security Copilot agents span across three key personas:

- **Administrators** - Discover and determine which agents to install, configure plugins, and review usage and success metrics of the agents.
- **End users** (analysts or data security team; or IT admins) - Interact with agents or prompts by using workflow outputs and providing feedback.

  For more information on use cases across personas, see Use cases.
- **Developers** - Build and publish agents for internal use or broader organizational deployment. For more information on the developer documentation, see Build custom agents.
-

Work with agents, plugins, and connectors

Microsoft Security Copilot uses these capabilities to automate and extend security operations. Understanding how these components work together helps you choose the right approach for each scenario.

Expand table

Scenario	Recommended approach	Guidance
Getting started / investigation	Use **agents** or start a chat session with **prompts**	- Agents: Begin with prebuilt or recommended agents - Prompts: Start a session from **All history** → **New session**
Ad-hoc analysis	Use **prompts** and **promptbooks**	Prompting in Security Copilot
Repeatable workflow & integration	Use **agents**, **plugins**, and **connectors**	Plugins and Connectors
Automation	Use **agents**	- Review key agent terminology - Explore Discover agents - Review Use cases
Advanced workflows	Build **custom agents**, **promptbooks**, and **custom plugins**	- Agents: Build custom agents - Promptbooks: Build your own promptbooks - Plugins: Build custom plugins

Combine agents and prompts

You can use agents together with prompts as part of your workflow:

- Start with an agent for structured workflows
- Use prompts to explore results or investigate further
- Return to agents for repeatable execution

This flexibility allows you to adapt Security Copilot to your workflow and operational needs.

Agents

An agent is an AI-driven security assistant or workflow that can autonomously execute and orchestrate tasks on behalf of security teams. Each agent has a defined goal such as triaging

phishing alerts, generating a threat intelligence briefing, or remediating vulnerabilities.

You can discover Microsoft-built agents, deploy partner agents, or build your own custom agents.

Agents can be:

- **Interactive** – respond to user input in real time
- **Automated** – triggered by events or schedules

Plugins

When an agent needs information or needs to take action, it uses a plugin. A plugin connects Security Copilot to a security product or service and provides access to data and actions.

For example:

- Retrieve alerts from Microsoft Sentinel
- Query threat intelligence data

A Security Owner can enable or disable plugins, and developers can build custom plugins.

For more information, see Plugins overview.

Connectors

Connectors integrate Security Copilot with external systems and workflows.

Connectors can:

- Trigger agents
- Run prompts
- Start automation workflows

For more information, see Connectors overview.

Custom agents

If you are a developer, you can build and publish Security Copilot custom agents tailored to your organization's needs.

Security Copilot supports building, testing, and publishing agents for internal or organizational use.

For more information, see the Agent development overview developer content documentation.

Promptbooks

A promptbook is a reusable, multi-step workflow made up of natural language prompts. It guides Security Copilot through structured investigations such as analyzing incidents or generating reports.

Promptbooks can:

- Include multiple prompts
- Reference plugins or skills
- Be reused across workflows

Analysts can run promptbooks manually or as part of agent-driven workflows.

For more information, see Promptbooks.

Next steps

- Discover agents

Last updated on 05/01/2026

Microsoft Security Copilot agents overview

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

Microsoft Security Copilot agents automate repetitive tasks and reduce manual workloads. They enhance security and IT operations across cloud, data security and privacy, identity, and network security. These agents handle high-volume, time-consuming tasks by pairing data and code with an AI language model. They respond to user requests and system events, helping teams work more efficiently and focus on higher-impact tasks.

Agents fit naturally into existing workflows. You don't need special training or other licensing to use them. Agents utilize SCUs to operate just like other features in the product. They integrate seamlessly with Microsoft Security solutions and the broader supported partner ecosystem. Agents learn based on feedback and keep you in control on the actions it takes. They handle resource-intensive tasks like threat intelligence briefings, and Conditional Access optimization. With Microsoft Security Copilot agents, you can scale up your teams, people, and processes.

Microsoft Security Copilot agents offer significant benefits for security teams and IT operations by automating routine tasks and freeing up valuable time for teams to concentrate on strategic initiatives and complex problem-solving. This leads to improved operational efficiency, enhanced security and giving teams the ability to respond more swiftly to emerging threats. With Security Copilot agents, organizations can achieve greater scalability and resilience in their security and IT processes.

Agent terminology

This article defines key agent terminology used during agent set up.

Trigger

A trigger is an event or condition that tells an agentic system to initiate an action or series of actions. You can set the agent to run automatically at specific intervals or choose to run it manually when needed.

Permissions

In the context of agents, permissions are the level of authorization an AI agent is given by an administrator during configuration that enables it to access specific information or carry out its tasks.

These permissions might include the ability to read data from other solutions, for example Microsoft Defender External Attack Surface Management or Microsoft Threat Intelligence. By giving the agent permissions to access data from these solutions, the agent can generate insights or provide recommendations based on the information it's able to gather.

Identity

An agent needs an identity to authenticate and securely access resources when it runs. During the agent setup process, you're given two types of identity to choose from:

- **Create an agent identity**

  > **Note**
  >
  > Currently, this option is only available for Microsoft-built agents.

  This option creates a dedicated identity for the agent using the Microsoft Entra Agent ID capability. Microsoft Entra Agent IDs are identities created specifically for AI agents. The user setting up the agent grants the agent ID permissions needed for the agent to run successfully. Using Agent IDs keeps access scoped, secure, and easier to manage. For more information, see What are agent identities?.

- **Connect with an existing user account**

  This option lets the agent use your credentials to run. It inherits your access and permissions while it's active.

Plugins

Plugins extend the capabilities of Security Copilot. A plugin is a component that extends what an agent can do by giving it access to capabilities in Microsoft and non-Microsoft services and public websites through APIs. Having access to plugins adds more context to the output of an agent.

Get started

To get started with Microsoft Security Copilot and harness agentic AI to transform security and IT operations, follow these steps in the journey to discover and use agents.

1. **Discover agents** - Browse Microsoft and partner-built agents
    - Standalone or embedded experience in the Microsoft security products
    - Security Store
2. **Configure and manage agents** - Setup agents
3. **Create your own agents** tailored for your use case - Build custom agents
4. **Extend agents with plugins** - Plugins
5. **Connect external systems** - Connectors

Last updated on 03/18/2026

Discover agents in Microsoft Security Copilot

You can easily discover Microsoft Security Copilot agents from both the **standalone** and **embedded** experiences. Finding the right agent is the first step to automating your security workflows.

On the **standalone** experience:

1. Toggle menu to navigate to the agent library.
2. Select **Agents**.
3. The agent library opens.
4. Browse through the library to peruse the various agents.

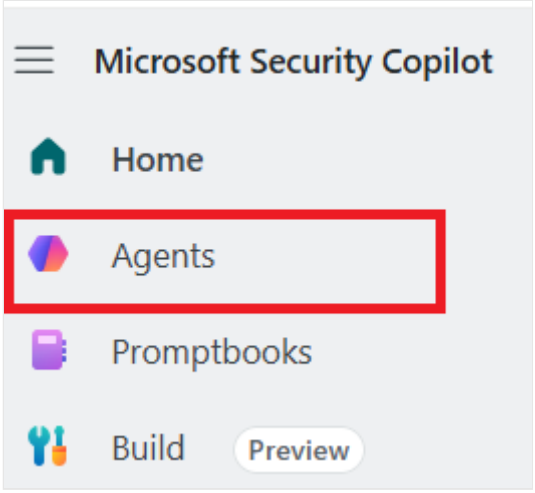

There are Microsoft agents and partner agents available for you to choose from. Depending on your role, you can either set them up or access the agent to run it. This intuitive access ensures that you can quickly integrate agents into your workflows.

In the **embedded** experiences, you'll see agents within the portal and can explore their capabilities.

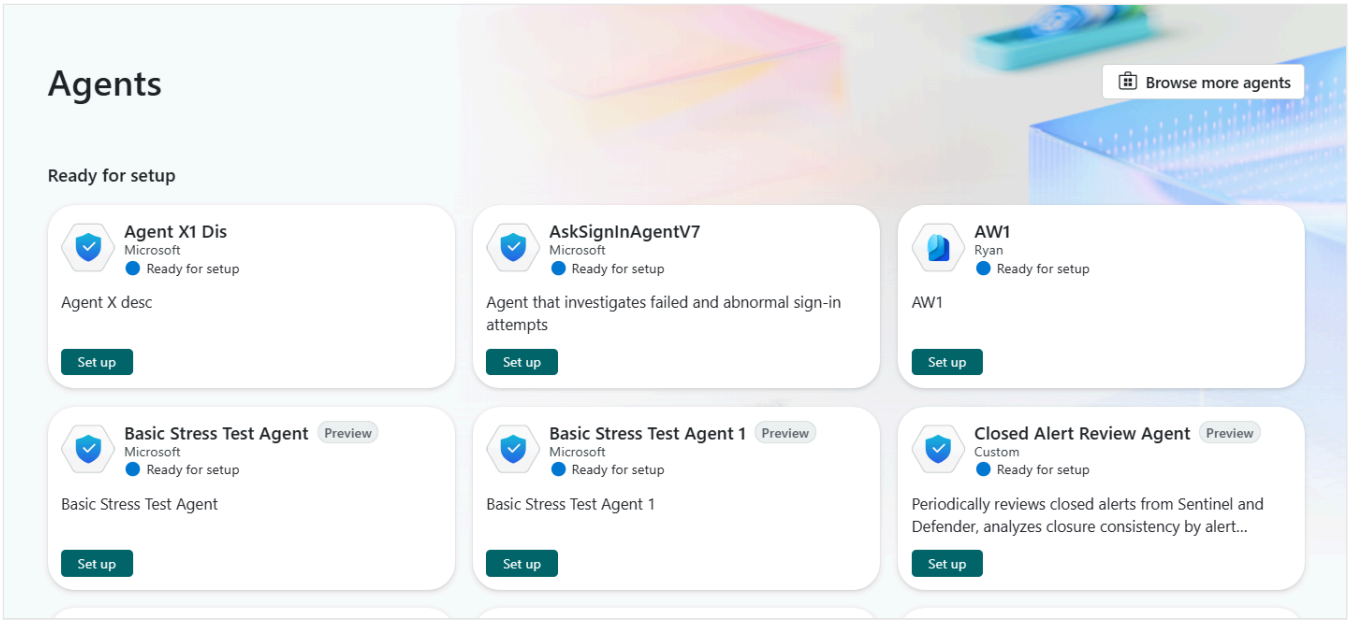

Agents in the Microsoft security ecosystem

Agents are available across the integrated Microsoft security products. For more information on the full list of agents, see:

- Identity and access in Microsoft Entra
- Security operations in Microsoft Defender
- Endpoint management in Microsoft Intune
- Data security in Microsoft Purview
- Analyze incidents and generate hunting queries in Microsoft Sentinel

Discover agents from Security Store

Browse and add agents directly from the integrated Security Store within Security Copilot. The Security Store provides a growing ecosystem of Microsoft and partner-built agents, making it easy to find solutions that fit your security needs. For more information, see What is Security Store.

Next steps

Set up agents

See also

- Partner-built agents
- Build custom agents

Last updated on 03/18/2026

Microsoft Security Copilot agents

This page provides information about agents available in the standalone and embedded experiences of Security Copilot.

Agents in the standalone experience

Security Copilot, accessed through https://securitycopilot.microsoft.com, is considered as the standalone experience.

Threat intelligence briefing agent

Generating a threat intelligence report can be a cumbersome and resource intensive task. It requires intelligence gathering and can take hours or days to complete.

The Threat Intelligence Briefing Agent generates timely, relevant threat intelligence reports with detailed technical analysis based on the latest threat actor activity and both internal and external vulnerability information. The agent correlates Microsoft threat data from Defender External Attack Surface Management (EASM) and real-time customer signals to add critical context to threat information in a matter of minutes, saving analyst teams hours or even days spent on intelligence gathering and correlation.

Expand table

Attribute	Description
Identity	Requires connection to an existing user account and creation of a new agent identity
License	Defender EASM Standard
Permissions	Read data from Agents, Microsoft Defender External Attack Surface Management, Microsoft Threat Intelligence, and Microsoft Threat Intelligence Agents
Products	Security Copilot
Plugins	Microsoft Defender External Attack Surface Management Microsoft Threat Intelligence Microsoft Threat Intelligence Agent
Role-based access	Security Copilot Owner and Security Copilot Contributor roles can see the report generated by the Threat intelligence briefing agent within the Microsoft Security Copilot Agents page
Trigger	Runs every 7 days or can be triggered manually

For more information, see Threat intelligence briefing agent

Agents in the embedded experience

Accessing Security Copilot embedded experiences in other Microsoft security products is considered as the embedded experience. Agents are available across the integrated Microsoft security products. For more information on the full list of agents, see:

- Identity and access in Microsoft Entra
- Security operations in Microsoft Defender
- Endpoint management in Microsoft Intune
- Data security in Microsoft Purview
- Analyze incidents and generate hunting queries in Microsoft Sentinel

---

Last updated on 03/18/2026

Threat Intelligence Briefing Agent (standalone experience)

> **Note**
>
> This article discusses the Threat Intelligence Briefing Agent standalone experience in Security Copilot. To learn more about the embedded experience in the Microsoft Defender portal, read Microsoft Security Copilot Threat Intelligence Briefing Agent in Microsoft Defender.

Threat intelligence analysts face several challenges in delivering insightful, actionable, contextualized intelligence. The task of developing threat intelligence briefings involves collecting information from various threat feeds, tools, and portals; filtering and correlating this information; and analyzing and mapping organizational risks. These activities happen before analysts can even start developing the report itself and generating insights for when they deliver the briefing. By then, as these processes can take anywhere from hours to days, the threats facing the organization have already evolved, which can render the briefing obsolete.

The Threat Intelligence Briefing Agent was developed in response to these pain points. The Threat Intelligence Briefing Agent in the Microsoft Security Copilot standalone portal generates threat intelligence briefings based on the latest threat actor activity and both internal and external vulnerability information—in a matter of minutes. The agent can help security teams save time by creating a customized, relevant report that provides CISOs, security managers, and analysts with key situational awareness and a solid foundation for defense work.

The agent leverages dynamic automation and deep generative AI along with its wealth of threat intelligence knowledge and signals. When building the briefing, the agent dynamically chooses the next step based on the outcome of the previous step, allowing it to decide in real-time what threat intelligence to include and prioritize. The agent then translates this technical information into a digestible report that can be consumed by various audiences.

The Threat Intelligence Briefing Agent is best suited for customers who have turned on Microsoft Defender External Attack Surface and Microsoft Defender for Endpoint, as the agent relies on signals and insights from these first-party integrations to deliver accurate and context-rich reports.

Prerequisites

Products

Microsoft Security Copilot is needed to run this agent.

Plugins

The following plugin is required to run this agent:

- Microsoft Threat Intelligence

The following plugin is optional but can add more context to the output:

- Microsoft Defender External Attack Surface Management

User account permissions configuration

> **Important**
>
> **Identity and permissions requirement**: This agent requires connection to an existing user account or creation of a new agent identity (recommended). The agent can read data from Defender External Attack Surface Management and Defender Vulnerability Management. You must configure the user account with the appropriate permissions outlined in the following section before setting up the agent.

Permission overview

The user account connected to the agent must have these permissions:

Required permissions:

- **Microsoft Defender for Endpoint**: Access to Defender Vulnerability Management data
- **Security Copilot Contributor**: Access to Security Copilot platform and agent management

Optional permissions:

- **Exposure Management (read)**: Access to Microsoft Security Exposure Management insights, including External Attack Surface Management data

Role-based access:

- Owners and contributors can see the report generated by the Threat Intelligence Briefing Agent within the Microsoft Security Copilot agent library page

Configure permissions

Step 1: Create custom role in Microsoft Defender XDR

> **Important**
>
> Microsoft recommends that you use roles with the fewest permissions. Using lower permissioned accounts helps improve security for your organization.

1. Sign in to the Microsoft Defender portal as Security Administrator or higher.
2. Navigate to **Permissions** > **Microsoft Defender XDR** > **Roles**.
3. Select **Create custom role**.
4. On the **Basics** tab:
   - **Role name**: `Threat Intel Agent - Read Only`
   - **Description**: `Read-only access for Threat Intelligence Briefing Agent`
   - Select **Next**
5. On the **Choose permissions** page:
   - Select **Security posture**
   - Select **custom permissions**
   - Under **Posture management**, select **Vulnerability management - Read**
   - Select **Apply** > **Next**
6. On the **Assign users and data sources** page:
   - Select **Add assignment**
   - **Assignment name**: `Threat Intel Agent Assignment`
   - **Employees**: Select the user account for the agent
   - **Data sources**: Select **Microsoft Defender for Endpoint**
   - Select **Next** > **Submit**

Step 2: Assign Security Copilot Contributor role

1. Sign in to Microsoft Security Copilot.

2. Select the home menu > **Role assignment** > **Add members**.
3. Search for and select the user account, then assign **Security Copilot Contributor** role.
4. Select **Add**.

Step 3 (Optional): Add External Attack Surface Management permissions

If your organization uses Microsoft Defender External Attack Surface Management:

1. In the Microsoft Defender portal, go to **Permissions** > **Microsoft Defender XDR** > **Roles**.
2. Find the `Threat Intel Agent - Read Only` role and select **Edit**.
3. Navigate to **Choose permissions** > **Security posture** > **Select custom permissions**.
4. Under **Posture management**, add **Exposure Management - Read**.
5. In **Data sources**, add **Microsoft Security Exposure Management**.
6. Save the changes.

> **Important**
>
> After setting up permissions, activate the Microsoft Defender XDR Unified role-based access control (RBAC) model for the role to take effect.

> **Tip**
>
> Consider using a dedicated service account for running agents to maintain separation of duties and enhance security monitoring.

Trigger

This agent runs at the set time interval when turned on, or manually when you want to run it.

Set up the agent

1. To run the Threat Intelligence Briefing Agent, go to the **Agents** page in the Microsoft Security Copilot standalone portal.

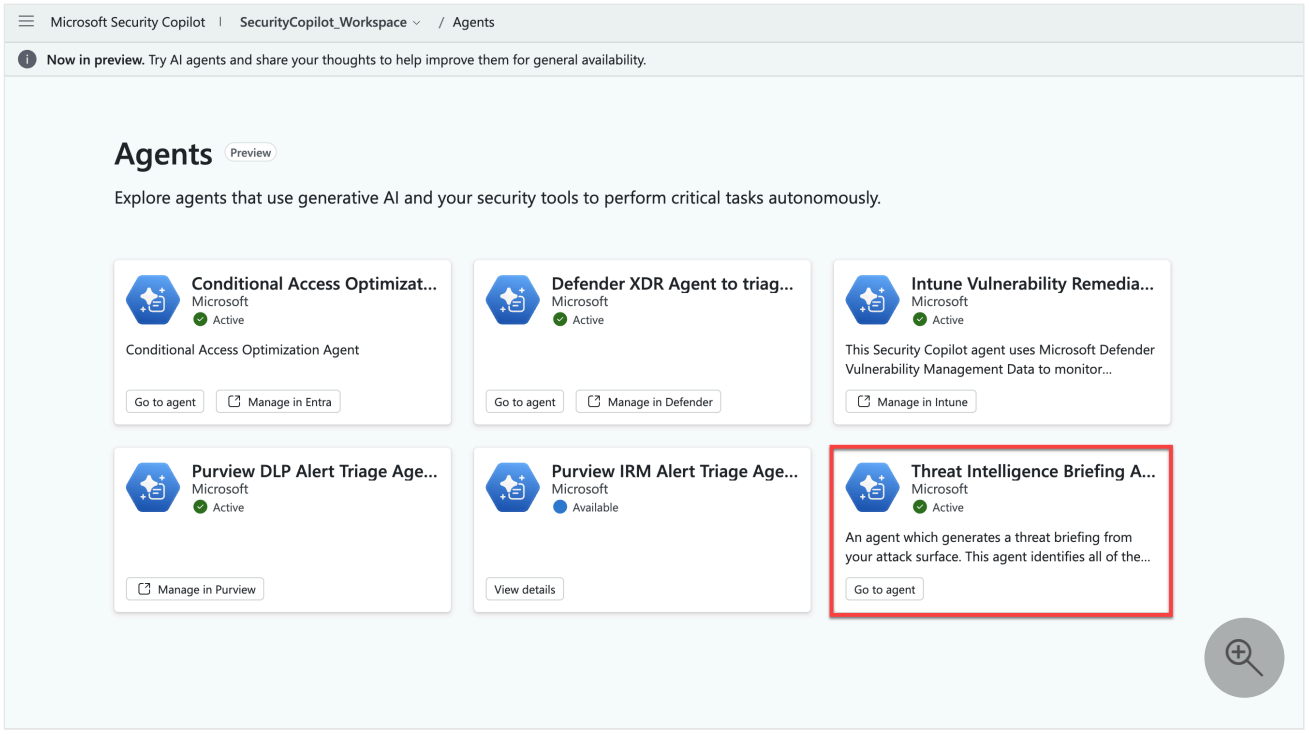

2. Choose the **Threat Intelligence Briefing Agent** and select **Set up**.

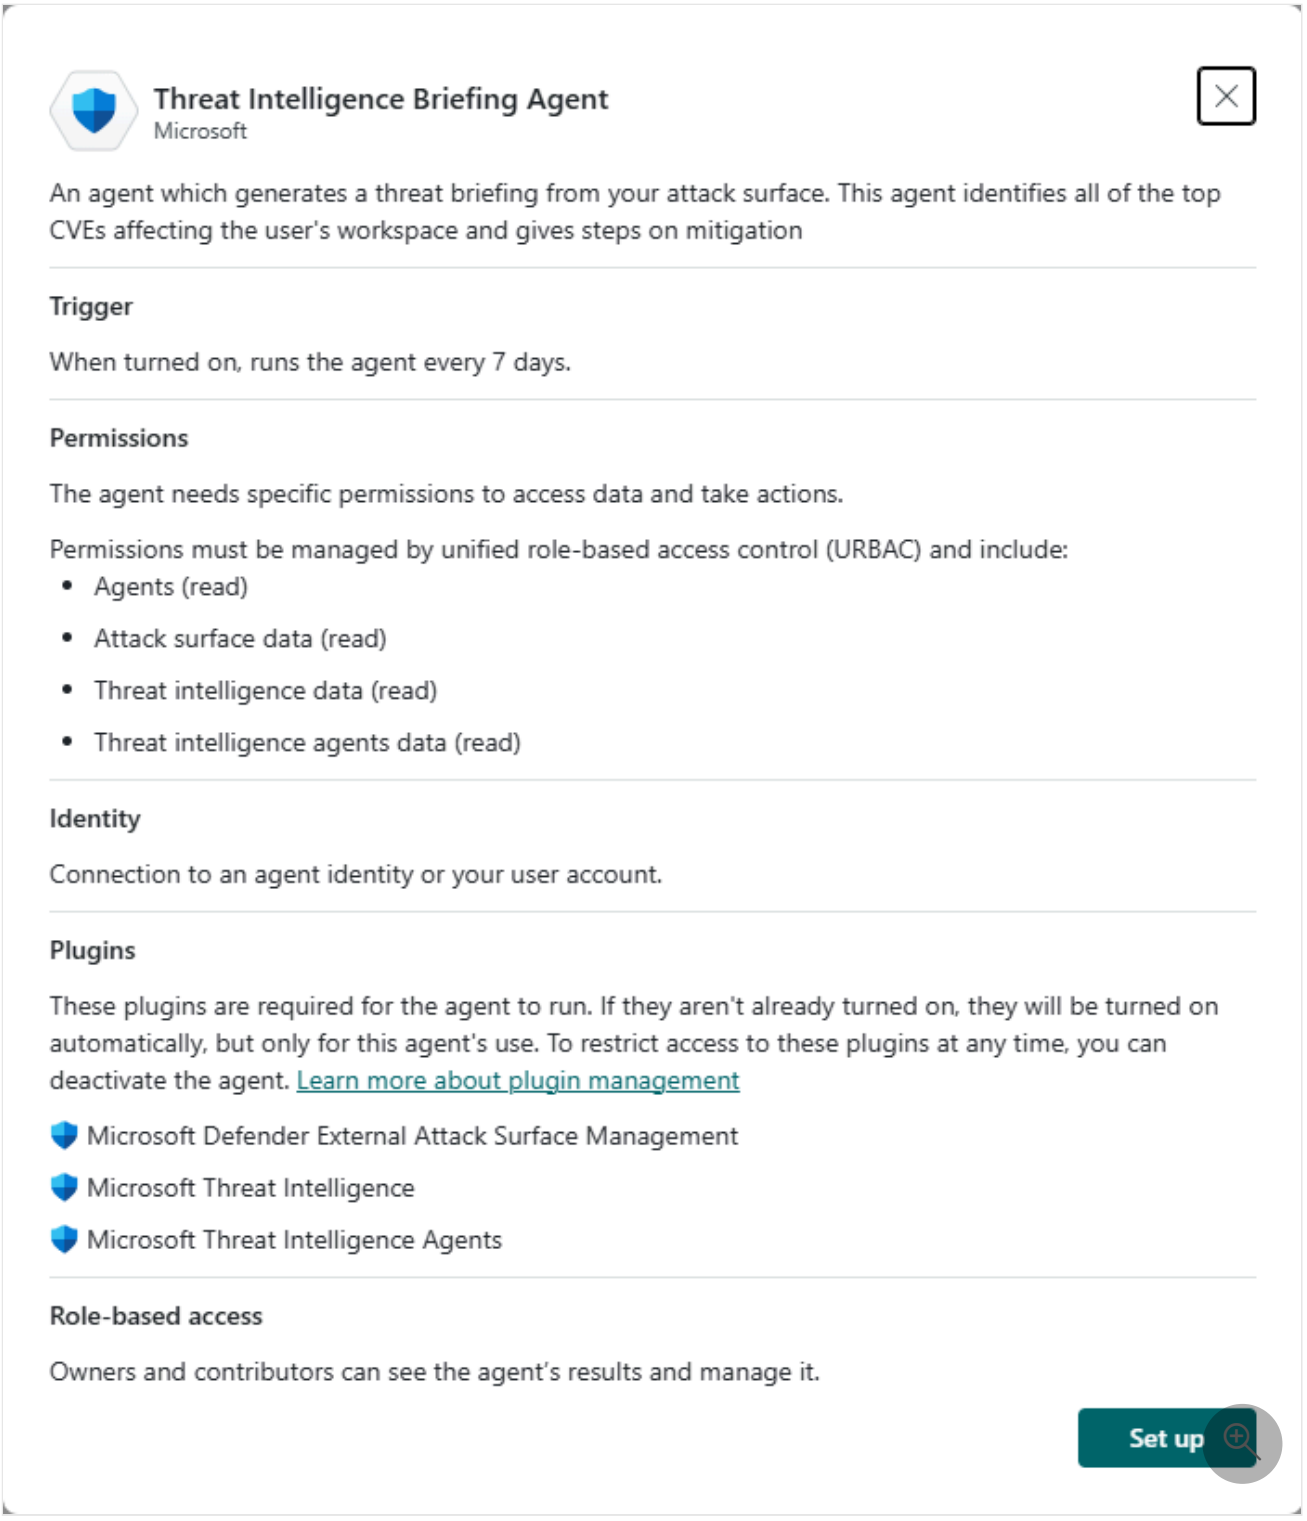

3. Select an identity for the agent. You have the option of choosing to create an agent identity or assign an existing user account. After this, wait for the agent to finish setting up.

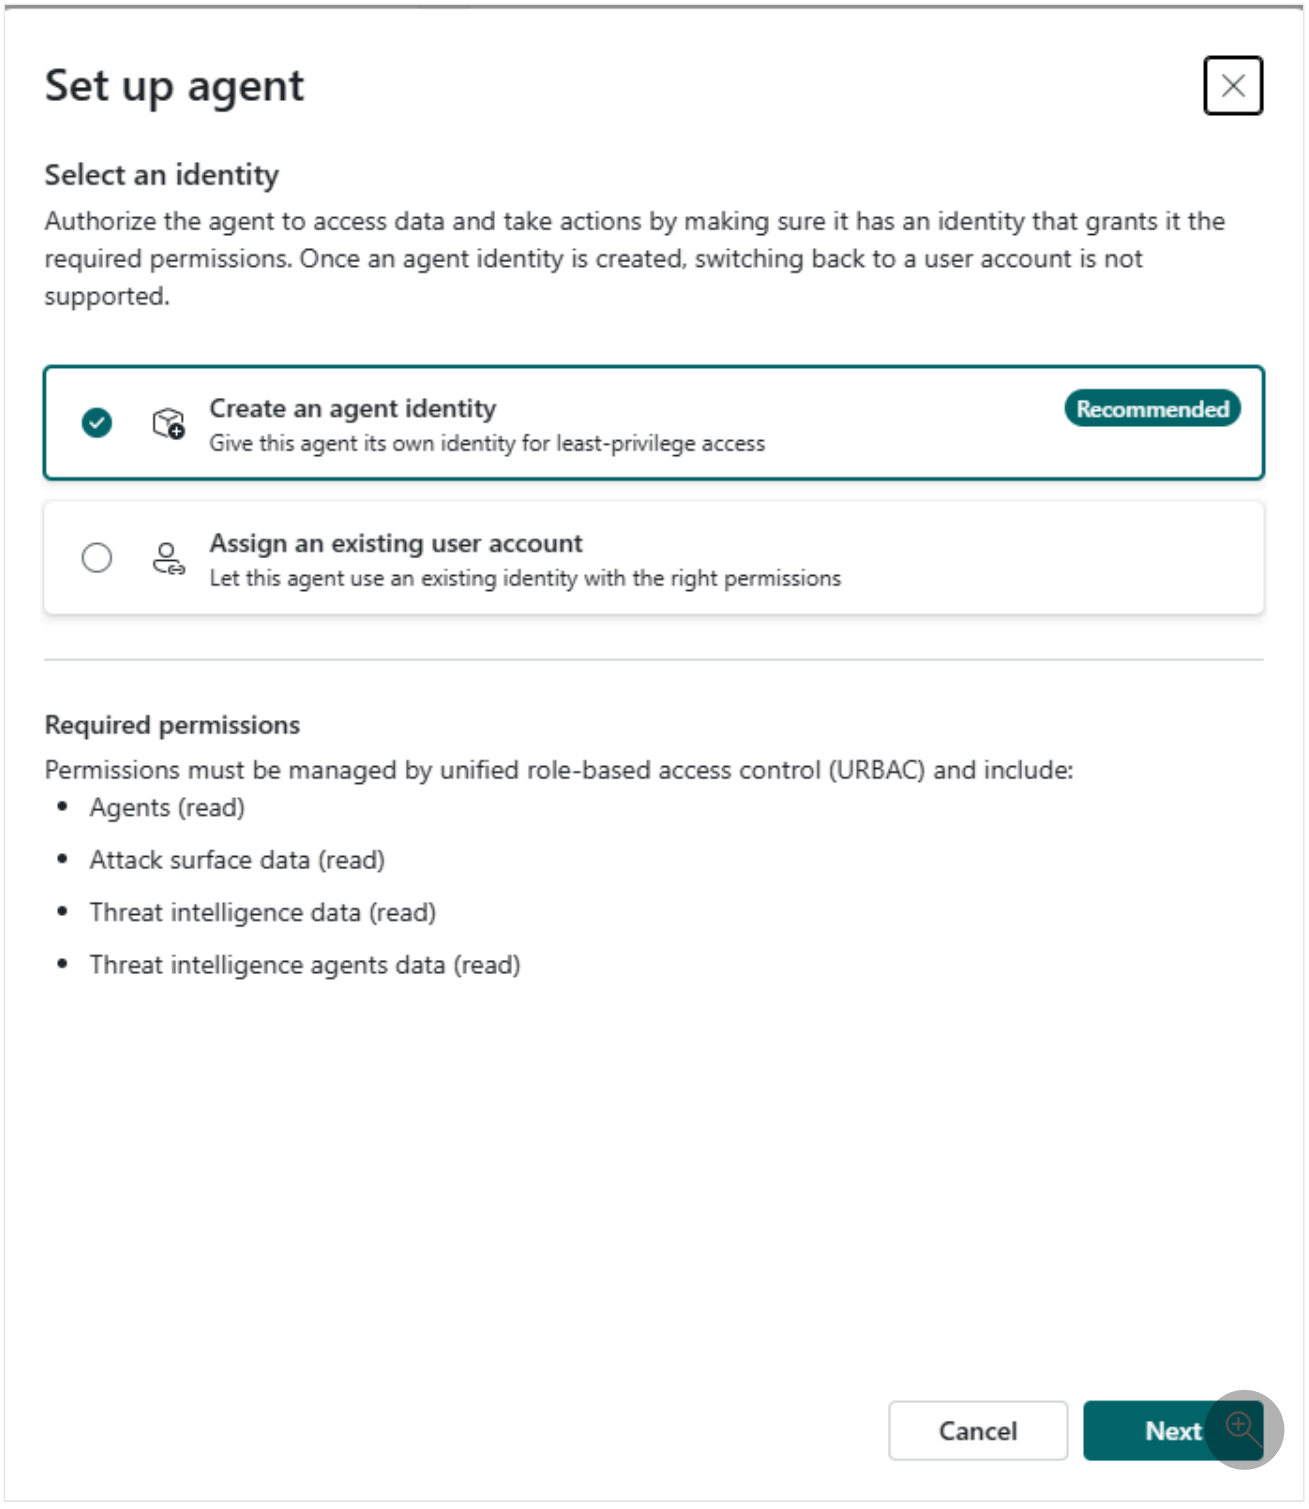

4. Specify the input parameters to customize the output, then select **Next**. You can edit these parameters later by selecting the three dots in the upper right section of the agent overview page.

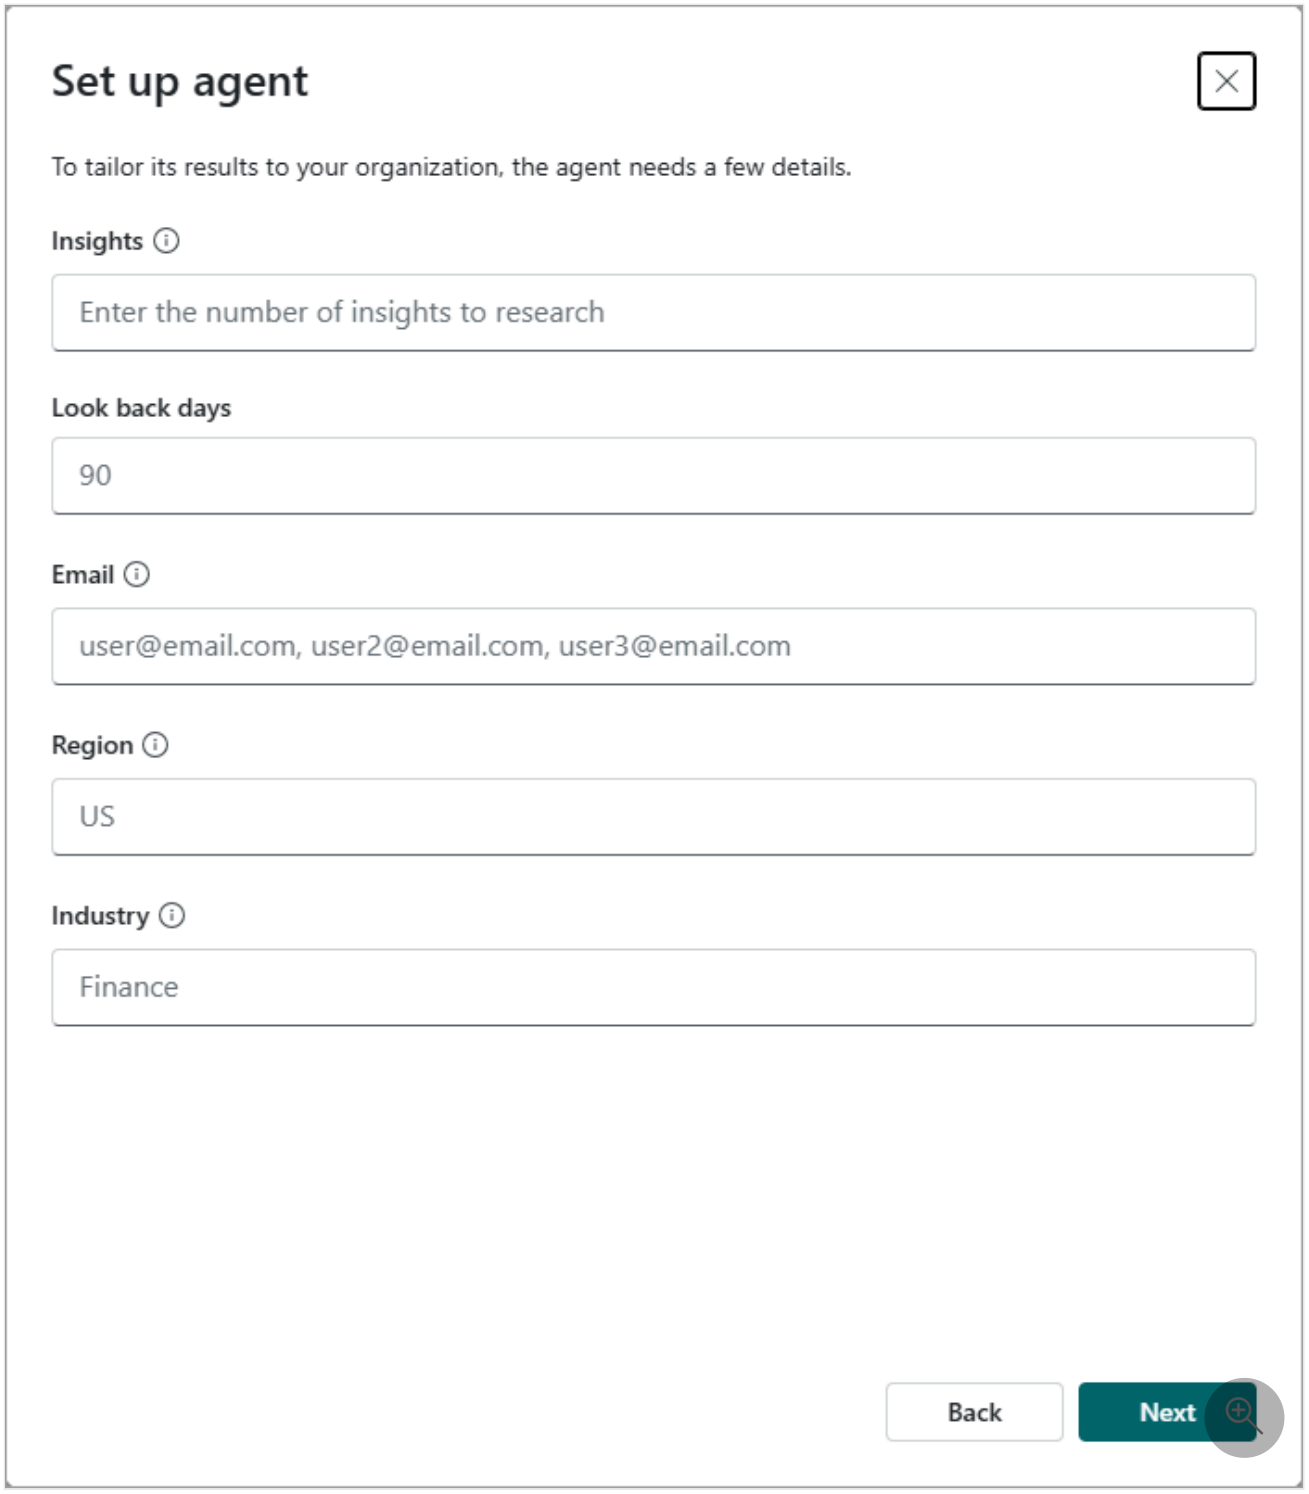

- Insights to research – the number of vulnerabilities the agent researches for active threats
- Look back days – how far back the agent researches threats against your vulnerabilities
- Email – email address of user or distribution group that the briefing is sent to
- Region – scope of geographical area the agent checks for threats
- Industry – sector or industry that the agent checks for threats

5. After the agent is created, select **Return to agents** to go back to the **Agents** page, or select **Go to agent** to go to the Threat Intelligence Briefing Agent overview page.

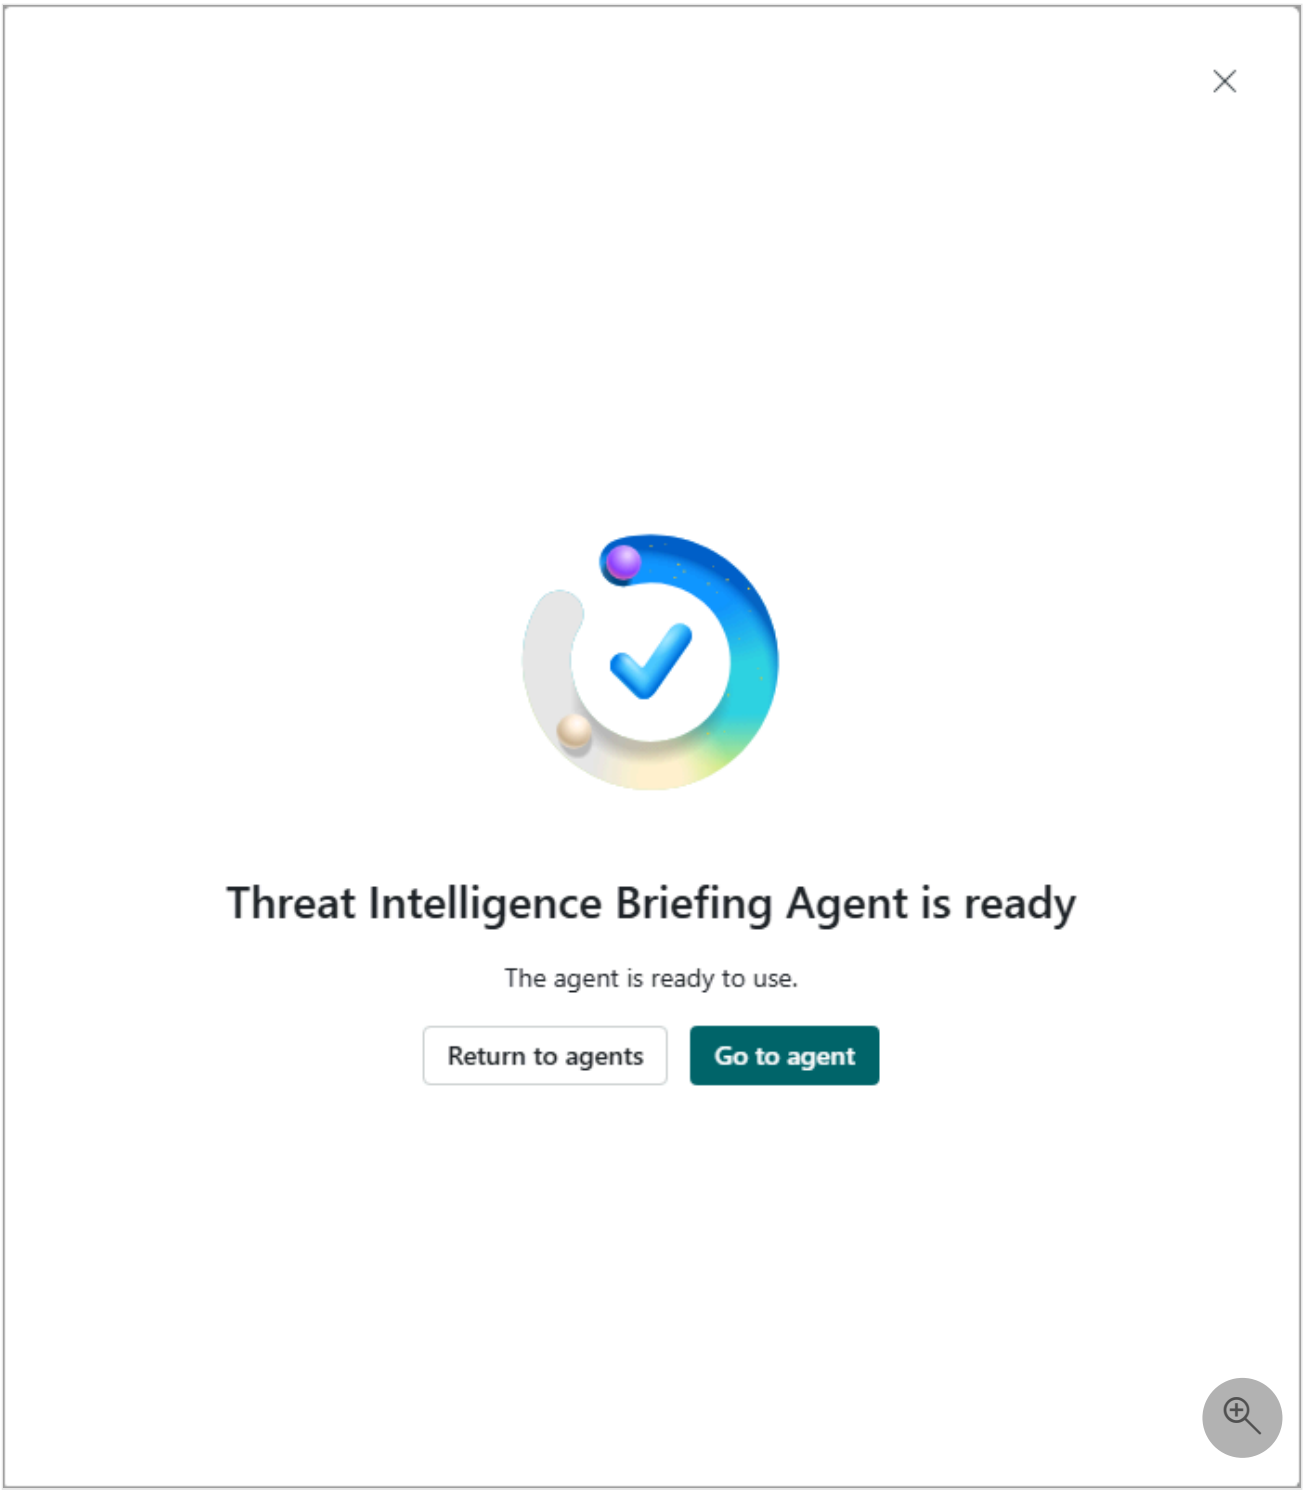

6. To run the agent, go to the upper right of the agent overview page and select **Run**. Select **On the trigger** to schedule the agent to run at the set time or select **One time** to run the report on demand.

Assess and provide feedback on the agent's output

The generated reports appear in the Threat Intelligence Briefing Agent page under **Activity**. The page displays the name of the report, the start time, method of generation, and current status.

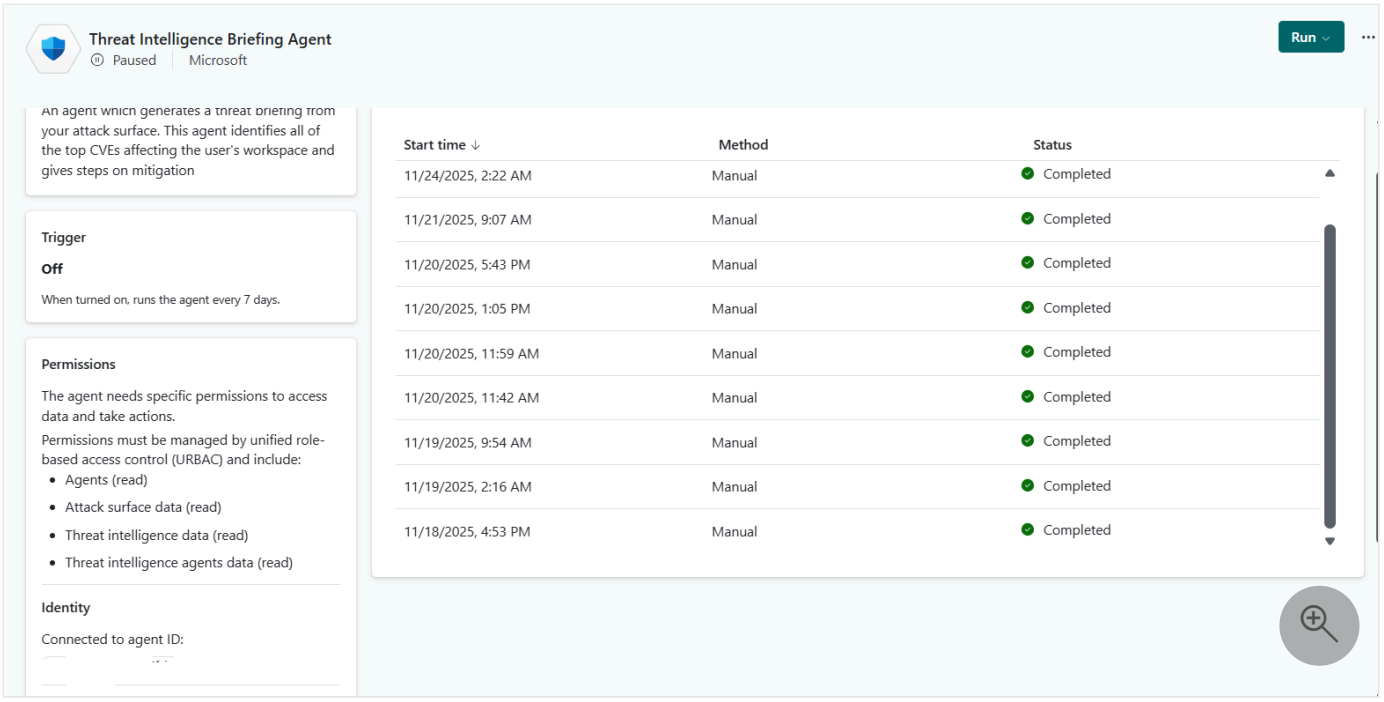

Select one of the reports to assess the agent's output.

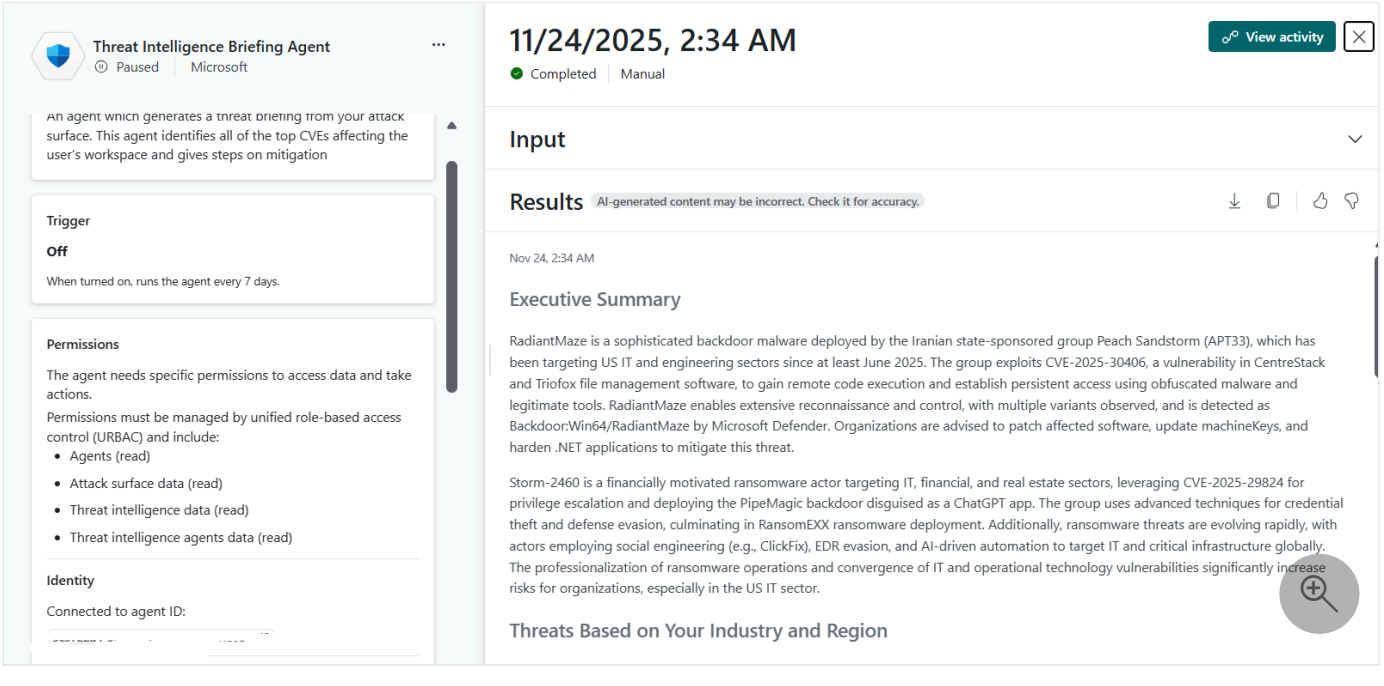

The threat intelligence briefing contains a relevant summary of threat information and detailed technical analysis, including any actively exploited vulnerability and its possible organizational impact.

The Threat Intelligence Briefing Agent dynamically chooses the next step based on the outcome of the previous step as it builds the briefing. You can view the agent's progress toward producing the threat briefing by selecting **View activity**.

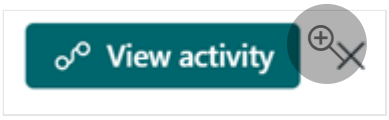

You'll see details of the activity, providing you with transparency on the steps the agent takes to produce the output.

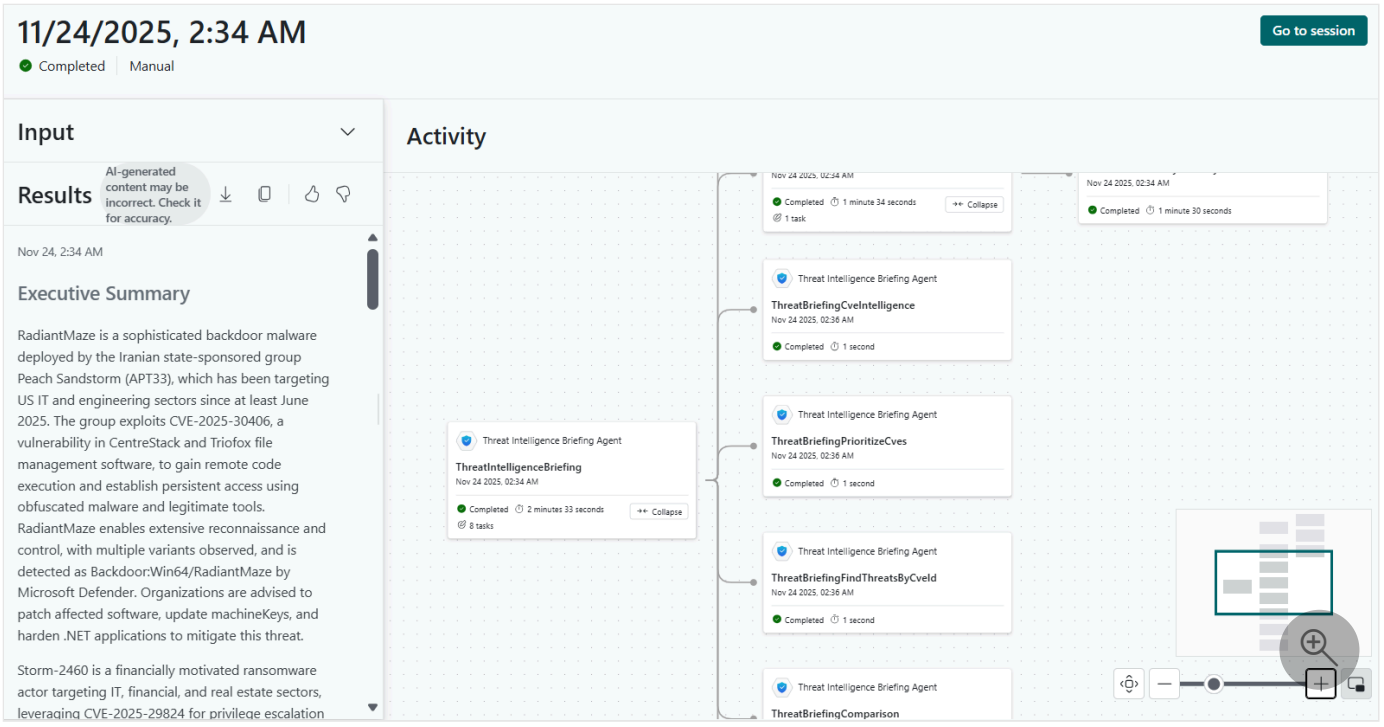

You can provide feedback about the briefing by selecting the thumbs up or thumbs down button. You can elaborate in the text box that appears after. Select **Submit** to give your feedback.

See also

- Microsoft Security Copilot agents
- Application card for Microsoft Security Copilot

Last updated on 05/11/2026

Deploy AI agents in Microsoft Defender

Applies to: Microsoft Defender XDR, Microsoft Sentinel in the Microsoft Defender portal

Security Store in the Microsoft Defender portal offers various agents that help you perform your security tasks efficiently. These agents include Microsoft Security Copilot agents published by Microsoft and partners. These agents integrate with Microsoft Defender and carry out various security operations (SOC) tasks, such as incident triage, investigation, threat hunting, and threat intelligence.

This article explains how to discover and deploy AI agents in Microsoft Defender.

> **Note**
>
> To learn more about publishing agents to Security Store, see Publish agents to Microsoft Security Store.

Prerequisites

To purchase, deploy, and use agents from Security Store, you need:

- Access to a Security Copilot workspace provisioned with SCU capacity.
- For partner-published agents, you need the Azure subscription contributor or owner role.

Discover and deploy agents in the Microsoft Defender portal

To discover and deploy agents in the Microsoft Defender portal:

1. Select **Security Copilot** > **Security Store**.
2. Browse or search for the agent you want to deploy.
3. Select the agent to view its details, including its capabilities, requirements, and setup instructions.
4. To purchase and deploy the agent:

- Select **Get agent** to begin the deployment process if you have sufficient permissions. For more information, see Prerequisites.
- Select **Copy link** to copy the agent's details page URL and share it with a security administrator, if you don't have permissions to deploy agents.

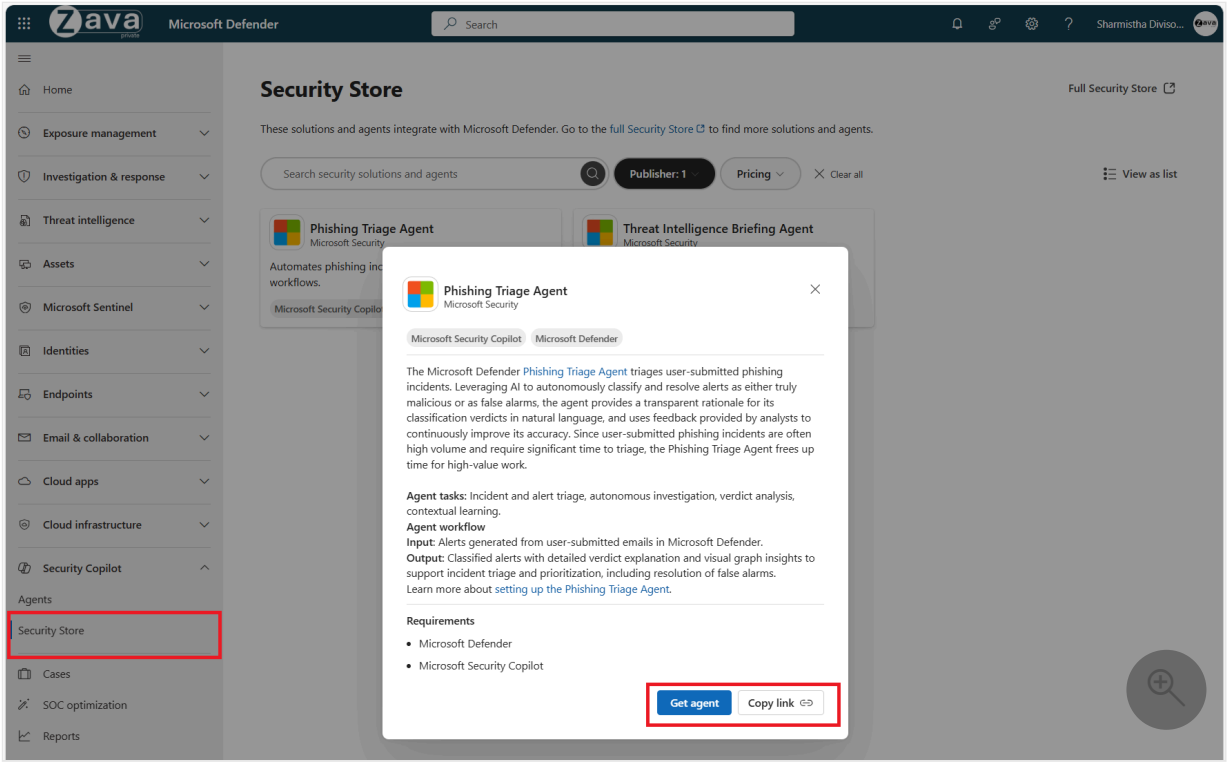

- For partner-published agents, complete the purchase and deploy on the Security Store website, as described in the Microsoft Security Store documentation.

  You can manage centralized purchases for partner-published agents through public offers, or through private offers, as described in How to Purchase SaaS Solutions (Private Offers).

5. After purchasing the agent, select **Security Copilot** > **Agents**, find your agent in the **Ready for setup** section, and then select **Set up** to begin agent setup.

   For more information on setting up, managing, and running partner-published agents, see Manage Security Copilot agents.

   For more information on Microsoft Security Copilot agents, see Microsoft Security Copilot agents in Microsoft Defender.

   After setup, the agent appears in the **Agents in use** section.

Microsoft Security Copilot agents in Microsoft Defender

This section details the Microsoft Security Copilot agents that are available in the Microsoft Defender portal.

- Security Alert Triage Agent (Preview)
- Threat Intelligence Briefing Agent
- Threat Hunting Assistant
- Security Analyst Agent
- Dynamic Threat Detection Agent

Security Alert Triage Agent (Preview)

> **Note**
>
> The Security Alert Triage Agent is the same agent as the **Phishing Triage Agent** with extended capabilities to triage a broader set of alert types. The Security Alert Triage Agent is available only to customers who are part of the preview. If you aren't part of the preview, the Phishing Triage Agent continues to be available in Security Store.

The Security Alert Triage Agent is an autonomous agent that helps security teams triage alerts at scale across multiple workloads. The agent applies AI-driven, dynamic reasoning to deliver clear verdicts for supported security workloads. This is the same agent as the Phishing Triage Agent, which has demonstrated measurable improvements in triage accuracy and efficiency. The agent can now triage a broader set of alerts in Microsoft Defender, including email and collaboration alerts (generally available), and cloud and identity alerts (preview). The agent operates autonomously, provides transparent rationale for its classification verdicts in natural language, and continuously learns and improves its accuracy based on feedback from analysts.

Expand table

Attribute	Description
Identity	Create a new agent identity or connect to an existing user account
License	Depends on alert types: • **Email and collaboration alerts:** Microsoft Defender for Office P2 • **Cloud alerts:** Microsoft Defender for Cloud and Microsoft Defender for Containers

Attribute	Description
	• **Identity alerts:** Entra ID P2 license, Microsoft Defender for Identity, and Microsoft Defender for Cloud Apps
Permissions	The agent requires these permissions to operate, depending on the alert types you want to triage: • **All alert types:** Security Copilot (read), Security data basics (read), Alerts (manage) • **Email and collaboration alerts:** Email & collaboration metadata (read), Email & collaboration content (read) • **Cloud and identity alerts:** No additional permissions required beyond the base permissions
Plugins	The agent automatically activates these Security Copilot plugins: • Microsoft Defender XDR • Microsoft Threat Intelligence • Security Alert Triage Agent
Products	• Security Copilot • Microsoft Defender for Office 365 Plan 2 (for email and collaboration alerts) • Microsoft Defender for Cloud (for cloud alerts) • Microsoft Defender for Identity and Microsoft Defender for Cloud Apps (for identity alerts)
Role-based access	Security Administrator Microsoft Entra role is required to set up and manage the agent Users with the same permissions as the Security Alert Triage Agent can view the agent's activity and results, and provide feedback on the agent's classification verdict.
Trigger	Runs automatically when a new alert is detected for the configured alert types, including user-reported emails (for email and collaboration alerts), cloud security alerts (for cloud alerts), and identity alerts.

Threat Intelligence Briefing Agent

The Threat Intelligence Briefing Agent provides security operations teams with regular, customized threat intelligence briefings. The agent autonomously gathers and synthesizes relevant threat intelligence data from various sources, delivering concise and actionable insights to help analysts stay informed about emerging threats and trends.

Expand table

Attribute	Description
Identity	Create a new agent identity or connect to an existing user account

Attribute	Description
License	**Not Applicable**
Permissions	**Required permissions:** • Microsoft Defender for Endpoint • Security Reader **Optional permissions:** • Exposure Management (read)
Products	Security Copilot
Plugins	The following plugins are required to run this agent: • Microsoft Threat Intelligence • Microsoft Threat Intelligence agents The following plugin is optional but can add more context to the output: • Microsoft Defender External Attack Surface Management
Role-based access	The **Security Administrator** role is required to set up and manage the agent. Users with the same permissions as the Threat Intelligence Briefing Agent can view the agent's activity and results.
Trigger	Runs at the set time interval that you configured during setup, or manually when you want to run it

Configure Defender for Endpoint permissions for the agent identity

When running the Threat Intelligence Briefing Agent with an agent identity, you must also configure the following Defender for Endpoint role permissions and Device Group access. Without these permissions, the Exposure Report may show as "not available" or return zero CVEs even when vulnerabilities exist in your environment.

Step 1 – Update the agent role permissions

1. Sign in to the Microsoft Defender portal.
2. Navigate to **Settings** > **Endpoints** > **Permissions** > **Roles**.
3. Locate the custom role assigned to the Threat Intelligence Briefing Agent.
4. Edit the role and confirm that the following permissions are enabled:

    - **Advanced Hunting** – Read

- **Vulnerability Management** – Read
- **Machine Configuration** – Read
- **Device Inventory** – Read

5. Save any changes if updates are made.

Step 2 – Grant Device Group access to the agent

1. In the Microsoft Defender portal, go to **Settings** > **Endpoints** > **Device Groups**.
2. For each Device Group that contains production endpoints:
   a. Open the Device Group.
   b. Select the **User Access** section.
   c. Add the Threat Intelligence Briefing Agent identity.
   d. Assign **Read** access.
3. Save the changes.

> ⓘ **Important**
>
> Allow time for permission updates to synchronize across Microsoft Defender services before running the agent.

Threat Hunting Assistant

The Threat Hunting Assistant revolutionizes threat hunting by enabling you to investigate threats using natural language from start to finish. It not only generates KQL queries but also interprets results, surfaces insights, and guides you through full hunting sessions. These capabilities empower you to hunt threats faster, more accurately, and with greater confidence.

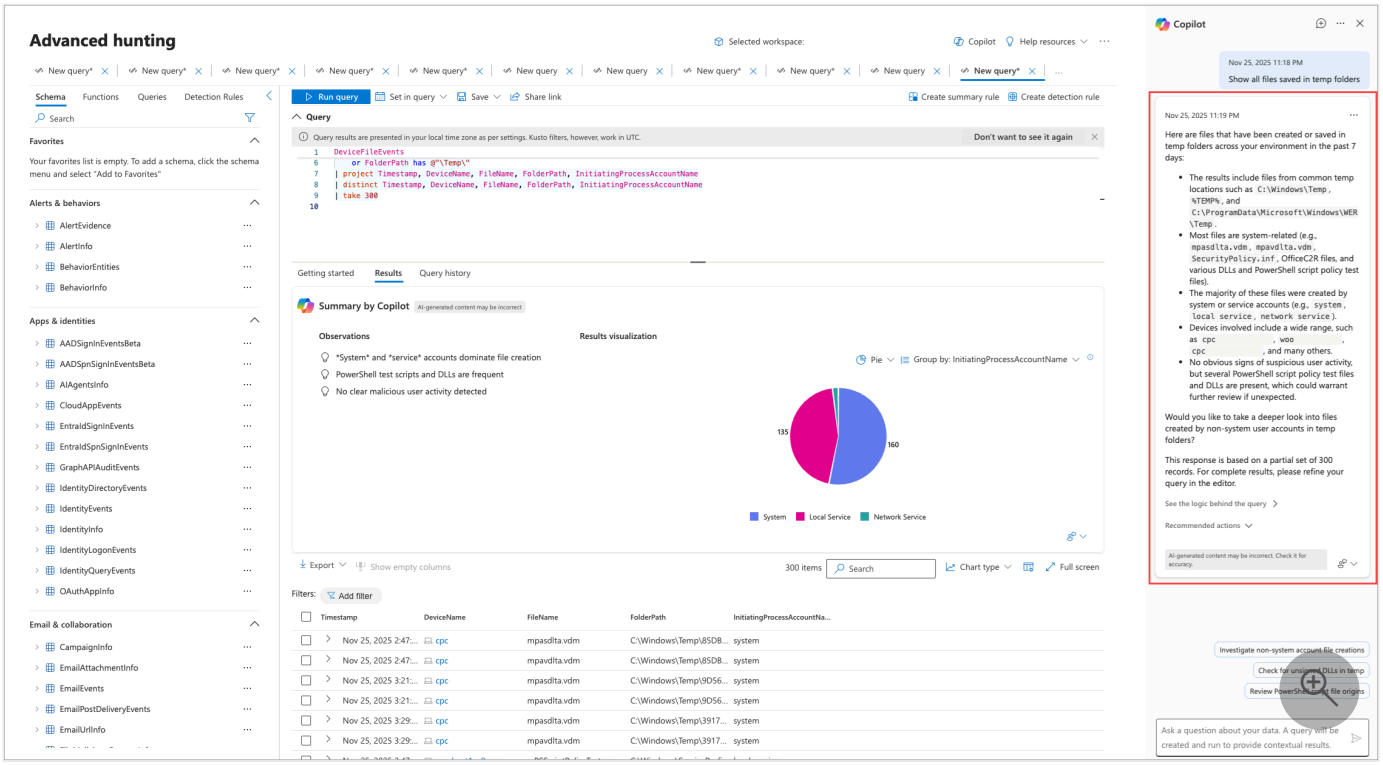

Security Analyst Agent

The Security Analyst Agent helps security analysts quickly identify, assess, and prioritize risks by performing ready-to-use or custom analyses on security data. The agent provides actionable and prioritized insights, recommendations, and reports to uncover top vulnerabilities and risks. It supports data from Microsoft Defender, Sentinel Log Analytics, or Sentinel Data Lake, and can perform complex analysis tasks such as anomaly detection, clustering, risk scoring, and forecasting without requiring code or queries.

Expand table

Attribute	Description
Identity	Tied to your user identity; each user configures the agent independently
License	**Not Applicable**
Permissions	Read access to Microsoft Defender XDR, Microsoft Sentinel Log Analytics Workspace, or Microsoft Sentinel Data Lake, depending on the data source you choose
Products	• Security Copilot • Microsoft Defender XDR (optional data source) • Microsoft Sentinel (optional data source)
Role-based access	Users with read access to the selected data sources can configure and use the agent.

Attribute	Description
Trigger	Runs on demand when you enter a security analysis prompt in the agent chat or select **Analyze with copilot** from advanced hunting query results

Dynamic Threat Detection Agent

The Dynamic Threat Detection Agent in the Defender portal is an always-on, adaptive backend service that uncovers hidden threats across Defender and Microsoft Sentinel environments. It uses AI to identify gaps and uncover false negatives by correlating alerts, events, anomalies, and threat intelligence. When the agent identifies a gap, it generates a dynamic alert with the full context in the alert details, including natural language explanations, mapped MITRE ATT&CK techniques, and tailored remediation steps.

Last updated on 08/05/2026

Microsoft Entra agents

Microsoft Entra agents can automate many identity and access management operations in your organization to help reduce manual workloads. These agents work seamlessly with Microsoft Security Copilot to automate repetitive tasks, provide suggestions, and help administrators focus on higher-value strategic work.

Microsoft Entra agents analyze your identity environment, apply best practices, and take automated actions to improve your identity and access security posture and operational efficiency. They integrate directly with Microsoft Entra services, using your organization's identity data and configuration to provide contextual, actionable insights.

What are Microsoft Entra agents?

Microsoft Entra agents are AI-powered tools that operate in your organization's identity environment to automate and optimize identity and access management tasks. The agents are grounded in the concepts and tasks for a specific product area, like Conditional Access. These agents can:

- **Automate routine tasks** - Handle time-consuming, repetitive identity and access management operations
- **Provide suggestions** - Analyze your environment and suggest improvements based on Microsoft best practices and Zero Trust principles
- **Operate autonomously** - Run on schedules or triggers to continuously monitor and optimize your identity infrastructure
- **Integrate seamlessly** - Work within your organization's existing Microsoft Entra workflows
- **Learn and adapt** - Improve suggestions over time, based on your environment and feedback

Each agent works a little differently, but at their core, they first analyze your current environment within the boundaries of the agent's capabilities. If the agent identifies a gap, opportunity, or potential issue, it can take action on your behalf. Each agent provides the context, reasoning, and activity history for how it came up with the suggestion.

Administrators can configure the agent to run automatically or trigger the agent to run manually.

Because each of the agents perform a specific set of tasks, they need a specific set of configurations to operate within the boundaries of that task. The administrator also needs certain

Microsoft Entra roles to set up and manage the agent.

- **Agent identity**: A unique agent identity is created when the agent is turned on. Learn more about agent identities.
- **Roles**: Specific Microsoft Entra built-in roles are needed to turn on, view, and interact with the agent. Not all roles can perform the same tasks with an agent.
- **Permissions**: The agent identity is granted specific read and write permissions needed to perform its tasks. These permissions can't be changed or removed.
- **Role-based access**: The administrator needs specific roles to set up, manage, and use the agent.

Available Microsoft Entra agents

The following agents are currently available for Microsoft Entra. Due to the fast pace at which these agents are released and updated, each agent might have features at various stages of availability. Preview features are added frequently.

Conditional Access Optimization Agent

The Conditional Access Optimization Agent ensures comprehensive user protection by analyzing your Conditional Access policies and recommending improvements. The agent evaluates your current policy configuration against Microsoft best practices and Zero Trust principles.

Expand table

Attribute	Description
Identity	A unique agent identity for authorization is created when the agent is turned on. The agent uses this identity to scan your tenant's Conditional Access policies and configurations for gaps, overlap, and misconfigurations.
Licenses	Microsoft Entra ID P1
Permissions	AuditLog.Read.All CustomSecAttributeAssignment.Read.All DeviceManagementApps.Read.All DeviceManagementConfiguration.Read.All GroupMember.Read.All LicenseAssignment.Read.All NetworkAccess.Read.All Policy.Create.ConditionalAccessRO Policy.Read.All

Attribute	Description
	RoleManagement.Read.Directory User.Read.All
Plugins	Microsoft Entra
Products	Microsoft Entra Conditional Access
Role-based access	Security Administrator to configure the agent Conditional Access Administrator to use the agent
Trigger	Runs every 24 hours or triggered manually

Identity Risk Management Agent (Preview)

The Identity Risk Management Agent in Microsoft Entra ID Protection helps administrators investigate potential risks, learn about potential effects, and take decisive action to protect their organization's critical assets.

Expand table

Attribute	Description
Identity	Uses Microsoft Entra Agent ID for authorization
Licenses	Microsoft Entra Agent ID
Permissions	Application.Read.All Policy.Read.All Group.ReadWrite.All GroupMember.Read.All User.Read.All Policy.ReadWrite.ConditionalAccess CustomSecAttributeAssignment.Read.All IdentityRiskyUser.Read.All AuditLog.Read.All
Plugins	Microsoft Entra
Products	Security Copilot Microsoft Entra ID Protection
Role-based access	Security Administrator
Trigger	Runs every 24 hours, triggered manually, or continuous monitoring

Discover agents in the Security Store

Security Store is embedded in the Microsoft Entra admin center, providing a centralized place to discover, purchase, and deploy Microsoft and partner-built agents and solutions. You can browse available agents and solutions, view details and requirements, and start the deployment process directly from the Entra portal.

For more information, see Discover and deploy agents and solutions in Microsoft Entra.

Getting started with Microsoft Entra agents

Prerequisites

- You must have available security compute units (SCU).
  - In order to purchase security compute units, you need to have an Azure subscription. Create your free Azure account.
- Review Privacy and data security in Microsoft Security Copilot

Setup process

1. Enable Security Copilot using the Security Copilot setup guide.
2. Sign in to the Microsoft Entra admin center using the least privileged role required for the agent you want to configure.
3. Browse to **Agents** and select **View details** for the agent you want to configure.

Agents in the Microsoft ecosystem

While this article focuses on Microsoft Entra agents, similar agents are available across other Microsoft security products. For more information, see Microsoft Intune, Microsoft Defender, and Microsoft Purview.

Related content

- Microsoft Security Copilot in Microsoft Entra
- Security Copilot scenarios in Microsoft Entra
- Microsoft Security Copilot agents overview

Last updated on 04/01/2026

Troubleshoot Conditional Access policies for Microsoft Security Copilot

Overview

Generative artificial intelligence (AI) services like Microsoft Security Copilot can bring value to your organization when used appropriately.

Apply Conditional Access policy to these generative AI services by following our recommendation to target all resources. These policies might include those for all users, risky users, sign-ins, device compliance, and users with insider risk.

Some organizations target these services directly by using the underlying service principals and custom security attributes in their Conditional Access policies:

- 43d7b169-1d9e-4d32-8cd8-06c5974ed90c - Security Copilot Agent Management
- bb5ffd56-39eb-458c-a53a-775ba21277da - Security Copilot Portal
- bb3d68c2-d09e-4455-94a0-e323996dbaa3 - Security Copilot API
- b0cf1501-8e0f-4fbb-b70a-52ca5ea7bda6 - Security Copilot Logic Apps Connector

In these cases, admins create, assign, and target these underlying service principals with custom security attributes.

Required roles

Custom security attributes are security sensitive and only delegated users can manage them. Assign one or more of the following roles to the user who manages or reports on these attributes.

Expand table

Role name	Description
Attribute Assignment Administrator	Assign custom security attribute keys and values to supported Microsoft Entra objects.
Attribute Assignment Reader	Read custom security attribute keys and values for supported Microsoft Entra objects.
Attribute Definition Administrator	Define and manage the definition of custom security attributes.
Attribute Definition Reader	Read the definition of custom security attributes.

Assign the appropriate role to the users who manage or report on these attributes at the directory scope. For detailed steps, see Assign Microsoft Entra roles.

> **Important**
>
> By default, Global Administrator and other administrator roles do not have permissions to read, define, or assign custom security attributes.

Create custom security attributes

Follow the instructions in the article Add or deactivate custom security attributes in Microsoft Entra ID to add the following **Attribute set** and **New attributes**.

- Create an **Attribute set** named *SecurityCopilotAttributeSet*.
- Create **New attributes** named *SecurityCopilotAttribute* with **Allow multiple values to be assigned** set to **No** and **Only allow predefined values to be assigned** set to **Yes**. Add the following predefined value:
  - MFARequired

> **Note**
>
> Conditional Access filters for applications only work with custom security attributes of type "string". Custom security attributes support creating the Boolean data type, but Conditional Access Policy only supports "string".

Assign custom security attributes to applications

1. Sign in to the Microsoft Entra admin center as at least a Conditional Access Administrator and Attribute Assignment Administrator.
2. Browse to **Entra ID** > **Enterprise apps**.
3. Select the apps to apply a custom security attribute to:
   a. 43d7b169-1d9e-4d32-8cd8-06c5974ed90c - Security Copilot Agent Management
   b. bb5ffd56-39eb-458c-a53a-775ba21277da - Security Copilot Portal
   c. bb3d68c2-d09e-4455-94a0-e323996dbaa3 - Security Copilot API
   d. b0cf1501-8e0f-4fbb-b70a-52ca5ea7bda6 - Security Copilot Logic Apps Connector
4. Under **Manage** > **Custom security attributes**, select **Add assignment**.
5. Under **Attribute set**, select the attribute set you created.
6. Under **Attribute name**, select the attribute you created.

7. Under **Assigned values**, select **Add values**, choose the value you created from the list, then select **Done**.
8. Select **Save**.

Targeting custom security attributes in Conditional Access policy

1. Sign in to the Microsoft Entra admin center as at least a Conditional Access Administrator and Attribute Definition Reader.
2. Browse to **Entra ID** > **Conditional Access**.
3. Select **New policy** or select an existing policy to update.
4. When configuring your **Target resources**, select the following options:
   a. Select what this policy applies to **Cloud apps**.
   b. Include **Select resources**.
   c. Select **Edit filter**.
   d. Set **Configure** to **Yes**.
   e. Select the **Attribute** you created.
   f. Set **Operator** to **Contains**.
   g. Set **Value** to one of your custom attributes.
   h. Select **Done**.

Which policy is causing issues?

It's sometimes hard for an admin to check which policy to update when there's an issue. Use the guidance in Troubleshooting sign-in problems with Conditional Access to check which policies apply, which policies don't apply, and run sign-in diagnostics to avoid ongoing issues.

Next steps

Understand authentication in Microsoft Security Copilot

Last updated on 03/24/2026

Security Copilot agents in Intune overview

Security Copilot agents in Intune are AI-powered assistants that enhance enterprise security. They automate tasks for endpoint protection, identity management, threat intelligence, and device configuration. They help IT teams quickly address vulnerabilities, policy gaps, and emerging threats.

Agents are built on Microsoft Security Copilot's generative AI and automation capabilities. They observe, reason, and act with oversight and review from your administrators. Each agent is tailored to a specific use case, operates within the Intune admin center, and uses role-based access controls.

Available agents

Microsoft Intune includes specialized Security Copilot agents, each designed for a specific security scenario. The following agents are available:

Change Review Agent

> The *Change Review Agent* evaluates the effect of approval requests in Intune and makes recommendations for the actions you can take.
>
> Learn more

Device Offboarding Agent

> The *Device Offboarding Agent* identifies stale or misaligned devices across Intune and Microsoft Entra ID, providing actionable insights and requiring admin approval before offboarding any devices.
>
> Learn more

Policy Configuration Agent

> With the *Policy Configuration Agent*, you import documents or write instructions in plain language. The agent uses this information to find matching settings in the Intune settings catalog and recommends values for these settings. You can then use the agent to create a policy with those settings and their values.
>
> Learn more

Vulnerability Remediation Agent

> The *Vulnerability Remediation Agent* uses Defender data to monitor vulnerabilities and prioritize remediation with AI-driven risk assessments.
>
> Learn more

Getting started with agents

Prerequisites

Before you begin, make sure you have:

- Security compute units (SCU) available
- Reviewed the Privacy and data security in Microsoft Security Copilot to understand how your data is handled.

Setup process

1. Enable Security Copilot using the Security Copilot setup guide.

2. Sign in to the Microsoft Intune admin center using the least privileged role required for the agent you want to configure.
3. Browse to **Agents** and select **View details** for the agent you want to configure.

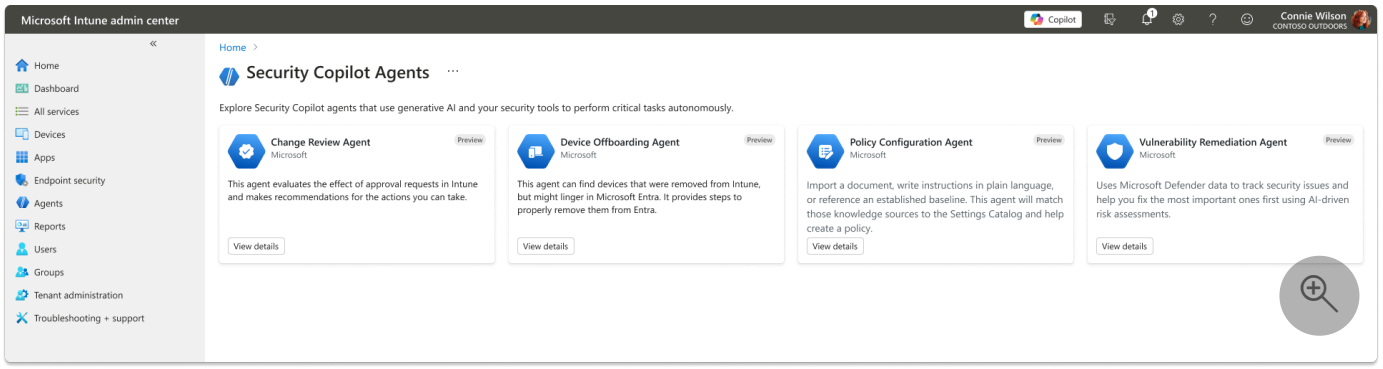

Agents in the Microsoft ecosystem

While this article focuses on Intune agents, similar agents are available across other Microsoft security products. For more information, see:

- Microsoft Entra
- Microsoft Defender
- Microsoft Purview

Related content

- Microsoft Security Copilot agents overview

Help shape the future of Intune agents

Join our **Intune Agents Feedback Forum** to share insights and influence upcoming capabilities in Microsoft Intune.

Sign up and learn more: https://aka.ms/IntuneAgentsForum

Last updated on 04/01/2026

Microsoft Security Copilot agents in Purview (preview)

Summarize this article for me

In preview, Microsoft Security Copilot agents are available in Microsoft Purview to help triage Microsoft Purview Data Loss Prevention (DLP) and Microsoft Purview Insider Risk Management alerts. These Microsoft Purview agents are AI-powered assistants that can help you triage alerts by working seamlessly with Microsoft security products.

Available agents

Triage Agent in Insider Risk Management

The Triage Agent in Insider Risk Management helps security teams by evaluating alerts based on user risk, file risk, and activity risk. The agent then sorts the triaged alerts into four categories. These categories are presented in the insider risk management solution on the **Alerts** tab.

Expand table

Attribute	Description
Identity	Runs as the administrator who turned on the agent. Agent authentication expires after 90 days and must be renewed.
License	Both the standard per seat licensing model and the pay-as-you-go billing model Microsoft Purview Insider Risk Management with Microsoft 365 E3/E5/A5/F5/G5 . For information on Security Copilot licensing in E5 see, Learn about Security Copilot in Microsoft 365 E5.
Permissions	Access policy configurations and settings in Insider Risk Management Read activities and events in Microsoft Purview Read file content and metadata involved in Insider Risk Management alerts Store user feedback and apply feedback when evaluating Insider Risk Management alerts
Plugins	Microsoft Purview
Products	Security Copilot Insider Risk Management
Role-based access	**View activity**: Insider Risk Management Analysts, Insider Risk Management Investigators, or Insider Risk Management role group

Attribute	Description
	Manage: All roles needed to view activity, plus the Purview Content Analyst role in the Purview Agent Management role group
Trigger	Runs on a selected schedule or on one alert at a time

Alert Triage Agent in Data Loss Prevention (preview)

The Alert Triage Agent in Data Loss Prevention (DLP) helps security teams by evaluating alerts based on the sensitivity risk, exfiltration risk, and policy risk. The agent then sorts the triaged alerts into four categories. These categories are presented in the DLP solution on the **Alerts** page.

For devices, you must set up evidence collection for file activities on devices and enable evidence collection in the DLP policy rule configuration.

Expand table

Attribute	Description
Identity	Runs as the administrator who turned on the agent. Agent authentication expires after 90 days and must be renewed.
License	Both the standard per seat licensing model and the pay-as-you-go billing model Microsoft Purview Data Loss Prevention with Microsoft 365 E3/E5/A5/F5/G5
Permissions	Access policy configurations and settings in DLP Read activities and events in Microsoft Purview Read file content and metadata involved in DLP alerts Store user feedback and apply feedback when evaluating DLP alerts Read file content and metadata involved in DLP alerts Store user feedback and apply feedback when evaluating DLP alerts
Plugins	Microsoft Purview
Products	Security Copilot Data Loss Prevention
Role-based access	**View activity**: Insider Risk Management Analysts, Insider Risk Management Investigators, or Insider Risk Management role group **Manage**: All roles needed to view activity, plus must be assigned the Purview Content Analyst role in the Purview Agent Management role group
Trigger	Runs on a selected schedule or on one alert at a time

Last updated on 02/24/2026

Partner-built agents

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

Organizations are increasingly using AI-powered agents to enhance their security operations. While Microsoft Security Copilot agents provide robust solutions, integrating partner agents can offer other benefits and provides you with the flexibility to use tools you're already familiar with. These agents offer unique capabilities, from privacy breach response to network supervision and alert triage, ensuring you can address diverse security challenges effectively.

There's a growing ecosystem of partner-built agents in Security Copilot. To find the most up-to-date list of agent offerings and how to deploy them, see Security Store.

Developers also have the option of building custom agents. For more information, see Agents development overview.

See also

- Microsoft Security Copilot agents overview

Last updated on 03/18/2026

Microsoft Security Store in Microsoft Security Copilot

09/30/2025

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

Security Store is Microsoft Security's new security-optimized storefront that enables organizations to find, try, buy, and deploy Microsoft and partner-built security solutions and agents. These solutions and agent integrate directly with Microsoft Security products such as Microsoft Defender, Microsoft Sentinel, Microsoft Entra, Purview, Intune and Security Copilot, offering a unified experience that simplifies procurement and accelerates time to value.

Security Store is integrated into the standalone experience of Microsoft Security Copilot, offering users a streamlined way to discover and acquire security agents directly within the portal.

Analysts and other security stakeholders can explore available agents both Microsoft and non-Microsoft, without leaving the Security Copilot portal. After you find an agent you'd like to use, you can proceed to the Security Store to purchase or subscribe to the agent or solution that you want to deploy. Security Store is where you manage billing as well.

After getting the agent or solution, you can then configure and operate the agents within Microsoft Security Copilot, maintaining a clear separation between purchasing and operational management.

Discover and get agents

1. Sign in to Security Copilot (https://securitycopilot.microsoft.com) and go to the **Home menu** > **Security Store**.

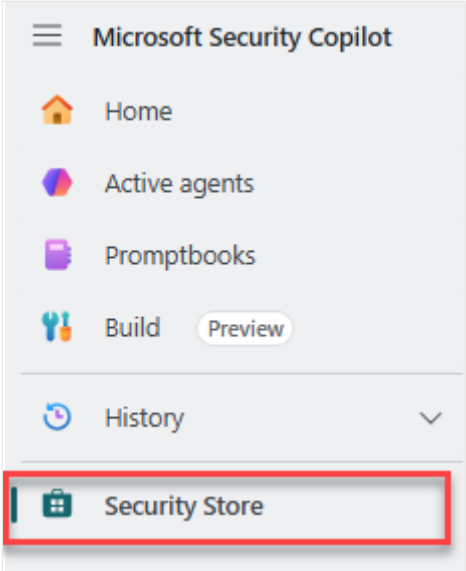

2. Search for an agent and refine your results by using filters such as publisher, product, or pricing. Agents can also be viewed in a list format.

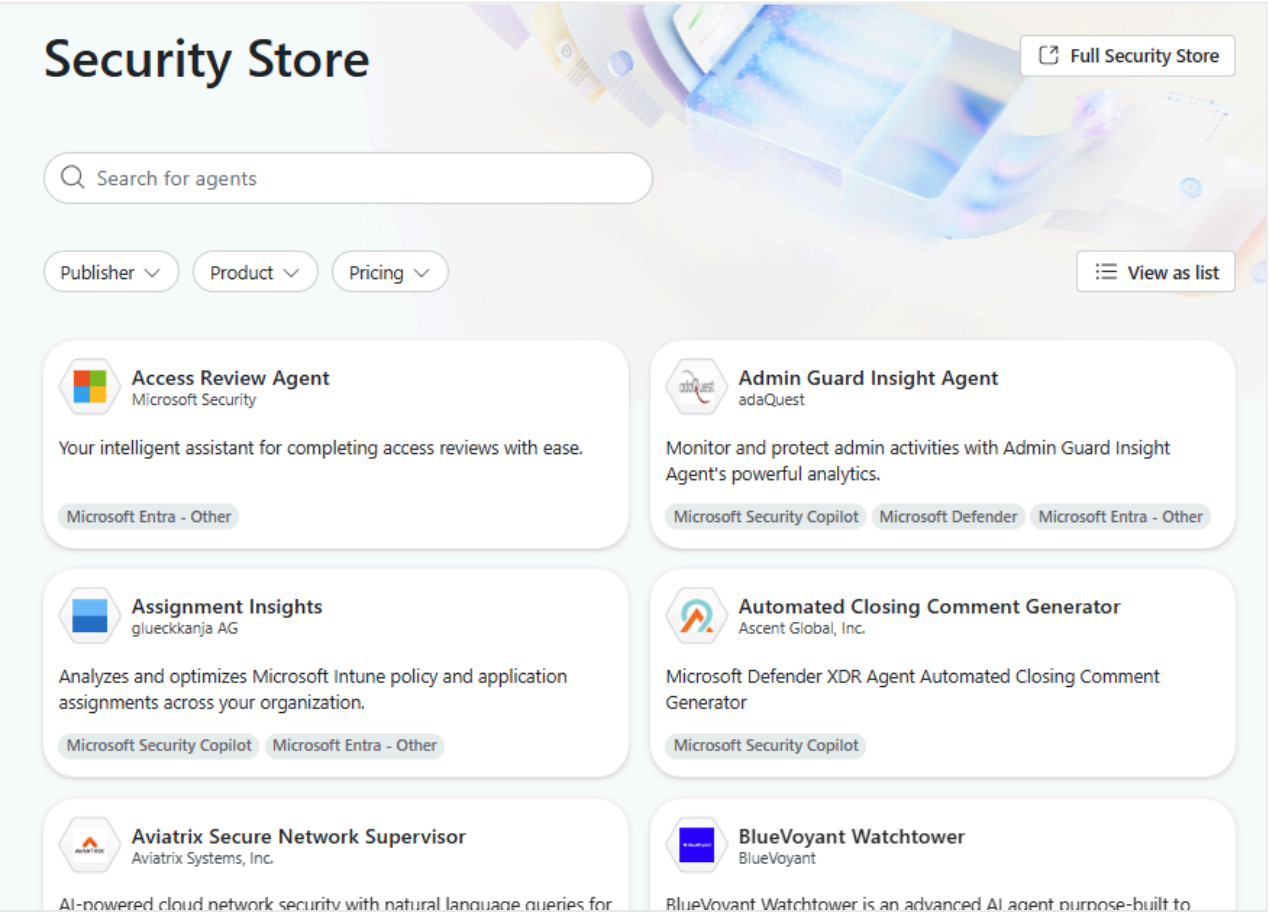

3. Select the agent you'd like to use and then select **Get agent**.

   For Microsoft agents, you'll be routed to the **Active agents** page.

   For Non-Microsoft agents, you'll be routed to Security Store where you can purchase or subscribe to agents. For more information, see Get Agents in Security Store.

> **Important**

Purchases and subscriptions done in Security Store are handled separately from the Security Copilot platform. Removing an agent from Security Copilot doesn't end a subscription, you'll need to manage the billing in Security Store.
Using agents within Security Copilot incurs Security Compute Units (SCUs), which are billed independently from any subscription fees.

After getting an agent and setting it up, you'll see the agent in the **Active agents** section.

If the agent has a dependent plugin, the plugin is enable for the agent to use (but not configured). To configure the plugin, go to **Manage sources** > find the plugin > configure the plugin. After completing the configuration steps, you can then use the agent successfully.

If the agent doesn't rely on a dependent plugin, you can configure it using the standard setup process.

Tip

If you're a developer and want to create agents and publish to Security Store, see **Custom agent overview**.

Setup and Manage Security Copilot agents

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

This article provides guidance on how Owners and Contributors can effectively set up and manage agents.

This topic helps you learn about:

- Setting up an agent
- Providing feedback directly to an agent to improve its performance
- Controlling an agent's activity by running or pausing it
- Editing its settings
- Managing its memory to ensure it produces results that are relevant and accurate
- Browsing for agents in Security Store

These management controls give you the flexibility to customize agents to fit your workflows while ensuring you maintain full control over their capabilities and behavior.

> **Note**
>
> Depending on how your administrator configured access determines if that agent is available in your workspace.

Set up Microsoft-built agents

You must set up an agent in order to use it. Follow these setup steps for agents built by Microsoft.

1. Sign in to Security Copilot (https://securitycopilot.microsoft.com).
2. In the left navigation pane, select **Agents**. The agents library is displayed.
3. Choose an agent that you want to use and select **Set up** to proceed to the agent details page.
4. Select an identity for the agent. You have the option of choosing to create an agent identity or assign an existing user account.

> **Tip**
>
> It's recommended that you create an agent identity. For more information, see What are agent identities?.

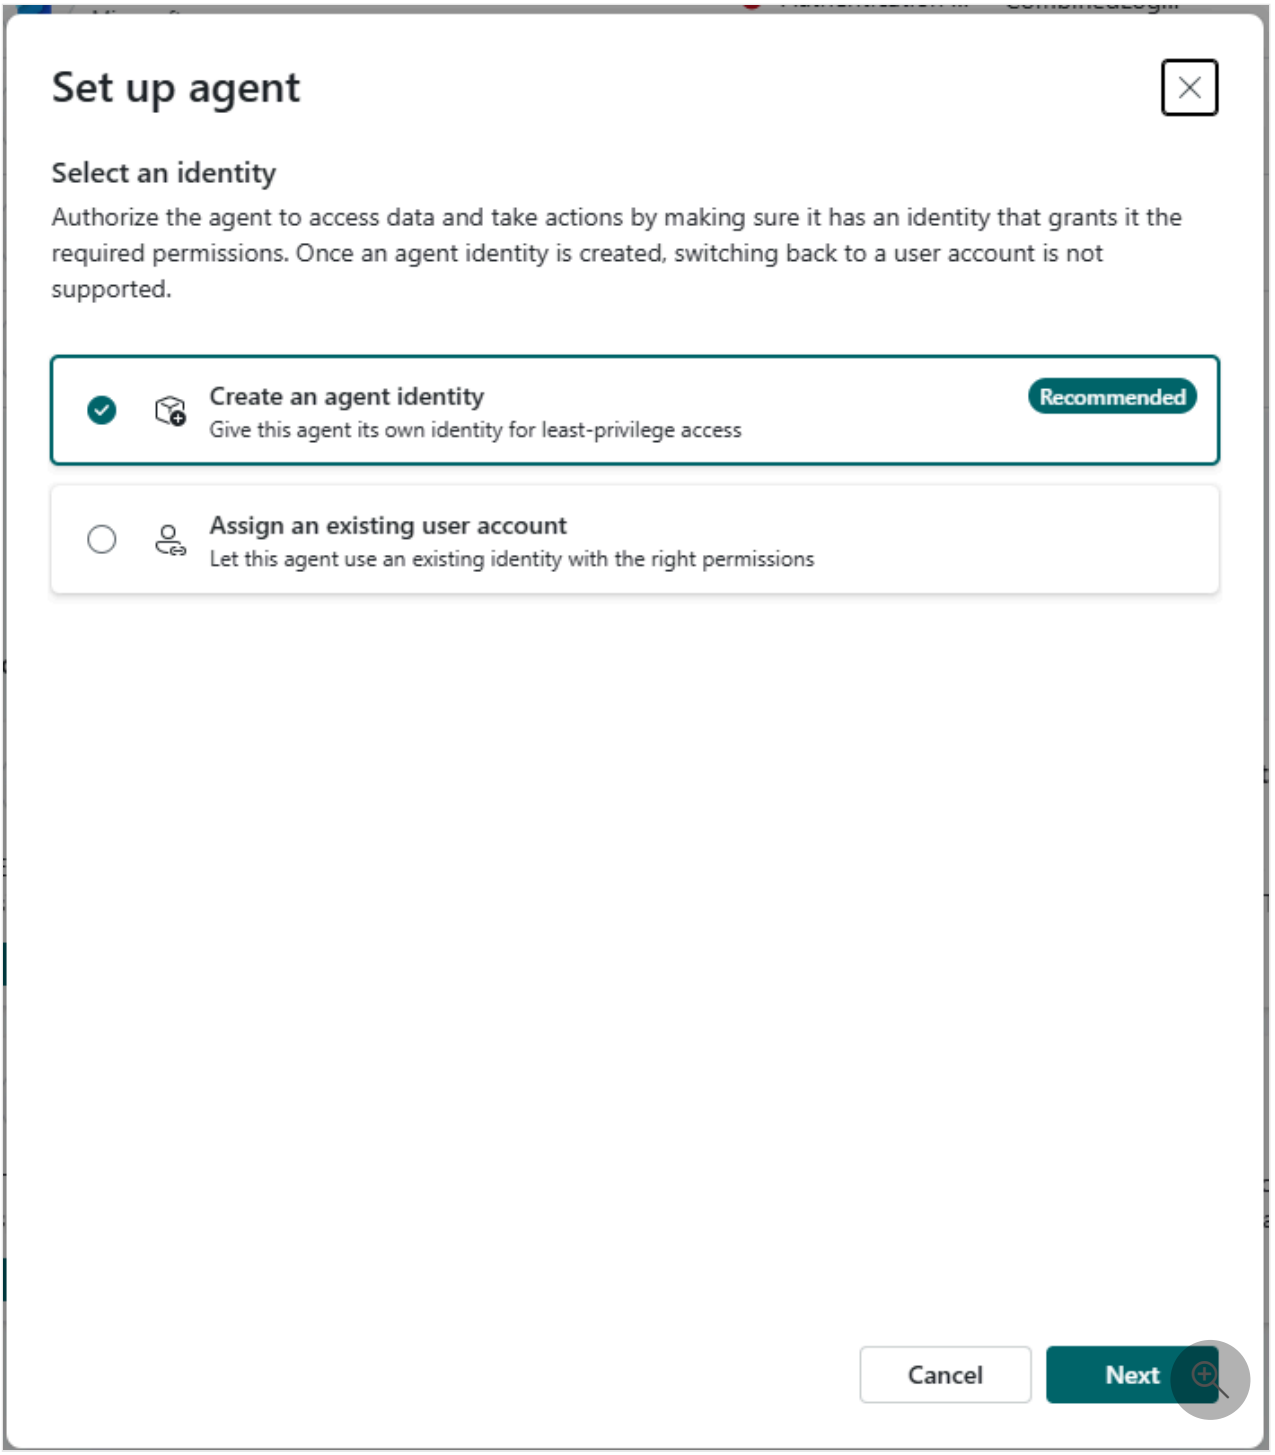

5. Follow the setup instructions and add input parameters, if necessary. Parameters vary by agent. For more information, see Agent terminology.

Expand table

Parameter	Description
Trigger	An event or condition that tells an agentic system to initiate an action or series of actions.
Permissions	The level of authorization an AI agent is granted that enables it to access specific information or carry out tasks.
Identity	The credentials that the agent uses when it runs. For some agents, you have the option of creating an agent ID or connect using an existing user account.
Plugins	A component that extends what an agent can do by providing access to capabilities in Microsoft and non-Microsoft services and public websites through APIs.
Products	The Microsoft products that are required for the agent to run.
Role-based access	The required roles or permissions required to turn on and/or run the agent.

6. After completing the setup, you'll be taken to the agent page.

7. Select **Run** to view the agent output or manage the agent.

Set up partner-built agents

> **ⓘ Important**
>
> Microsoft recommends that you use roles with the fewest permissions. Using lower permissioned accounts helps improve security for your organization. Global Administrator is a highly privileged role that should be limited to emergency scenarios when you can't use an existing role.

When setting up a partner-built agent that accesses a Microsoft tool (skill) and accesses data from a Microsoft product such as Microsoft Intune, Microsoft Entra, Microsoft Sentinel, Defender, or Defender Threat Intelligence, a Global Administrator within the tenant must approve to allow access of required permissions for the agent.

Once the approval is granted, Security Copilot owners and contributors can finish the agent setup including setting up the identity, trigger, and other values.

For agents that don't need Microsoft product permissions, approval isn't required.

For Security Copilot Owner and Contributors

1. During the agent setup process, a banner is shown stating that agent requires consent from the Global Administrator and the Set up button disabled.

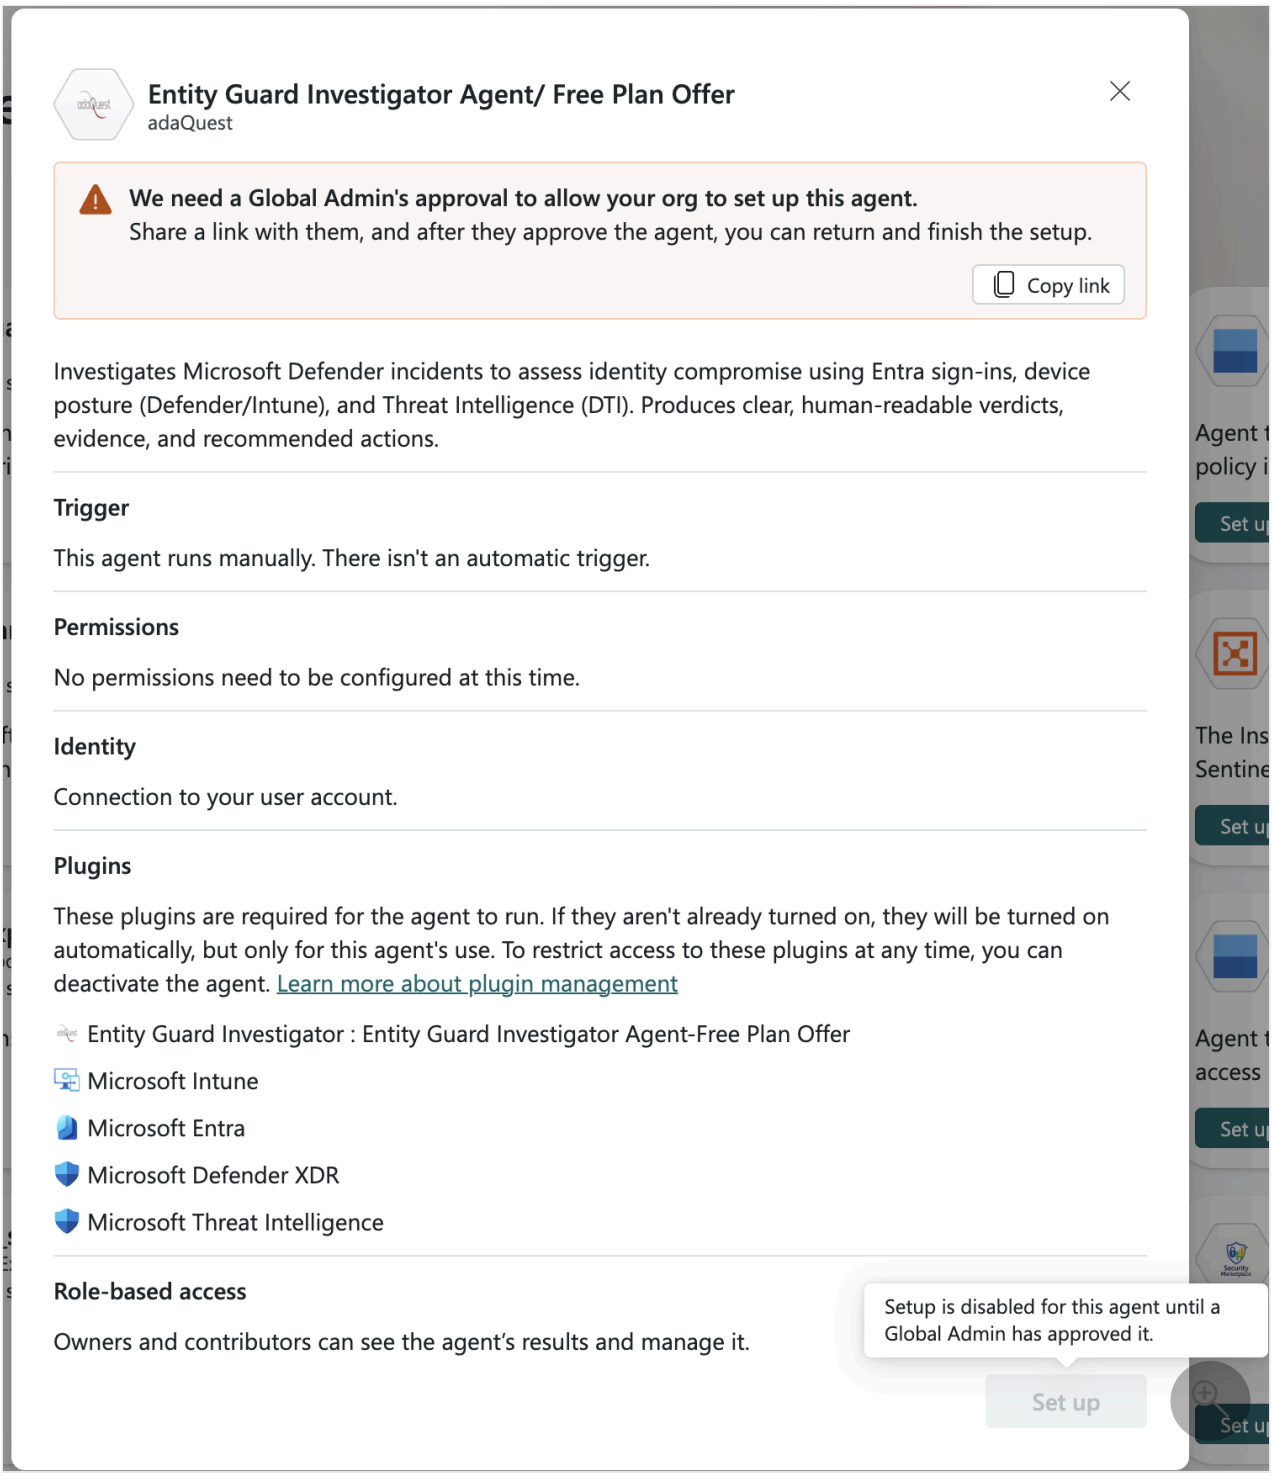

2. Select **Copy Link** to share it with the Global Administrator. After the Global Administrator completes the approval steps, you can proceed with the agent setup steps.

3. Select **Set up** to proceed to the agent details page.

4. Assign an existing user account and sign in, then select **Next**.

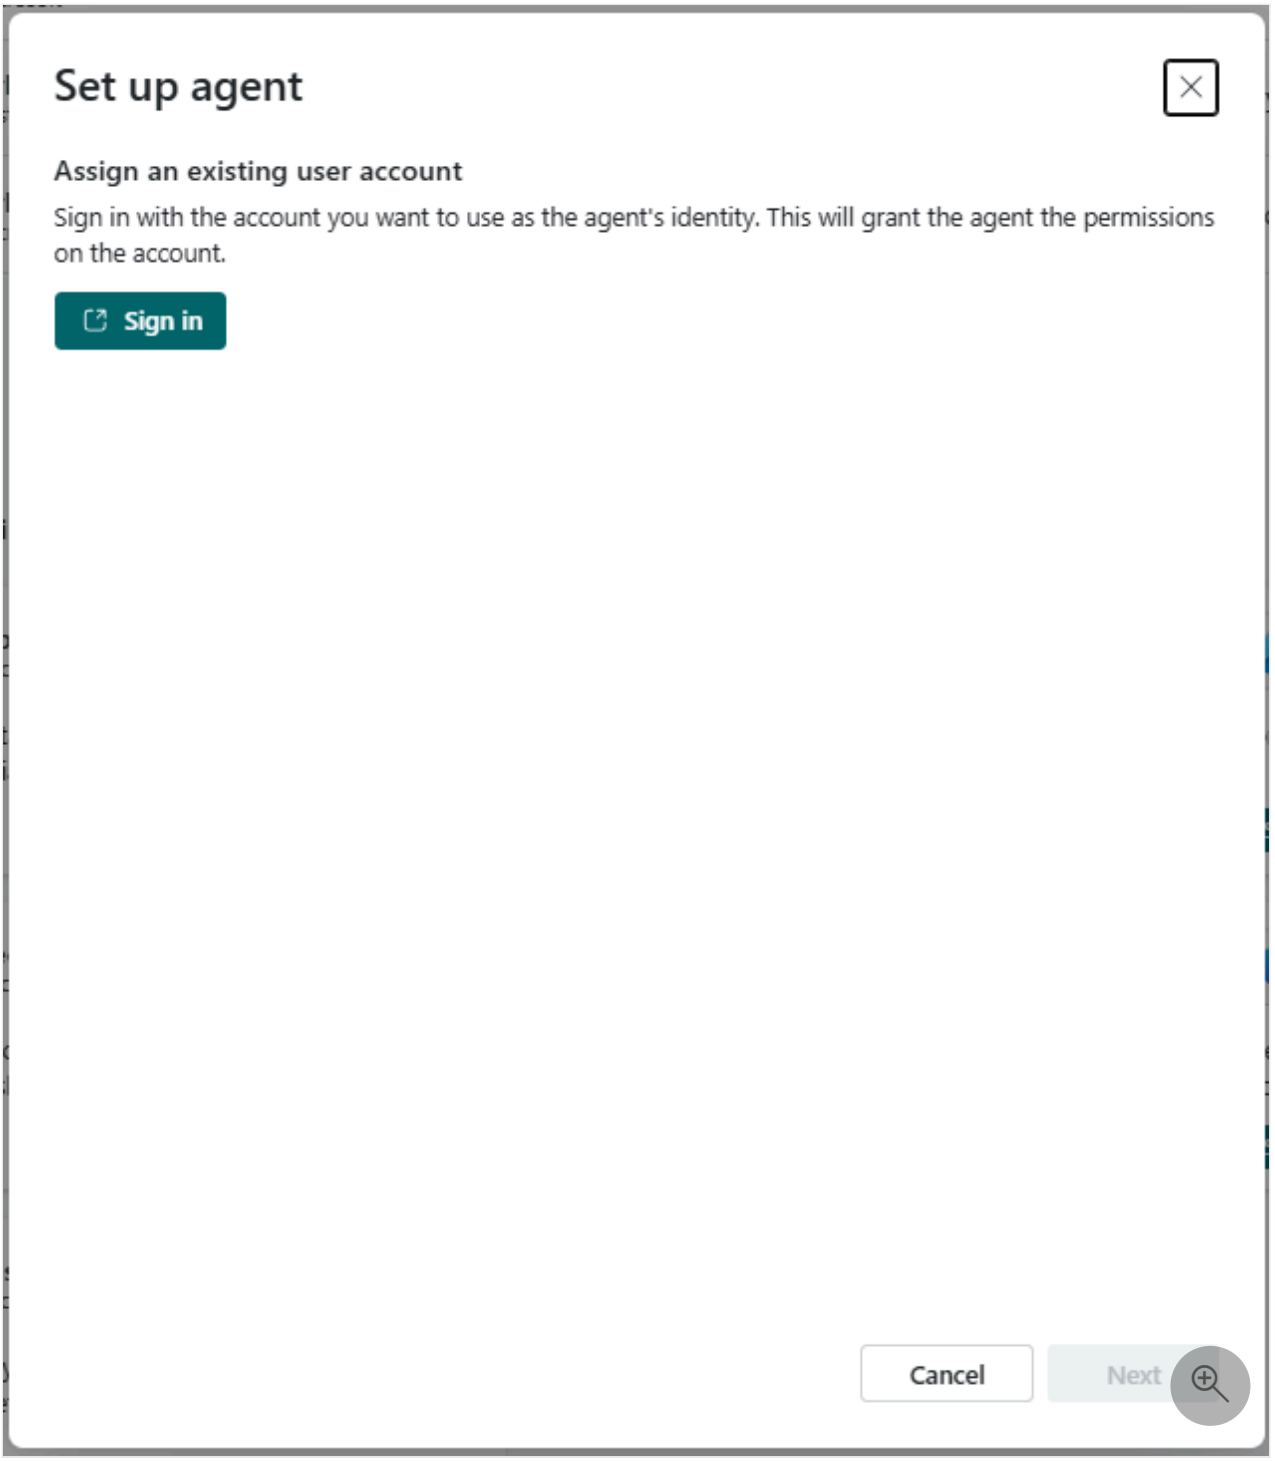

5. Follow the setup instructions and add input parameters, if necessary. Parameters vary by agent. For more information, see Agent terminology.
6. After completing the setup, you'll be taken to the agent page.
7. Select **Run** to view the agent output or manage the agent.

For Global Administrators

1. Open the link provided by the Security Copilot Owner/Contributor to approve the agent setup. You can also navigate to **Active agents** > select **Set up** on the agent.

2. The agent modal displays the agent details and the message to approve. Review the required permissions and other details about the agent such as description, trigger, and others.

3. Start the approval process by selecting **Start approval**. Follow the prompts to grant the necessary permissions.

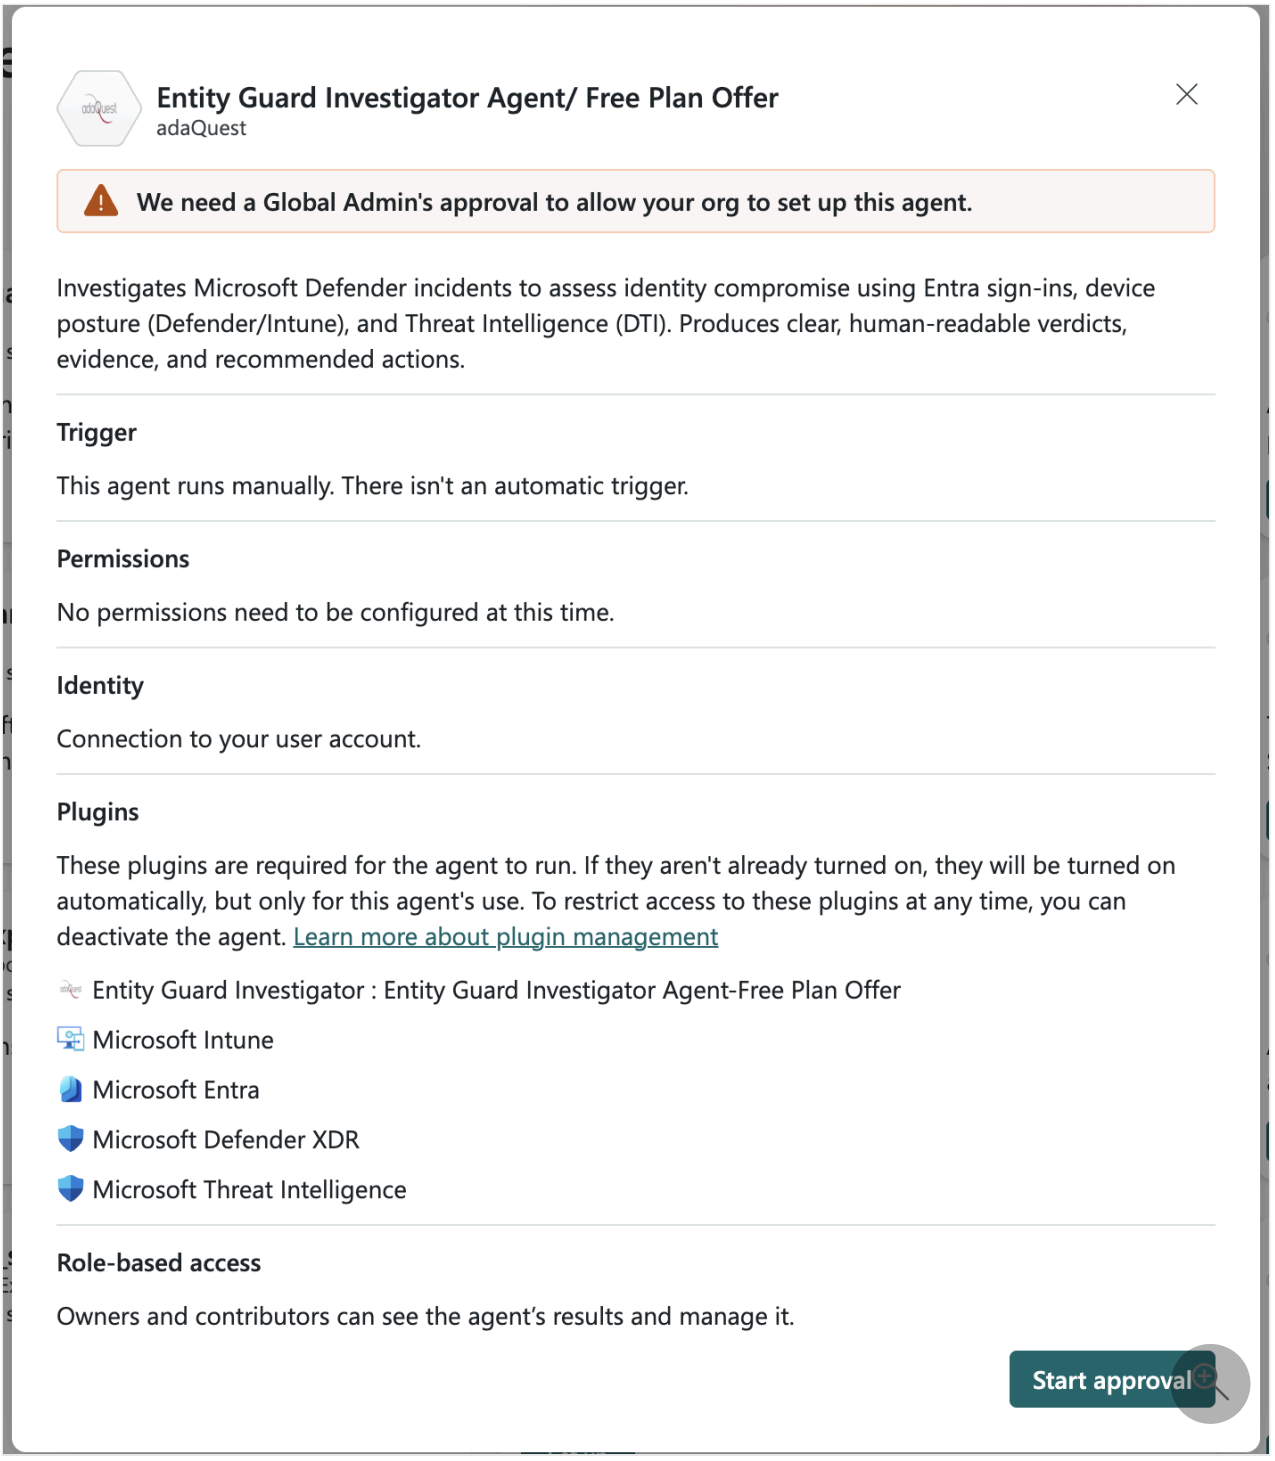

You need to approve the use of the agent for use in your org. Depending on the agent, you might also need to approve agents that use Microsoft-built skills or plugins to access data and perform actions.

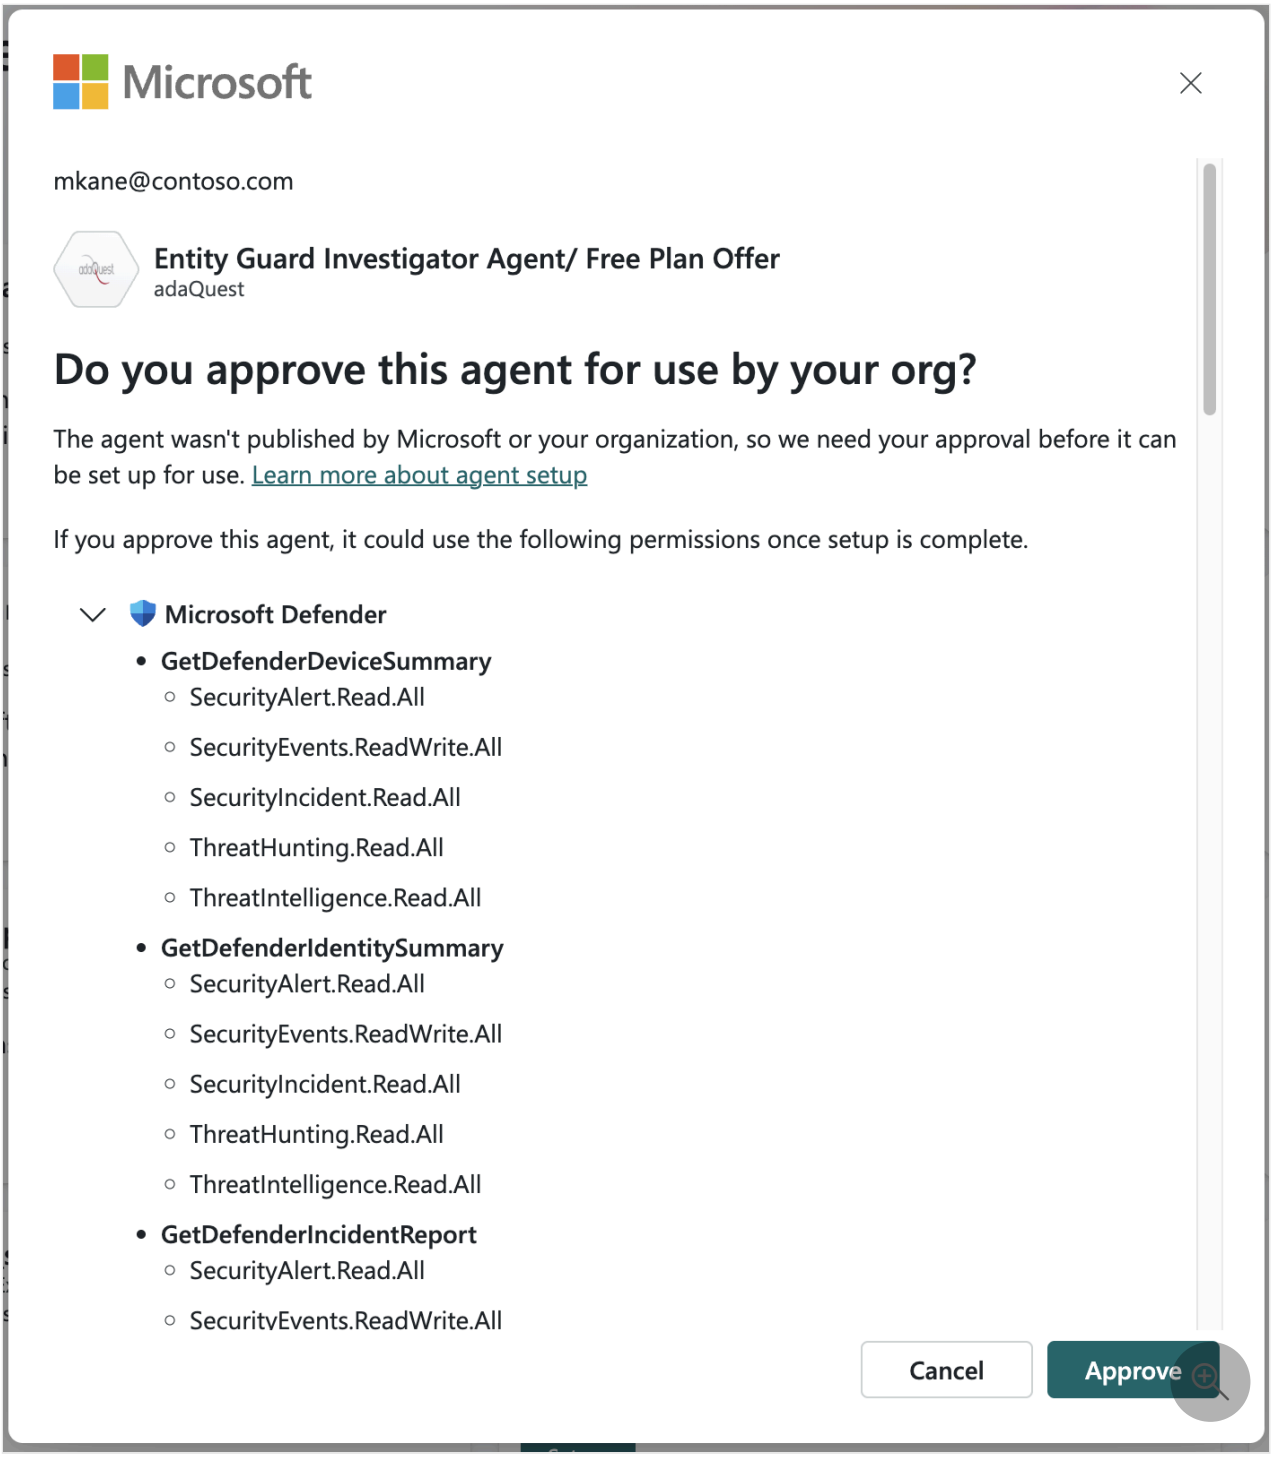

4. Once the approval is complete, the Security Copilot owner or contributor can finish the setup of the agent. You also have the option of setting up the agent for users in your org.

Provide feedback

Owners and Contributors can provide feedback to the agent to help it improve. When you provide feedback related to the agent, the agent takes your input into consideration during its operations. This feedback is stored in the agent's memory, allowing it to tailor its responses and actions based on your insights. You can also manage the agent's memory to further refine the considerations it takes when forming outputs. For more information, see Managing the agent's memory.

If your feedback is to Microsoft, your inputs may be considered in future plans. This feedback helps improve the overall product by informing the development team about user experiences and suggestions for enhancements.

1. In the agent results pane, select the feedback button.

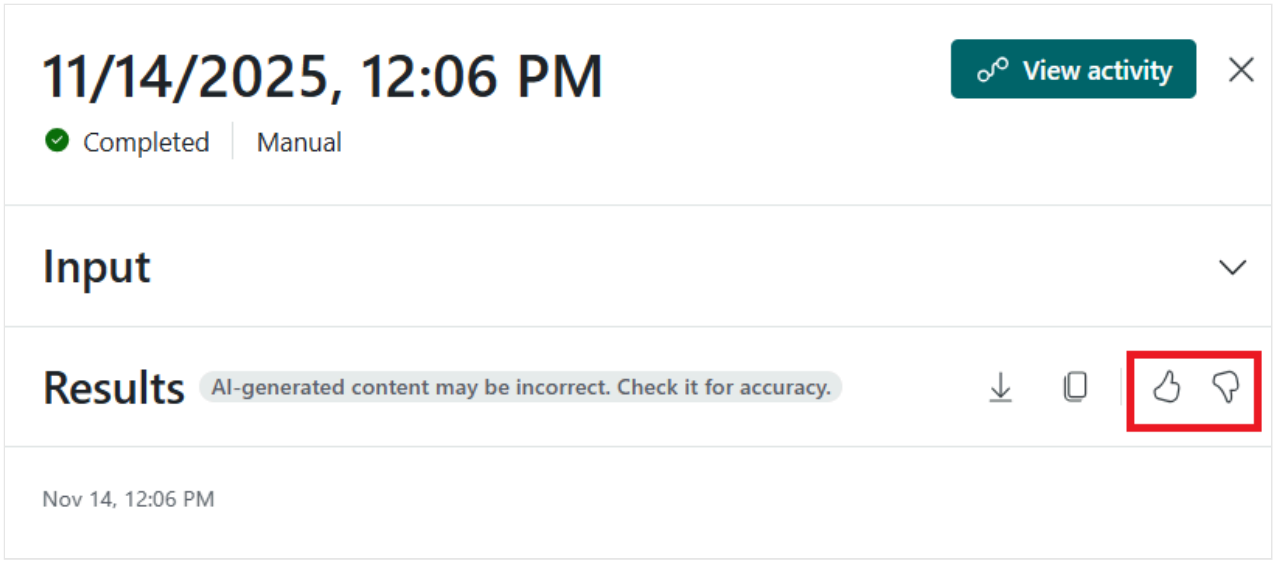

2. Select whether your feedback is for the agent or to Microsoft from the drop-down menu.

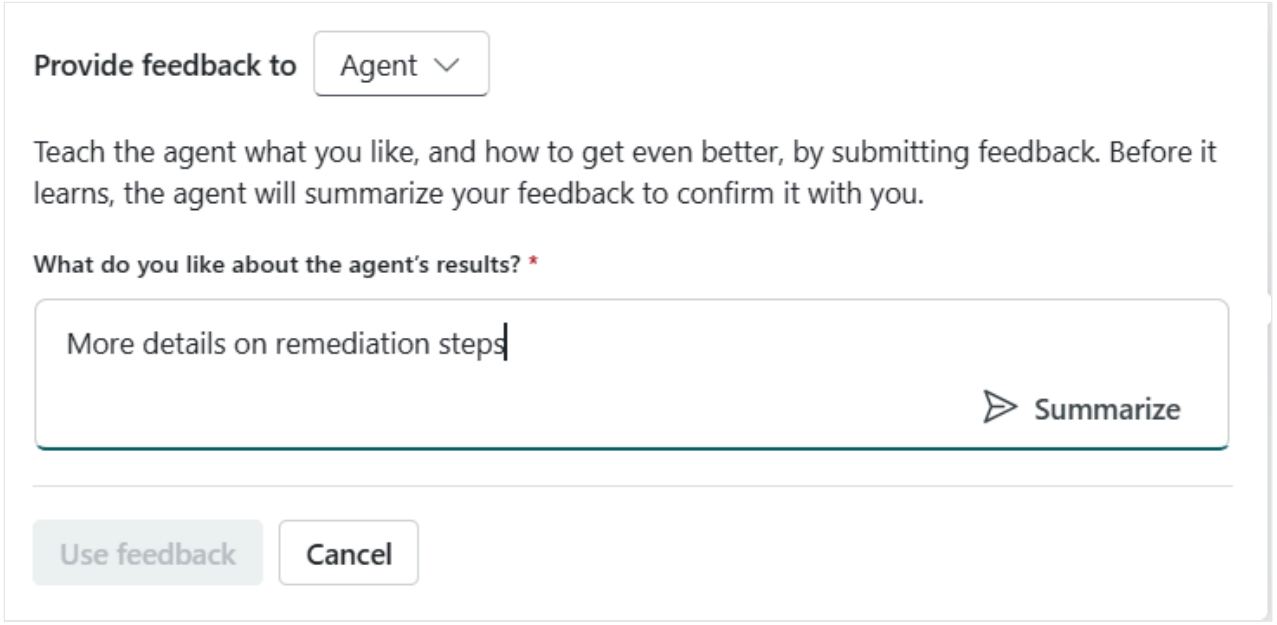

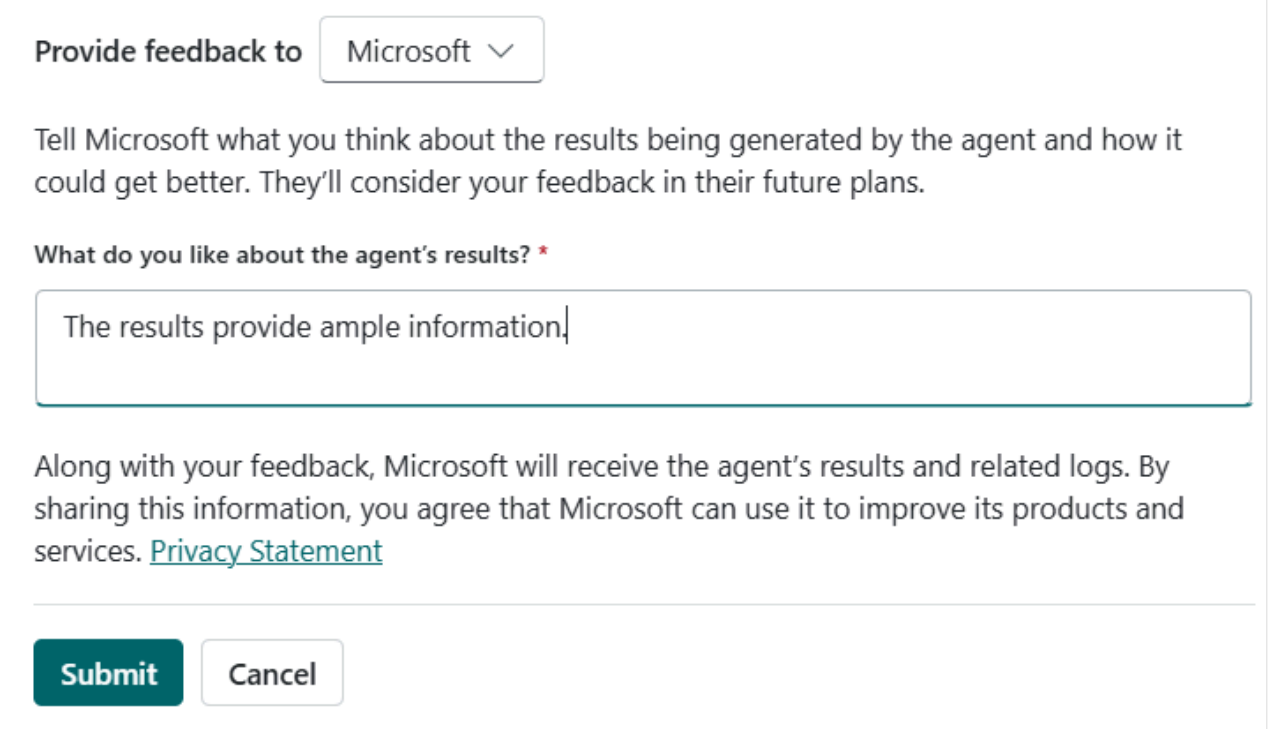

3. Provide details related to your feedback and select **Submit**.

Manage the agent's memory

After providing feedback to the agent, you have the option of tweaking its memory. Managing the agent's memory is crucial for tailoring its performance. You can control what information the agent retains and uses, ensuring that it considers relevant data when forming outputs. This management helps in refining the agent's responses and improving its overall effectiveness.

1. From the agent page, select the overflow menu (...) then **Manage memory**. Your previously entered feedback to the agent is listed.

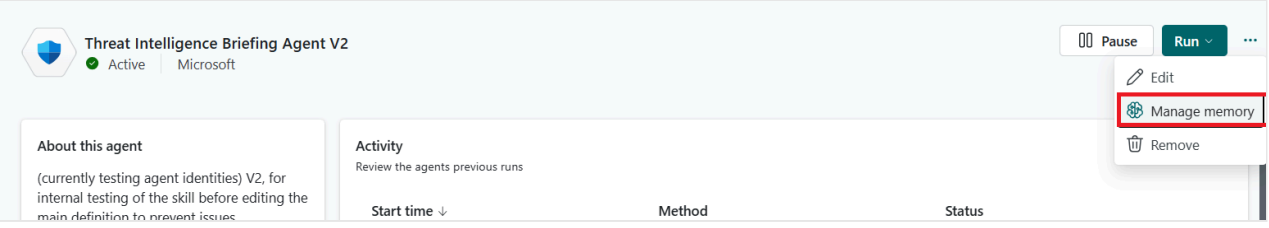

2. Select the feedback you want to review.

3. You can choose to remove the feedback from memory by selecting **Reject feedback**.

Run or pause an agent

You have the flexibility to configure the agent to run automatically or manually. If you prefer the agent to operate in the background, you can set it to run automatically. Alternatively, you

can choose to run it manually when needed. If the agent isn't required to be active, you can pause it to stop its operations temporarily.

1. Go to the agent you want to configure.
2. Select whether to pause or run the agent.

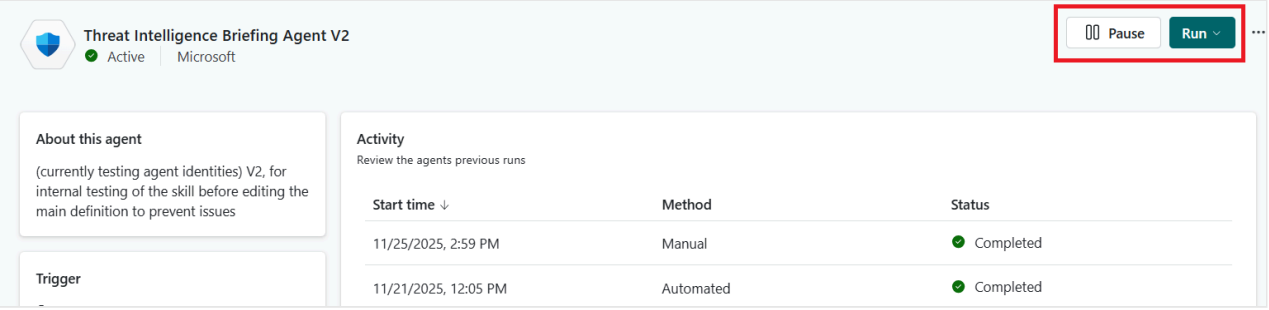

If the agent trigger is already turned on, it automatically runs as scheduled and the **On the trigger** button is greyed out. You can also choose to manually run the agent by selecting **One time**.

Edit an agent

Owners and Contributors can edit an agent to modify its configuration settings. This includes adjusting its identity or other available parameters to better suit your needs. Editing allows you to customize the agent's behavior and ensure it aligns with your specific requirements.

1. Go to the agent you want to edit.
2. Select **Edit**.

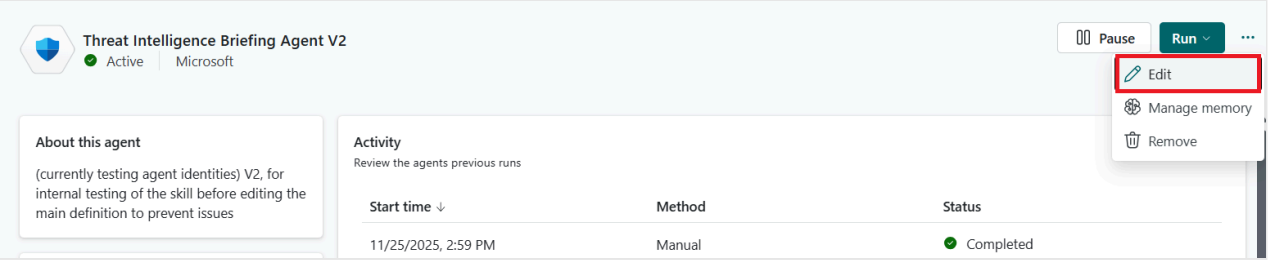

3. Update the agent details and select **Finish**.

Browse more agents in Security Store

In the Agents page, you can easily access agents in Security Store. For more information, see Security Store documentation.

Select **Browse more agents** or select **Security Store** from the menu. For more information, see Microsoft Security Store in Microsoft Security Copilot.

Last updated on 03/23/2026

Microsoft Security Copilot agent development overview

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

Microsoft Security Copilot agents are AI-powered systems designed to act on behalf of an individual, team, or business or operational domain such as Security Operations, Compliance, IT Administration, or Identity Governance to execute and orchestrate security related tasks. Agents interact with their environment to fulfill user-defined objectives by combining reasoning, planning, and executing actions to achieve specific goals.

Microsoft Security Copilot agents automate repetitive and time consuming tasks, reducing manual workloads for security professionals with precision and efficiency. Agents help in responding to security threats, user requests and system events, managing triage, investigation, and remediation of incidents. This enables organizations to operate more efficiently and focus human expertise on critical, high-impact issues.

Extend Security Copilot with custom agents

You can extend Security Copilot by building custom agents tailored to your organization's specific security and operational needs.

Security Copilot empowers developers to build, test, and publish agents to Security Copilot.

Custom agents use AI-driven automation to streamline tasks by combining the following components:

- **Tools (Skills)** - Functions or actions the agent can perform.
- **Triggers** - Conditions or events that initiate the agent.
- **Orchestrators** - Logic that determines how tasks are executed.
- **Instructions** - System-level directives that an agent must follow.
- **Feedback** - Store response in memory to guide subsequent runs.

Custom agents integrate with various security workflows, helping your organization maintain a proactive and adaptive security posture. The solution supports a workflow driven automation that responds to events or schedules, executing a series of actions based on the agent structure and guidance from Large Language Models (LLMs).

Develop custom agents

Security Copilot is a comprehensive platform to easily create, manage, and deploy a variety of agents. Agents use AI to execute security processes in the organization enabling you to achieve more. These agents can range from simple, prompt-and-response agents to more advanced, fully autonomous agents and they can be published across various channels.

The process of creating a Security Copilot agent involves a structured workflow that includes building an agent manifest, testing agent functionality, and publishing the validated agent.

The process for build, test, and publish is as follows:

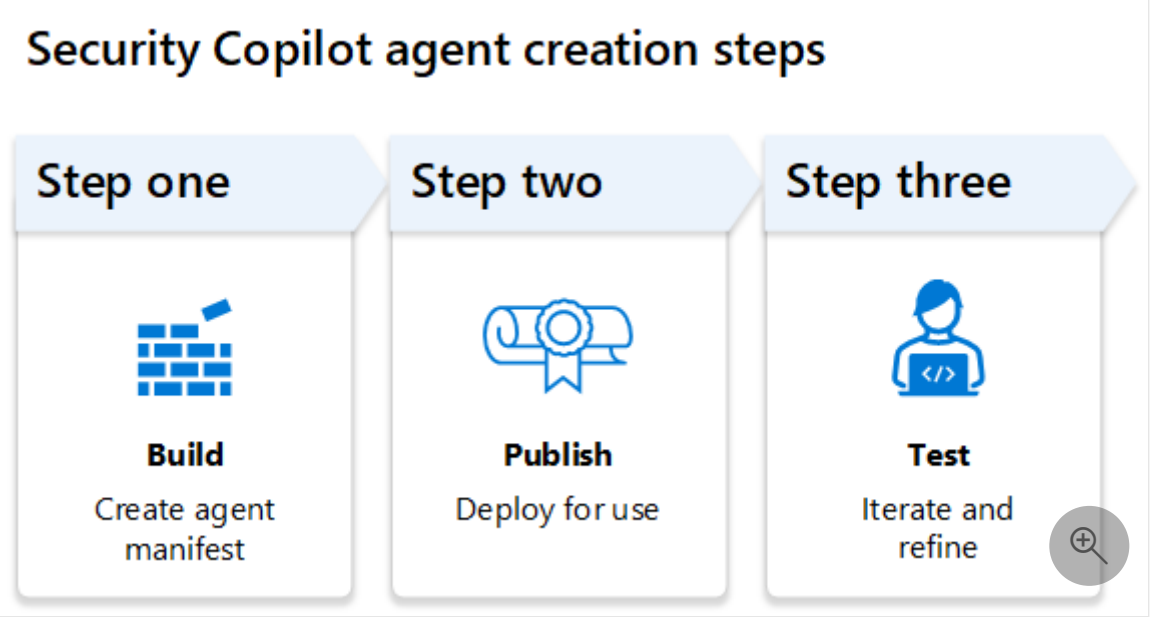

Developers can choose from these different experiences to get started to build agents. See personas.

Expand table

Feature	Types
Standalone experience	The standalone experience supports three different ways of creating an agent: 1. NL2Agent: Build agents simply by describing what you want in natural language. 2. Create an agent from scratch using agent builder: Configure agents using the agent builder interface in the Security Copilot platform.

Feature	Types
	3. Upload YAML: Build a YAML in your Integrated Development Environment (IDE) of choice and upload it to Security Copilot.
Model Context Protocol (MCP)	MCP tools: Create agents using natural language in an MCP compatible IDE using MCP tools.
Upload Plugins	YAML Manifest: Build and upload your agent manifest.

Whatever experience you choose to build your custom agent, natural language or the agent builder, the experience converges in a YAML manifest file that is deployed as an agent to Security Copilot or Security Store for partners.

- For testing the agent, see Test agent.
- For publishing the agent, see Publish agent.

Security Copilot personas

Security Copilot agents span across three personas:

Developers - Agents created by developers for use within an organization or agents created for the purpose of being available across organizations.

Administrators - Determine which agents to install in an organization, setup/initiate an agent, and review usage and success metrics of the agents.

End users (analysts or data security team; or IT admins) - Interact with the agent by using the output of their workflows and providing feedback on their workflows.

Role-Based Access Control in Security Copilot

Access to the Security Copilot platform requires assignment of specific Security Copilot roles, which are distinct from Microsoft Entra and Azure Role-Based Access Control (RBAC) roles. Security Copilot role controls what other activities you have access to on the platform, such as configuring settings, assigning permissions, and performing tasks. Security Copilot roles are defined and managed within Copilot and only grant access to Security Copilot features. For more information on authentication to other Microsoft products, see Authentication in Security Copilot.

Agent development and publishing permissions

Security Copilot uses RBAC to manage agent lifecycle operations. The following rules apply:

Copilot contributor:

- Build and test agents.
- Publish agents at user scope.

Copilot owner:

- Build and test agents.
- Publish agents at workspace scope (a scoped, tenant-bound environment where users, automations, and agents operate).

> **Note**
>
> While publishing agents, selecting **Anyone in this workspace**, makes the agents visible in the Agents page.

Agents

The Agents page in Security Copilot is the central location where you can manage and run your custom agents.

Once deployed, authorized users can access and install agents from the Agents page that a security analyst can then run.

> **Tip**
>
> Check out the **Security Copilot agents overview** to learn about discovering Microsoft and partner agents that are available in the standalone experience.

Agents can significantly enhance the efficiency of the security operations in your organization. Automating routine tasks and gathering insights through agents allows security teams to focus on high-priority investigations and strategic initiatives.

On the Agents page:

- **Ready for setup**: Displays all Microsoft and custom agents that are not yet configured. To set up an agent, see Manage Security Copilot agents.
- **Agents in use**: Indicates agents that have been configured and are ready to run.

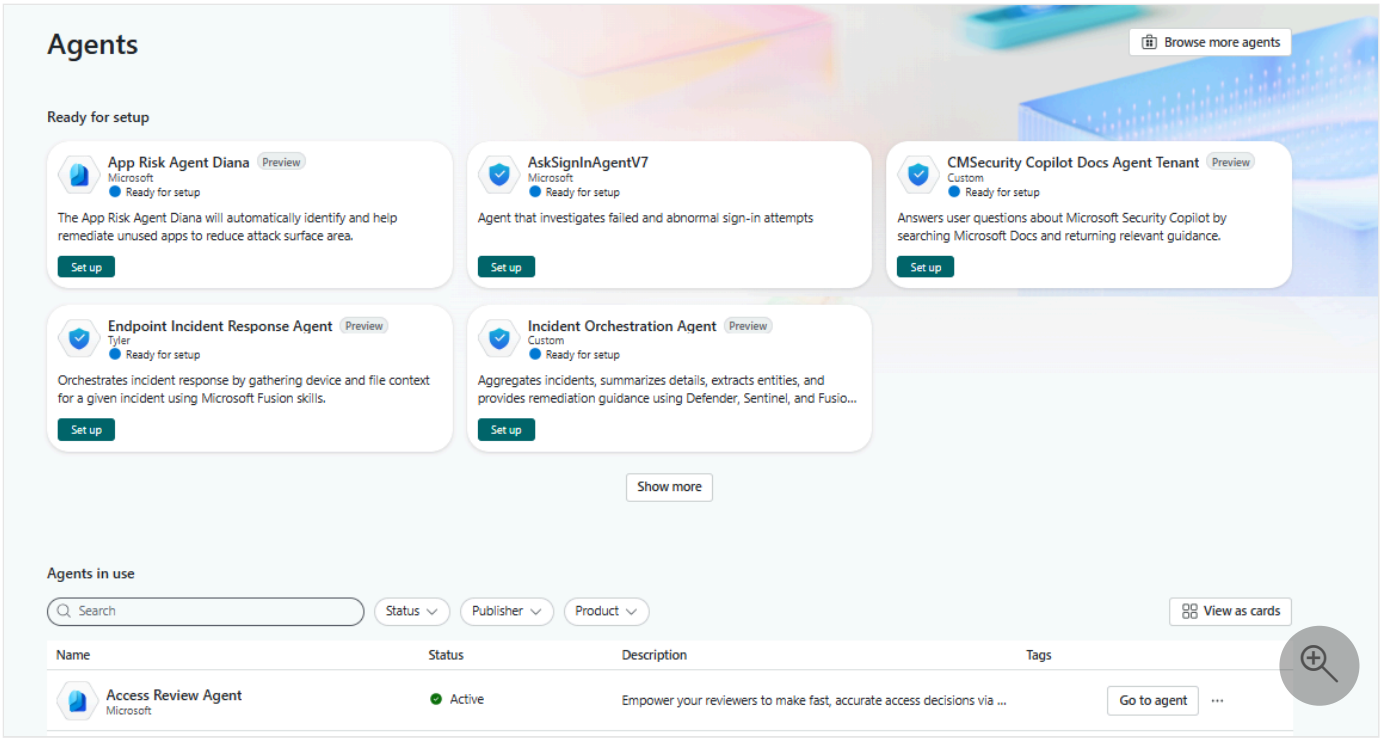

Custom agents that you develop and publish in user scope or workspace scope appear under **Build** > **My agents**. For more information, see Edit agents.

> **Note**
>
> The Security Copilot custom agent development experience is targeted at developers. For Security Copilot users such as security and compliance professionals who use Security Copilot to identify, summarize, triage, and remediate issues, see Security Copilot Admin and user library.

See also

- Planning guide
- Create an agent using MCP

Last updated on 05/26/2026

Security Copilot agent error messages

Learn about the error messages that you might encounter when using agents.

Expand table

Error message	Messages returned	Issue	Resolution
CapacityNotFound	Capacity not found. Check your capacity configuration.	Possible causes of this could include: 1. Disassociating capacity from the active workspace. 2. Not attaching a capacity to the workspace.	Ensure capacity is provisioned correctly. For more information, see Manage usage.
CapacityNotValid	Capacity isn't valid. Check your capacity configuration.	Possible causes of this could include: 1. Disassociating capacity from the active workspace. 2. Not attaching a capacity to the workspace.	Ensure capacity has been provisioned correctly. For more information, see Manage usage.
FailedToCompleteInAllottedTime	Run failed to complete in allotted time.	Agents might have gone over the maximum time required to complete the run. The current maximum time is set to 30 minutes in the system. This issue can happen in situations such as a high load or where certain processes within the agent are taking longer than expected.	Try again and if it persists, contact the agent owner.
FailedToParseFetchSkillResult	Failed to parse fetch skill result. Fetch skill must return valid JSON.	Agent returns an unexpected invalid output.	Contact the agent owner.
FailedToRetrieveFetchSkillResult	Failed to get results of fetch skill.	Agent returns an unexpected invalid output.	Contact the agent owner.

Error message	Messages returned	Issue	Resolution
FailedToStartExecution	Failed to start fetch skill.	Transient failure due to interprocess communication.	Try again and if it persists, contact support.
FailedToStartRun	Run failed to start.	Transient failure due to interprocess communication.	Try again and if it persists, contact support.
FetchSkillEvaluationFailed	Fetch skill evaluation resulted in error.	Agent underlying skill had an issue.	Contact the agent owner.
ForbiddenIdentityConfiguration	Received 403 Forbidden invoking skill. Update your identity configuration.	An internal error occurred during an interprocess communication during a prompt/evaluation using the identity associated with the agent. It wasn't able to communicate with the service due to insufficient permissions. Each tenant/service logs should be able to identify if an identity was blocked from accessing the service.	Check with your sign-in logs administrator whether any access is preventing the agent from using the identity. Other steps: 1. Go to the agent. 2. Edit the agent. 3. Try to sign in with the identity into the services that the permissions identified in the agent manifest. OR 4. Try choosing another identity that has full access to all the permissions listed in the manifest.

Error message	Messages returned	Issue	Resolution
InternalServerError	An error occurred while processing your request.	Transient failure due to interprocess communication or a service outage.	Try again and if it persists, contact support.
InvalidIdentityConfiguration	The identity configuration is invalid.	Failed to get a token from underlying services in Security Copilot. While rare, this issue can occur when one of the identity services doesn't complete the request or when it is under high load.	Try again and if it persists, contact support.
MissingIdentityConfiguration	IdentityId is missing.	A required identity wasn't configured for the agent to run.	Steps: 1. Go to the agent. 2. Edit the agent. 3. Set up an identity that has full access to all the permissions listed in the manifest.
ProcessSkillEvaluationFailed	Process skill evaluation resulted in error.	Agent underlying skill had an issue.	Contact the agent owner.
SecurityCopilotEvaluationError	An error occurred in a child evaluation.	Agent underlying skill had an issue.	Contact the agent owner.
Throttled	Skill invocation is throttled due to capacity limits. Try again later.	Possible causes of this error could include: 1. Provisioned capacity exceeded. 2. During a high load situation, the system might be throttling the process.	Review your provisioned capacity and its current usage and adjust if needed. For more information, see Manage usage.

Error message	Messages returned	Issue	Resolution
			If the provisioned capacity and usage are looking normal, retry again. If the issues persists, contact support.
TriggerNotEnabled	N/A	An agent won't run because a trigger is misconfigured.	Try to have the workspace owner reinstall the agent or get help from the agent owner
TriggerNotFound	N/A	An agent won't run because a trigger is missing.	Try to have the workspace owner reinstall the agent or get help from the agent owner.
UnauthorizedIdentityConfiguration	Unauthorized identity configuration. Update your identity configuration.	An internal error happened during an interprocess communication during a prompt/evaluation using the identity associated with the agent. It wasn't able to communicate with the service due to an invalid token. Each tenant/service logs should be able to identify if an identity was blocked from accessing the service.	Check with your sign-in logs administrator whether any access is preventing the agent from using the identity. Other steps: 1. Go to the agent. 2. Edit the agent. 3. Try to sign in with the identity into the services that the permissions identified in the agent manifest. OR

Error message	Messages returned	Issue	Resolution
			4. Try choosing another identity that has full access to all the permissions listed in the manifest.

Last updated on 01/12/2026

Plugins overview - Microsoft Security Copilot

A plugin is a collection of related tools that users can enable in Security Copilot when relevant to creating new prompting, promptbook, or agentic capabilities. A tool (also referred to as a skill) is a resource accessed outside of the agent instructions and LLM that a Security Copilot agent can invoke to help complete its outcome.

For agents, plugins extend what an agent can do by giving it access to resources outside of an agent and LLM. An example is a plugin that uses an API to connect to an external service such as reputation lookup, where underlying tool definitions for the plugin could return a reputation score, related malware, threat intelligence, etc.

Security Copilot comes with many default plugins and supports several non-Microsoft plugins. You can also extend Security Copilot's capabilities by adding or creating your own plugin. The Security Copilot platform enables developers and users to write plugins that can be invoked to perform specialized tasks.

> **Note**
>
> Products that integrate with Security Copilot as plugins need to be purchased separately.

For more information on how to develop plugins that use the OpenAI schema, see Plugins for Microsoft Copilot documentation.

Preinstalled plugins

Get familiar with the plugins Security Copilot can use to source information or take action when it's responding to your prompts. Depending on which services your organization uses, any of the plugins in the following lists might be available to you.

To find out which plugins Security Copilot can use when you interact with it, select the plugin button. Check for plugins that are toggled on in the list that opens. Security Copilot automatically uses the available plugins without any extra setup from you.

Microsoft plugins

Security Copilot uses the on-behalf-of authentication flow to provide access to other Microsoft services that your organization already has access to. For more information, see Understand authentication.

For the full list of built-in plugins, see Microsoft plugins.

Non-Microsoft plugins

For more information on non-Microsoft plugins, see Other plugins.

Websites

- Public web

Custom plugins

You can create new plugins to extend what Copilot can do by following the steps in Create new plugins.

To add and manage your custom plugins to Security Copilot, follow the steps in Manage custom plugins.

Related content

Use the following Security Copilot articles for guidance:

- Manage plugins
- Custom plugins

Last updated on 06/04/2026

Use plugins in Microsoft Security Copilot

Security Copilot includes many preinstalled plugins for Microsoft security services and other commonly used services and websites. Additionally, depending on how your administrator has configured plugins, you might be able to to add, delete, update, and manage both preinstalled and custom plugins.

> **Note**
>
> Microsoft processes and stores data from these plugins according to our Privacy and Data policies.

Watch the following video to learn more about plugins:

https://learn-video.azurefd.net/vod/player?id=f57f9efb-d800-4436-85cb-02f34811c608&locale=en-us&embedUrl=%2Fcopilot%2Fsecurity%2Fuse-plugins

Control plugins

To personalize the way Security Copilot interacts with you, there are two ways to configure the plugins on your list:

- Turn any plugin on or off
- Apply filters
- Personalize plugin settings

Enable or disable plugins

Copilot uses on-behalf-of authentication to access security-related data through active Microsoft plugins. If a plugin requires its own authentication setup though, the plugin determines the authentication instead of on-behalf-of. For more information, see Understand authentication.

Plugin sources are separated into these plugin categories:

- **Microsoft**
- **Non-Microsoft**
- **Websites**
- **Custom**

Select which plugins Security Copilot uses as a data source by turning service toggles on or off.

1. Select the **Security Copilot sources** icon.
2. Select the plugin you'd like to use by turning the toggle on or off. Custom plugins you don't have access to won't appear. Preinstalled plugins you don't have access to will indicate they're *Restricted*.

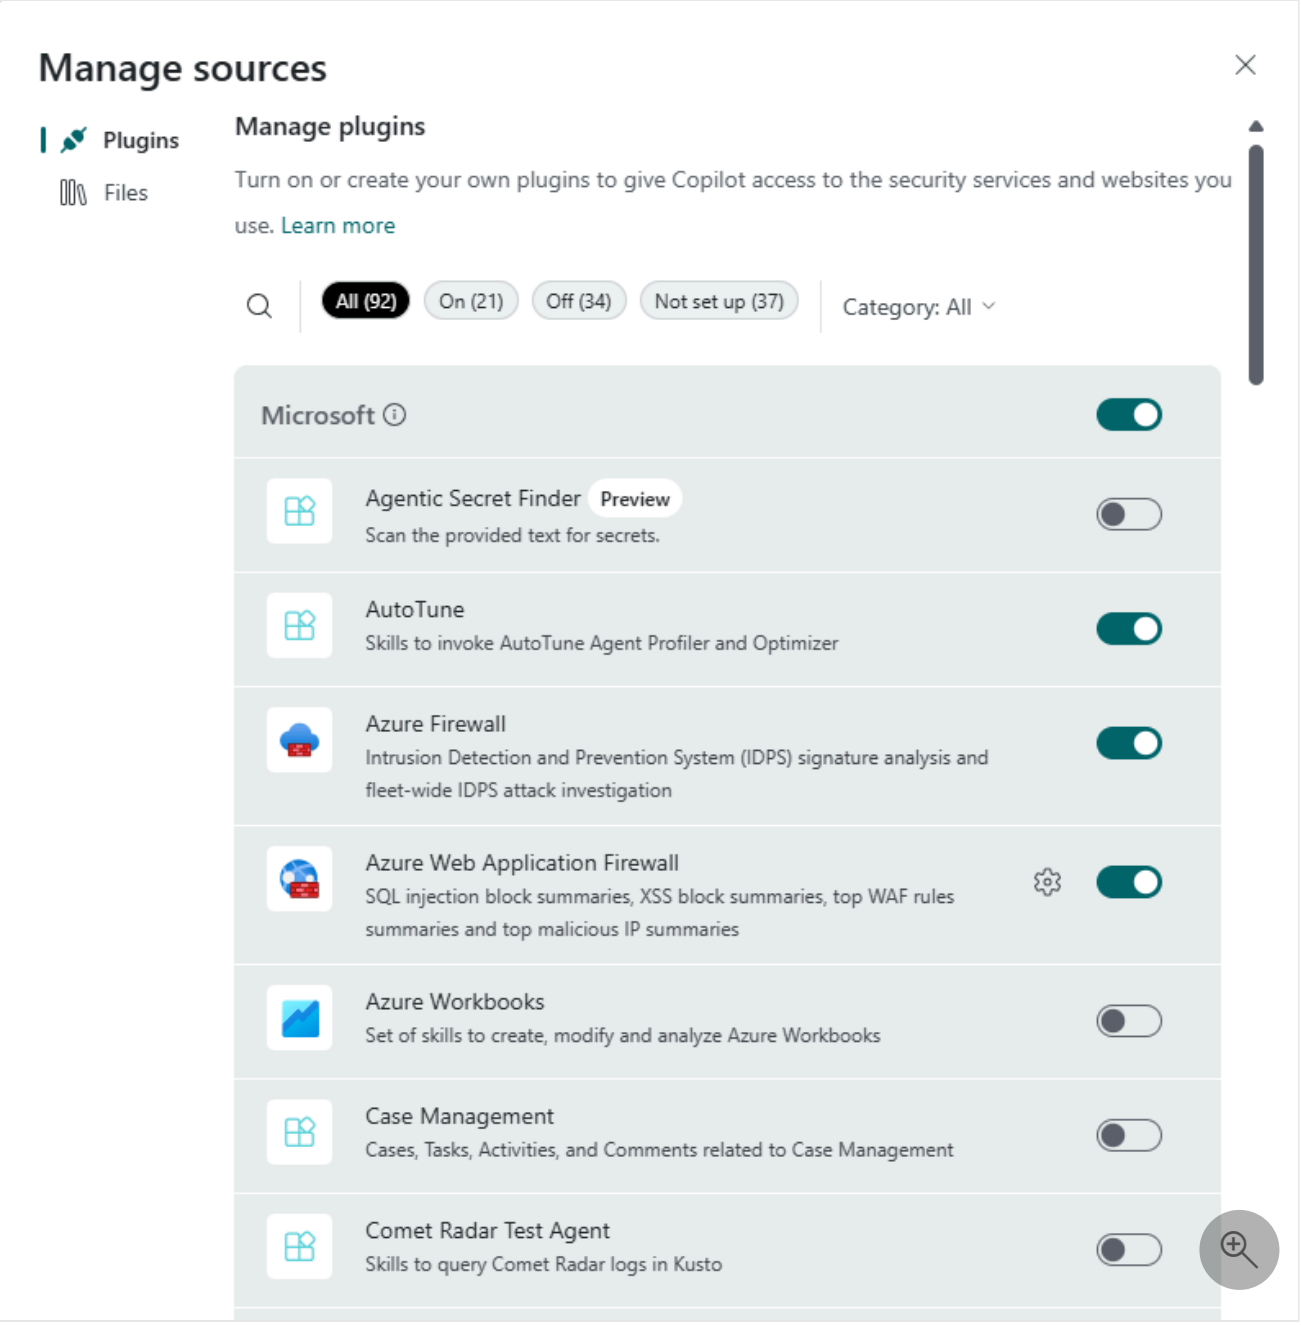

Apply filters

You can apply filters to improve discoverability and quickly find the plugin that suits your needs. By applying filters, you can narrow down plugins based on **state**, such as active or inactive, and by **category**. You can also apply multiple filters to better refine your search, helping you see the most relevant plugins for your workflow. For example, you can select all active custom plugins.

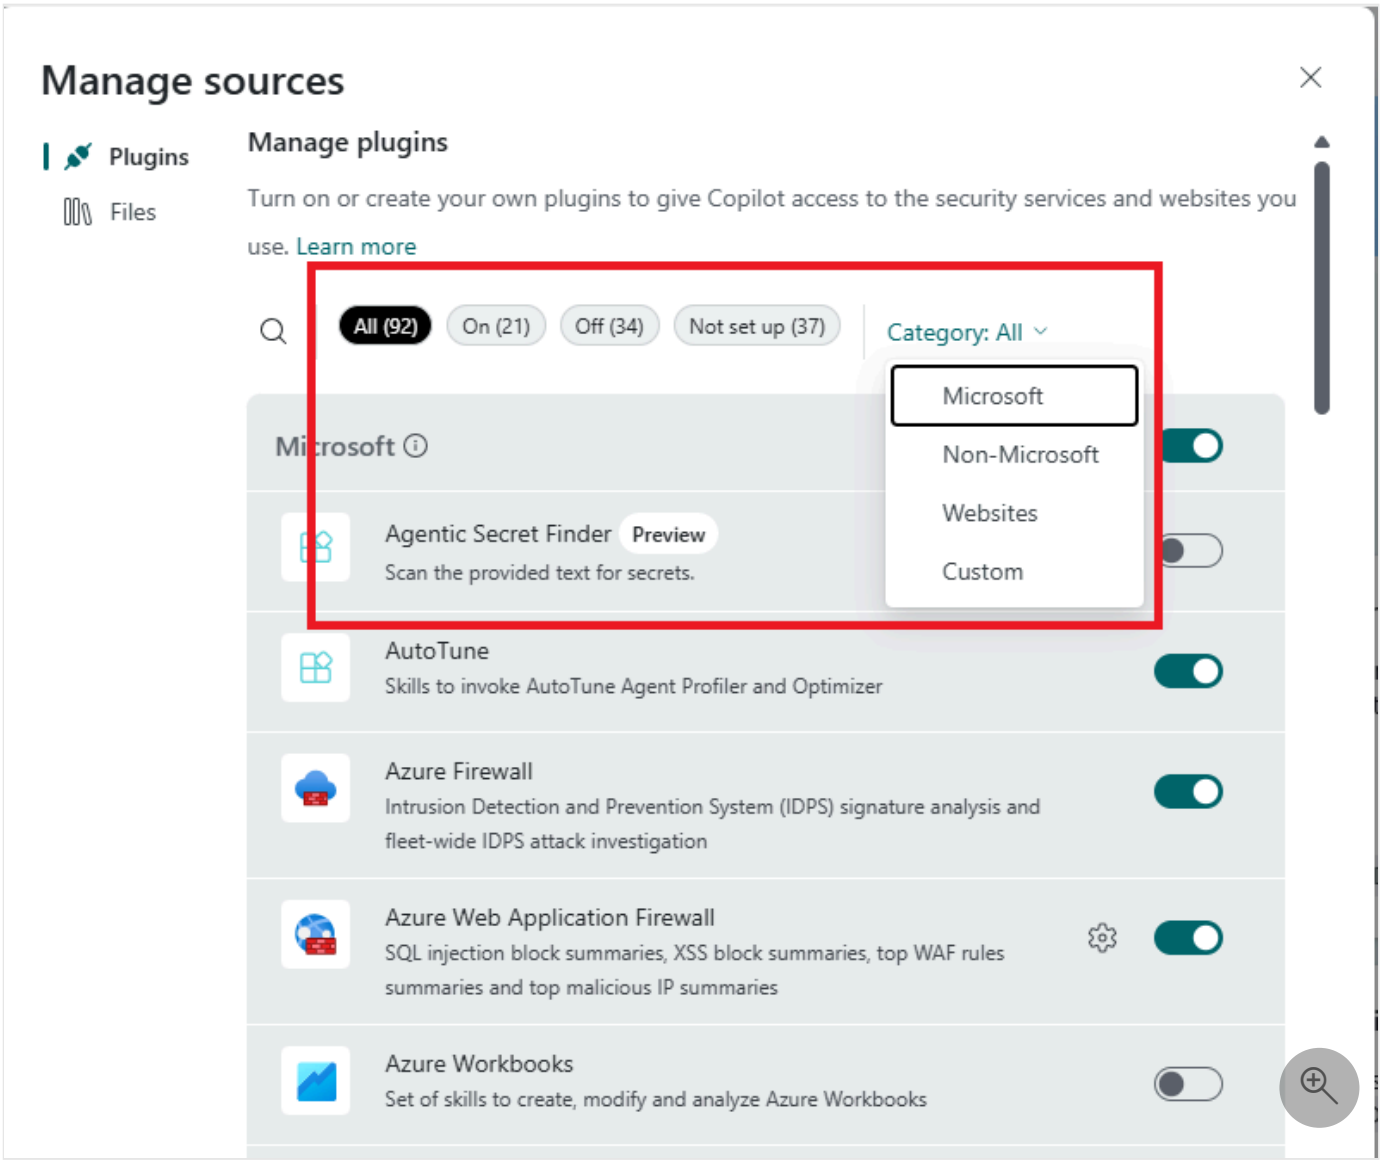

Personalize plugin settings

Personalize Security Copilot plugin experiences by configuring the plugin setting. For example, a user who wants the Microsoft Sentinel plugin to use a particular workspace that they designate as the default, personalizes the settings to specify the workspace.

Currently personalization settings are available for the following Microsoft plugins:

- Microsoft Sentinel
- Microsoft Defender External Attack Surface Management

1. Select the **Security Copilot sources** icon.

2. Select the gear icon.

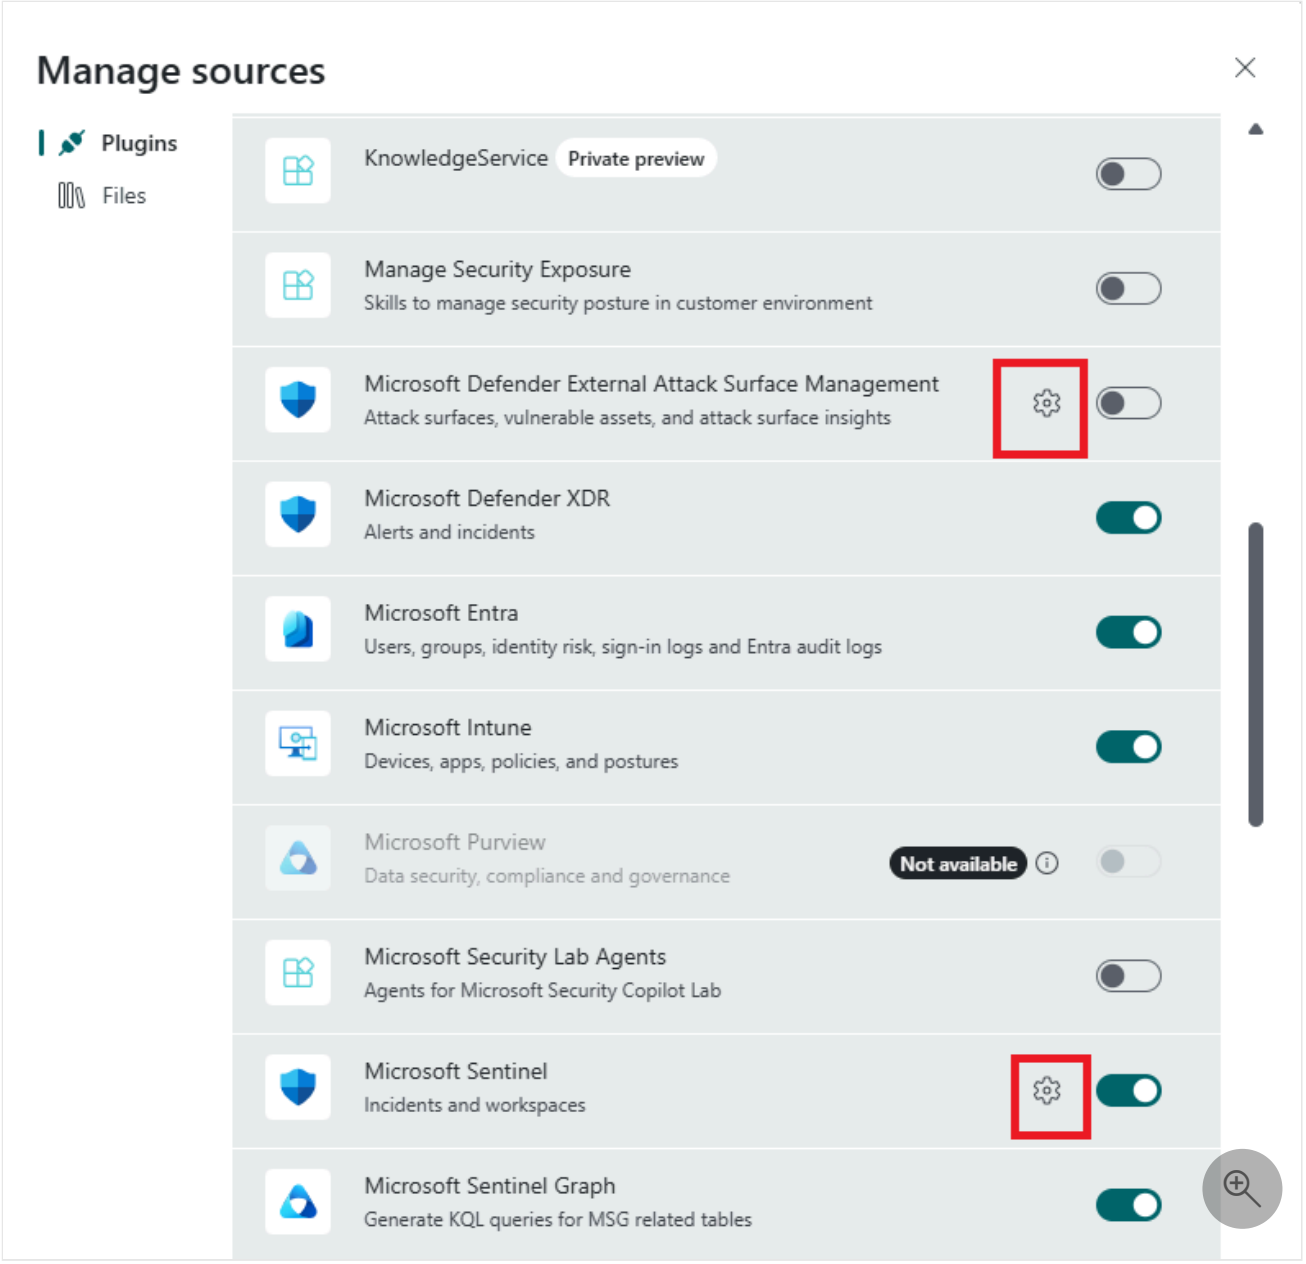

3. Provide setting details, then select **Save**.

Last updated on 06/01/2026

Azure AI Search plugin in Microsoft Security Copilot (Preview)

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

Azure AI Search is a service that allows you to search, retrieve information, and extract insights from your content at scale and securely. It is commonly used for document search, data exploration, and chat-style Copilot apps over proprietary data.

There are two primary workloads in Azure AI Search: indexing and querying. Indexing loads your content into the search service and makes it searchable, while querying allows you to search the content from your index.

The Microsoft Security Copilot integration for Azure AI Search allows you to connect to one of your Azure AI Search indexes from within Copilot, allowing you to reference your content in Copilot to generate contextualized, relevant, and specific responses.

Requirements

To use the Azure AI Search plugin in Copilot, make sure that the Azure AI Search index that you would like to connect to Copilot is configured as follows:

- The text field must be searchable.
- The title field must be filterable.
- The vector field must use **text-embedding-ada-002**.

One way you can get started with setting up an index in the required way is to use integrated vectorization. To get started with the integrated vectorization wizard, see Quickstart: Integrated vectorization (preview).

These steps are required to get your index ready for integration into Copilot.

Connect Azure AI Search to Copilot

To create the connection:

1. Open Copilot, making sure you're authenticated.

2. Select the Copilot sources icon in the prompt bar.

3. Go to **Microsoft** > **Azure AI search** and select **Set up**.

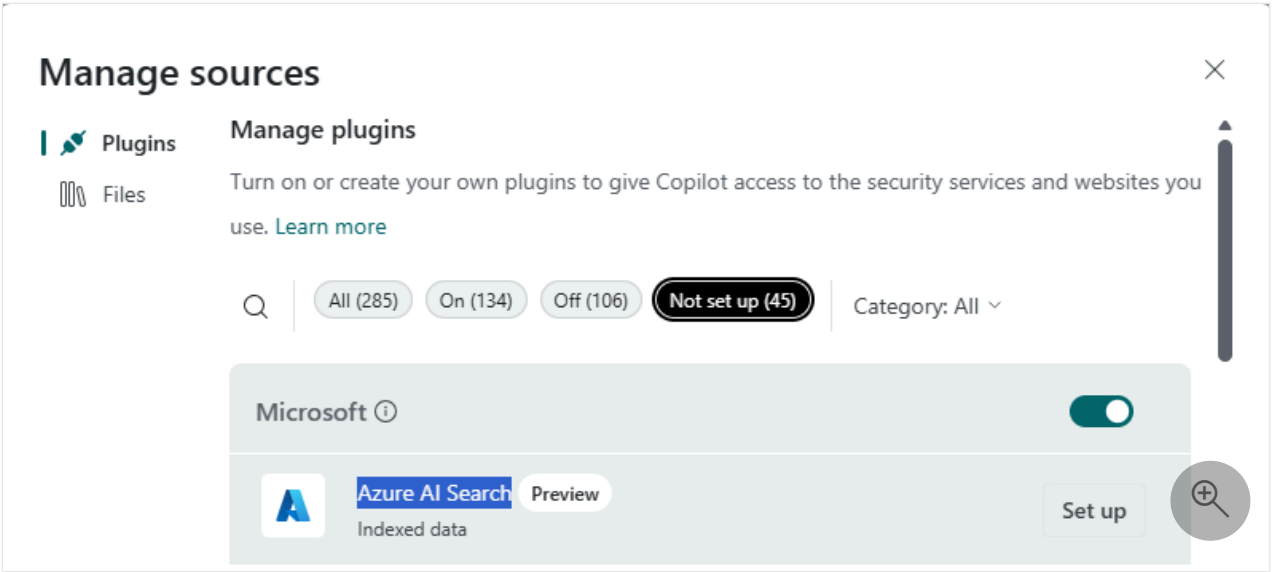

4. In **Azure AI Search settings**, enter the following information in the corresponding fields:

- **Configuration level** - specify the parameters if it applies only to you or for everyone in the workspace
- **Azure AI Search instance** – the name of your search instance
- **Index** – the name of the index within your instance to search
- **Vector** – the name of the field label in the index containing the vector of embeddings
- **Text** – the name of the field label in the index containing the text to search
- **Title** – the title of each document to display as a source
- **Value** – the API key for authentication

> ⓘ Note
>
> Using private links are not supported.

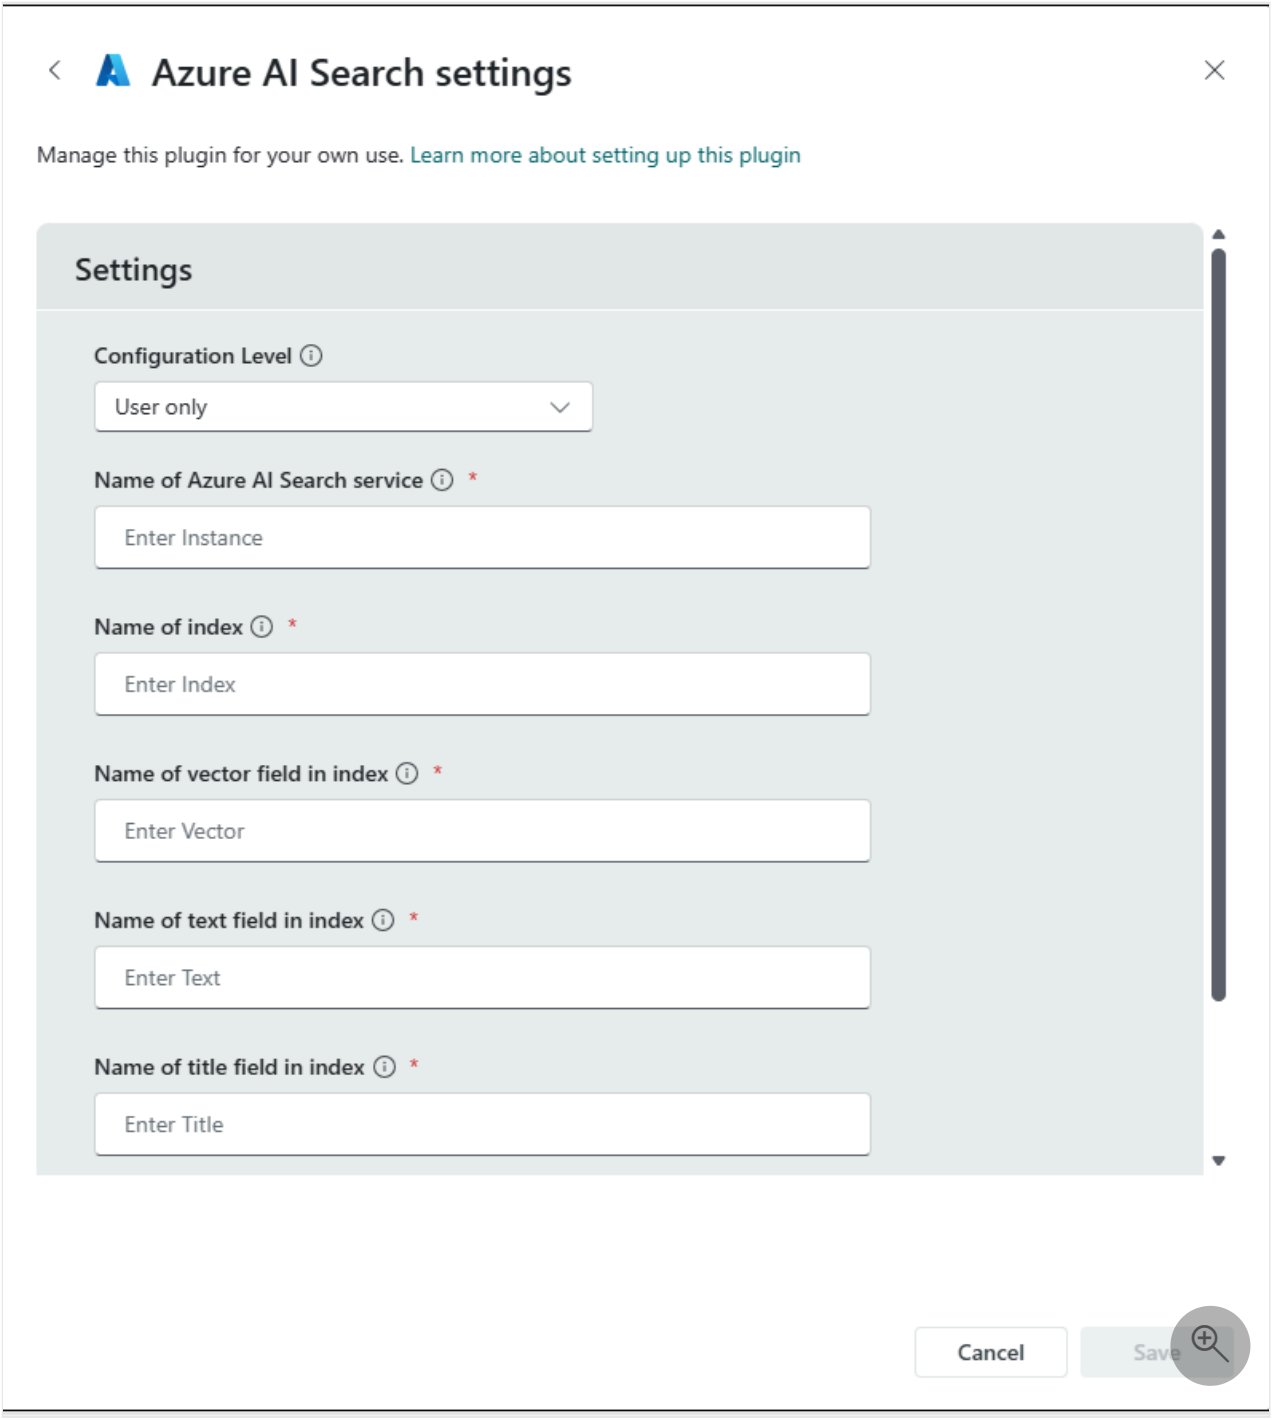

You can find the required information for the plugin setup in your Azure AI Search page in the Azure portal.

5. Select **Save** to complete the setup.

6. Close the **Azure AI Search settings** window.

> **Note**
>
> Currently, Copilot does not validate your credentials when you save your settings. If they are not correct, you will see an error later when Copilot attempts to run the Azure AI Search plugin.

This plugin isn't exclusive to knowledge base scenarios, but it's well-suited for this. The indexing and vectorizing capabilities provided by Azure AI Search are what enable apps like Copilot to search through and retrieve information across large volumes of data and content.

For more information about prompting a knowledge base in Copilot, read Prompting a knowledge base connected using Azure AI Search.

Known limitations

- Only one index at a time can be connected using the plugin. To query a different index, you need to update the plugin settings.
- The plugin supports hybrid search (keyword and vector search), so it can only support queries on indexes that were configured according to the required settings in Requirements.
- To invoke the plugin, you need to explicitly mention "Azure AI Search" in the prompt.
- Copilot displays the titles of the referenced documents from the index, but they aren't hyperlinked.

See also

- Plugins overview
- Manage plugins
- Responsible AI

Last updated on 01/12/2026

Azure Firewall integration in Microsoft Security Copilot

Security Copilot is a generative AI-powered security solution that helps increase the efficiency and capabilities of security personnel to improve security outcomes at machine speed and scale. It provides a natural language, assistive copilot experience that helps support security professionals in end-to-end scenarios such as incident response, threat hunting, intelligence gathering, and posture management. For more information about what it can do, see What is Microsoft Security Copilot?

Know before you begin

If you're new to Security Copilot, familiarize yourself with it by reading these articles:

- What is Microsoft Security Copilot?
- Microsoft Security Copilot experiences
- Get started with Microsoft Security Copilot
- Understand authentication in Microsoft Security Copilot
- Prompting in Microsoft Security Copilot

Security Copilot integration in Azure Firewall

Azure Firewall is a cloud-native and intelligent network firewall security service that provides best-in-class threat protection for your cloud workloads running in Azure. It's a fully stateful firewall as a service with built-in high availability and unrestricted cloud scalability.

The Azure Firewall integration in Security Copilot helps analysts perform detailed investigations of the malicious traffic intercepted by the IDPS feature of their firewalls across their entire fleet by using natural language questions.

You can use this integration in two different experiences:

- Security Copilot portal (standalone experience)

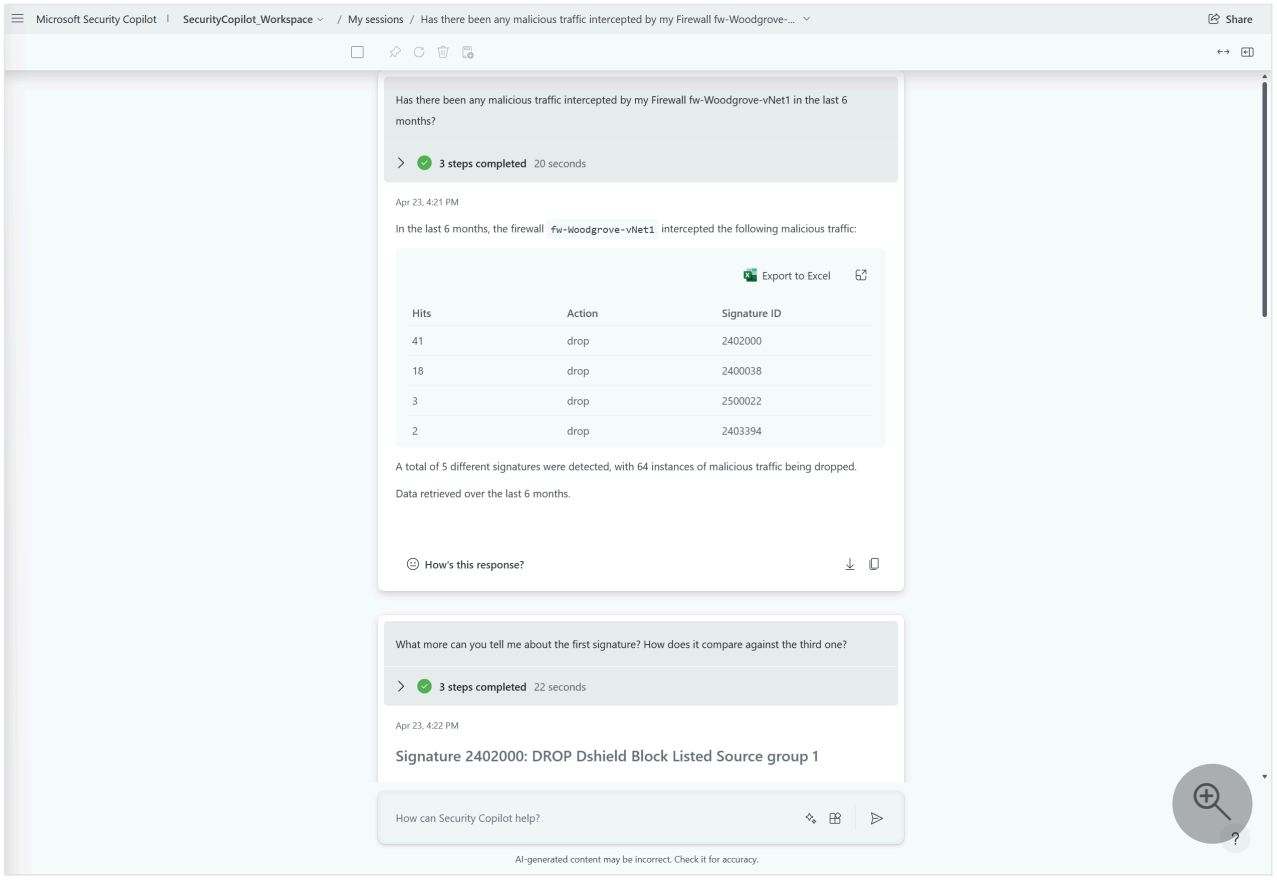

- Azure Copilot (embedded experience) in the Azure portal:

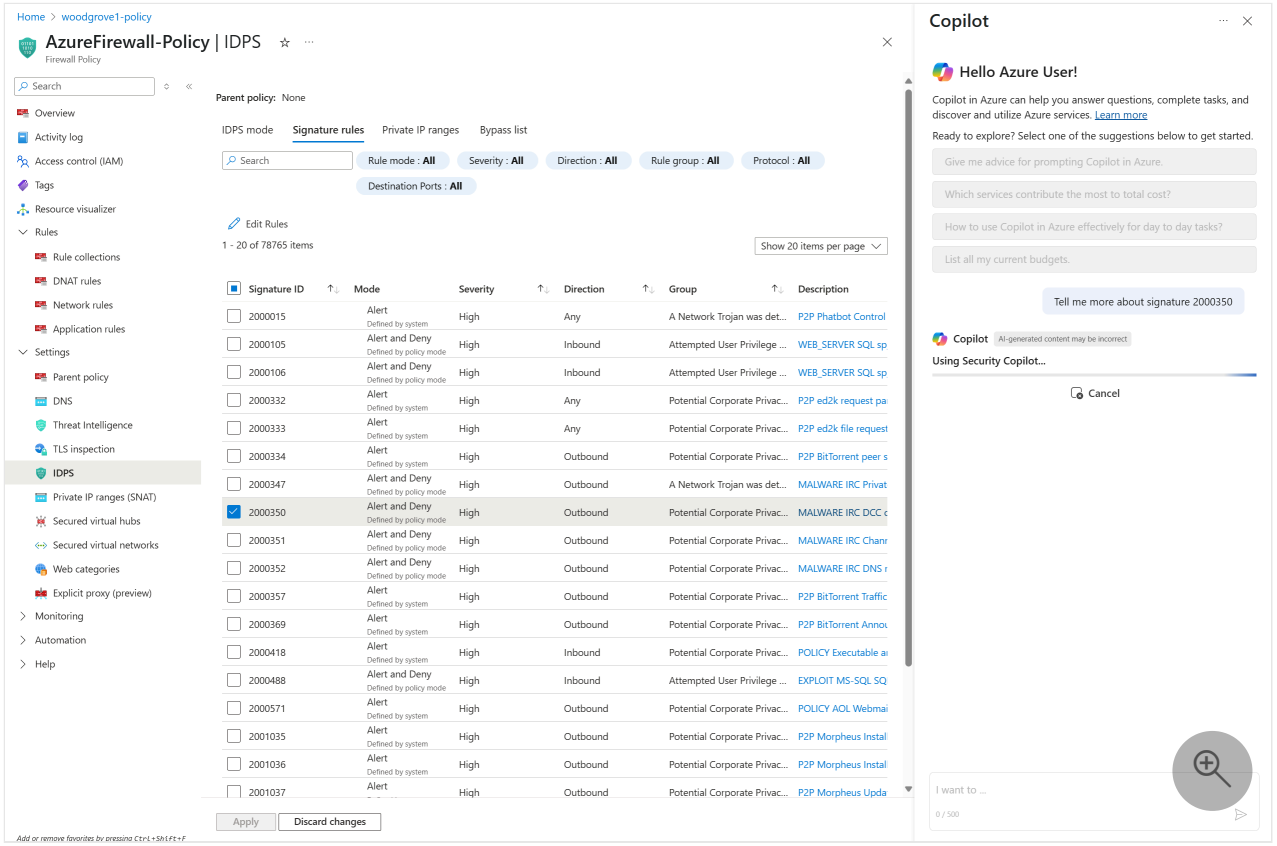

For more information, see Microsoft Security Copilot experiences and Azure Copilot capabilities.

Key features

Security Copilot has built-in system features that get data from the different plugins that are turned on.

To view the list of built-in system capabilities for Azure Firewall, use the following procedure on the Security Copilot portal:

1. In the prompt bar, select the **Prompts** icon.
2. Select **See all system capabilities**.
3. The **Azure Firewall** section lists all the available capabilities that you can use.

Enable the Azure Firewall integration in Security Copilot

1. Ensure your Azure Firewall is configured correctly:
    - Azure Firewall Structured Logs – to use Azure Firewalls with Security Copilot, configure them with resource specific structured logs for IDPS and send these logs to a Log Analytics workspace.
    - Role Based Access Control for Azure Firewall – the users using the Azure Firewall plugin in Security Copilot must have the appropriate Azure Role-based access control roles to access the Firewall and associated Log Analytics workspaces.
2. Go to Security Copilot and sign in with your credentials.
3. Make sure that the Azure Firewall plugin is turned on. In the prompt bar, select the **Sources** icon. In the **Manage sources** pop-up window that appears, confirm that the **Azure Firewall** toggle is turned on. Then, close the window. No other configuration is necessary. As long as structured logs are sent to a Log Analytics workspace and you have the right Role-based access control permissions, Copilot finds the data it needs to answer your questions.

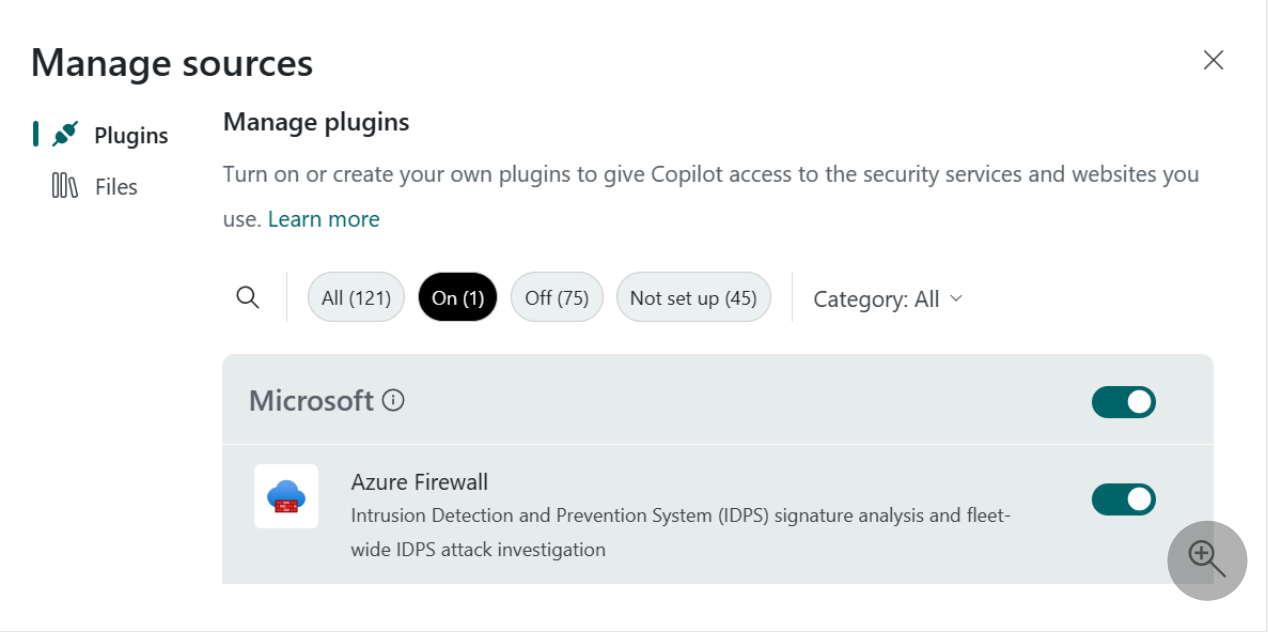

4. Enter your prompt in the prompt bar on either the Security Copilot portal or via the Azure Copilot experience in the Azure portal.

> **Important**
>
> Use of Azure Copilot to query Azure Firewall is included with Security Copilot and requires security compute units (SCUs). You can deploy SCUs and increase or decrease them at any time. For more information on SCUs, see Get started with Microsoft Security Copilot. If you don't have Security Copilot properly configured but ask a question relevant to the Azure Firewall capabilities through the Azure Copilot experience, you see an error message.

Sample Azure Firewall prompts

To get information from Azure Firewall, use prompts. This section lists the ones that work best today. It's continuously updated as new capabilities are launched.

Retrieve the top IDPS signature hits for a given Azure Firewall

Get **log information** about the traffic intercepted by the IDPS feature instead of constructing KQL queries manually.

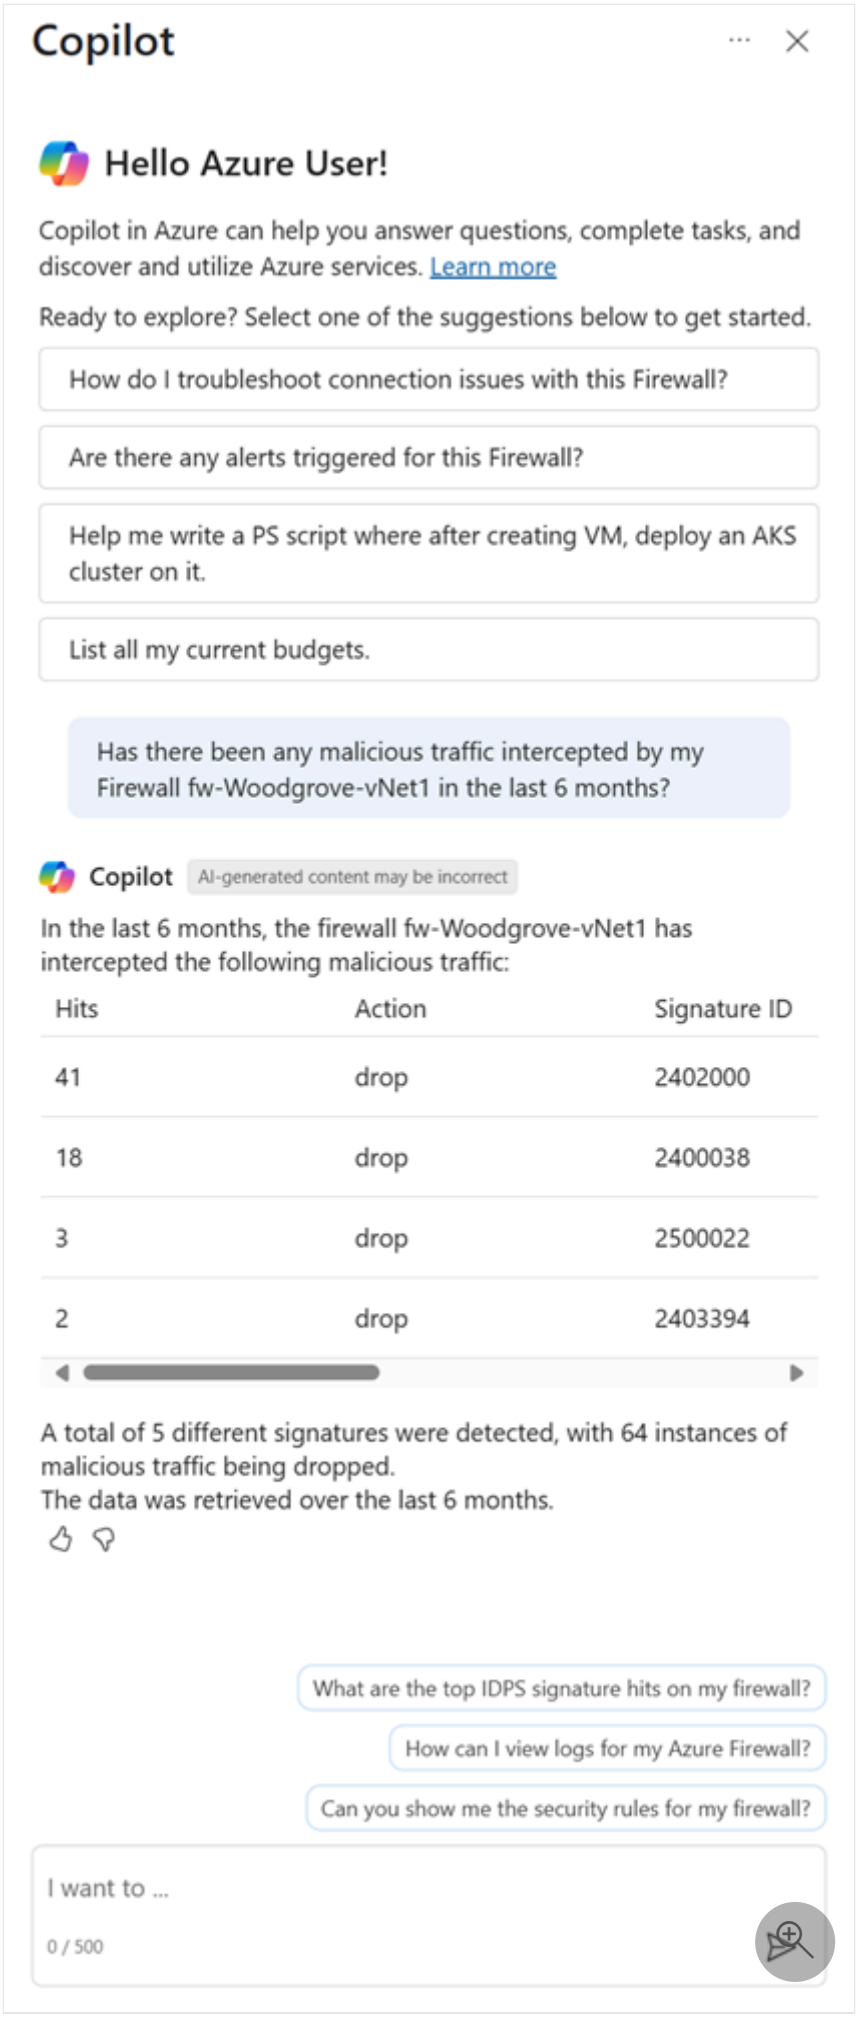
Copilot
Hello Azure User!
Copilot in Azure can help you answer questions, complete tasks, and discover and utilize Azure services. Learn more
Ready to explore? Select one of the suggestions below to get started.
How do I troubleshoot connection issues with this Firewall?
Are there any alerts triggered for this Firewall?
Help me write a PS script where after creating VM, deploy an AKS cluster on it.
List all my current budgets.
Has there been any malicious traffic intercepted by my Firewall fw-Woodgrove-vNet1 in the last 6 months?
Copilot
AI-generated content may be incorrect
In the last 6 months, the firewall fw-Woodgrove-vNet1 has intercepted the following malicious traffic:
Hits
Action
Signature ID
41
drop
2402000
18
drop
2400038
3
drop
2500022
2
drop
2403394
A total of 5 different signatures were detected, with 64 instances of malicious traffic being dropped.
The data was retrieved over the last 6 months.
What are the top IDPS signature hits on my firewall?
How can I view logs for my Azure Firewall?
Can you show me the security rules for my firewall?
I want to ...
0 / 500

Sample prompts:

- Has my Firewall `<Firewall name>` intercepted any malicious traffic?
- What are the top 20 IDPS hits from the last seven days for Firewall `<Firewall name>` in resource group `<resource group name>`?
- Show me in tabular form the top 50 attacks that targeted Firewall `<Firewall name>` in subscription `<subscription name>` in the past month.

Enrich the threat profile of an IDPS signature beyond log information

Get **additional details** to enrich the threat information and profile of an IDPS signature instead of compiling it yourself manually.

Copilot

What more can you tell me about the first signature? How does it compare against the third one?

Copilot AI-generated content may be incorrect

Signature 2402000: DROP Dshield Block Listed Source group 1
Description:

- **Category:** Misc Attack
- **Severity:** 2
- **Details:** This signature identifies traffic from sources that are blocklisted by Dshield, a community-driven project that tracks and reports on malicious activity. The traffic is dropped to prevent potential threats from reaching the network.

Signature 2500022: COMPROMISED Known Compromised or Hostile Host Traffic group 12
Description:

- **Category:** Misc Attack
- **Severity:** 2
- **Details:** This signature detects traffic from hosts that are known to be compromised or hostile. The traffic is dropped to mitigate the risk of further compromise or malicious activity within the network.

Comparison

- **Category:** Both signatures fall under the "Misc Attack" category.
- **Severity:** Both signatures have a severity level of 2.
- **Purpose:** Signature 2402000 focuses on blocking traffic from sources listed by Dshield, while Signature 2500022 targets traffic from known compromised or hostile hosts.
- **Action:** Both signatures result in the traffic being dropped to protect the network.

In summary, both signatures are designed to block potentially harmful traffic, but they target different sources of threats. Signature 2402000 deals with blocklisted sources, whereas Signature 2500022 addresses known compromised or hostile hosts.

References

- MDTI: SocGholish

What other signatures have been detected by my firewall?

How can I improve my firewall's security settings?

Can you provide more details on the second signature?

I want to ...

0 / 500

Sample prompts:

- Explain why IDPS flagged the top hit as high severity and the fifth hit as low severity.
- What can you tell me about this attack? What are the other attacks this attacker is known for?
- I see that the third signature ID is associated with CVE `<CVE number\>`. Tell me more about this CVE.

> **Note**
>
> The Microsoft Threat Intelligence plugin is another source that Security Copilot might use to provide threat intelligence for IDPS signatures.

Look for a given IDPS signature across your tenant, subscriptions, or resource groups

Perform a **fleet-wide search** (over any scope) for a threat across all your firewalls instead of searching for the threat manually.

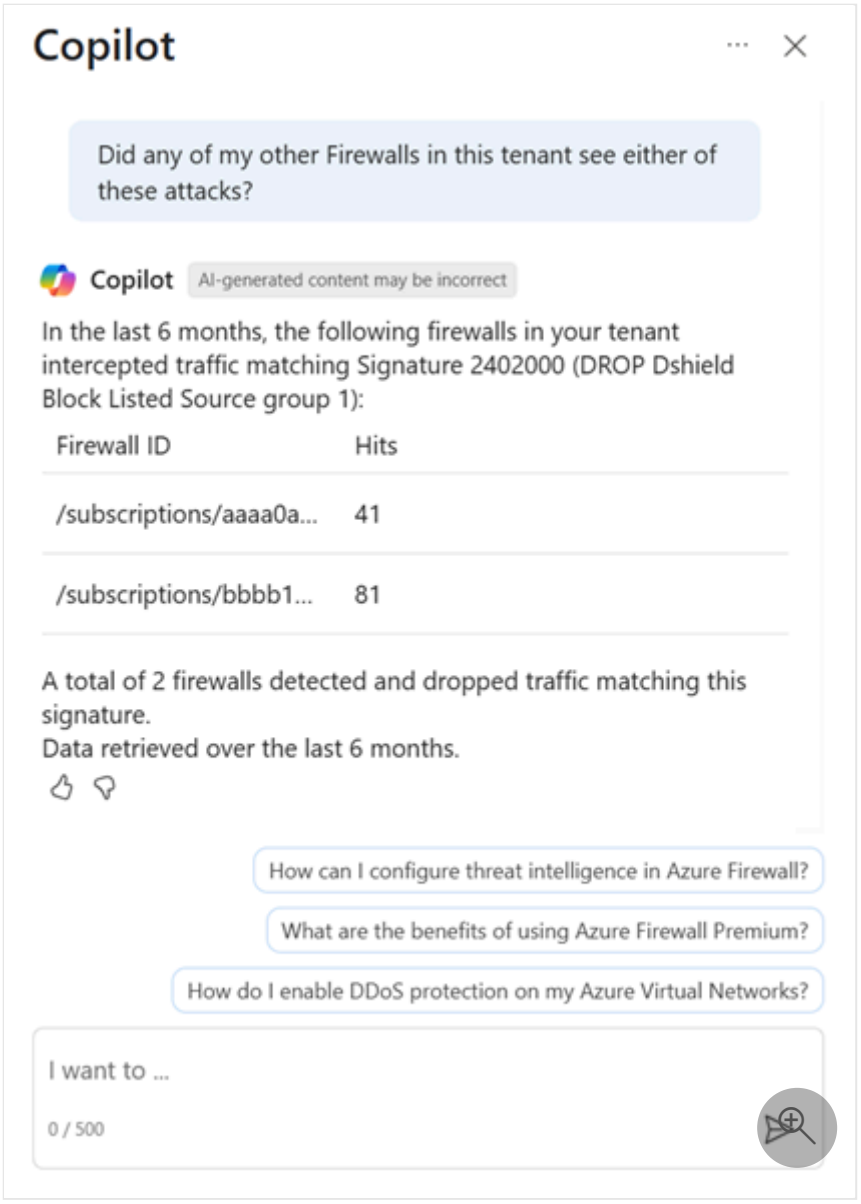

Sample prompts:

- Was signature ID `<ID number\>` only stopped by this one firewall? What about others across this entire tenant?
- Was the top hit seen by any other firewall in the subscription `<subscription name>`?
- Over the past week did any firewall in resource group `<resource group name\>` see signature ID `<ID number>`?

Generate recommendations to secure your environment by using Azure Firewall's IDPS feature

Get **information from documentation** about using Azure Firewall's IDPS feature to secure your environment instead of having to look up this information manually.

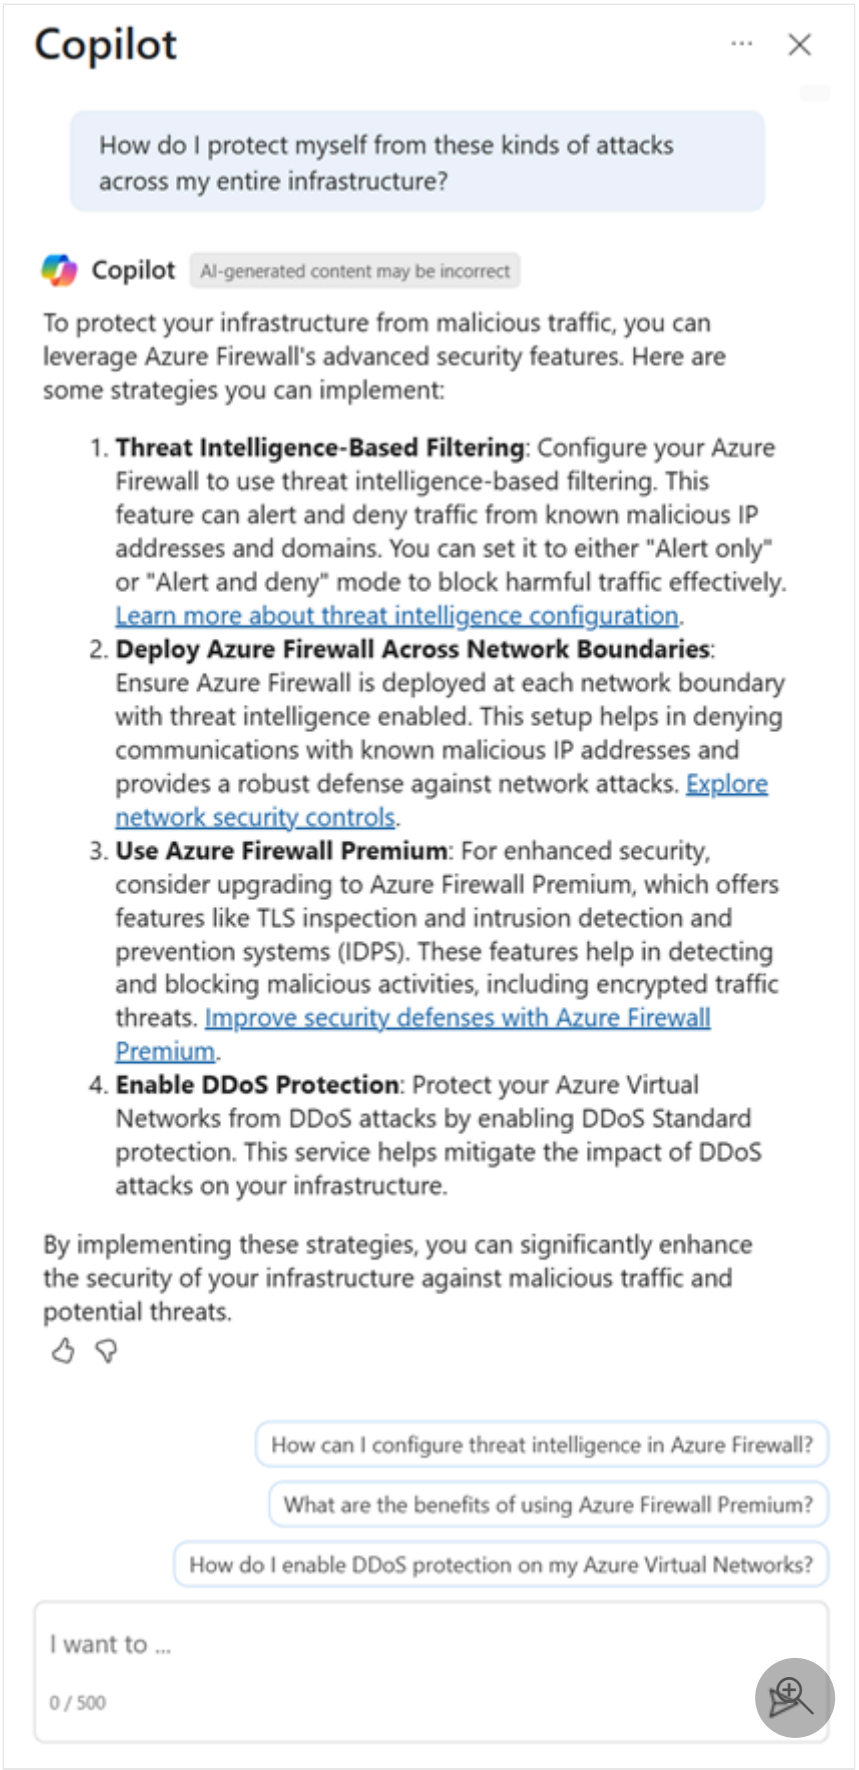

Sample prompts:

- How do I protect myself from future attacks from this attacker across my entire infrastructure?

- If I want to make sure all my Azure Firewalls are protected against attacks from signature ID `<ID number\>`, how do I accomplish this?
- What is the difference in risk between alert only and alert and block modes for IDPS?

> **Note**
>
> Security Copilot might also use the *Ask Microsoft Documentation* capability to provide this information and when using this capability via the Azure Copilot experience, the *Get Information* capability might be used to provide this information.

Provide feedback

Your feedback is vital to guide the current and planned development of the product. The best way to provide this feedback is directly in the product.

Through Security Copilot

Select **How's this response?** at the bottom of each completed prompt and choose any of the following options:

- **Looks right** - Select if the results are accurate, based on your assessment.
- **Needs improvement** - Select if any detail in the results is incorrect or incomplete, based on your assessment.
- **Inappropriate** - Select if the results contain questionable, ambiguous, or potentially harmful information.

For each feedback option, you can provide additional information in the subsequent dialog box. Whenever possible, and especially when the result is **Needs improvement**, write a few words explaining how the outcome can be improved. If you entered prompts specific to Azure Firewall and the results aren't related, include that information.

Through Azure Copilot

Use the **like** and **dislike** buttons at the bottom of each completed prompt. For either feedback option, you can provide additional information in the subsequent dialog box. Whenever possible, and especially when you dislike a response, write a few words explaining how the outcome can be improved. If you entered prompts specific to Azure Firewall and the results aren't related, include that information.

Privacy and data security in Security Copilot

When you interact with Security Copilot either through the Security Copilot portal or the Azure Copilot experience to get information, Copilot pulls that data from Azure Firewall. The prompts, the data retrieved, and the output shown in the prompt results are processed and stored within the Copilot service. For more information, see Privacy and data security in Microsoft Security Copilot.

Related content

- What is Microsoft Security Copilot?
- Microsoft Security Copilot experiences
- Get started with Microsoft Security Copilot
- What is Microsoft Azure Copilot?
- Azure Copilot Capabilities

Last updated on 03/29/2026

Azure Web Application Firewall integration in Microsoft Security Copilot

06/09/2025

Microsoft Security Copilot is a cloud-based AI platform that provides natural language copilot experience. It can help support security professionals in different scenarios, like incident response, threat hunting, and intelligence gathering. For more information, see What is Microsoft Security Copilot?

Azure Web Application Firewall (WAF) integration in Microsoft Security Copilot enables deep investigation of Azure WAF events. It can help you investigate WAF logs triggered by Azure WAF in a matter of minutes and provide related attack vectors using natural language responses at machine speed. It provides visibility into your environment's threat landscape. It allows you to retrieve a list of most frequently triggered WAF rules and identify the top offending IPaddresses in your environment.

Microsoft Security Copilot integration is supported on both Azure WAF on Azure Application Gateway and Azure WAF on Azure Front Door.

Know before you begin

If you're new to Microsoft Security Copilot, you should familiarize yourself with it by reading these articles:

- What is Microsoft Security Copilot?
- Microsoft Security Copilot experiences
- Get started with Microsoft Security Copilot
- Understand authentication in Microsoft Security Copilot
- Prompting in Microsoft Security Copilot

Microsoft Security Copilot integration in Azure WAF

This integration supports the standalone experience and is accessed through https://securitycopilot.microsoft.com. This is a chat-like experience that you can use to ask questions and get answers about your data. For more information, see Microsoft Security Copilot experiences.

Key features

Azure Web Application Firewall integration in Microsoft Security Copilot provides several powerful capabilities to help you analyze and understand your security posture. These features use AI to translate complex WAF logs into actionable insights through natural language responses.

- Providing a list of top Azure WAF rules triggered in the customer environment and generating deep context with related attack vectors.

  This capability provides details about Azure WAF rules that are triggered due to a WAF block. It provides an ordered list of rules based on trigger frequency in the desired time period. The analysis processes Azure WAF logs and connects related logs over a specific time period. The result is an easy-to-understand natural language explanation of why a particular request was blocked.

- Providing a list of malicious IP addresses in the customer environment and generating related threats.

  This capability provides details about client IP addresses blocked by the Azure WAF. The analysis processes Azure WAF logs and connects related logs over a specific time period. The result is an easy-to-understand natural language explanation of which IP addresses the WAF blocked and the reason for the blocks.

- Summarizing SQL injection (SQLi) attacks.

  This capability provides details about SQL injection (SQLi) attacks that were blocked by Azure WAF. By analyzing Azure WAF logs and correlating related data over a specific time period, this skill delivers an easy-to-understand natural language explanation of why SQLi requests were blocked.

- Summarizing Cross-site scripting (XSS) attacks.

  This Azure WAF skill helps you understand why Azure WAF blocked Cross Site Scripting (XSS) attacks to web applications. The skill analyzes Azure WAF logs and connects related incidents over a specific time period. The result is an easy-to-understand natural language explanation of why an XSS request was blocked.

Enable the Azure WAF integration in Copilot for Security

To enable the integration, follow these steps:

1. Ensure that you have at least Copilot contributor permissions.
2. Open https://securitycopilot.microsoft.com/ .
3. Open the Security Copilot menu.
4. Open **Sources** in the prompt bar.
5. On the Plugins page, set the Azure Web Application Firewall toggle to **On**.
6. Select the Settings on the Azure Web Application Firewall plugin to configure the Log Analytics workspace for Azure Front Door WAF or the Azure Application Gateway WAF.
7. To start using the skills, use the prompt bar.

Sample Azure WAF prompts

You can create your own prompts in Microsoft Security Copilot to perform analysis on the attacks based on WAF logs. This section shows some ideas and examples.

Before you begin

- Be clear and specific with your prompts. You might get better results if you include specific device IDs/names, app names, or policy names in your prompts.

  It might also help to add WAF to your prompt. For example:
  - Was there any SQL injection attack in my regional WAF in the last day?
  - Tell me more about the top rules triggered in my global WAF

- Experiment with different prompts and variations to see what works best for your use case. Chat AI models vary, so iterate and refine your prompts based on the results you receive For guidance on writing effective prompts, see Creating your own prompts.

The following example prompts might be helpful.

Summarize information about SQL injection attacks

- Was there a SQL injection attack in my global WAF in the last day?
- Show me IP addresses related to the top SQL injection attack in my global WAF
- Show me all SQL injection attacks in regional WAF in the last 24 hours

Summarize information about cross-site scripting attacks

- Was any XSS attack detected in my Application Gateway WAF in the last 12 hours?

- Show me list of all XSS attacks in my Azure Front Door WAF

Generate a list of threats in my environment based on WAF rules

- What were the top global WAF rules triggered in the last 24 hours?
- What are top threats related to the WAF Rule in my environment? <enter rule ID>
- Was there any bot attack in my regional WAF in the last day?
- Summarize custom rule blocks triggered by Azure Front Door WAF in the last day.

Generate a list of threats in my environment based on malicious IP addresses

- What was the top offending IP in regional WAF in the last day?
- Summarize list of malicious IP addresses in my Azure Front Door WAF in the last six hours?

Provide feedback

Your feedback on the Azure WAF integration with Microsoft Security Copilot helps with development. To provide feedback in Copilot, select **How's this response?** At the bottom of each completed prompt and choose any of the following options:

- Looks right - Select if the results are accurate, based on your assessment.
- Needs improvement - Select if any detail in the results is incorrect or incomplete, based on your assessment.
- Inappropriate - Select if the results contain questionable, ambiguous, or potentially harmful information.

For each feedback item, you can provide more information in the next dialog box that appears. Whenever possible, and when the result is Needs improvement, write a few words explaining what can be done to improve the outcome.

Limitation

If you migrate to Azure Log Analytics dedicated tables in the Application Gateway WAF V2 version, the Microsoft Security Copilot WAF Skills aren't functional. As a temporary workaround, enable Azure Diagnostics as the destination table in addition to the resource-specific table.

Privacy and data security in Microsoft Security Copilot

To understand how Microsoft Security Copilot handles your prompts and the data that's retrieved from the service (prompt output), see Privacy and data security in Microsoft Security Copilot.

Related content

- What is Microsoft Security Copilot?
- Microsoft Security Copilot experiences
- Get started with Microsoft Security Copilot

Security Copilot with Microsoft Sentinel

Applies to: Microsoft Sentinel with Defender XDR in the Microsoft Defender portal, Microsoft Sentinel in the Azure portal, Security Copilot

Microsoft Security Copilot is a platform that helps you defend your organization at machine speed and scale. Microsoft Sentinel's vast security data provides an excellent source for Copilot to help analyze incidents and generate hunting queries.

Together with other Security Copilot sources you enable, your Microsoft Sentinel incidents and data provide wider visibility into threats and their context for your organization.

Know before you begin

If you're new to Security Copilot, you should familiarize yourself with it by reading these articles:

- What is Microsoft Security Copilot?
- Microsoft Security Copilot experiences
- Get started with Microsoft Security Copilot
- Understand authentication in Microsoft Security Copilot
- Prompting in Microsoft Security Copilot

Security Copilot integration with Microsoft Sentinel

Use Microsoft Sentinel data with Security Copilot in both the standalone Security Copilot portal and the embedded experience in the Microsoft Defender portal after you onboard Microsoft Sentinel. For more information, see Microsoft Security Copilot experiences.

Key features

Microsoft Sentinel data integrates with Security Copilot in the Defender portal as follows:

- When you also have Microsoft Defender XDR, Copilot in Microsoft Defender XDR benefits from unified incidents integrated with Microsoft Sentinel.
- In the standalone experience, Microsoft Sentinel provides the following plugins to integrate with Security Copilot:
  Microsoft Sentinel (Preview)
  Natural language to KQL for Microsoft Sentinel (Preview).

Enable Security Copilot integration with Microsoft Sentinel

To maximize your Security Copilot integration with Microsoft Sentinel do the following:

- configure a default Microsoft Sentinel workspace for Security Copilot
- connect your Microsoft Sentinel workspace to Microsoft Defender XDR

Configure a default Microsoft Sentinel workspace

Increase your prompt accuracy by configuring a Microsoft Sentinel workspace as the default.

1. Navigate to Security Copilot at https://securitycopilot.microsoft.com/ .
2. Open **Sources** in the prompt bar.
3. On the **Manage plugins** page, set the toggle to **On**
4. Select the gear icon on the Microsoft Sentinel (Preview) plugin.

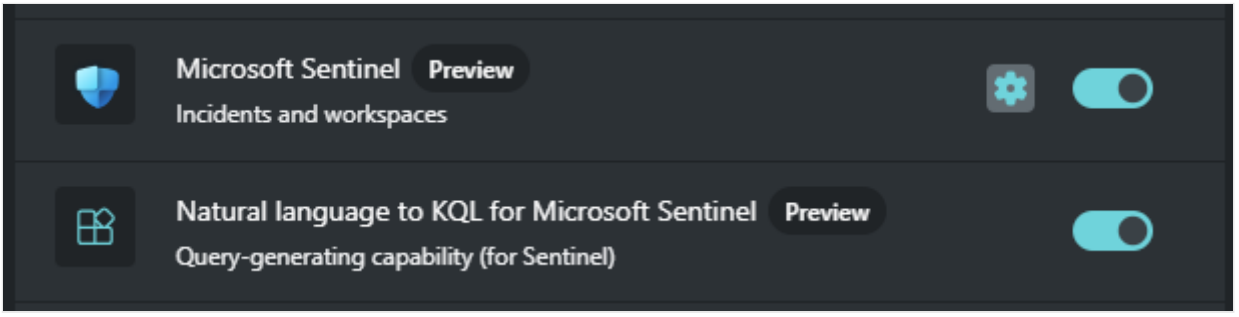

5. Configure the default workspace name.

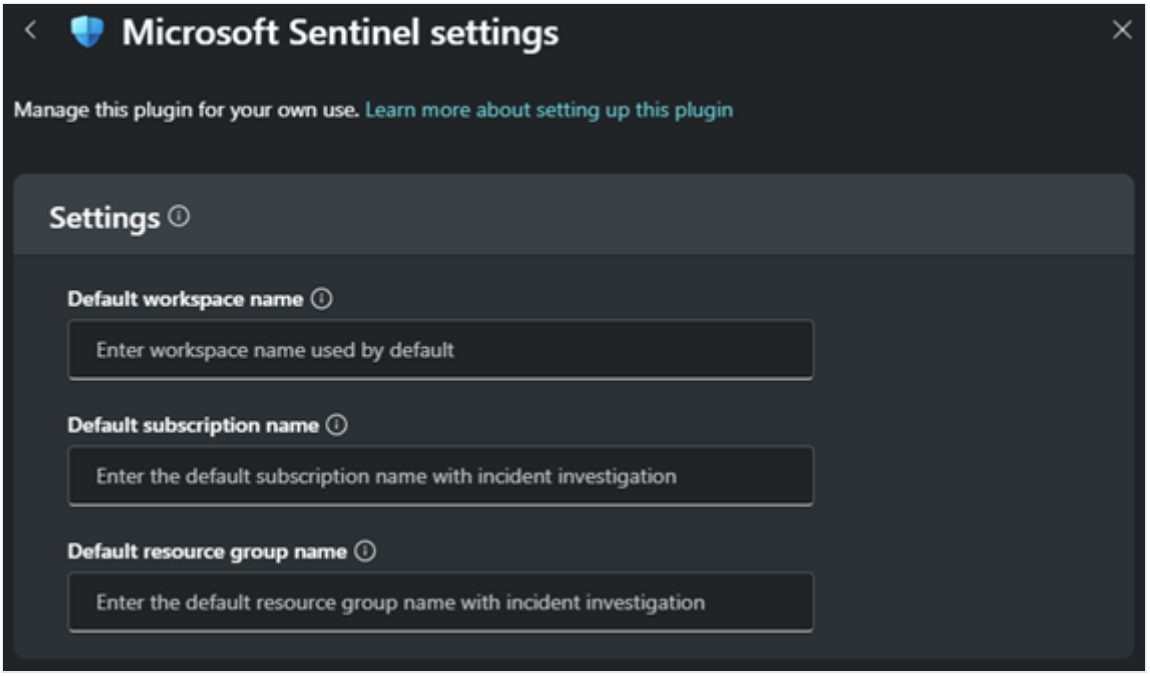

> **Tip**
>
> Specify the workspace in your prompt when it doesn't match the configured default.
>
> Example: `What are the top 5 high priority Sentinel incidents in workspace "soc-sentinel-workspace"?`

Integrate Microsoft Sentinel with Copilot in Defender

Use the Microsoft Defender portal with your Microsoft Sentinel data for an embedded Security Copilot experience. Microsoft Sentinel's unique data sources flowing into Microsoft Defender XDR unified incidents allow Copilot in Defender to maximize its capabilities.

For example:

- The SAP (Preview) solution is installed in your workspace for Microsoft Sentinel.
- The near real-time rule **SAP - (Preview) File Downloaded From a Malicious IP Address** triggers an alert, creating a Microsoft Sentinel incident.
- Microsoft Sentinel was onboarded to the Defender portal.
- Microsoft Sentinel incidents are now unified with Defender XDR incidents.
- Use Copilot in Microsoft Defender for incident summary, guided responses and incident reports.

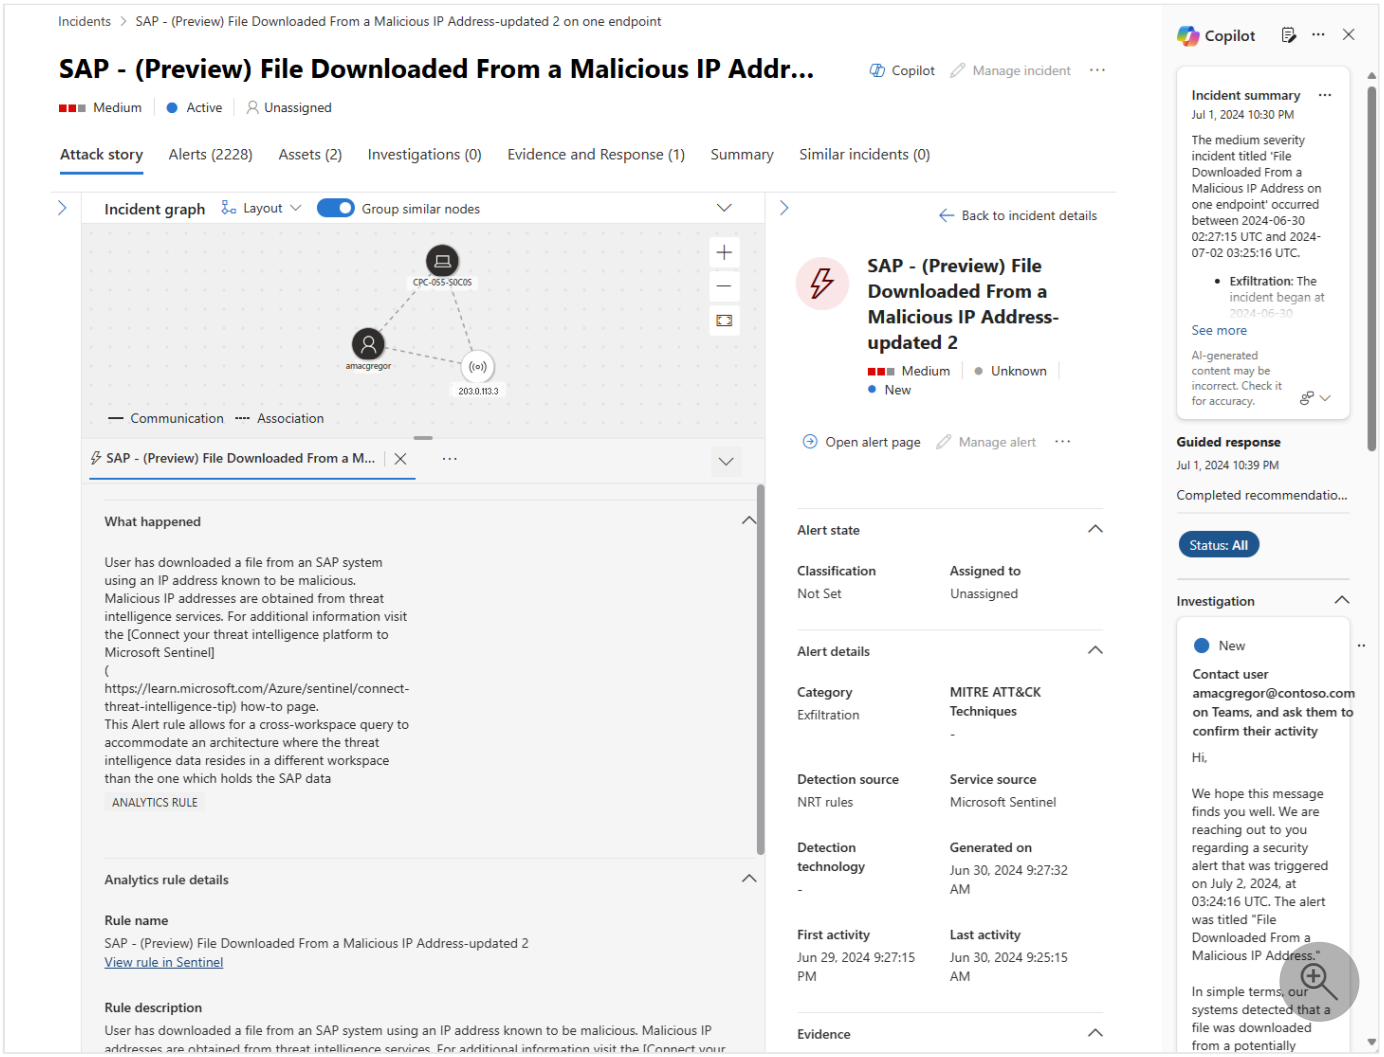

For more information, see the following resources:

- Integrate Microsoft Defender XDR
- Microsoft Sentinel in the Microsoft Defender portal
- Copilot in Microsoft Defender

Integrate Microsoft Sentinel with Security Copilot in advanced hunting

The Natural language to KQL for Microsoft Sentinel (Preview) plugin generates and runs KQL hunting queries using Microsoft Sentinel data. This capability is available in the standalone experience and the advanced hunting section of the Microsoft Defender portal.

> **Note**
>
> In the unified Microsoft Defender portal, you can prompt Security Copilot to generate advanced hunting queries for both Defender XDR and Microsoft Sentinel tables. Not all Microsoft Sentinel tables are currently supported.

For more information, see Security Copilot in advanced hunting.

Sample Microsoft Sentinel prompts

Consider the **Microsoft Sentinel incident investigation** promptbook as a starting point for creating effective prompts. This promptbook delivers a report about a specific incident, along with related alerts, reputation scores, users, and devices.

Expand table

Guidance	Prompt
Nudge Copilot to provide human readable information instead of responding with object IDs.	`Show me Sentinel incidents that were closed as a false positive. Supply the Incident number, Incident Title, and the time they were created.`
Copilot knows who you are. Use the "me" pronoun to find incidents related to you. The following prompt targets incidents assigned to you.	`What Sentinel incidents created in the last 24 hours are assigned to me? List them with highest priority incidents at the top.`
When you narrow a prompt response down to a single incident, Copilot knows the context.	`Tell me about the entities associated with that incident.`
Copilot is good at summarizing. Describe a specific audience you want the prompts and responses summarized for.	`Write an executive report summarizing this investigation. It should be suited for a nontechnical audience.`

For more prompt guidance and samples, see the following resources:

- Using promptbooks
- Prompting in Microsoft Security Copilot
- Rod Trent's Security Copilot Prompt Library

Provide feedback

Your feedback is vital to guide the current and planned development of the product. The best way to provide this feedback is directly in the product. Select **How's this response?** at the bottom of each completed prompt and choose any of the following options:

- **Looks right** - Select if the results are accurate, based on your assessment.
- **Needs improvement** - Select if any detail in the results is incorrect or incomplete, based on your assessment.

- **Inappropriate** - Select if the results contain questionable, ambiguous, or potentially harmful information.

For each feedback option, you can provide more information in the next dialog box that appears. Whenever possible, and especially when the result is **Needs improvement**, write a few words explaining what can be done to improve the outcome. If you entered prompts specific to Azure Firewall and the results aren't related, then include that information.

Privacy and data security in Security Copilot

To understand how Security Copilot handles your prompts and the data that's retrieved from the service (prompt output), see Privacy and data security in Microsoft Security Copilot.

Related articles

- Microsoft Copilot in Microsoft Defender
- Microsoft Defender XDR integration with Microsoft Sentinel

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 08/07/2026

Microsoft Security Copilot and Chat in Microsoft Defender

Applies to: Microsoft Defender XDR, Microsoft Sentinel in the Microsoft Defender portal

This article provides an overview of Microsoft Security Copilot in Microsoft Defender, including key capabilities, access steps, and links to detailed guidance.

> **Note**
>
> Microsoft Defender XDR provides a unified XDR experience for Microsoft Defender for Endpoint, Microsoft Defender for Identity, Microsoft Defender for Office 365, Microsoft Defender for Cloud Apps, and Microsoft Defender for Vulnerability Management. Learn more about this pre- and post-breach defense suite in What is Microsoft Defender XDR?

Security Copilot prerequisites

If you're new to Security Copilot, you should familiarize yourself with it by reading the following articles:

- What is Security Copilot?
- Security Copilot experiences
- Get started with Security Copilot
- Understand authentication in Security Copilot
- Prompting in Security Copilot
- Application card for Microsoft Copilot in Microsoft Defender

Microsoft Security Copilot integration in Microsoft Defender

Microsoft Security Copilot brings together the power of AI and human expertise to help security teams respond to attacks faster and more effectively. Copilot in Defender is available to users who have provisioned access to Security Copilot. You can access Copilot in two ways:

- Security Copilot is embedded in the Microsoft Defender portal to help provide security teams with enhanced capabilities to investigate and respond to incidents, hunt for threats,

and protect their organization with relevant threat intelligence.

- Defender Chat experience (preview) is an open prompt chat assistant built into Microsoft Defender. It helps SOC analysts investigate threats, explore incidents, and answer security questions in plain language, without needing to navigate multiple screens or write complex queries.

Copilot in Defender operates using Microsoft's AI principles. For more information, see the Application card for Microsoft Copilot in Microsoft Defender.

Defender Chat experience (preview)

Get started

To open the Defender chat experience from anywhere in the Defender portal, select the **Copilot** button in the top navigation bar. The chat panel slides open on the right side of the screen and stays in context while you continue working. A welcome screen appears with a greeting and an input field ready for your first question.

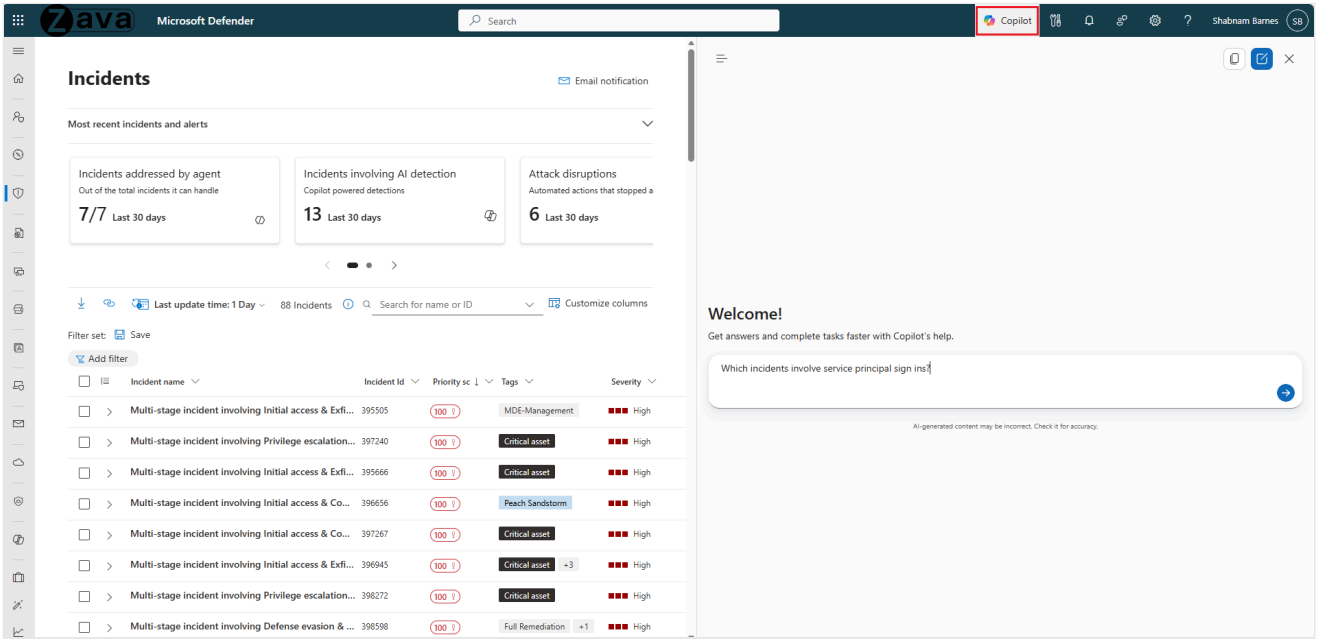

To close the panel, select **Close** in the header or select the Copilot button again. Your conversation is preserved and you can reopen the panel and pick up where you left off.

Page context awareness

Defender Chat responds based on the page you're currently viewing in the Defender portal and can answer questions based on that context.

If you ask a question such as "Which users are involved in *this* incident?", the chat understands which incident, alert, device, or entity you're referring to based on your current page without needing to provide IDs or names.

Chat conversation capabilities

Interactive conversations

The chat remembers the full context of your conversation, so you can ask follow-up questions naturally. For example, you can start with *Show me high-severity incidents from the past week*, then follow up with *Tell me more about the first one*, and the chat understands what you mean.

Step-by-step plans

For complex or multi-step requests, the chat might first present a proposed plan outlining the steps it intends to take. You can Approve or Reject the plan before any actions are taken. This keeps you in control, especially for investigations that require multiple data lookups.

For example: If you ask *Investigate incident 12345 and summarize the key findings*, the chat might propose the following plan:

1. Retrieve incident details
2. Fetch associated alerts
3. Collect evidence and impacted entities
4. Summarize findings

After you approve the plan, the chat executes each step and shows its progress in real time.

Clarifying questions

If your request is ambiguous, the chat might ask a clarifying question and offer quick-select options (up to four suggestions) to help you get to the right answer faster. Select an option or type your own response.

Conversation history

Your conversations are saved automatically. Use the Conversations panel on the left side of the chat to:

- Resume a previous conversation

- Start a new session
- Delete a conversation
- Clear all conversations

> **Note**
>
> - Conversations aren't synced across devices or shared with other users. The last ten conversations are stored locally in your browser.

Working with responses

Responses are formatted with structured tables, bullet points, and section headers for readability. You can:

- Copy a response: Select the *copy* icon on any message to copy it to your clipboard
- Export tables: Select *Export* on any table to export it to Excel for further analysis
- Stop generation: Select *Stop* to interrupt a response that's taking too long or heading in the wrong direction
- Retry: If something goes wrong, select *Retry* to attempt the response again

Provide feedback

All Copilot in Defender capabilities have an option for providing feedback. Reviewing and providing feedback helps improve future responses. To provide feedback, use the 👍 / 👎 buttons on any response.

Privacy and data security

Copilot continuously evolves using data that's stored, processed, and shared depending on the settings defined by your administrator. Microsoft ensures that your data is always protected and secure when using Copilot. To learn more about data security and privacy in Copilot, see Privacy and data security in Copilot.

Plugins in Security Copilot

Copilot uses preinstalled Microsoft plugins like Microsoft Defender, Defender Threat Intelligence, and Natural Language to KQL for Microsoft Sentinel and Defender plugins to generate relevant

information, provide more context to incidents, and generate more accurate results. Ensure that plugins are turned on in Copilot to allow access to relevant data and to generate requested content from other Microsoft services in your organization.

Related content

- Learn how to summarize incidents
- Use guided responses when responding to incidents
- Run script analysis
- Analyze files
- Generate device summaries
- Generate identity summaries
- Generate KQL queries
- Create incident reports
- Use threat intelligence
- Get started with Security Copilot
- Privacy and data security in Copilot
- Security Copilot Responsible AI FAQs
- Security Copilot embedded experiences

> **Tip**
>
> Do you want to learn more? Engage with the Microsoft Security community in our Tech Community: **Microsoft Defender XDR Tech Community** .

Note: The author created this article with assistance from AI. Learn more

Last updated on 07/17/2026

Security Copilot in Microsoft Entra

Microsoft Security Copilot is a platform that brings together the power of AI and human expertise to help administrators and security teams respond to attacks faster and more effectively. Microsoft Entra is one of the Microsoft plugins that enable the Security Copilot platform to generate accurate and relevant information. Through this plugin, Security Copilot can help you investigate and resolve identity risks, assess identities and access with AI-driven intelligence, and complete complex tasks quickly. Security Copilot gets insights from your Microsoft Entra users, groups, sign-in logs, audit logs, and more, while also providing contextualized insights and recommendations in security best practices.

You can explore sign-ins and risky users, get contextualized insights on how to resolve incidents, learn how to protect accounts using natural language. Built on top of real-time machine learning, Security Copilot can help you find gaps in access policies, generate identity workflows, and troubleshoot faster. Security Copilot can assist with scenarios across different products in the Microsoft Entra admin center, such as an incident investigation that includes the Microsoft Entra ID Protection risky user reports and Microsoft Entra sign-in logs.

Security Copilot experiences

Security Copilot in Microsoft Entra includes the standalone Security Copilot experience, embedded experiences in the Microsoft Entra admin center, specialized agents that can perform tasks, and natural language prompts that can be used in several end-to-end scenarios.

For more information, see:

- Microsoft Security Copilot experiences
- Microsoft Security Copilot agents
- Microsoft Security Copilot scenarios

Get started

In the Security Copilot platform, Microsoft Entra is a plugin that provides access to your organization's identity data and insights. To use Security Copilot in Microsoft Entra, you need to onboard to Security Copilot and turn on the Microsoft Entra plugin.

1. Onboard to Security Copilot by following the Get started with Microsoft Security Copilot guide.

- This guide contains guidance around subscription requirements, billing, and capacity.
- Take a moment to understand authentication in Microsoft Security Copilot.

2. In the Security Copilot standalone experience, select the **Sources** icon from the prompt bar and turn on the Microsoft Entra plugin for Security Copilot, if it's not already on.

Once you're all set up in Security Copilot, you can start using natural language prompts to help remediate identity-based incidents. You can always check the **Promptbook library** in the standalone Security Copilot experience for more examples.

- *Give me all user details for karita@woodgrovebank.com and extract the user Object ID.*
- *Does karita@woodgrovebank.com have any registered devices in Microsoft Entra?*
- *List the recent risky sign-ins for karita@woodgrovebank.com.*
- *Can you give me sign-in logs for karita@woodgrovebank.com for the past 48 hours? Put this information in a table format.*
- *Get Microsoft Entra audit logs for karita@woodgrovebank.com for the past 72 hours. Put information in table format.*

Access Copilot chat in the Microsoft Entra admin center

Copilot chat is available directly from the left side navigation in the Microsoft Entra admin center.

1. Sign in to the Microsoft Entra admin center as at least a Security Reader.

2. Select **Copilot chat** from the side navigation menu.

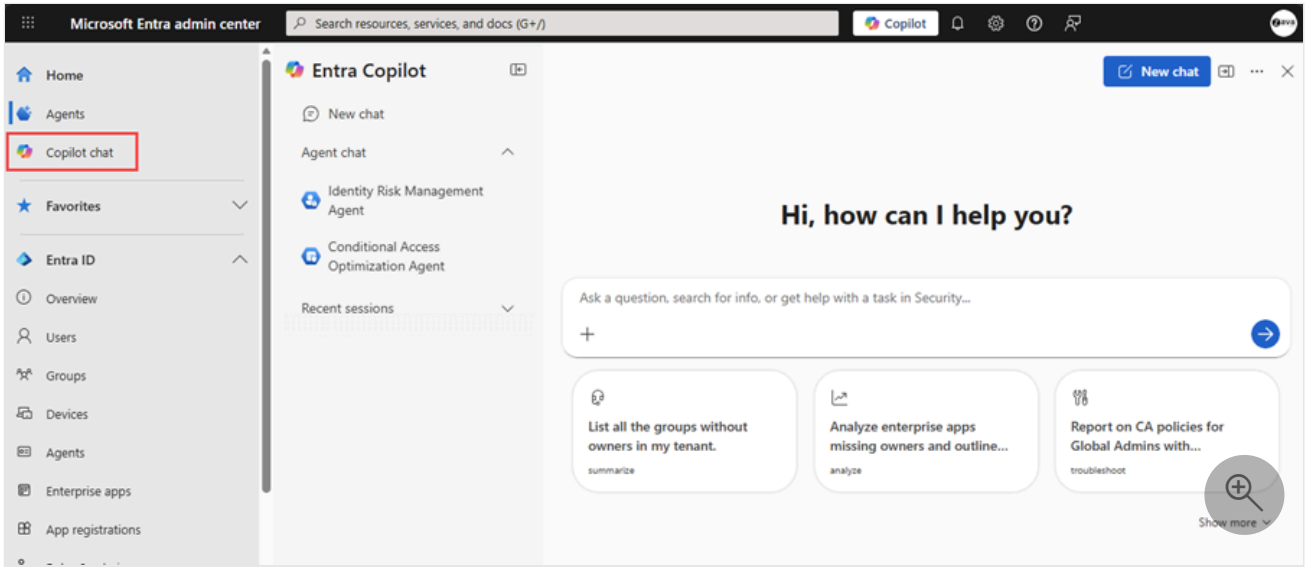

3. Enter a question or task in the chat prompt bar, or select one of the suggestion cards to get started.

The **Entra Copilot** panel opens with the following options:

- **New chat**: Start a new conversation with Copilot.
- **Agent chat**: Select a specialized agent, such as the **Identity Risk Management Agent** or the **Conditional Access Optimization Agent**, to get targeted assistance.
- **Recent sessions**: View and resume previous Copilot chat sessions.

The Microsoft Entra Copilot chat experience is designed to meet you where you are in your workflow and stay with you when needed. You can expand or collapse the chat history pane on the left and switch between "sidecar mode" and "fullscreen mode" in the upper-right corner.

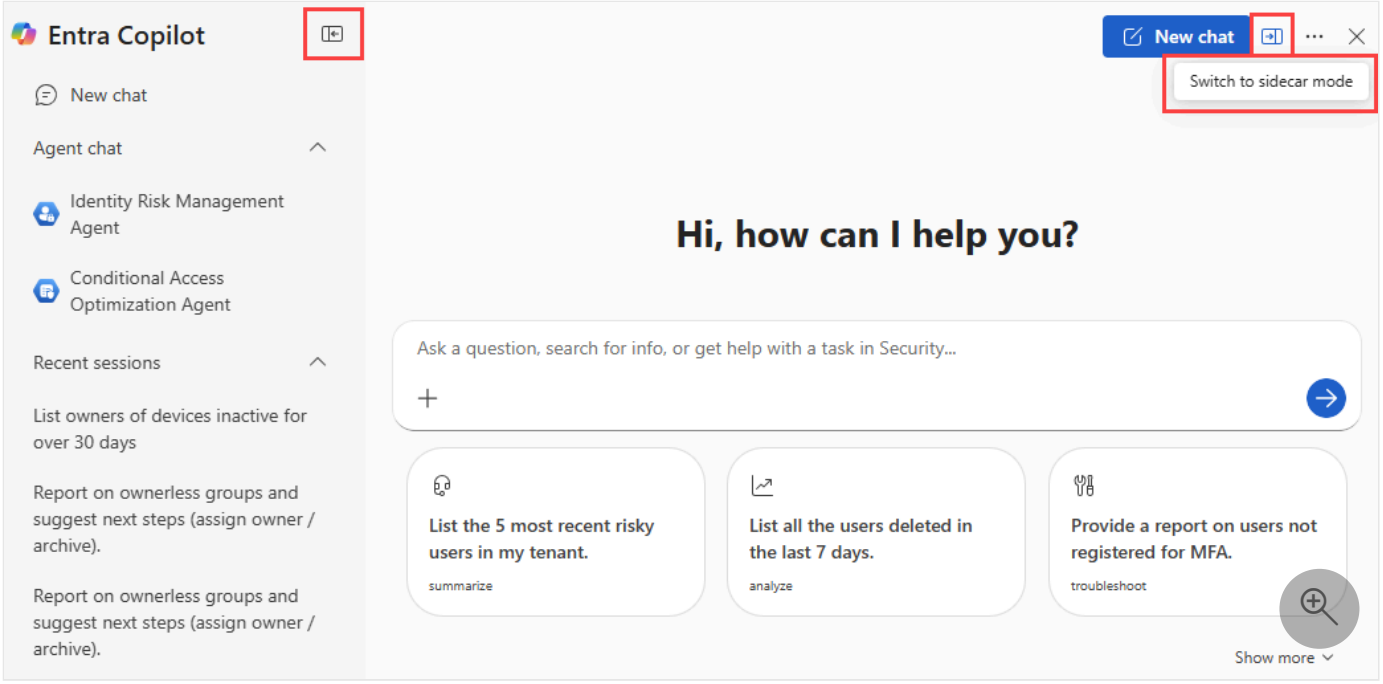

In addition, when you navigate around the Microsoft Entra admin center from the Copilot chat experience, the sidecar remains open in the new page, allowing you to continue your conversation with Copilot without interruption.

Provide feedback

Copilot in Microsoft Entra uses AI and machine learning to process data and generate responses for each of the key features. However, AI-generated content might be incorrect. Your feedback on the generated responses helps improve the accuracy of Copilot and Microsoft Entra over time.

All key features have an option for providing feedback, but the exact steps vary based on the feature. Look for the thumbs up and thumbs down icons to let us know if the response was helpful or not.

Privacy and data security in Security Copilot

To understand how Security Copilot handles your prompts and the data that's retrieved from the service(prompt output), see Privacy and data security in Microsoft Security Copilot.

Related content

- Responsible AI FAQs
- Investigate security incidents using the Microsoft Entra skills in Microsoft Security Copilot.
- Investigate risky apps using the Microsoft Entra skills in Microsoft Security Copilot.

ⓘ **Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 05/13/2026

Security Copilot in Microsoft Intune

Article • 03/03/2025

Microsoft Security Copilot is a cloud-based AI platform that provides a natural language Copilot experience. It can help support security professionals in different scenarios, like incident response, threat hunting, and intelligence gathering. For more information about what it can do, go to What is Microsoft Security Copilot?

Know before you begin

If you're new to Security Copilot, you should familiarize yourself with it by reading these articles:

- What is Microsoft Security Copilot?
- Microsoft Security Copilot experiences
- Get started with Microsoft Security Copilot
- Understand authentication in Microsoft Security Copilot
- Prompting in Microsoft Security Copilot

Security Copilot integration in Microsoft Intune

If you use Microsoft Intune in the same tenant as Security Copilot, then you can use Security Copilot to get insights about your Intune data.

There are Intune capabilities built into Security Copilot, and you can use prompts to get more information, including:

- Information about your devices, apps, compliance & configuration policies, and policy assignments managed in Intune
- Managed device attributes and hardware details
- Issue with specific devices and compare a working & non-working device

This article shows you how to access your Microsoft Intune data in Security Copilot and includes sample prompts.

Key features

There are three areas to use Copilot in Intune:

- Policy and setting management

- Device details and troubleshooting
- Device query

Security admin focus

Security Copilot has a Security Operations Center (SOC) or security admin focus. So, if you're a SOC analyst or security admin, then you can use Security Copilot to get the security posture of devices that Intune manages.

For example, there's a user or device that is showing signs of malicious intent. Also, you notice some events are happening after the malicious intent, like an unknown device enrolling in Intune. Maybe someone is trying to use stolen credentials to enroll and get access. You need to get more information.

In Security Copilot, you can use the Intune capabilities to get more information, like:

- Ask about a specific device, get all the properties about that device, including the device name, device ID, and device manufacturer.
- Determine when the device is enrolled in Intune.
- Find the primary user of a device
- Determine the type of device, like a laptop or mobile phone.
- Check the compliance status, especially if a device is noncompliant, and why it's noncompliant.

In Microsoft Defender, you can use this information, including the device type, to determine your next steps. For example, you might take different actions based on the type of device (laptop vs. mobile phone vs. tablet). Security Copilot can also give you a link to the device in Microsoft Defender, so you can run any Defender actions.

What you need to know

- When an admin submits a prompt, Copilot can only access the data that the admin has permissions to, which includes the RBAC roles and Intune scope tags assigned to them.

  If you want your admins to access all your Intune data in Security Copilot, then use the following role in Microsoft Entra ID:
  - Intune Service Administrator (also known as Intune Administrator)

  For more information on roles and authentication, go to:
  - Roles and authentication in Microsoft Security Copilot
  - Role based access control (RBAC) in Intune
  - Use RBAC and scope tags for distributed IT in Intune

- You can access your Intune data in the Security Copilot portal and Copilot in the Intune admin center. For more information on Copilot in Intune vs. Security Copilot, and other common questions, go to the Microsoft Copilot in Intune FAQ.

Enable the Security Copilot integration in Intune

To use the Intune capabilities in Security Copilot, enable the Intune plugin.

1. Go to Security Copilot and sign in with your credentials.
2. In the prompt bar, select **Sources** (right corner).

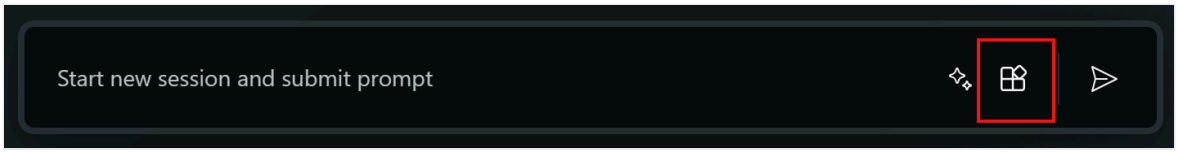

3. In **Manage sources**, turn on Microsoft Intune:

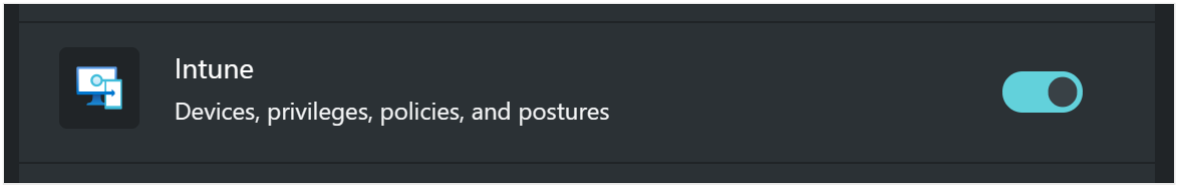

> **Note**
>
> Some roles can enable or disable plugins. For more information, go to Manage plugins in Microsoft Security Copilot.

Use the built-in features

In Security Copilot, there are built in system features that are helpful for Intune admins. For a walkthrough of Security Copilot, go to Navigating Microsoft Security Copilot.

This section describes some of the features that are helpful for Intune admins.

System capabilities

Capabilities are built-in features that can get data from the different plugins that you enable, including Microsoft Intune. When you use a prompt to ask something about

your Intune data, like apps assigned to a user or device details, your prompts use these Intune capabilities.

To view the list of Intune built-in system capabilities for Intune, use the following steps:

1. In the Security Copilot portal prompt bar, select the Copilot prompts icon > **See all system capabilities**.

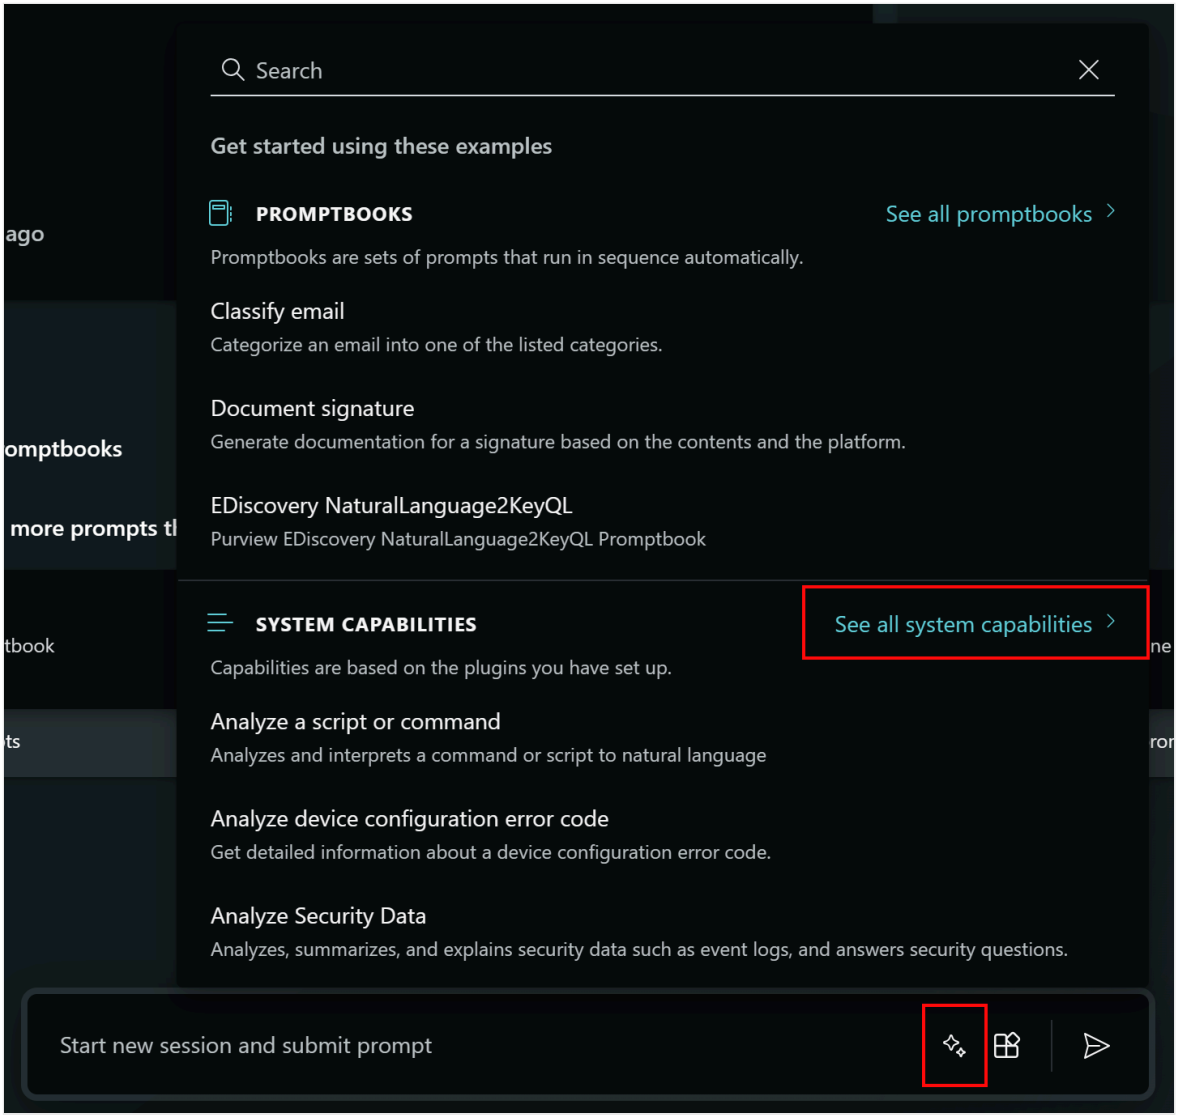

2. In the Microsoft Intune section, there's a list of all the built-in capabilities for Intune. You can select any of the capabilities and get more information about that capability.

Sessions

When you use prompts in the Microsoft Intune admin center or in the Security Copilot portal, the sessions are saved. To see the saved sessions, use the following steps:

1. In the Security Copilot portal, go to the menu at the top left > **My sessions**.

2. When you select a session, your previous prompts and results are shown. Every session also has a session ID in the URL. You can share this session ID with others to review the same prompt session.

   For example, your session ID is something like `https://securitycopilot.microsoft.com/sessions/023d1c61-f3c7-4702-8924-075a1058900d`.

Sample Intune prompts

You can create your own prompts in Security Copilot to get information about your Intune data. This section lists some ideas and examples.

Before you begin

- Be clear and specific with your prompts. You might get better results if you include specific device IDs or names, app names, or policy names in your prompts.

  It might also help to add **Intune** to your prompt, like:
  - **According to Intune, how many devices were enrolled this week?**
  - **Tell me about Intune devices for (user name).**

- Experiment with different prompts and variations to see what works best for your use case. Chat AI models vary, so iterate and refine your prompts based on the results you receive.

  You can also save your prompts in a promptbook for future use. For more information, go to:
  - Prompting in Microsoft Security Copilot
  - Using promptbooks in Microsoft Security Copilot

General information about your Intune data

Get **general information** about your Intune data, like the number of devices, apps deployed, platform versions of your devices, and more.

Sample prompts:

- What apps are added to Intune?
- What Intune apps are assigned the most?
- How many devices were enrolled in Intune in the last 24 hours?
- Tell me about Intune devices for Jon Smith.

Policy targets

Get details on **policy targets**, like the groups that have a specific app assigned or how many users have a specific app assigned.

Sample prompts:

- How many users is ContosoApp assigned to?
- Which groups are ContosoApp assigned to?
- How many apps are assigned to the device ID *Enter the device ID* in Intune?
- Why is the "Allow Microsoft Store App to auto update" policy applying to DeviceA?
- Tell me about Intune devices for user UserA.
- Why is PolicyA applying on DeviceB?

Specific devices

Get information about a **specific device**, like its group memberships and the apps assigned to it.

Sample prompts:

- What devices are used by UserA@contoso.com?
- What groups is DeviceA in?
- Tell me about DeviceA.
- Who is the primary user for DeviceA?
- Is ContosoApp installed on DeviceA?
- Show me discovered apps on DeviceA.

Similarities and differences

Get the **similarities and differences** between two devices, like the compliance policies, hardware, and device configurations assigned to both devices.

Sample prompts:

- What is the hardware configuration difference between the DeviceA and DeviceB devices?
- What are the similarities in compliance policies between the DeviceA and DeviceB devices?
- What is the difference in device configuration profile between the DeviceA and DeviceB devices?
- Compare installed applications on DeviceA and DeviceB.

Provide feedback

Your feedback on the Intune integration with Security Copilot helps with development. To provide feedback, in Security Copilot, use the feedback buttons at the bottom of each completed prompt.

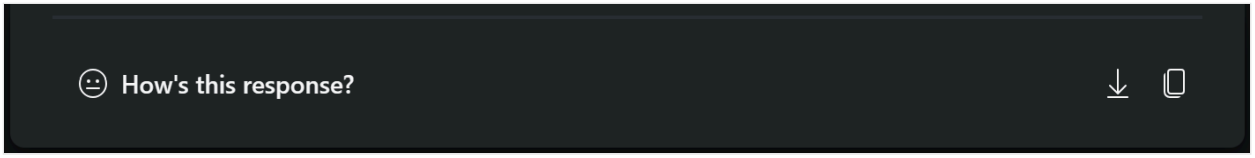

Whenever possible, and when the result isn't what you expect, write a few words explaining what can be done to improve the outcome. If you entered Intune-specific prompts and the results aren't Intune related, then include that information.

Privacy and data security in Security Copilot

For more information about data privacy in Security Copilot, go to Privacy and data security in Microsoft Security Copilot.

When you interact with the Security Copilot to get Intune data, the Security Copilot pulls that data from Intune. The prompts, the Intune data that's retrieved, and the output shown in the prompt results is processed and stored within the Security Copilot service.

When you use Security Copilot to get Intune data, Security Copilot also has access to the data and permissions defined by the RBAC roles and Intune scope tags assigned to you.

Related articles

- What is Microsoft Security Copilot?
- Privacy and data security in Microsoft Security Copilot
- Use Microsoft Copilot in Intune

Feedback

Was this page helpful? Yes No

Microsoft Security Copilot integration in Defender EASM

Microsoft Defender External Attack Surface Management (Defender EASM) continuously discovers and maps your digital attack surface to provide an external view of your online infrastructure. This visibility helps security and IT teams identify unknowns, prioritize risk, eliminate threats, and extend vulnerability and exposure control beyond the firewall. Attack surface insights are generated by analyzing vulnerability and infrastructure data to showcase the key areas of concern for your organization.

Microsoft Security Copilot (Security Copilot) integration in Defender EASM helps you interact with Microsoft-discovered attack surfaces. Identifying attack surfaces helps your organization quickly understand its externally facing infrastructure and relevant, critical risks. It provides insight into specific areas of risk, including vulnerabilities, compliance, and security hygiene.

For more information about Security Copilot, see What is Security Copilot? For information about the embedded Security Copilot experience, see Query your attack surface with Defender EASM by using Azure Copilot.

Know before you begin

If you're new to Security Copilot, it's a good idea to familiarize yourself with the solution by reading these articles:

- What is Microsoft Security Copilot?
- Security Copilot experiences
- Get started with Security Copilot
- Understand authentication in Security Copilot
- Prompting in Security Copilot

Security Copilot in Defender EASM

Security Copilot can surface insights from Defender EASM about your organization's attack surface. You can use Security Copilot built-in features. To get more information, use prompts in Security Copilot. The information can help you understand your security posture and mitigate vulnerabilities.

This article introduces you to Security Copilot and includes sample prompts that can help Defender EASM users.

Key features

The EASM Security Copilot integration can help you:

- **Get a snapshot of your external attack surface and generate insights into potential risks**: You can get a quick view of your external attack surface by analyzing internet-available information combined with the Defender EASM proprietary discovery algorithm. It provides an easy-to-understand natural language explanation of the organization's externally facing assets, such as hosts, domains, webpages, and IP addresses. It highlights the critical risks associated with each.

- **Prioritize remediation efforts based on asset risk and Common Vulnerabilities and Exposures (CVEs) list items**: Defender EASM helps security teams prioritize their remediation efforts by helping them understand which assets and CVEs pose the greatest risk in their environment. It analyzes vulnerability and infrastructure data to showcase key areas of concern, providing a natural language explanation of the risks and recommended actions.

- **Use Security Copilot to surface insights**: You can use Security Copilot to ask about insights by using natural language and extract insights from Defender EASM about your organization's attack surface. Query details like the number of Secure Sockets Layer (SSL) certificates that aren't secure, ports that are detected, and specific vulnerabilities that affect the attack surface.

- **Expedite attack surface curation**: Use Security Copilot to curate your attack surface by using labels, external IDs, and state modifications for a set of assets. This process speeds up curation, so you can organize your inventory faster and more efficiently.

Enable Security Copilot integration

To set up Security Copilot integration in Defender EASM, complete the steps described in the next sections.

Prerequisites

To enable integration, you need to have these prerequisites:

- Access to Microsoft Security Copilot
- Permissions to activate new connections

Connect Security Copilot to Defender EASM

1. Access Security Copilot and ensure that you're authenticated.

2. Select the **Security Copilot plugin** icon on the upper-right side of the prompt input bar.

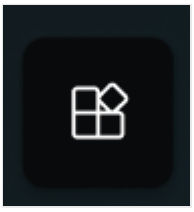

3. Under **Microsoft**, locate **Defender External Attack Surface Management**. Select **On** to connect.

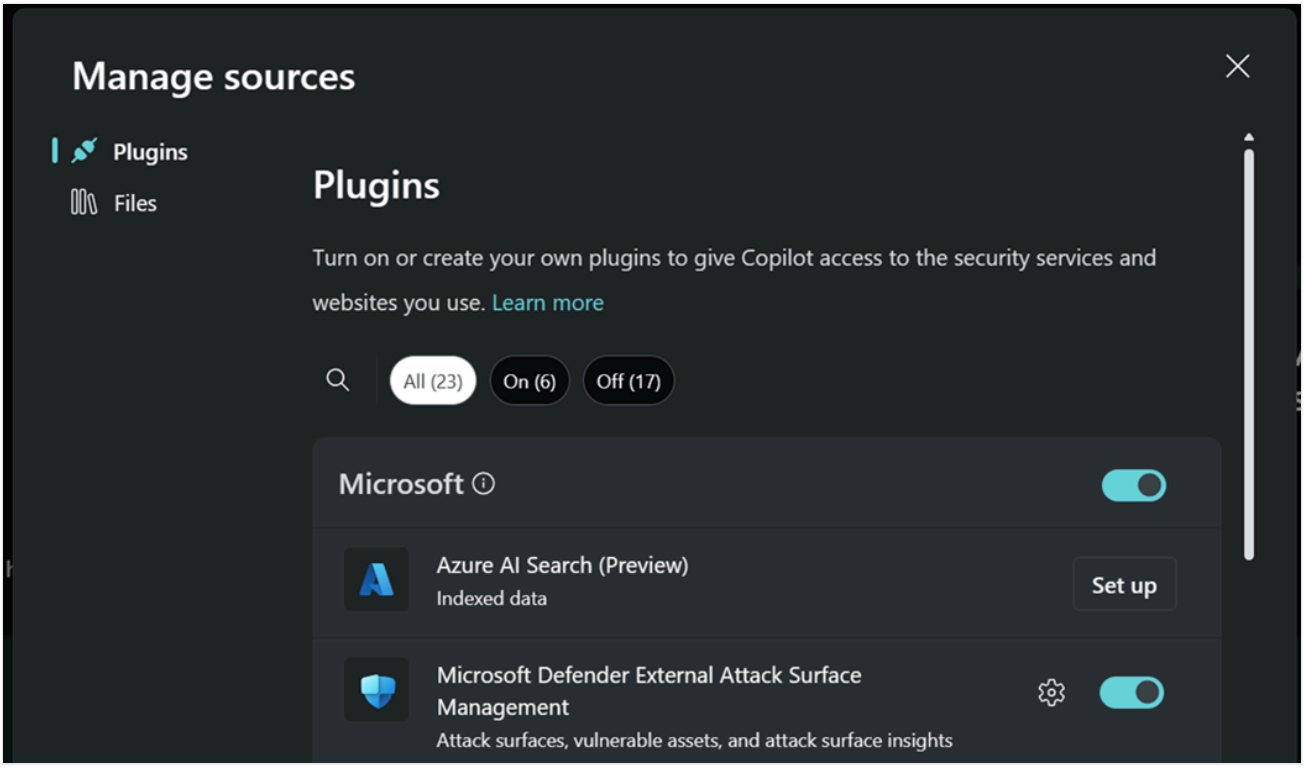

4. If you want Security Copilot to pull data from your Defender EASM resource, select the gear icon to open the plugin settings. Enter or select values by using the values from your resource's **Essentials** section on the **Overview** pane.

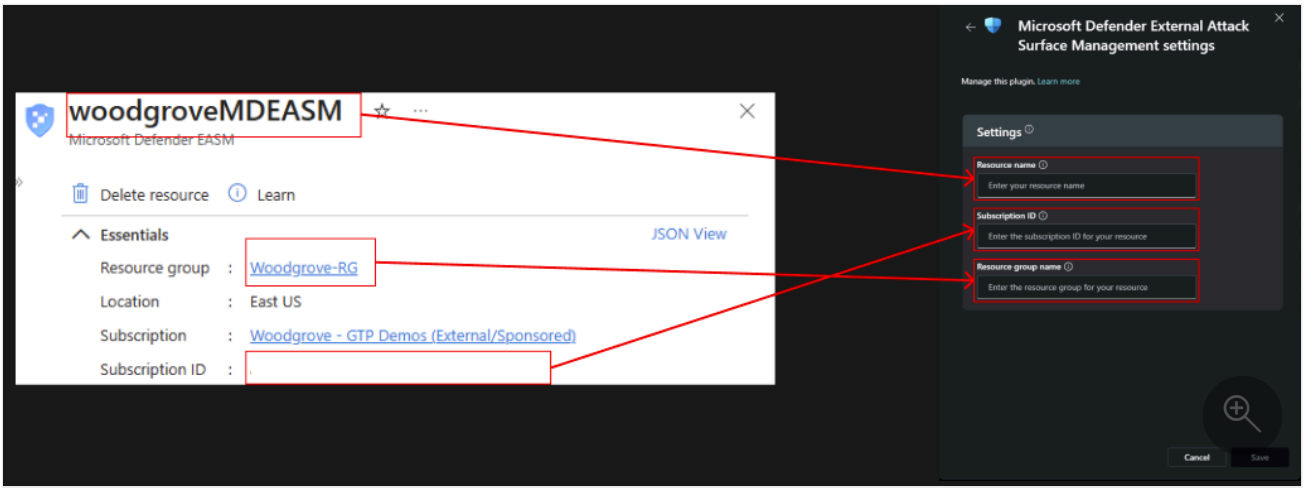

> **Note**
>
> You can use your Defender EASM skills even if you haven't purchased Defender EASM. For more information, see Plugin capabilities reference.

Sample Defender EASM prompts

Security Copilot primarily uses natural language prompts. When you query information from Defender EASM, you submit a prompt that guides Security Copilot to select the Defender EASM plugin and invoke the relevant capability.

For success with Security Copilot prompts, we recommend the following approaches:

- Ensure that you reference the company name in your first prompt. Unless otherwise specified, all future prompts then provide data about the initially specified company.
- Be clear and specific with your prompts. You might get better results if you include specific asset names or metadata values (for example, CVE IDs) in your prompts.

  It might also help to add **Defender EASM** to your prompt, like in these examples:
  - **According to Defender EASM, what are my expired domains?**
  - **Tell me about Defender EASM high-priority attack surface insights.**
- Experiment with different prompts and variations to see what works best for your use case. Chat AI models vary, so iterate and refine your prompts based on the results you receive.
- Security Copilot saves your prompt sessions. To see previous sessions, in Security Copilot, on the menu, select **My sessions**.

For a walkthrough of Security Copilot, including the pin and share features, see Navigate Security Copilot.

For more information on writing Security Copilot prompts, see Security Copilot prompting tips.

Plugin capabilities reference

Expand table

Capability	Description	Inputs	Behaviors
Get attack surface summary	Returns the attack surface summary for either the customer's Defender EASM resource or a specific company name.	**Example inputs**: - Get attack surface for LinkedIn. - Get my attack surface. - What is the attack surface for Microsoft? - What is my attack surface? - What are the externally facing assets for Azure? - What are my externally facing assets? **Optional inputs**: - `CompanyName`	If your plugin is configured to an active Defender EASM resource and no other company is specified: - Returns an attack surface summary for the customer's Defender EASM resource. If another company name is provided: - If no exact for match for company name is found, returns a list of possible matches. - If there's an exact match, returns the attack surface summary for the company name.
Get attack surface insights	Returns the attack surface insights for either the customer's Defender EASM resource or a specific company name.	**Example inputs**: - Get high-priority attack surface insights for LinkedIn. - Get my high-priority attack surface insights. - Get low priority attack surface insights for Microsoft. - Get low priority attack surface insights. - Do I have high-priority vulnerabilities in my external attack surface for Azure? **Required inputs**: - `PriorityLevel` (the priority	If your plugin is configured to an active Defender EASM resource and no other company is specified: - Returns attack surface insights for the customer's Defender EASM resource. If another company name is provided: - If no exact for match for company name is found, returns a list of

Capability	Description	Inputs	Behaviors
		level must be high, medium, or low; if not provided, it defaults to high) **Optional inputs**: - `CompanyName` (the company name)	possible matches. - If there's an exact match, returns the attack surface insights for the company name.
Get assets affected by a CVE	Returns the assets affected by a CVE for either the customer's Defender EASM resource or a specific company name.	**Example inputs**: - Get assets affected by CVE-2023-0012 for LinkedIn. - Which assets are affected by CVE-2023-0012 for Microsoft? - Is Azure's external attack surface impacted by CVE-2023-0012? - Get assets affected by CVE-2023-0012 for my attack surface. - Which of my assets are affected by CVE-2023-0012? - Is my external attack surface impacted by CVE-2023-0012? **Required inputs**: - `CveId` **Optional inputs**: - `CompanyName`	If your plugin is configured to an active Defender EASM resource and no other company is specified: - If plugin settings aren't filled out, fail graciously and remind customers. - If plugin settings are filled out, returns the assets affected by a CVE for the customer's Defender EASM resource. If another company name is provided: - If no exact for match for company name is found, returns a list of possible matches. - If there's an exact match, returns the assets affected by a CVE for the specific company name.
Get assets affected by a CVSS	Returns the assets affected by a Common Vulnerability Scoring System (CVSS) score for either the customer's Defender EASM resource or a specific company name.	**Example inputs**: - Get assets affected by high-priority CVSS scores in LinkedIn's attack surface. - How many assets have critical CVSS score for Microsoft? - Which assets have critical CVSS scores for Azure? - Get assets affected by high-priority CVSS scores in my attack surface.	If your plugin is configured to an active Defender EASM resource and no other company is specified: - If plugin settings aren't filled out, fail graciously and remind customers. - If plugin settings are filled out, returns the assets affected by a

Capability	Description	Inputs	Behaviors
		- How many of my assets have critical CVSS scores? - Which of my assets have critical CVSS scores? **Required inputs**: - `CvssPriority` (the CVSS priority must be *critical*, *high*, *medium*, or *low*) **Optional inputs**: - `CompanyName`	CVSS score for the customer's Defender EASM resource. If another company name is provided: - If no exact for match for company name is found, returns a list of possible matches. - If there's an exact match, returns the assets affected by a CVSS score for the specific company name.
Get expired domains	Returns the number of expired domains for either the customer's Defender EASM resource or a specific company name.	**Example inputs**: - How many domains are expired in LinkedIn's attack surface? - How many assets are using expired domains for Microsoft? - How many domains are expired in my attack surface? - How many of my assets are using expired domains for Microsoft? **Optional inputs**: - `CompanyName`	If your plugin is configured to an active Defender EASM resource and no other company is specified: - Returns the number of expired domains for the customer's Defender EASM resource. If another company name is provided: - If no exact for match for company name is found, returns a list of possible matches. - If there's an exact match, returns the number of expired domains for the specific company name.
Get expired certificates	Returns the number of expired SSL certificates for either the customer's Defender EASM resource or a specific company name.	**Example inputs**: - How many SSL certificates are expired for LinkedIn? - How many assets are using expired SSL certificates for Microsoft? - How many SSL certificates	If your plugin is configured to an active Defender EASM resource and no other company is specified: - Returns the number of SSL certificates for

Capability	Description	Inputs	Behaviors
		are expired for my attack surface? - What are my expired SSL certificates? **Optional inputs**: - `CompanyName`	the customer's Defender EASM resource. If another company name is provided: - If no exact for match for company name is found, returns a list of possible matches. - If there's an exact match, returns the number of SSL certificates for the specific company name.
Get SHA1 certificates	Returns the number of SHA1 SSL certificates for either the customer's Defender EASM resource or a specific company name.	**Example inputs**: - How many SSL SHA1 certificates are present for LinkedIn? - How many assets are using SSL SHA1 for Microsoft? - How many SSL SHA1 certificates are present for my attack surface? - How many of my assets are using SSL SHA1? **Optional inputs**: - `CompanyName`	If your plugin is configured to an active Defender EASM resource and no other company is specified: - Returns the number of SHA1 SSL certificates for the customer's Defender EASM resource. If another company name is provided: - If no exact for match for company name is found, returns a list of possible matches. - If there's an exact match, returns the number of SHA1 SSL certificates for the specific company name.
Translate natural language to a Defender EASM query	Translates any natural language question into a Defender EASM query and returns the assets that match the query.	**Example inputs**: - What assets are using jQuery version 3.1.0? - Get the hosts with port 80 open in my attack surface. - Find all the page, host, and ASN assets in my inventory	If your plugin is configured to an active Defender EASM resource: - Returns the assets matching with the translated query.

Capability	Description	Inputs	Behaviors
		that have an IP address that is IP *X*, IP *Y*, or IP *Z*. - Which of my assets have a registrant email of `<name@example.com>`?	

Switch between resource data and company data

Even though we added resource integration for our skills, we still support pulling data from prebuilt attack surfaces for specific companies. To improve the Security Copilot accuracy in determining when a customer wants to pull from their attack surface or from a prebuilt, company attack surface, we recommend using **my**, **my attack surface**, and so on, to convey that you want to use your resource. Use **their**, ***specific company name***, and so on to convey that you want to use a prebuilt attack surface. Although this approach improves the experience in a single session, we strongly recommend using two separate sessions to avoid any confusion.

Provide feedback

Your feedback on Security Copilot generally, and the Defender EASM plugin specifically, is vital to guide current and planned development of the product. The optimal way to provide this feedback is directly in the product, using the feedback buttons at the bottom of each completed prompt. Select **Looks right**, **Needs improvement**, or **Inappropriate**. We recommend that you choose **Looks right** when the result matches expectations, **Needs improvement** when it doesn't, and **Inappropriate** when the result is harmful in some way.

Whenever possible, and especially when the result you select is **Needs improvement**, please write a few words to explain what we can do to improve the outcome. This request also applies when you expect Security Copilot to invoke the Defender EASM plugin, but a different plugin is engaged instead.

Privacy and data security in Security Copilot

When you interact with Security Copilot to get Defender EASM data, Copilot pulls that data from Defender EASM. The prompts, the data that's retrieved, and the output that's shown in the prompt results are processed and stored in the Security Copilot service.

For more information about data privacy in Security Copilot, see Privacy and data security in Security Copilot.

Related content

- What is Microsoft Security Copilot?
- Privacy and data security in Microsoft Security Copilot
- Query your attack surface with Defender EASM using Azure Copilot

---

Last updated on 08/19/2026

Security Copilot in Microsoft Surface Management Portal (Public Preview)

Article • 03/03/2025

Microsoft Security Copilot is a cloud-based AI platform that provides a natural language Copilot experience that can be used to provide information and resolve issues. Surface Management Portal provides a centralized platform to manage and monitor Surface devices at scale. Surface Management Portal integration with Microsoft Copilot for Security helps Surface device administrators be more efficient, and resolve hardware issues faster and at scale.

Copilot for Surface Management Portal provides visibility into your environment's warranty coverage landscape, and end-of-servicing timelines. Additionally, it offers insights into various aspects of device management, including compliance policies, malware protection, and other threats to reduce risks.

Know before you begin

If you're new to Security Copilot, you should familiarize yourself with it by reading these articles:

- What is Microsoft Security Copilot?
- Microsoft Security Copilot experiences
- Get started with Microsoft Security Copilot
- Understand authentication in Microsoft Security Copilot
- Prompting in Microsoft Security Copilot

Security Copilot integration in Microsoft Surface Management Portal

If you use Microsoft Surface Management Portal in the same tenant as Security Copilot, then you can use Security Copilot to get insights from your Surface Management Portal data.

There are Surface Management Portal capabilities built into Security Copilot, and you can use prompts to get more information, including:

- Information about your device fleet, including summary on warranties, and end of service and support.

- Details about issues with specific devices, so you can troubleshoot more quickly and easily.
- Information to help you learn about safeguarding and improving the lifetime of your Surface devices.

This article shows you how to access your Microsoft Surface Management Portal data in Security Copilot.

Key features

Copilot for Surface Management Portal brings the capabilities of Security Copilot to the Microsoft Surface Management Portal admin center, enabling you to respond to Surface device issues and escalations quickly. Bringing AI to Surface Management Portal allows teams to immediately understand warranty coverage of their device fleet and determine remediation steps in a timely manner. As an administrator, you can compare the latest Surface for Business devices with their current Surface devices. You can also learn more about protecting your organizations devices from malware and other threats, so that you can reduce risks.

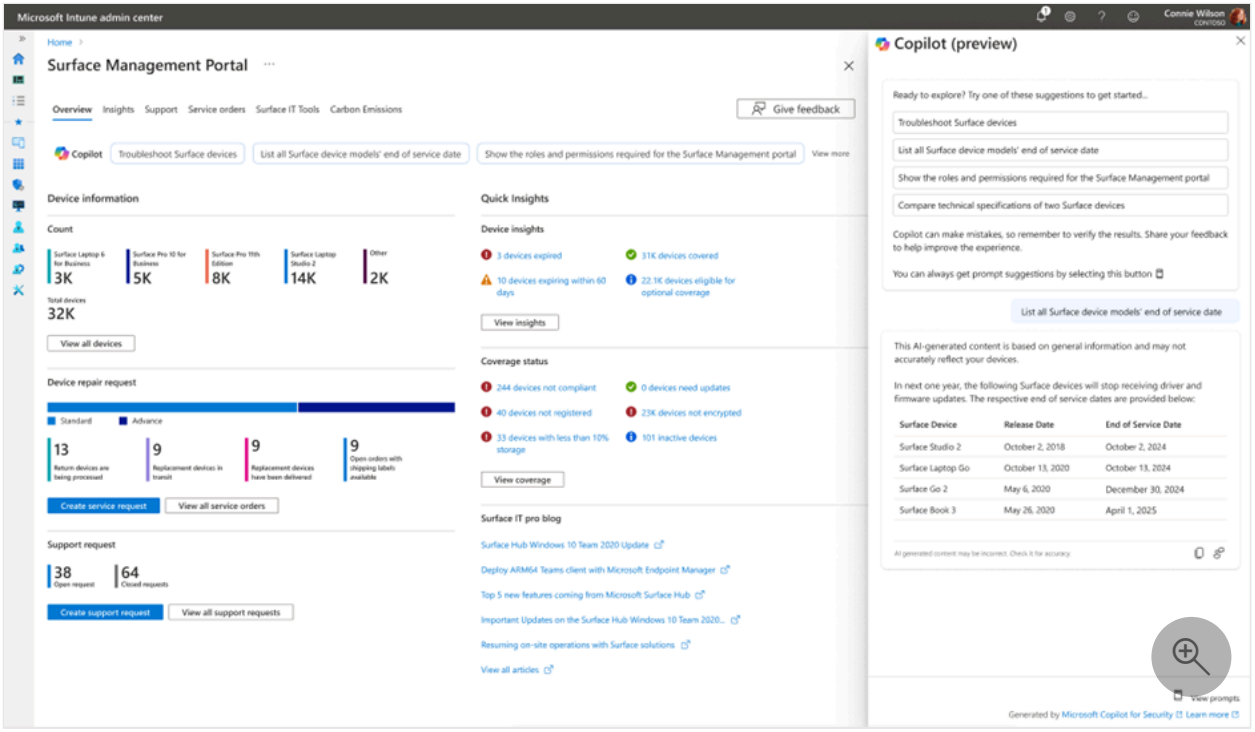

Enable the Security Copilot Integration in Surface Management Portal

Copilot for Surface Management Portal is included with Security Copilot. There aren't any other licensing requirements needed to use Copilot for Surface Management Portal. Ensure that you have provisioned sufficient **security compute units** to leverage this

capability. For more information, see Get started with Microsoft Security Copilot and Manage usage of security compute units in Security Copilot.

Enable integration of Security Copilot:

1. Sign in to the Security Copilot portal.
2. In the prompt bar, select **Sources** at the right corner.

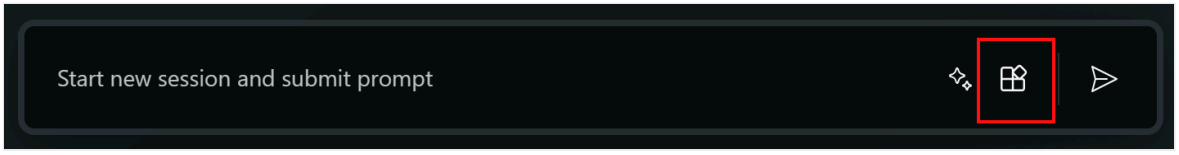

3. In **Manage Sources**, activate the **Surface Management Portal** plugin.

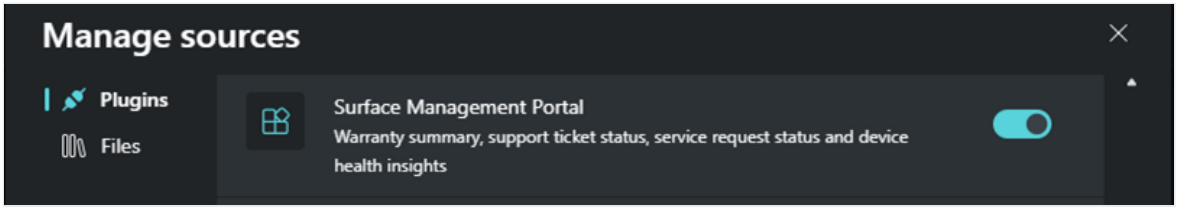

4. Close the plugin pane.

> **Note**
>
> Only certain roles can enable or disable plugins. For more information, see Manage plugins in Microsoft Security Copilot.

Sample Surface Management Portal prompts

Once you have set up Security Copilot, you can start using Copilot for Surface Management Portal by selecting one of the available prompts.

Sample prompts:

- Troubleshoot Surface devices.
- List all my Surface device models' end of service date.
- Compare technical specifications of two Surface devices.
- Generate a warranty coverage report for all devices.
- What are the best practices for protecting Surface devices?

> **Note**

These capabilities are available for public preview.

Provide feedback

Your feedback is vital to guide the current and planned development of the product. The best way to provide feedback is to use the feedback buttons at the bottom of each completed prompt, directly within the product.

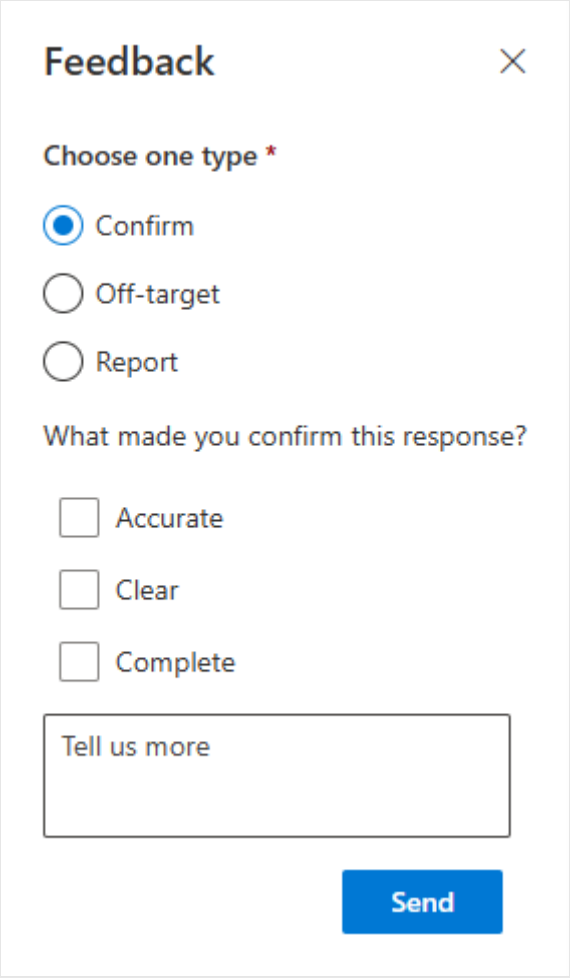

Privacy and data security in Security Copilot

To understand how Security Copilot handles your prompts and the output data that it retrieves from the service, see Privacy and data security in Microsoft Security Copilot.

Responsible AI FAQ in Microsoft Security Copilot

At Microsoft, we take our commitment to responsible AI seriously. Security Copilot is being developed in accordance with our AI principles. To learn more about how

Microsoft approaches responsible AI for Security Copilot, see Responsible AI FAQ.

Feedback

Was this page helpful? Yes No

Use Microsoft Security Copilot for threat intelligence

Applies to: Microsoft Defender XDR, Microsoft Sentinel in the Microsoft Defender portal

Microsoft Security Copilot is a cloud-based AI platform that provides a natural language copilot experience. It can help support security professionals in different scenarios, like incident response, threat hunting, and intelligence gathering. For more information about what it can do, read What is Microsoft Security Copilot?.

Security Copilot customers receive access to Microsoft Threat Intelligence for each authenticated Copilot user. To ensure that you have access to Copilot, see the Security Copilot purchase and licensing information.

Once you have access to Security Copilot, the key features discussed in this article become accessible in either the Security Copilot portal or the Microsoft Defender portal.

Know before you begin

If you're new to Security Copilot, you should familiarize yourself with it by reading these articles:

- What is Microsoft Security Copilot?
- Microsoft Security Copilot experiences
- Get started with Microsoft Security Copilot
- Understand authentication in Microsoft Security Copilot
- Prompting in Microsoft Security Copilot

Security Copilot integration with Microsoft Threat Intelligence

Security Copilot delivers information about threat actors, indicators of compromise (IOCs), tools, and vulnerabilities, as well as contextual threat intelligence from Microsoft Threat Intelligence. You can use the prompts and promptbooks to investigate incidents, enrich your hunting flows with threat intelligence information, or gain more knowledge about your organization's or the global threat landscape.

- Be clear and specific with your prompts. You might get better results if you include specific threat actor names or IOCs in your prompts. It might also help if you add **threat intelligence** to your prompt, like:
  - Show me threat intelligence data for Aqua Blizzard.
  - Summarize threat intelligence data for "malicious.com."
- Be specific when referencing an incident (for example, "incident ID 15324").
- Experiment with different prompts and variations to see what works best for your use case. Chat AI models vary, so iterate and refine your prompts based on the results you receive.
- Copilot saves your prompt sessions. To see the previous sessions, from the Security Copilot Home menu, go to **My sessions**.

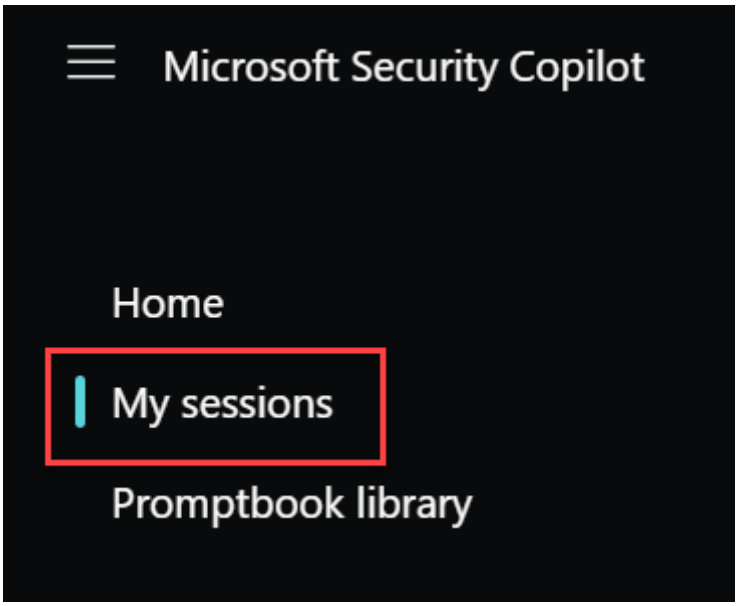

> **Note**
>
> For a walkthrough on Copilot, including the pin and share feature, read Navigate Microsoft Security Copilot.

Learn more about creating effective prompts

Key features

Security Copilot lets security teams understand, prioritize, and take action on threat intelligence information immediately.

You can ask about a threat actor, attack campaign, or any other threat intelligence that you want to know more about, and Copilot generates responses based on threat analytics reports, intel

profiles and articles, and other Microsoft Threat Intelligence content.

You can also select any of the built-in prompts that are available in the Defender portal to do the following actions:

- **Summarize** the latest threats related to your organization.
- **Prioritize** which threats to focus on based on your environment's highest exposure level.
- **Ask** about the threat actors targeting a specific industry.

Learn more about using Copilot in Defender for threat intelligence

Turn on the Security Copilot integration

1. Go to Microsoft Security Copilot and sign in with your credentials.

2. Make sure that the Microsoft Threat Intelligence plugin is turned on. In the prompt bar, select the **Sources** icon.

   In the **Manage sources** pop-up window that appears, under **Plugins**, confirm that the **Microsoft Threat Intelligence** toggle is turned on, then close the window.

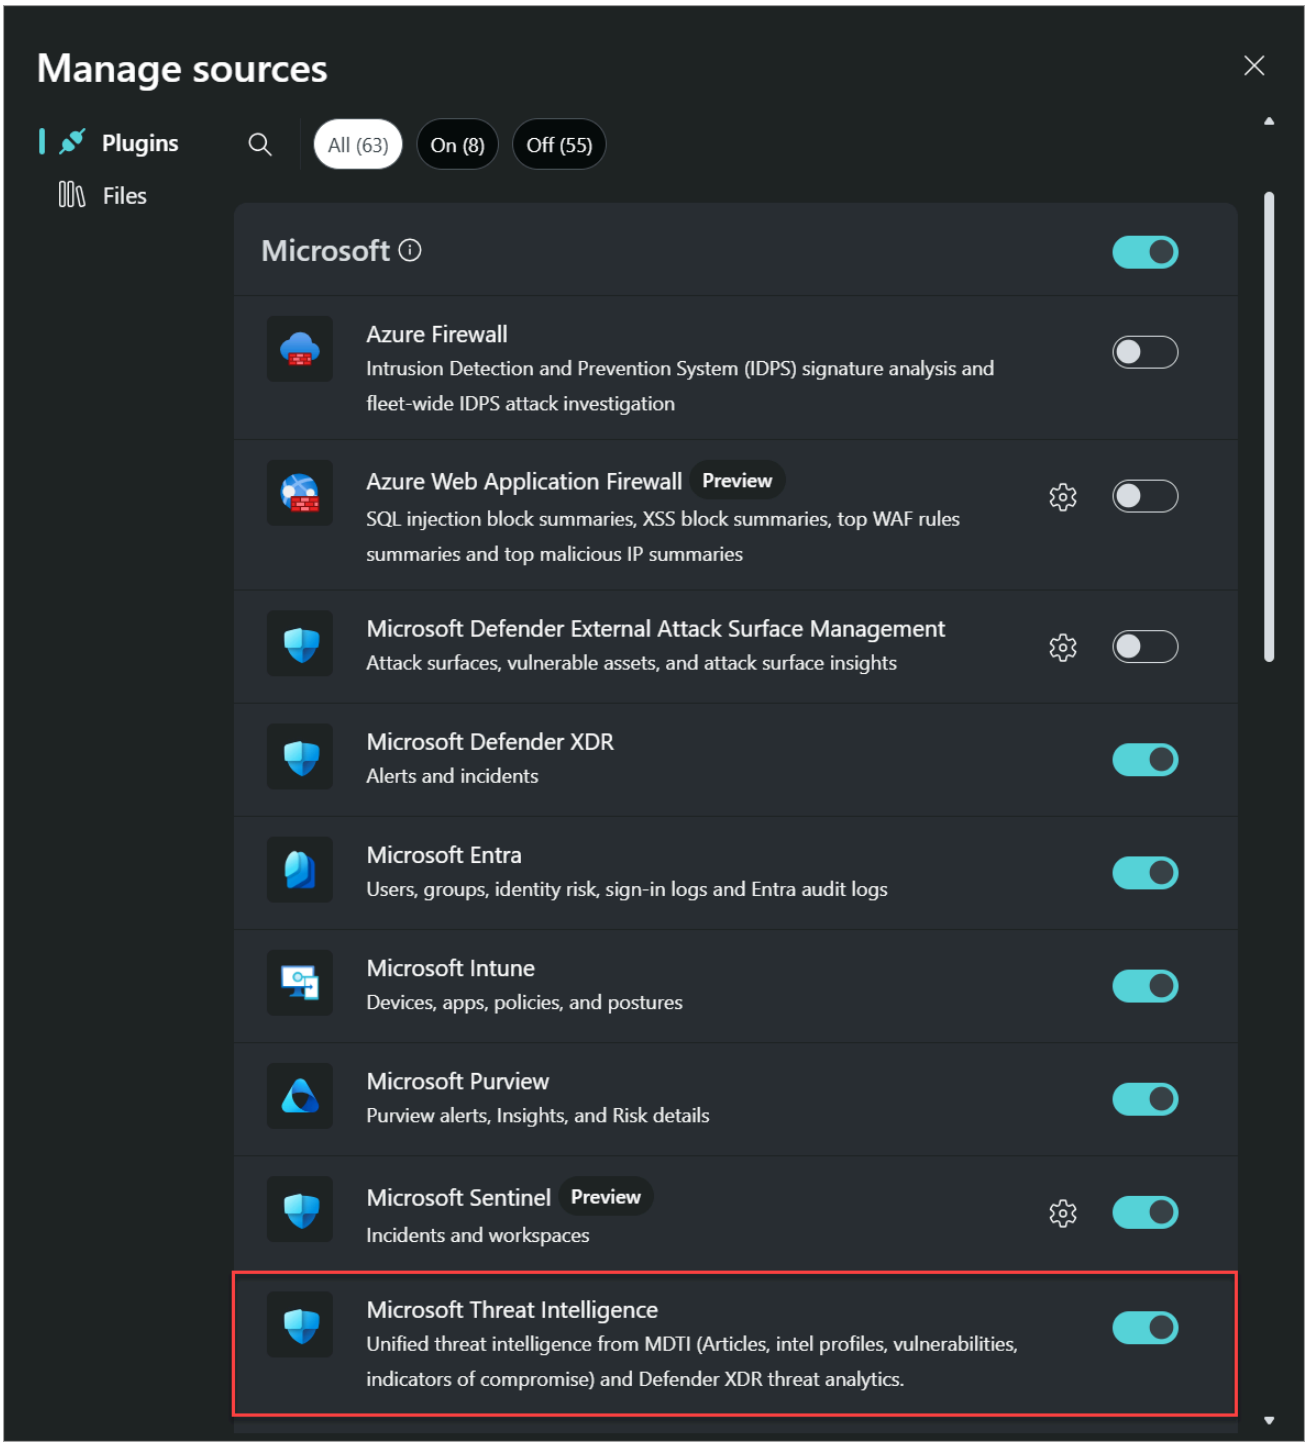

> **Note**
>
> Some roles can turn the toggle on or off for plugins like Microsoft Threat Intelligence. For more information, read Manage plugins in Microsoft Security Copilot.

3. Enter your prompt in the prompt bar.

Built-in system features

Security Copilot has built-in system features that can get data from the different plugins that are turned on.

To view the list of built-in system capabilities for Microsoft Threat Intelligence:

1. In the prompt bar, select the **Prompts** icon.

2. Select **See all system capabilities**. The *Microsoft Threat Intelligence* section lists all the available capabilities that you can use.

Copilot also has the following promptbooks that also deliver information from Microsoft Threat Intelligence:

- **Check impact of an external threat article** – Analyzes an external or third-party article to extract related IOCs, summarize the intelligence, and generate hunting queries so you can assess the potential impact of the threat reported in the article to your organization.
- **Threat actor profile** – Generates a report profiling a known threat actor, including suggestions to defend against their common tools and tactics.
- **Threat Intelligence 360 report based on MDTI article** – Analyzes a Microsoft Threat Intelligence article to extract related IOCs, summarize the intelligence, and generate hunting queries so you can assess the potential impact of the threat reported in the article to your organization.
- **Vulnerability impact assessment** – Generates a report summarizing the intelligence for a known vulnerability, including steps on how to address it.

To view these promptbooks, in the prompt bar, select the **Prompts** icon then select **See all promptbooks**.

Sample Microsoft Threat Intelligence prompts

You can use many prompts to get information from Microsoft Threat Intelligence. This section lists some ideas and examples.

General information about threat intelligence trends

Get threat intelligence from threat articles and threat actors.

Sample prompts :

- Summarize the recent threat intelligence.

- Show me the latest threat articles.
- Get threat articles related to ransomware in the last six months.

Threat actor mapping and infrastructure

Get information on threat actors and the tactics, techniques, and procedures (TTPs), sponsored states, industries, and IOCs associated with them.

Sample prompts:

- Tell me more about Silk Typhoon.
- Share the IOCs associated with Silk Typhoon.
- Share the TTPs associated with Silk Typhoon.
- Share threat actors associated with Russia.

Vulnerability data by CVE

Get contextual information and threat intelligence on Common Vulnerabilities and Exposures (CVEs), which are derived from Microsoft Threat Intelligence articles, threat analytics reports, and data from Microsoft Defender Vulnerability Management and Microsoft Defender External Attack Surface Management.

Sample prompts:

- Share the technologies that are susceptible to the vulnerability CVE-2021-44228.
- Summarize the vulnerability CVE-2021-44228.
- Show me the latest CVEs.
- Show me threat actors associated with CVE-2021-44228.
- Show me the threat articles associated with CVE-2021-44228.

Indicator data in relation to threat intelligence

Get detailed information about an indicator (for example, IP addresses, domains, and file hashes) based on the numerous data sets available in Microsoft Threat Intelligence, including reputation scores, WHOIS information, domain name system (DNS), host pairs, and certificates.

Sample prompts:

- What can you tell me about the domain <*domain name*>?
- Show me indicators related to <*domain name*>.
- Show me all resolutions for <*domain name*>.

- Show me host pairs related to *<domain name>*.
- Show me the reputation of the host *<host name>*.
- Show me all resolutions for IP address *<IP address>*.
- Show me the open services in *<IP address>*.

Malware information

Get detailed information about a specific malware detection, which is derived from the Microsoft Security Intelligence threat encyclopedia.

Sample prompt:

- What can you tell me about Ransom:Win32/WannaCrypt?

Provide feedback

Your feedback on the Microsoft Threat Intelligence integration in Security Copilot helps with development. To provide feedback, in Copilot, select **How's this response?** at the bottom of each completed prompt and choose any of the following options:

- **Looks right** - Select this button if the results are accurate, based on your assessment.
- **Needs improvement** - Select this button if any detail in the results is incorrect or incomplete, based on your assessment.
- **Inappropriate** - Select this button if the results contain questionable, ambiguous, or potentially harmful information.

For each feedback button, you can provide more information in the next dialog box that appears. Whenever possible, and when the result is **Needs improvement**, write a few words explaining what can be done to improve the outcome. If you entered prompts specific to threat intelligence and the results aren't related, then include that information.

Privacy and data security in Security Copilot

When you interact with Security Copilot to get Microsoft Threat Intelligence data, Copilot pulls that data from Microsoft Threat Intelligence. The prompts, the data retrieved, and the output shown in the prompt results are processed and stored within the Copilot service. Learn more about privacy and data security in Microsoft Security Copilot

Related content

- What is Microsoft Security Copilot?
- Privacy and data security in Microsoft Security Copilot
- Microsoft Threat Intelligence in Microsoft Defender XDR

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 08/02/2026

Security Copilot in Microsoft Purview Overview

Summarize this article for me

Microsoft Security Copilot is a cloud-based AI platform that can assist security and compliance professionals in protecting their organization's data. Security and compliance professionals can use Security Copilot to identify, summarize, triage, and remediate issues within the following Microsoft Purview solutions:

- Data Security Posture Agent (DSPM)
- Microsoft Purview Data Loss Prevention (DLP)
- Microsoft Purview Insider Risk Management
- Microsoft Purview Communication Compliance
- Microsoft Purview eDiscovery

For more information about what Security Copilot can do and the different scenarios it supports, read What is Microsoft Security Copilot?.

Know before you begin

If you're new to Security Copilot, you should familiarize yourself with it by reading these articles:

- What is Microsoft Security Copilot?
- Microsoft Security Copilot experiences
- Get started with Microsoft Security Copilot
- Understand authentication in Microsoft Security Copilot
- Prompting in Microsoft Security Copilot
- Configure Owner settings
- Learn about Security Copilot in Microsoft 365 E5

Security Copilot integration in Microsoft Purview

When you sign up for Security Copilot in the same tenant as Microsoft Purview, you can use both the Security Copilot embedded and standalone experiences. Security Copilot capabilities, like summarizing DLP or insider risk management alerts, are embedded into Microsoft Purview features.

Copilot in Microsoft Purview embedded experiences is a set of capabilities that are embedded in Microsoft Purview features. For more information, see standalone and embedded

experiences.

Copilot in Microsoft Purview standalone experience is a chat-like experience that you can use to ask questions and get answers about your data. For more information, see standalone and embedded experiences.

Key features in the embedded experience

You can open Security Copilot in Microsoft Purview by selecting the **Copilot** icon in the top navigation bar. It's available across all Microsoft Purview solutions.

The embedded experience in Purview can help you:

- *Get AI-generated summaries from Microsoft Purview product documentation*.
- *Use suggested or customized prompts in DSPM*. Learn more about the data security posture in your organization. For more information, see Use Microsoft Security Copilot with Data Security Posture Management.
- *Summarize alerts in DLP*. For more information on this and how to access Security Copilot in DLP, see Investigate a DLP alert.
- *Get insights into policies*. Security copilot can help you understand what your policies are doing in your organization, and where they're active. For more information, see Get Insights with Microsoft Security Copilot.
- *Summarize alerts in insider risk management*. For more information on this and how to access Security Copilot in Insider Risk Management, see Investigate insider risk management activities.
- *Summarize policy matches based on trainable classifiers for communication compliance*. Communication compliance also provides an interactive prompt experience to dig deeper into the summary. For more information on this and the steps to access Copilot in Communication Compliance, see Investigate and remediate communication compliance alerts.
- *Get contextual summary for eDiscovery cases*. For more information, see Group and view documents in a review set in eDiscovery (Premium).
- *Gain insights in activity explorer data and generate filters from natural language prompts (preview)*. For more information, see Microsoft Security Copilot in activity explorer (preview).

Key features in the standalone experience

The Security Copilot in Microsoft Purview standalone experience has many built in capabilities. You can use these capabilities to get insights from your Purview data and make connections between datapoints. This information can help you understand your information security and compliance posture and triage alerts.

- *DLP and Insider Risk Management data and user risk promptbook*. Security Copilot in Microsoft Purview analyzes data from DLP and insider risk management and, by running multiple, sequential prompts contained in a promptbook, presents you with integrated results. For more information, see Microsoft Copilot in Microsoft Purview prompts and promptbooks.

System capabilities of Security Copilot

In the standalone experience, there are built-in capabilities (prompts) that are available once the Microsoft Purview plugin is enabled.

Security Copilot in Microsoft Purview brings three types of capabilities:

- Summarize Microsoft Purview alerts.
- Triage Microsoft Purview alerts.
- Drill down into your Microsoft Purview data.

Enable the Microsoft Purview source in Microsoft Security Copilot

> **Important**
>
> Security Copilot in Purview must be enabled for both the standalone and embedded experiences to work.

Security Copilot in Purview is enabled by default. To enable or disable the Microsoft Purview source in Microsoft Security Copilot, follow these steps:

1. Ensure that you have permissions.
2. Open https://securitycopilot.microsoft.com/.
3. Open **Sources** in the prompt bar.

4. On the **Manage plugins** page, set the **Purview** toggle to **On** to enable or **Off** to disable.

Review the Microsoft Purview system capabilities

1. Select the capabilities control in the prompt bar.

2. Select **See all system capabilities** to see all the system capabilities that are available for Microsoft Purview. Here are a few:

    - Get Data Risk Summary
    - Get User Risk Summary
    - Summarize Purview Alert
    - Triage Purview Alerts
    - Zoom into Purview Data and User Risk

Sample prompts

For guidance on writing effective prompts, see Prompting in Microsoft Security Copilot. Here are some examples:

- *Show me the top five DLP alerts from the past 24 hours.*
- *Summarize the DLP alert with ID <12345>.*
- *What's the risk profile of the user that's associated with the DLP alert <12345>.*
- *Show me the Insider Risk Management alerts for <user> .*
- *What items did user <user> exfiltrate in the past 30 days.*

Provide feedback

Your feedback is vital to guide the current and planned development of the product. The best way to provide this feedback is directly in the product. Select **How's this response?** at the bottom of each completed prompt and choose any of the following options:

- **Looks right** - Select if the results are accurate, based on your assessment.
- **Needs improvement** - Select if any detail in the results is incorrect or incomplete, based on your assessment.
- **Inappropriate** - Select if the results contain questionable, ambiguous, or potentially harmful information.

For each feedback option, you can provide more information in the next dialog box that appears. Whenever possible, and especially when the result is **Needs improvement**, write a few words explaining what can be done to improve the outcome. If you entered prompts specific to Microsoft Purview and the results aren't related, then include that information.

Privacy and data security in Microsoft Security Copilot

To understand how Microsoft Security Copilot in Purview handles your prompts and the data that's retrieved from the service (prompt output), see the Privacy and data security in Microsoft Security Copilot.

Related articles

- Privacy and data security in Microsoft Security Copilot
- Microsoft Security Copilot Responsible AI FAQ

Last updated on 02/24/2026

Microsoft Security Copilot in Defender for Cloud

Microsoft Defender for Cloud integrates Microsoft Security Copilot and Microsoft Copilot for Azure into its experience. These integrations let you ask security-related questions, receive responses, and automatically trigger the necessary skills to analyze, summarize, remediate, and delegate recommendations by using natural language prompts.

Both Security Copilot and Copilot for Azure are cloud-based AI platforms that provide a natural language copilot experience. These platforms help security professionals understand the context and effect of recommendations, remediate or delegate tasks, and address misconfigurations in code.

Defender for Cloud's integration with Security Copilot and Copilot for Azure on the recommendations page lets you enhance your security posture and mitigate risks in your environments. It streamlines the process of understanding and implementing recommendations, making your security management more efficient and effective.

Learn more about Security Copilot

If you're new to Security Copilot, you should familiarize yourself with it by reading these articles:

- What is Microsoft Security Copilot?
- Microsoft Security Copilot experiences
- Get started with Microsoft Security Copilot
- Understand authentication in Microsoft Security Copilot
- Prompting in Microsoft Security Copilot

Security Copilot integration in Defender for Cloud

Defender for Cloud integrates Copilot directly into the Defender for Cloud experience. With this integration, you can analyze, summarize, remediate, and delegate your recommendations by using natural language prompts.

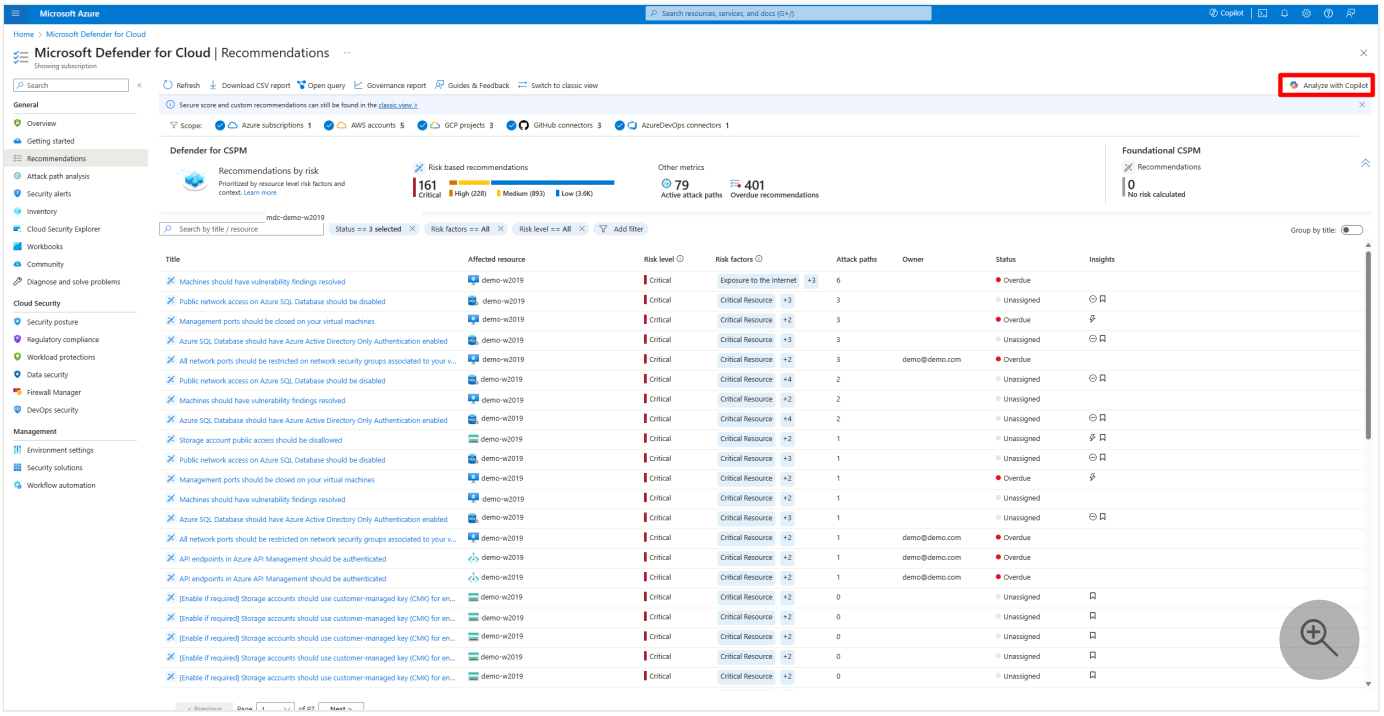

When you open Copilot, you can use natural language prompts to ask questions about the recommendations. Copilot provides responses in natural language, helping you understand the context of the recommendation. It explains the effect of implementing the recommendation and provides steps for implementation.

Some sample prompts include:

- Show critical risks for publicly exposed resources
- Show critical risks to sensitive data
- Show resources with high severity vulnerabilities

Copilot can assist with refining recommendations, providing summaries, remediation steps, and delegation. It enhances your ability to analyze and act on recommendations.

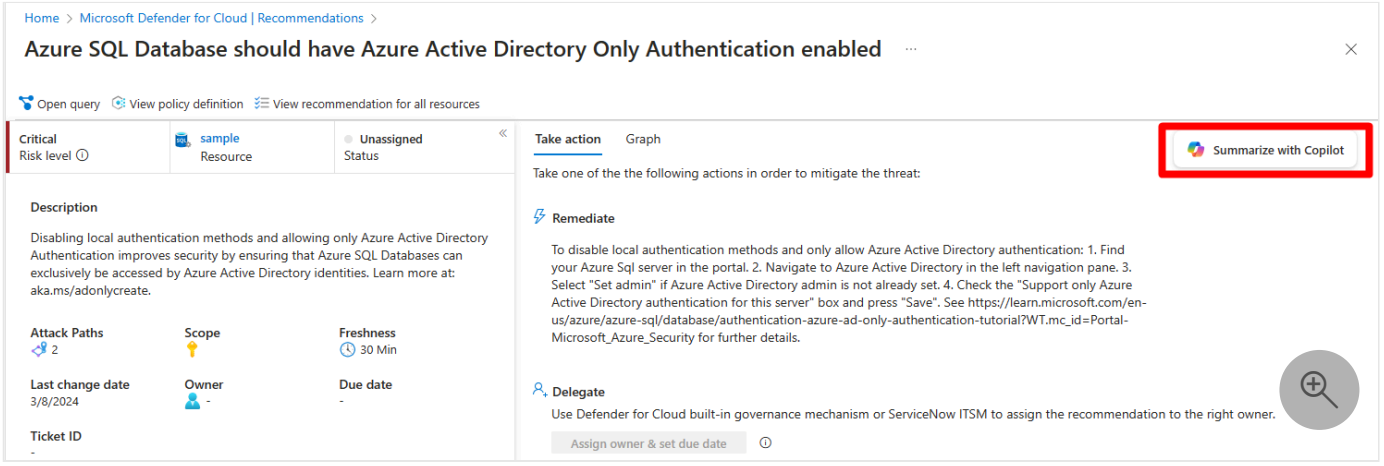

Key features

The following section provides information about the available features in Security Copilot in Defender for Cloud.

Data processing workflow

When you use Security Copilot in Defender for Cloud, the following data processing workflow takes place:

1. A user enters a prompt in the Copilot interface.
2. Copilot for Azure receives the prompt.
3. Copilot for Azure evaluates the prompt and the active page to determine the skills needed to resolve the prompt.
4. If the prompt is security related and the skill is available, Security Copilot executes the skills and sends back a response to Copilot in Azure for presentation.
5. If a security-related prompt is received but the skill is unavailable, Azure Copilot searches all of its available skills to find the most relevant skills to resolve the prompt. It then sends a response to the user.

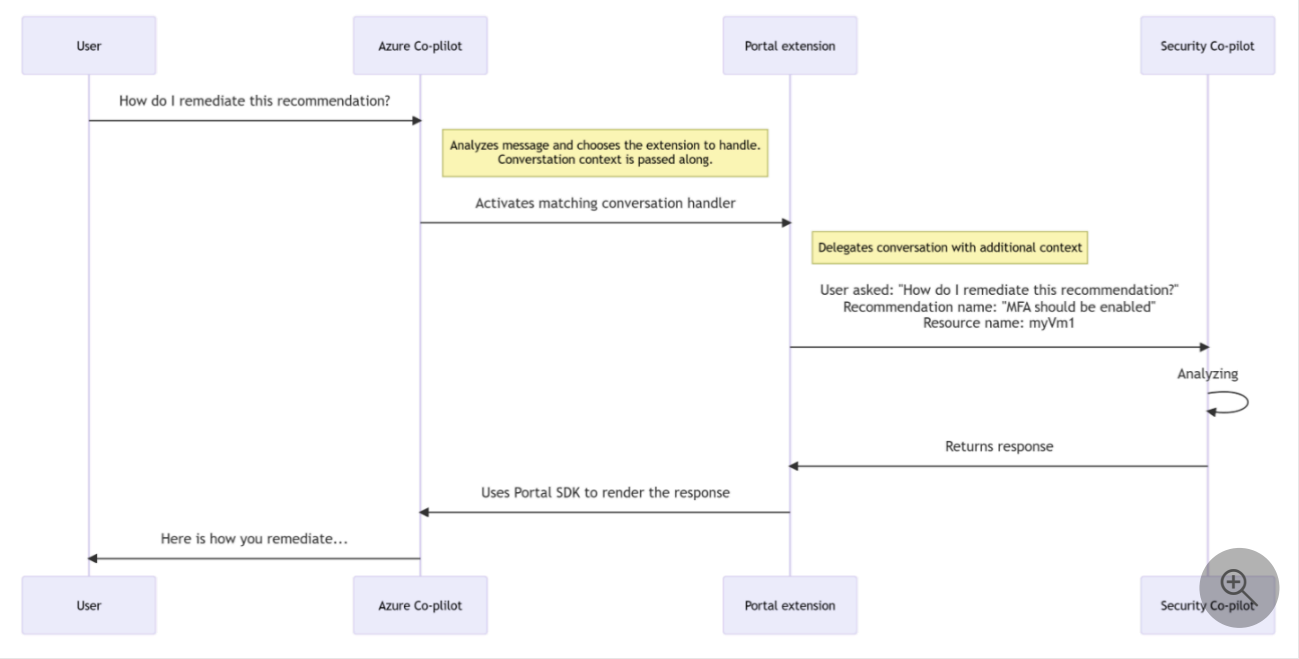

Check out the Security Copilot FAQs.

Enable the Security Copilot integration in Defender for Cloud

Security Copilot in Defender for Cloud doesn't rely on any of the available plans in Defender for Cloud. Security Copilot is available for all users when you:

1. Enable Defender for Cloud on your environment.
2. Have access to Azure Copilot.
3. Have Security Compute Units assigned for Security Copilot.

To enjoy the full range of Security Copilot's capabilities in Defender for Cloud, enable the Defender for Cloud Security Posture Management (DCSPM) plan on your environments. The DCSPM plan includes extra security. These features include Attack path analysis, Risk prioritization, and others. You can use Security Copilot to navigate and manage each feature. Without the DCSPM plan, you can still use Security Copilot in Defender for Cloud, but with limited capacity.

Monitor your usage

Security Copilot has a usage limit. When the usage in your organization nears its limit, you're notified when you submit a prompt. To avoid a disruption of service, contact the Azure capacity owner or contributor to increase the Security Compute Units (SCU) or limit the number of prompts.

Learn more about usage limits.

Provide feedback

Your feedback on the Defender for Cloud integration with Security Copilot helps with development. Provide feedback in Copilot, with the **How's this response** after each completed prompt, and select any of the following options:

- **Looks right** - Select this option if the results are accurate, based on your assessment.
- **Needs improvement** - Select this option if any detail in the results is incorrect or incomplete, based on your assessment.
- **Inappropriate** - Select this option if the results contain questionable, ambiguous, or potentially harmful information.

For each feedback option, you can provide more information in the next dialog box that appears. Whenever possible, and when the result is **Needs improvement**, write a few words explaining what can be done to improve the outcome. If you entered prompts specific to Defender for Cloud and the results aren't related, include that information.

Privacy and data security in Security Copilot

When you interact with Security Copilot to get Defender TI data, Copilot pulls that data from Defender for Cloud. The prompts, the data retrieved, and the output shown in the prompt results are processed and stored within the Copilot service.

Learn more about privacy and data security in Microsoft Security Copilot.

Related content

- Analyze recommendations with Microsoft Security Copilot
- Summarize recommendations with Microsoft Security Copilot
- Remediate recommendations with Microsoft Security Copilot
- Delegate recommendations with Microsoft Security Copilot
- Remediate code with Microsoft Security Copilot

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 06/16/2026

Non-Microsoft plugins for Microsoft Security Copilot

> **Important**
>
> Some information in this article relates to a prereleased product that may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

Microsoft Security Copilot supports many plugins, including Microsoft and non-Microsoft plugins. This article describes several examples of non-Microsoft plugins you can use with Security Copilot. Each plugin has a specific function. Some, but not all, require authentication to configure integration with Security Copilot.

> **Note**
>
> Microsoft will be processing and storing data from these plugins according to our Privacy and Data policies.

Find and use other plugins

1. Sign in to Microsoft Security Copilot.
2. Access **Manage plugins** by selecting the **Sources** button from the prompt bar.
3. Scroll down to **Non-Microsoft**, and select the plugin you want to use. Here are some examples:

    - AbuseIPDB
    - Aviatrix
    - Censys
    - CheckPhish
    - CIRCL Hash Lookup
    - CrowdSec CTI
    - CyberArk Privilege Cloud
    - Cybersixgill
    - Cyware Intel Exchange
    - Cyware Respond
    - Darktrace

- Forescout Risk and Exposure Management
- Forescout Vedere Labs
- GreyNoise Enterprise and GreyNoise Community
- HP Workforce Experience Platform
- Intel 471 Threat Intelligence
- IPGeolocation
- IPinfo
- Jamf
- Netskope
- OpenText Core Threat Detection and Response
- Pure Signal Scout
- Quest Security Guardian
- Red Canary
- ReversingLabs
- Saviynt
- ServiceNow SIR
- SGNL
- Shodan
- Silverfort
- Splunk
- Tanium
- TufinMate
- UrlScan
- Valence Security
- Whoisfreaks

See also

Manage plugins in Microsoft Security Copilot

Last updated on 03/18/2026

1Password

> **Important**
>
> Some information in this article relates to a prerelease product that may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

> **Note**
>
> This article contains information about third-party plugins. This guidance is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

The 1Password plugin for Microsoft Copilot for Security leverages audit data collected by Microsoft Sentinel to provide security teams with insights into 1Password usage, potential security events, and anomalous behaviors. This plugin does not retrieve secrets or credentials from 1Password, but rather analyzes the audit logs and events to enhance security monitoring and incident response capabilities.

Prerequisites

- 1Password Business or Enterprise subscription with audit logging capability.
- Microsoft Sentinel deployed and configured in your Azure environment.
- 1Password Events Reporting integration set up to send audit logs to Microsoft Sentinel.
- Microsoft Sentinel 1Password data connector configured and actively ingesting 1Password audit data.
- Administrative access to Microsoft Security Copilot.
- Linked Microsoft Sentinel workspace in Security Copilot.
- Familiarity with Kusto Query Language (KQL) for custom queries.

Know before you begin

Integration with Microsoft Security Copilot works by connecting to your Microsoft Sentinel workspace where 1Password events are stored.

You'll need to take the following steps before using the plugin.

Step 1: Configure 1Password Audit Logging

To enable audit logging in 1Password, follow the Events Reporting documentation or complete these steps:

1. Sign in to 1Password.com as an owner or administrator.
2. Click **Integrations** in the sidebar.
3. Click **Directory** at the top of the page.
4. Find the **Microsoft Sentinel** integration and select **Set Up**.
5. Enter a name for your integration (for example, "Microsoft Sentinel Connector").
6. Choose between:
    - **Send events from all vaults** to report events for your entire account.
    - **Choose vaults** to select specific vaults for event reporting.
7. Select **Add Integration**.
8. Save the bearer token that's displayed - you'll need this to configure the Microsoft Sentinel connector.
9. Follow the 1Password Sentinel integration guide to activate the serverless connector.
10. Configure the appropriate log retention policies.
11. Verify that user activities, authentication events, and admin actions are being logged.

For detailed configuration options and troubleshooting, refer to the Events Reporting setup guide.

Step 2: Set up Microsoft Sentinel Connector for 1Password

1. Access your Microsoft Sentinel workspace.
2. Navigate to the Data Connectors section.
3. Locate and select the 1Password data connector.
4. Follow the configuration wizard to connect to your 1Password environment.
5. Validate that audit logs are being successfully ingested into Microsoft Sentinel.

Step 3: Install and configure the plugin

1. Sign in to Microsoft Security Copilot.
2. Access **Manage Plugins** by selecting the **Sources** button from the prompt bar.
3. Next to **1Password**, select **Set up**.
4. Upload the 1Password plugin package or specify its repository location.
5. Configure the plugin to connect to your Microsoft Sentinel workspace.

Configuration guide

Required parameters

Configure the following parameters for the plugin to function correctly:

1. **Microsoft Entra Tenant ID**
    - The unique identifier of your Microsoft Entra environment
    - Format: `xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx`
2. **Subscription ID**
    - The unique identifier of the Azure Subscription of Microsoft Sentinel
    - Format: `xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx`
3. **ResourceGroup Name**
    - This is the resource group name that security copilot will use for sentinel.
4. **Workspace Name**
    - This is the workspace name that security copilot will use for sentinel.

Configuration methods

1. **Plugin Configuration Portal**
    - Configure settings directly in the Microsoft Copilot for Security plugin interface
    - Specify workspace connection parameters
    - Test connectivity before saving

Sample 1Password prompts

The following section provides example prompts to try.

1. Natural Language Queries

Use conversational language to analyze 1Password audit data:

- "Show me all failed login attempts in 1Password over the past week"
- "Who accessed sensitive vaults in 1Password outside of business hours?"
- "Has anyone created new shared vaults in the last month?"

2. Structured Commands

For more precise control, use structured command syntax:

- `analyze-events -type:"item.view" -timeframe:"last 7 days" -user:"admin"`
- `detect-anomalies -dataset:"authentication" -baseline:"30 days"`
- `report-activity -vault:"Finance" -action:"create,delete,modify"`

3. Workflow Integration

Include 1Password audit analysis as part of larger security workflows:

- "Investigate this user account and check for unusual 1Password activity"
- "Analyze login patterns across our SSO providers including 1Password"
- "Generate a compliance report showing all access to regulated data in 1Password"

Response Formatting

The plugin formats responses to provide clear security insights:

- **Event Timelines**: Chronological view of related security events
- **User Activity Profiles**: Aggregated view of user behaviors
- **Anomaly Highlighting**: Clear identification of outlier events
- **Visual Analytics**: Charts and graphs for pattern recognition

Advanced usage

Security Incident Investigation

The plugin enables advanced investigation of security incidents involving 1Password usage:

Example Investigation Workflow

1. Receive alert about suspicious access to sensitive vault
2. Query 1Password audit logs for all actions by the flagged user
3. Analyze access patterns and compare to historical baseline
4. Correlate with other security telemetry (network logs, endpoint data)
5. Generate comprehensive incident timeline and recommended actions

Troubleshoot the 1Password plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*. Make sure the plugin is turned on. This error may occur if the lookback period is too long, causing the query to attempt to retrieve an excessive amount of data. If the issue persists, sign out of Copilot for Security, and then sign back in.

Data Ingestion Issues

Expand table

Issue	Possible Causes	Resolution Steps
Missing Audit Data	Connector misconfiguration, Ingestion pipeline issues	1. Verify Microsoft Sentinel connector status 2. Check 1Password audit logging settings 3. Review data connector logs 4. Validate Log Analytics agent functionality
Data Delay	Ingestion latency, High data volume, Processing bottlenecks	1. Check ingestion pipeline status 2. Review Microsoft Sentinel health metrics 3. Optimize data collection rules 4. Consider dedicated capacity for critical data

Query Issues

Expand table

Issue	Possible Causes	Resolution Steps
Query Timeout	Complex queries, Large data volume, Resource constraints	1. Optimize query complexity 2. Add appropriate filters 3. Use time-based partitioning 4. Consider materialized views
Schema Mismatch	Custom field mappings, Schema evolution, Parser errors	1. Review schema definitions 2. Update field mappings 3. Check for data format changes 4. Update custom parsers if needed

Plugin Access Issues

Expand table

Feature	Common Problems	Troubleshooting Steps
Workspace Access	Permission issues, Misconfigured workspace ID	1. Verify workspace access permissions 2. Check workspace connection string 3. Review Microsoft Copilot for Security configuration 4. Validate Microsoft Entra ID permissions
KQL Functions	Function registration failure, Syntax errors, Execution timeout	1. Check function registration status 2. Validate KQL syntax 3. Review function permissions 4. Optimize function logic

Performance Optimization

- Use efficient KQL patterns and avoid cross-joins on large datasets
- Implement appropriate time filters to limit data processing
- Consider materialized views for frequently run analytical queries
- Monitor query performance and optimize resource-intensive operations

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Provide feedback

To provide feedback, contact https://support.1password.com/.

See also

Other plugins for Microsoft Copilot for Security Manage plugins in Microsoft Copilot for Security

Last updated on 03/18/2026

AbuseIPDB

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

AbuseIPDB is a project managed by Marathon Studios Inc. Their mission is to help make the Web safer by providing a central repository for Webmasters, System Administrators, and other interested parties to report and identify IP addresses that have been associated with malicious activity online. You can use the AbuseIPDB plugin with Microsoft Security Copilot.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot works with an API key. You'll need to take the following steps before using the plugin.

1. Get your AbuseIPDB API key. If you don't have one yet, follow these steps:
   a. Go to the AbuseIPDB website.
   b. Select on the **Sign Up** button, located at the top right corner of the page.
   c. Fill in the required information, such as your email address, username, and password.
   d. Complete any other steps for verification, if necessary.
   e. Once registered, sign-in to your AbuseIPDB account.

2. Access the API key.
   a. After logging in, go to the AbuseIPDB Homepage page.
   b. Navigate to the API tab and select on **Create key**.
   c. After clicking on Create Key a modal appears. You'll be asked to input a name; you can pick any name for your key. For example, enter "developer_key" and select the create button.
   d. The newly generated API key should populate on the same API tab. You need this same key-in later steps for the connector.

3. Sign in to Microsoft Security Copilot.
4. Access **Manage Plugins** by selecting the **Sources** button from the prompt bar.
5. Next to **AbuseIPDB**, select **Set up**.
6. In the AbuseIPDB plugin settings pane, in the **Value** field, paste your API Key, and then select **Save**.

Sample AbuseIPDB prompts

After the AbuseIPDB plugin is configured, you can use the following capabilities with Security Copilot.

The following table provides examples you can try:

Expand table

Capability	Input	Example prompts
`CheckIPAddress` Accepts an IP address (v4 or v6) and provides information about the queried IP, including version, origin country/region, usage type, ISP, and domain. Abusive reports are included.	Required: - `ipAddress` Optional: - `maxAgeInDays` (default 30, min 1, max 365) - `verbose`	- Tell me about IP address 118.25.6.39 using AbuseIPDB database - What does the abuseipdb database say about the IP address 180.126.219.127? - I'm curious about any abuseipdb records for the IP address 180.126.219.127. Can you look that up for me for the last 10 days?
`CheckNetworkBlock` This capability takes a subnet in CIDR notation (v4 or v6) and returns details about the queried network, including its netmask, possible hosts count, and address space description. URL-encode the network due to forward slashes in CIDR notation.	Required: - `network` Optional: - `maxAgeInDays` (default 30, min 1, max 365)	- Check the network block 1.1.1.1/24 for any reported IP addresses within the past 10 days. - I'm trying to find out if there are any reported IP addresses on the network block 1.1.1.1/24 in the past 10 days. Can you assist me with that? - Could you help me check if any IP addresses are

Capability	Input	Example prompts
		reported on the network block 1.1.1.1/24 within the last 10 days?
`ListBlockListedIpAddresses` This capability returns a list IP addresses most reported by AbuseIPDB users with or without other filters. Entries include IP, abuse score, and last report timestamp, sorted by score and timestamp.	Optional: - `confidenceMinimum` (default 100, min 25, max 100) - `limit` (default 10,000, min 1, restrictions) - `plaintext` - `onlyCountries` - `exceptCountries` - `ipVersion` (default 4,6 mixed)	- List blocklisted IP addresses from AbuseIPDB - According to AbuseIPDB, are there any blocklisted IP addresses from China?
`ReportIPAddress` Reporting IP addresses is the core feature of AbuseIPDB. To report an IP address, use this skill. It will return the updated abuseConfidenceScore for the IP address we specify.	Required: - `ip` - `categories` (restrictions 30, these nos. belong to a mapping) Optional: - `comment` - `timestamp`	- Could you report the IP address 180.126.219.126 on AbuseIPDB for category 18 and add a comment "This category is separate from DDoS attacks."? - Submit the IP address 180.126.219.126 to AbuseIPDB, tagging it as category 1,2. - Make an entry for the IP address 180.126.219.126 in AbuseIPDB, designating it as category 18.

Provide feedback

To provide feedback, contact AbuseIPDB.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Last updated on 01/15/2026

Aviatrix

08/25/2025

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

Aviatrix and Microsoft Security Copilot have partnered to bring together an AI enabled plugin to allow customers to leverage Microsoft Defender Threat Intelligence with Aviatrix to gain insight into new threats and mitigate them through firewall policy enforcement. Integration with Aviatrix across Microsoft Security Copilot, Microsoft Defender Threat Intelligence, and Microsoft Sentinel are outlined and enabled through Microsoft LogicApps in this document.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot works with an API key. You need to take the following steps before using the plugin.

1. Sign in to Microsoft Security Copilot.
2. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.
3. Next to **Aviatrix**, select **Set up**.
4. Enter your Aviatrix controller username and password.
5. Save your changes.

Sample Aviatrix prompts

After the Aviatrix plugin is configured, you can type "Aviatrix" in your Security Copilot prompt bar, followed by an action. The following list provides example prompts to try:

- `Aviatrix show me my firewall policies.`
- `Aviatrix show me log data for x.x.x.x.`
- `Aviatrix show me all blocked ports.`
- `Aviatrix show me all allowed rules.`

Aviatrix allows you to gain insight by allowing users to easily query our detailed resource, firewall, and log data. Be creative with your prompts using natural language.

Troubleshoot the Aviatrix plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*. Make sure the plugin is turned on. This error may occur if the lookback period is too long, causing the query to attempt to retrieve an excessive amount of data. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Provide feedback

To provide feedback, contact Aviatrix.

See also

Other plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Censys

> **Important**
>
> Some information in this article relates to a prereleased product that can be substantially modified before commercial release. Microsoft makes no warranties, express or implied, with respect to the information provided here.

This plugin allows Security Copilot users to enrich investigations using threat intelligence from the Censys Platform. Currently the following functionalities are supported:

- Generating and executing Censys queries using natural language
- Searching the Censys Global Index
- Aggregating search results by specific fields
- Retrieving detailed information about IP addresses
- Looking up certificate information by SHA256 hash
- Getting web property information by domain and port

Prerequisites

- Access to the Censys Platform account
- A Personal Access Token (PAT) from Censys
- Your Censys Organization ID

> **Note**
>
> This article contains information about non-Microsoft plugins. This guidance is provided to help complete integration scenarios. However, Microsoft doesn't provide troubleshooting support for non-Microsoft plugins. Contact the vendor for support.

Know before you begin

Integration with Security Copilot works with a combination of a Personal Access Token (PAT) and Organization ID. You need to take the following steps before using the plugin.

Authentication Setup

1. Sign in to your Censys Account.

2. Navigate to the **Personal Access Tokens** tab and select **Create New Token**.
3. Give your token a descriptive name (for example, "Security Copilot Token") and select **Create**.
4. Locate your Organization ID from the Censys Platform URL (appears as `?org=ORG_ID`).
5. Combine your PAT and Organization ID with a colon separator: `PAT:ORG_ID`

Plugin Configuration

1. When you're asked to set up authentication, enter the combined PAT and Organization ID in the following format into the **Value** field:

```plain
PAT:ORG_ID
```

2. Select **Save** to complete setup.

Available Skills

The Censys Plugin for Microsoft Security Copilot exposes the following skills:

- Query Generation
  - Creating Censys queries from natural language
  - Converting user intent into CenQL syntax
- Global Index Search
  - Executing queries against the Censys database
  - Retrieving and formatting search results
- Data Aggregation
  - Aggregating results by specified fields
  - Generating statistical summaries
- Asset Information
  - IP address lookups
  - Certificate information retrieval
  - Web property analysis

With the Censys plugin for Microsoft Security Copilot, you can invoke interactions with Censys in the context of a natural conversation. Here's an example:

1. A user can ask about specific IP addresses or domains of interest

2. The user can then use follow-up prompts such as "What other assets are associated with this IP?"
3. The user can further investigate by asking for aggregated data about similar assets or certificates

Example Prompts

Expand table

Skill	Prompt
Generate Query	`Generate a Censys query to find all the nginx servers`
Search Global Index	`Query Censys for host.location.country: "Netherlands"`
Aggregate Results	`What are the top 10 ASNs in the world?`
IP Lookup	`What does Censys know about IP 1.1.1.1?`
Multiple IP Lookup	`Get information about 1.1.1.1 and 2.2.2.2`
Certificate Lookup	`Get the certificate 14e641e9f3eb7cf4bdf6b0bc364bd8529feafdb78904c45a49ce25d7daecc824`
Web Property Lookup	`Get the web property google.com:443`
Multiple Properties	`Look up google.com:443 and censys.io:443`

Microsoft Security Copilot maintains context between queries, allowing for natural conversation flows. For example, after looking up an IP address, you can ask follow-up questions about related certificates or web properties without needing to specify the IP again.

Troubleshoot the Censys plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*, check:

1. Ensure the plugin is turned on
2. Verify your PAT and Organization ID are correctly formatted

3. Check that your PAT hasn't expired
4. Confirm you have sufficient API credits in your Censys account
5. If issues persist, sign out of Security Copilot and sign back in

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set:

- Ensure you're using supported prompt formats
- Try using direct skill invocation with "/"
- Check if other plugins might be intercepting similar queries

Provide feedback

To provide feedback, contact the Censys team.

See also

Other plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Last updated on 02/19/2026

CheckPhish

08/25/2025

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

This plugin allows users to utilize the CheckPhish AI to analyze URLs for potential phishing threats, tech support scams, cryptojacking, and other security risks. Utilizing CheckPhish's advanced machine learning models, this plugin provides detailed insights including:

1. The likelihood of a URL being malicious
2. Classification of the threat type
3. Timestamps for when the URL was last scanned
4. Detailed threat analysis
5. Historical data on URLs to track behavior over time

For more information, see CheckPhish.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot works with an API key. You need to take the following steps before using the plugin.

1. Get your CheckPhish API key. If you don't have one yet, follow these steps:
   a. Registering on the CheckPhish platform.
   b. Find your API key in your user profile.
   c. Copy the API key.

2. Sign in to Microsoft Security Copilot.

3. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.

4. Next to **CheckPhish**, select **Set up**.

5. In the **Value** field, paste your API Key, and then select **Save**.

Sample CheckPhish prompts

After the CheckPhish plugin is configured, you can use the following capabilities with Security Copilot. The following table lists example prompts to try.

Expand table

Capability	Example prompts
URL Scan - Analyze a URL for potential threats Required input: - URL	- `Ask CheckPhish to check [URL]` - `What does CheckPhish say about [URL]?`

Troubleshoot the CheckPhish plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*. Make sure the plugin is turned on. This error may occur if the lookback period is too long, causing the query to attempt to retrieve an excessive amount of data. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Provide feedback

To provide feedback, contact CheckPhish.

See also

Other plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

CIRCL Hash Lookup

08/25/2025

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

The Computer Incident Response Center Luxembourg (CIRCL) is a government-driven initiative designed to provide a systematic response facility to computer security threats and incidents. CIRCL Hash Lookup integration enables you to validate suspicious files in the form of hashes, either MD5, SHA-1, or SHA-256, in Security Copilot. You can use this plugin to get information about a file and verify whether it's allowlisted or blocklisted by trusted security platforms.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

You don't have to sign up for an account to use CIRCL Hash Lookup. You also don't need an API key. You'll need to take the following steps before using the plugin.

1. Sign in to Microsoft Security Copilot.
2. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.
3. Verify that **CIRCL Hash Lookup** is turned on.

Sample CIRCL Hash Lookup prompts

Try one or more of the capabilities listed in the following table:

Expand table

Capability	Example prompts
`lookupMD5HashCIRCL` This endpoint allows you to lookup information on CIRCL about a specific MD5 hash Required input: hash	- `lookup this MD5 hash using CIRCL` - `lookup information on this MD5 hash using CIRCL`
`lookupSHA1HashCIRCL` This endpoint allows you to lookup information on CIRCL about a specific SHA-1 hash Required input: hash	- `lookup this SHA-1 hash using CIRCL` - `lookup information on this SHA-1 hash using CIRCL`
`lookupSHA256HashCIRCL` This endpoint allows you to lookup information on CIRCL about a specific SHA-256 hash Required input: hash	- `lookup this SHA-256 hash using CIRCL` - `lookup information on this SHA-256 hash using CIRCL` Alternatively, you can type the name of a specific capability, like `lookupMD5HashCIRCL` in the prompt bar, instead of using one of the example prompts.

Troubleshoot the CIRCL Hash Lookup plugin

Errors occur

If you get errors, such as *Couldn't complete your request*, or *An unknown error occurred*, make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts are not invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use. For example, if you're using multiple sources that provide hash information like CIRCL, there might be conflicts. You can either use the product name `CIRCL` in your prompts, or type the name of a specific capability, like `lookupMD5HashCIRCL` instead.

Provide feedback

To provide feedback, contact CIRCL.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

circl.lu

CrowdSec Cyber Threat Intelligence

08/25/2025

CrowdSec Threat Intelligence is an open-source, collaborative security stack that enables you to analyze behaviors, respond to attacks, and share signals across the community. CrowdSec Threat Intelligence provides information about IP addresses and verification or identification of potentially aggressive IP addresses. You can use the CrowdSec Cyber Threat Intelligence (CrowdSec CTI) plugin with Microsoft Security Copilot.

This plugin allows users to enhance their IP investigations with threat intelligence sourced from CrowdSec and get insights such as:

- Curated IP and IP range reputations
- Background noise level assessment
- Detailed records of malicious behaviors
- MITRE techniques associated with the IP
- Countries/Regions targeted by the attacker
- Classification of the attacker
- Historical activity and aggressiveness metrics (covering the past 24h, 7 days, 30 days, and overall)

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot works with an API key. You'll need to take the following steps before using the plugin.

> **Note**
>
> Depending on which account you have, you might have a limit of up to 50 queries per day. It depends on your licensing for CrowdSec.

1. Get your CrowdSec API key. If you don't have one yet, follow these steps:

    a. Go to the CrowdSec website and create your free account.

b. In your personal account settings, go to **API Keys** and select **+ New key**. You can follow the [steps here] (https://doc.crowdsec.net/u/cti_api/integration_securitycopilot/)

2. Sign in to Microsoft Security Copilot .

3. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.

4. Next to **CrowdSec Threat Intelligence**, select **Set up**.

5. In the **Value** field, paste your API Key, and then select **Save**.

Sample CrowdSec CTI prompts

After the CrowdSec CTI plugin is configured, you can use it by taking one of the following steps:

- Access the capability directly by typing `LookupIpAddressSmokeDataset` in the prompt bar; or
- Prompt Security Copilot to use the CrowdSec Threat Intelligence API on an IP address

The following table summarizes how this capability works.

Expand table

Capability	What it does
`LookupIpAddressSmokeDataset` Example Prompt(s): - What can CrowdSec tell me about this IP: [IP] - According to CrowdSec what are the top targeted countries/regions by this IP: [IP] - Inputs: [IP] Required Input: IP Address	Searches CrowdSec's dataset for an IP address to know more about: - What it does in terms of observed behaviors, targeted protocols, and exploited vulnerabilities. - In what categories it belongs, such as proxy/VPN, CDN exit node, and Legit security scanner. - What it targets, in terms of countries/regions or services. - Existing cross-references, such as lists - How virulent it is. - For how long it's been reported by users. - The confidence level of the information.

Troubleshoot the CTI plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*, make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use. You can either use the product name `CrowdSec` in your prompts, or type the name of a specific capability, like `LookupIpAddressSmokeDataset` instead.

Provide feedback

To provide feedback, contact CrowdSec through Discourse or using the support or feedback action directly in your CrowdSec Console.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

CyberArk Privilege Cloud

08/25/2025

CyberArk Privilege Cloud is a SaaS solution that provides a simplified path to securely store, rotate and isolate credentials, monitor sessions, and quickly deliver scalable risk reduction to the business. You can use the CyberArk plugin with Microsoft Security Copilot to:

- List privileged accounts in Security Copilot using natural language.
- Leverage Copilot generative AI abilities to interact with privileged accounts data.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft doesn't provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot requires an OAuth service user account for authentication. You'll need to take the following steps before using the plugin.

1. Sign in to your CyberArk Identity Security Platform Shared Services (ISPSS) with a user account that has permissions to manage users in CyberArk Identity.

2. Configure a new OAuth service user account by following these steps:

   a. From the service picker, choose **Identity Administration**.

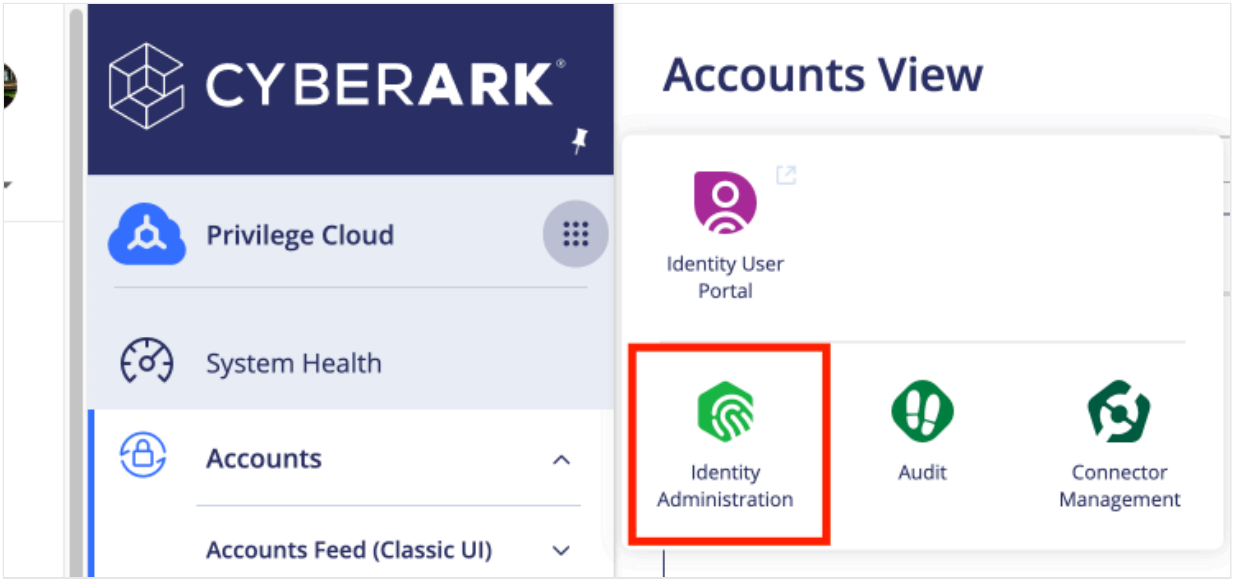

   b. Expand **Core Services**, and then select **Users**.

c. In the **Users** screen, select **Add User**, and then fill in the required details.

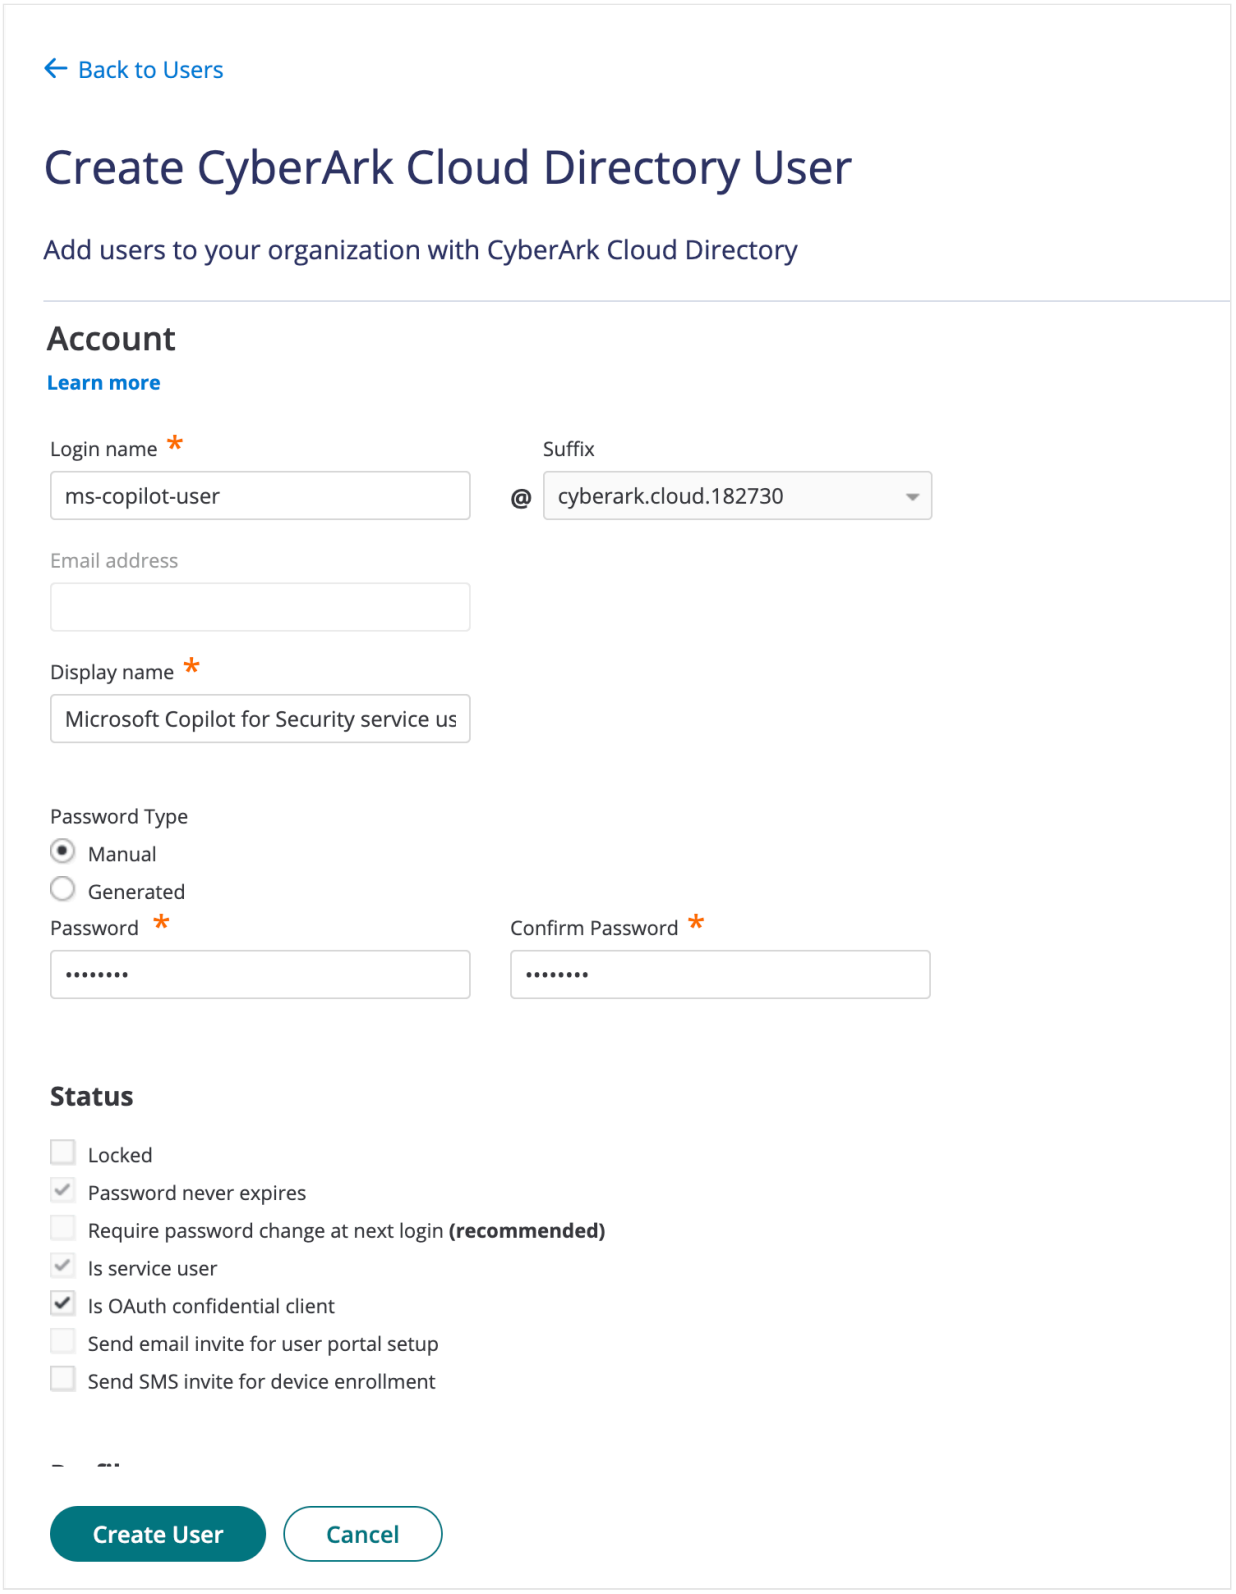

d. Under **Status** mark the checkbox next to **Is OAuth confidential client**, and then select **Create User**.

e. In the menu, select **Roles**, and then assign the new user account to a *Privilege Cloud Users* role.

3. Configure CyberArk Privilege Cloud to work with the Security Copilot plugin. Follow these steps:

a. In the service picker, select **Privilege Cloud**.

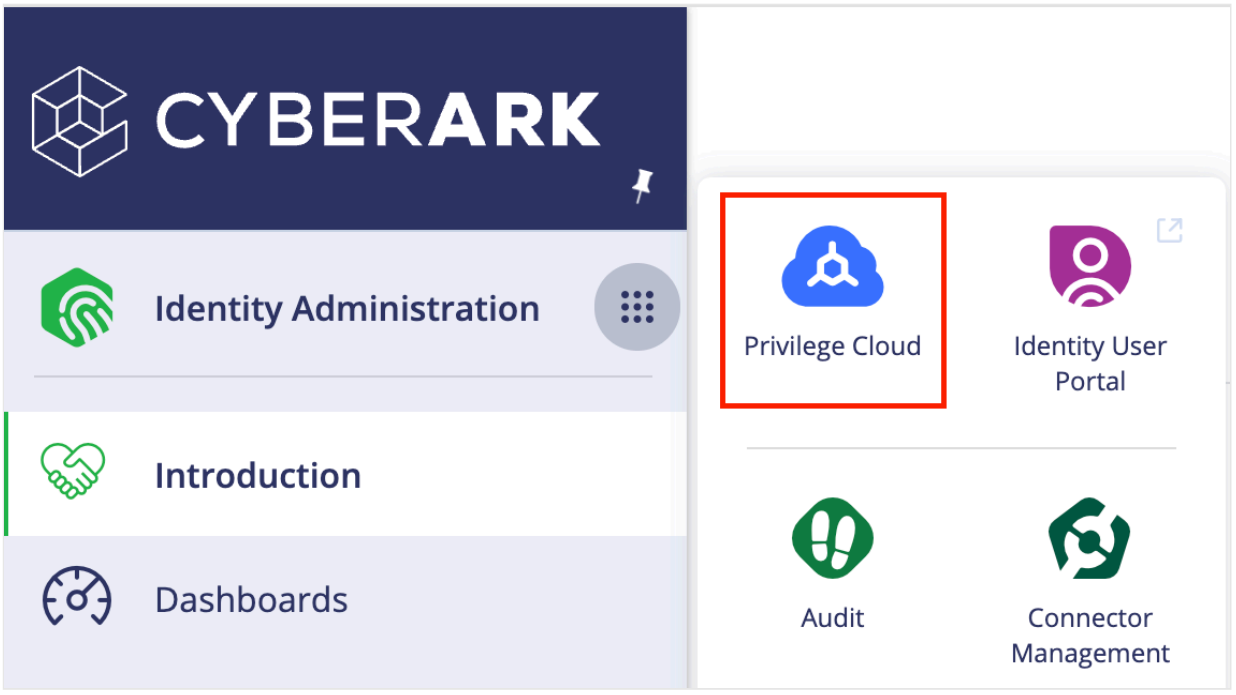

b. In the menu, expand **Policies**, select **Safes**, and then choose the safe you want to grant permissions for.

c. In the **Safe** pane, select the **Members** tab. Select **Add members**, and select the user account you created earlier, and then select **Next**.

d. Assign appropriate permissions, and then select **Add**.

4. Sign in to Microsoft Security Copilot.

5. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.

6. Next to **CyberArk**, select the toggle to enable it.

7. Provide your CyberArk Privilege Cloud instance URL and user account information.

8. Save your changes.

Sample CyberArk prompts

After the CyberArk plugin is configured, you can use it by typing `CyberArk` in your Security Copilot prompt bar, followed by an action. You can use the plugin to list privileged accounts and get account details. The following table lists some example prompts to try and their expected output:

Expand table

Prompt example	Expected output
`Find an account with the username test_user`	Get all account details about an account with the username “test_user”.
`Get all account details about an account with the username “test_user”`	Get a positive or negative answer whether an account with the username “test_user” and an address “10.0.1.103” is managed by CyberArk.
`Get all accounts and show the results in a table`	Get details about all accounts, displayed in a table format.

Provide feedback

To provide feedback, contact CyberArk.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Cybersixgill

08/25/2025

Cybersixgill offers real-time threat intelligence solutions to help security teams detect and respond to imminent threats from the clear, deep, and dark web.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft doesn't provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Prerequisites

- The plugin requires the use of an API key that can be obtained from Cybersixgill developer portal.
- Once logged into the developer portal, obtain and safely copy your API client ID and secret.

> **Note**
>
> If you aren't a Cybersixgill customer, send email to cfs@cybersixgill.com to get your three-month free trial.

Know before you begin

You need to take the following steps to use the plugin.

1. Sign in to Microsoft Security Copilot.
2. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.
3. Select **Settings** to configure your **Cybersixgill** plugin.
4. Select **Save** and begin prompting in Security Copilot.

Sample Cybersixgill prompts

After the Cybersixgill plugin is configured, you can try one or more of the capabilities listed in the following table:

Expand table

Capability	Example prompts
Check if my organizational credentials were compromised in the dark web	- `According to Cybersixgill, tell me if credentials of my organization were compromised in the dark web in the last month` - `According to Cybersixgill, tell me if the following xxx@xxx.com credentials were exposed in the dark web in the last month`
Search about recent breach/ransomware events	- `According to Cybersixgill, Tell me about the last 10 breach events from forum_breach related to the UK. Summarize it in a table and provide victim name with sector` `Tell me about ransomware attacks in the last 90 days targeting the healthcare sector using Cybersixgill database`
Get details about specific CVE	- `According to Cybersixgill, tell me if for CVE-2023-34362 there is POC, metasploit and known to be exploited in the wild with any associated APTs. include references`
Get comprehensive profiles of APTs, Actors, and related Malware and IOCs	- `According to Cybersixgill, tell me about related malware and IOCs for APT39` - `According to Cybersixgill, tell me about associated APTs with Raccoon Stealer?` - `According to Cybersixgill, I'm a CTI analyst at a european MSSP. A customer of mine is worried about the APT MuddyWater. Tell me about it` - `According to Cybersixgill, tell me to which apt cve-2023-36884 is connected to`
Lookup on specific IOCs or IPs	- `According to Cybersixgill, tell me about the following IP [IP]` - `According to Cybersixgill, give me IOCs for racoon stealer`

Troubleshoot the Cybersixgill plugin

Errors occur

If you get errors, such as *Couldn't complete your request*, or *An unknown error occurred*, make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use. For example, if you're using multiple sources that provide information like Cybersixgill, there might be conflicts. Try using the product name `Cybersixgill` in your prompts.

Provide feedback

To provide feedback, contact Cybersixgill.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Cyware Intel Exchange

08/25/2025

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

Cyware Intel is an automated Threat Intelligence Platform (TIP) for ingestion, enrichment, analysis, prioritization, actioning, and bidirectional sharing of threat data.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot requires an API Key. You'll need to take the following steps before using the plugin.

1. Get your Cyware Respond API key. If you don't have one yet, follow these steps:
    a. Go to the Cyware website and create an account.
    b. Go to the Admin panel, and then go to **Open APIs**.
    c. Create a new OpenAPI credential set.
    d. Copy your OpenAPI token generated.
2. Sign in to Microsoft Security Copilot.
3. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.
4. Next to **Cyware Intel Exchange**, select **Set up**.
5. Provide your Cyware Respond Instance URL and API Token.
6. Select **Save and Test**. Then select **Save**.

Sample Cyware Respond prompts

After the Cyware Intel Exchange plugin is configured, you can use the following capabilities with Security Copilot. The following table lists example prompts to try.

Expand table

Capability	Example prompts
Querying Threat Intelligence	- `Search Cyware Intel Exchange for the malware Ryuk.` - `Can you search for the above indicator on Cyware Intel Exchange`.
Convert Natural Language query to Cyware query language query	- `Can you give me a CQL query to find all malwares associated with the threat actor APT2?` - `Generate CQL query for searching an indicator on Cyware Intel ExchangeGenerate CQL query for searching an indicator on Cyware Intel Exchange.`

Troubleshoot the Cyware Intel Exchange plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*, make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts are not invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Provide feedback

To provide feedback, contact Cyware.

See also

Other plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Cyware Respond

08/25/2025

Cyware offers threat intelligence management, security collaboration, and orchestrated response solutions. Cyware Respond is an end-to-end incident management and threat response automation platform. You can use the Cyware Respond plugin with Microsoft Security Copilot to find specific types of incidents, actions, applications, critical software assets, malware, vulnerabilities, and more.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot requires an API Key. You'll need to take the following steps before using the plugin.

1. Get your Cyware Respond API key. If you don't have one yet, follow these steps:
    a. Go to the Cyware website and create an account.
    b. Go to the Admin panel, and then go to **Open APIs**.
    c. Create a new OpenAPI credential set.
    d. Copy your OpenAPI token generated.
2. Sign in to Microsoft Security Copilot.
3. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.
4. Next to **Cyware Respond**, select **Set up**.
5. Provide your Cyware Respond Instance URL and API Token.
6. Select **Save and Test**. Then select **Save**.

Sample Cyware Respond prompts

After the Cyware Respond plugin is configured, you can use it by typing `Cyware Respond` in your Security Copilot prompt bar, followed by an action. The following table provides several examples you can try:

Expand table

Capability	Example prompts
Incident Management	

Querying incidents	`"Search for incidents related to 'ransomware' in Cyware Respond."`
	
	`"Find incidents mentioning 'data breach' in Cyware Respond."`

Retrieving incident details	`"Get details for the above incident from Cyware Respond."`
	
	`"Show more information about the incident in Cyware Respond with UUID <uuid>."`

Action Management	

Querying actions	`"Look up actions tagged with 'isolate' from this week in Cyware Respond."`
	
	`"Find all 'network scan' actions performed yesterday in Cyware Respond."`

Action details	`"Retrieve details of the last action we discussed from Cyware Respond."`
	
	`"Can you provide more info on the previously mentioned action in Cyware Respond?"`

Application Management	

Querying applications	`"Search for applications with 'firewall' in their name in Cyware Respond."`
	
	`"Find security tools updated recently in Cyware Respond."`

Application details	`"Get details for the above application from Cyware Respond."`
	
	`"Show configuration details for the previously discussed security application in Cyware Respond."`

Software asset management	
Querying software assets	`"List critical software assets in Cyware Respond."`

Capability	Example prompts
	`"Search for software needing updates in Cyware Respond."`

Software details	`"Provide details for the last mentioned software asset in Cyware Respond."`
	
	`"Can you fetch the version info of the software we talked about earlier in Cyware Respond?"`

Campaign management	

Querying campaigns	`"Find all active campaigns related to 'phishing' in Cyware Respond."`
	
	`"Search for recent campaigns targeting remote employees in Cyware Respond."`

Campaign details	`"Get campaign details for the previously mentioned campaign in Cyware Respond."`
	
	`"Show more details about that cyber threat campaign from Cyware Respond."`

Threat intelligence management	

Querying threat intelligence	`"Search for threat intel on 'DDoS attacks' in Cyware Respond."`
	
	`"Retrieve updates on 'APT groups' from this week in Cyware Respond."`

Threat intel details	`"Show details of the last threat intel we discussed in Cyware Respond."`
	
	`"Can you provide more info on the threat actor mentioned earlier in Cyware Respond?"`

Malware Management	

Querying malware	`"Find all malware detections from the past week in Cyware Respond."`
	
	`"Search for 'spyware' detections in the marketing department in Cyware Respond."`

Malware details	`"Give me the details of the previously mentioned malware in Cyware Respond."`
	
	`"I need more information on that 'ransomware' we identified in Cyware Respond."`

Vulnerability management	

Capability	Example prompts

Querying vulnerabilities	`"Search for high-severity vulnerabilities in Cyware Respond."`
	
	`"List all vulnerabilities discovered in the network infrastructure in Cyware Respond."`

Vulnerability details	`"Get details of the vulnerability we discussed last time in Cyware Respond."`
	
	`"Show mitigation steps for the above-mentioned vulnerability in Cyware Respond.`

Troubleshoot the Cyware plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*, make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts are not invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Provide feedback

To provide feedback, contact Cyware.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Darktrace

Darktrace offers cybersecurity AI services to provide preemptive visibility into security posture, real-time detection, and autonomous response to known and unknown threats. You can use the Darktrace plugin with Microsoft Security Copilot to proactively detect and investigate threats across your digital ecosystem.

- Query Darktrace's rich and high fidelity security data using natural language, streamlining the investigation process and enabling faster threat response.
- Leverage Copilot's generative AI capabilities to interact with Darktrace's security alerts, device information, and email security insights, enhancing the overall understanding of your security landscape.
- Gain comprehensive visibility into your organization's security posture, including device connections, external endpoints, Cyber AI Analyst incidents, and email-related threats.

By integrating Darktrace's unparalleled AI-driven threat detection and mitigation with Microsoft Security Copilot, this plugin improves the capabilities of security teams across the board. Obtain real-time insights into possible security incidents, privileged accounts, and unusual user activities through an easy-to-use, natural language interface.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

The Darktrace integration with Security Copilot uses OAuth Client Credentials flow to authenticate and requires a Client Application to be configured in the Darktrace Active AI Security Portal. Please note that Security Copilot integration is currently only available to Early Adopters of the ActiveAI Security Portal. If you would like to trial Security Copilot, please contact your Darktrace account team. You'll need to take the following steps before using the plugin.

1. Sign into your Darktrace ActiveAI Security Portal, go to the Client Applications page, and create your CCA.

    - Your Darktrace API URL
    - Client ID
    - Client Secret

- Your Scopes
- Token Endpoint

2. Sign in to Microsoft Security Copilot.
3. Access **Manage plugins** by selecting the **Sources** button from the prompt bar.
4. Next to **Darktrace**, select the toggle to enable it.
5. In the settings pane for the plugin, provide the information you saved from step 1. Leave AuthorizationContentType field as default, and in the TokenEndpoint field enter `https://auth.login.darktrace.com/oauth/token`.
6. Save your changes.

Sample Darktrace prompts

After the Darktrace plugin is configured, you can use it by including `Darktrace` in your Security Copilot prompt along with an action. The following table lists example prompts to try.

Expand table

Capability	Example prompts

Get Model Breaches	`What were the top 5 high-scoring Darktrace alerts in the last week?`
	
	`Show me all the model breaches involving device ID 500 in the past month from Darktrace.`
	
	`Retrieve the model breach information for pbid 1234 and explain what occurred.`
	
	`Display Darktrace model breaches for subnet ID 250 with a minimum score of 0.7 in the last 72 hours.`

Get Model Breach Comments	`What are the latest comments on the Darktrace model breach with pbid 4567?`
	
	`Has anyone added any valuable information to the model breach with pbid 1234 in Darktrace?`

Get AI Analyst Incident Groups	`Show me the Darktrace AI Analyst incidents with a score higher than 90 from the past week.`
	
	`List all critical incidents involving device ID 1500 in the last month from Darktrace.`
	
	`Retrieve incidents associated with subnet ID 300 in the past 7 days from`

Capability	Example prompts

	`Darktrace and display them in the German language.`
	
	`Are there any Darktrace AI Analyst incidents with the unique identifier "abcd-1234-efgh-5678" in the system?`

Get AI Analyst Incident Events	`Give me a detailed breakdown of the Darktrace incident with groupid "g04a3f36e-4u8w-v9dh-x6lb-894778cf9633".`
	
	`Show me more information about this Darktrace incident in Spanish.`
	
	`What are the Darktrace AI Analyst events for device ID 1000 that are part of critical incidents?`

Get AI Analyst Incident Comments	`What are the recent comments on the Darktrace AI Analyst event with incident_id "04a3f36e-4u8w-v9dh-x6lb-894778cf9633"?`
	
	`Has anyone from the security team added any context to this Darktrace incident?`

Get Device Information	`Can you provide detailed information about device ID 1234 from Darktrace, including any tags associated with it?`
	
	`What is the current IP address of device ID 9 in the Darktrace system?`
	
	`Retrieve the device information for the entity with MAC address "AA:BB:CC:DD:EE:FF" from Darktrace.`

Search Devices	`Find all devices in subnet 10.0.1.0/24 with a "Respond" tag and sort them by last seen using Darktrace.`
	
	`Has Darktrace seen a laptop with the IP address 8.8.8.8?`
	
	`Show me a list of devices that could be owned by "sarah" in Darktrace, sorted by last seen activity.`

Format Model Breaches as Table	`Get me high scoring model alerts from Darktrace in the past week, format the results as a table, and give me definitions of any complex terminology.`
	
	`Display all Darktrace model breaches involving device ID 250 in the past month in a table format, and include a column with descriptions of the breach categories.`
	
	`Present Darktrace model breaches for subnet ID 100 with a minimum score of 0.8 in the last 72 hours as a table, and include a column with the involved devices' hostnames.`

Capability	Example prompts

Analyze AI Analyst Incident Trends	`What are the common themes among the Darktrace AI Analyst incidents with a score higher than 90 from the past month?`
	
	`Analyze the Darktrace AI Analyst incidents involving device ID 1500 in the last month and identify any patterns or trends in the security events.`
	
	`Are there any recurring issues or attack vectors in the Darktrace AI Analyst incidents associated with subnet ID 200 in the past 7 days?`

Investigate AI Analyst Incident Event	`Can you provide a detailed analysis of the Darktrace AI Analyst event with incident_id "04a3f36e-4u8w-v9dh-x6lb-894778cf9633" and suggest possible mitigations?`
	
	`Investigate the Darktrace AI Analyst event with incident_id "04a3f36e-4u8w-v9dh-x6lb-894778cf9633" in depth and provide a report on its potential impact on our network.`

Device Owner Identification	`Can you identify the owner of the device with IP address 8.8.8.8 in Darktrace based on its activity patterns and associated user accounts?`
	
	`Determine the likely owner of the device with MAC address "AA:BB:CC:DD:EE:FF" in Darktrace by analyzing its usage patterns and associated services.`

Troubleshooting the Darktrace plugin

Timestamp Support

To filter data to a specific time frame, this plugin may require timestamps in Epoch/UNIX format. To retrieve the relevant time frame in the correct format, use a service such as https://epochconverter.com or https://unixtime.org.

Provide feedback

To provide feedback, contact Darktrace.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Last updated on 02/19/2026

Forescout Risk and Exposure Management

08/25/2025

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

Forescout Risk and Exposure Management (REM) platform provides data across all devices in a customer environment as discovered by the Forescout platform. REM gives a single view of device risk and vulnerabilities, including a timeline to track events and changes to a device risk value. Combining the contextual data from the Forescout REM platform with Microsoft Security Copilot allows security teams to quickly assess the risk posed across their environments and prioritise work by getting clear information on their exposure to different vulnerabilities or configurations which exist.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot works with an API key. You need to take the following steps before using the plugin.

1. Get your Forescout Risk and Exposure Management API key. If you don't have one yet, follow these steps:
   a. Login to Forescout Cloud.
   b. In the Forescout Cloud Console, select **Integrations** under the **Administration** menu.
   c. Click the **Generate API Key** button next to the category of your application - IoT/OT. The Generate API Key configuration screen appears.
   d. Provide a meaningful name for the API key such as **Microsoft Security Copilot**.
   e. Select a time for the API key to expire or select **Never Expires**. If the key is set to expire, select users to receive Email notifications about the API key generation and expiry date.

f. Click the **Generate** button and copy the API key that appears on screen. This API key is unique and non-retrievable once the window is closed. Save this key in a secure location for use in Microsoft Security Copilot.
g. Copy the API Endpoint URL that appears below the key.

2. Sign in to Microsoft Security Copilot.
3. Access **Manage Plugins** by selecting the **Sources** button from the prompt bar.
4. Next to **Forescout Risk and Exposure Management**, select **Set up**.
5. In the Forescout Risk and Exposure Management plugin settings pane, under **Forescout Cloud Data Exchange API Endpoint URL** enter the API endpoint URL.
6. In the **Value** field, paste your API Key, and then select **Save**.

Sample Forescout Risk and Exposure Management prompts

After the Forescout Risk and Exposure Management plugin is configured, you can use the following capabilities with Security Copilot.

The following table lists example prompts to try.

Expand table

Capability	Example prompts
Get REM assets	- `Using Forescout, tell me about my top 5 riskiest devices.` - `Tell me about the risks for X.` - `Show me all devices with CVE-1234-5678.`

Troubleshoot the Forescout Risk and Exposure Management plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*. Make sure the plugin is turned on. This error may occur if the lookback period is too long,

causing the query to attempt to retrieve an excessive amount of data. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Provide feedback

To provide feedback, contact Forescout Risk and Exposure Management.

See also

Other plugins for Microsoft Security Copilot Manage plugins in Microsoft Security Copilot

Forescout Vedere Labs

08/25/2025

Forescout Vedere Labs research team provides a threat intelligence feed containing IP, URL, and File hash indicators for all activity seen and monitored by Forescout, including information on Known Exploited Vulnerabilities and Vedere Labs own reported CVEs. With the extensive research conducted, this provides indicators and CVE details across IT, OT, IoT and IoMT, allowing anyone to benefit from this research. The research helps security teams speed up threat hunting efforts in combination with Microsoft Security Copilot. Additionally, this service also allows for lookups against domain names to check for use of Domain Generation Algorithms (DGA) or data exfiltration techniques.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Prerequisites

Forescout Vedere Labs Threat Feed API Key.

Know before you begin

Integration with Security Copilot requires an API key for authorization. Navigate to Forescout Vedere Labs and register for a free API Key to start taking advantage of the vulnerabilities and indicators provided by this feed. You'll need to take the following steps before using the plugin.

1. Sign in to Microsoft Security Copilot.
2. Access **Manage Plugins** by selecting the **Sources** button from the prompt bar.
3. Next to **Forescout Vedere Labs**, select **Set up**.
4. In the Forescout Vedere Labs settings pane, input your API key in the **Value** field, and then select **Save**.

Sample Forescout Vedere Labs prompts

After the Forescout Vedere Labs plugin is configured, you can use the following capabilities with Security Copilot.

The following table provides examples you can try:

Expand table

Capability	Example prompts

Get Indicators	`Show me vedere labs file indicators for the past 8 hours`
	
	`Tell me about any indicators for [IP] or [File Hash]`

Get Exploited CVEs	`Show me all known exploited vulnerabilities in the last 7 days`
	
	`give me the latest 5 KEVs according to Vedere Labs from the past 7 days with a confidence score of at least 6`

Get Vedere Labs CVEs	`Show me the most recent Vedere Labs CVE`

Lookup Domain	`Has "example.com" got any indicators of malicious use`
	
	`Check if "dga.com" is using any techniques related to malware`

Provide feedback

To provide feedback, contact Forescout Vedere Labs.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

GreyNoise Enterprise and GreyNoise Community

08/25/2025

GreyNoise Enterprise is a cybersecurity platform that collects and analyzes Internet-wide scan and attack data to provide valuable insights into potential threats. GreyNoise integration enables you to use the GreyNoise database to enhance your organization's security posture, identify emerging threats, and prioritize response efforts. You can use either the GreyNoise Enterprise or GreyNoise Community plugin with Security Copilot to get information about IP addresses, scanning activity, and attacker behaviors.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot works with a GreyNoise account and an API key. Depending on the plan you choose, you might have a limit on how many queries you can run using your API Key. You'll need to take the following steps before using the plugin.

1. Get your GreyNoise credentials and API key. If you don't have them yet, follow these steps:

   a. Go to the GreyNoise website and create your account.

      You can start with a free account to obtain a community API Key, or purchase a subscription to obtain an enterprise API Key.

   b. In GreyNoise Visualizer, in the upper right corner, select your name, and then select **My API Key**.

   c. Copy your API Key.

2. Sign in to Microsoft Security Copilot.

3. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.

4. Take one of the following steps:

- If you purchased a GreyNoise enterprise subscription, turn **GreyNoise Enterprise Plugin** on.
- If you're using a free GreyNoise account, turn **GreyNoise Community Plugin** on.

5. Select the Settings icon, and in the **Value** field, paste your API Key, and then select **Save**.

Sample GreyNoise prompts

After you set up your GreyNoise plugin, you can use it by typing a capability name in your prompt bar in Security Copilot. For example, you can type `LookupIpAddressNoise`.

The following table summarizes capabilities and what they do.

Expand table

Capability	Example prompts	What it does
`Lookup IP Address Noise` (works with either GreyNoise Enterprise or GreyNoise Community) Required Input: IP Address (v4 or v6)	- `Tell me about Ip address "118.25.6.39" using the GreyNoise database` - `Use the GreyNoise database to provide info on "118.25.6.39"` - `What does the GreyNoise database say about the IP address 180.126.219.127?` - `I'm curious about any GreyNoise records for the IP address 180.126.219.127. Can you look that up for me?` - `Can you provide me with information on any GreyNoise reports for the IP address 180.126.219.127?` - `I'd like to know if there are any GreyNoise entries for the IP address 180.126.219.127. Can you check that for me?` - `Could you give me an overview of the GreyNoise record for the IP address 180.126.219.127?`	Retrieves noise information about the provided IP address. Returns the following kinds of information: - IP address classification, such as malicious - Noise, such as whether the IP address is likely involved in some form of malicious activity - Riot, such as whether the IP is part of a known benign service or infrastructure - Name associated with the IP - Last seen (when the IP was last active) - Link: A link to visualize the IP's activity on GreyNoise - Success or error message depending, on whether the lookup was successful

Capability	Example prompts	What it does
`Lookup IP Context` (requires GreyNoise Enterprise) Required Input: IP Address (v4 or v6)	`Find the GreyNoise IP Context for IP 183.221.243.13`	Provides context about IPs that GreyNoise observed scanning the internet. Returns a comprehensive set of information including classification (malicious, benign, etc.), last seen timestamp, associated actors, tags, and metadata.
`Lookup IP Quick` (requires GreyNoise Enterprise) Required Input: IP Address (v4 or v6)	`Use GreyNoise to do a quick check of IP 183.221.243.13`	Provides a quick way to check if an IP is "noise" or not. Returns a boolean indicating whether the IP is present in the dataset or not.
`Lookup Multiple IPs` (requires GreyNoise Enterprise) Required Input: IP Address (v4 or v6)	`Lookup Multiple IPs using GreyNoise 183.221.243.13 and 8.8.8.8`	Provides a quick way to check information on multiple IPs. Returns an array of context information for each IP address, similar to the LookupIpContext endpoint.
`Lookup IP Riot` (requires GreyNoise Enterprise) Required Input: IP Address (v4 or v6)	`Use GreyNoise to check the Riot information on IP 183.221.243.13`	Provides information about IPs commonly added to allowlists. Returns a boolean indicating whether the IP is part of the RIOT dataset or not, along with some basic context information if it is.
`Lookup GNQL` (requires GreyNoise Enterprise) Required Input: GNQL Query	`Use GreyNoise to check the GNQL information on tags:"RDP Scanner"`	Allows you to use the GreyNoise Query Language (GNQL) to make complex queries against the GreyNoise dataset. Returns an array of results that match the GNQL query.
`Look up CVE` (requires GreyNoise Enterprise) Required Input: CVE	`Use the GreyNoise Query Language (GNQL) stats to query against the GreyNoise dataset for CVE information`	Allows you to use the GreyNoise Query Language (GNQL) stats to query against the GreyNoise dataset for CVE information Returns an array of results that match the GNQL CVE query.

Troubleshoot the GreyNoise plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*, make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use. To prioritize and target GreyNoise, try disabling other custom plugins.

Provide feedback

To provide feedback, contact GreyNoise.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

HP Workforce Experience Platform

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

The HP Workforce Experience Platform (WXP) plugin for Security Copilot allows users to query data about their fleet and get responses back quickly and efficiently. This plugin can answer questions about PC warranty, Windows 11 readiness, resource consumption. It can also answer device age/stability/health, security, blue screens, and app performance including hangs and crashes.

> **Note**
>
> This article contains information about non-Microsoft plugins. This is provided to help complete integration scenarios. However, Microsoft doesn't provide troubleshooting support for non-Microsoft plugins. Contact the vendor for support.

Prerequisites

This section lists the requirements for the integration

- HP Workforce Experience
- Security Copilot

Know before you begin

The HP Workforce Experience (WXP) plugin for Security Copilot uses OAuth 2.0 for authentication with the WXP service. This requires user credentials for authentication.

User credentials

1. Sign in into HP Workforce Experience Platform (https://workforceexperience.hp.com/) as an IT Admin.
2. Go to Account Management and select the **Users** tab.
3. Add a user (whoever needs access to the plugin in Security Copilot).

4. When adding a user, an IDP of either HP ID or AAD can be selected. This user can have a role of either Report Admin or IT Admin.

> ⓘ **Note**
>
> 1. This user should be a part of only one customer account in HP Workforce Experience Platform.
> 2. Irrespective of the IDP selected in the user creation flow, when the OAuth flow is triggered from Security Copilot, an HP ID is required.

Client Credentials

A key part of the authentication process is client credentials. Client credentials for HP WXP can be obtained from https://developers.hp.com/.

1. Sign up on https://developers.hp.com/ if you don't have an account already present.
2. Send an email to workforceexperience@hp.com requesting to enable API credentials request access in the PI APIs
3. Alternatively, create a case from https://workforceexperience.hp.com/ . Include the email address used for signup in the developer portal. Additionally request to enable them in the dev portal.
4. Once the case is generated, it takes upto 24 hours for enablement.
5. Go to https://developers.hp.com/hp-proactive-management.
6. Select "Get Credentials" against HP Proactive Insights Analytics.
7. In the redirect URL, enter https://securitycopilot.microsoft.com/auth/v1/callback.
8. Leave "Client ID" empty.
9. Under Attributes -> Scopes, select the checkbox of "Read".
10. Select "Create".
11. The API credentials are generated.
12. To access the credentials, go to your profile by clicking on your email on the top right corner.
13. Select My API Credentials.

Microsoft Security Copilot Integration Configuration

1. Go to https://securitycopilot.microsoft.com/.
2. Select the Sources icon.
3. Under the "Non-Microsoft" section, find the plugin "HP Workforce Experience".

> **Note**
>
> You might see two plugins – one for US and one for EU. Use the plugin based on the region your customer company is created in HP Workforce Experience Platform.

4. Select Setup. A popup window appears.
5. Enter the ClientId and ClientSecret.
6. Enter your user credentials when prompted.

Sample HP Workforce Experience Platform prompts

1. For some prompts, data would be picked up from Microsoft Intune plugins. To ensure the data is picked up through the HP WXP plugin, in the prompt explicitly request to get the data from HP Workforce Experience Platform.

Expand table

Prompt example	Expected output
What device warranties expire soon?	A list of devices with warranties expiring soon, if applicable. *Example:* `"The following devices have warranties that have already expired: [table with serial numbers, warranty status, and warranty end dates]"`
Which devices aren't compatible with Windows 11?	A list of devices not compatible with Windows 11, if applicable. *Example:* `"There are currently no devices that aren't compatible with Windows 11 based on the available data."`
List the top 10 devices that have the least disk space available	A list of devices with the least amount of disk space in your fleet. *Example:* `"Here are the top 10 devices with the least disk space available: [table containing device names, device models, and disk space used]"`
What is the DEX score for device [serial number]?	The DEX score of the mentioned device. *Example:*

Prompt example	Expected output
	`"Your DEX score is [XX.X]. Here are the details of the device contributing to this score: [table containing device name, device model, device manufacturer]"`
When was the last application crash on one of my devices?	Information on the last application that crashed on a device in the fleet. *Example:* `"The last application crash on one of your devices occurred on [Month Date Year] at [time]. The application that crashed was [application name] on the device named [device name]. Here are the details of the last crash: [table including device name, application name, and crash date and time]"`

Troubleshoot the HP Workforce Experience Platform plugin

Errors occur

If you encounter errors occur, such as: *Couldn't complete your request*, or *An unknown error occurred*, make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Last updated on 02/19/2026

Intel 471 Threat Intelligence

08/25/2025

Intel 471 delivers relevant and timely insights into the cyber underground.

You can use the Intel 471 Threat Intelligence plugin with Microsoft Security Copilot to provide:

- Rich context about threats through indicators of compromise to facilitate responding to active threats found in the networks and CTI analysis.
- Relevant and timely intelligence information about the activity groups and adversary scenarios.
- Visibility of potential third-party breaches to improve the risk posture and change the odds of a successful attack.
- Vulnerability intelligence, including the precursors to such activity such as an increase in interest levels among activity groups, proof-of-concept (POC) code being developed, traded, or sold, and more.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot requires Intel 471 API credentials, which are available as part of Intel 471's paid subscriptions. For more information, please contact sales@intel471.com. You'll need to take the following steps before using the plugin.

1. Get your Intel 471 user account and password.
2. Sign in to Microsoft Security Copilot.
3. Access **Manage Plugins** by selecting the **Sources** button from the prompt bar.
4. Next to **Intel 471 Threat Intelligence**, select **Set up**.

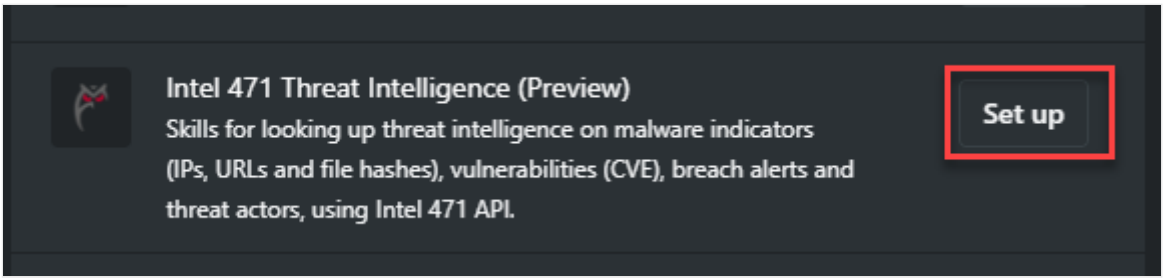

5. In the Intel 471 Threat Intelligence settings pane, provide your Intel 471 user account and password.

6. Save your changes.

Sample Intel 471 Threat Intelligence prompts

After the Intel 471 Threat Intelligence plugin is configured, you can use natural language prompts such as the examples provided in the following image. You can also choose them from the prompts menu located in the prompt input.

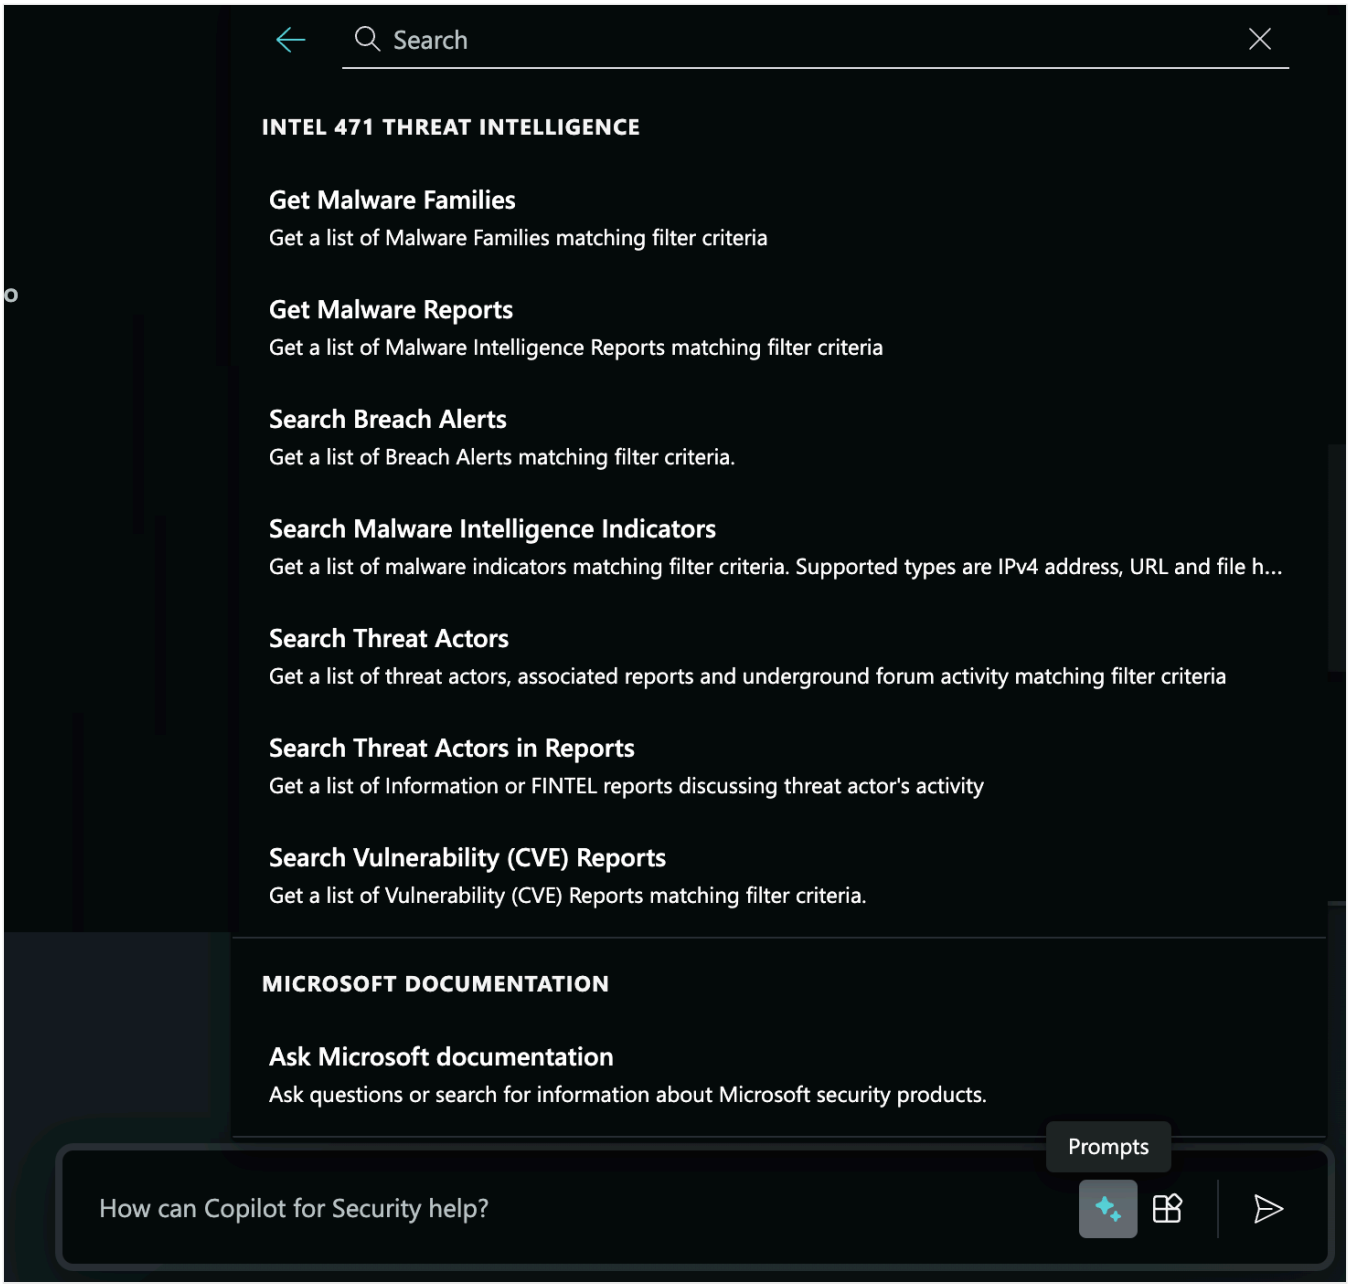

The following table provides examples you can try:

Expand table

Capability Prompt Suggestion	Example natural language prompt
Malware Intelligence	- Search for Malware Intelligence Indicator with sha256 hash cdf2489960aa1acefedf4f1ab4165365db481591a17e31e7b5fca03cc0ccff0f - Search for reports on operaloader malware family
Vulnerability Intelligence	- Search for vulnerability report on CVE-2023-6553

Capability Prompt Suggestion	Example natural language prompt
Breach Intelligence	- Find a breach alert for victim TTEC - What is the source of the information for breach of TTEC by actor netsher?
Adversary Intelligence	- Find information regarding activity group with handle netnsher. Use Intel 471 plugin

Provide feedback

To provide feedback, contact Intel 471 Threat Intelligence.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

IPGeolocation

08/25/2025

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

IPGeolocation is a powerful API service providing comprehensive geolocation data, time zone information, currency details, security insights (such as VPN, proxy, and bot detection), and hostname resolution based on IP addresses. Designed for developers, security analysts, and digital service providers, our platform helps enhance security, optimize user experiences, and improve decision-making processes with real-time geolocation intelligence. IPGeolocation seamlessly integrates with your existing systems to improve workflow efficiency.

By integrating the IPGeolocation API into Security Copilot, users can enhance their system's ability to detect threats, tailor content delivery, and make informed decisions with accurate location and security data based on an IP address.

> **Note**
>
> This article contains information about non-Microsoft plugins. This article is provided to help complete integration scenarios. However, Microsoft doesn't provide troubleshooting support for non-Microsoft plugins. Contact the vendor for support.

Know before you begin

Integration with Security Copilot works with an API key. You need to take the following steps before using the plugin.

1. Get your IPGeolocation API key. If you don't have one yet, follow these steps:
   a. First, Sign Up on IPGeolocation using any preferred method (GitHub, Google, or Email). If you signed up with email, follow the instructions to verify your account and sign in.
   b. From your dashboard, copy your API key.

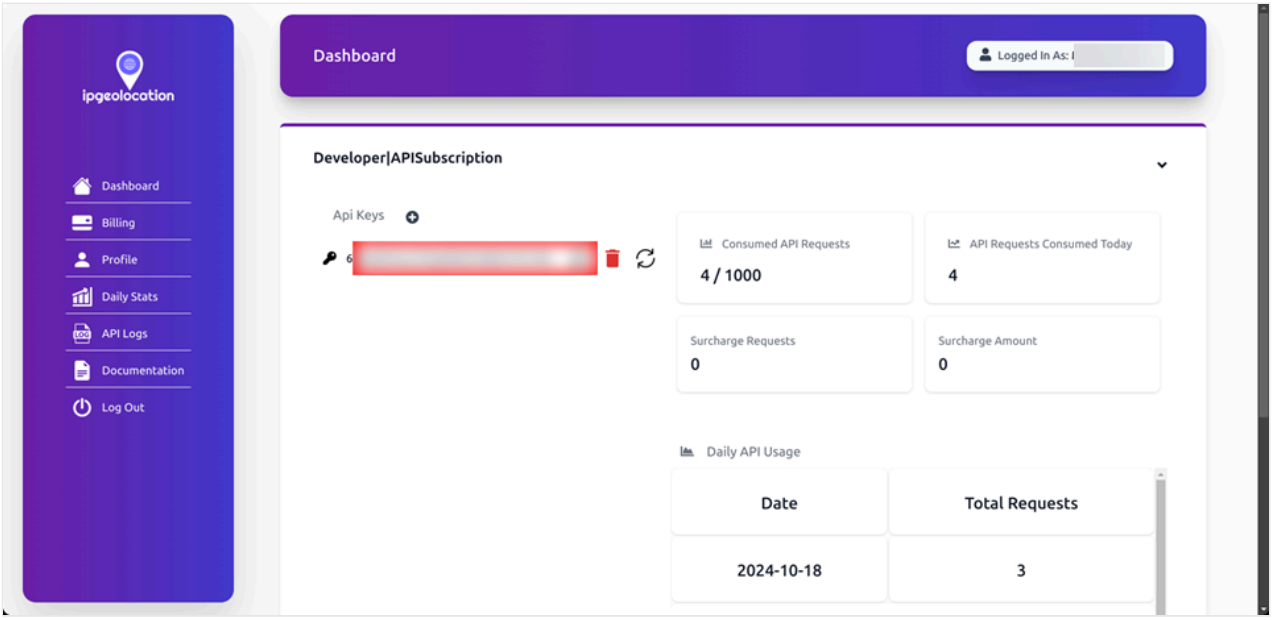

2. Sign in to Microsoft Security Copilot.

3. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.

4. Next to **IPGeolocation**, select **Set up**.

5. In the **Value** field, paste your API Key, and then select **Save**.

Sample IPGeolocation prompts

> **Note**
>
> Security, domain lookup (in place of IP Address), and hostname lookup features are available only with paid subscriptions. All other information can be accessed with a free or developer plan, which allows up to 1,000 API requests per day. Paid plans don't have any request limits. For more details, visit IPGeolocation Pricing.

After the IPGeolocation plugin is configured, you can use the following capabilities with Security Copilot. The following table lists example prompts to try.

Expand table

Capability	Example prompts
IP Address Details (Geolocation, Security, Time Zone, and Hostname Info)	`Give me complete details including threat score and security data, geolocation, time zone info, and hostname for IP address 49.12.212.42.`

Capability	Example prompts
Time Zone and Currency Information	`Provide the time zone and currency associated with the IP address 49.12.212.42.`
ASN Lookup Information	`What ASN is associated with the IP address 49.12.212.42?`
Location Insights for a Specific IP Address	`Give me all available location insights for the IP address '8.8.8.8' utilizing IPGeolocation.`
Security Insights for an IP Address	`Is the IP address 146.70.250.27 a VPN or an attacker IP address according to IPGeolocation?`
Lookup Domain	`Provide all available insights, including location, security, time zone, and other relevant data, for the domain example.com using ipgeolocation.`

Troubleshoot the IPGeolocation plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred* | Make sure the plugin is turned on. This error may occur if the lookback period is too long, causing the query to attempt to retrieve an excessive amount of data. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Provide feedback

To provide feedback, contact IPGeolocation.

See also

Other plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

IPinfo

08/25/2025

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

IPinfo is an internet data company specializing in IP geolocation, IP to Privacy Detection (VPN, Tor, Proxy, etc.), ASN data, company data, carrier metadatadata, WHOIS data, and other internet data services. IPinfo's internet data can help organizations conduct cybersecurity research, threat intelligence and network administration, among many other things.

With IPinfo's broad internet data resources, the IPinfo plugin for Microsoft Security Copilot allows users to get diverse metadata information on internet-connected devices, create robust threat intel models, and run effective security control measurements.

IPinfo offers access to various internet data based on different service tiers. The IPinfo plugin will provide internet metadata based on the service tier of your IPinfo access token. However, our free IP geolocation data is accessible to users through a free account that you can start today. The free tier lets users have 50,000 IP to Geolocation/ASN lookups monthly. Through Copilot, you can use the "IPinfo" plugin to get more advanced data services.

Explore the plugin features from this video, Microsoft Security Copilot Demo.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot works with an API key. You need to take the following steps before using the plugin.

1. Sign up for IPinfo and get your access token. If you do not have an IPinfo account yet, follow these steps:

a. Go to the IPinfo website and create your free account. To access IPinfo's premium data, see the API offerings on the pricing page.
b. After signing up, go to the account dashboard and scroll down to the token section.
c. From there, copy your IPinfo access token.

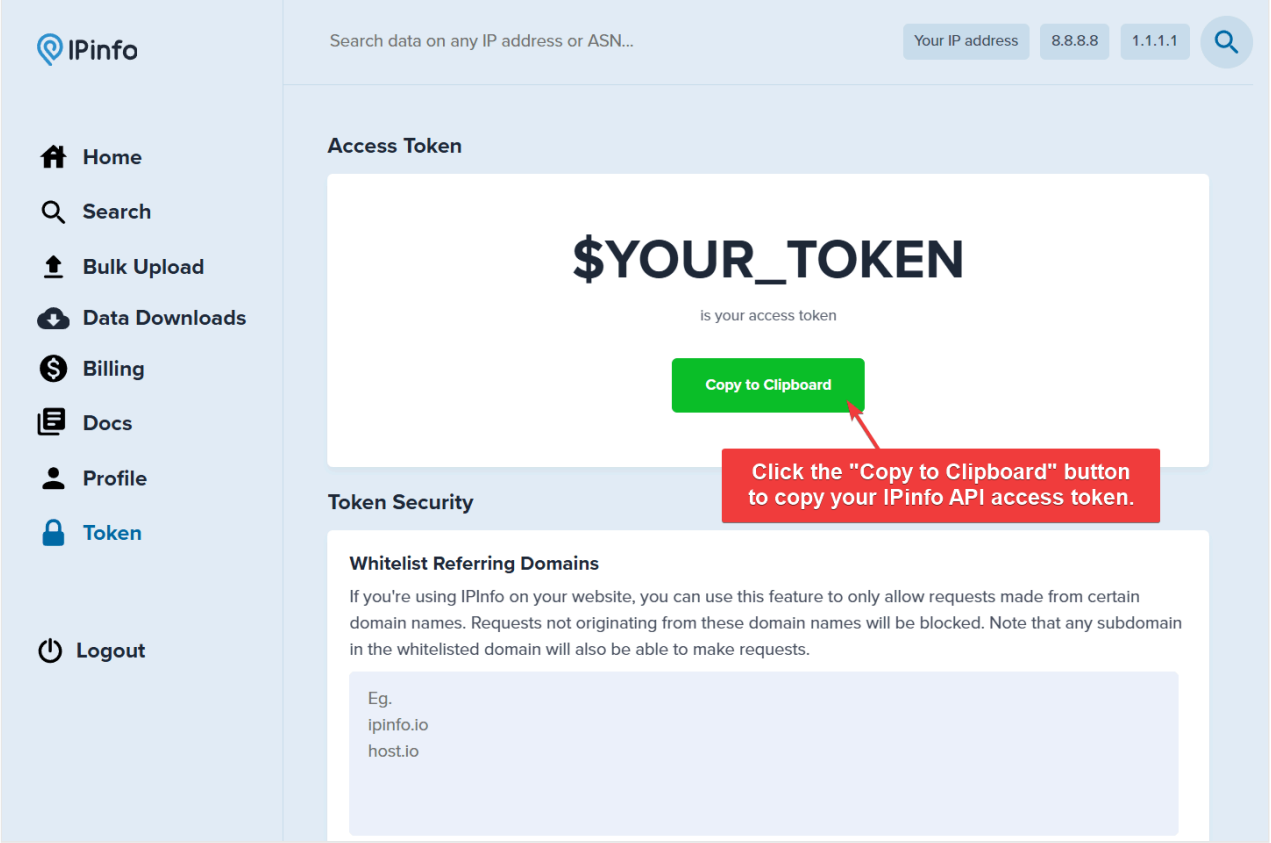

2. Sign in to Microsoft Security Copilot.
3. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.
4. Next to **IPinfo**, select **Set up**.
5. In the **Value** field, paste your API Key, and then select **Save**.

Sample IPinfo prompts

After the IPinfo plugin is configured, you can type "IPinfo" in your Security Copilot prompt bar, followed by an action. The following table lists example prompts to try.

Expand table

Capability	Description	Example prompts
getInformationByIp	Returns information about the provided IP. Supports both IPv4 and IPv6 addresses. Returns all the IP metadata available on the IP address from IPinfo.	`Look up the IP address` 8.8.8.8 `using IPinfo. Give me the IP insights on the IP address` 8.8.8.8 `using IPinfo.`

Capability	Description	Example prompts
	The free tier provides IP location information. For ASN, company, privacy detection, and other IP metadata, you have to subscribe to a paid tier.	`Give me information on the IP` 8.8.8.8 `using IPinfo.`
getLocationByIp	Returns a location for the selected IP.	`Give me all available location insights for the IP address '8.8.8.8' utilizing IPinfo.`
getPrivacyInformationByIp	Returns privacy information about the selected IP, such as VPN, proxy, Tor, relay and hosting service/data center/bot.	`Is the IP address` 171.239.211.111 `a VPN or an anonymous IP address according to IPinfo?`
getCompany	Returns information of the organization that is operating an IP address such as name, organization type and their official domain.	`Using IPinfo, tell me about the company that operates the IP address 60.29.162.215.`
getAsn	Returns information about an ASN, such as the assigned prefixes, related domain, and more.	`Using IPinfo, provide me with insights into ASN 52210. List out the ranges they operate.`
getRanges	Our IP address ranges API returns a list of all of the IP addresses ranges assigned to or operated by a company, identified by the provided domain name.	`What are the IP ranges operated by vercel.com according to IPinfo?`
getDomains	Our Hosted Domains, or Reverse IP API returns a list of all of the domains hosted on the provided IP address.	`What domains are hosted on the IP address` 208.80.154.224 `according to IPinfo?`
getAbuse	IP to abuse contact information such as email address, postal/ZIP code, city, state, country/region, name, network, and phone number.	`Please provide me with the abuse contact for the IP address` 2a15:1e89:8000:: `using IPinfo.`

Troubleshoot the IPinfo plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*. Make sure the plugin is turned on. This error may occur if the lookback period is too long, causing the query to attempt to retrieve an excessive amount of data. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Provide feedback

To provide feedback, contact IPinfo.

See also

Other plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Jamf

08/25/2025

Jamf Pro provides comprehensive MDM data. You can use the Jamf plugin with Microsoft Security Copilot to streamline how your security analysts access data, simplify the process of gathering crucial information, and facilitate seamless collaboration between IT and Security teams.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot requires an API Key. You must have the *Analyst Viewer* or *Admin* role assigned in Jamf to get your API key. You'll need to take the following steps before using the plugin.

1. Sign into your Jamf Pro instance.
2. Create your API client. See Jamf Pro Documentation 11.4.0: API Roles and Clients.
3. Assign an API role to the client that has the *Read Computers* privilege.

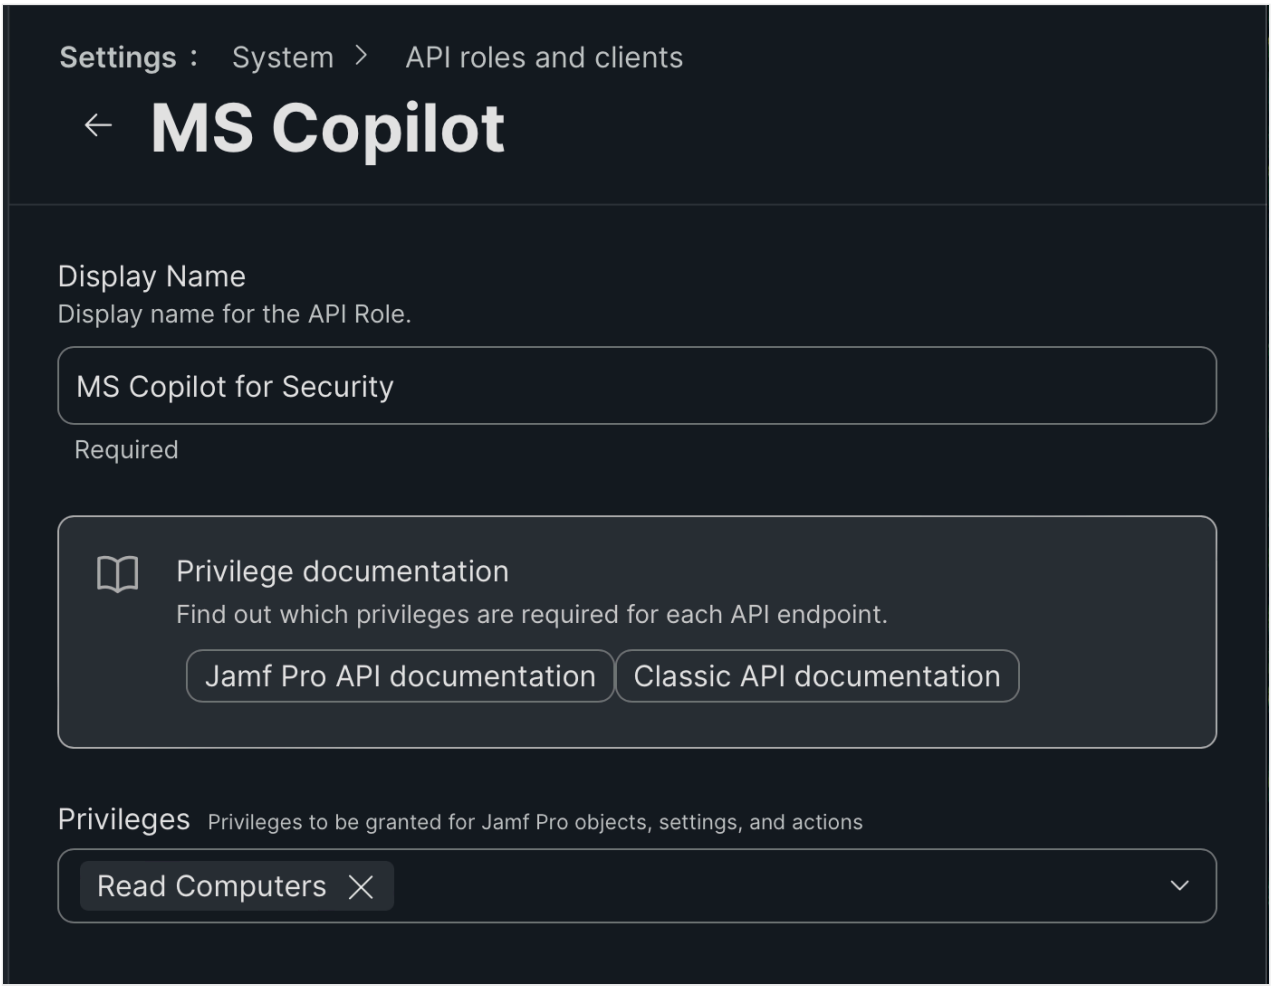

4. Securely copy the client ID, client secret, and required scopes (scope number).

5. Sign in to Microsoft Security Copilot.

6. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.

7. Next to "**Inventory Insights by Jamf**, select the toggle to enable it.

8. Provide the following information:

    - Jamf Pro Instance URL: `https://your-instance.jamfcloud.com`
    - ClientId: (Paste the copied client ID)
    - ClientSecret: (Paste the copied client secret)
    - TokenEndpoint: `https://your-instance.jamfcloud.com/api/oauth/token`
    - Scopes: `api-role:2` (whatever API scope number is associated with the role you created)
    - AuthorizationContentType: `application/x-www-form-urlencoded`

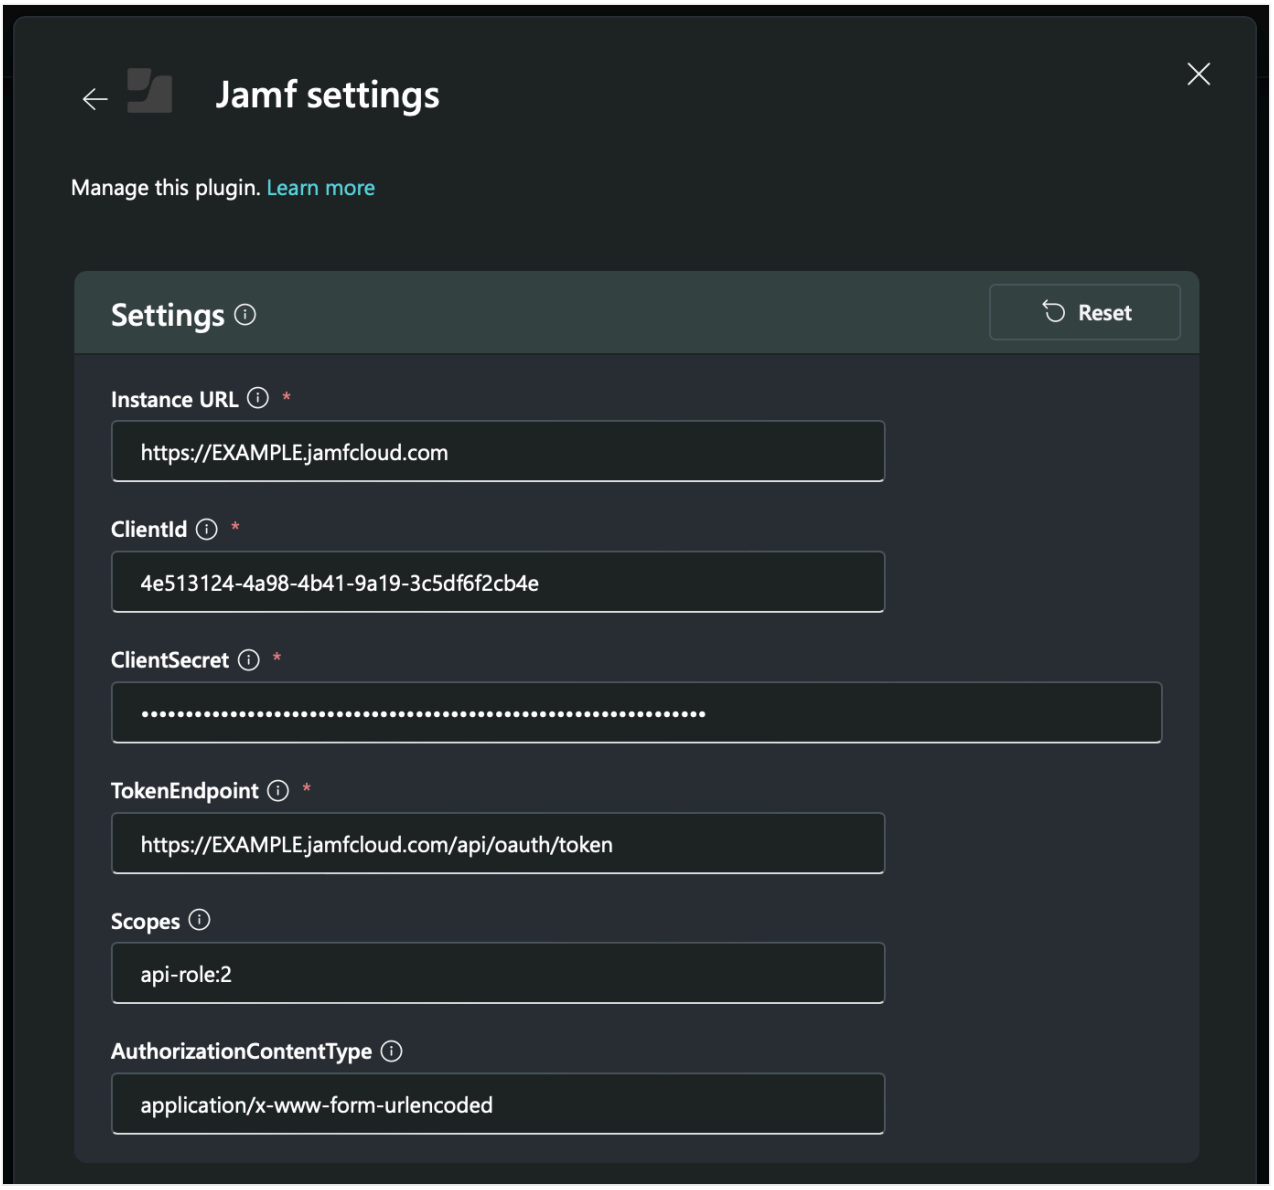

9. Save your changes.

Sample Jamf prompts

After the Jamf plugin is configured, you can use it by typing `Jamf` in your Security Copilot prompt bar, followed by an action.

Here are some examples to try:

`Does the mac H2WFP7BPF6NV have any suspicious applications installed?`

`What are all the user accounts on mac H2WFP7BPF6NV and what are their privilege levels?`

`Get complete inventory details for mac serial number H2WFP7BPF6NV.`

`Show hardware and security configuration for the mac with serial number H2WFP7BPF6NV.`

`List all installed applications on mac H2WFP7BPR6NV.`

`Check Filevault status for mac H2WFP7BPF6NV.`

`What are the firewall settings on mac H2WFP7BPQ6FV?`

Provide feedback

To provide feedback, contact Jamf.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Netskope

08/25/2025

Netskope One is a cloud-native platform that offers converged security and networking services to enable your Secure Access Services Edge (SASE) and Zero Trust transformation. In addition to using the built-in Netskope plugin with Microsoft Security Copilot, you can incorporate other Netskope custom plugins. This article describes how to set up and use the built-in plugin for Security Copilot.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot requires an API token. You'll need to take the following steps before using the plugin.

1. Get your Netskope API token. If you don't have one, follow the steps in Netskope's REST API v2 Overview.
2. Sign in to Microsoft Security Copilot.
3. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.
4. Next to **Netskope Reporting for Threat and Data Protection**, select **Set up**.
5. In the **Value** field, paste Netskope API token.
6. Select **Save**.

Sample Netskope Reporting prompts

After the Netskope Reporting plugin is set up, you can use the following six capabilities with Security Copilot:

- `Get_audit_events`
- `Get_data_alert`
- `Get_data_application`

- `Get_data_infrastructure`
- `Get_data_network`
- `Get_data_page`

The following table provides examples of prompts to try:

Expand table

Capability	Example prompt
Retrieve a list of application events for a specific time frame and domain	`show me Netskope DLP Alerts from 1724996122 and 1724997122`
Retrieve a list of security assessment alerts for a specific app	`show me Netskope application events from 1724996122 and 1724997122 for user example@hotmail.com`
Retrieve a list of page events for a specific user	`show me Netskope page events for the last 15 minutes from user <IP>`
Retrieve a list of page events for a specific user, domain, and traffic type	`show me Netskope page events that occurred in the last 15 minutes from user <IP address> to domain <domain> that is user generated with the traffic type being web`
Retrieve a list of alerts for the last 90 minutes that aren't yet acknowledged	`show me Netskope alerts for the last 90 minutes and only show alerts that are not hacked`

Troubleshoot the Netskope plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred* | Make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Timestamp Support

To filter data to a specific time frame, this plugin requires timestamps in Epoch/UNIX format. To retrieve the relevant time frame in the correct format, use a service such as https://epochconverter.com or https://unixtime.org.

Provide feedback

To provide feedback, contact Netskope.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

OpenText Core Threat Detection and Response

Summarize this article for me

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft doesn't provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Your OpenText Core Threat Detection and Response subscription enables you to use this plugin to interact with Security Copilot, gain data insights produced by the product, and take appropriate actions on the risky entities and alerts occurring in your organization.

- Summarize risky activity across the organization.
- Retrieve top risky users, devices, and rare processes.
- Examine specific risky entity and rare process activities.

Prerequisites

1. Sign in to your OpenText Core Threat Detection and Response account.
2. Make a note of the <Product_URL> used to sign into your account. Save this URL for plugin configuration.
3. Generate an API token from your account settings. Save the token securely for plugin configuration.

Know before you begin

1. Open your API token file and locate the **access token** field. Keep the file open. You need to copy and paste the access token value later in this procedure.
2. Sign in to Microsoft Security Copilot.
3. Select the **Sources** button.
4. From the **Manage sources** dialog box, scroll down to the **Non-Microsoft** source section.

5. Navigate to the **OpenText Core Threat Detection and Response** plugin and select **Set up**.
6. In the **OpenText Core Threat Detection and Response settings** dialog box do the following:
    a. In the **Instance URL** box, enter the <Product_URL> for your tenant instance. The URL must be of the format `https://<domain name>/`.
    b. In the **Value** box, enter the access token value of the API token you downloaded. **Note:** Don't include the quotation marks when you copy the value of the API access token.
    c. Select **Save**. The Non-Microsoft plugins area displays the OpenText Core Threat Detection and Response plugin.
7. Configuration is complete. You can begin to submit prompts.

Sample OpenText Core Threat Detection and Response prompts

The skills listed describe the full interaction of this version of Security Copilot plugin with OpenText Core Threat Detection and Response.

Use natural language prompts. Examples are provided below.

Skills & Prompts

- **Top Risky Users**

  Prompt: What are the top five riskiest users on <date>?

- **Top Risky Devices**

  Prompt: What are the top five riskiest devices on <date>?

- **Top Rare Processes**

  Prompt: What are the top five riskiest rare processes executed on <date>? Include alert IDs.

- **Summarize Risky Activity**

  Prompt: Summarize the risky activity across the organization on <date>.

- **Entity Investigation**

Prompt: Summarize the risky behaviors of <username|hostname> on <date>.

- **Alert Details**

  Prompt: What are the details of the alert with ID <alert_id>?

- **Insider Threat Summary**

  Prompt: Can you summarize the insider threats to the above investigation and provide a conclusion and recommendation?

Troubleshoot the OpenText Core Threat Detection and Response plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*. Make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Prompts are failing

If the prompts fail to invoke, ensure that you are using a supported prompt. You can refer to the promptbook shared with you to understand the valid prompts. If you're using the correct prompts and yet the prompts fail to invoke, then the plugin service might not be running. Wait for some time for the service to restart, and then try again. If you still face issues, contact OpenText Support.

Provide feedback

To provide feedback, contact https://www.opentext.com/products/core-threat-detection-and-response.

See also

Other plugins for Microsoft Security Copilot Manage plugins in Microsoft Security Copilot

Last updated on 02/23/2026

Pure Signal Scout

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

Pure Signal Scout The Team Cymru's Pure Signal Scout plugin enables Security Copilot customers to gather detailed threat intelligence on IP addresses and domains in real-time. Scout delivers unparalleled speed, accuracy, and visibility, helping security analysts make faster, better-informed decisions. Scout's powerful features ensure that security teams can stay ahead of evolving cyber threats. It provides insights into malicious IP activity, domain information, open ports, communication, and more. With AI-enriched data from sources like Passive DNS and cryptographic certificates, it helps security teams detect and block harmful IPs, streamlining investigations and improving response times.

Know before you begin

Integration with Security Copilot works with an API key that you must retrieve before using the plugin. You can create a new user account or use your existing account to get the key.

Set up your Pure Signal Scout account and API key

1. Go to Scout Insight trial and fill out the required details.
2. Verify your email and set a password.
3. Sign into the Pure Signal Scout portal.
4. Go to the API Keys page.
5. Select **Create**.
6. (Optional) Add a description for the key.
7. Select **Create Key** to generate the key.

8. If the **Create** button is disabled, your organization has reached the maximum number of keys. In this case, follow these steps:

    a. Next to an old key, select **Revoke**.

    b. Select **Create Key** to start generating a new key.

Use your existing account and get your API key

1. Sign into the Pure Signal Scout portal.
2. Go to the API Keys page.
3. Select **Create**.
4. (Optional) Add a description for the key.
5. Select **Create Key** to generate the key.
6. If the **Create** button is disabled, your organization has reached the maximum number of keys. In this case, follow these steps:

    a. Next to an old key, select **Revoke**.

    b. Select **Create Key** to start generating a new key.

Configure the Pure Signal Scout Plugin in Security Copilot

1. Sign in to Microsoft Security Copilot.
2. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.
3. Next to **Pure Signal Scout Plugin**, select **Set up**.

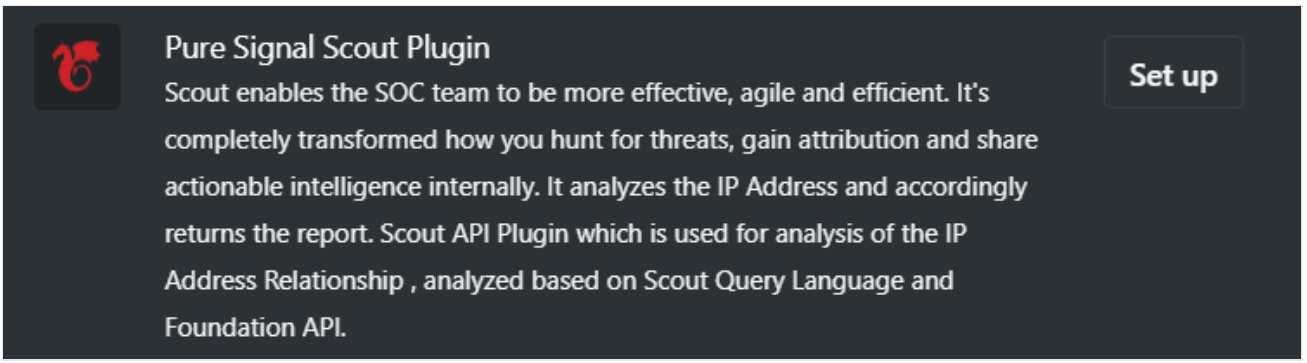

4. In the **Value** field, paste the Pure Signal Scout API key, and then select **Save**.

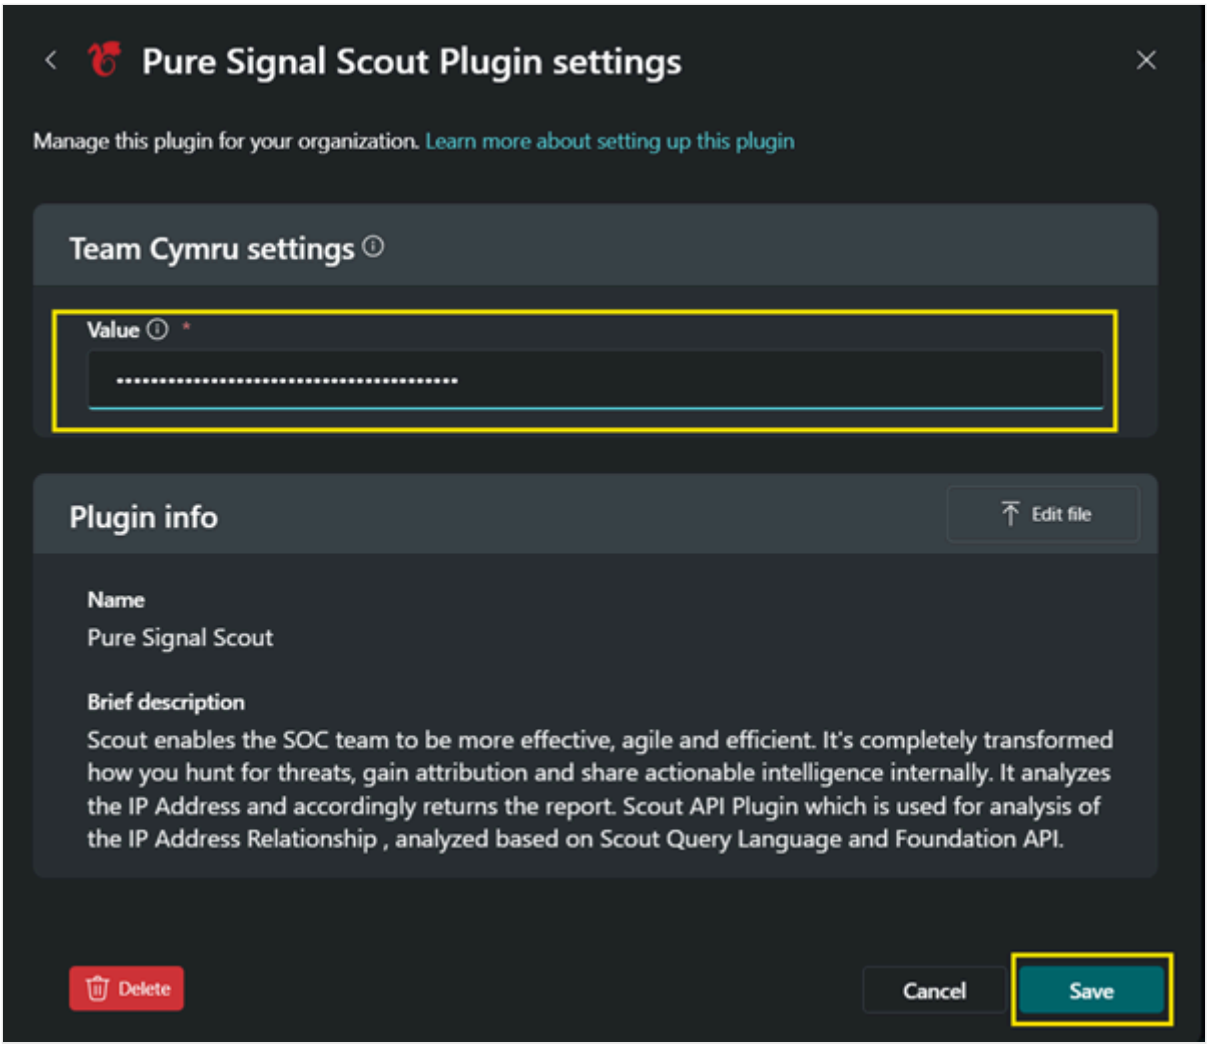

Sample Pure Signal Scout prompts

Expand table

Capability	Description	Input Parameters	Example Prompts
`ScoutFoundationAPI`	Accepts IP addresses and enables bulk analysis, offering insights into whether the set contains suspicious, malicious, or informational IPs. The Scout Foundation API also provides AS information, country codes, and key tags associated with IP addresses.	Required: ips (max 10 IP Addresses)	- `List down the malicious and suspicious IPs from these 8.8.8.8 175.155.2.48 185.220.101.101 192.42.116.175 188.165.200.97 185.220.101.88 178.20.55.182 104.182.36.17 with help of Pure Signal Scout plugin.` - `Using the Pure Signal Scout plugin, find if the IP Address 185.220.100.240`

Capability	Description	Input Parameters	Example Prompts
			`is malicious ?` - `List down key tag associated with IP Address 12wd85.220.100.240 using Scout plugin.`
`ScoutIPDetailsAPI`	The Scout IP Details API provides comprehensive information about a specific IP address, covering details such as identity, network communication history, passive DNS data, open ports, X.509 certificates, TLS/SSL fingerprints, and WHOIS records. This endpoint enables users to retrieve a full report on an IP address's behavior and relationships over time by specifying start and end dates or selecting a range of days.	Required: `IP AddressOptional: start_date` : `start_date` can't be more than 90 days back from current date & 30 days before end date `end_date :` can't be in the future Days: min: 1, max: 30 (Relative offset in days from current time in UTC. It can't exceed the maximum range of 30 days.) `size:` default: 100, min: 1, max: 1000 (The size of the response, in records, to return) `sections: identity` comms pdns open_ports x509 fingerprints whois summary proto_by_ip	- `Using Pure Signal Scout to find what open ports are available on this IP address 47.156.224.38?` - `Are there any unusual communication patterns for this IP address 47.156.224.38? Check with the Pure Signal Scout plugin.` - `What are the most frequent destinations for this IP address 47.156.224.38? Find using Scout plugin.` - `Using Scout plugin find what are the connections between this IP address 47.156.224.38 and specific ASNs?` - `Has this IP address 175.155.2.48 been seen in any honeypot data? Find using Scout.` - `What are the potential threats associated with this IP address 175.155.2.48?` - `What are the country/region origins of IPs communicating with this one IP 47.156.224.38?`

Capability	Description	Input Parameters	Example Prompts
`ScoutSearchAPI`	The Scout Search API provides detailed information about domains and supports advanced search queries using the Scout Query Language. It returns results that may include country codes, autonomous system (AS) information, tags, WHOIS data, open ports, passive DNS (PDNS), communication details, service information, X.509 certificates, and fingerprints. This API allows users to build a query using various formats, including IP address, domain names, websites, or advanced queries with specific selectors (for example, pdns.domain). Full details on the available search selectors can be found at Scout documentation.	Required: query Optional: `start_date:` can't be more than 90 days back from current date & 30 days before end date `end_date:` can't be future Days: min: 1, max: 30 (Relative offset in days from current time in UTC. It can't exceed the maximum range of 30 days.) size: default: 100, min: 1, max: 5000 (The size of the response, in records, to return)	- `Using Pure Signal Scout to find what is the WHOIS information for this domain 10crypto.top ?` - `Using the Pure Signal Scout plugin can you show me the historical DNS data for this domain akamai.com?` - `What are the most recent WHOIS updates for this subtitleseeker.com? Use Scout plugin.` - `Using Scout to find what is the reverse WHOIS information for this domain 10crypto.top?` - `Use Scout to run the query pdns.domain="*ngrok.io" and give the certificate details about top 10 associated IPs.` - `What organization is associated with subtitleseeker.com in the WHOIS data? Use Scout.`

Troubleshoot the Pure Signal Scout plugin

Errors occur

If you encounter errors like "Couldn't complete your request," try these steps to troubleshoot:

1. Start a new session to refresh context and attempt the prompt again in a new session.
2. Specify a Detailed Prompt: Refer to the best prompt guidelines for Security Copilot here and crafting detailed prompts according to Pure Signal Scout plugin skills and capabilities.

3. Adjust the Size Parameter: If the issue persists, reduce the size parameter in your prompt to control response size, for example, "Use size as 10." This can also be adjusted in the IP Details API or Scout Search API to help minimize response payloads.

If the issue continues, sign out of Security Copilot, then sign back in and try again.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities or if they seem to be activating a different plugin, it may be due to other plugins or custom plugins that provide similar capabilities as Pure Signal Scout. For instance, if you have multiple plugins that offer threat intelligence or domain information, conflicts might arise. To prioritize and specifically Pure Signal Scout, consider disabling other custom plugins. Alternatively, you can either use the product name **Pure Signal Scout** in your prompts or specify a particular skill.

Provide feedback on the Pure Signal Scout plugin

To provide feedback, contact Pure Signal Scout.

Related content

Use the following Security Copilot articles for guidance:

Other plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Last updated on 05/29/2026

Quest Security Guardian

08/25/2025

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

Quest Security Guardian is an Active Directory security tool designed to reduce your attack surface. From a simplified unified workspace, Security Guardian reduces alert fatigue by prioritizing your most exploitable vulnerabilities and Active Directory configurations that demand attention. The solution spotlights what happened if you're exposed and how to fix the problem.

The Security Guardian Plugin exposes all of Security Guardian's threat indicators, findings, and alerts to Microsoft Security Copilot. This allows all the data to be combined with other security findings (increasing signal strength) and ultimately acted on by Security Analysts to identify and fix vulnerabilities faster.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

You'll need to take the following steps before using the plugin.

1. Sign in to Microsoft Security Copilot.
2. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.
3. Next to **Quest Security Guardian**, select the toggle to enable it.

   Provide the following information:

   - Tenant Id: The ID of the Microsoft Entra ID Organization that the Microsoft Sentinel workspace is in.

- Subscription Id: The ID of the Azure Subscription that the Microsoft Sentinel workspace is in.
- Resource Group Name: The name of the Resource Group that the Microsoft Sentinel workspace is in.
- Workspace Name: The name of the Microsoft Sentinel workspace.

4. Save your changes.

Sample Quest Security Guardian prompts

After the Quest Security Guardian plugin is configured, you can use it to:

- Access the plugin capabilities by selecting the **Plugin button** from the prompt bar and selecting **Quest Security Guardian**.
- Prompt Security Copilot using any of the following example prompts.

Example prompts:

- `Generate a report using Security Guardian and summarize the top 5 trends the indicators show and then summarize how these impact our company and suggest a couple of ways to resolve or mitigate them.`

Troubleshoot the Quest Security Guardian plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*. Make sure the plugin is turned on. This error may occur if the lookback period is too long, causing the query to attempt to retrieve an excessive amount of data. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Provide feedback

To provide feedback, contact Quest Security Guardian.

See also

Other plugins for Microsoft Security Copilot Manage plugins in Microsoft Security Copilot

ReversingLabs

08/25/2025

ReversingLabs provides empowers SOC teams to understand the file-based threats in an environment. Analysts can use Security Copilot to summarize complex file reputation information and file analysis reports for quicker triage and response time.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot requires a Spectra Intelligence user account. You'll need to take the following steps before using the plugin.

1. Get your ReversingLabs user account and password.
2. Sign in to Microsoft Security Copilot.
3. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.
    - If you use, **ReversingLabs Spectra Analyze**, select "Set Up" to configure and enable it.
    - If you use, **ReversingLabs Spectra Intelligence**, select "Set Up" to configure and enable it.
4. In the ReversingLabs settings pane, provide your ReversingLabs user account and password.
5. Save your changes.

Sample ReversingLabs Spectra Intelligence prompts

After the ReversingLabs Spectra Intelligence plugin is configured, you can use it by typing `ReversingLabs` in your Security Copilot prompt bar, followed by an action. The following table provides several examples you can try:

Expand table

Capability Prompt Suggestion	Description	Example natural language prompt
`GetFileHashReputation`	This capability is used to retrieve file hash reputation information.	`What is the reputation of hash a6e728c3331f46763f643f7192959716034767e5?`
`GetDetailedFileAnalysisResults`	Retrieve additional analysis details for the supplied file hash.	`Get the MITRE ATT&CK techniques from the detailed file analysis of a6e728c3331f46763f643f7192959716034767e5.`
`GetDynamicAnalysisReport`	Retrieve the full sandbox dynamic analysis report for the supplied file hash.	`Get the dynamic analysis report for a6e728c3331f46763f643f7192959716034767e5 and summarize any new network connections created.`

Sample ReversingLabs Spectra Analyze prompts

After the ReversingLabs Spectra Analyze plugin is configured, you can use it by typing `ReversingLabs` in your Security Copilot prompt bar, followed by an action. The following table provides several examples you can try:

Expand table

Capability Prompt Suggestion	Description	Example natural language prompt
`GetFileClassification`	This capability is used to retrieve file hash classification information from a Spectra Analyze appliance.	`What is the classification of hash a6e728c3331f46763f643f7192959716034767e5?`
`GetTiCoreReport`	Retrieve the TitanumCore analysis report for the provided file hash from a Spectra Analyze appliance.	`Get the TiCore report for hash a6e728c3331f46763f643f7192959716034767e5`

Frequently Asked Questions (FAQ)

Why are prompts failing to invoke ReversingLabs?

If prompts fail to invoke, make sure you're using a supported prompt (see the preceding examples). Otherwise, invoke ReversingLabs directing by using `/`.

Why are prompts failing to return data?

If prompts fail to return data, check your Spectra Intelligence usage limits.

Provide feedback

To provide feedback, contact ReversingLabs.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Saviynt

08/25/2025

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

Saviynt offers comprehensive insights into identity-related risks. You can use the Saviynt plugin with Microsoft Security Copilot to effectively enrich incidents and identify risky user access.

Some of the capabilities it provides include:

- Obtaining user profile details.
- Retrieving vendor or third-party organization details.
- Accessing comprehensive user access details across applications.
- Identifying dependencies of the access they possess.
- Checking if the identity holds privileged roles.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Prerequisites

Saviynt Identity Cloud API Access: Refer to Saviynt API Reference for details.

Know before you begin

Integration with Security Copilot requires a bearer token for API access authorization. You'll need to take the following steps before using the plugin.

1. Follow the steps below to obtain the authorization token:
   a. Login to the Saviynt Tenant using your service account credentials.
   b. Navigate to **Admin** > **Settings** > **Webservice Auth Configuration**.
   c. Select **Generate Token**.

d. Copy the value for the access_token field in a safe place.

> **Note**
>
> Alternatively, you can use Postman to generate the authorization token using the {{url}}/ECM/api/login API. Refer to Saviynt's API documentation for more details.

2. Sign in to Microsoft Security Copilot.
3. Access **Manage Plugins** by selecting the **Sources** button from the prompt bar.
4. Next to **Saviynt**, select **Set up**.
5. In the Saviynt settings pane:
   - In the **InstanceURL** field, enter the Saviynt tenant URL (for example, "https://saviynttenant.saviyntcloud.com")
   - In the **Value** field, paste your access_token, and then select **Save**.

Sample Saviynt prompts

After the Saviynt plugin is configured, you can use the following capabilities with Security Copilot.

Expand table

Capability	Example prompts

Users	`Get user details for <First Name> <Last Name> using the search criteria FirstName and LastName`
	
	`Get user details for <username>`
	
	`Get risk associated with user <username>`

Organizations	`Get the organization details for <organization name>`
	
	`Get a list of organizations that have not been certified for the last six months`

Dependent Entitlements	`Get dependent entitlements for "<entitlement name>" for <application>`
Access Details	`Get the access details for user <username>`

Capability	Example prompts
Saviynt Role details	`List the users who have access to Sav Role <Sav Role Name>`

Provide feedback

To provide feedback, contact Saviynt.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

ServiceNow plugin for Microsoft Security Copilot

08/18/2025

> **Important**
>
> Some information in this article relates to a prereleased product, which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

ServiceNow is an enterprise app ecosystem designed to connect and automate business processes, delivered as a SaaS application. Many security operation centers (SOCs) use ServiceNow as part of their incident management flow, and often extend the core functionality with customizations, integrations, and security-specific applications.

The Copilot for Security plugin for ServiceNow enables connectivity between a Security Copilot session and a ServiceNow incident queue. Users can import a ServiceNow incident into Copilot for Security. Users can easily correlate data from Microsoft security products such as Defender for Endpoint, enrich with external threat intelligence, and persist the results of their investigation in ServiceNow. These capabilities, combined with the narrative prompt and generative AI powered by Azure OpenAI help speed incident resolution. These capabilities also help up-skill team members, and provide more comprehensive views into any security incident.

Know before you begin

Prerequisites

- Supported Products (Cloud):
  - ServiceNow IT Service Management (ITSM) package – Incident Response application
  - ServiceNow Security Operations (SecOps) package – Security Incident Response application
- Access to your ServiceNow instance, with permissions to create new users
- Permissions to set up plugins
- Available ServiceNow API quota

ServiceNow setup

There are two connectivity methods to ServiceNow:

- OAuth authorization code grant flow
- HTTP Basic auth

Choose the one that suits your needs.

OAuth authorization code grant flow (ServiceNow inbound integration) - Recommended

> **Note**
>
> Before you begin, ensure you have the ServiceNow admin permission and the Security Copilot Owner or Contributor permission.

1. Make sure the OAuth plugin is active and the OAuth activation property is set to true. Follow the links to activate both if they aren't activated by default.

2. Access your ServiceNow instance and create a new Application Registry object by following these steps:

   a. Navigate to **System OAuth** > **Application Registry**, select **New**.

   b. Select **Create an OAuth API endpoint for external clients**.

   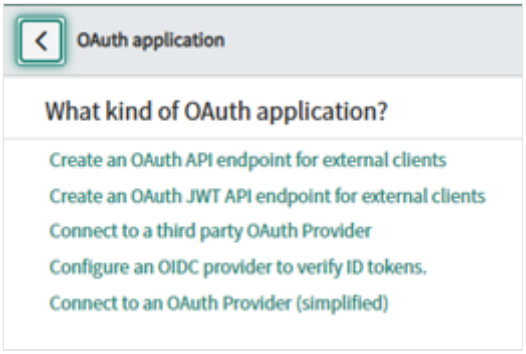

   c. Enter a name that identifies the Security Copilot integration.

   d. Enter the redirect URI corresponding to the Security Copilot instance, for example `https://securitycopilot.microsoft.com/auth/v1/callback`.

   e. Assign an auth scope corresponding to your access control needs. The user account scope would suffice.

   f. Save the Application Registry object.

   g. From the newly created object, note the client ID and client secret.

h. This OAuth setup should only be done once per ServiceNow instance. The resulting OAuth Application Registry object can be used multiple times by different users.

HTTP Basic auth

1. Access your ServiceNow instance and locate the option to create a new user:
   a. Navigate to **Discovery** > **Credentials** in the ServiceNow platform.
   b. Select **New** to create a new credential record.
   c. Select **Basic Auth** as the credential type. Enter the following details:
      i. **Name:** Descriptive name for the credential.
      ii. **Username:** Username for authentication.
      iii. **Password:** Password for authentication.
      iv. Save the credentials.

2. Assign the `itil` and `rest_api_explorer` roles to the username:
   a. Navigate to **User Administration** > **Users**.
   b. Search for and select the username created.
   c. In the **Roles** related list, select **Edit**.
   d. Add `itil` and `rest_api_explorer` roles from the available roles list.
   e. Save the role assignments.

3. Note the credentials and the ServiceNow instance URL.

Security Copilot connection

1. Sign in to Microsoft Security Copilot.

2. Access **Manage Plugins** by selecting the **Sources** button from the prompt bar.

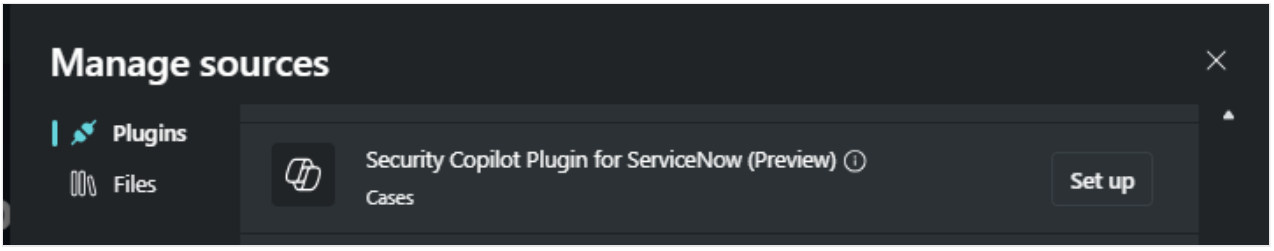

3. Next to **ServiceNow**, select **Set up**.

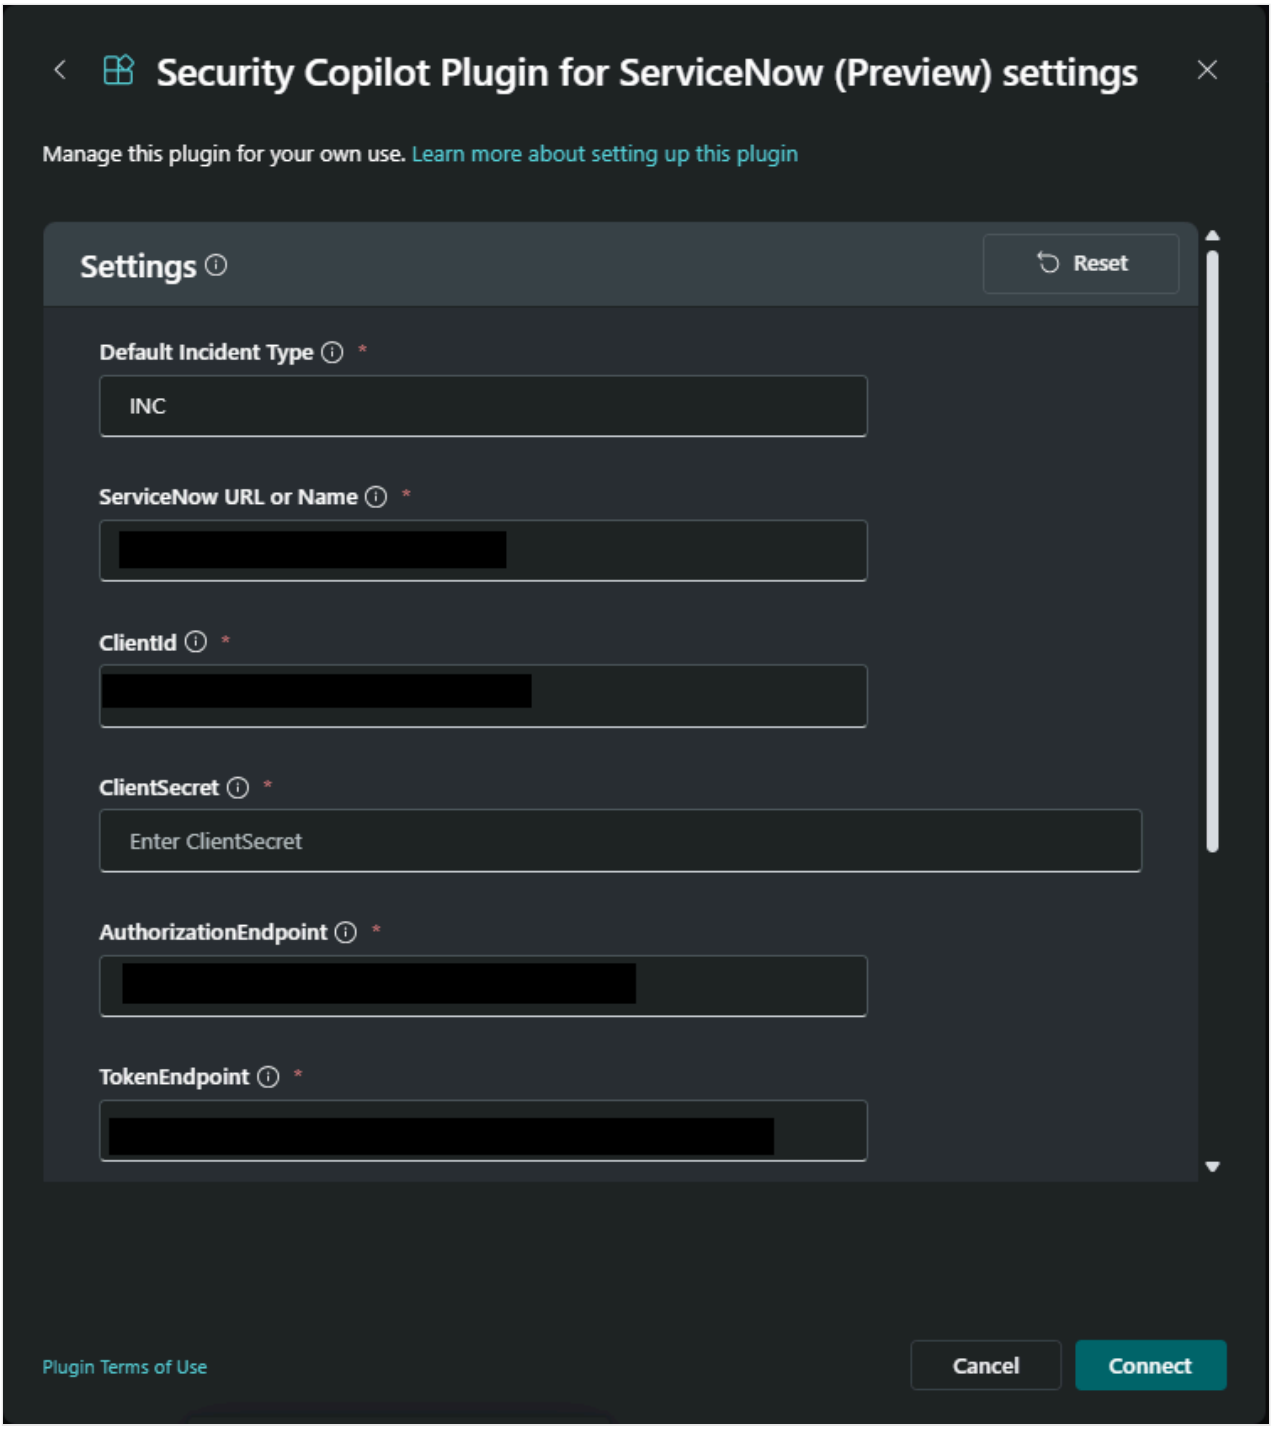

> **Note**
>
> Be sure to scroll all the way down to reach other fields such as Scopes.

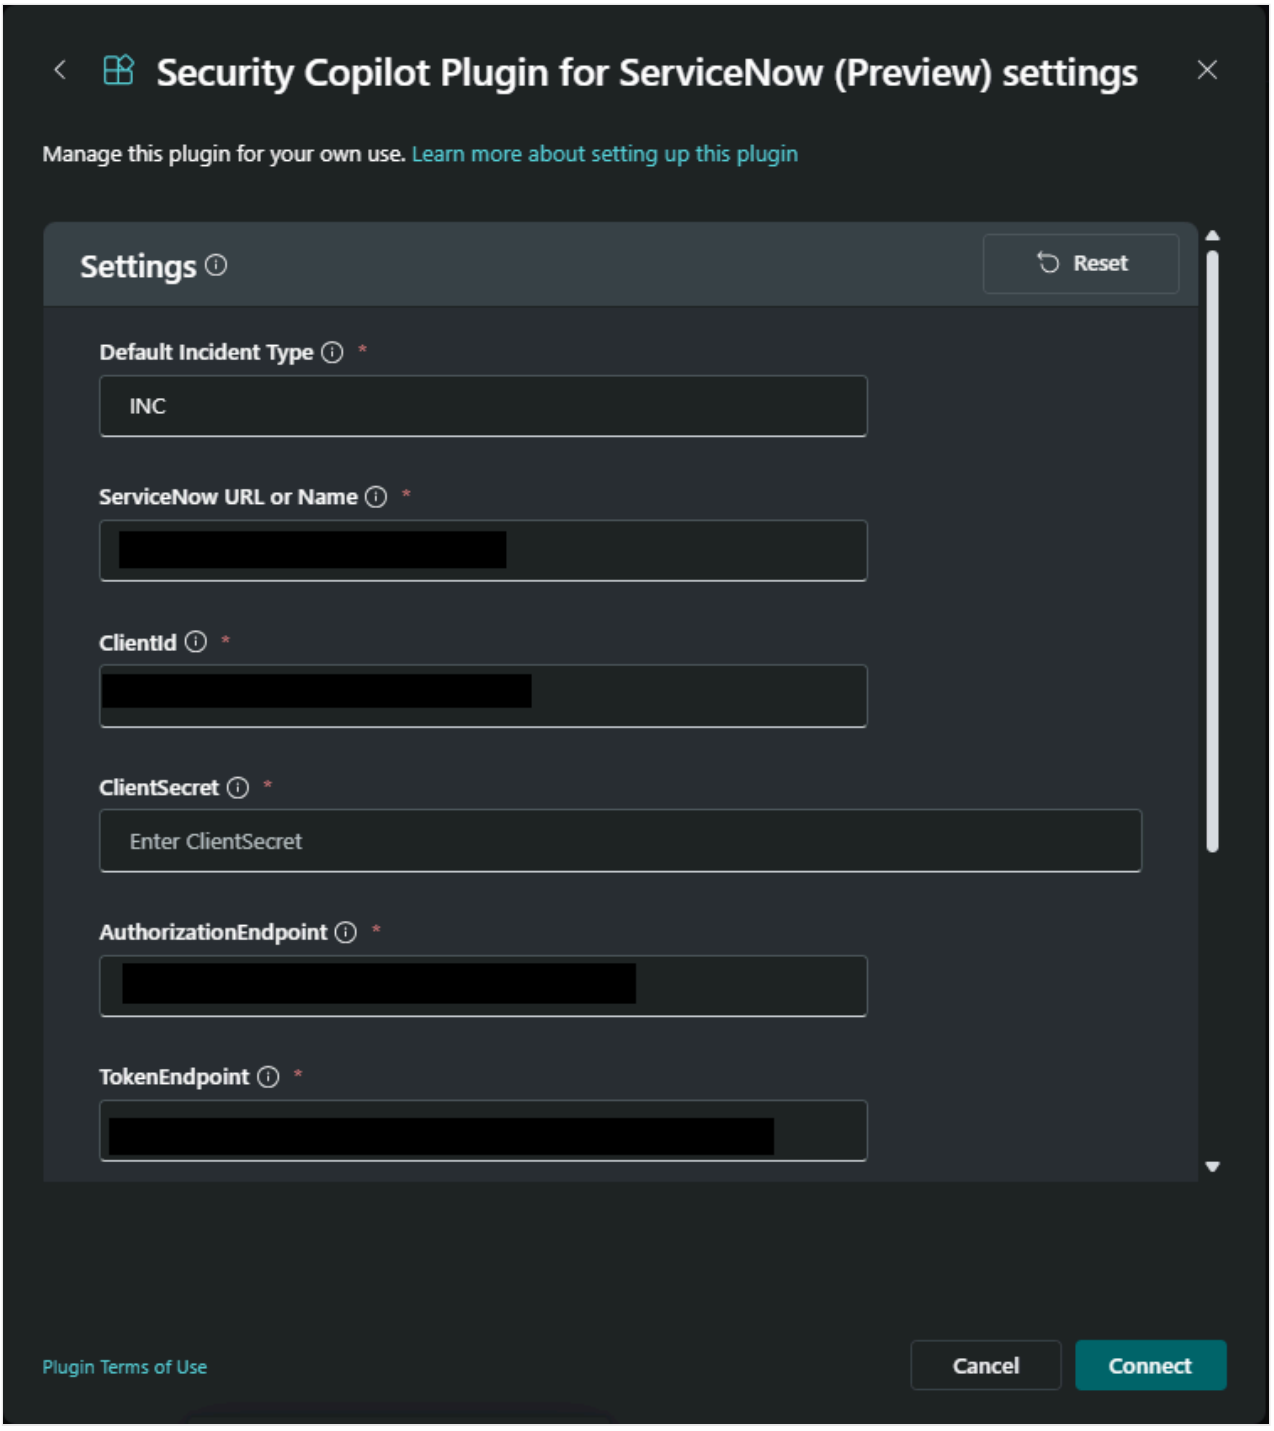

4. Grant consent to the Security Copilot application so that you can access your ServiceNow instance.

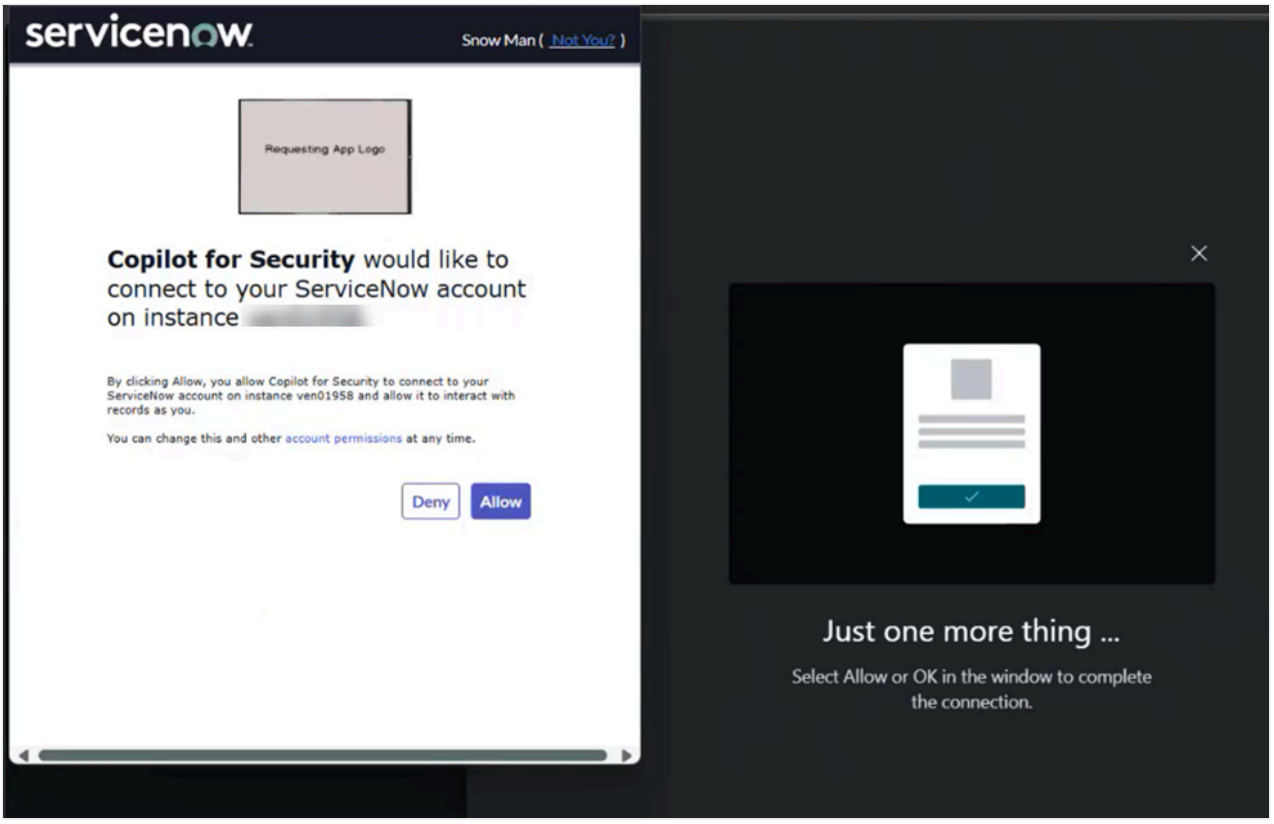

5. OAuth authorization_code instructions
   a. Enter auth parameters corresponding to the Application Registry object you created earlier.
   b. The instance URL, client ID, and client secret correspond to values of the same name in the Application Registry object. To determine what to enter for each value, refer to the following table:

Expand table

Setting name	Description	Example
Default Incident Type	Incident type to default to when the type isn't provided within prompts. Must be one of: INC or SIR	INC or SIR
Instance	Set to the ServiceNow instance URL	`https://xyz.service-now.com/`
ClientId	Set to client ID in ServiceNow Application Registry object	
ClientSecret	Set to client secret in ServiceNow Application Registry object	
AuthorizationEndpoint	Set to `https://<service-now-instance-domain>/oauth_auth.do`.	`https://xyz.service-now.com/` oauth_auth.do

Setting name	Description	Example
TokenEndpoint	Set to `https://<service-now-instance-domain>/oauth_token.do.`	`https://xyz.service-now.com/` oauth_token.do
Scopes	Set to the scopes defined in the ServiceNow Application Registry object. Comma-separate multiple scopes.	user account
AuthorizationContentType	Should be left unchanged from default application/x-www-form-urlencoded	application/x-www-form-urlencoded
Resource	Resource ID	Can be left blank

6. Select **Save** or **Connect** to complete the setup.

> ⓘ **Note**
>
> Currently, the system doesn't validate your credentials when you save your settings. If they aren't correct, you'll see an error later when Security Copilot attempts to invoke the ServiceNow plugin.

7. Close the plugins window.

Sample ServiceNow prompts

Security Copilot operates primarily with natural language prompts. When you're ready to load an incident into your Security Copilot session or search for related ServiceNow incidents, you'll submit a prompt that guides Security Copilot to select the ServiceNow skill set and invoke the proper skill.

When you word your prompt, ensure you mention ServiceNow as the preferred source of your incident. Security Copilot can connect to several systems that each provide incidents, and it benefits from guidance on which source you prefer.

Example prompts for incident queries:

- "Load ServiceNow incident INC12345"
- "What ServiceNow incidents refer to IP address 10.0.0.1?"
- "Show me recent high severity ServiceNow incidents"
- "Show details on the third incident"
- "What is the MDTI reputation score for those IP addresses?"
- "Write a link to this Copilot investigation to ServiceNow incident INC12345"

- "Write the following text to that ServiceNow incident: Investigated with Security Copilot and deployed new detection logic."
- "Write a summary of this investigation to ServiceNow incident INC12345"
- "Summarize this investigation and write it as a comment on the ServiceNow incident."

Appending comments to ServiceNow incidents

If enabled with appropriate permissions, the ServiceNow connector can append comments to ServiceNow incidents. The comment text can be provided by the user, or sourced from Security Copilot features, such as session sharing links or pinned prompt investigation summaries.

Example prompts include:

- "Write a link to this Copilot investigation to ServiceNow incident INC12345"

After loading an incident:

- "Write the following text to that ServiceNow incident: Investigated with Security Copilot and deployed new detection logic."

After pinning a few prompts:

- "Write a summary of this investigation to ServiceNow incident INC12345" or if the session is already loaded "Summarize this investigation and write it as a comment on the ServiceNow incident."

> **Note**
>
> The current version of the plugin appends comments to the ServiceNow incident. Future versions may provide the option of writing to the work notes instead.

Troubleshoot the ServiceNow plugin

Ensure the relevant IPs for your region from the following list are allowed. For more information, see Egress IP addresses for Microsoft Security Copilot.

Steps to add Allowed IPs:

1. Sign in to your ServiceNow Instance:
    - Use an account with administrative privileges.

2. Navigate to the IP Access Control Settings. ServiceNow documentation on the steps: IP Address Access Control :

    - In the left navigation pane, search for "IP Access Control" or go to: **System Security > IP Access Control**

    Review the IP addresses from this list and select the appropriate one for your region and add it to the allowlist with direction Type as inbound:

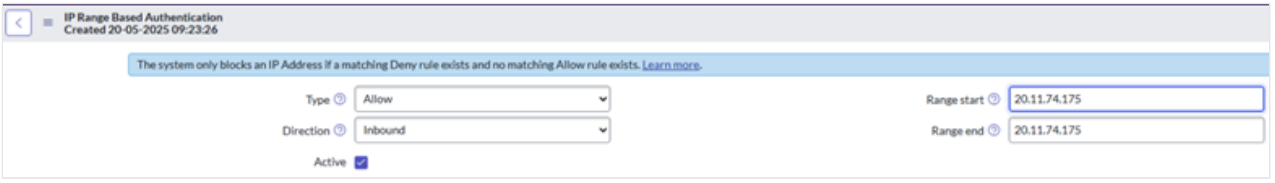

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*, make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use. You can either use the product name ServiceNowSIR in your prompts, or type the name of a specific capability, like <> instead.

Provide feedback

To provide feedback, contact ServiceNow product management.

ServiceNow SIR

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

In today's rapidly evolving threat landscape, organizations need cutting-edge security solutions that combine AI-driven intelligence with automated response capabilities. The integration between ServiceNow Security Incident Response (SIR) and Microsoft Security Copilot redefines security operations, enabling dynamic decision-making, real-time insights, and autonomous response capabilities. By implementing this game changer of a solution, organizations can strengthen their security posture, reduce cyber risks, and optimize security operations for the future.

ServiceNow Security Incident Response and Microsoft Security Copilot in action!

To seamlessly integrate ServiceNow Security Incident Response with Microsoft Security Copilot, enabling AI-driven, autonomous security operations that enhance threat detection, accelerate incident response, and improve overall cybersecurity resilience through intelligent AI to AI communication.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Scope

1. Enrich incidents in SIR with threat intelligence and user/device information to get the contextual awareness for incident remediation
2. Share the summary of the security incident from SIR along with corelated insights on user and device information from CMDB to Microsoft

Value Proposition

1. Accelerated Investigation and Response: Analyze threats with AI-assisted insights and swiftly transition to remediation using contextual intelligence
2. Improved Operational Efficiency: Reduce the Mean Time to Resolve (MTTR) with natural language enable AI assistance

3. Stronger Security Posture: Use AI-enhanced threat intelligence to proactively identify and mitigate risks before they escalate.

Integration Approach

1. AI to Ai communication between SIR and Microsoft Security Copilot.
2. Customers would be able to ask quick queries on ServiceNow Now Assist Panel on the threat intel, user and device information from Microsoft.
3. Similarly, customers would be able to ask for summary of security incidents along with corelated insights of users and device from CMDB on Microsoft security copilot

Know before you begin

Register the application in ServiceNow

1. Navigate to **System oAuth > Application Registry > New > Create an OAuth API endpoint for external clients**.
2. Set the Name: Microsoft Security Copilot ServiceNowSIR.
3. Set Redirect URL: https://securitycopilot.microsoft.com/auth/v1/callback.
4. **Submit**.

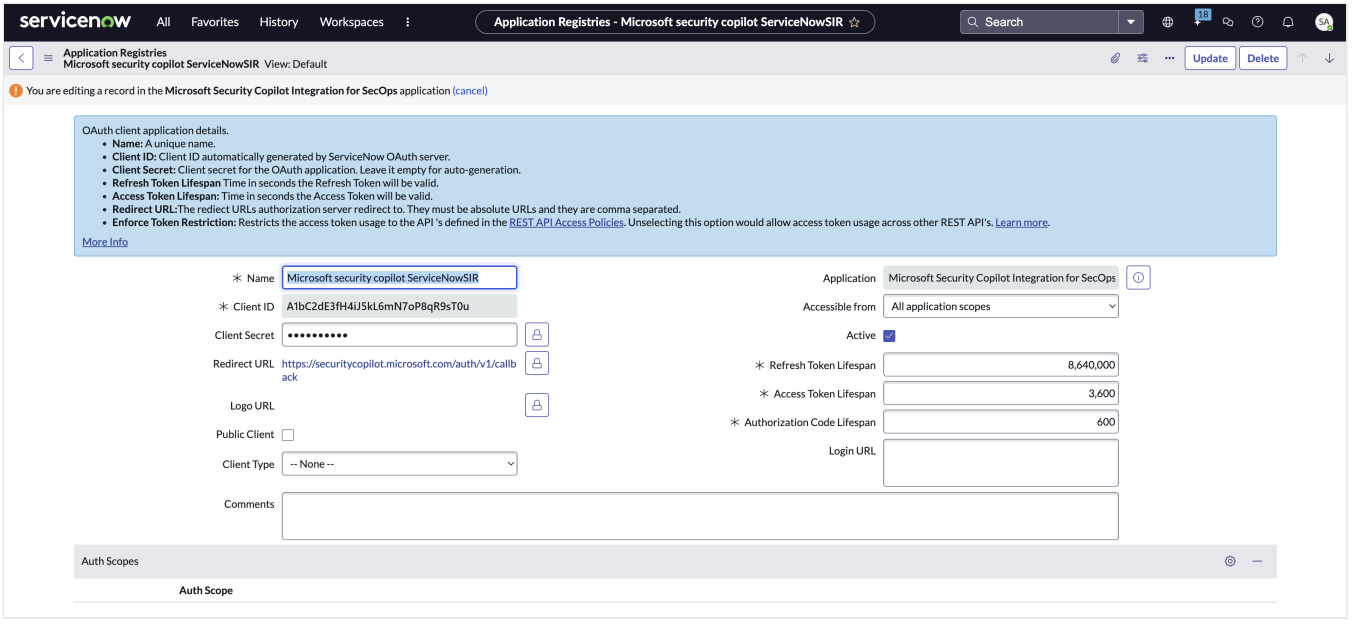

Set up the ServiceNowSIR plugin in Microsoft Security Copilot

1. Set the ServiceNow Instance URL.
2. Set the **ClientId** and **ClientSecret** generated in the previous section.
3. Set the **AuthorizationEndpoint** and the **TokenEndpoint**.

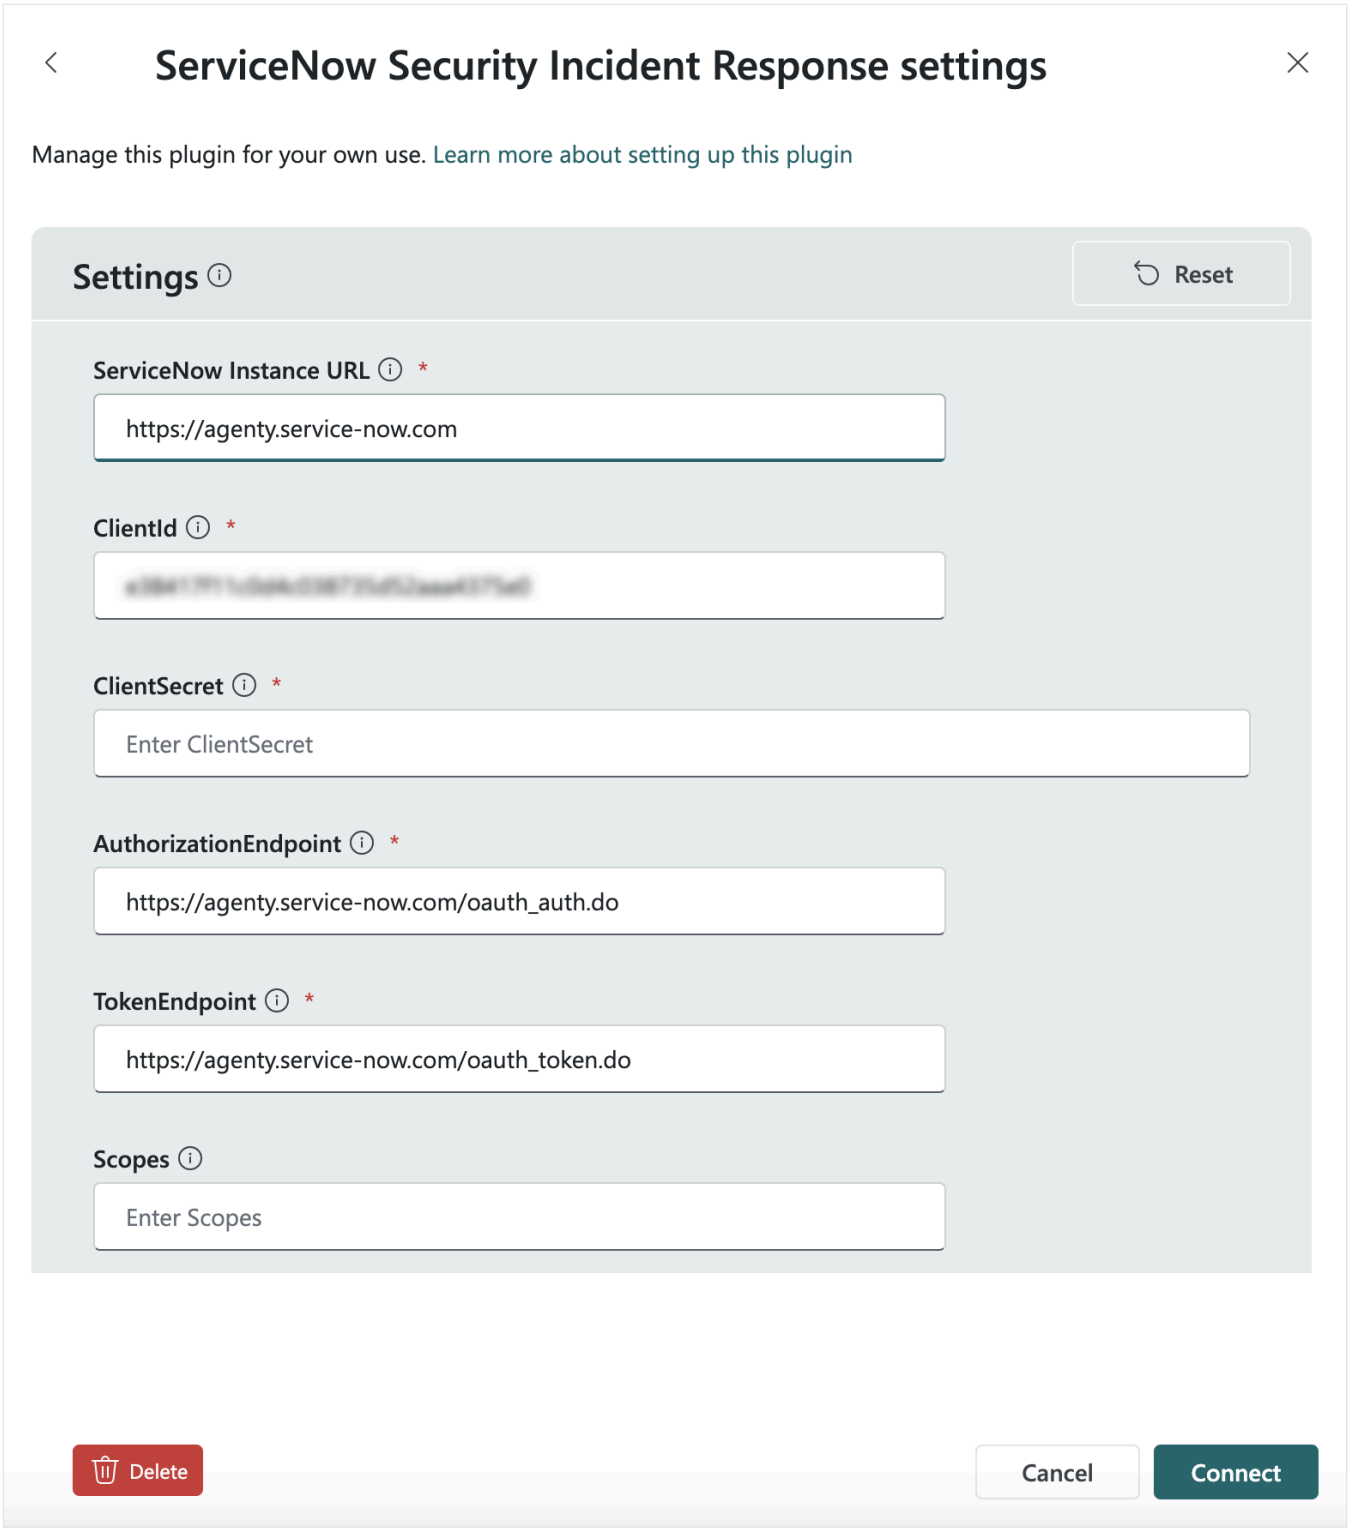

Sample ServiceNowSIR prompts

After the ServiceNowSIR is configured, you can use it by taking one of the following steps:

Expand table

Skill	Inputs	Example Prompts
Get security incident details	• security incident number (required)	• Load ServiceNow security incident SIR0001624 • Get ServiceNow the description and short description of security incident SIR0001624
Get security incident summarization	• security incident number (required)	• Get ServiceNow security incident SIR0001624 summary
Get correlation insights for a security incident	• security incident number (required) • variable (required) • name (required) • lookback_period (optional)	• Get ServiceNow security incident SIR0010081 correlation insights with variable cmdb_ci, name DatabaseServer2 and lookback_period=365 • Get ServiceNow security incident SIR0010081 correlation insights with variable affected_user, name Alice Brown and lookback_period=365 • Get ServiceNow security incident SIR0010081 correlation insights with variable threat_intelligence and name `https://keyloggerapp.xxx.com` and lookback_period=365
Get user by name	• name (required)	• Get ServiceNow user Alice Brown • Get the department and location of ServiceNow user Alice Brown
Get CMDB CI by name	• name (required)	• Get ServiceNow configuration item DatabaseServer2 • Get ServiceNow the business unit of configuration item DatabaseServer2 • Get ServiceNow the location and Manufacturer of configuration item DatabaseServer2

Troubleshoot the ServiceNowSIR plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*, make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use. You can either use the product name ServiceNowSIR in your prompts, or type the name of a specific capability, like <> instead.

Provide feedback

To provide feedback, contact ServiceNowSIR product management.

Last updated on 02/19/2026

SGNL

08/25/2025

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

SGNL provides a dynamic access platform that provides a foundation for Zero Standing Privilege (ZSP) initiatives to protect against threats to user sessions and credentials. You can use the SGNL Access Intelligence plugin with Microsoft Security Copilot to understand and identify fine-grained access decisions and trends across your organization.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot requires your SGNL tenant ID and ID token. Make sure your SGNL account has sufficient permissions to read the SGNL Access Decision Logs. You'll need to take the following steps before using the plugin.

1. Sign in to Microsoft Security Copilot.
2. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.
3. Next to **SGNL**, select **Set up**.
4. Paste your SGNL tenant ID in the Username field, and your ID token in the password field. Then select **Save**.

Sample SGNL prompts

After the SGNL plugin is configured, you can use it to summarize logs and identify who accesses sensitive assets. The following table lists some capabilities and example prompts you can try:

Expand table

Capability	Example prompt
Summarize SGNL logs	`Summarize the logs from SGNL for today`
Identify principals who accessed sensitive assets	`Which principals have accessed sensitive assets today`
Identify who accessed a specific asset	`Which users have accessed <Asset>`
Summarize a user's access for a specific time period	`Tell me about <user> access over the last 7 days`

Troubleshoot the SGNL plugin

Errors occur

If you encounter errors occur, such as: *Couldn't complete your request*, or *An unknown error occurred*, make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Provide feedback

To provide feedback, contact SGNL.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Shodan

08/25/2025

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

Shodan is a search engine that allows users to find specific types of devices connected to the internet using various filters. It provides a global view of how certain devices are connected and can be used to discover which devices are connected to the internet, where they're located, and who is using them. You can use the Shodan plugins with Microsoft Security Copilot to get enhanced visibility of their internet-facing assets and better detect threats and vulnerabilities.

If you do not have a Shodan membership and/or account, you can use the "Shodan InternetDB" plugin to retrieve IP information of open ports, hostnames, open ports, vulnerabilities. If you have a Shodan membership, you can use the "Shodan" plugin with your API key to get access to advanced capabilities that you're subscribed to.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot requires a Shodan membership and an API Key. You'll need to take the following steps before using the plugin.

1. Get your Shodan API key. If you don't have one yet, follow these steps:

    a. Go to your Shodan portal and sign in.

    b. Select **Account**, and on the **Account Overview** tab, next to **API Key**, select **Show**.

    c. Copy your API key.

2. Sign in to Microsoft Security Copilot.

3. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.
4. Next to **Shodan**, select the toggle to enable it.
5. In the Shodan Plugin settings pane, in the **Value** field, paste your API key.
6. Save your changes.

Sample Shodan prompts

After the Shodan plugin is configured, you can use it by typing `Shodan` in your Security Copilot prompt bar.

The following table lists capabilities and example prompts to try:

Expand table

Capability	Example prompts

`Check Shodan InternetDB IP address`
(uses Shodan Internet DB Plugin)
 	
Ask Shodan InternetDB about an IP address
 	
Required:
- `ip`
	`Use the Shodan InternetDB database to provide info on "118.25.6.39"`
	
	`Check IP address "118.25.6.39" using Shodan InternetDB database`

`CheckShodanHostIP`
(requires a Shodan membership)
 	
Accepts an IP address (v4 or v6) and provides information about the queried IP, including related country/region, last updated dates, hostnames, and ISP.
 	
Required:
- `ip`
 	
Optional:
- `history`
- `minify`
	`Check IP Address 8.8.8.8 using Shodan`
	
	`Use Shodan to check IP address 8.8.8.8`
	
	`What does Shodan say about IP address 8.8.8.8?`

`GetShodanHostCount`
 	
Behaves like `GetShodanHostSearch`, except that it doesn't return any host results; instead, it returns the total number of results that matched the query and any facet information that was requested. This method doesn't consume query credits.
	`What does Shodan know about the host count for port:22?`
	
	`Use Shodan to look up the host count for port:22`

Capability	Example prompts
Required: - `query` Optional - `facets`	

`GetShodanHostSearch`
(requires a Shodan membership ↗)
 	
Searches Shodan using the same query syntax as the website and uses facets to get summary information for different properties. This method might use API query credits depending on usage. If any of the following criteria are met, your account is deducted one query credit:
- The search query contains a filter.
- Accessing results past the first page using the *page*. For every 100 results past the first page, one query credit is deducted.
 	
Required:
- `query`
 	
Optional:
- `facets`
	`Search for hosts running port:22 using Shodan.`
	
	`Use Shodan to look up the hosts running port:22.`

`GetShodanHostSearchFacets`
 	
Returns a list of facets that can be used to get a breakdown of the top values for a property.
	`List all search facets from Shodan records.`
	
	`What are all the Shodan search facets?`

`GetShodanHostSearchFilters`
 	
Returns a list of search filters that can be used in the search query.
	`List all filters that can be used when searching Shodan records.`
	
	`What are the Shodan search filters?`

`GetShodanHostSearchTokens`
 	
Enables you to determine which filters are being used by the query string and what parameters were provided to the filters.
 	
Required:
- `query`
	`Use Shodan to break down Raspbian port:22 into tokens.`
	
	`Get the Shodan host search tokens for Raspbian port:22.`

`GetShodanPorts`
 	
Returns a list of port numbers that the crawlers are looking for.
	`List all ports that Shodan is crawling on the Internet.`
	
	`Get all Shodan ports.`

Capability	Example prompts

`GetShodanProtocols`
 	
Returns an object containing all the protocols that can be used when launching an Internet scan.
	`List all protocols that can be used when performing on-demand Internet scans via Shodan.`
	
	`What protocols can be used with Shodan?`

`GetShodanScans`
(requires a Shodan membership ↗)
 	
Returns a list of all the on-demand scans that are currently active on the account.
	`Get list of all the created scans via Shodan.`
	
	`What are all the scans created by Shodan?`

`GetShodanScansID`
(requires a Shodan membership ↗)
 	
Checks the progress of a previously submitted scan request. Possible values for the status are:
- `SUBMITTING`
- `QUEUE`
- `PROCESSING`
- `DONE`
 	
Required:
- `id`
	`Get the status of the scan request DQdcm6QYgENbGj0R using Shodan.`
	
	`What does Shodan say about the scan request DQdcm6QYgENbGj0R?`

`GetShodanAlertIDInfo`
(requires a Shodan membership ↗)
 	
Returns information about a specific network alert.
 	
Required:
- `id`
	`Get the details for the network alert 0DC55K0N2HHZS3D1 using Shodan.`
	
	`What does Shodan say about the network alert 0DC55K0N2HHZS3D1?`

`GetShodanAlertsInfo`
(requires a Shodan membership ↗)
 	
Returns a list of all the network alerts that are currently active on the account.
	`Get a list created alerts using Shodan.`
	
	`What are all the created alerts in Shodan?`

`GetShodanAlertTriggers`
 	
Returns a list of all the triggers that can be enabled on network alerts.
	`Get a list of available triggers using Shodan.`
	
	`What are all the available triggers in Shodan?`

Capability	Example prompts

`GetShodanNotifiers`
 	
Returns a list of all the notifiers that the user has created.
	`Get a list of my notifiers using Shodan.`
	
	`What are all the notifiers in Shodan?`

`GetShodanNotifierProvider`
 	
Returns a list of all the notification providers that are available and the parameters to submit when creating them.
	`Get a list of notifier providers on Shodan.`
	
	`What are all the notifier providers in Shodan?`

`GetShodanBulkData`
 	
Returns a list of datasets that can be downloaded in Shodan.
	`Get a list of datasets that can be downloaded in Shodan.`
	
	`What are all the datasets that can be downloaded in Shodan?`

`GetShodanBulkDataDataset`
 	
Required:
- `dataset`
 	
Returns a list of files that users can download from the dataset.
	`Get a list of files that can be downloaded in Shodan raw-daily dataset.`
	
	`What are all the files that can be downloaded from Shodan raw-daily dataset?`

`GetShodanDomain`
 	
Required:
- `domain`
 	
Optional:
- `history`
- `type`
 	
Get all the subdomains and other DNS entries for the given domain.
	`Check domain google.com using Shodan.`
	
	`Use Shodan to check domain google.com`

`GetShodanDomainResolve`
Optional:
- `hostnames`
 	
Look up the IP address for the provided list of hostnames.
	`Check the IP of google.com, facebook.com using Shodan.`
	
	`Use Shodan to check IP of google.com, facebook.com`

`GetShodanDomainReverse` Optional: - `ips`	`Check the hostnames of 74.125.227.230,204.79.197.200 using Shodan.`

Capability	Example prompts
Look up the hostnames that have been defined for the given list of IP addresses.	`Use Shodan to check the domains of 74.125.227.230,204.79.197.200`

`GetShodanHTTPHeaders`
 	
Shows the HTTP headers that your client sends when connecting to a webserver.
	`Check the HTTP headers using Shodan.`
	
	`Use Shodan to check the HTTP headers.`

`GetShodanMyIP` Get your current IP address as seen from the Internet.	`Get my current IP address using Shodan.` `Use Shodan to get my current IP address.`

Troubleshoot the Shodan plugin

Errors occur

If you encounter errors occur, such as: *Couldn't complete your request*, or *An unknown error occurred*, make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Provide feedback

To provide feedback, contact Shodan.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Silverfort

08/25/2025

The Silverfort plugin for Microsoft Security Copilot enables seamless integration of Silverfort's advanced CEF data within the Microsoft Sentinel environment to bring in Identity protection insights. This plugin empowers security teams to enhance their threat detection and response capabilities through intuitive natural language queries and detailed insights.

The Silverfort plugin leverages KQL-based queries to extract and analyze data from Silverfort's security logs within the Microsoft Sentinel workspace. Users can customize their queries using a range of input parameters to retrieve targeted information, enabling more efficient threat investigation and proactive defense measures.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Prerequisites

To use the Silverfort plugin you must first configure Common Event Format (CEF) and Syslog through Azure Monitor Agent (AMA) forwarding using the following guide. This plugin queries data stored in the CommonSecurityLog table in a log Analytics workspace enabled for Microsoft Sentinel. For more information, see Ingest syslog and CEF messages to Microsoft Sentinel with the Azure Monitor Agent.

After setting up a forwarder, proceed to:

- Assign a public IP address
- Allow Syslog inbound traffic on port 514 in network settings

You're now ready to set up your Syslog Server in Silverfort to send the events and have the newly configured AMA forward the information to Microsoft Sentinel.

1. Enter the IP Address of your AMA Forwarder in the Server IP field.
2. Enter port 514 in the Port field.
3. Select TCP Protocol in the Protocol Field.

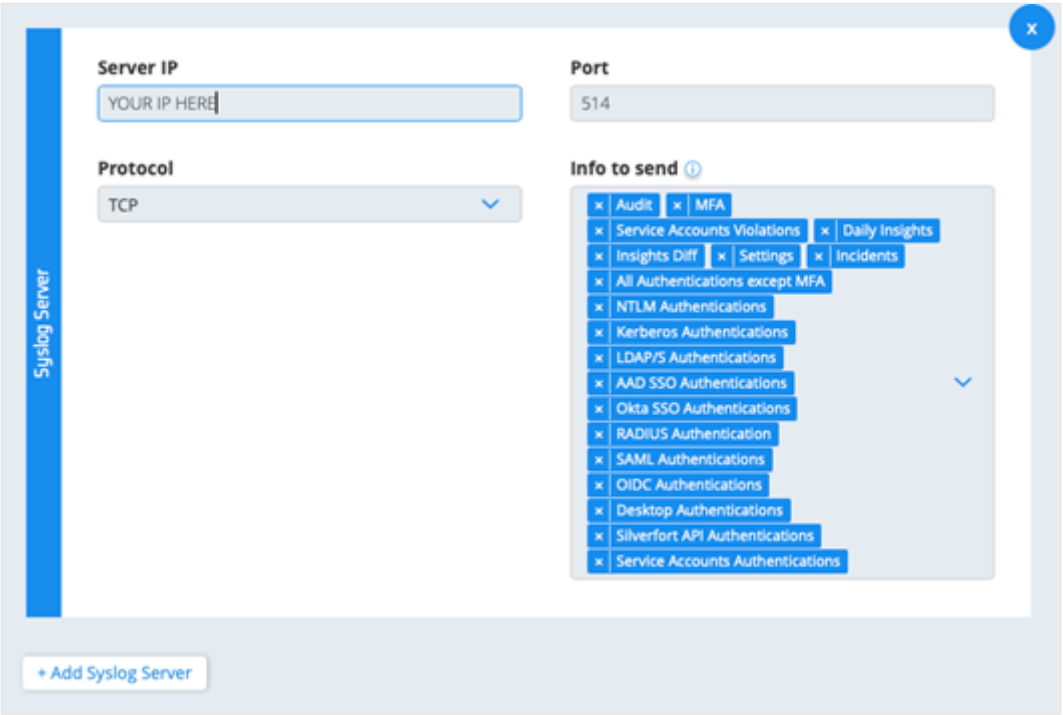

4. Verify that all the information in the fields are present except the Splunk application.
5. Select **Save All**.

Know before you begin

You'll need to take the following steps before using the plugin.

1. Sign in to Microsoft Security Copilot.
2. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.
3. Next to **Silverfort**, select the toggle to enable it.

   Provide the following information:

   - TenantId: The ID of the Microsoft Entra ID Organization that the Microsoft Sentinel workspace is in.
   - WorkspaceName: The name of the Microsoft Sentinel workspace.
   - SubscriptionId: The ID of the Azure Subscription that the Microsoft Sentinel workspace is in.
   - ResourceGroupName: The name of the Resource Group that the Microsoft Sentinel workspace is in.
4. Save your changes.

Sample Silverfort prompts

After the Silverfort plugin is configured, you can use it by taking one of the following steps:

- Access the plugin capabilities by selecting the **Plugin button** from the prompt bar and selecting **Silverfort**.
- Prompt Security Copilot using any of the following example prompts.

The following table lists example prompts to try:

Expand table

Capability	Example prompts

QuerySilverfortInformation
Queries information related to CEF data within Microsoft Sentinel based on time, risk, indicator, source IP, source hostname, and others.
	`Provide a count of Silverfort risk requests that have the Silverfort policy name 'mypolicy' in the last week.`
	
	`How many Silverfort MFA subtype requests have there been in the last week that have an MFA response of 'Blocked'?`
	
	`How many requests in the last week have a Silverfort policy action of 'MFA'?`
	
	`Give me the top 10 Silverfort requests with Criticalrisk where the source username is "john.doe@something.com".`

QuerySilverfortIncidents
Queries for all Silverfort related Incidents within the time period specified
	`Give me all Silverfort incidents in the last month.`
	
	`Give me all Silverfort incidents in the last week with status ongoing.`

Troubleshoot the Silverfort plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred* | Make sure the plugin is turned on. This error may occur if the lookback period is too long, causing the query to attempt to retrieve an excessive amount of data. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Provide feedback

To provide feedback, contact Silverfort.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Splunk

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

This plugin allows Security Copilot users to make calls to the Splunk REST API. Currently the following functionalities are supported:

- Performing normal and one-shot ad-hoc SPL queries.
- Creating, retrieving, and dispatching saved searches in Splunk.
- Retrieving and viewing information about fired alerts from saved searches in Splunk.
- Getting information about currently running search jobs in Splunk.

Prerequisites

- Access to an installation of Splunk
- Ensure to allow Security Copilot's egress IPs to contact your Splunk instance. For more information, see Security Copilot IP address ranges. Follow the steps to allow the below IPs based on the Splunk instance type you're using. For example, for Splunk Cloud, use the guidance here: Splunk Cloud Platform Admin Manual.
- One of the following authentication methods in Splunk
  - Splunk authentication token (preferred)
  - Splunk username and password for basic authentication

Documentation for setting up a Splunk authentication token can be found here. In addition, there are other considerations you may need to take in mind if you're running Splunk Cloud. Those considerations are documented here.

> **Note**
>
> This article contains information about non-Microsoft plugins. This is provided to help complete integration scenarios. However, Microsoft doesn't provide troubleshooting support for non-Microsoft plugins. Contact the vendor for support.

Know before you begin

Integration with Security Copilot works with either an API key or basic authentication. You need to take the following steps before using the plugin.

API Key Authentication

API Key Authentication is the preferred method of authentication. To set up authentication via API Key, you need to have the following pieces of information:

- The URL for accessing the REST API
- The Splunk authentication token for the Splunk user account you're using to access the API. Documentation for setting up a Splunk authentication token can be found here. In addition, there are other considerations you may need to take in mind if you're running Splunk Cloud. Those considerations are documented here.

1. When you're asked to set up authentication, select the **APIKey** option.

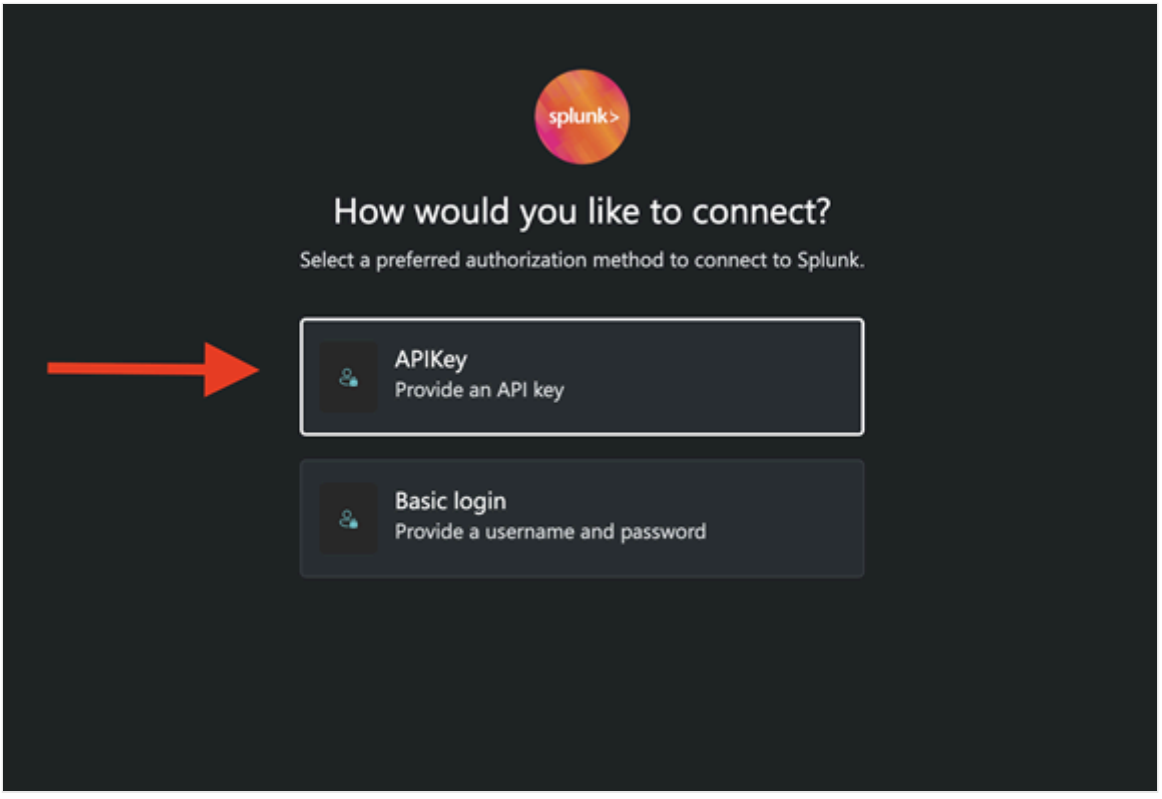

2. Add the Splunk API URL to the field for **Splunk Instance API URL**. Add the Splunk authentication token in the **Value** field.

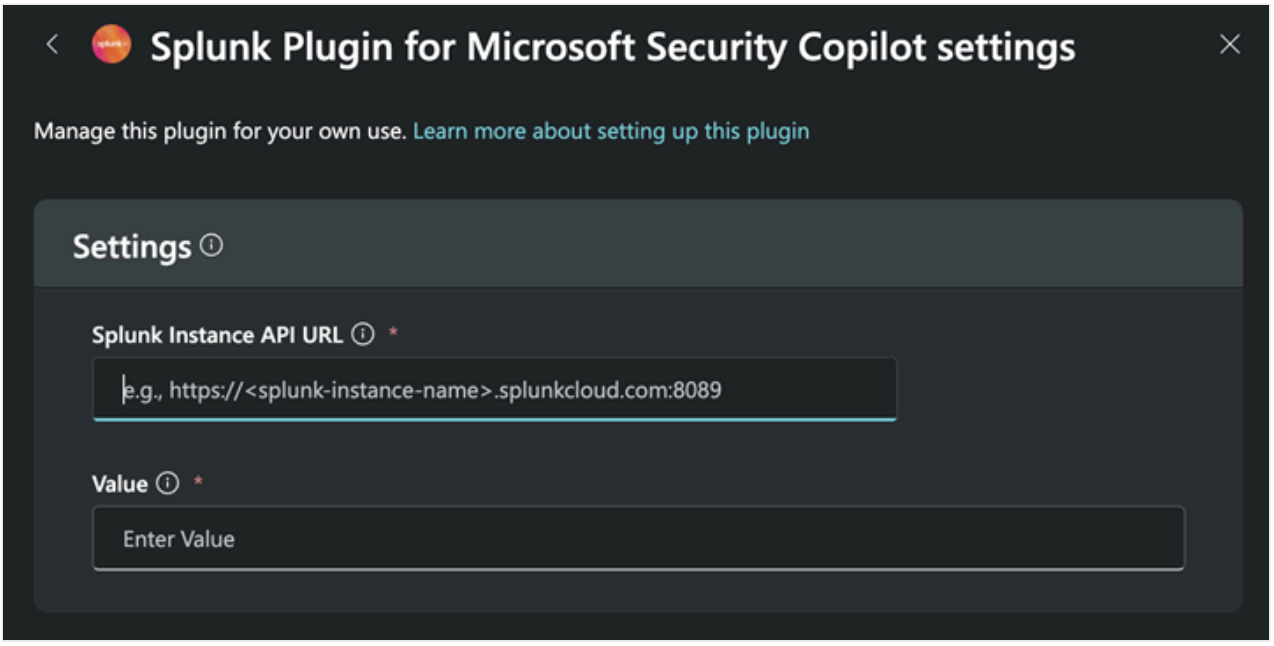

3. Select **Save** to complete setup.

Basic Authentication

To set up authentication using Basic Authentication, you need to have the following pieces of information:

- The URL for accessing the REST API
- The username and password for the Splunk user account you're using to access the API.

1. When you're asked to set up authentication, select the **Basic login** option.

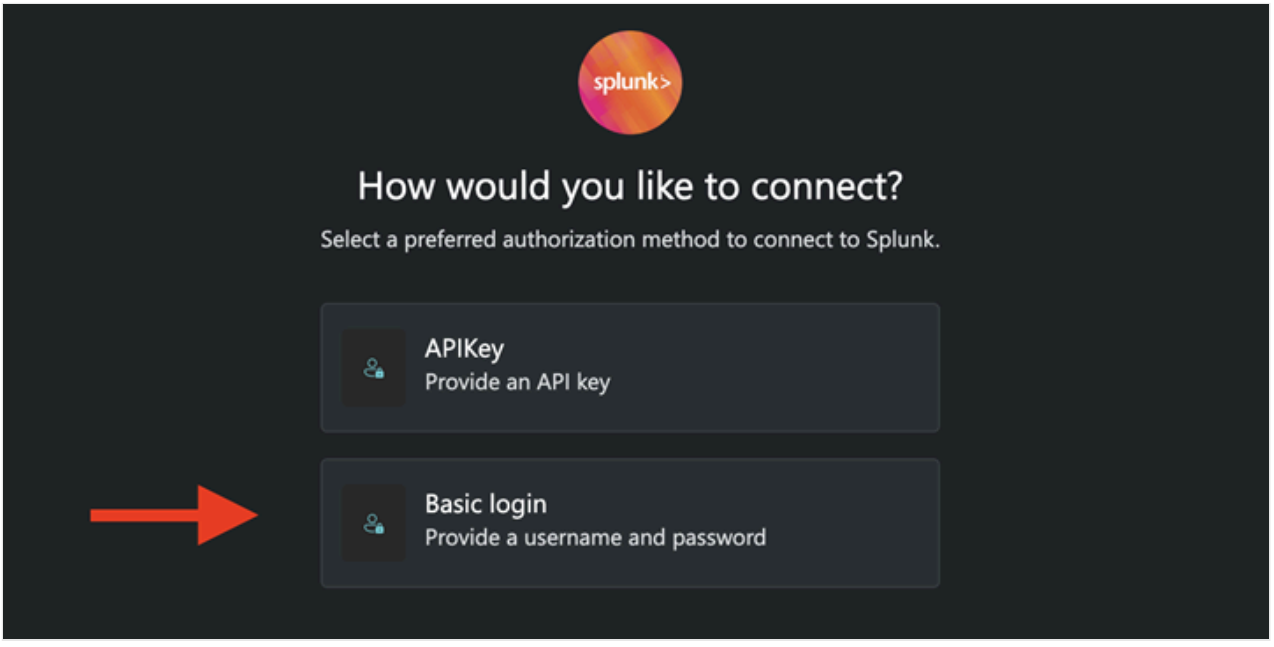

2. Add the Splunk API URL to the field for **Splunk Instance API URL**. Add the Splunk username in the **Username** field. Add the Splunk password in the **Password** field.

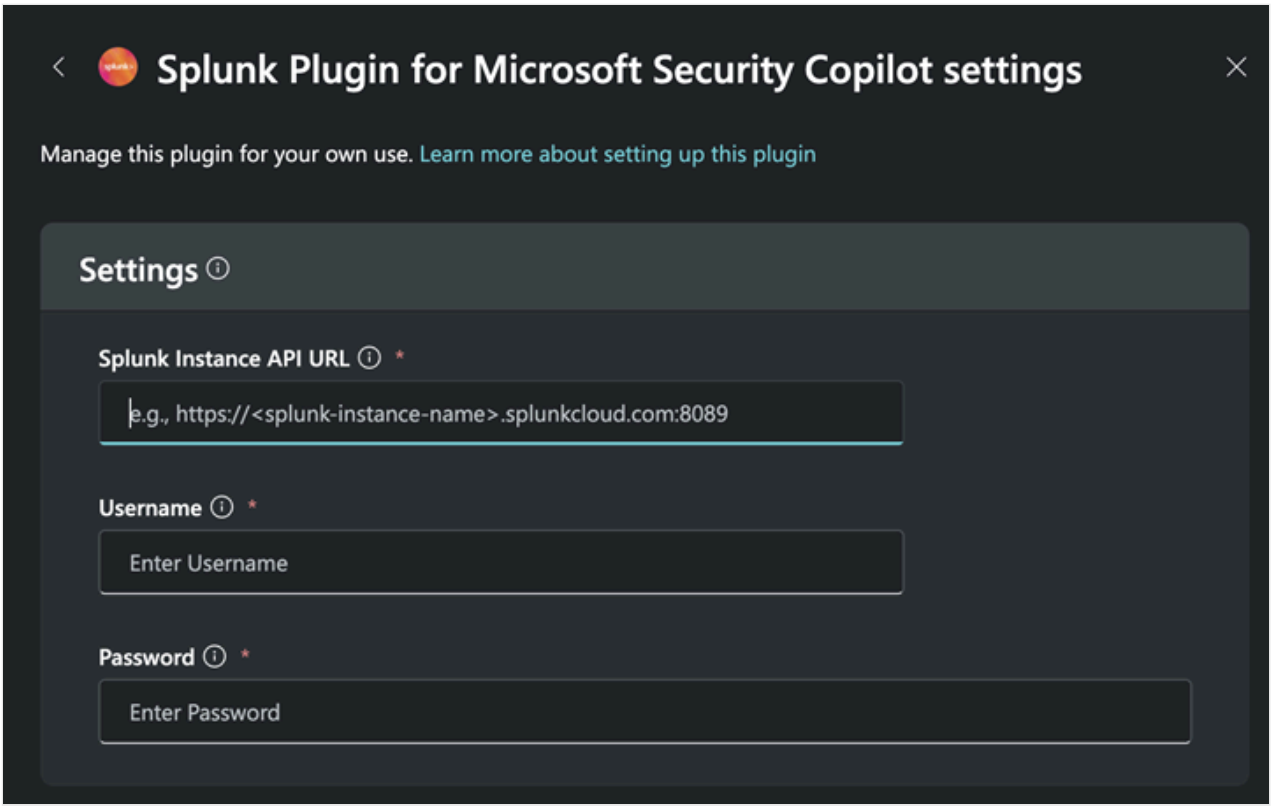

3. Select **Save** to complete setup.

Available Skills

The Splunk Plugin for Microsoft Security Copilot exposes the following skills:

- Ad-hoc searches
  - Creating search jobs
  - Getting information about search jobs
  - Retrieving results from search jobs
  - Running one-shot searches
- Saved searches
  - Retrieving saved searches
  - Creating saved searches
  - Dispatching a saved search
- Fired alerts from saved searches
  - Retrieving fired alerts
  - Retrieving fired alert details

With the Splunk plugin for Microsoft Security Copilot, you can invoke interactions with Splunk in the context of a natural conversation. Here's an example:

1. A user can use the public web to research data on a recently announced vulnerability / CVE.

2. The user can then use a follow-up prompt such as "Save this CVE number as a search in Splunk across all indexes". Security Copilot will maintain the context from the previous prompt in the most recent prompt.
3. The user can then modify the saved search within Splunk to incorporate more advanced SPL techniques or to create visualizations.

Sample Splunk prompts

Expand table

Skill	Prompt
Create a search job	`Run the following search in Splunk in normal mode: index=notable "System Network Configuration Discovery". Ensure to begin the query with the word "search".`
Get the search job results	`Get the search job results for SID 1740764708.5591 from Splunk`
Run a oneshot search	`Run the following search in Splunk in oneshot mode: index=notable "System Network Configuration Discovery"`
Create a saved search	`Save the following search in Splunk: index=notable "System Network Configuration Discovery". Name the search "Network Config Discovery report".`
Retrieve saved searches	`Get all of the saved searches for the copilot user from Splunk`
Dispatch a saved search	`Dispatch the saved search "Top Mitre Techniques" in Splunk`
Retrieve fired alerts	`Get the list of fired alerts from Splunk`
Retrieve fired alert details	`Tell me about the fired alert Apache_HTTP_StatusCode_Alert_Test`
Get current search jobs	`Get all of the current search jobs in Splunk`

Microsoft Security Copilot will often understand and get context from answers that are returned. As a result, you can use natural conversation in a chain of prompts. For example: if you use a prompt such as `Dispatch the saved search "Cloud Alerts" in Splunk`, the search job ID is returned. Security Copilot will have that search job ID in its current context, and you can follow up with `Get the search job results` rather than having to manually specify a search job ID.

Troubleshoot the Splunk plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*. Make sure the plugin is turned on. This error may occur if the lookback period is too long, causing the query to attempt to retrieve an excessive amount of data. If the issue persists, sign out of Security Copilot, and then sign back in. Also, make sure that the authentication mechanism has the appropriate permissions within Splunk (ensure that the Splunk user you're authenticating as with bearer authentication, has permissions to invoke API calls). Finally, if you're connecting to Splunk enterprise, make sure that the SSL you're using for the REST API endpoint isn't using a self-signed certificate.

If you're creating a search job and the error is an HTTP status 400 with a status code of BadRequest, there are two potential causes:

- The keyword `search` is missing from the SPL query being sent to the Splunk API. This can be addressed by adding `Make sure the SPL query starts with the word "search".` to the prompt.
- The execution mode isn't specified or is incorrectly specified. This can be addressed by adding one of the following statements (or some flavor of it) to your prompt:
  - `Run this search in normal mode.`
  - `Run this search in oneshot mode.`

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Provide feedback

To provide feedback, contact the Splunk partner engineering team.

See also

Other plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Last updated on 02/19/2026

Tanium

Tanium provides a converged endpoint management (XEM) reference platform to manage complex security and technology environments. Tanium protects endpoints from cyber threats by integrating workflows across IT, Risk, Compliance, and Security into a single platform. Tanium delivers comprehensive visibility across devices, a unified set of controls, real-time remediation, and a common taxonomy to protect critical information and infrastructure at scale.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Microsoft Security Copilot requires a Tanium instance URL and API token. You'll need to take the following steps before using the plugin.

1. Sign in to Tanium Console to retrieve information you need to configure the Tanium plugin.
    a. Select Modules > Connect > Overview. The Connect overview page appears.
    b. Select Settings and navigate to **Microsoft Copilot for Security**.
    c. Select **Tanium Instance URL Copy** to copy the Tanium instance URL to the clipboard. Paste the URL in a text editor, such as Notepad.
    d. Select **Generate** to generate an API token, and copy the token value to the clipboard. Paste the token value in a text editor.
2. Sign in to Microsoft Security Copilot.
    a. Select **Sources** from the prompt bar. The **Manage sources** modal appears.
    b. Go to the **Tanium** plugin, and then select **Set up**.
    c. On the Tanium settings, copy the Tanium Instance URL and paste it in **Tanium Instance URL**.
    d. Copy the token value and paste it in **Value**.

e. Select **Save**.

Sample Tanium prompts

After the Tanium plugin is configured, you can use it to retrieve information about endpoints (devices) in your organization. The following table lists some capabilities and example prompts you can try:

Expand table

Capability	Example prompts
Get Logged In User Retrieves the user that is currently logged into an endpoint Requires Tanium Core Platform	`Using Tanium, return the user currently logged into the endpoint with the hostname hostname so that I can investigate possible unauthorized endpoint use. Return a Tanium Console Question Results URL so that I can view more real-time information for this endpoint.`
Get Real-time Data from Endpoints Retrieves real-time data from endpoints based on a Tanium sensor. For more information on supported sensors Requires Tanium Core Platform, sensor-dependent	`Using Tanium, return the computer name and IP address of endpoints. Display the results in a table, alphabetically sorted by computer name, and return a Tanium Console Question Results URL so that I can view the real-time list of endpoints.`
Count Endpoints Having Package Version Retrieves the total count of endpoints that have the given software package Requires Asset, SBOM	`Using Tanium, return the total number of endpoints with a software package for software-name, so that I can start cataloging which computers have the software installed. Display the results in a table, alphabetically sorted by host name, and return a Tanium Console Question Results URL so that I can view the real-time list of endpoints.`
List Endpoints Having Package Retrieves up to 10 endpoints that have the given software package Requires Asset, SBOM	`Using Tanium, return the endpoints with a software package for software-name so that I can start cataloguing which computers might have an out-of-date version. Display the results in a table, alphabetically sorted by host name, and return a Tanium Console Question Results URL so that I can view the real-time list of endpoints.`
List Process SHA-256 Hashes and Versions Retrieves the SHA-256 file hash and version for a given process	`Using Tanium, return the SHA-256 hash value and process version for the running process process-name, so that I can find other instances of this process based on the hash value.`

Capability	Example prompts
Requires Asset, SBOM, Threat Response	
Get Vulnerability Test Results Returns whether an endpoint is vulnerable to a given CVE, and the reason why it is vulnerable Requires Tanium Comply	`Using Tanium, examine whether endpoint <hostname> is vulnerable to <cve-id>, and return the reasons that this endpoint is vulnerable, along with a suggested plan of action to remediate the intrusion.`
List Endpoints Vulnerable To CVE Retrieves up to 10 endpoints vulnerable to a given CVE ID Requires Tanium Comply	`Using Tanium, return the endpoints vulnerable to cve-id, so that I can remediate the vulnerability on these endpoints. Display the results in a table, sorted alphabetically by host name, and return a Tanium Console Question Results URL so that I can view the real-time list of endpoints.`
View Endpoint Processes Retrieves a URL to the Threat Response Live Connection page for the requested endpoint, which contains a list of running processes Requires Direct Connect, Threat Response	`Using Tanium, return a Threat Response Live Connection URL for the endpoint with the hostname hostname, so that I can review the running processes and identify potential vulnerabilities.`
List Service Module Details Retrieves running service module information for an endpoint, including name, caption, and image path Requires Incident Response	`Using Tanium, return information for the service modules running on the endpoint with the hostname hostname, so that I can review the list for unexpected service modules. Display the results in a table, alphabetically sorted by service module name, and return a Tanium Console Question Results URL so that I can view the real-time list of service modules.`
List Service Process Details Retrieves running service process information for an endpoint, including name, process ID, and file path Requires Incident Response	`Using Tanium, return information for the service processes running on the endpoint with the hostname hostname, so that I can review the list for unexpected service processes. Display the results in a table, alphabetically sorted by service process name, and return a Tanium Console Question Results URL so that I can view the real-time list of service processes.`
List WMI Event Consumers Retrieves Windows Management Instrumentation (WMI) event consumers running on an endpoint	`Using Tanium, return the WMI event consumers running on the endpoint with the hostname hostname so that I can ensure only expected event consumers are running, and return a Tanium Console Question Results URL so that I can view the real-time list of event consumers.`

Capability	Example prompts
Requires Incident Response	

List File Details
Retrieves details for a file by name, including the endpoints on which it is installed, the file path, and file size
 	
Requires Index
	`Using Tanium, return information for the file named file-name so that I can determine if it is running on unintended endpoints. Display the results in a table, alphabetically sorted, and return a Tanium Console Question Results URL so that I can view the real-time list.`
	
	or
	
	`Using Tanium, return information for the file named file-name installed on the endpoint with the hostname hostname, so that I can determine if it is running on unintended endpoints. Display the results in a table, alphabetically sorted, and return a Tanium Console Question Results URL so that I can view real-time information.`

List Child Processes for Process File
Returns all child processes running on an endpoint based on a given process file name
 	
Requires Threat Response
	`Using Tanium, list the child processes of process-name so that I can analyze resource usage. Display the results in a table, alphabetically sorted by process name, and return a Tanium Console Question Results URL so that I can view the real-time list of endpoints.`
	
	or
	
	`Using Tanium, list the child processes of process-name that are running on the computer with the hostname hostname, so that I can analyze resource usage. Display the results in a table, alphabetically sorted by process name, and return a Tanium Console Question Results URL so that I can view the real-time list of endpoints.`

List Endpoints with Process Command Retrieves up to 10 endpoints running the given command line command Requires Threat Response	`Using Tanium, return the endpoints running the command line command process-command, so that I can ensure this process is not running on unexpected endpoints. Display the results in a table, sorted alphabetically by host name, and return a Tanium Console Question Results URL so that I can view the real-time list of endpoints.`
List Endpoints with Process Name Retrieves up to 10 endpoints running the given process Requires Threat Response	`Using Tanium, return the endpoints running a process called process-name, so that I can ensure this process is not running on unexpected endpoints. Display the results in a table, sorted alphabetically by host name, and return a Tanium Console Question Results URL so that I can view the real-time list of endpoints.`

Capability	Example prompts
List Endpoints with Process MD5 Hash Retrieves up to 10 endpoints running the given process matching the provided MD5 hash value Requires Threat Response	`Using Tanium, return all endpoints that are running a process with the MD5 hash value md5-hash-value, so that I can ensure this process is not running under a different file name. Display the results in a table, sorted alphabetically by host name, and return a Tanium Console Question Results URL so that I can view the real-time list of endpoints.`
List File Operations Retrieves historical file operation information from endpoints, including endpoint name, file path, and the file operation type, such as create or delete Requires Threat Response	`Using Tanium, return file operation information for the endpoint named hostname running on the file path "_partial-file-path" over the past time-frame so that I can determine if any malicious file behavior is occuring on the endpoint. Display the results in a table, alphabetically sorted, and return a Tanium Console Question Results URL so that I can view the real-time list.`  or  `Using Tanium, return file operation information for files running on the file path "_partial-file-path" over the past time-frame so that I can determine if there is any malicious file creation or deletion. Display the results in a table, alphabetically sorted, and return a Tanium Console Question Results URL so that I can view the real-time list.`
List Processes Connected To IPv4 Address Retrieves the processes running on an endpoint with the given IPv4 address Requires Threat Response	`Using Tanium, return the processes running on the endpoint with the IPv4 address ipv4-address, so that I can analyze any potential security intrusions and resource usage. Display the results in a table, sorted alphabetically by process name, and return a Tanium Console Question Results URL so that I can view the real-time list of endpoints.`
List Process Ran As User Retrieve the processes running on an endpoint as a given user Requires Threat Response	`Using Tanium, return the processes running as the user user-name, so that I can determine whether there are issues with unauthorized access. Display the results in a table, sorted alphabetically by computer name, and return a Tanium Console Question Results URL so that I can view the real-time list of endpoints.`  or  `Using Tanium, return the processes running as the user user-name on the endpoint with the hostname hostname, so that I can determine whether there are issues with unauthorized access. Display the results in a table, sorted alphabetically by process`

Capability	Example prompts
	`name, and return a Tanium Console Question Results URL so that I can view the real-time list of endpoints.`

Troubleshoot the Tanium plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*, make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts are not invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Provide feedback

To provide feedback, contact Tanium.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Last updated on 02/19/2026

TufinMate

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

Tufin Orchestration Suite is a centralized firewall management platform that allows you to define, implement, and enforce comprehensive security policies across complex hybrid networks. Tufin makes it easier to identify and modify risky firewall configurations, optimize firewall usage, evaluate the impact of vulnerabilities on the network, and achieve continuous compliance with industry regulations.

TufinMate for Security Operation Centers is an extension of the Tufin Orchestration Suite platform, accessed from Security Copilot. It empowers security analysts to reduce incident response times by providing them with context on permitted network access paths, firewall rule compliance, permissiveness, and last hit information.

Use the TufinMate plugin with Microsoft Security Copilot to answer questions about how traffic is permitted to traverse your network. See where access is permitted, identify permissive or unused rules, verify the compliance of an access flow, and receive instructions on how to initiate a change request.

The TufinMate plugin helps your team:

- Assess the priority of a vulnerability based on access flows to or from the relevant asset.
- Better understand the network impact radius of a security incident by verifying whether access to sensitive assets is permitted or blocked by firewalls.
- Identify overly permissive or unused firewall rules that may be relevant to a security incident.
- Verify the compliance of a firewall device or access path.
- Gather the information needed to submit a network change request.

Prerequisites

The TufinMate Management App must be installed on your Tufin Orchestration Suite cluster. For information on how to install the application, see the TufinMate User Guide.

Know before you begin

Integration with Security Copilot requires an API Key. Follow these instructions generate a key during setup.

1. Sign in to Microsoft Security Copilot.
2. On the prompt bar, select the **Sources** button to access **Manage plugins**.
3. Next to **TufinMate**, select the toggle to enable it.
4. The plugin asks you for a value (an access identifier for API authentication). You need to add a unique API Key generated by the TufinMate Management App. To retrieve it:
   a. Sign into Tufin Orchestration Suite and open the TufinMate Management Application.
   b. From the Available Integrations screen, within the card for Security Copilot, select **Manage Integration**.
   c. Under Integrate with Security Copilot, select **Generate API Key**.

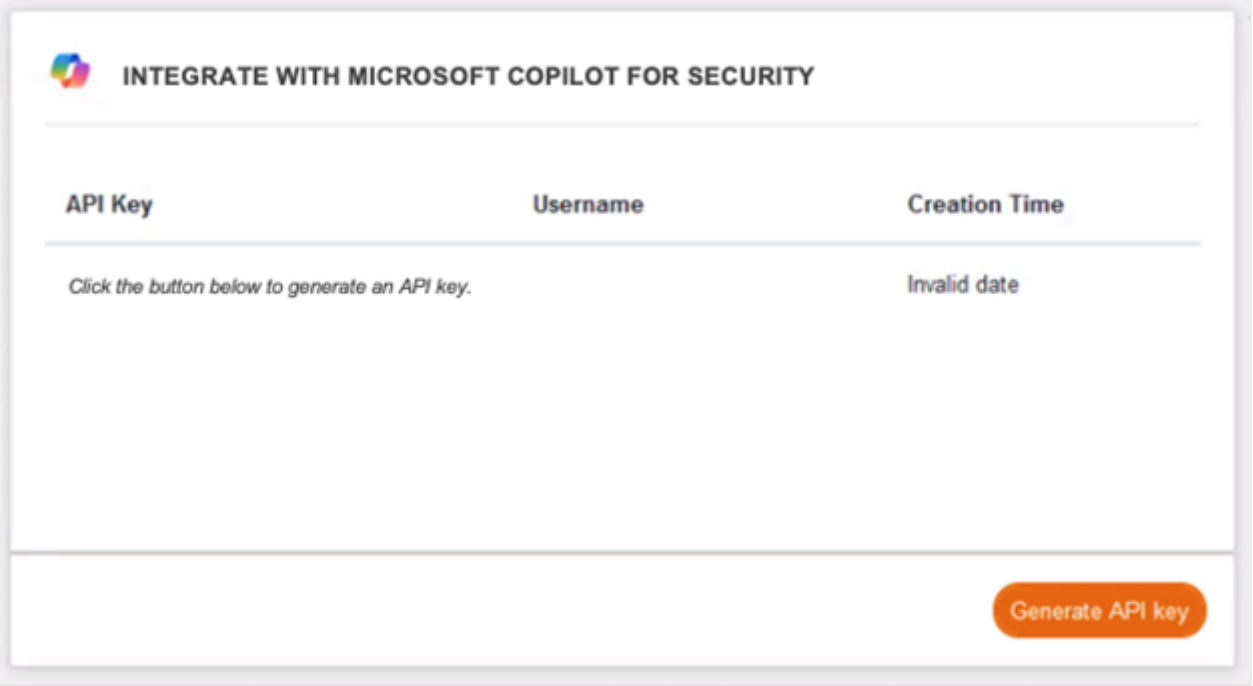

5. Copy the API Key that was generated. Note, you can't generate the same API key twice.
6. Navigate back to Security Copilot and add the API key.

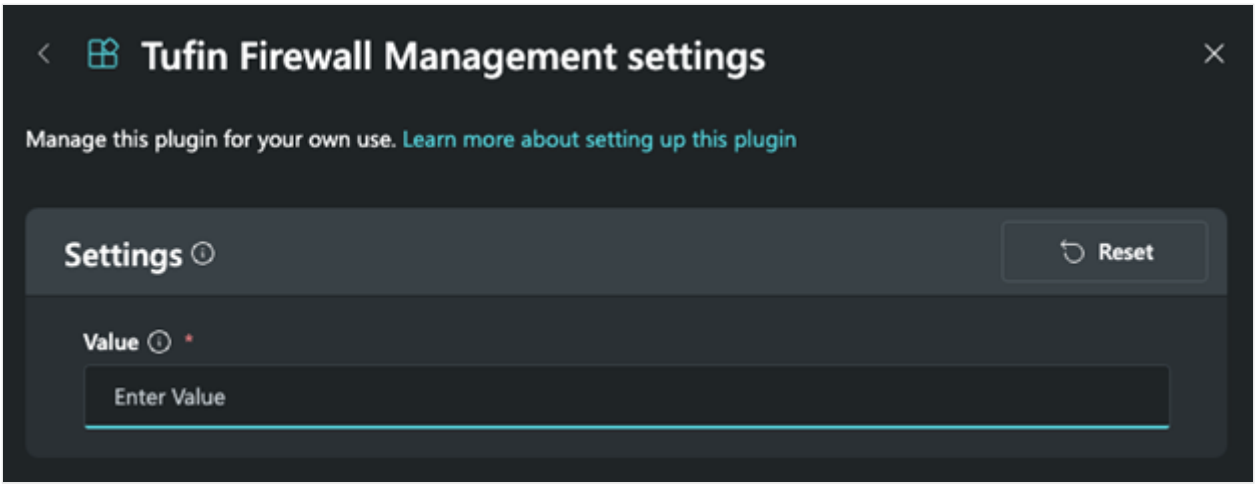

7. Select **Save**.

8. Use the sample prompts listed in the following section to get started using the TufinMate plugin.

Sample TufinMate prompts

After the TufinMate plugin is configured, you can use it by typing TufinMate in your Security Copilot prompt bar, followed by an action. The following table lists example prompts to try to get started using the TufinMate plugin.

Expand table

Capability	Example prompts

Network Access Status	`Can {{source}} access {{destination}} using any service?`
	
	`Tell me if traffic is permitted between {{source}}.`

Device Details	`Can you provide some details about device id {{device ID}} using the Tufin Firewall Management plugin?`

Last Hit Information	`Does the device {{insert device ID or device name or device IP}} have any rules with no hits in the past 90 days?`
	
	`Are there any rules on the path between <<source>> and {{destination}} with no hits in the past 30 days?`

Firewall Rule Permissiveness	`Does the device {{insert device ID or device name or device IP}} have any highly permissive rules?`
	
	`Are there any rules on the path between <<source>> and {{destination}} that are highly permissive?`

Capability	Example prompts

Compliance Violations	`Does the device {{insert device ID or device name or device IP}} have any rules with unified security policy violations?`
	
	`Are there any rules on the path between <<source>> and {{destination}} with compliance violations?`

Change Request Guidance	`How could I use Tufin to request traffic be permitted?`
Ticket Type Question	`Which ticket type should I use to request {{Insert one of the following: Access, Rule Decommission, Block Access, Recertification, Rule Modification}}?`

Troubleshoot the TufinMate plugin

Contact Tufin Support

To submit a support request or manage your support ticket, visit the Tufin User Portal.

Generate logs for TufinMate plugin

From the TufinMate Management App, click on the Global Settings icon on the left sidebar menu. Use the log information setting to modify the level of logs generated. Use the provided prompt to retrieve the logs.

Related content

Use the following Security Copilot articles for guidance:

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Last updated on 05/29/2026

UrlScan

08/25/2025

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

UrlScan.io is a free online service and tool that allows users to scan and analyze URLs (Uniform Resource Locators) or website links to determine potential security threats and risks associated with those URLs. It helps users assess the safety and trustworthiness of a website or a specific web page.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot works with an API key. You'll need to take the following steps before using the plugin.

1. Get your URL Scan API key. If you don't have one yet, follow these steps:

   a. Go to the urlscan.io website and create your account. Then sign in.

   b. Go to your profile page, and then select **+ Create API key**.

   c. Return to your profile page, and copy your API Key.

2. Sign in to Microsoft Security Copilot.

3. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.

4. Next to **UrlScan**, select **Set up**.

5. In the **Value** field, paste your API Key, and then select **Save**.

Sample UrlScan prompts

After the UrlScan plugin is configured, you can use it by taking one of the following steps:

- Access the capability directly by typing `UrlScan` in the prompt bar; or
- Prompt Security Copilot to use the UrlScan API on a website

The following table summarizes how this capability works.

Expand table

Capability	Example prompts	What it does
`UrlScan` Required Input: URL	- `Can you scan https://www.microsoft.com/en-us/ for threats.` - `Use Url scan to check if https://www.microsoft.com/en-us/ is malicious`	Scans a URL for potential security threats and risks.

Troubleshoot the UrlScan plugin

Errors occur

If you counter errors, such as *Couldn't complete your request*, or *An unknown error occurred*, make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts are not invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Provide feedback

To provide feedback, contact UrlScan.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Valence Security

08/25/2025

> **Important**
>
> Some information in this article relates to a prereleased product which may be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

Valence combines SaaS security posture management (SSPM) and advanced remediation with business user collaboration to find and fix Software as a Service (SaaS) security risks. The Valence SaaS Security Platform provides visibility and remediation capabilities for business-critical SaaS applications. You can use the Valence Security plugin with Microsoft Security Copilot to:

- Investigate user activity across multiple SaaS platforms; and
- Create reports to understand a specific user's SaaS security posture

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Know before you begin

Integration with Security Copilot requires a Valence Security subscription and an API Key. You'll need to take the following steps before using the plugin.

1. Get your API key. If you don't have one, follow these steps:

   a. Go to the Valence Security web site and sign in.

   b. Go to API Keys Settings and generate your API key.

   c. Copy your API key.

2. Sign in to Microsoft Security Copilot.

3. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.

4. Next to **Valence Security**, select **Set up**.

5. In the **Value** field, paste your API Key, and then select **Save**.

Sample Valence Security prompts

After the Valence Security plugin is configured, you can use it to investigate or report on user activities across SaaS platforms. The following table provides some example prompts you can try in your Security Copilot prompt bar:

Expand table

Capability	Example prompts

Investigate user activities across SaaS platforms	`"Which high privileged actions did <USER> take over the past day?"`
	
	`"Show me recent alerts for <USER>"`

Create a report about user activities in SaaS platforms	`"Summarize <USER> security SaaS posture"`

Troubleshoot the Valence Security plugin

Errors occur

If you encounter errors, such as *Couldn't complete your request*, or *An unknown error occurred*, make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts are not invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use.

Provide feedback

To provide feedback, contact Valence Security.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Whoisfreaks

08/25/2025

Whoisfreaks helps elevate your cyber-security strategy with domain and IP intelligence services. Designed for analysts, researchers, and brand owners. Our platform provides unparalleled insights and monitoring capabilities to protect your digital assets. Stay ahead of threats, ensure brand integrity, and make informed decisions with real-time data you can trust. Our platform seamlessly integrates with existing systems, enhancing workflow efficiency and effectiveness.

> **Note**
>
> This article contains information about third-party plugins. This is provided to help complete integration scenarios. However, Microsoft does not provide troubleshooting support for third-party plugins. Contact the third-party vendor for support.

Prerequisites

- The plugin requires the use of an API key that can be obtained from the Whoisfreaks portal
- Once logged into the Whoisfreaks portal, obtain and safely copy your API key.

Know before you begin

You'll need to take the following steps before using the plugin.

1. Sign in to Microsoft Security Copilot.
2. Access **Manage Plugins** by selecting the **Plugin** button from the prompt bar.
3. Select **Settings** to configure your **Whoisfreaks** plugin.
4. Select **Save** and begin prompting in Security Copilot.

Sample Whoisfreaks prompts

After the Whoisfreaks plugin is configured, you can try one or more of the capabilities listed in the following table:

Expand table

Capability	Example prompts

Get Whoisfreaks Live, historical, and reverse WHOIS data	`Give me ownership details of ipgeolocation.io using whoisfreaks.`
	
	`Which domains are acquired by <Name>?`

Get Whoisfreaks Live, historical and reverse DNS information	`give me historical DNS data for <domain name>`
Get Whoisfreaks Live SSL Certificates Info for a specific domain	`When does whoisfreaks.com' SSL certificate expire?`
Get Whoisfreaks IP WHOIS info	`Give me WHOIS details of 8.8.8.8`
Get Whoisfreaks ASN WHOIS info	`'AS13335' AS number belongs to which organization?`
Get Whoisfreaks Domain Availability	`Is 'whoisfreaks.pk' a registered domain or not?`

Troubleshoot the Whoisfreaks plugin

Errors occur

If you get errors, such as *Couldn't complete your request*, or *An unknown error occurred*, make sure the plugin is turned on. If the issue persists, sign out of Security Copilot, and then sign back in.

Prompts aren't invoking the correct capabilities

If prompts aren't invoking the correct capabilities, or prompts are invoking some other capability set, you might have custom plugins or other plugins that have similar functionality as the capability set you want to use. For example, if you're using multiple sources that provide information like Whoisfreaks, there might be conflicts.

Provide feedback

To provide feedback, contact Whoisfreaks.

See also

Non-Microsoft plugins for Microsoft Security Copilot

Manage plugins in Microsoft Security Copilot

Preview plugin terms for Microsoft Security Copilot

08/25/2025

Preview terms

This service is in "Preview," and is subject to the Preview terms set in the Product Terms. Microsoft makes no commitments about the functionality, accuracy, or quality of any Security Copilot Plugins (defined in the following section) during the Preview period. We don't recommend using Preview Plugins for production purposes, as data may be lost, deleted, or corrupted.

Microsoft Security Copilot Plugin Terms of Service

These Microsoft Terms of Service ("Terms") explain your and Microsoft's rights and obligations with respect to transporting your data to and from a third-party product ("Third Party Product") and importing that data into Security Copilot ("Microsoft Service"), via the Microsoft Security Copilot plugin ("Security Copilot Plugin"). In addition to these Terms, your use of the Microsoft developed third party Plugin and your access to (and export or import of) data from/to a Third-Party Product may be subject to other, third-party terms and privacy policies. Microsoft isn't a party to your agreement with a Third-Party Product, nor your agreement with any other third party. Third-party applications and services might involve storing, transmitting, and processing your organization's Customer Data on third-party systems that are outside of the Microsoft Service and therefore aren't covered by Microsoft compliance and data protection commitments.

Note that you, as the data controller for your data, are responsible for all data ingested into Microsoft Services, including obtaining applicable consents from individuals whose data is ingested. Only after such data is successfully ingested into the Microsoft Services will such data be considered "Customer Data," for purposes of the Product Terms, Privacy Statement, and Microsoft's data processing and other obligations. You can also find more about Microsoft's commitments to data protection and privacy by visiting our Trust Center.

You're responsible for configuring, and using a Plugin to access and extract any data from a Third-Party Product, and doing so in compliance with any applicable Third-Party Product terms and privacy policies, as well as with applicable law. Microsoft isn't responsible for, and disclaims all liability related to: (i) your use of a third party Plugin, (ii) your interactions with any Third Party Product (including related APIs or other platforms), and (iv) your access and use of any data outside of the Microsoft Services, including any errors, data loss, or other losses that occur

during or as a result of exporting any data into Microsoft Services. Please note that in the event that a Plugin is deleted or disabled, you may lose data in the Plugin. Microsoft isn't responsible for such data loss. If a Plugin becomes disabled or unavailable for any reason you're responsible for manually processing data subject rights.

Connectors overview in Microsoft Security Copilot

The connectors in Security Copilot are a wrapper around the API that allows the developers and users to call out to the Microsoft Security Copilot platform to perform specialized tasks.

For more information, see Microsoft Security Copilot connector.

Creating connectors

The following connectors are currently supported:

- Logic Apps connector
- Copilot Studio connector

Last updated on 12/03/2025

Microsoft Security Copilot - Copilot Studio connector

The Copilot Studio connector enables you to access Security Copilot. This topic provides an introduction to the connector action and its properties to get started integrating Security Copilot in your automation workflows.

The connector exposes the following connector actions:

- Submit a Security Copilot prompt - Submit a natural language prompt to create a new Security Copilot investigation. After completion, the evaluation result will then be returned to your workflow.
- Fetch a Security Copilot prompt status - Submit a natural language prompt to pull the status of a Security Copilot evaluation. After completion, the evaluation result will then be returned to your workflow.

> **Tip**
>
> This connector is also available in all the Microsoft Power Platform experiences such as Microsoft Power Automate and Power Apps.

For more information, see Power Platform Connector Page.

Prerequisites

- Tenant - Ensure that your tenant administrator sets up access to Microsoft Security Copilot before using the connector.
- Ensure the following authentication type is met:
  - **User Authentication** - This connector only supports delegated permissions via OAuth Authorization Code flow. The user that establishes connection to the connector when designing the Copilot Studio workflow must have access to Microsoft Security Copilot. Learn more about Microsoft Security Copilot Authentication.
- Data access for enhanced interaction - It's important that the authenticated user can access data from various remote security products, including tasks like reading Defender incident reports and gathering multifactor authentication (MFA) details, among others.
- Create a connection in Power Apps - You'll need to create a connection in Power Apps for the Copilot Studio action to connect to Security Copilot. Ensure that you create the

connection under same Power Apps environment as Copilot Studio on which you are creating the agent. For more information, see Manage connections.

Submit a Security Copilot prompt

- Create and set up a new agent in Copilot Studio.
- Select **Add action** to configure the initial trigger steps.

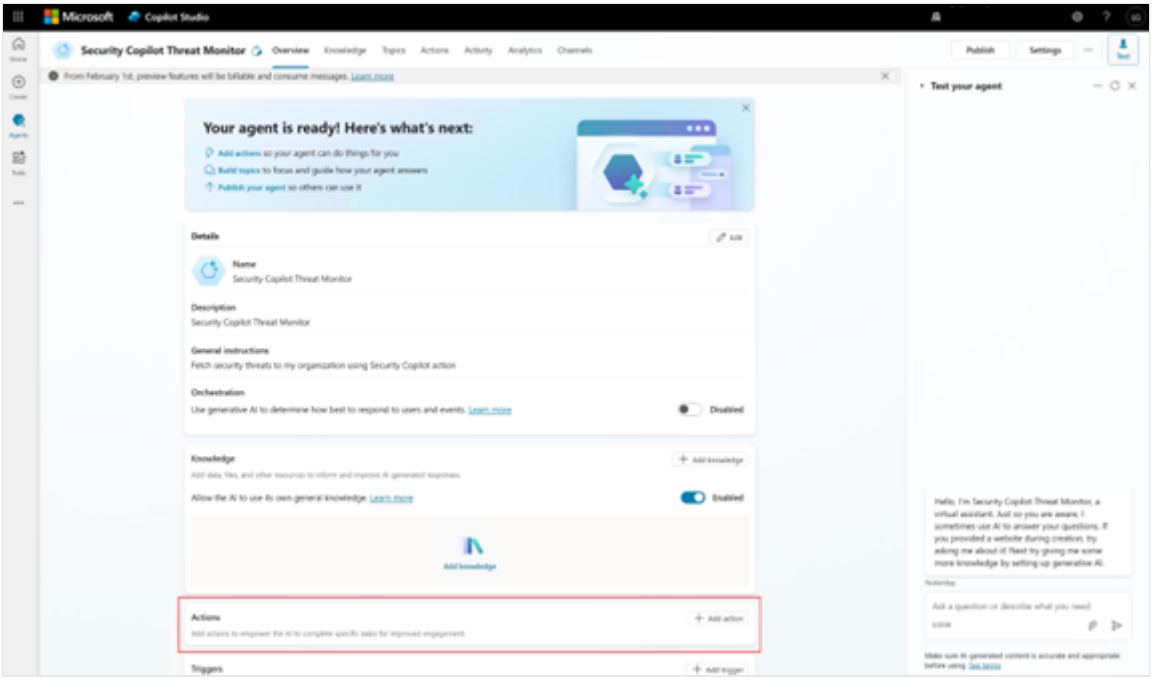

- Search for Microsoft Security Copilot to see the related actions you can use.
- Select **Submit a Security Copilot prompt**.

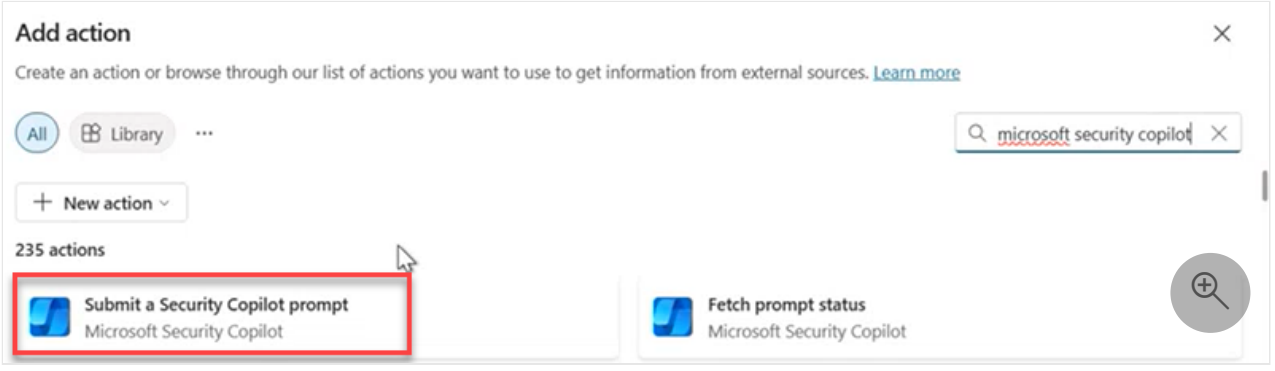

- Configure the action's connections, then select **Next**.
- Configure your action's options, inputs, and outputs, then select **Add action**.

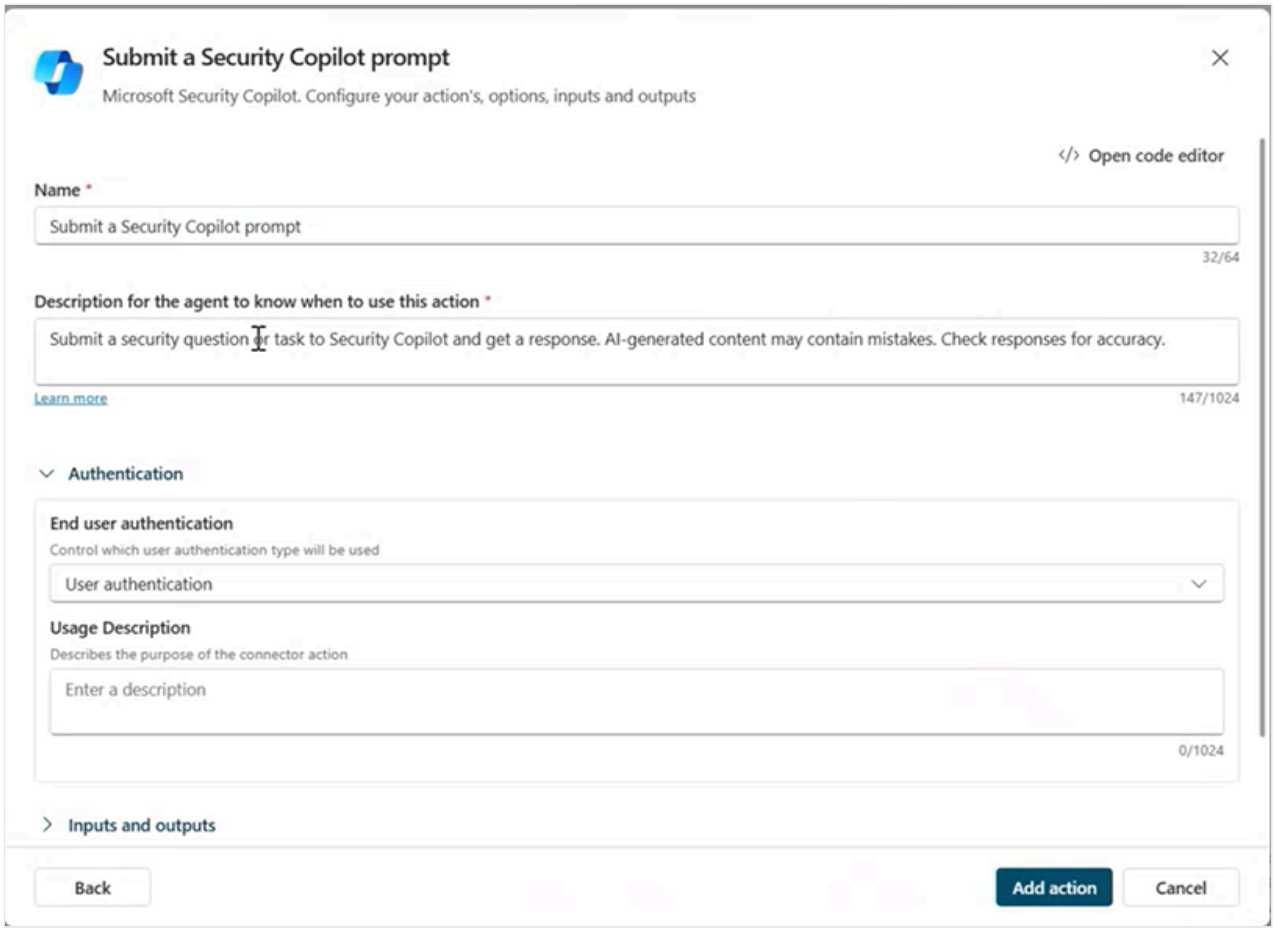

Proceed with adding the **Fetch prompt status** action to the agent.

Fetch a Security Copilot prompt status

- Using the same agent, select **Add action** to configure the initial trigger steps.
- Search for Microsoft Security Copilot to see the related actions you can use.
- Select **Fetch prompt status**.

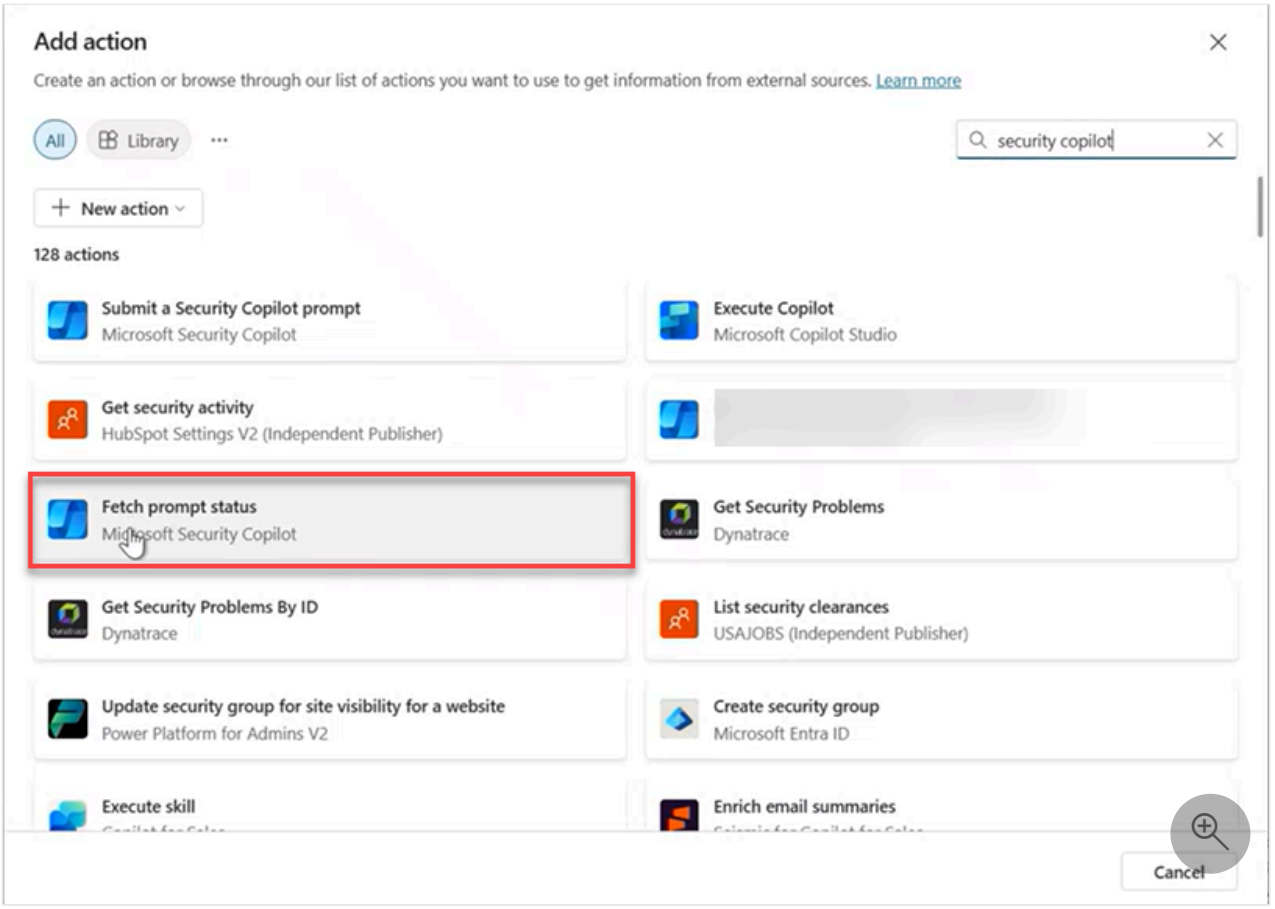

- Configure your action's options, inputs, and outputs, then select **Add action**.

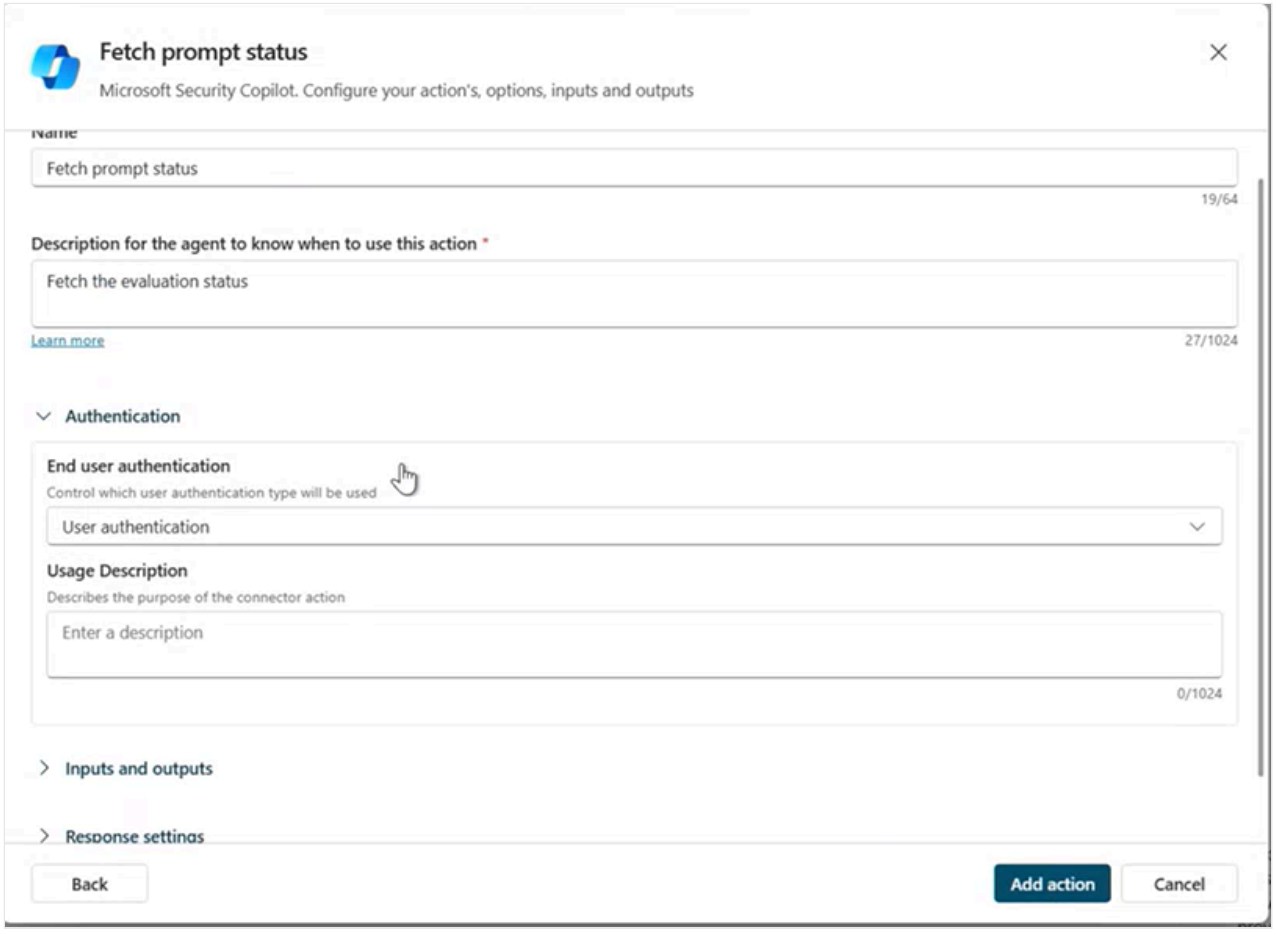

The two actions are now added to the agent.

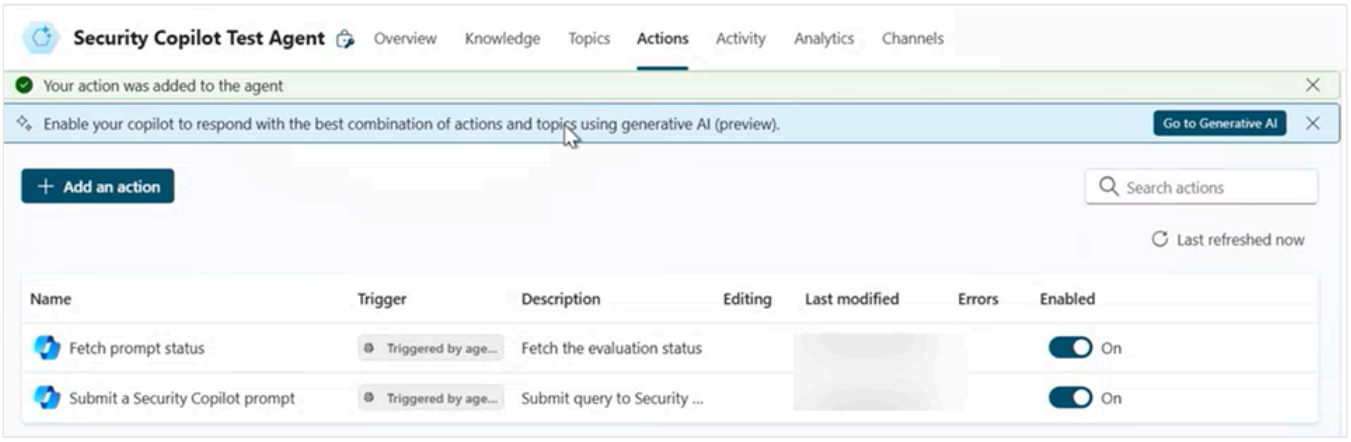

Test the agent

To test the agent, you'll need to turn on the Generative option in the settings page. For more information, see Turn on generative orchestration for an agent.

- Navigate to **Settings** > **Generative AI** > **Generative**.

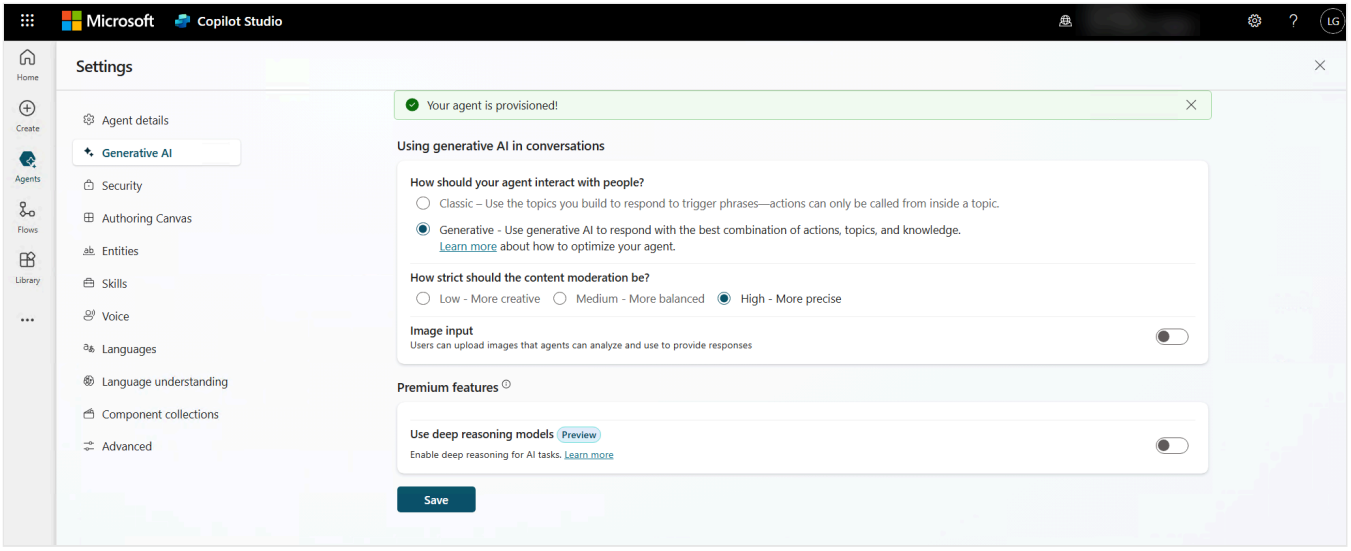

You'll need to configure this setting otherwise you won't be able to test out the action from the natural language chat experience.

You can test your agent in Copilot Studio by entering the prompt **Submit a Security Copilot, include the Session Id, Evaluation Id, and Prompt Id in the response*.

> **ⓘ Important**
>
> Make sure to include Session Id, Evaluation Id, and Prompt Id in the prompt.

You can also test **Fetch the status of the prompt** and you should receive the responses with the details.

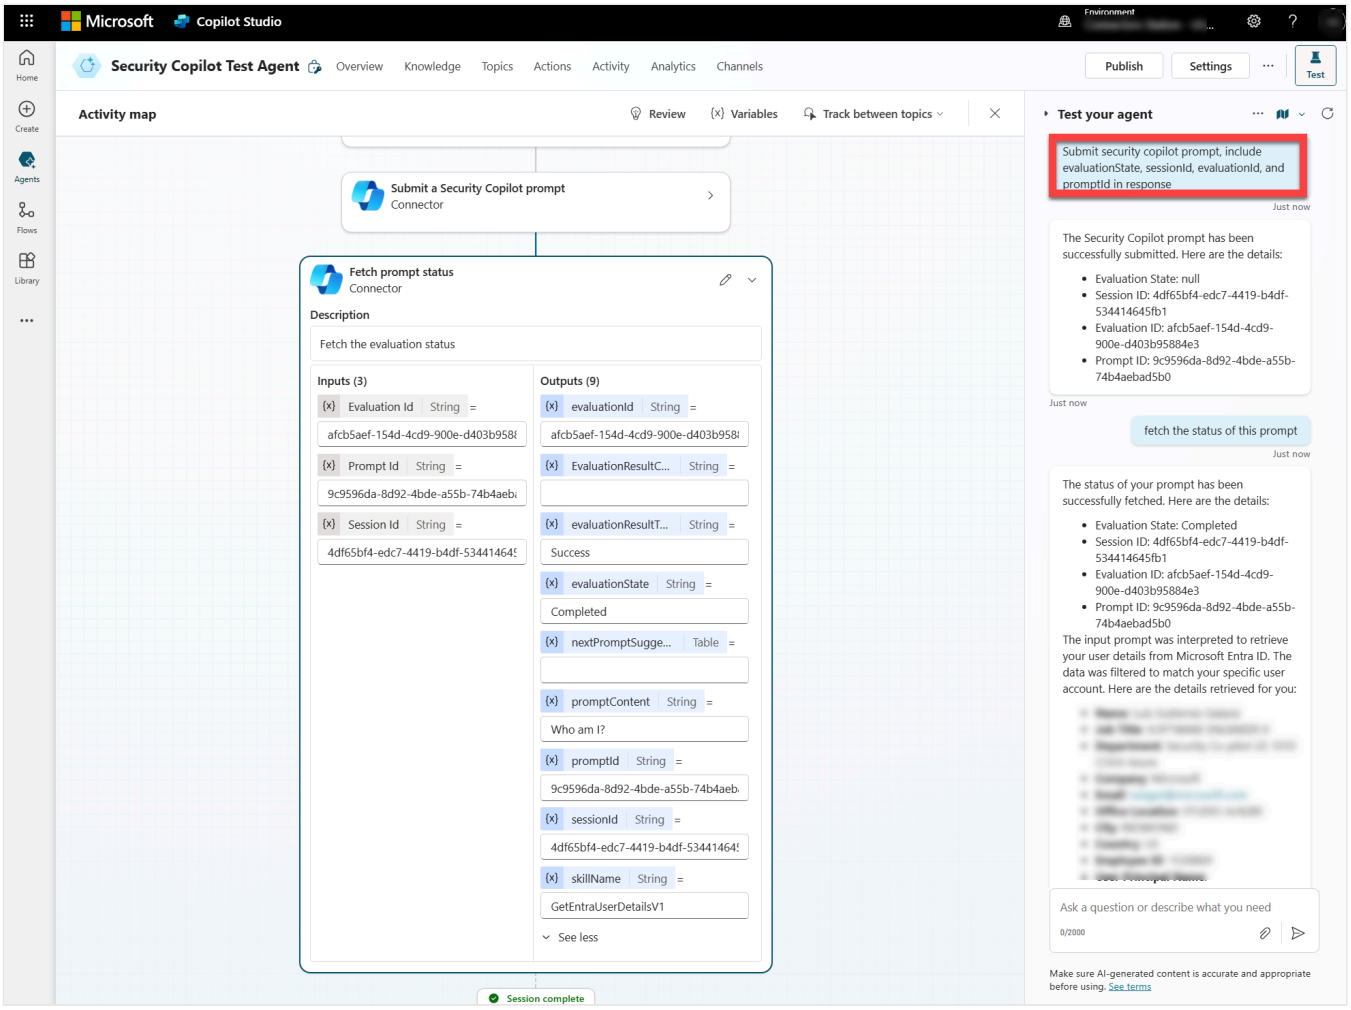

Last updated on 01/15/2026

Microsoft Security Copilot Logic Apps Connector

The Microsoft Security Copilot Logic Apps connector allows you to call into Copilot from an Azure Logic Apps workflow. This document provides an introduction to the connector action and its properties to get started integrating Security Copilot in your automation workflows.

The connector exposes 2 connector actions:

- Submit a Security Copilot prompt - Submit a natural language prompt to create a new Security Copilot investigation. After completion, the evaluation result will then be returned to your workflow.
- Submit a Security Copilot promptbook - Given a promptbook, invoke a new Security Copilot promptbook evaluation and return the output to your Azure Logic Apps workflow.

> **Tip**
>
> This connector is also available in all the Microsoft Power Platform experiences such as Microsoft Power Automate and Power Apps.

For more information, see Power Platform Connector Page.

Prerequisites

- Tenant - Ensure that your tenant administrator sets up access to Microsoft Security Copilot before using the connector.
- User Authentication - This connector only supports delegated permissions via OAuth Authorization Code flow. The user that establishes connection to the connector when designing the Azure Logic Apps workflow must have access to Microsoft Security Copilot. Learn more about Microsoft Security Copilot Authentication.
- Data access for enhanced interaction - It's important that the authenticated user can access data from various remote security products, including tasks like reading Defender incident reports and gathering multifactor authentication (MFA) details, among others.

Submit a Security Copilot prompt

- Create and set up a new Logic Apps workflow in the Azure portal.

- The connector is available in both "Standard" and "Consumption" plan types.

- Configure the initial trigger step and proceed to search for the Copilot action "Submit a Security Copilot prompt".

- Fill in the parameters on the Copilot action.

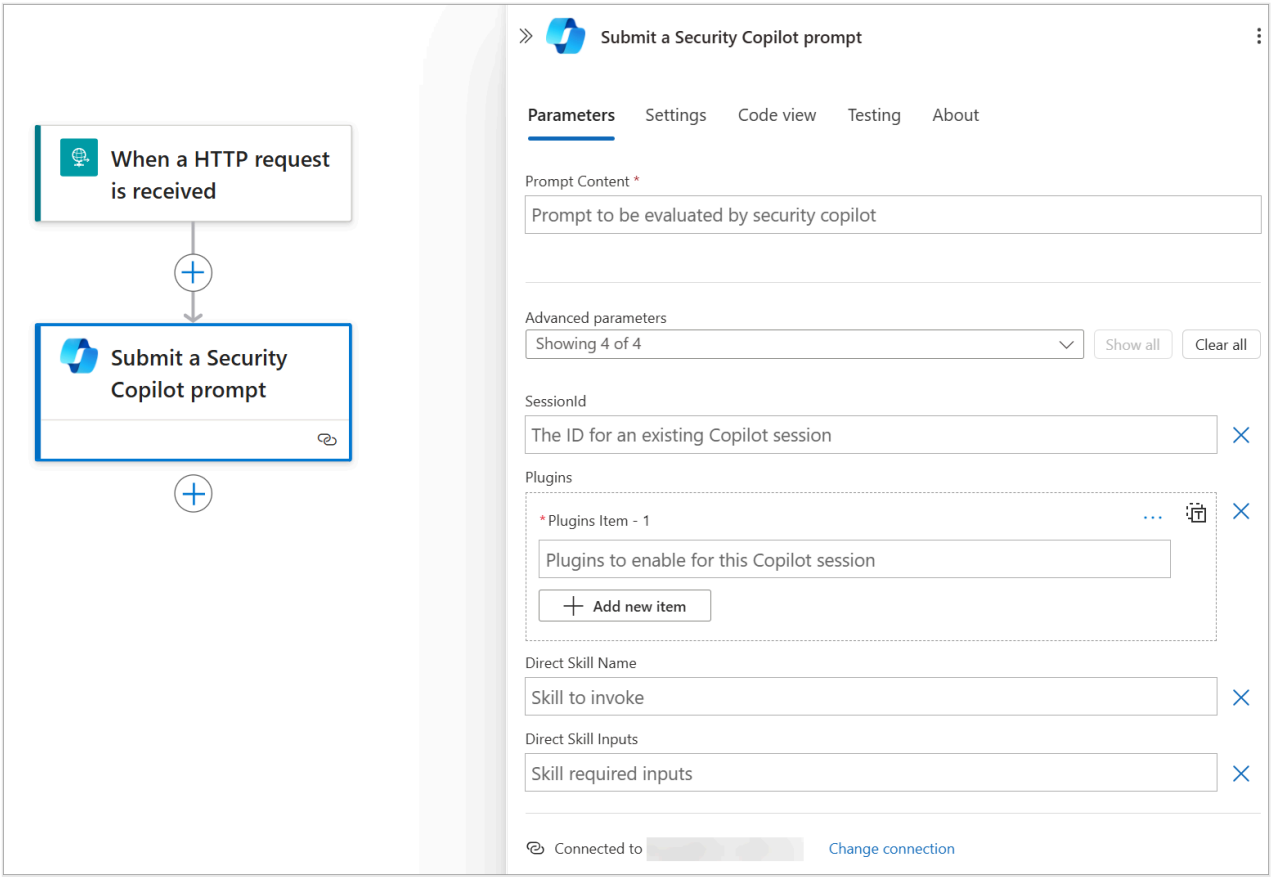

- Reference the following descriptions for the Copilot Action Parameters

  - Prompt Content (Required): Enter the natural language prompt to be evaluated.

  - Session ID (Optional): Enter an existing Copilot sessionId for maintaining continuity in your tasks, allowing to ask prompts within the same investigation. If omitted, Copilot creates a new investigation.

  - Plugins (Optional): Enter a list of plugin names for your investigation, scoping down the potential plugins that Copilot can pick to execute and answer your natural language prompt. Useful to avoid plugin collisions.

  - Direct Skill Name (Optional): Enter a specific Copilot skill to execute. Convenient way to bypass the standard Copilot planner if you know a specific skill to trigger.

  - Direct Skill Inputs (Optional): Enter a JSON body containing key/value pairs for required skill parameters. Useful when providing a Direct Skill Name to execute.

Submit a Security Copilot promptbook

- Create a new Azure Logic Apps workflow in the Azure portal.
  - The connector is available in both "Standard" and "Consumption" plan types.
- After setting up the Azure Logic App and configuring the initial trigger step, proceed to search for the Security Copilot action.

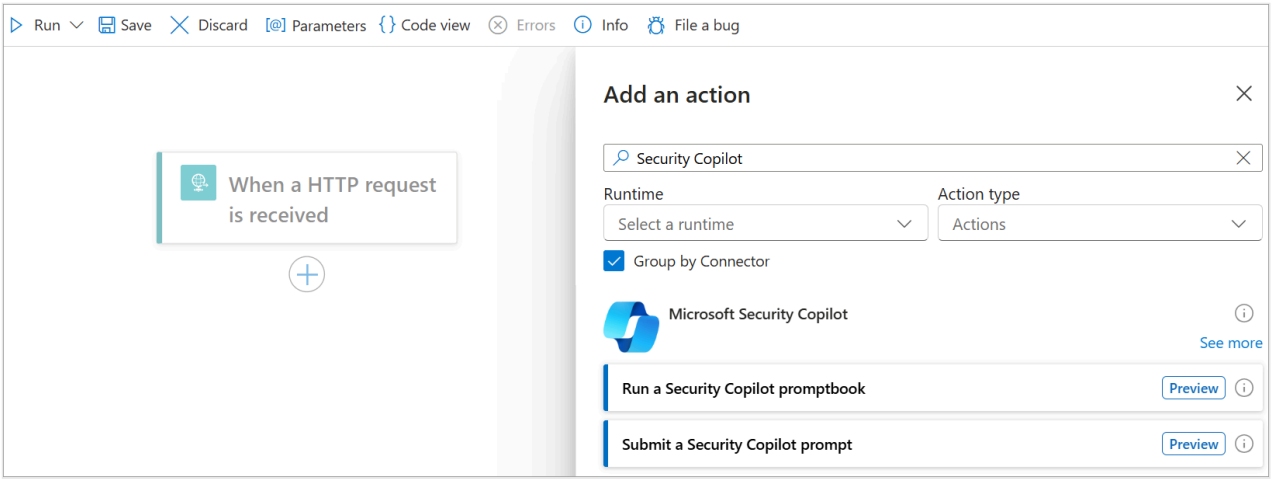

- After choosing the new Security Copilot action, proceed to fill in the following information.

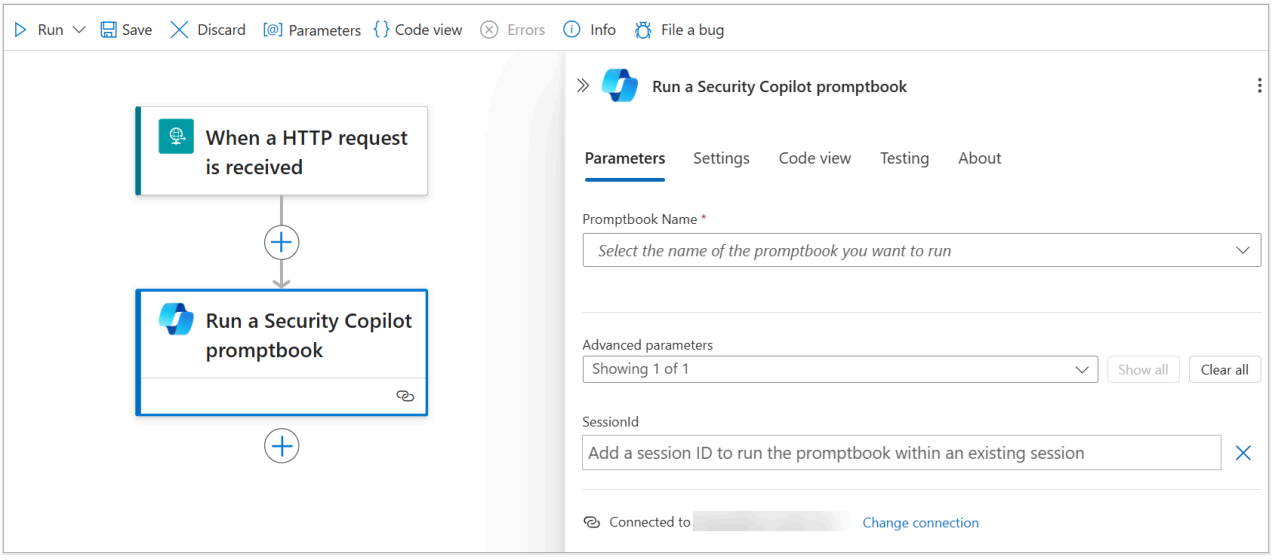

- Promptbook Name (required): Select a promptbook from the dropdown list.
- Promptbook Inputs (required): Inputs are displayed dynamically for promptbooks with input fields. Enter the specific inputs for each input field.

> **Note**
>
> You will need to ensure that required fields, such ash `<SENTINEL_INCIDENT_ID>`, `<THREATACTORNAME>`, or `<DEFENDER_INCIDENT_ID>` are filled in the input fields.

Otherwise, the Azure Logic App workflow will not run successfully.

- Session ID (optional): Enter an existing Security Copilot sessionId for maintaining continuity in your tasks, allowing to ask prompts within the same session. If omitted, Security Copilot will create a new session.

Last updated on 01/15/2026

Trigger Security Copilot agents from Azure Logic Apps

Building on how organizations already use Logic Apps to automate security workflows, a new connector action for Security Copilot in Logic Apps flows allows teams to easily invoke agents they create as part of repeatable workflows. This automated workflow brings deeper AI-driven investigation, context, and decision support into tasks such as incident triage, threat intelligence analysis, and policy validation.

This article explains how to trigger Microsoft Security Copilot agents within the automated Azure Logic Apps workflows by configuring prerequisites, adding the agent action in Logic Apps, and verifying the agent run in Security Copilot.

Using the Security Copilot Logic Apps agent connector, you can:

- Trigger previously configured Microsoft Security Copilot agents from Logic App
- Chain agent execution with other Logic Apps actions
- Pass inputs dynamically from earlier workflow steps including advanced parameter support

Before you begin

Make sure you have the following resources and permissions:

- Access to **Microsoft Security Copilot** with permission to create and run agents.
- An agent that is already created and configured (agent that is set up and ready to run) in Security Copilot.

  > **Note**
  >
  > When you create your own agent, you must upload the plugin containing the agent to **Anyone in this Workspace** to be available for triggering from Logic Apps.

- Access to **Azure Logic Apps** (Consumption or Standard).
- Permission to create connections in Azure Logic Apps.

Step 1: Set up an agent in Security Copilot

Logic Apps can only trigger agents that are fully configured and available in your Security Copilot workspace.

1. Create a new agent or select an existing agent in Security Copilot. The following table provides reference to documentation for creating agents:

Expand table

Type	Description
Plugins	Build and upload a custom plugin
Agents	Build custom agents using natural language or the agent builder form

2. Complete the agent **Set up** process. For more information, see Setup agents.

3. Configure the **agent run identity** during setup.

    - This identity is used when the agent executes.
    - It is separate from the identity used to sign in to the Logic Apps connector.

4. If you create your own **custom agent**, you must upload as a plugin that contains the agent and set sharing to **Anyone in this workspace**, so that the agent is available to trigger from Azure Logic Apps.

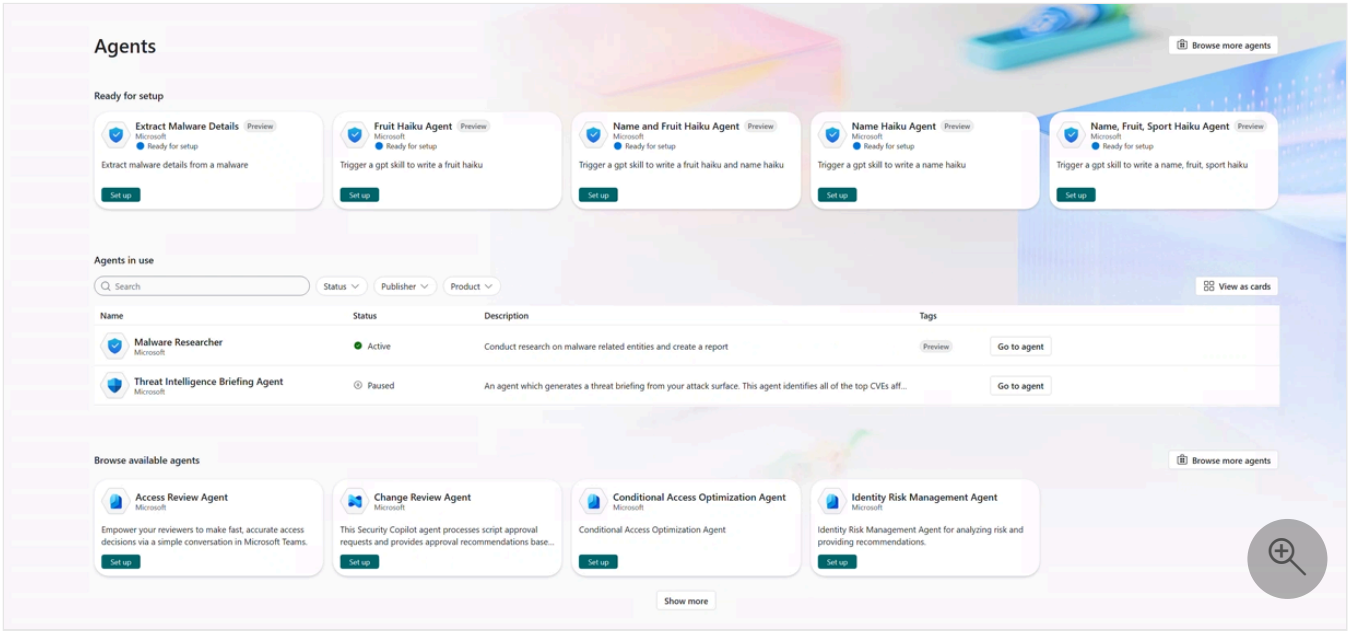

The illustration shows the agents that are set up and ready to use in the **Agents in use** section.

> **Note**
>
> Currently, agents must be created and set up in the **default workspace**. Agents created in non-default workspace aren't available to trigger from Logic Apps. So, if an agent isn't shared with the workspace or is missing required setup, it does not appear in Logic Apps.

Step 2: Add an action in Azure Logic Apps

1. In the Azure portal, create a new Logic App or open an existing one.
2. Open the **Logic app designer**.
3. Add a **trigger** as the first step to start any workflow in Azure Logic Apps. For information on adding a trigger, see Add a trigger.
4. Select **Add an action** in the Logic app designer.

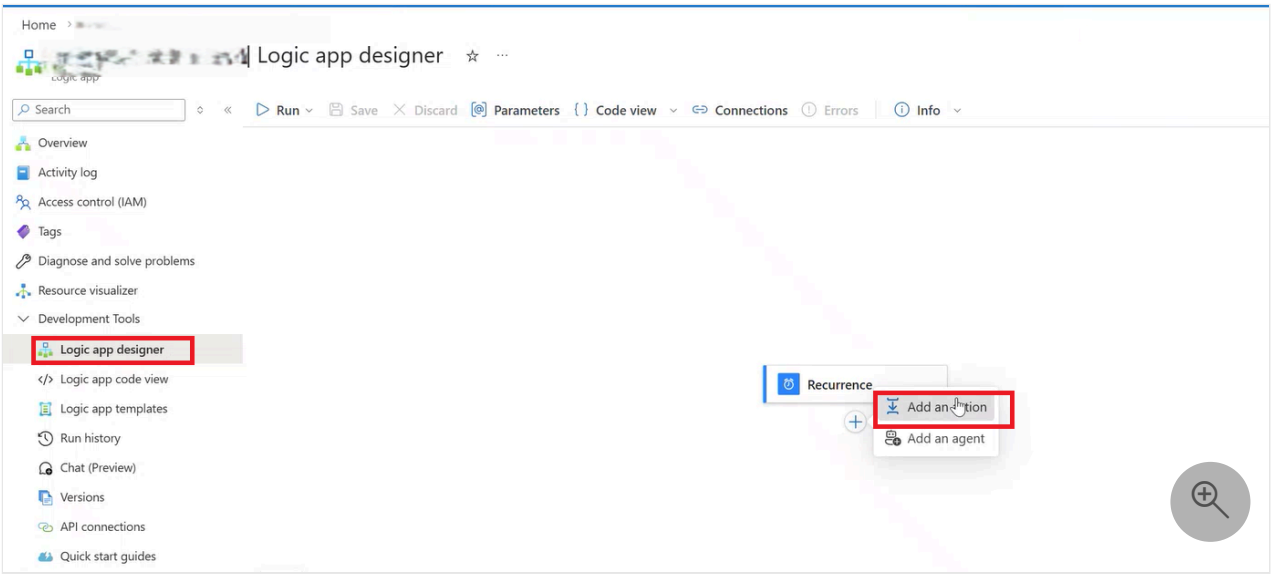

Step 3: Add the Security Copilot agent action

1. Search for **Security Copilot** connector. The Security Copilot actions are displayed.
2. Select the action **Execute a Security Copilot Agent**.

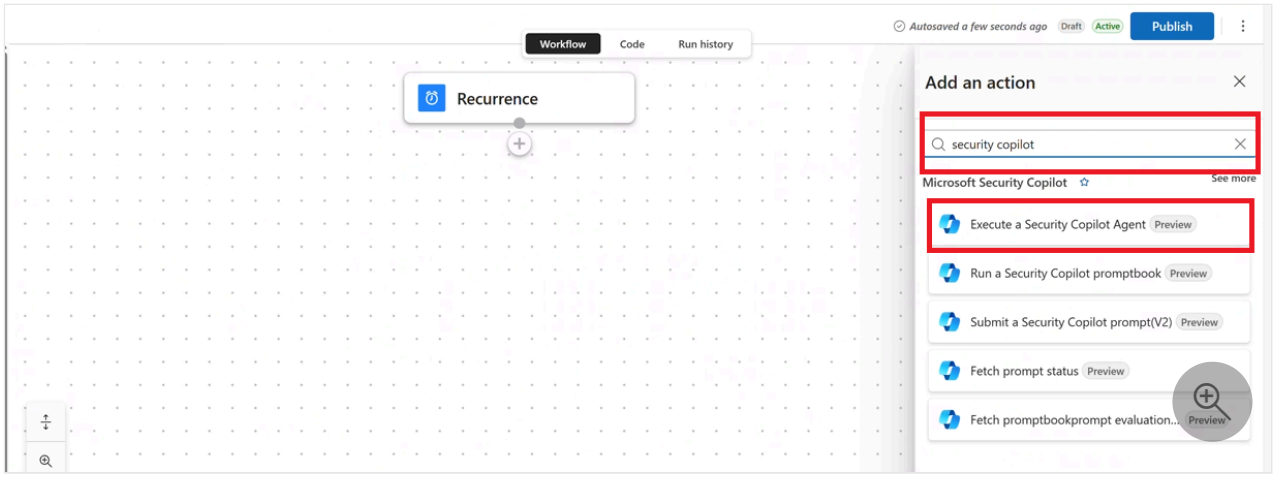

Step 4: Create the connector connection

Create a connection when you add the agent action for the first time.

1. The **Create connection** window is displayed.

2. Sign in with an identity that has access to Security Copilot to create a new connection.

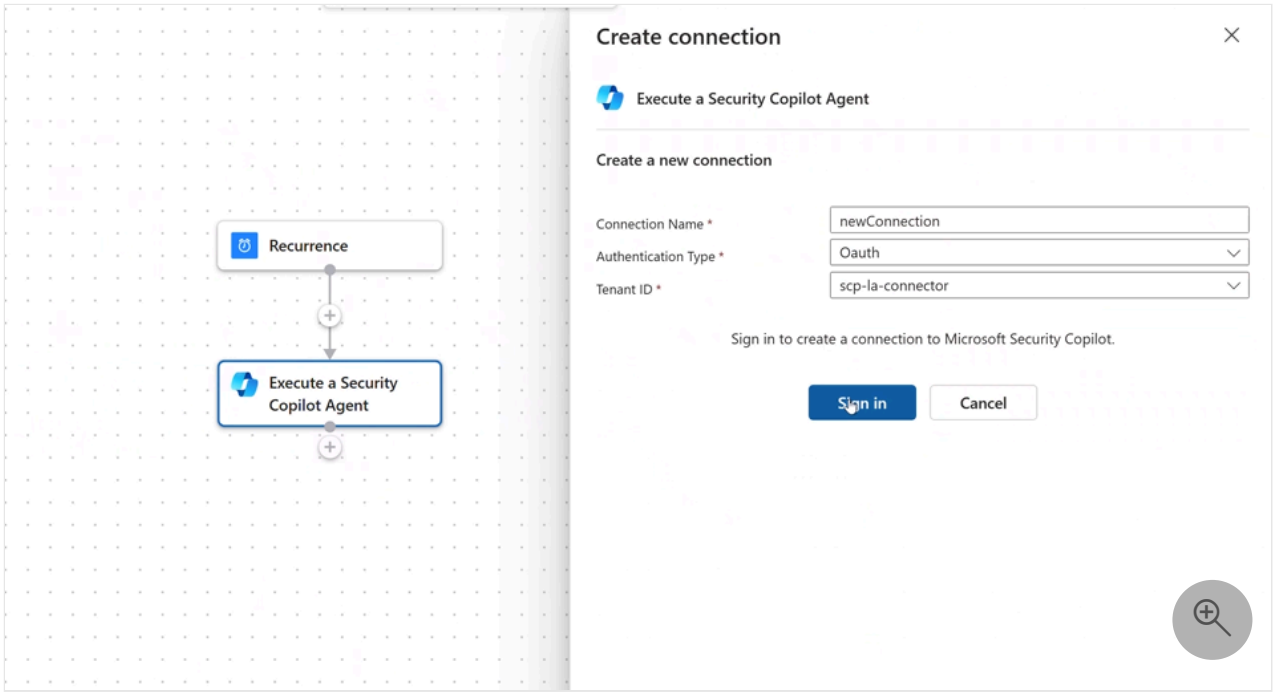

> **Important**
>
> The connection sign-in identity is used only to access Security Copilot. The agent itself runs using the agent run identity configured during agent setup.

Step 5: Configure the agent parameters

1. In the **Agent Name** field, select the agent you want to run.

2. Select the trigger configuration (the default trigger is typically sufficient).

3. Provide required parameters.

    - Fill in all the required inputs.
    - Find optional inputs under **Advanced parameters**.

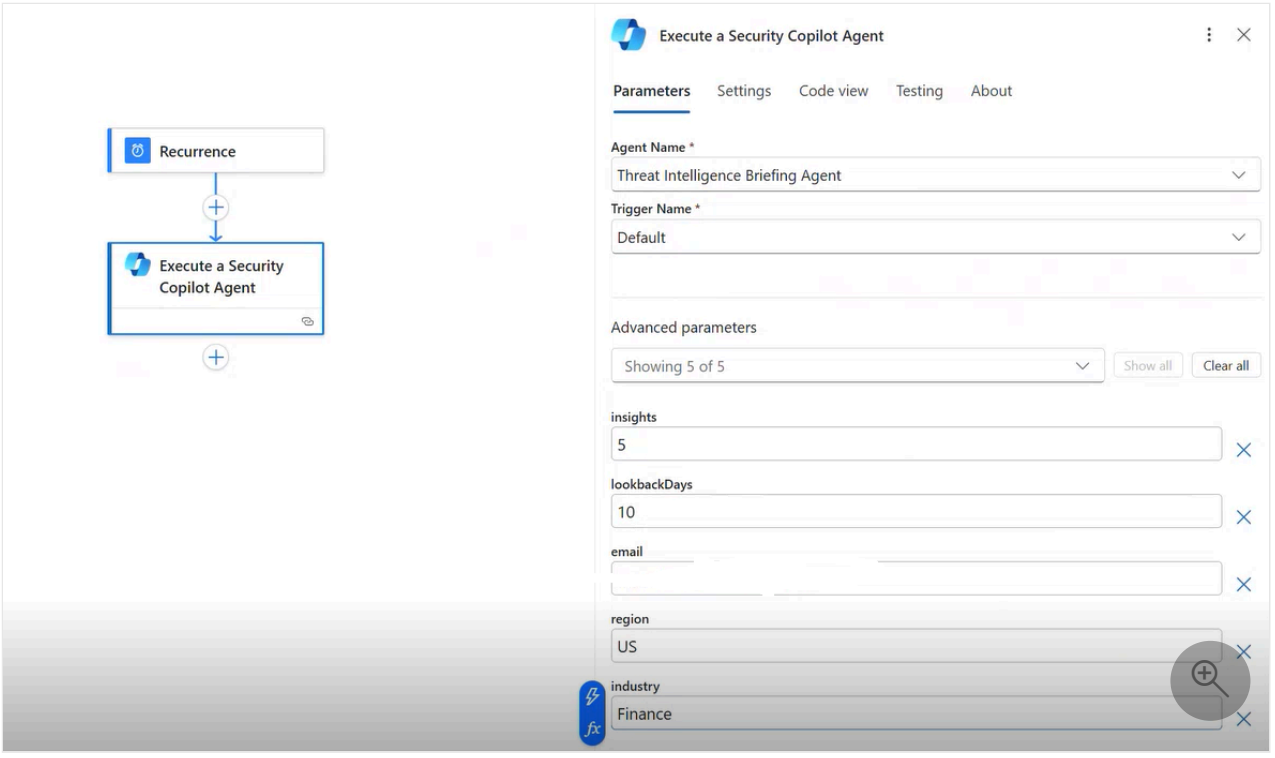

4. (Optional) You can pass in outputs of previous steps as agent inputs.

5. Select **Save** to save the workflow.

6. **Run** the Logic App manually or wait for the trigger to fire.

7. Select **Show runs** to see the **Runs history**. Confirm that the agent action completed successfully.

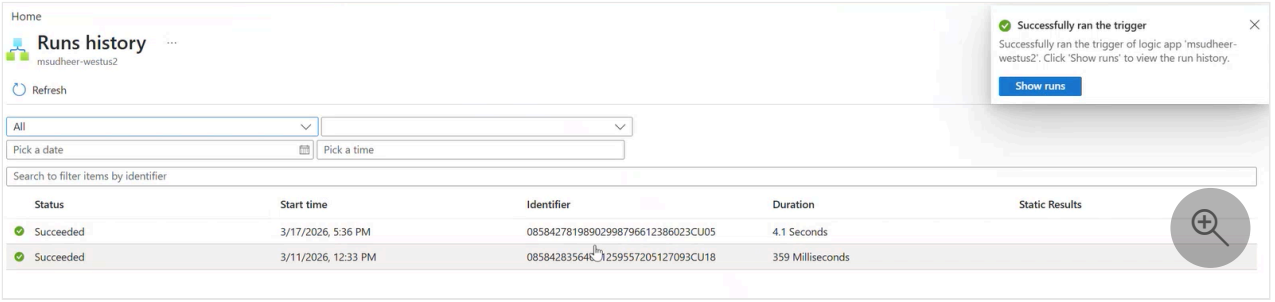

8. You can also see the response in the Status as **Succeeded**.

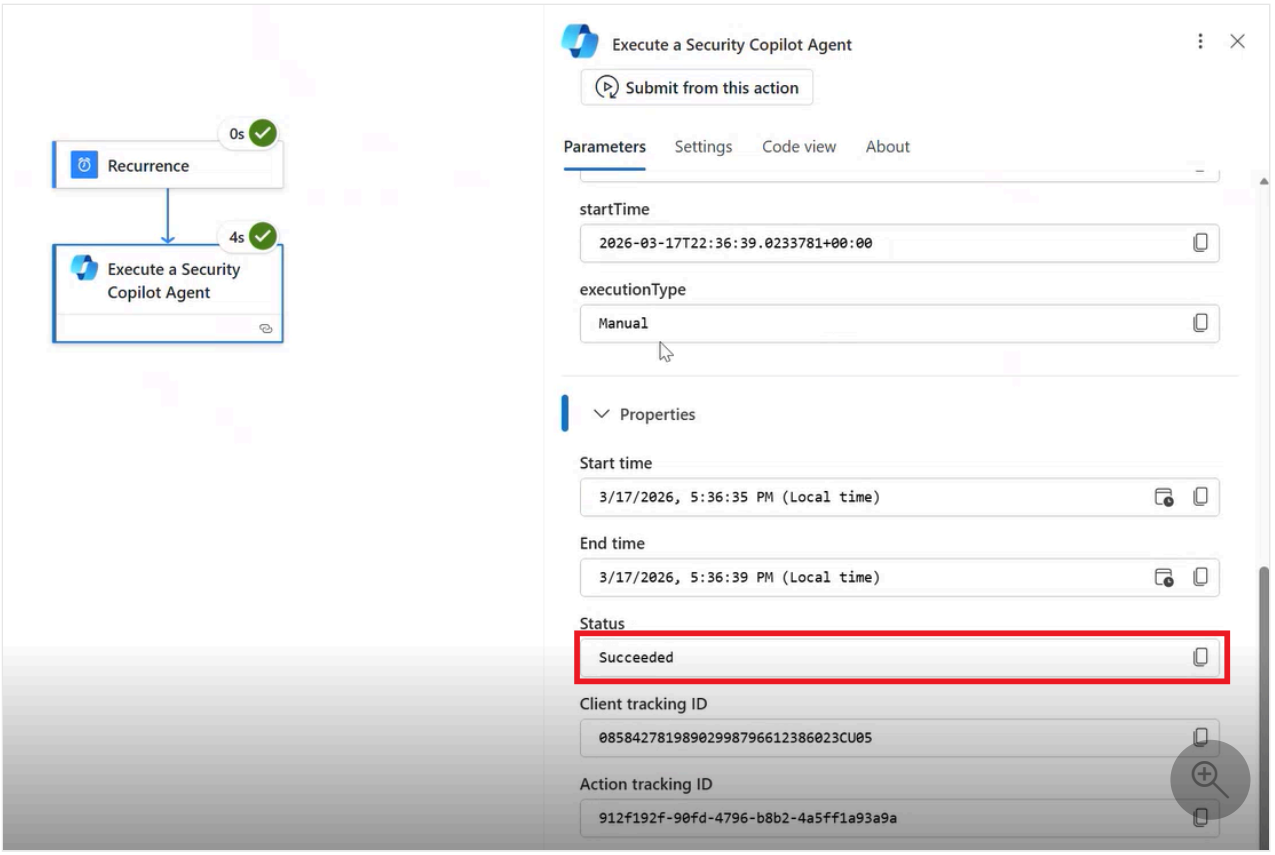

Step 6: Verify the agent run in Security Copilot

1. Navigate to **Security Copilot**.
2. Open the agent you triggered.
3. Confirm that the agent run appears in the agent run history.

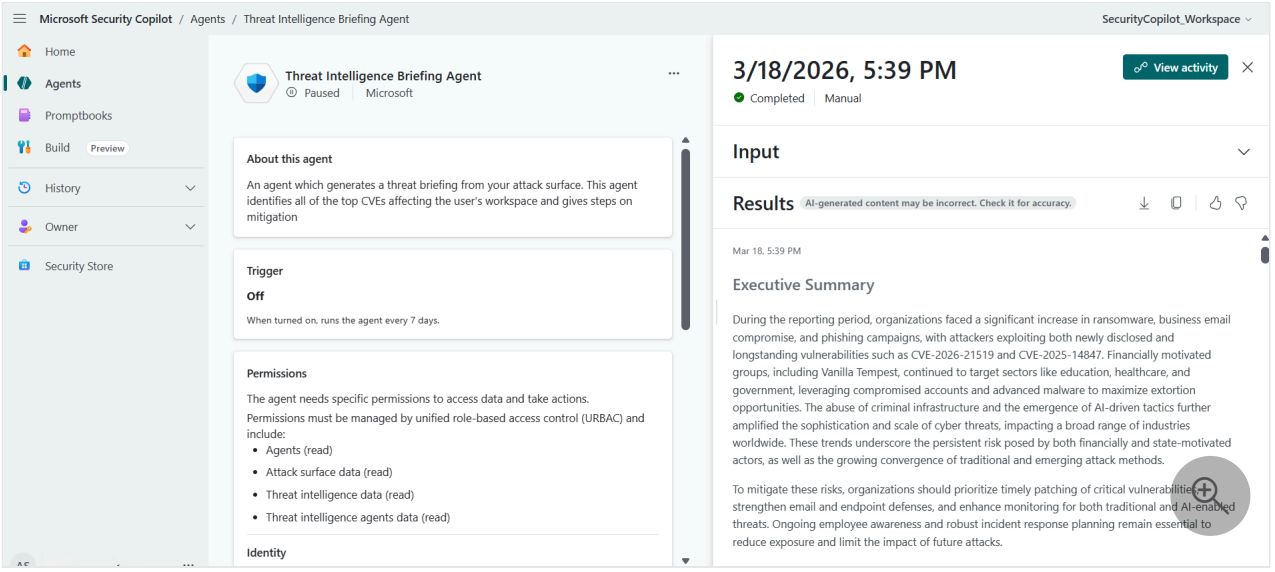

You can also edit the Logic Apps action to select a different agent.

Troubleshooting

Agent doesn't appear in Logic Apps

- Make sure you fully set up the agent in Security Copilot.
- For custom agents, confirm you share the plugin that contains the agent with **Anyone in this workspace** so the agent is available to trigger from Logic Apps.

Permission or sign-in errors

- Verify the Logic Apps connection identity has access to Security Copilot.
- Verify the agent run identity has the required permissions.

Missing required inputs

- Check that you provide all required agent inputs in the action configuration.

Last updated on 04/10/2026

Security Copilot Logic Apps Connector error codes

08/25/2025

Learn about the error codes that you might encounter when using the Security Copilot Logic Apps Connector.

Expand table

Code	Errors	Details
403	Forbidden	- You might not have the right permissions assigned to you. For more information, see Understand authentication in Microsoft Security Copilot.
429	Too many requests	- Ensure you have sufficient capacity. For more information, see Manage usage of security compute units in Security Copilot.

Use promptbooks in Microsoft Security Copilot

Security Copilot comes with prebuilt promptbooks, a series of prompts that have been put together to accomplish specific security-related tasks. They can function in a similar way as security playbooks—ready-to-use workflows that can serve as templates to automate repetitive steps—for instance, regarding incident response or investigations. Each prebuilt promptbook requires a specific input (for example, a code snippet or a threat actor name).

View the promptbook library

You can find the different promptbooks by going to the promptbook library or by selecting the Prompts icon at the prompt bar. You can then search for a promptbook or select **See all promptbooks** to view all.

Use different promptbooks

Available promptbooks include:

- Check impact of an external threat article
- Incident investigation
- Microsoft User analysis
- Suspicious script analysis
- Threat actor profile
- Threat Intelligence 360 report based on MDTI article
- Vulnerability impact assessment

Watch the following video to learn more about promptbooks:

https://learn-video.azurefd.net/vod/player?id=7bb51c3d-ba80-430e-b3bc-2eed3942cf70&locale=en-us&embedUrl=%2Fcopilot%2Fsecurity%2Fusing-promptbooks

Check impact of an external threat article

Using this promptbook, you can analyze any external threat intelligence articles to extract indicators from it and gather any relevant Microsoft Defender Threat Intelligence article.

To run the **Check impact of an external threat article** promptbook:

1. Make sure that you have turned on the **Microsoft Threat Intelligence** plugin by following the steps in Enable the Security Copilot integration in Defender TI.

2. In Security Copilot, select the Prompts button in the prompt bar and start to type a few letters of the promptbook named **Check impact of an external threat article** until the promptbook appears in the list.

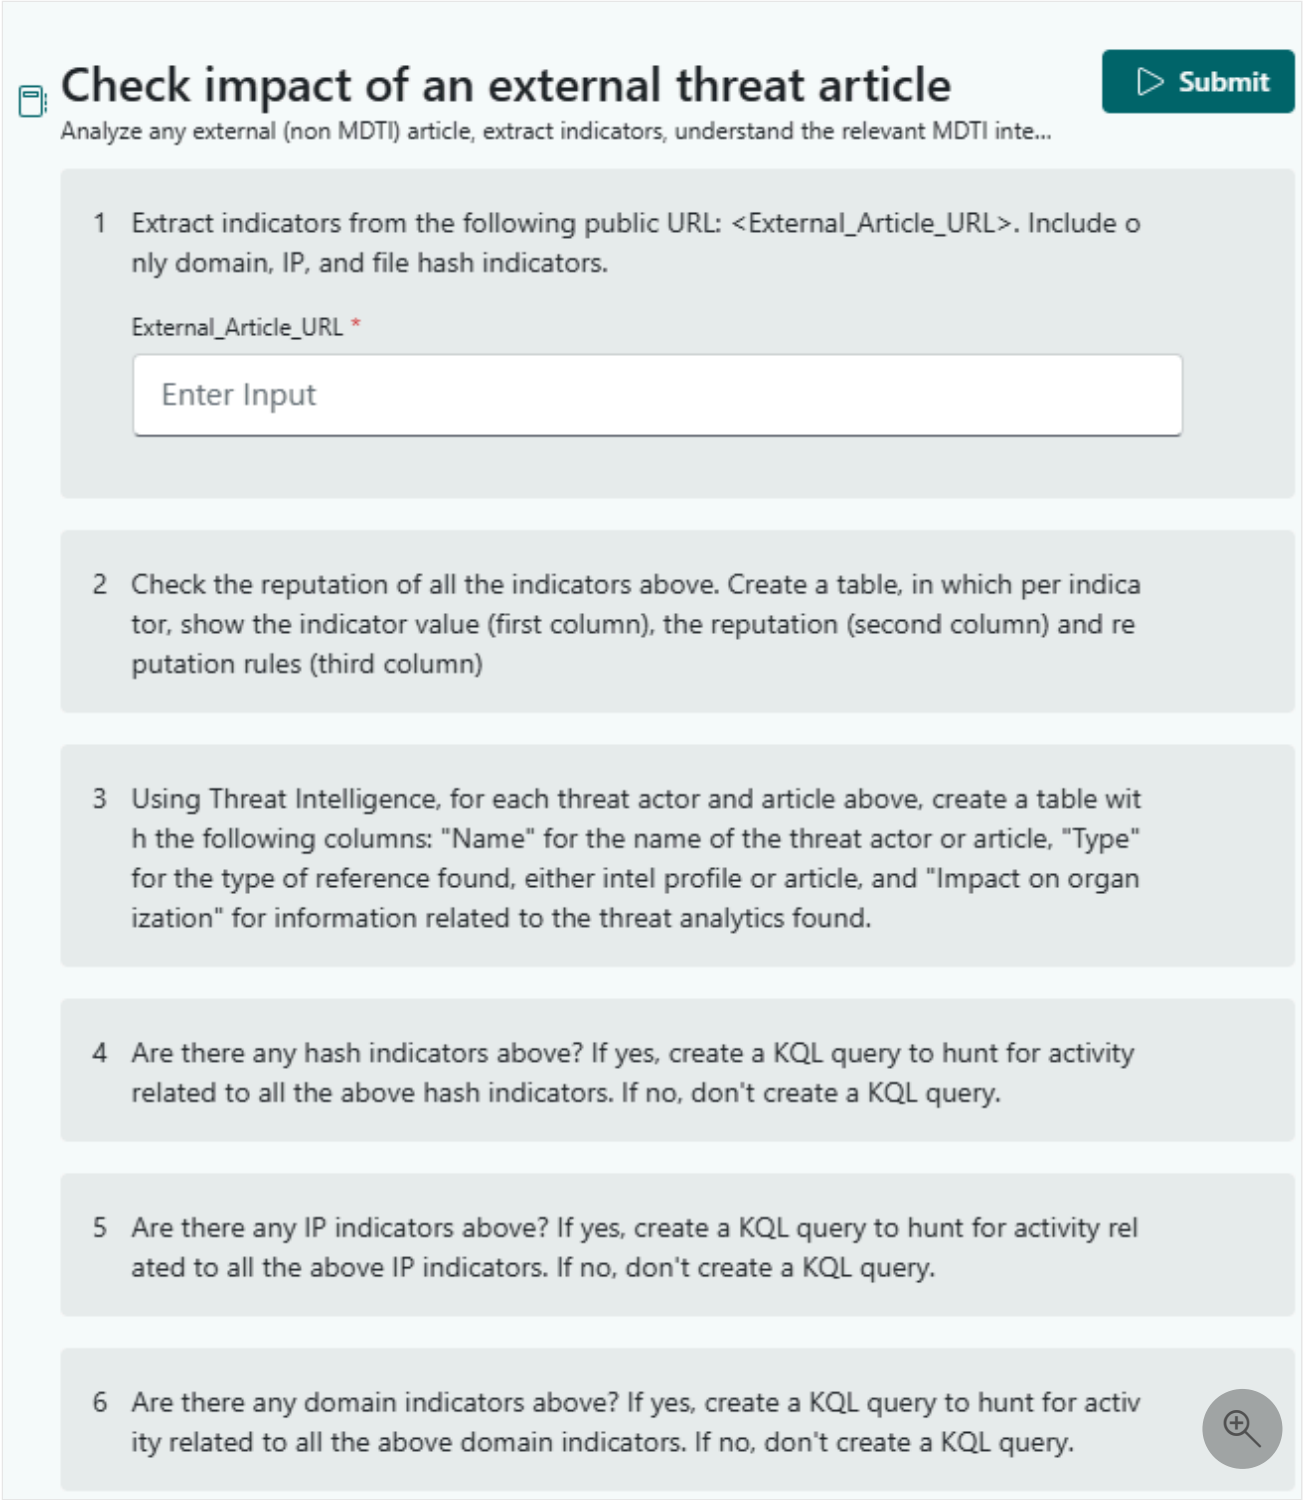

3. Provide the URL of the external threat intelligence article you want to analyze and extract indicators from.

4. Next, select **Submit**.

5. Wait for Security Copilot to run the URL through the different prompts. If you see a round progress indicator in place of the response, then the promptbook is still running. Security Copilot generates responses for each of the prompts, building on each response until it gets to the last prompt.

6. Read the responses by Security Copilot. You can use the generated KQL queries to aid in investigation.

Incident investigation

You can run the incident investigation promptbook after supplying an incident number to either the Microsoft Sentinel or Microsoft Defender XDR plugin. Use the appropriate promptbook for the plugin you want to use. The incident investigation promptbooks contain several prompts to generate an executive report for a nontechnical audience that summarizes the investigation. Each prompt builds on the previous prompt.

To run the Microsoft Sentinel **Incident investigation** promptbook:

1. Select the Prompts button in the prompt bar and start to type "incident investigation" until the promptbooks appear in the list.

2. Select **Microsoft Sentinel incident investigation**. (To use the Microsoft Defender XDR plugin instead, select **Microsoft Defender XDR incident investigation**.)

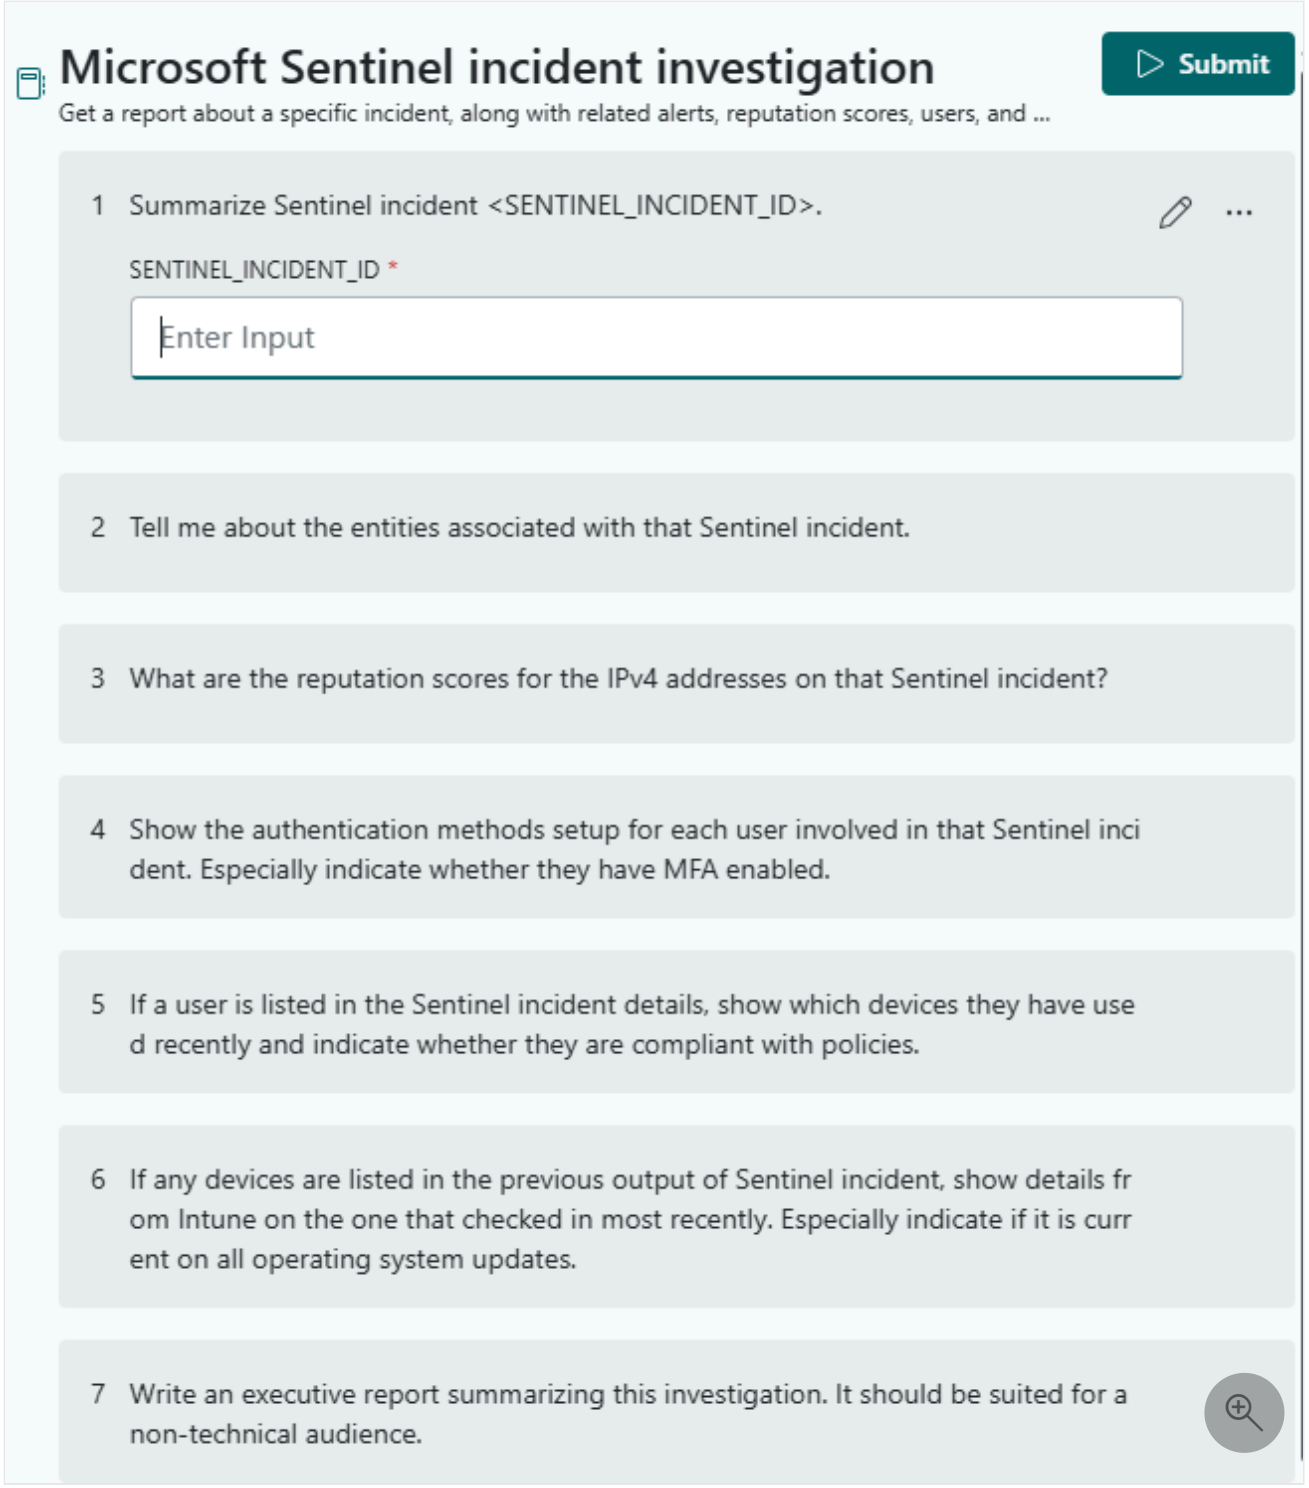

3. Supply the incident number that you want to investigate in the input box that says *SENTINEL_INCIDENT_ID*.

4. Next, select **Submit**.

5. Wait for Security Copilot to run the incident number through the different prompts. If you see a round progress indicator in place of the response, then the promptbook is still running. Security Copilot generates responses for each of the prompts, building on each response until it gets to the last prompt.

6. Read the responses by Security Copilot. The last prompt by Security Copilot generates an executive report summarizing the investigation based on the responses. Review and verify whether the responses are accurate and meet your needs.

Microsoft User analysis

The Microsoft User analysis promptbook can be used by an IT Admin to analyze and get detailed insights about a user and associated devices across multiple Microsoft 365 products. This includes sign-in and authentication data from Microsoft Entra ID, device information from Intune, unusual activity details from Microsoft Purview, and a Microsoft Defender summary highlighting important detections.

To get a comprehensive response from this promptbook, you first need to activate or make sure that you have the following roles:

- Security Reader role for Microsoft Entra ID, Intune, and Defender at a minimum
- Insider Risk Management Investigator or Analyst role for Microsoft Purview

To run the **Microsoft User analysis** promptbook:

1. Go to the Promptbook library and look for the **Microsoft User Analysis** promptbook.

2. Select **Start new session**.

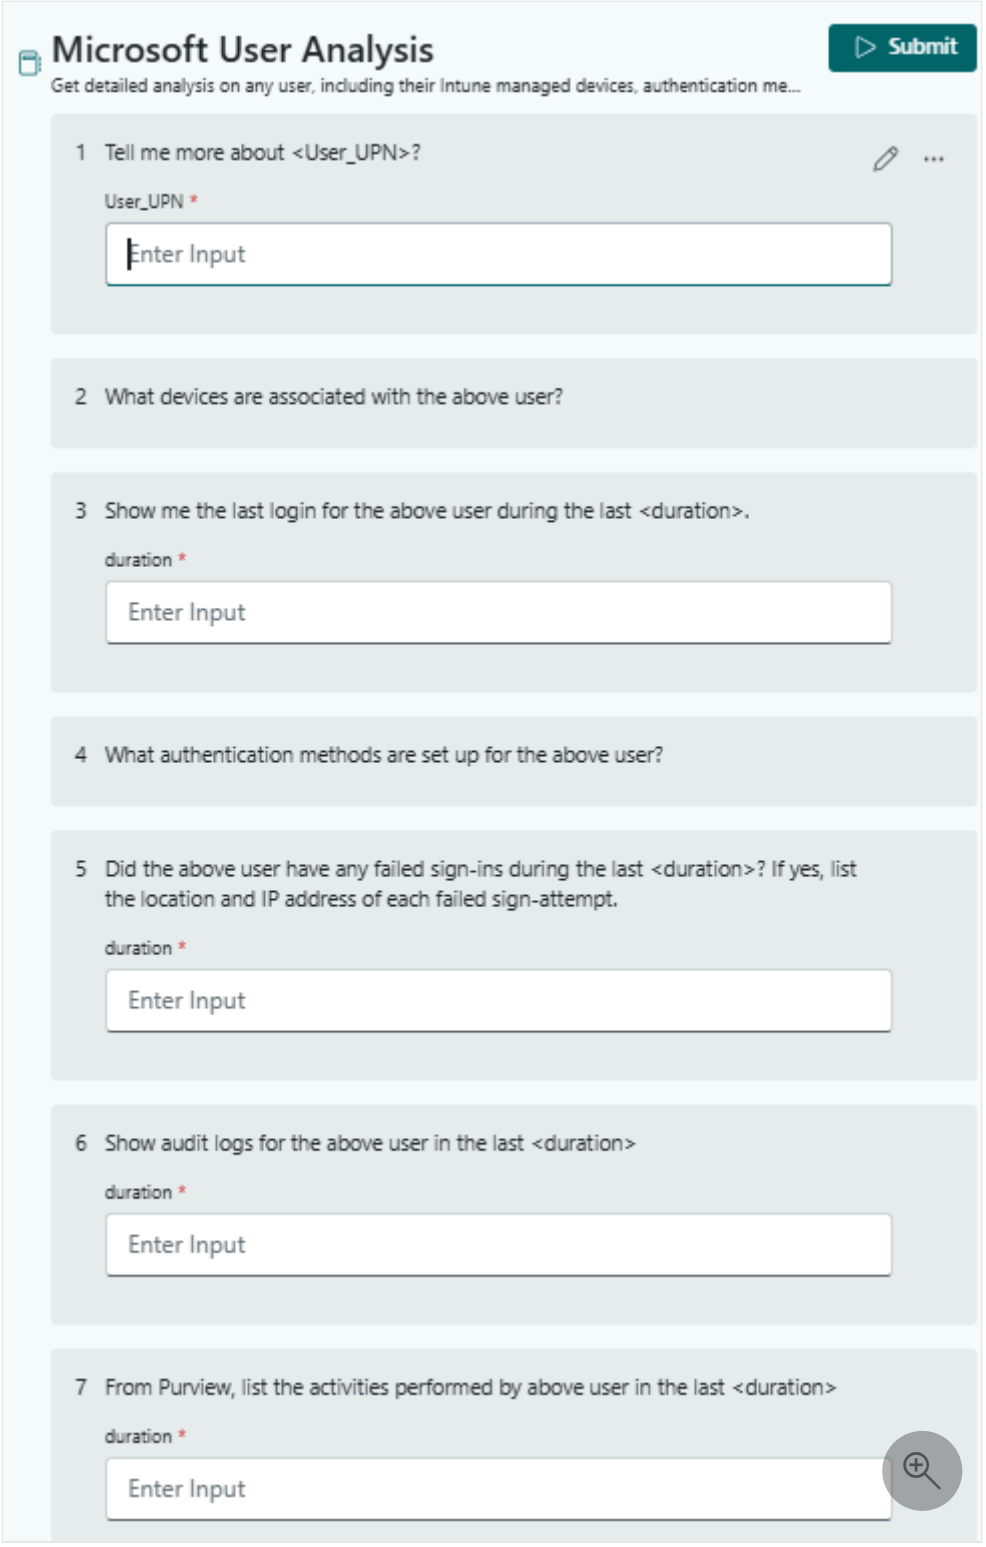

3. You need the following inputs:

    - the user principal name or UPN of the user
    - the time range you want Security Copilot to look up the information.

4. Next, select the **Submit** button in the upper right corner of the dialog box.

5. Wait for Security Copilot to run your inputs through the different prompts. If you see a round progress indicator in place of the response, then the promptbook is still running.

Security Copilot generates responses for each of the prompts and builds on each until it gets to the last prompt.

6. Read the response from Security Copilot. Using the response from the prompts, you can deduce faster whether the user being investigated has been performing suspicious activities so you can focus on your next steps in securing your system.

Suspicious script analysis

The suspicious script analysis promptbook is useful when you're investigating a PowerShell or Windows command-line script. For example, if a PowerShell script was involved in a critical incident in your network, you can copy the body of the script and run the promptbook to find out more about it.

To run the **Suspicious script analysis** promptbook:

1.Select the Prompts button in the prompt bar and start to type "suspicious script analysis" until the promptbooks appear in the list.

1. Select **Suspicious script analysis**.

2. Paste the script string that you want analyzed in the input box saying *SNIPPET*.

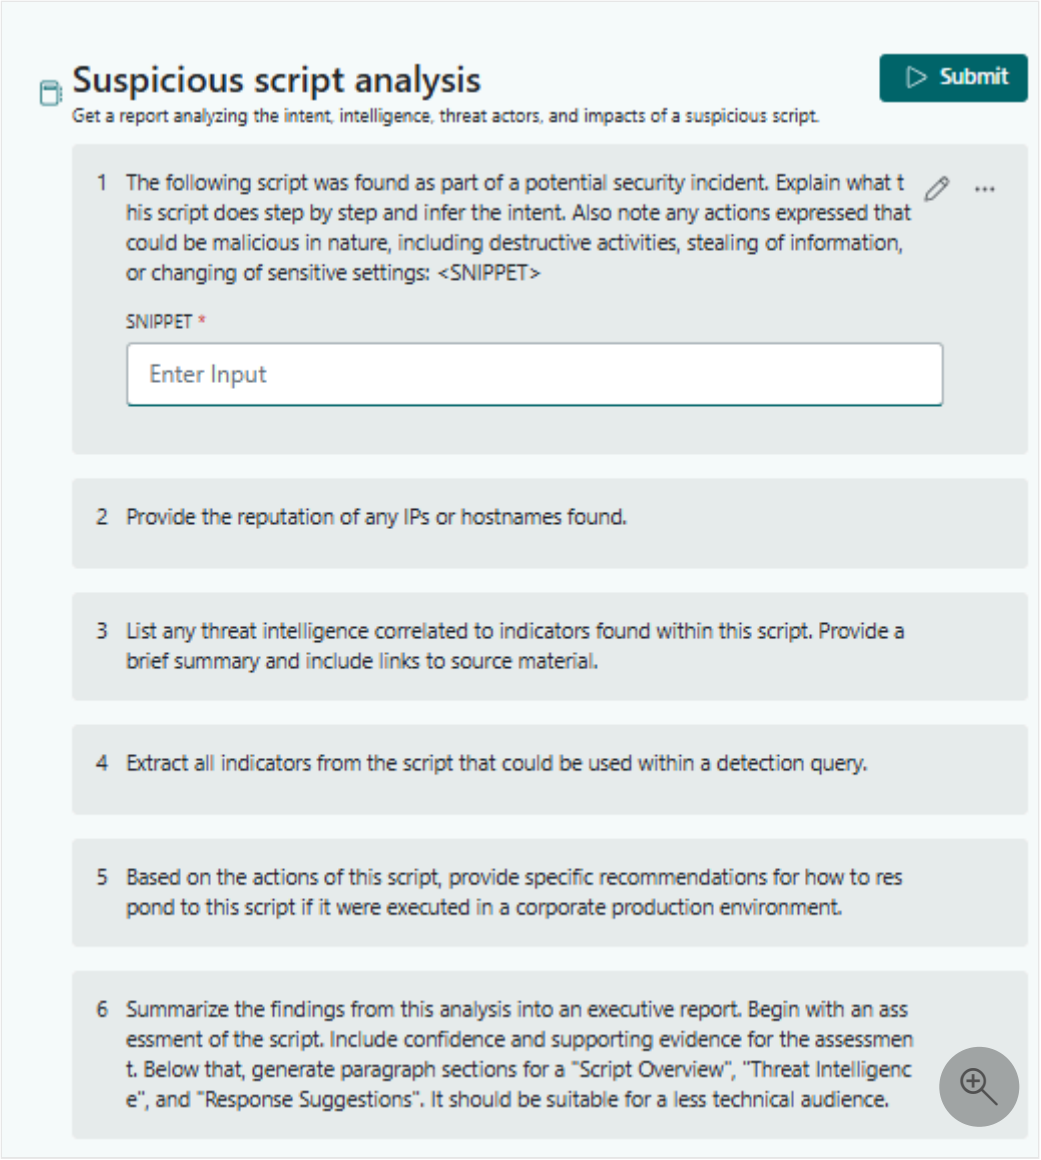

3. Next, select **Submit**.

4. Wait for Security Copilot to run the script content through the different prompts. If you see a round progress indicator in place of the response, then the promptbook is still running. Security Copilot generates responses for each of the prompts, building on each response until it gets to the last prompt.

5. Read the responses by Security Copilot. The last prompt by Security Copilot generates a full report of what the script does, any related threat activities, and recommended next steps based on the assessment about file intent. Review and verify whether the responses are accurate and meet your needs.

Threat actor profile

The threat actor profile promptbook is a quick way to get an executive summary about a specific threat actor. The promptbook looks for any existing threat intelligence articles about the actor,

including known tools, tactics, and procedures (TTPs) and indicators, including remediation suggestions. It then summarizes the findings into a report for less technical readers.

To run the **Threat actor profile** promptbook:

1. Select the Prompts button in the prompt bar and start to type "threat actor profile" until the promptbooks appear in the list.
2. Select **Threat actor profile**.
3. Type the name of the threat actor in the input box that says *Threat actor name*.

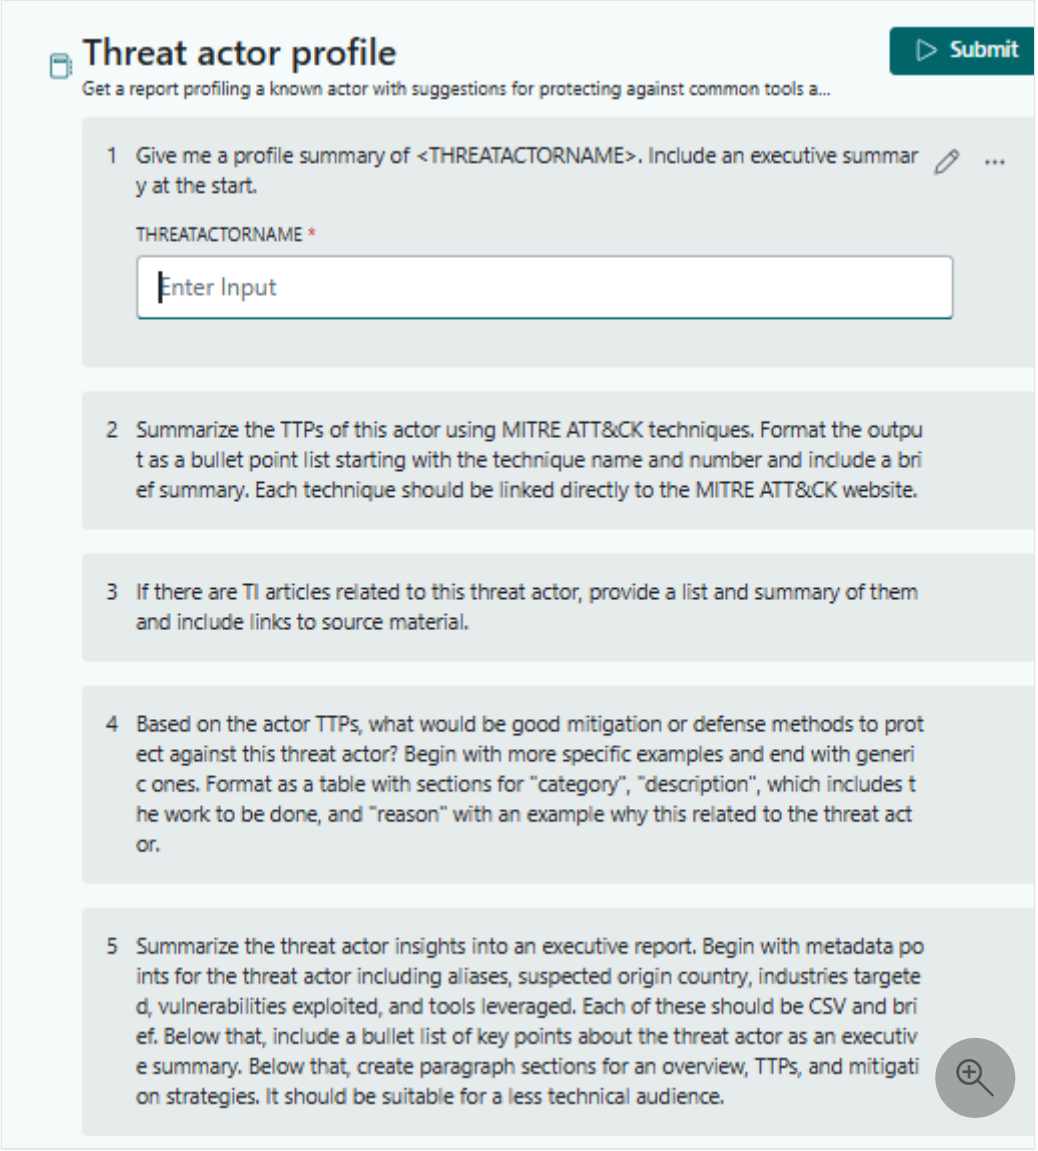

4. Next, select **Submit**.
5. Wait forSecurity Copilot to run the threat actor name through the different prompts. If you see a round progress indicator in place of the response, then the promptbook is still running. Security Copilot generates responses for each of the prompts and builds on each until it gets to the last prompt.

6. Read the response by Security Copilot. The last prompt by Security Copilot generates an easily readable report that includes pertinent information about the identified threat actor. Review and verify whether the responses are accurate and meet your needs.

Threat Intelligence 360 report based on MDTI article

By using this promptbook, you can get a detailed report about whether the threats discussed in a given Microsoft Defender Threat Intelligence article are affecting the organization, including relevant indicators and hunting queries.

To run the **Microsoft Threat Intelligence** promptbook:

1. Make sure that you have turned on the **Microsoft Threat Intelligence** plugin by following the steps in Enable the Security Copilot integration in Defender TI.
2. In Security Copilot, select the Prompts button in the prompt bar and start to type the promptbook name until the promptbook appears in the list.
3. Select the promptbook named **Threat Intelligence 360 report based on MDTI article**.
4. Type the name of the Defender Threat Intelligence article in the input box that says *MDTI Article Name*.

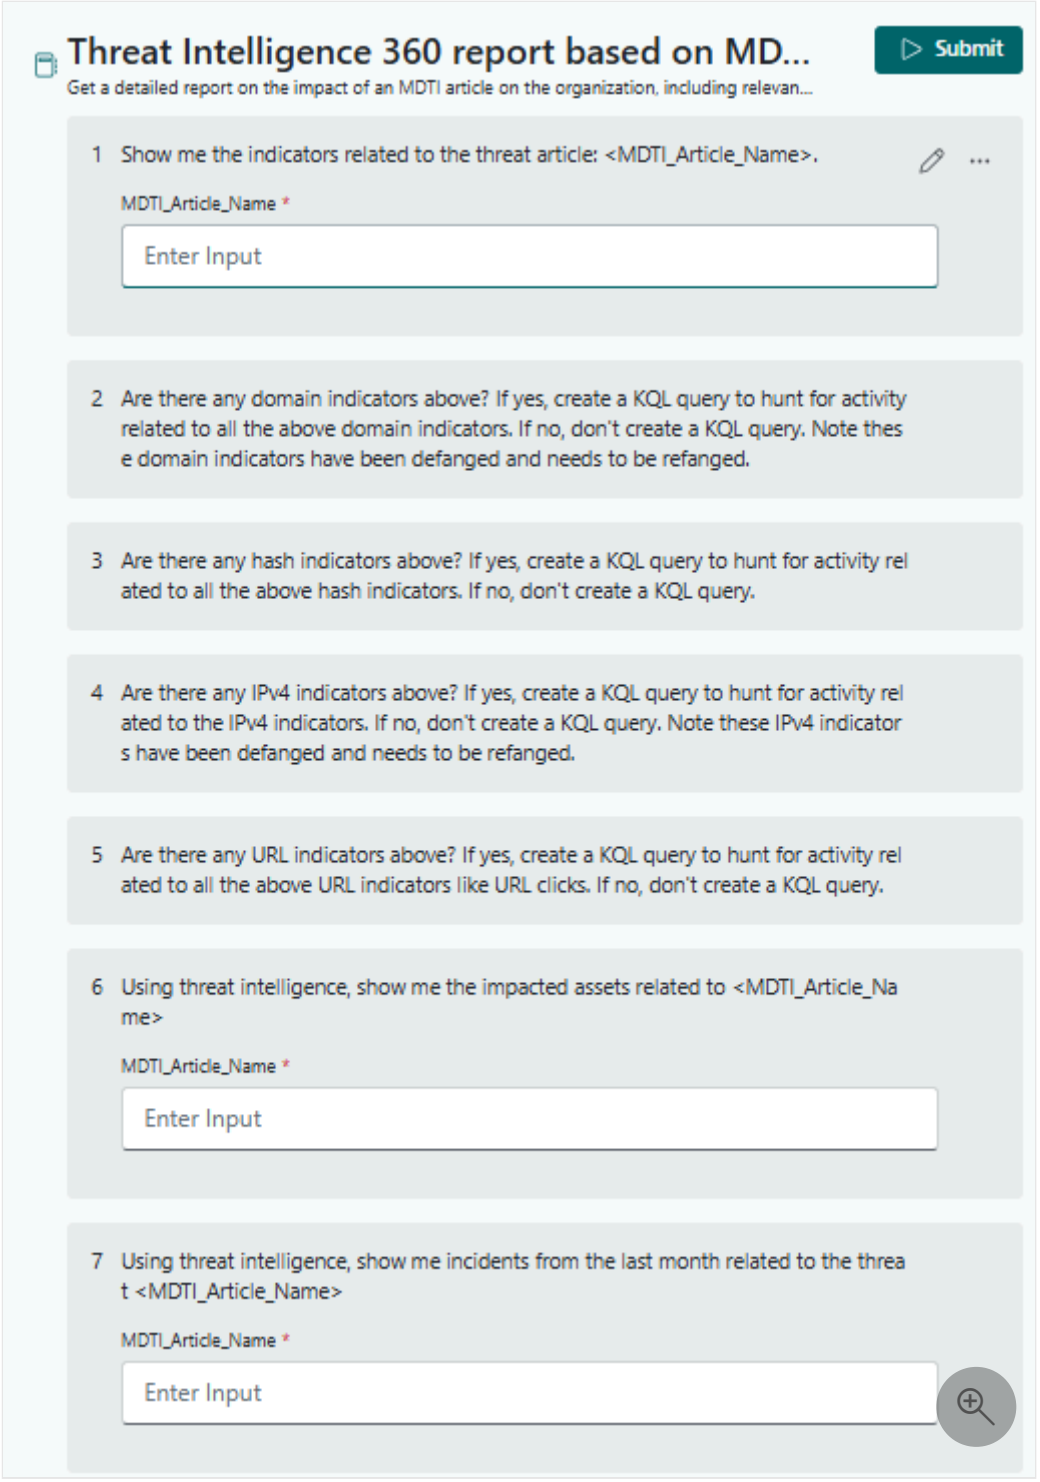

5. Next, select **Submit**.

6. Wait for Security Copilot to run the article through the different prompts. If you see a round progress indicator in place of the response, then the promptbook is still running. Security Copilot generates responses for each of the prompts and builds on each until it gets to the last prompt.

7. Read the response by Security Copilot.

Vulnerability impact assessment

The vulnerability impact assessment promptbook accepts a CVE number or known vulnerability name to find out if the vulnerability has been publicly disclosed or exploited and whether it has been used by threat actors in their campaigns. It can then provide recommendations to address or mitigate the threat and summarize these findings in an executive summary.

To run the **Vulnerability impact assessment** promptbook:

1. Select the Prompts button in the prompt bar and start to type "vulnerability impact assessment" until the promptbooks appear in the list.
2. Select **Vulnerability impact assessment**.
3. Type the CVE number or common vulnerability name that you want to learn about in the input box saying *CVEID*.

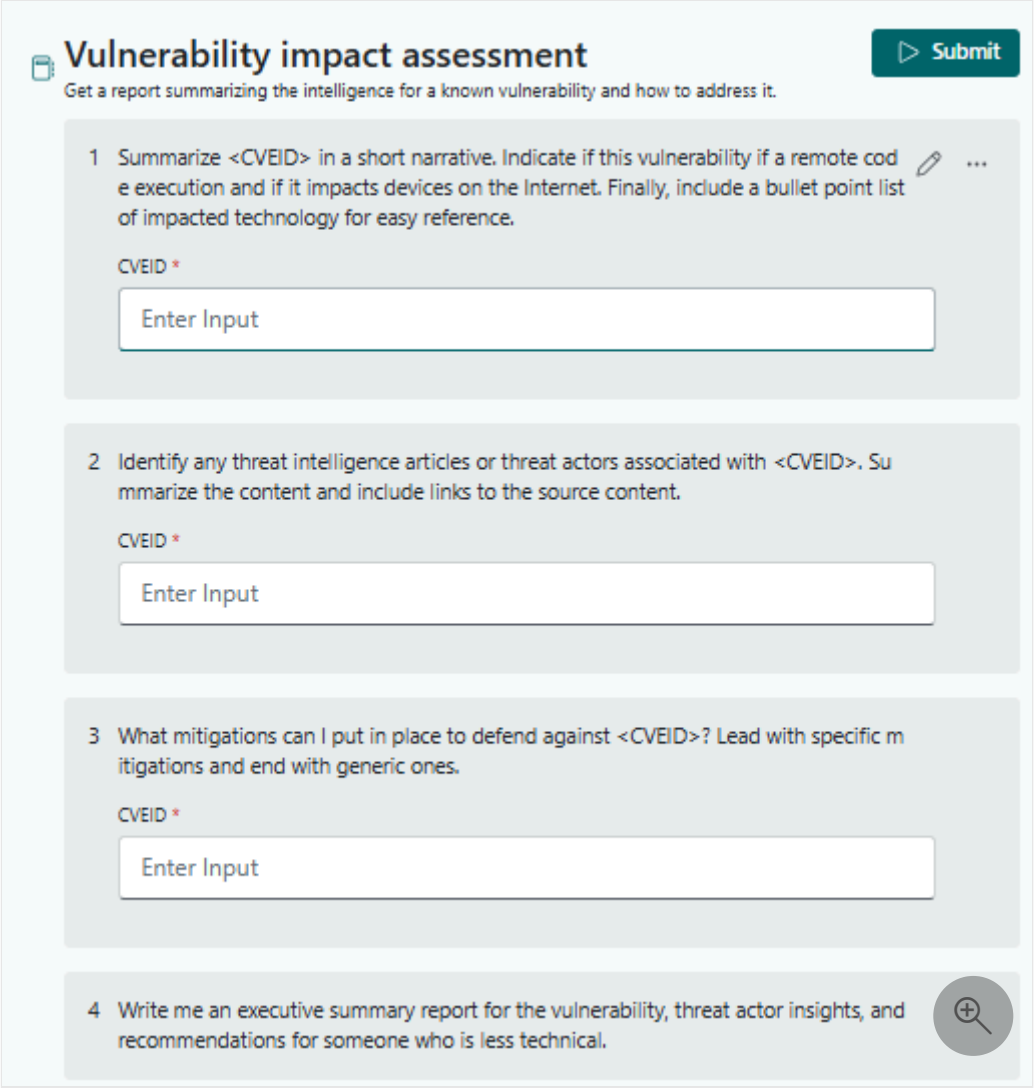

4. Next, select **Submit**.
5. Wait for Security Copilot to run the vulnerability name or CVE through the different prompts. If you see a round progress indicator in place of the response, then the

promptbook is still running. Security Copilot generates responses for each of the prompts and builds on each until it gets to the last prompt.

6. Read the response from Security Copilot. The last prompt generates an easily readable report about the vulnerability. The report includes details about known exploitation activities, including mitigation suggestions. Review and verify whether the responses are accurate and meet your needs.

Related content

Use the following Security Copilot articles for guidance:

- Build your own promptbooks
- Create effective prompts
- Security Copilot use cases for security and IT roles

Last updated on 06/01/2026

Security Copilot use cases for security and IT roles

10/24/2025

Security Copilot empowers users to accomplish the following key use cases as they apply to the different personas or roles in a security operations center or IT team.

Investigate and remediate security threats

Gain context for incidents to quickly triage complex security alerts into actionable summaries and remediate quicker with step-by-step response guidance.

- **SOC** - Receive actionable step-by-step guidance for incident response, including directions for triage, investigation, containment, and remediation.
- **Identity admins** - Streamline incident resolution by quickly summarizing critical information like user roles, sign-logs, and risk factors to help analysts understand the scope and specifics of potential compromises.
- **CISO** - Get up-to-date, summarized threat intelligence from Microsoft and open source, providing contextual insights on relevant exposures and threat actors, tooling, and techniques.

Read the following related resources:

- Use case: Incident response and remediation
- Use case: Triage incidents based on enrichment from threat intelligence
- Incident investigation promptbook
- Triage and investigate incidents with guided responses from Microsoft Copilot in Microsoft Defender

Build KQL queries or analyze suspicious scripts

Eliminate the need to manually write query-language scripts or reverse-engineer malware scripts with natural language translation to allow every team member to perform technical tasks.

- **TI analysts** - Build Kusto Query Language (KQL) queries to hunt threats across your organization more quickly and easily.
- **Data security admins** - Simplify eDiscovery investigation by translating inquiries from natural to keyword query language (KeyQL). Prompt in natural language to make your

search iterations faster and more accurate. Strengthen team expertise and enable more assertive evidence search.
- **IT admins** - Construct and run KQL queries to get device details from single and multiple devices.
- **Cloud security admins** - Fix issues in Infrastructure-as-Code (IaC) by submitting pull requests to developers with step-by-step remediation guidance and necessary code from Copilot.

Read the following related resources:

- Suspicious script analysis promptbook
- Use case: Investigate an incident and associated suspicious entities
- Generate KQL queries with Security Copilot in Microsoft Defender

Understand risks and manage security posture of the organization

Get a broad picture of your environment with prioritized risks to uncover opportunities to improve posture more easily.

- **Cloud security admins** - Understand your comprehensive multicloud risks and receive actionable step-by-step guidance for risk remediation, including AI-generated scripts. Facilitate risk remediation across teams with delegation and pull request creation capabilities.
- **IT admins** - Get a comprehensive view of your environment with prioritized risks and AI-generated summaries to identify overlapping settings, prevent policy conflicts, and minimize vulnerabilities during policy creation or updates.
- **Data security admins** - Assess and manage organization's data security posture by reviewing and tackling prioritized data security risks with a centralized data security dashboard to reduce investigation time.
- **TI analysts** - Get summarized threat intelligence relevant to an artifact to quickly contextualize an incident and extract MITRE techniques, tactics, and procedures (TTPs) to understand associated threat activity.

Read the following related resources:

- Check impact of an external threat article promptbook
- Threat Intelligence 360 report based on an MDTI article promptbook

Troubleshoot IT issues faster

Synthesize relevant information rapidly and receive actionable insights to identify and resolve IT issues quickly.

- **IT admins** - Reduce mean time from discovery to response of IT incidents by analyzing error codes and summarizing device context.
- **Identity admins** - Simplify sign-in log troubleshooting by automating data gathering and correlation into summaries of relevant information (for example, failed MFA attempts and policy changes). Get recommendations to resolve access issues efficiently, such as re-registering devices or adjusting policies.

Define and manage security policies

Define a new policy, cross-reference it with others for conflicts, and summarize existing policies to manage complex organizational context quickly and easily.

- **IT admins** - Reduce the risk of operational disruptions and vulnerabilities by analyzing conflicting or misconfigured policies and get guidance for recommended policy settings.
- **Data security admins** - Streamline DLP policy tuning and policy planning of data landscape to improve coverage and controls.

Configure secure lifecycle workflows

Build groups and set access parameters with step-by-step guidance to ensure a seamless configuration to prevent security vulnerabilities.

- **Identity admins** - Quickly configure workflows based on user roles, groups, and access package parameters, providing step-by-step guidance to ensure comprehensive and accurate setup.

Develop reports for stakeholders

Get a clear and concise report that summarizes the context and environment, open issues, and protective measures prepared for the tone and language of the report's audience.

- **SOC** - Summarize investigations in Copilot into exportable, natural language reports to simplify communication to security stakeholders.
- **CISO** - Get a clear and concise report that covers threats, threat actors, analysts' work, and protective measures, tailored for the board of directors in their native languages.

Read the following related resources:

- Check impact of an external threat article promptbook

- Incident investigation promptbook
- Threat Intelligence 360 report based on an MDTI article promptbook
- Summarize an incident with Microsoft Copilot in Microsoft Defender

Use case: Triage incidents based on enrichment from threat intelligence

The article covers how an analyst uses Security Copilot to quickly triage an incident using sample prompts.

Roles mentioned: SOC analyst and identity admin for incident triage (main), CISO and other stakeholders for incident report summary (recipient)

Incident triage scenario

As a security operations center (SOC) analyst, you review alerts and incidents assigned to you. Your duty is to identify if real action needs to be taken. You leverage information within the alerts associated with the incident to guide your process. You often collect contextual information to further understand the next steps you should take. With the enrichment of involved entities and full understanding of the underlying alerts, you determine whether to escalate or remediate the incident.

In this detailed example, an analyst uses Security Copilot to quickly triage an incident. If the incident is a real threat, the goal is to either gather new indicators of compromise or link entities to finished intelligence. In this case, the threat intelligence is summarized by Security Copilot to show the connection to a known threat actor and inform the severity assessment.

Steps to triage incidents using threat intelligence

1. **Start your day with Security Copilot. Retrieve the latest Microsoft Defender XDR incident that's assigned to you and summarize the alerts associated with it.**

   Prompt used:

   What is the latest active Defender incident assigned to me, `angus.macgregor@contoso.com`*? Summarize it, including the alerts associated with it.*

   Response:

What is the latest active Defender incident assigned to me, angus.macgregor@contoso.com? Summarize it, including the alerts associated with it.

3 steps completed 36s

Chose M365 9s

Processed your message 6s

Prepared your response. 22s

Oct 6, 7:16 AM | 38.41s

The latest active Defender incident assigned to you, angus.macgregor@contoso.com, is incident **18727** with the display name "Possible attempt to access Primary Refresh Token (PRT) on one endpoint". It was created on 2023-10-05 and last updated on 2023-10-06. The incident has a medium severity and is classified as unknown.

There is one alert associated with this incident:

- Alert ID: cccccccc-dddd-eeee-3333-444444444444_1
- Status: New
- Severity: Medium
- Service Source: Microsoft Defender for Endpoint
- Detection Source: Microsoft Defender for Endpoint
- Title: Possible attempt to access Primary Refresh Token (PRT)
- MITRE Techniques: T1003, T1528
- Created: 2023-10-05T21:39:00.3084336Z
- Last Updated: 2023-10-06T06:03:05.2866667Z

The alert description states that a possible attempt to access Primary Refresh Token (PRT) has been detected. A Primary Refresh Token (PRT) is a key artifact of Azure AD authentication on Windows 10, Windows Server 2016 and later versions, iOS, and Android devices. It is a JSON Web Token (JWT) that's specially issued to Microsoft first party token brokers to enable single sign-on (SSO) across the applications used on those devices. Attackers can attempt to access this resource to move laterally into the organization or perform credential theft.

Recommended actions for this alert include validating the alert, scoping the incident, containing and mitigating the breach, and contacting your incident response team or Microsoft support for investigation and remediation services.

Looks like possible credential theft. You follow the recommended actions and begin to scope the incident and validate the alert.

2. **Focus in on specific entities to get more information about them.**

   Prompt used:

   Elaborate on the details of this Defender incident alert, including the entities involved.

   Response:

Elaborate on the details of this alert including the entities involved.

3 steps completed 25s

Chose Fusion 6s

Processed your message 2s

Prepared your response. 17s

Oct 6, 7:18 AM | 26.10s

The alert with ID 11112222-bbbb-3333-cccc-4444dddd555555_1 involves the following entities:

- Device:
  - Azure AD Device ID: cccccccc-2222-3333-4444-dddddddddddd
  - Device DNS Name: cpc-avance-uyoa0
  - MDE Device ID: H4iJ5kL6mN7oP8qR9sT0uV1wX2yZ3a
  - OS Platform: Windows 11
  - OS Build: 22621
  - Version: 22H2
  - Health Status: Inactive
  - Onboarding Status: Onboarded
  - RBAC Group ID: 37225
  - RBAC Group Name: Full Remediation - Cloud PCs
- User Account:
  - AdeleVance
  - Azure AD User ID: ffff5f5f-aa6a-bb7b-cc8c-dddddd9d9d9d
  - Domain Name: AzureAD
  - User Principal Name: adele.vance@contoso.com
  - User SID: S-1-12-1-3834083884-1079035900-3132596361-2383408245
- IP Address: 0.0.0.0

The alert was detected on the device with DNS name cpc-avance-uyoa0, which is running Windows 11. The user account involved in the alert is AdeleVance (adele.vance@contoso.com). The device has an inactive health status and is part of the "Full Remediation - Cloud PCs" RBAC group. The IP address associated with the alert is 0.0.0.0.

Now you have a user account and device to further investigate. In this case, you choose to understand more about the affected user before digging into the details of the attack on the device.

3. **Get more information about this user to guide next steps. What type of actions might be next for someone with her credentials?**

   Prompt used:

   Tell me more about the user entity.

   Response:

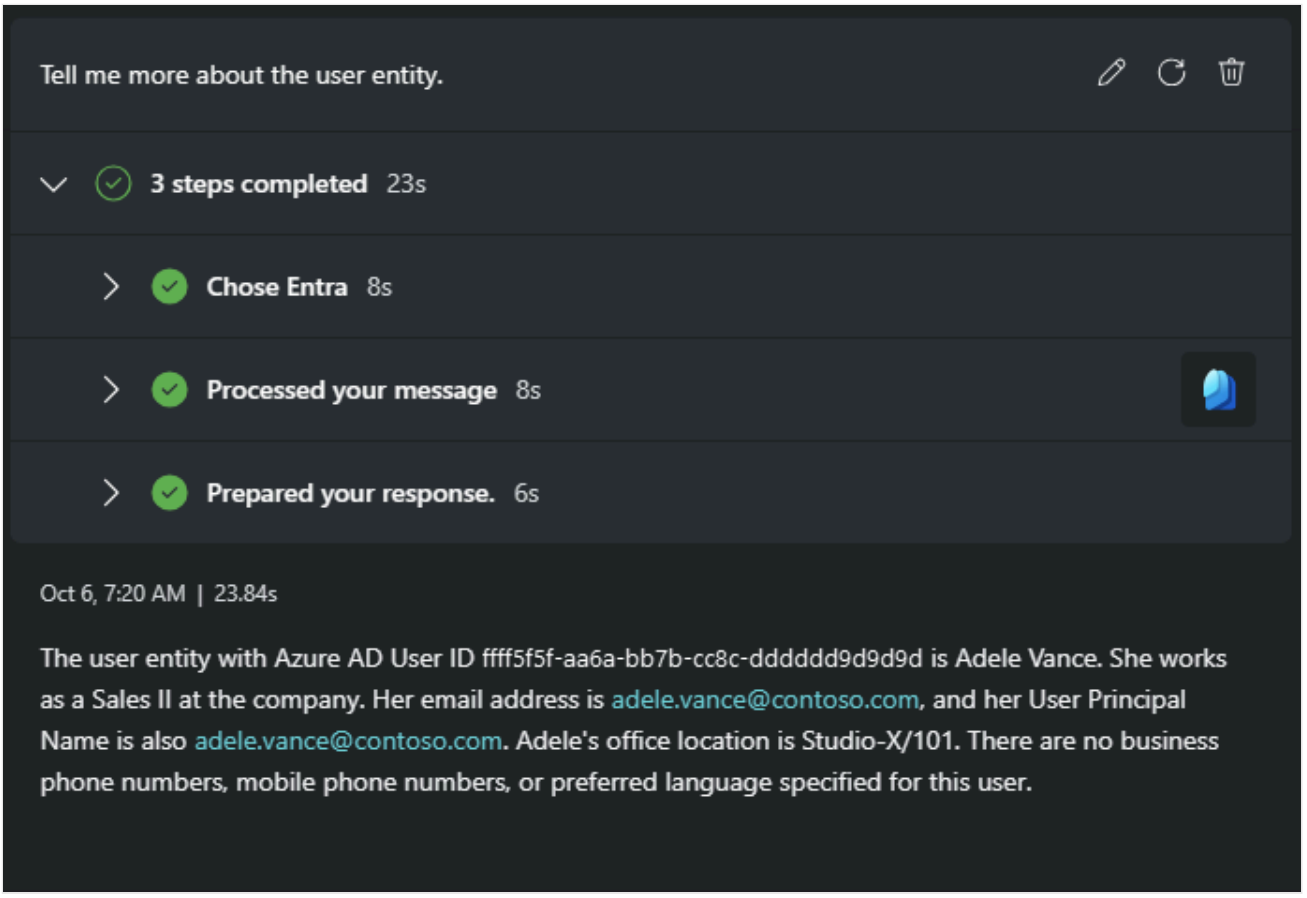

   You find out this user is works in Sales. If her credential was stolen, this could affect sales data. You remember that your Sentinel workspace has an SAP solution to help detect threats there. Is this Defender alert linked to a Microsoft Sentinel incident? Your first priority is to determine if there has been any suspicious activity by this user in SAP.

4. **Use a saved hunting query to correlate entities with Sentinel incidents.**

   You manually activate the suggested prompt for the **Natural language to Sentinel KQL** plugin to run your query.

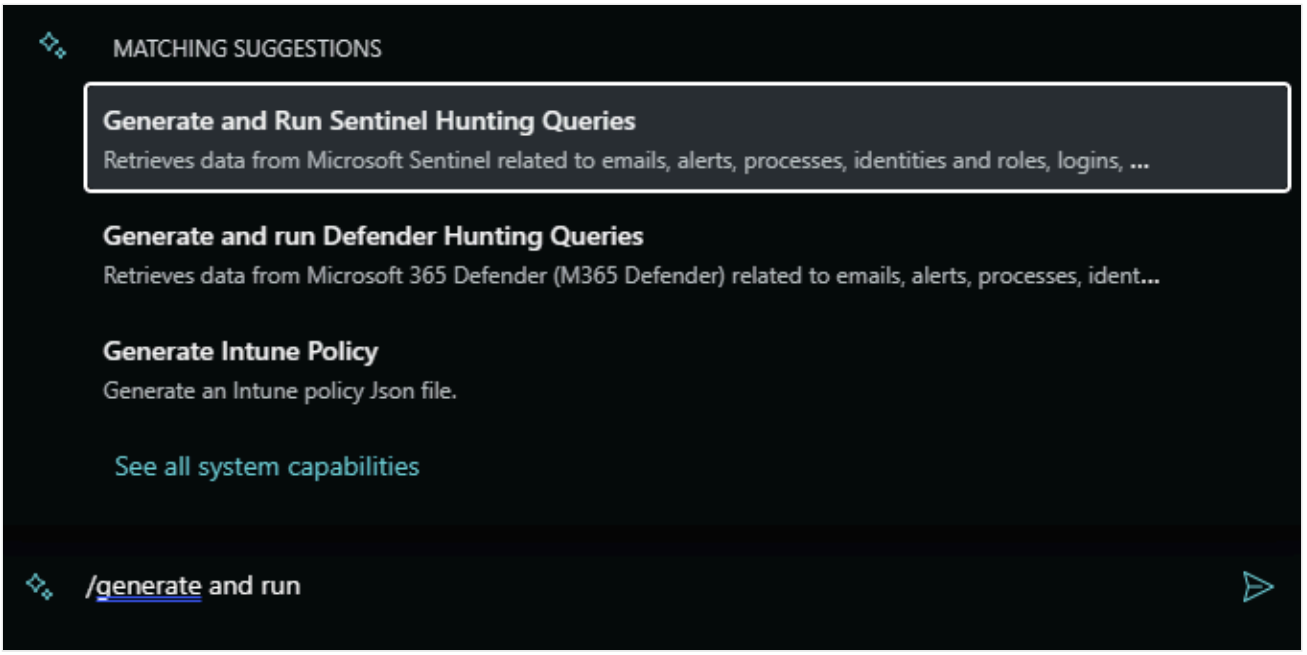

> **Tip**
>
> If the query needs to be adjusted slightly, edit the prompt and re-run it. For instance, your query projected `IncidentNames`, but those are simply GUIDs. You remember it's the `Title` field you really want. Simply edit the prompt and select the **Re-run prompt** option.

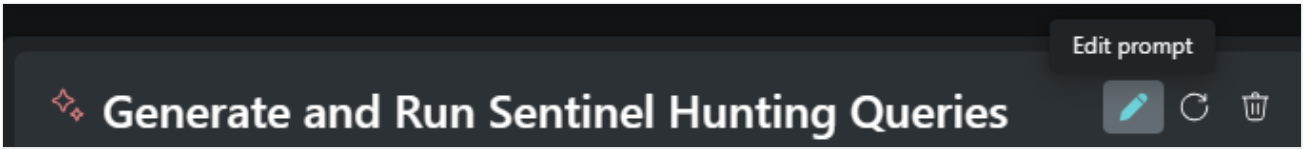

Adjusted prompt used:

Run the following KQL

```
    *SecurityAlert | where Entities has "`adele.vance@contoso.com`" and
TimeGenerated >= datetime(10/06/2023) | join kind=inner (SecurityIncident | mv-
expand SystemAlertId = AlertIds | extend SystemAlertId = tostring(SystemAlertId))
on SystemAlertId | summarize by IncidentNumber, Title*
```

Response:

Generate and Run Sentinel Hunting Queries

DataRequest:Run the following KQL SecurityAlert | where Entities has "adele.vance@contoso.com" and TimeGenerated >= datetime(10/06/2023) | join kind=inner (SecurityIncident | mv-expand SystemAlertId = AlertIds | extend SystemAlertId = tostring(SystemAlertId)) on SystemAlertId | summarize by IncidentNumber, Title

3 steps completed 38s

Chose NL2KQLSentinel 0s

Processed your message 27s

Prepared your response. 10s

Oct 6, 7:25 AM | 39.50s

Here are the security incidents related to the email address "adele.vance@contoso.com" since 2023-10-06:

Export to Excel

IncidentNumber	Title
33089	Possible attempt to access Primary Refresh Token (PRT) on one endpoint
33090	Possible attempt to access Primary Refresh Token (PRT)
33805	SAP - (Preview) File Downloaded From a Malicious IP Address

A total of 3 incidents were found, and all 3 of them are displayed above.

The Kusto query used to generate this output is:

```
SecurityAlert
| where Entities has adele.vance@contoso.com and TimeGenerated >= datetime(2023-10-06)
| join kind=inner (
    SecurityIncident
    | mv-expand SystemAlertId = AlertIds
    | extend SystemAlertId = tostring(SystemAlertId)
) on SystemAlertId
| summarize by IncidentNumber, Title

```

The SAP related incident is now your first priority.

5. **Pivot your investigation to the SAP incident associated with the user from the original alert.**

   Prompt used:

   Elaborate on Sentinel incident 33805 and give me details about the entities.

   Response:

Elaborate on Sentinel incident 33805 and give me details about the entities.

3 steps completed 45s

Oct 6, 7:27 AM | 46.47s

Sentinel incident 33805 has the title "SAP - (Preview) File Downloaded From a Malicious IP Address". The entities involved in this incident are as follows:

1. IP Address:
    - Address: 185.82.217.3
2. User Account:
    - AdeleVance
    - Domain Name: AzureAD
    - User Principal Name: adele.vance@contoso.com
    - SID: S-1-12-1-3834083884-1079035900-3132596361-2383408245
    - Display Name: Adele Vance
    - Department: Sales
    - Job Title: Sales II
    - Office Location: Studio-X/101
3. Host:
    - Host Name: CPC-avance-S0C0S
    - NetBIOS Name: CPC-avance-S0C0S
    - OS Family: Windows
    - FQDN: cpc-avance-s0c0s
    - Mdatp Device ID: H4iJ5kL6mN7oP8qR9sT0uV1wX2yZ3a
    - Risk Score: Medium
    - Health Status: Active
    - Onboarding Status: Onboarded
    - Last Seen: 2023-06-10T09:53:05.1134431Z
    - Last External IP Address: 20.172.15.250
    - Last IP Address: 10.1.0.9

These are the entities associated with incident 33805. The incident is related to a file downloaded from a malicious IP address, involving the user Adele Vance and the host CPC-avance-S0C0S.

A lot of information is returned from this prompt. The malicious IP and possible exfiltration of financial data stand out as important items to investigate further.

6. **Find out more about the IP address entity and examine how it was determined to be malicious.**

   Prompt used:

   Give me more details about the IP address and why is it malicious?

   Response:

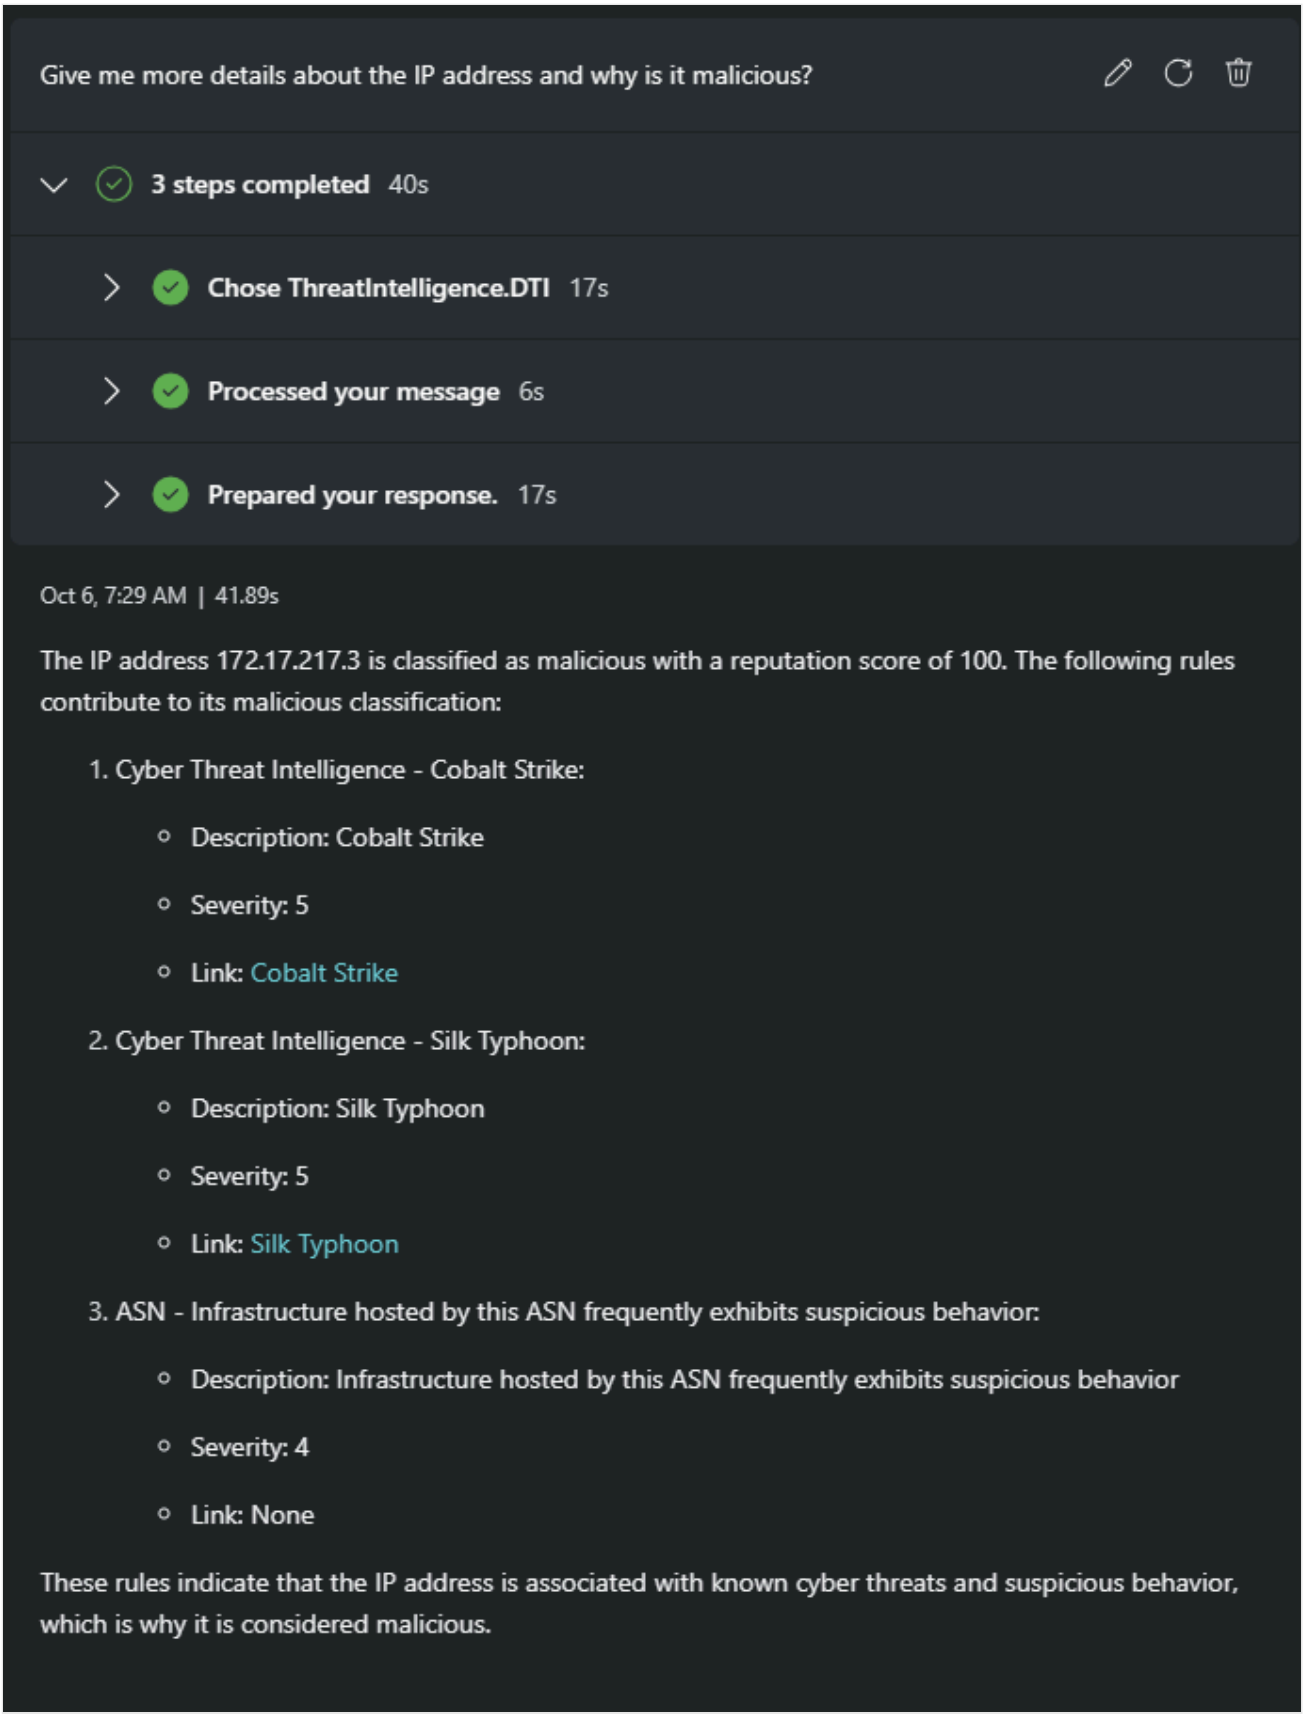

7. **Create a summary report**

   Save time in the escalation process with a summary for leadership and incident response teams.

   Prompt used:

Write a report based on this investigation. Lead with your assessment of the original Defender incident and whether the threat of credential theft is real. Conclude your assessment on how that threat relates to the Sentinel incident regarding the file downloaded to a malicious IP.

Response:

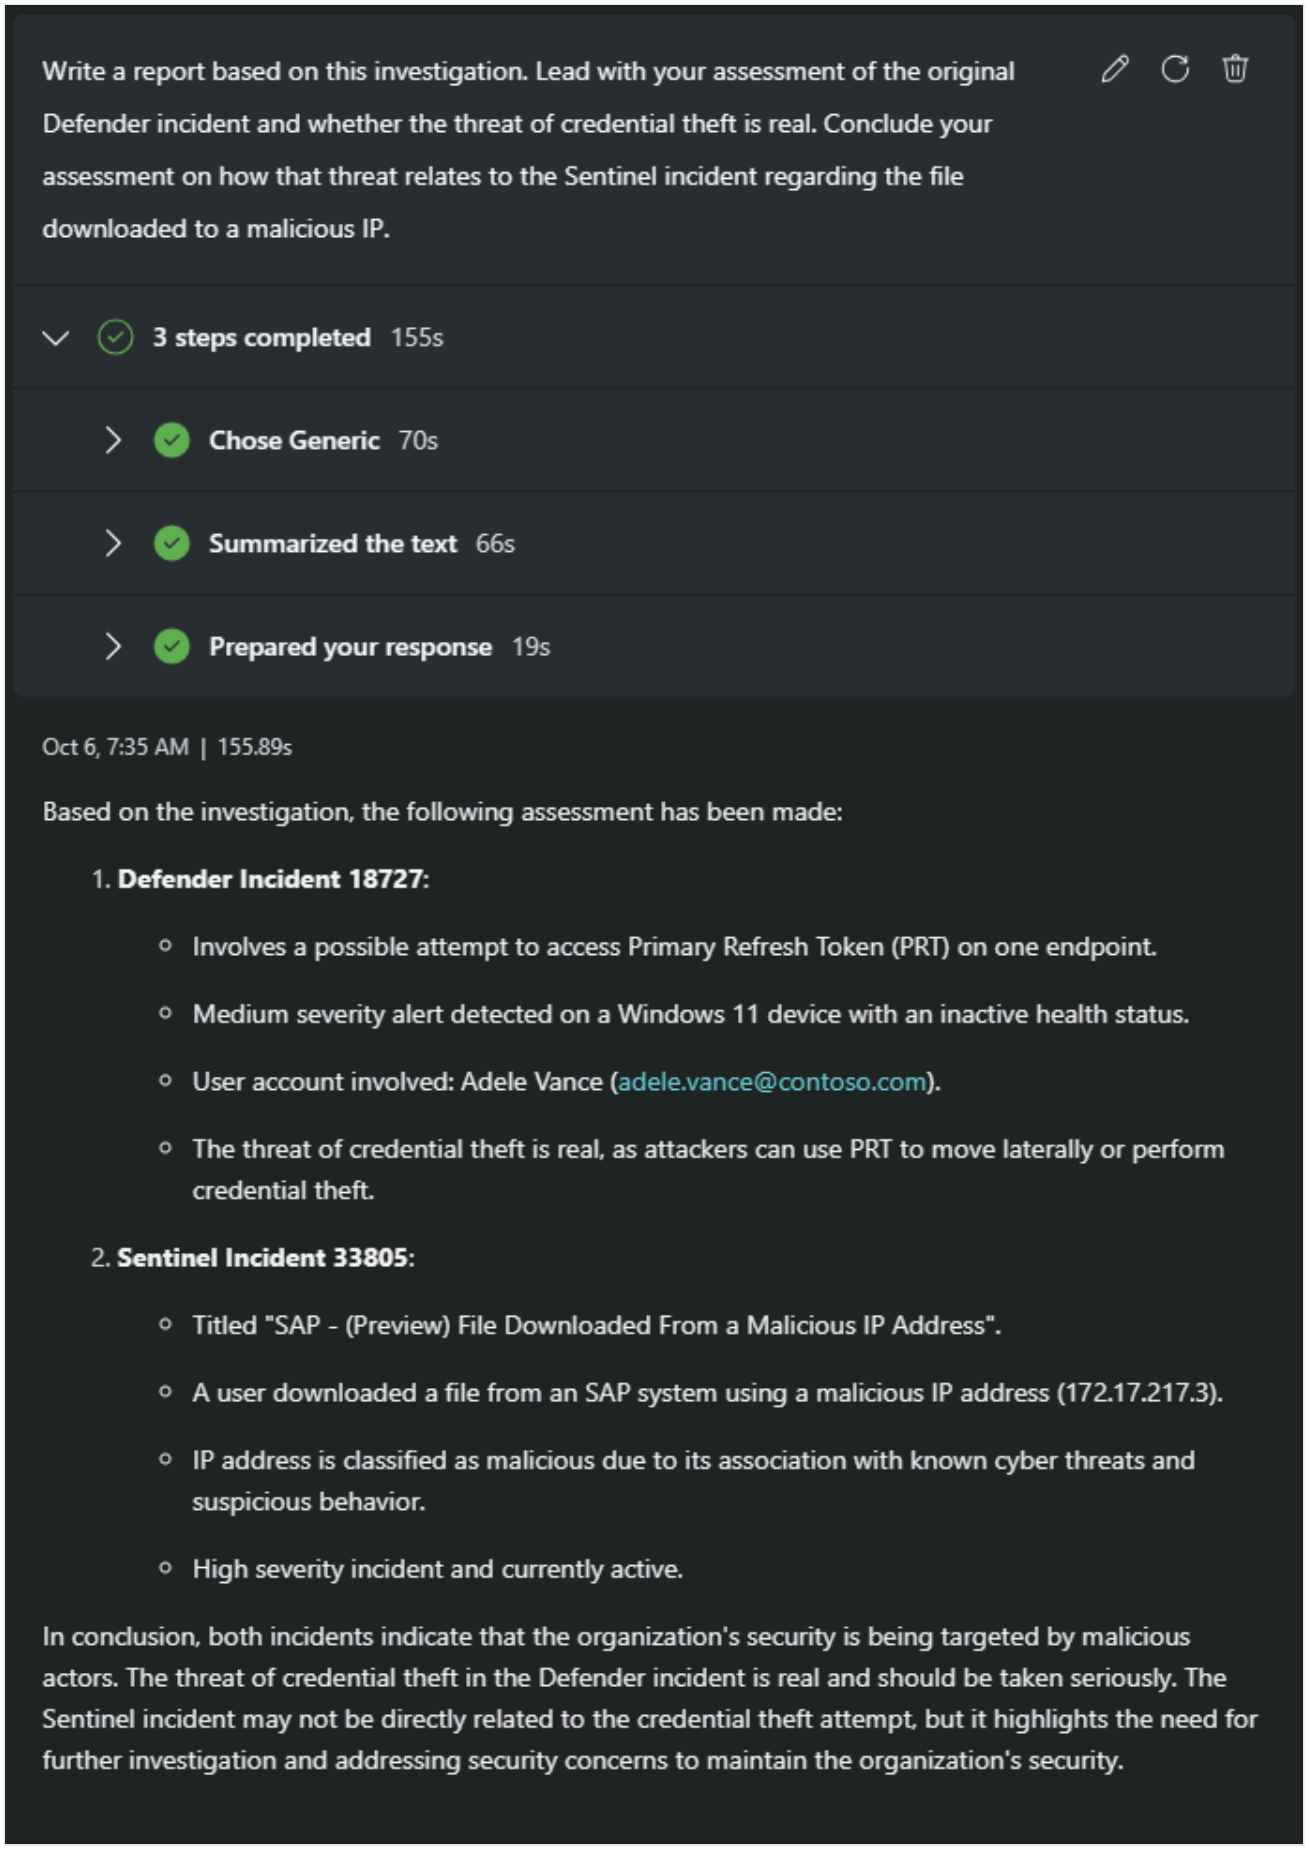

8. **Pin the most useful prompts responses and edit the session name.**

   You've reached your goal and determined the assigned Microsoft Defender XDR incident is a real threat. By linking it to a Microsoft Sentinel incident involving an exfiltrated SAP file,

you prepare to collaborate with your escalation team.

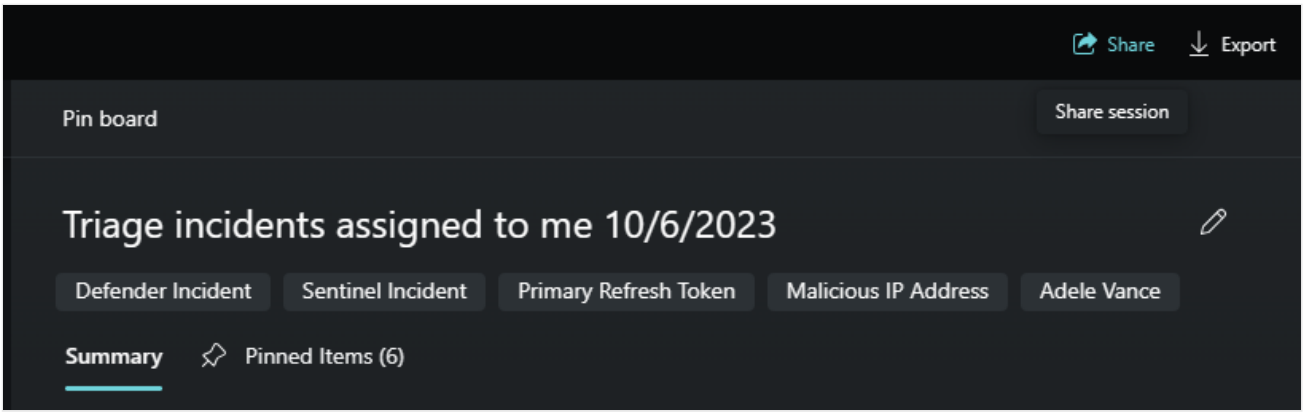

Conclusion

In this use case, Security Copilot helped triage an assigned incident quickly. You confirmed the alert required real action by investigating related incidents. The hunt resulted in finding an incident with an IP entity linked to finished intelligence about the threat actor and C2 tool used. With a succinct pin board, you shared the session and a summary report giving the escalation team the information they need to respond effectively.

Last updated on 06/04/2026

Use case: Investigate an incident and associated suspicious entities

Role mentioned: SOC TI analyst analyzing suspicious scripts

Incident investigation scenario

During an incident, security analysts are typically tasked with investigating alerts and gathering pertinent information that's associated with the incident. They conduct root cause analyses and correlate information from a range of sources to determine the potential impact to the organization.

Depending on the scenario, analysts may need to analyze logs, examine malware, reverse engineer files or scripts, and investigate URLs that were observed.

An essential component of an investigation involves understanding what remediation steps to take and effectively conveying significant discoveries to keep stakeholders informed about the current state of the incident.

In this example, Security Copilot is used to perform a comprehensive incident investigation by gathering contextual information from alerts, analyzing a suspicious script, and generating an assessment accompanied by a set of remediation steps.

Steps

1. Start investigating in Microsoft Defender XDR.

   Security Copilot is integrated in Microsoft Defender XDR. From an incident page, select the **Copilot** button to get a summary of an incident and get details such as the time and date of when an attack started, the entity or asset that started the attack, and assets involved in the attack.

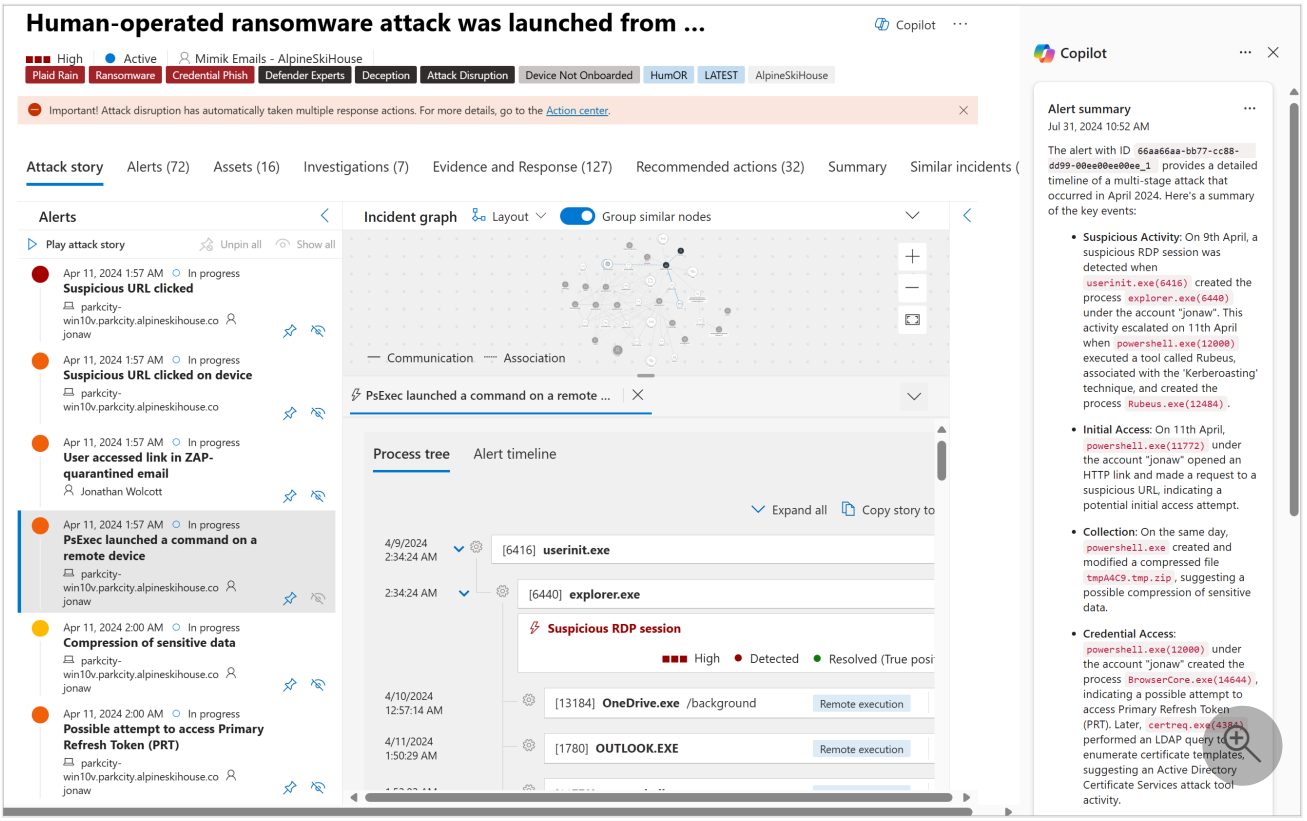

2. Analyze the suspicious script.

   Microsoft Defender XDR flags when a suspicious script executes. Use Security Copilot to explain what the suspicious script is doing.

   > **Note**
   >
   > Script analysis functions are continuously in development. Analysis of scripts in languages other than PowerShell, batch, and bash are being evaluated.

   With a click of a button, a description is displayed along with an overall summary of the script.

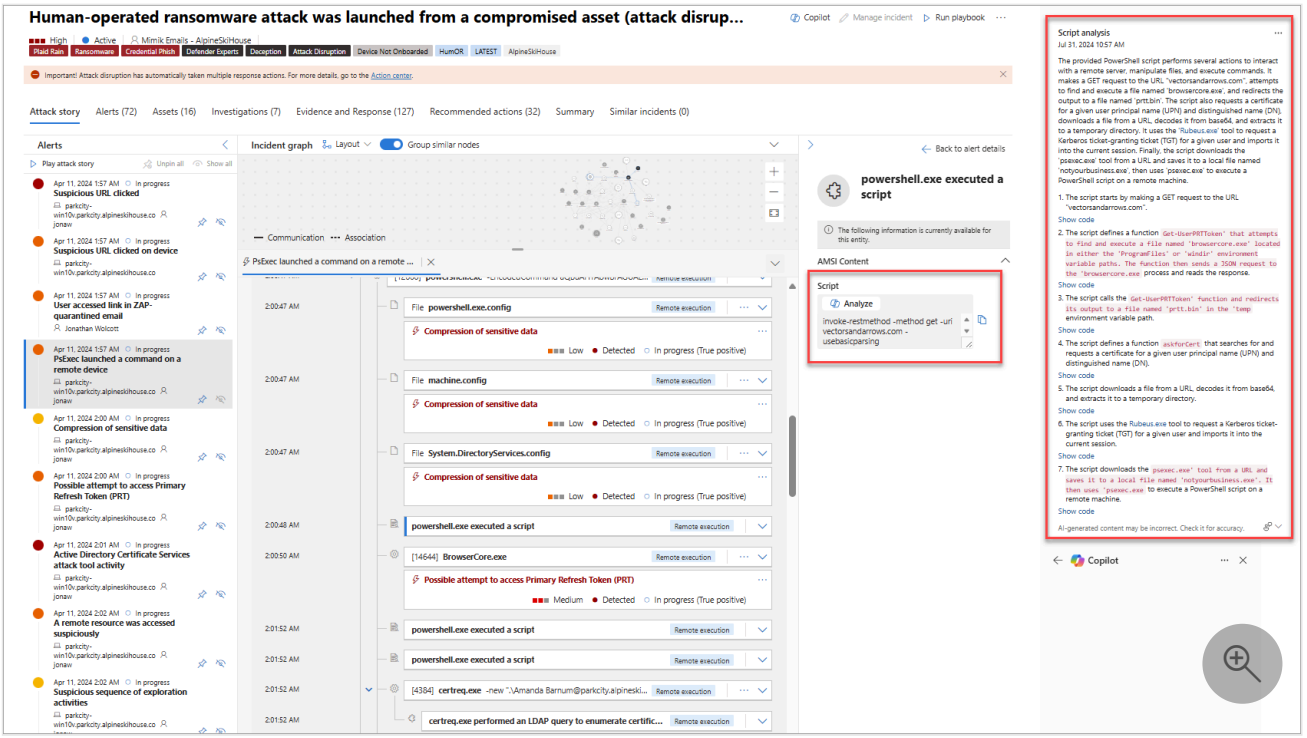

3. Extend the investigation in Security Copilot using natural language prompts and more plugins.

   Continue your investigation in the standalone experience of Security Copilot by selecting **Open in Security Copilot**.

The standalone experience allows you to extend the investigation by using natural language prompts.

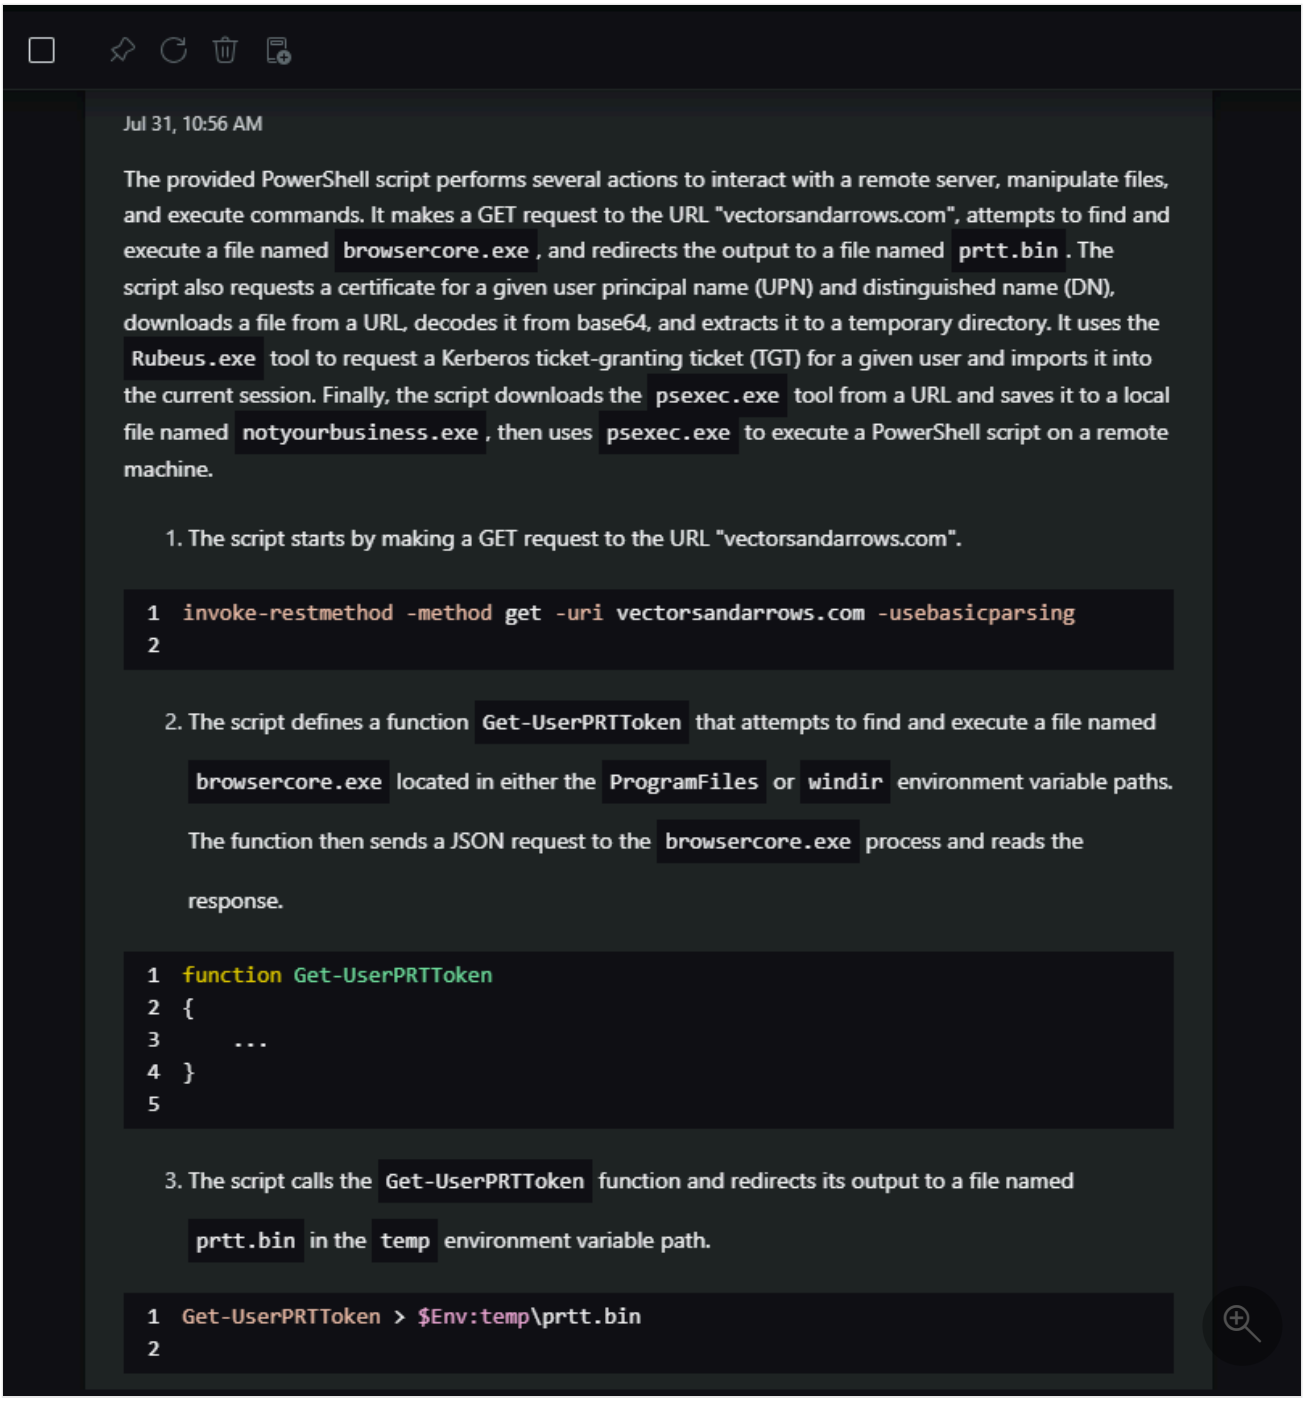

4. To gain a more comprehensive understanding of the incident, use Security Copilot to gather more information about the suspicious activity seen in the command line script.

   Prompt used:

   What can you tell me about the reputation of the indicators in the script? Are they malicious? If so, why?

   Response:

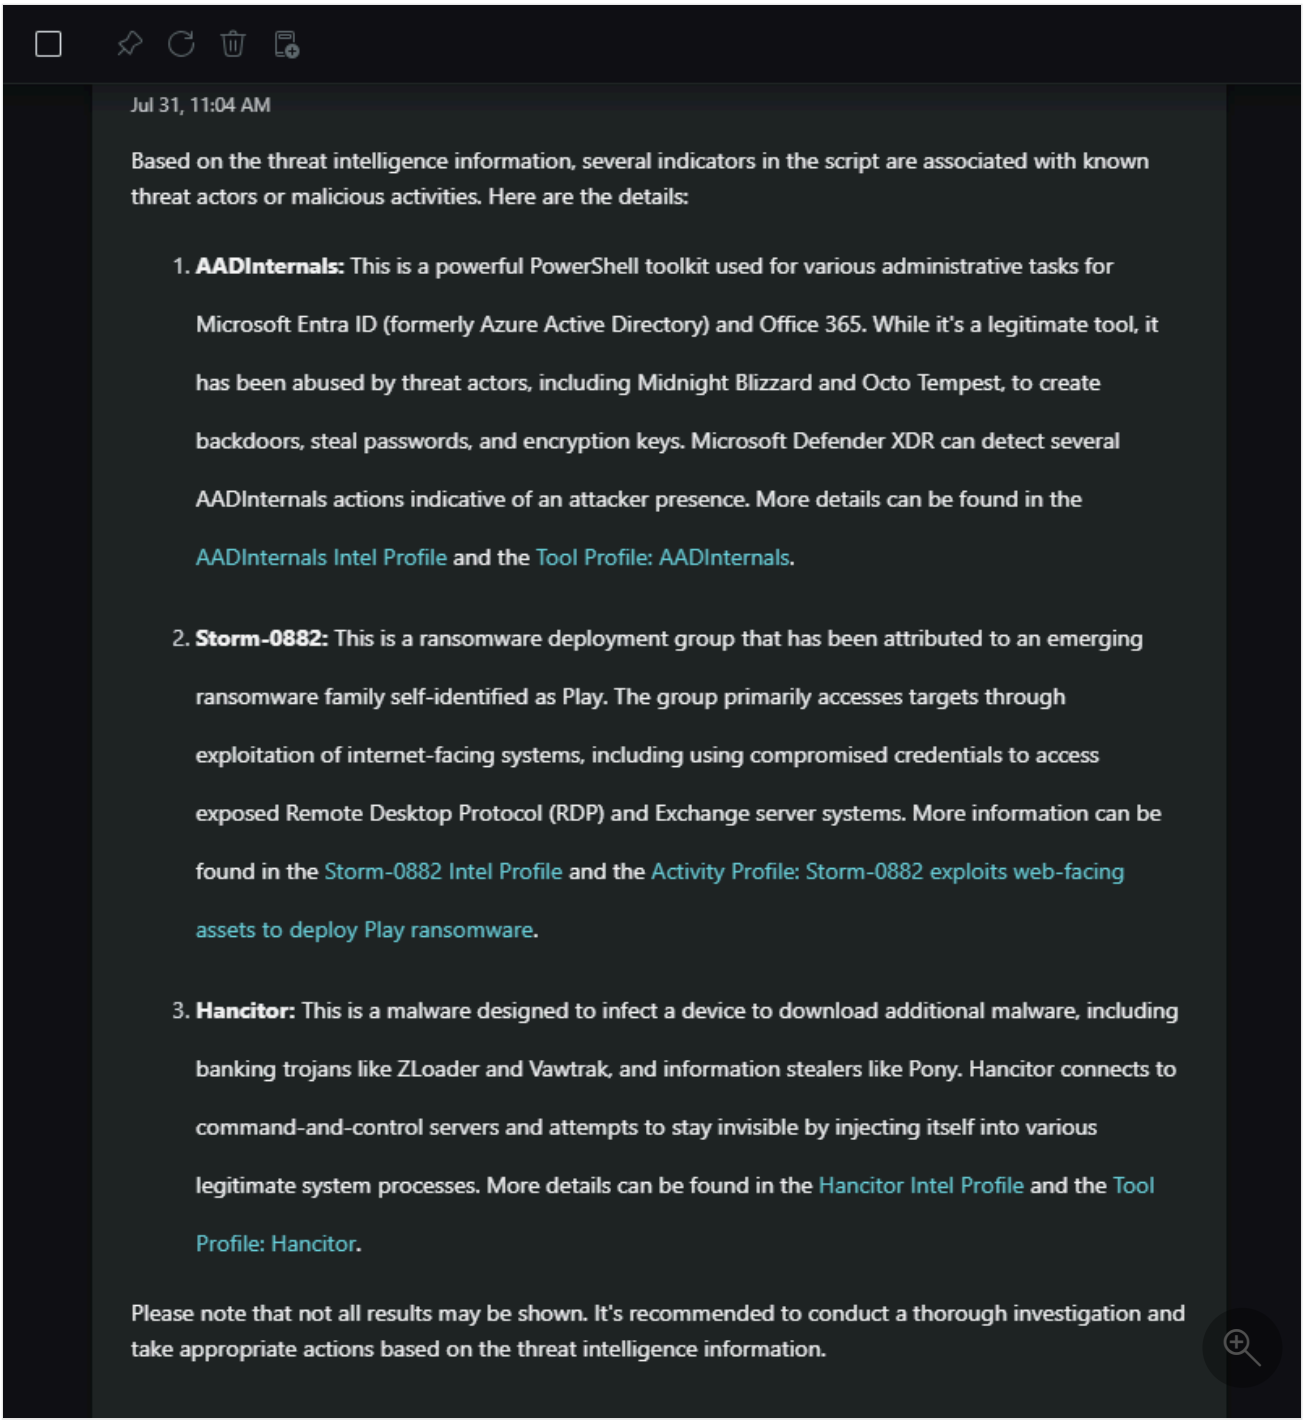

The response indicates that the several indicators in the script are associated with known threat actors. You can pin this response as a critical piece of information that can be used later.

5. Use Security Copilot to provide an assessment of the incident with supporting evidence and a set of recommendations.

   Prompt used:

   Summarize the findings from the investigation and conclude with a set of recommendations.

   Response:

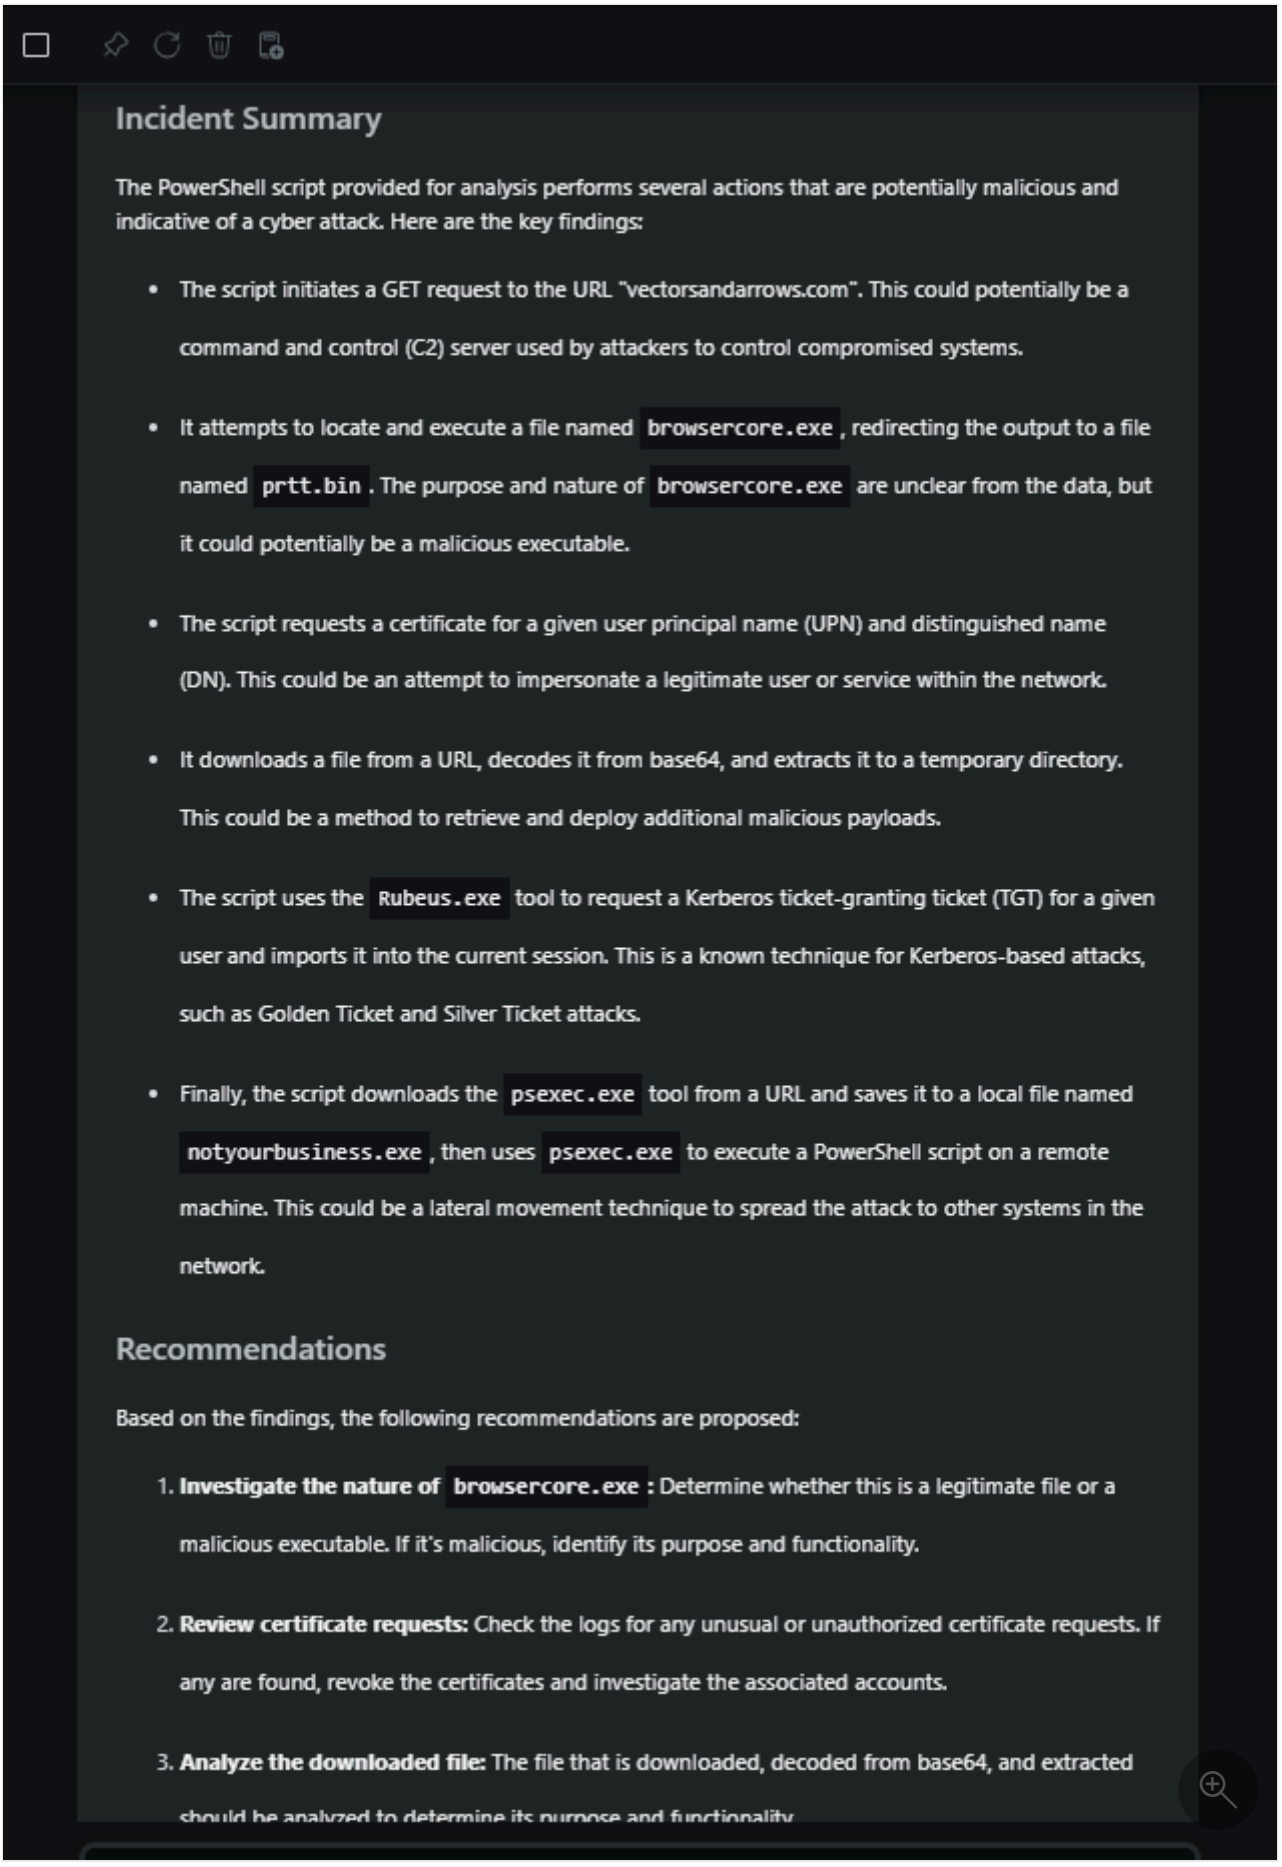

Incident Summary

The PowerShell script provided for analysis performs several actions that are potentially malicious and indicative of a cyber attack. Here are the key findings:

- The script initiates a GET request to the URL "vectorsandarrows.com". This could potentially be a command and control (C2) server used by attackers to control compromised systems.
- It attempts to locate and execute a file named `browsercore.exe`, redirecting the output to a file named `prtt.bin`. The purpose and nature of `browsercore.exe` are unclear from the data, but it could potentially be a malicious executable.
- The script requests a certificate for a given user principal name (UPN) and distinguished name (DN). This could be an attempt to impersonate a legitimate user or service within the network.
- It downloads a file from a URL, decodes it from base64, and extracts it to a temporary directory. This could be a method to retrieve and deploy additional malicious payloads.
- The script uses the `Rubeus.exe` tool to request a Kerberos ticket-granting ticket (TGT) for a given user and imports it into the current session. This is a known technique for Kerberos-based attacks, such as Golden Ticket and Silver Ticket attacks.
- Finally, the script downloads the `psexec.exe` tool from a URL and saves it to a local file named `notyourbusiness.exe`, then uses `psexec.exe` to execute a PowerShell script on a remote machine. This could be a lateral movement technique to spread the attack to other systems in the network.

Recommendations

Based on the findings, the following recommendations are proposed:

1. **Investigate the nature of `browsercore.exe`:** Determine whether this is a legitimate file or a malicious executable. If it's malicious, identify its purpose and functionality.
2. **Review certificate requests:** Check the logs for any unusual or unauthorized certificate requests. If any are found, revoke the certificates and investigate the associated accounts.
3. **Analyze the downloaded file:** The file that is downloaded, decoded from base64, and extracted should be analyzed to determine its purpose and functionality

Tip

You can export the response for future reference. You also have the option of sharing the entire session with other analysts. Other team members who are reviewing the incident can leverage the pin board to get a complete summary of the investigation steps, which saves them valuable time.

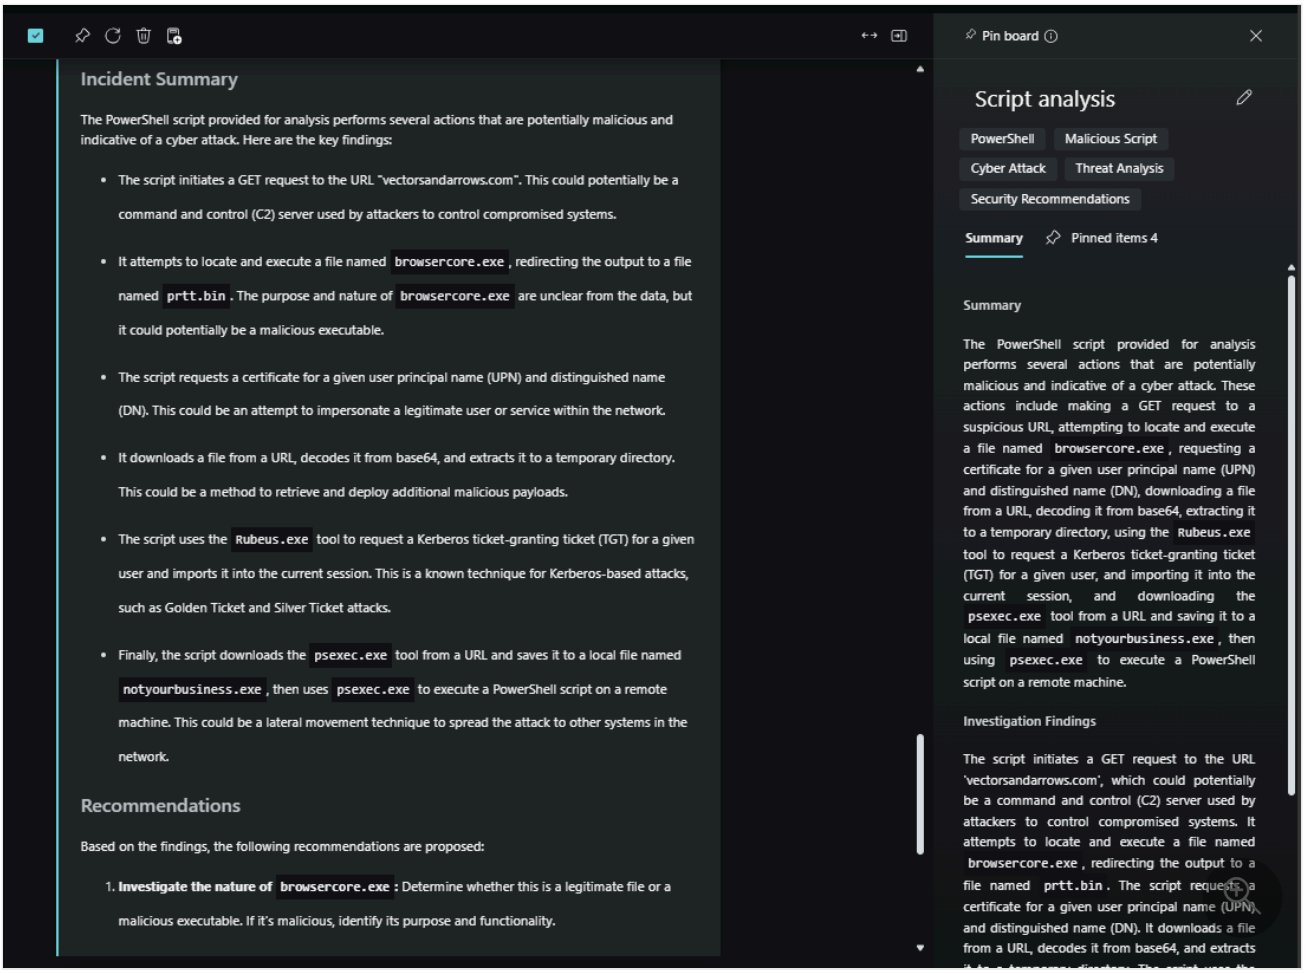

Conclusion

In this use case, Security Copilot helped conduct a thorough investigation of an incident. Using natural language, analysts are able to get an explanation of what the suspicious script is doing and verify the indicators in the script are associated with known threat actors.

In addition, Security Copilot generated an assessment through a summary report and provided a set of recommendations to contain the incident, which can also be used to improve skills.

Last updated on 05/29/2026

Use case: Incident response and remediation

10/24/2025

Role mentioned: Tier 1 SOC analysts investigating security incidents and resolving noncritical alerts

Scenario

Security Copilot can support Tier 1, entry-level, or junior analysts in assessing and managing incidents with greater efficiency and effectiveness.

This use case presents an abbreviated version of the incident analysis in Case Study: Harnessing Security Copilot in Defending Against Cyberthreats and extended here to include remediation steps.

You're a tier 1 analyst. An incident has been assigned to you by your team lead and your job is to determine whether the incident is indeed malicious and if it is, to resolve it as soon as possible.

Steps

1. In the Microsoft Defender XDR incident queue, you can see that the high-severity incident titled *Multistage incident involving Initial access & Lateral movement on multiple endpoints reported by multiple sources* correlates 25 alerts from various Microsoft security solutions. It involves three devices, four users, and one email account.

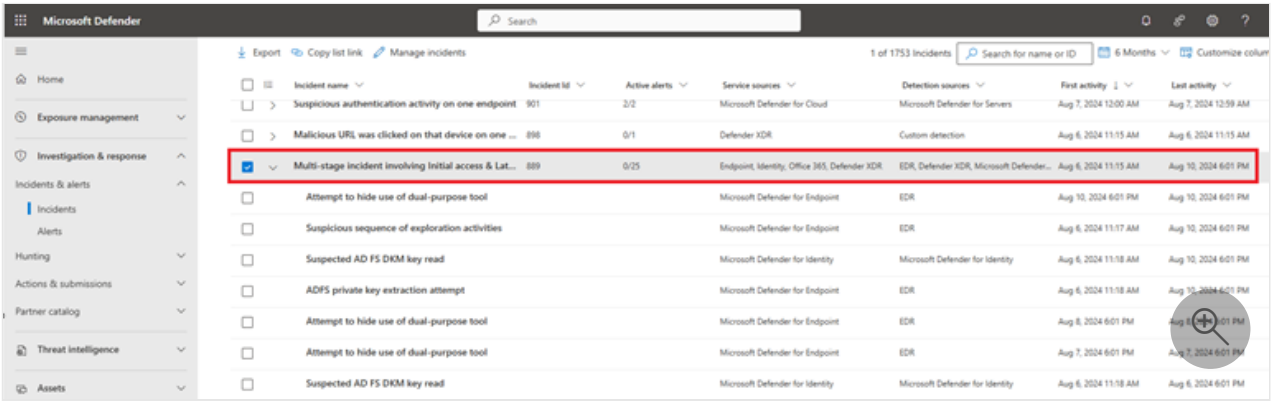

2. Open the incident page to see the attack story. You check the Security Copilot pane on the right side of the attack story and see that it automatically creates a summary of the incident, along with some steps under guided response.

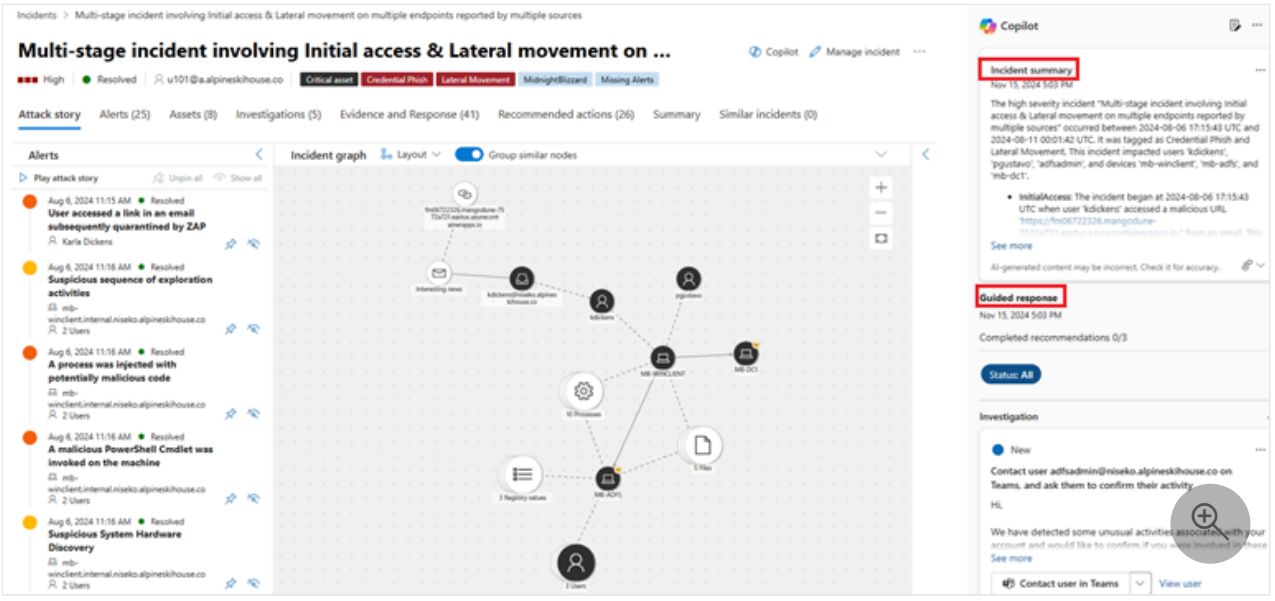

3. To get an overview of what transpired, you read the summary provided by Security Copilot. You copy the incident summary by selecting **Copy to clipboard** in the options menu and paste it in another document.

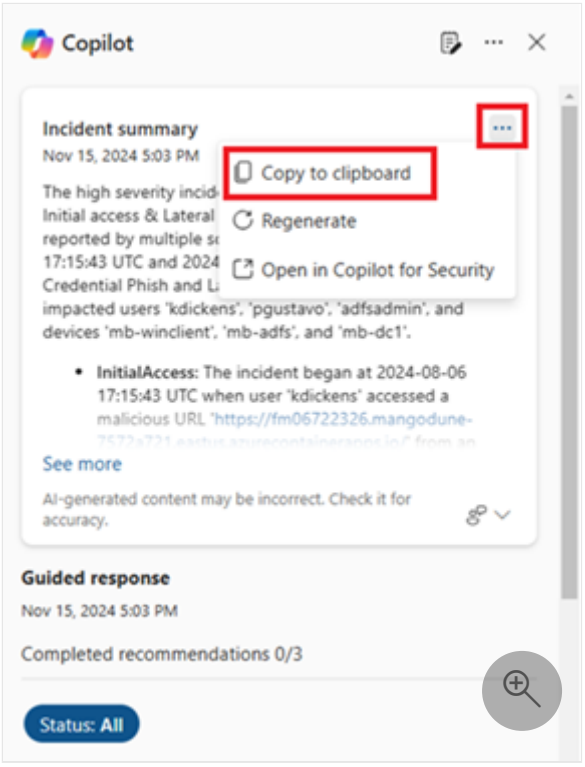

This is useful for the required report as it is already organized by attack phase and highlights the important aspects of the attack. Without Security Copilot, you would need to develop the report from scratch.

4. You then easily determine that the complex incident involved multiple attacks. You go through each of the attack phase as laid out in Initial access, discovery, and defense evasion. For this use case, you're interested in the suspected DCSync attack under the credential access activity.

5. You select **Open in Security Copilot** from the Copilot pane to ask about DCSync attacks.

   Prompt: *What's a DCSync attack? Can you provide its TTPs?*

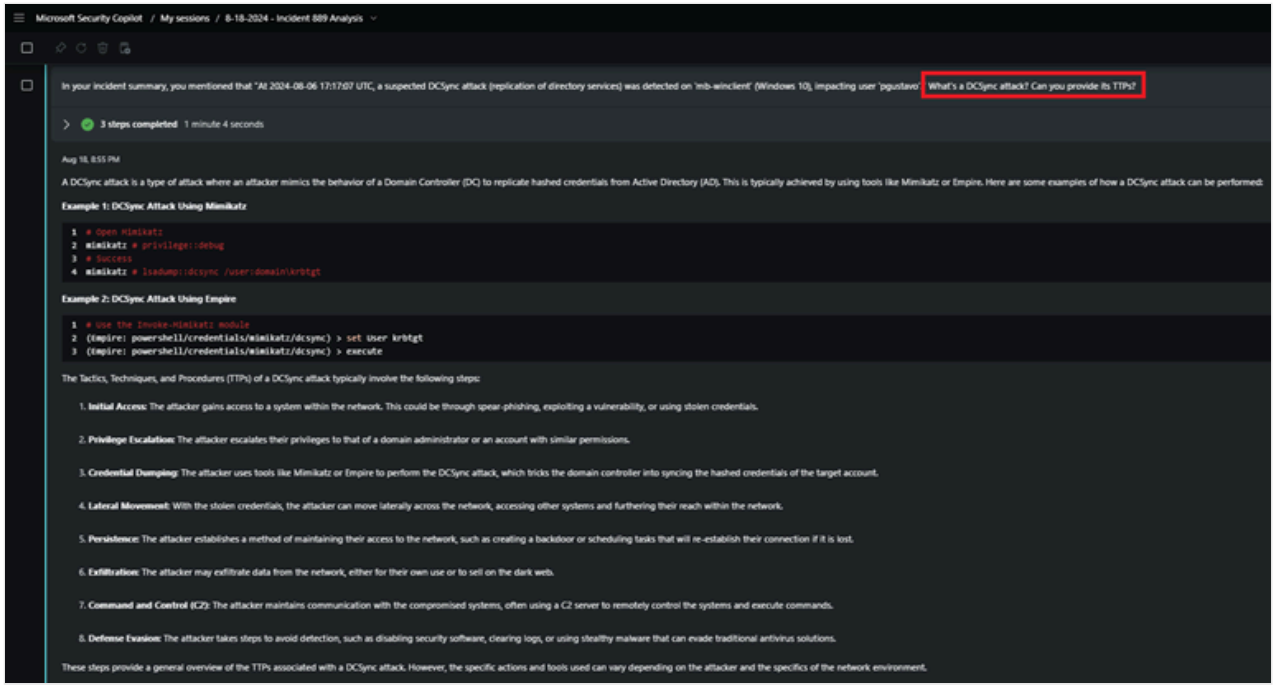

   Security Copilot tells you that in DCSync attacks, an attacker mimics the behavior of a domain controller to replicate hashed credentials from Active Directory, usually through Mimikatz or Empire.

6. Knowing this, you return to Defender XDR and look at the *Suspected DCSync attack* more closely. You see that the user sent a replication request to a domain controller. You ask the standalone Security Copilot whether this is considered typical behavior.

   Prompt: *Is it common for a regular user on a Windows 10 device to send a replication request to a DC?*

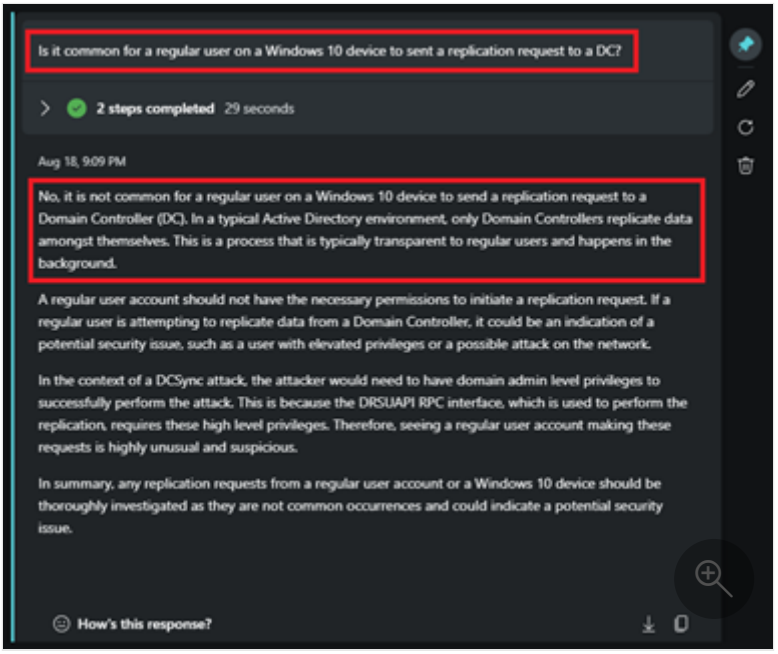

With Security Copilot's response you can confirm that this is a true security incident that needs further investigation by the Tier 2 team.

7. The next security alert in the incident summary is a suspicious service creation event that is classified under Persistence. You see that the same alert title was generated by Defender for Identity and Defender for Endpoint. However, when you check the alert, you see that the service path listed is encoded.

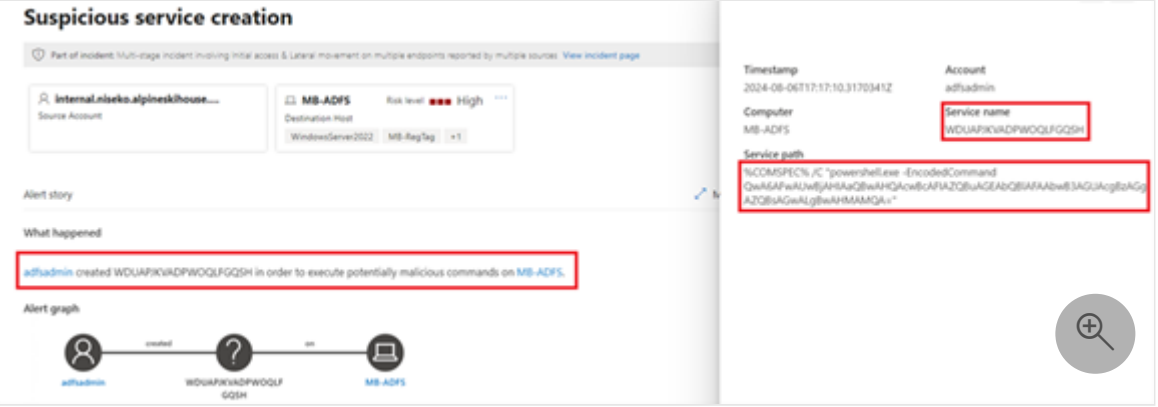

8. To investigate the script, you consult Security Copilot.

Prompt: *Can you decode then tell me more about this service path: %COMSPEC% /C "powershell.exe -EncodedCommand QwA6AFwAUwBjAHIAaQBwAHQAcwBcAFIAZQBuAGEAbQBlAFAAbwB3AGUAcgBzAGgAZQBsAGwALgBwAHMAMQA="*

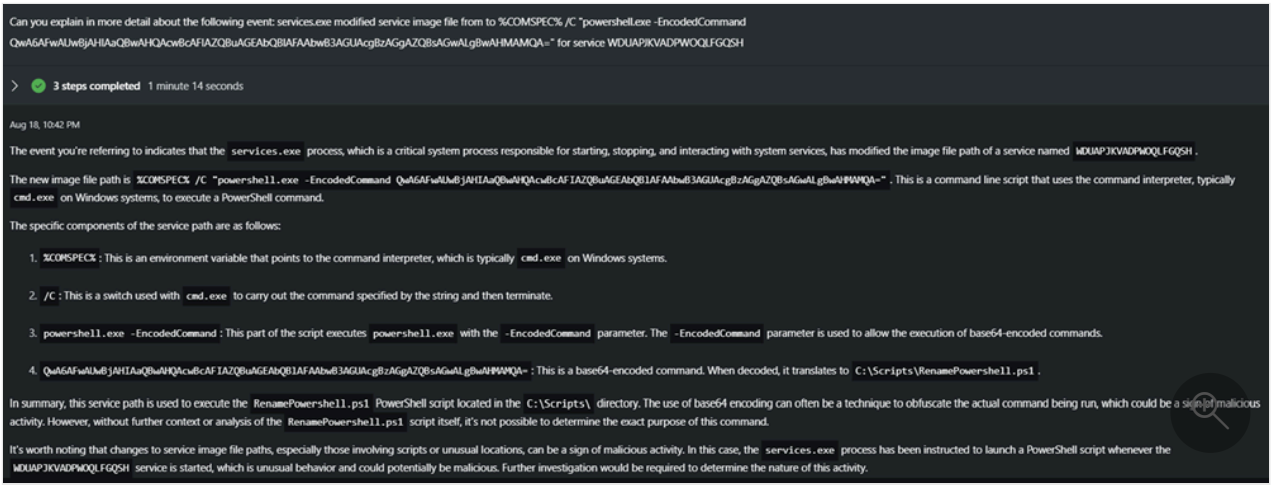

Copilot identifies the encoded command as Base64 encoding and decodes it for you. It appears that the service path is used to run a PowerShell script located in a certain directory. The use of encoding to obfuscate the command could be a sign of malicious activity.

9. As you gain more confidence that this is a true security incident, you can go back to review the **Guided response** in the Security Copilot pane in the incident page to determine which actions you can take to contain or mitigate the attack.

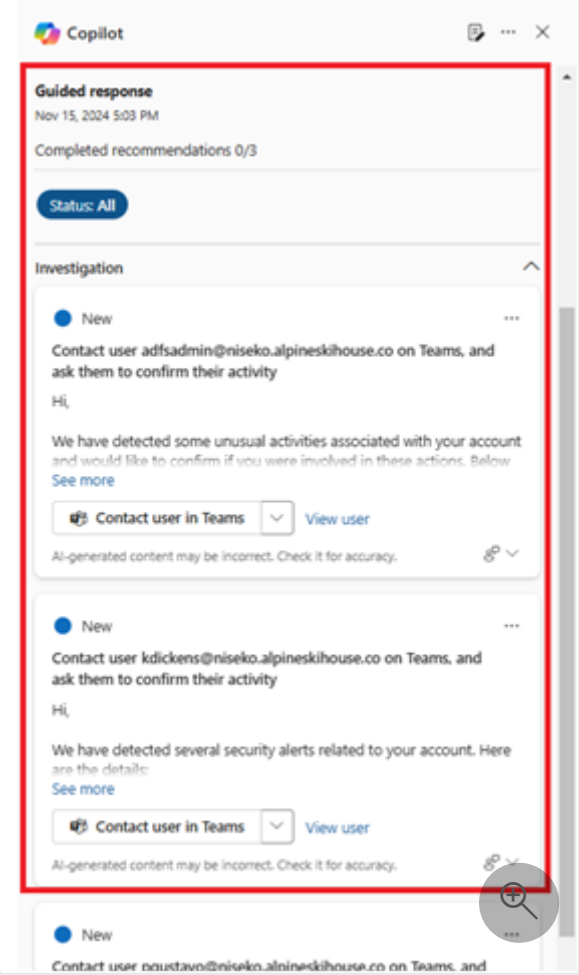

Conclusion

Using Security Copilot can provide several efficiency benefits to analysts tasked to investigate and remediate large incidents involving multiple alerts and entities. It helps in understanding the attack overview, in providing analysis for malicious scripts, in providing context for named attacks, and in performing the necessary steps to remediate the threat.

Best practices for effective prompts

As with any Microsoft Copilot, a prompt refers to the text-based, natural language input you provide in the prompt bar that instructs Security Copilot to generate a response. The quality of the response that Security Copilot returns depends in large part on the quality of the prompt used. In general, a well-crafted prompt with clear and specific inputs leads to more useful responses by Security Copilot.

Watch the following video to learn more about creating effective prompts:

https://learn-video.azurefd.net/vod/player?id=b7fe7b15-6342-4f0b-bed0-9e8ad118ad96&locale=en-us&embedUrl=%2Fcopilot%2Fsecurity%2Fprompting-tips

To see the different ways you can create prompts in Security Copilot, read Prompting in Security Copilot.

Elements of an effective prompt

Effective prompts give Security Copilot adequate and useful parameters to generate a valuable response. Security analysts or researchers should include the following elements when writing a prompt.

- **Goal** - specific, security-related information that you need
- **Context** - why you need this information or how you plan to use it
- **Expectations** - format or target audience you want the response tailored to
- **Source** - known information, data sources, or plugins Security Copilot should use

Every good prompt should have a goal. Whether it comes in the form of instructions or questions, it should indicate what you want out of your current session.

For Security Copilot, context can mean such information as the time frame, or that you plan to use the response for a report. Expectations can include whether you want the response to be in a table format, a list of action steps, a summary, or even a diagram. Source might be useful in specifying which Microsoft plugins you're referring to, if needed. Some plugins require more context to work effectively or supporting plugins to ensure a response when initial responses fail.

Other prompting tips

Some things to remember when coming up with your own prompts:

- **Be specific, clear, and concise** as much as you can about what you want to achieve. You can always start simply with your first prompt, but as you get more familiar with Security Copilot, include more details following the elements of an effective prompt.
  - Basic prompt: Pearl Sleet actor
  - Better prompt: Can you give me information about Pearl Sleet activity, including a list of known indicators of compromise and tools, tactics, and procedures (TTPs)?

- **Iterate.** Subsequent prompts are typically needed to either clarify what you need further, or try other versions of a prompt to get closer to what you're looking for. Like all LLM-based systems, Security Copilot can respond to the same prompt in slightly different ways.

- **Provide necessary context** to narrow down where Security Copilot looks for data.
  - Basic prompt: Summarize incident 15134.
  - Better prompt: Summarize incident 15134 in Microsoft Defender XDR into a paragraph that I can submit to my manager and create a list of entities involved.

- **Give positive instructions** instead of "what not to do". Security Copilot is geared toward action, so telling it what you want it to do for exceptions is more productive.
  - Basic prompt: Give me a list of unmanaged devices in my network.
  - Better prompt: Give me a list of high-risk unmanaged devices in my network. If they're named "test", remove them from the list.

- **Directly address Security Copilot as "You"**, as in, "You should ..." or "You must ...", as this is more effective than referring to it as a model or assistant.

While these guidelines can help you get started in creating prompts, it’s important to note that you’re not limited to forming prompts following the structure of the previous examples. What’s great about Security Copilot is that it's designed to respond to questions or instructions made in your own words (that is, using natural language).

You have the flexibility to adapt these guidelines to your specific needs.

Watch the following video to learn more about creating better prompts:

https://learn-video.azurefd.net/vod/player?id=97fec5b0-f2b5-4b3e-8f7f-70bb875a2dcb&locale=en-us&embedUrl=%2Fcopilot%2Fsecurity%2Fprompting-tips

Related content

Use the following Security Copilot articles for guidance:

- Prompting in Security Copilot
- Build your own promptbooks

Last updated on 06/02/2026

Access the Security Copilot audit log

In today's stringent regulatory environment, it's important for organizations to monitor and analyze how users interact with security products. Organizations may need to keep track of actions, transactions, and configuration settings on a platform to ensure that they're meeting compliance regulations and regulatory standards.

Security Copilot today provides access to audit logs through Microsoft Purview Unified Audit Log (UAL), Microsoft Purview Data Security Posture Management (DSPM) for AI, and the Office Management API to help you satisfy compliance and regulatory requirements. The UAL from Purview gives you visibility into information such as administrator events and activity metadata, whereas DSPM for AI provides insights into prompt/response pairs.

- **Admin events** - Privileged actions such as changes to tenant-level settings or administrative changes (for example, data sharing, agent, agent trigger, plugin and promptbook configurations).
- **Activity metadata** - Logs of user interactions within the Security Copilot platform (for example, a user asked a prompt at a specific time with information on the activity type).
- **Prompt Response Pairs** - Activity in Security Copilot can contain a prompt to a skill, which produces a response pair. The pair and its content are a core functionality of the platform and a critical point of inspection for audit purposes.

> **Note**
>
> Customer content such as that contained within the prompt and response pair is currently only included through Microsoft Purview DSPM for AI. For more information, please see Microsoft Purview for DSPM.

By keeping track of these interactions, you can potentially identify risks, ensure stable operation of important workflows, and protect production data.

Enable Security Copilot for Purview Audit and DSPM Solutions

To enable Purview DSPM for AI and UAL solutions for Security Copilot, the capability must be enabled within Security Copilot and configured within the respective Purview solutions. There are

two options to enable the capabilities within Security Copilot that are as follows:

- During the first run experience, a Security Administrator may opt into allowing Microsoft Purview to access, process, copy and store administrator actions, user actions, and Copilot responses. For more information on getting started with Security Copilot see Get started with Security Copilot.

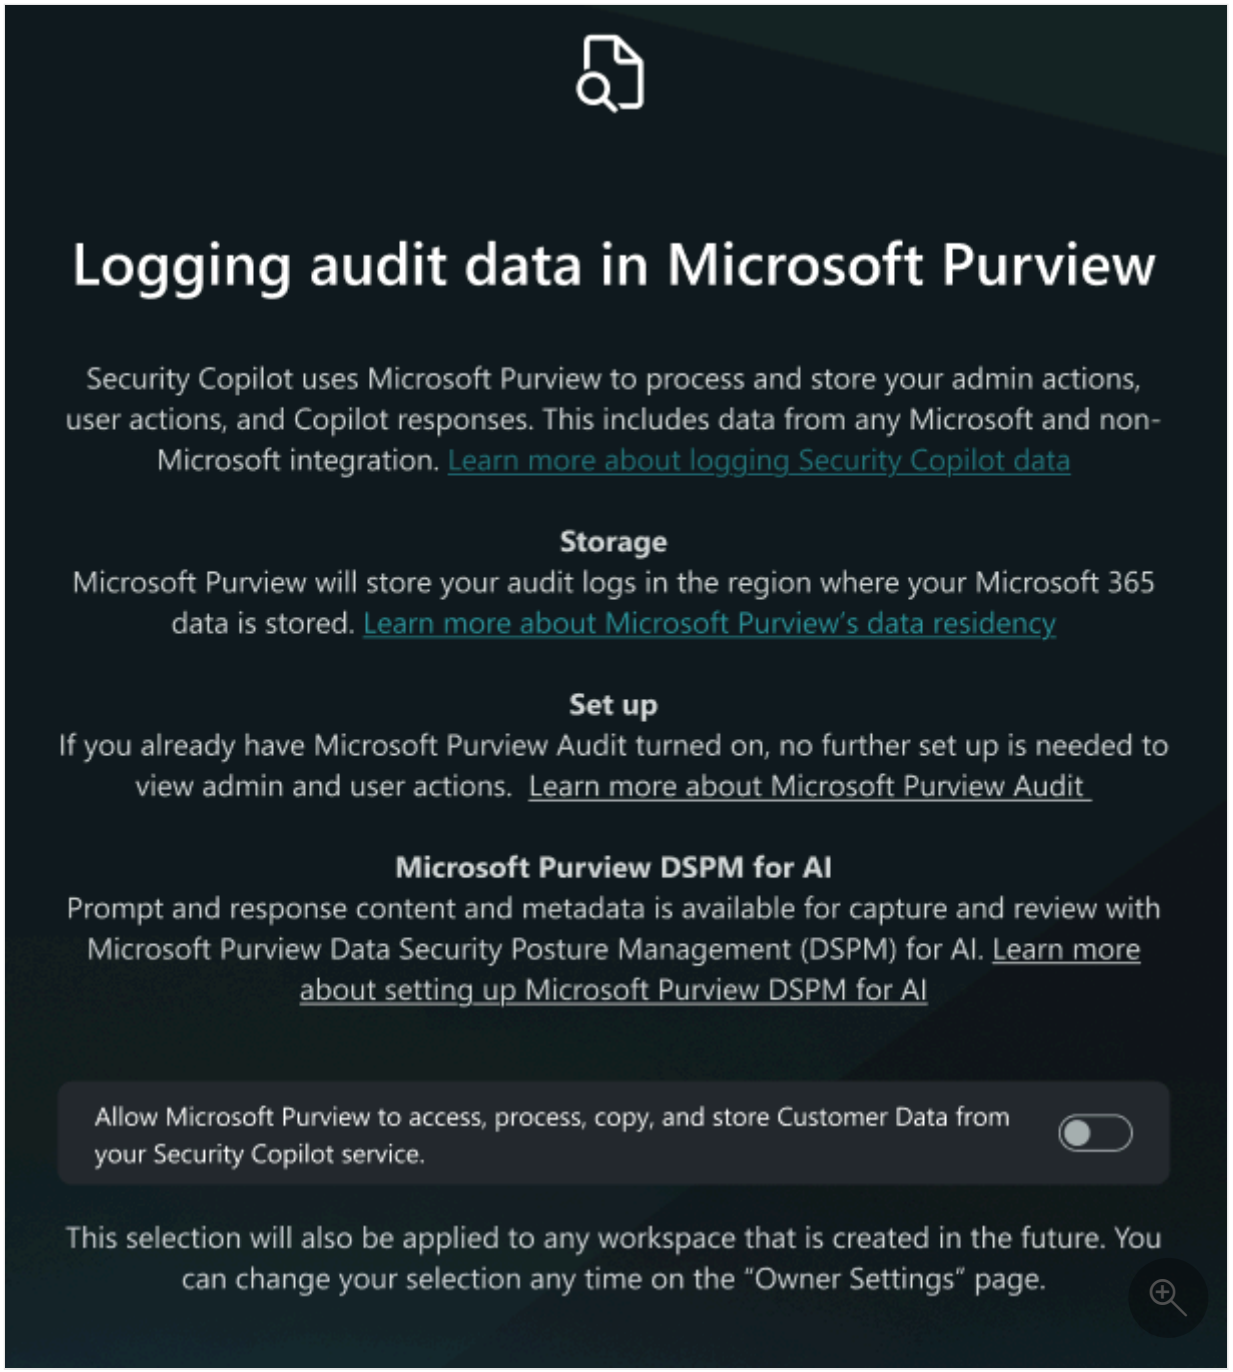

- A Security Administrator can also access this option through the Owner settings page. For more information on role types, see Understand authentication.

Update audit log settings

Use the following steps to update the audit log settings:

1. Sign in to Security Copilot (https://securitycopilot.microsoft.com).
2. Select the home menu icon.
3. Navigate to the **Owner** > **Owner settings** > **Logging audit data in Microsoft Purview**.

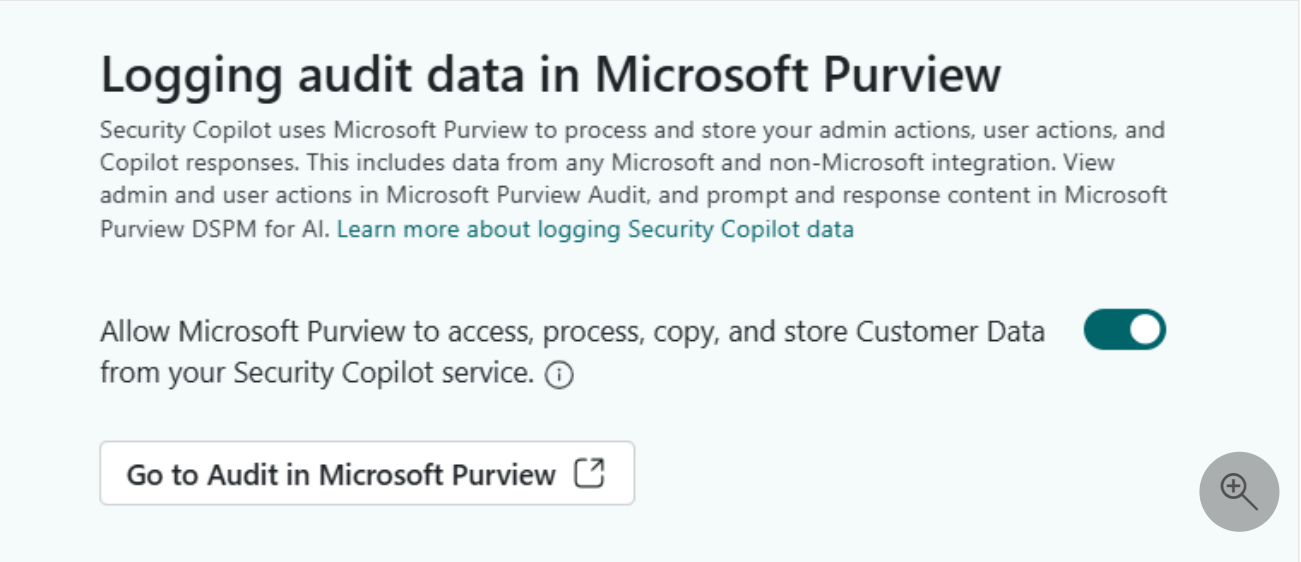

> **Important**
>
> Microsoft Purview will store your Customer Data in the region where your Microsoft 365 data is stored. For more information, see, Privacy and data security. The default retention period for audit logs is 180 days but can be extended using audit log retention policies. For more information, see Manage audit log retention policies.

4. You can turn the toggle on or off.

Access Security Copilot audit logs via Unified Audit Log in Microsoft Purview

Before you begin

This section gives an overview of the prerequisites to access the Unified audit log.

You'll need to:

- Verify that you have opted to allow Microsoft Purview access inside Security Copilot. For more information, see Enable the audit log capability.
- Verify that the audit log feature is active within Microsoft Purview. For more information, see Before you search the audit log.

> **Note**
>
> You'll require proper permissions to access the audit log in Microsoft Purview. For more information, see **Permissions in the Microsoft Purview portal**. These access rights might be different from those in Security Copilot.

Options to access the Unified Audit Log in Microsoft Purview

You can take the following actions to access the audit log in Microsoft Purview:

- Search the audit log - Article that provides instructions on how to access and search through the audit log event data to gain insight and further investigate user activities.
- Search through the audit log activities - Article that describes the activities that are captured in the audit log.
- Activities Operation Name References for Security Copilot - A listing of activity operation names for Security Copilot management, agent management, and AI interactions such as prompts.
- Search the audit log using a PowerShell script - Article that provides instructions on how you can run a PowerShell script to search the audit log to help you investigate security incidents and compliance issues.
- Monitor user activities and system events with Security Copilot and Microsoft Sentinel - Blog that provides insight on how to leverage sending Security Copilot audit logs to your cloud native SIEM to gain deeper insights into usage and take proactive measures to mitigate risks.

Access prompt response pair content in Microsoft Purview DSPM for AI

Before you begin

This section gives an overview of the prerequisites to access prompt response pair content via Purview DSPM for AI.

You'll need to:

- Verify that you have opted to allow Microsoft Purview access inside Security Copilot. For more information, see Enable the audit log capability.
- Verify that necessary setup steps are complete, see Learn | DSPM for AI Setup.

- Validate that any administrator requiring policy creation or the ability to view content captured in prompts and responses meets the minimum permission levels. The role and permission configurations are detailed on Learn | DSPM for AI
- Confirm that DSPM policies for Copilot content capture have been configured, see DSPM for AI Default Policies on Learn.

Once the prerequisties are met different ways to enact information protection, and viewing the content within prompt and response pairs are posbible.

- Review the capbilites of the Purview DSPM for AI solution Microsoft Purview data security and compliance protections for generative AI apps
- Protect the information accessed by different Copilot experiences Information protection considerations for Copilot
- View and search content within the Activity explorer events for Purview DSPM for AI Activity explorer events

---

Last updated on 05/29/2026

Connect your org knowledge base (Preview)

> **Important**
>
> Some information in this article relates to a prereleased product which might be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

Using the Azure AI search plugin or file upload, you can bring in business context like your organization's security knowledge base into Microsoft Security Copilot so you can directly reference them in your session.

An organization's knowledge base might include the repository of an organization's proprietary information, like: wikis, policy and compliance documents, investigation and response procedures templates, historical data, KQL (Kusto Query Language) libraries, or even information from public repositories like Microsoft Learn documentation, prescriptive guidelines, and frameworks.

Integrating this wealth of knowledge into Copilot allows Copilot to reason over the knowledge base or documents and generate responses that are more relevant, specific, and customized to your operational need.

Options for connecting a knowledge base

There are two ways to integrate a knowledge base into Copilot:

1. Azure AI Search plugin – set up the Azure AI Search plugin to connect to an Azure AI Search index containing your organization's knowledge base. Follow the steps in Prompting for a knowledge base connected using Azure AI Search.
2. File upload – upload your files and get Copilot to reason over them. Follow the steps in Prompting for a knowledge base uploaded as a file.

Prompting for a knowledge base connected using Azure AI Search

To connect a knowledge base using the Azure AI Search plugin, follow the steps in Azure AI Search.

Once connected, prompt Copilot to look for information related to the knowledge base you uploaded. Make sure to mention "**Azure AI Search**" in the prompt. For example:

- New staff can ask Copilot about the organization's processes within their Copilot workflows. This avenue can speed up the onboarding process.

  Sample prompt: Summarize my organization's multifactor authentication policies in Azure AI Search.

- While performing an incident investigation, SOC analysts can ask Copilot for investigation steps from their organization's processes and procedures.

  Sample prompt: Search my Azure AI Search index for investigation steps that I should follow for this incident involving multiple failed sign-in attempts by a user.

The Azure AI Search plugin performs a semantic and keyword search across the index's contents to find the relevant information so it's advisable to be as specific as you can in prompting. For more on the elements of an effective prompt and other prompting tips, read Create effective prompts.

Prompting for an uploaded file

To upload a file, follow the steps in Add a source by uploading a file.

You can use this option to directly add text documents as a source that Copilot can refer to when responding to a prompt that references it.

In the prompt bar, you need to mention "uploaded files" if you want Copilot to reason over your available files. You can also include the file name if you would like to guide Copilot to reason over a specific file. For example:

- You can ask Copilot to check a user account's actions against your organization's data handling policies to find out if any violations were committed.

  Sample prompt: Based on the Contoso Data Handling Policy file, list any actions taken by user Preston-123 that might be a violation of the data handling policies. Include a verdict to the action and level of confidence of the verdict. Cite the policy name and section applicable to the verdict.

Known limitations

For both the Azure AI Search plugin and file upload methods of connecting knowledge bases:

- Copilot reasons over text content only, not other data or media types.
- Searches over structured content like tables aren't well-supported yet.

Related content

Use the following Security Copilot articles for guidance:

- Azure AI Search plugin in Security Copilot
- Plugins overview
- Manage plugins
- Responsible AI

Last updated on 06/04/2026

Add a source by uploading a file (Preview)

> **Important**
>
> Some information in this article relates to a prereleased product which might be substantially modified before it's commercially released. Microsoft makes no warranties, express or implied, with respect to the information provided here.

You can upload a file in Microsoft Security Copilot to allow Copilot to refer to the contents of the file for responses when you prompt Copilot and mention your "uploaded files."

Examples of files you can upload are your organization's policy and compliance documents, investigation and response procedures, and templates. Integrating this wealth of knowledge into Copilot allows Copilot to reason over the knowledge base or documents and generate responses that are more relevant, specific, and customized to your operational needs.

> **Note**
>
> The file upload feature in Copilot doesn't support use cases such as evaluation or detonation of malicious files today.

Upload a file to Security Copilot

1. Check that the file you're about to upload is a common text file type like DOCX, MD, PDF, and TXT formats, and that each file doesn't exceed 3 MB. You can upload files up to 20 MB in total.
2. In the Copilot home page, select the sources icon in the prompt bar to open the **Manage sources** page.
3. Select **Files** to open the file upload section.

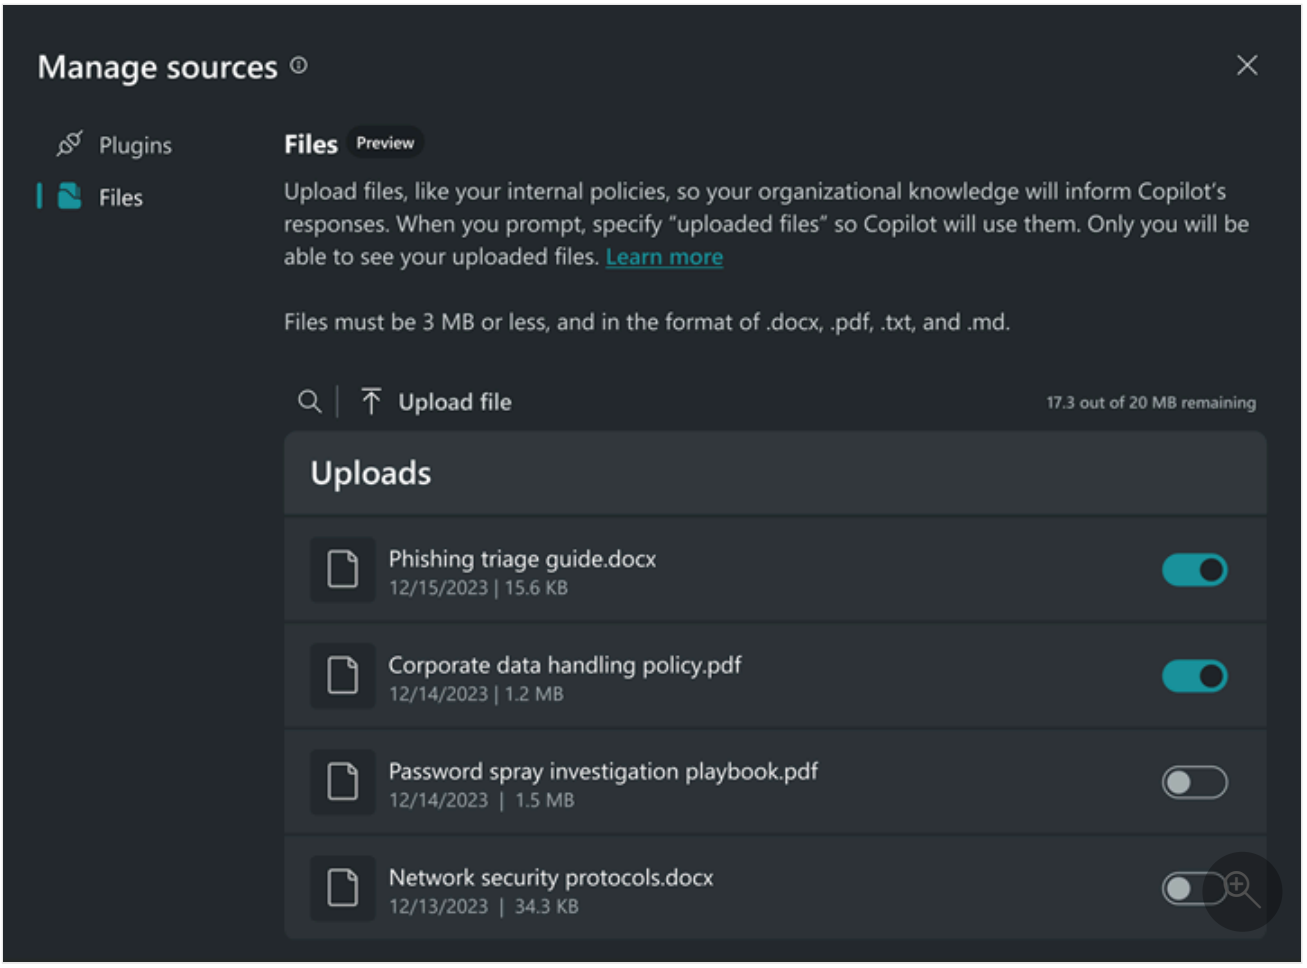

4. Select **Upload file** to look for your file.

5. Wait for the file to appear in **Uploads**. If an error message appears, correct the problem and try again.

6. To include the file as a source in your current session, toggle the button beside the file until the toggle is on the right (lit up). If you aren't going to use the file yet, or exclude the file as a source in your current session, keep the toggle button untoggled. (Toggle is on the left, in grey.)

7. To prompt using the uploaded file, read Prompting for an uploaded file.

Storage and availability

Uploaded files are always stored in the home Geo of the tenant. Uploaded files are stored in the Security Copilot service, not inside your Microsoft tenant boundary.

Uploaded files are only available to the user account that uploaded them, and not available to other users within or outside the tenant.

To learn more about how Security Copilot handles Customer Data, read Privacy and data security in Security Copilot.

Delete an uploaded file

To delete the file, select the trash icon next to the file name in **Manage sources** > **Files** > **Uploads**.

Restrict file upload (for owners)

Owners can choose who are allowed to upload files in owner settings. The options are:

- No one can upload files
- Contributors and Owners can upload files

By default, file upload is available for all users (Contributors and Owners).

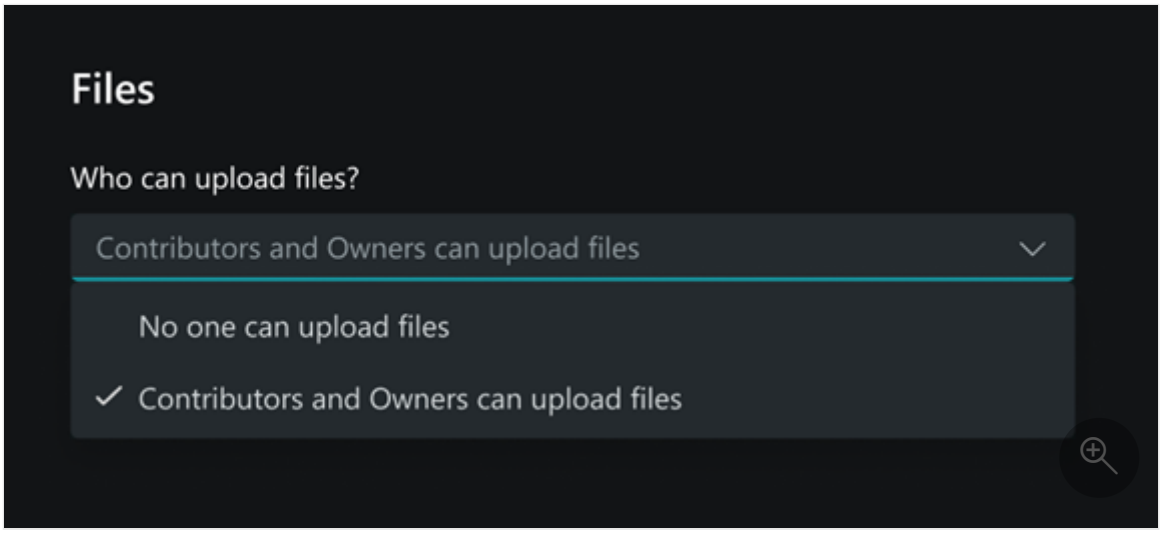

Related content

Use the following Security Copilot articles for guidance:

- Plugins overview
- Manage plugins
- Responsible AI

Last updated on 06/04/2026

Grant partners access to Microsoft Security Copilot

If you're working with a Microsoft Managed Security Solution Provider (MSSP), ensure they access your Security Copilot capabilities only when you grant them access.

There are several ways to allow a partner to manage your Security Copilot.

1. Azure Lighthouse
   Approve your MSSP to gain Security Copilot permissions for your Microsoft Sentinel workspaces and other supported Azure resources. The capacity plan (SCUs) used are the partner tenant's capacity plan.

2. GDAP (recommended)
   Approve your MSSP to gain Security Copilot permissions for your tenant. They assign a security group the permissions needed using Granular Delegated Admin Privileges (GDAP). The capacity plan (SCUs) used are the customer's capacity plan.

3. B2B collaboration
   Set up guest accounts for individuals from your partner to log into your tenant. The capacity plan (SCUs) used are the customer tenant's capacity plan.

Compare access methods

There are tradeoffs for these methods. Use the following table to help decide which method is best for your organization. It's possible to mix these methods for an overall partner strategy.

Expand table

Consideration	GDAP	B2B collaboration	Azure Lighthouse
How is **time-bound access** implemented?	Access is time-bound by default and built into the permission approval process.	Customers can implement Privileged Identity Management (PIM) with time-bound access, but they must maintain it themselves.	Privileged Identity Management (PIM) with time-bound access is possible for Azure role-based access control (Azure RBAC).

Consideration	GDAP	B2B collaboration	Azure Lighthouse
How is **least-privileged access** administered?	GDAP requires security groups. A list of least-privileged roles needed guides the setup.	Security groups are optional, and maintained by customer.	Azure Lighthouse requires Microsoft Entra security groups. For more information, see Tenants, users, and roles in Azure Lighthouse scenarios.
What **plugins are supported**?	A partial set of plugins are supported.	All plugins available for the customer are available to the partner.	Azure Lighthouse supports delegated access to Azure resources, which includes Microsoft Sentinel. For more information on supported workloads, see What is Azure Lighthouse.
What is the **standalone login** experience?	The MSSP uses **Service management** to seamlessly log into Security Copilot for the appropriate tenant.	Use the tenant switch selection from the Security Copilot setting.	The MSSP analyst, from their Security Copilot standalone portal, is able to prompt against any Microsoft Sentinel workspace that the analyst has permissions to access using the Azure Lighthouse delegated access.
What is the **embedded experience**?	Supported, with **Service management** links to facilitate access.	Supported normally.	Supported, with the Unified SecOps Platform integration between Microsoft Defender XDR and Microsoft Sentinel. Microsoft Sentinel incidents show under "Investigation & response" allowing for embedded Security Copilot experiences to be ran. For more information, see Microsoft Defender XDR integration with Sentinel.
Who's Security Copilot capacity plan (SCUs) are used when performing prompts?	Customer tenant's capacity plan.	Customer tenant's capacity plan.	Partner tenant's capacity plan.

Grant partner access using Azure Lighthouse

Azure Lighthouse allows an MSSP to set up least-privileged and time-bound access explicitly granted by the Security Copilot customer for delegated access to Azure resources. When performing Security Copilot prompts within the partner's tenant against Microsoft Sentinel

workspaces of the customer, the partner's capacity plan SCUs are used instead of the customer's capacity plan SCUs.

For more information, see Introduction to Azure Lighthouse.

Here's the current matrix of Security Copilot plugins that support Azure Lighthouse:

Expand table

Security Copilot plugin	Supports Azure Lighthouse
Microsoft Defender External Attack Surface Mgmt	No
Microsoft Defender Threat Intelligence	No
Microsoft Defender XDR	No
Microsoft Entra	No
Microsoft Intune	No
Microsoft Purview	No
Microsoft Sentinel	Yes
NL2KQL Defender	No
NL2KQL Sentinel	Yes

Onboard a customer to Azure Lighthouse

1. Gather tenant and subscription details

    - To onboard a customer's tenant, you must have an active Azure subscription, and you'll need the following information to create an Azure Resource Manager template.
        - MSSP's tenant ID
        - Customer's tenant ID
        - Azure Subscription IDs for each specific subscription in the customer's tenant that will be managed by the MSSP including the offering details.
        - For more information, see, Onboard a customer to Azure Lighthouse.
        - For more information If you require publishing a private or public managed service offering to Azure Marketplace, see Publish a Managed Service offer to Azure Marketplace.

2. Define roles and permissions

- As a service provider, you may want to perform multiple tasks for a single customer, requiring different access for different scopes. You can define as many authorizations as you need in order to assign the appropriate Azure built-in roles. To define authorizations in your template, you must include the ID values for each user, user group, or service principal in the managing tenant to which you want to grant access. You'll also need to include the role definition ID for each built-in role you want to assign. For more information, see, Onboard a customer to Azure Lighthouse - Azure Lighthouse.

- To use Microsoft Entra Privileged Identity Management (PIM) for just-in-time access, you need to create eligible authorizations. For more information, see Create eligible authorizations.

3. Create an Azure Resource Manager template

- To onboard your customer, you need to create an Azure Resource Manager template for your offer with the following information. The mspOfferName and mspOfferDescription values are visible to the customer in the Service providers page of the Azure portal once the template is deployed in the customer's tenant. You can create this template in the Azure portal, or by manually modifying the templates provided in our samples repo. For more information, see, Onboard a customer to Azure Lighthouse.

4. Deploy Azure Resource Manager template

- Once you have created your template, a user in the customer's tenant must deploy it within their tenant. A separate deployment is needed for each subscription that you want to onboard (or for each subscription that contains resource groups that you want to onboard). The deployment may be done by using PowerShell, by using Azure CLI, or in the Azure portal. For more information, see, Onboard a customer to Azure Lighthouse.

5. Confirm successful onboarding

- When a customer subscription has successfully been onboarded to Azure Lighthouse, users in the service provider's tenant will be able to see the subscription and its resources. You can confirm this within the Azure portal, using PowerShell or by using Azure CLI. For more information, see, Onboard a customer to Azure Lighthouse.

Grant partner access using GDAP

GDAP allows an MSSP to set up least-privileged and time-bound access explicitly granted by the Security Copilot customer. Only MSSPs registered as a Cloud Solution Partner (CSP) are allowed to manage Security Copilot. Access is assigned to an MSSP security group, which reduces the administrative burden for both the customer and the partner. An MSSP user is assigned the appropriate role and security group to manage the customer.

For more information, see Introduction to GDAP.

Here's the current matrix of Security Copilot plugins that support GDAP:

Expand table

Security Copilot plugin	Supports GDAP
Defender External Attack Surface Management	No
Microsoft Entra	Overall, no, but a few capabilities work.
Intune	Yes
MDTI	No
Defender XDR	Yes
NL2KQL Defender	Yes
NL2KQL Sentinel	No
Purview	Yes
Microsoft Sentinel	No

For more information, see Workloads supported by GDAP.

GDAP relationship

1. The MSSP sends a GDAP request to their customer. Follow the instructions in this article, Obtain permissions to manage customer. For best results, the MSSP should request **Security reader** and **Security operator** roles to access Security Copilot platform and plugins. For more information, see Understand authentication.

2. The customer approves the GDAP request from the partner. For more information, see Customer approval.

Security group permissions

1. The MSSP creates a security group and assigns it the approved permissions. For more information, see Assign Microsoft Entra roles.

2. The customer adds the roles the MSSP requested to the appropriate Security Copilot role (Copilot owner or copilot contributor). For example, if the MSSP requested **Security operator** permissions, the customer adds that role to the **Copilot contributor** role within Security Copilot. For more information, see Assigning Security Copilot roles.

MSSP Security Copilot access

1. The MSSP user account needs membership to the partner security group assigned and an approved role to access the delegated Microsoft Sentinel workspace of the customer in the partner tenant.

2. The MSSP can access Security Copilot standalone portal and use the example prompts below:

    - "List all Sentinel workspaces" to see all available Sentinel workspaces within the partner tenant and the delegated customer Microsoft Sentinel workspaces.

    - "Summarize incident <incident ID> from workspace <name of customer's Microsoft Sentinel workspace>" to see a summarization of the Microsoft Sentinel incident from the customer's Microsoft Sentinel workspace.

3. The MSSP can access their Security Copilot's Usage Monitoring Dashboard in the standalone experience to see the SCU cost associated with the prompt performed against the customer's Microsoft Sentinel workspace. This can be validated by seeing the session ID in the Usage Monitoring Dashboard and the session ID in the browser URL when viewing the Security Copilot session.

Grant partner access using B2B collaboration

B2B collaboration invites individual partner accounts as guests to the customer tenant to operate Security Copilot.

Set up a guest account for your partner

> ⓘ **Note**

To perform the procedures described in this option, you must have an appropriate role, such as User Administrator, or Billing Administrator, assigned in Microsoft Entra.

1. Go to the Microsoft Entra admin center and sign in.
2. Go to **Identity** > **Users** > **All users**.
3. Select **New user** > **Invite external user**, and then specify settings for the guest account.
    a. On the **Basics** tab, fill in the user's email address, display name, and a message if you want to include one. (You can optionally add a **Cc recipient** to receive a copy of the email invitation.)
    b. On the **Properties** tab, in the **Identity** section, fill in the user's first and last name. (You can optionally fill in any other fields you want to use.)
    c. On the **Assignments** tab, select **+ Add role**. Scroll down, and select either **Security Operator** or **Security Reader**.
    d. On the **Review + invite** tab, review your settings. When you're ready, select **Invite**.

    The partner receives an email with a link to accept the invitation to join your tenant as a guest.

Tip

To learn more about setting up a guest account, see **Invite an external user**.

B2B access to Security Copilot

After you have set up a guest account for your partner, you're ready to notify them that they can now use your Security Copilot capabilities.

1. Tell your partner to look for an email notification from Microsoft. The email contains details about their user account and includes a link they must select to accept the invitation.
2. Your partner accesses Security Copilot by visiting securitycopilot.microsoft.com and signing in using their email account.
3. The partner uses the tenant switch feature to ensure they're accessing the appropriate customer. For example, the following image shows a partner from Fabrikam using their

credentials to work in Security Copilot for their customer, Contoso.

Alternatively, set the tenant ID directly in the URL, for example, `https://securitycopilot.microsoft.com/?tenantId=aaaabbbb-0000-cccc-1111-dddd2222eeee`.

4. Share the following articles to help your MSSP get started using Security Copilot:

   - Navigate Microsoft Security Copilot
   - Get started with Microsoft Security Copilot
   - Investigate an incident's malicious script

Get technical support for partner access setup

Currently, if your MSSP or partner has questions and needs technical support for Security Copilot outside of the partner center, the customer organization should contact support on the MSSP's behalf.

Last updated on 05/22/2026

Export Activity API

This document describes the API contracts and expected output data schemas for the Security Copilot Admin Export APIs.

Export APIs provide workspace administrators with the ability to export prompts and prompt responses in a paginated format.

Authenticate and authorize Export Activity API requests

- Authorization: Bearer token authentication
- Required Permissions: Workspace Owner / Administrator

Authenticate with a Service Principal

1. **Create an app registration** (for example, use a name like `mysecuritycopilotexportapp`).
    - Use the Microsoft Entra ID quickstart: Register an app (portal).
    - If you prefer CLI, create the identity via the tutorial: Create a service principal (Azure CLI).
    - Add credentials to the app registration as needed: Add credentials (client secret/cert). For a step-by-step client-secret creation, see Create client secret (explicit steps) or the portal guide, Register an app and create a Service Principal (portal).

    > ⓘ **Note**
    >
    > Only existing Owners can do the specific action.

2. **Add the new Service Principal as an Owner in Security Copilot role assignments**

    See, Security Copilot roles and assignments. Security Copilot **supports assigning permissions to role-assignable groups (RAGs)**, so create a RAG and then add your Service Principal (SP) to it as follows:

    - Create a role-assignable group

- Add the SP to the group (pick either): Portal—Manage groups (add members) or API—Microsoft Graph: Add group member

3. **Retrieve a bearer token for your service principal.**

   Example using Azure CLI and a client secret:

```
# Login as service principal (supports no-subscription tenants)
az login --service-principal -u 94e67e0c-7c41-4f5b-b5ae-f5b5918e2382 -p <client-secret> --tenant 536279f6-15cc-45f2-be2d-61e352b51eef --allow-no-subscriptions

# Retrieve access token for Security Copilot (v1 resource pattern)
az account get-access-token --resource https://api.securitycopilot.microsoft.com
```

Authenticate as a User

Ensure your user account has the Security Copilot Owner role before retrieving a bearer token.

- Example using Azure CLI (v2 pattern with `/.default`):

Azure CLI

```
az account get-access-token --scope https://api.securitycopilot.microsoft.com/.default
```

Related content

- Sign in as SP (CLI)
- Get access token (CLI)
- .default scope explained
- Get access token (CLI) and .default scope.

Response behavior for non-Owners

For non-owners, the following response is returned for API calls:

JSON

```
{
  "message": "Your role doesn\u0027t have access to the info requested. Contact a Security Administrator to change your role or try again with a different account. Learn more about copilot",
  "code": "403",
  "traceId": "0HNF1M54NKVJ3:00000041",
```

```
  "error": {
    "message": "Your role doesn\u0027t have access to the info requested. Contact a
Security Administrator to change your role or try again with a different account.
Learn more about copilot",
    "copilotErrorId": "doesNotHaveAccessToSecurityCopilot",
    "code": "403",
    "innerError": {
      "message": null,
      "date": "2025-08-22T19:32:04.4726177Z",
      "correlationId": "aaaa0000-bb11-2222-33cc-444444dddddd"
    },
    "traceId": "0HNF1M54NKVJ3:00000041"
  }
}
```


API endpoints

Prompts Export API

HTTP

```
https://api.securitycopilot.microsoft.com/exports/prompts
```

`GET /exports/prompts`

Query parameters

N/A

Request example

HTTP

```
GET /exports/prompts?sessionCount=500&startDate=2024-01-01T00:00:00Z&endDate=2024-12-
31T23:59:59Z
Authorization: Bearer <token>
```

Response schema

Success response - `(200 OK)`

JSON


```
{
  "workspaceId": "string",
  "workspaceName": "string",
```

```
  "tenantId": "string",
  "prompts": [
    {
      // Prompt object schema (see Framework.Models.Prompt)
    }
  ],
  "sessionsContinuationToken": "string?",
  "totalCount": "integer?",
  "sessionCount": "integer"
}
```

Error codes and messages

Expand table

Status Code	Error Code	Error Message
400	Bad Request	Invalid parameters or missing workspace/tenant information
404	Not Found	Admin export APIs not enabled
500	Internal Server Error	Server error during export

Evaluations Export API

HTTP

```
https://api.securitycopilot.microsoft.com/exports/evaluations
```

`GET /exports/evaluations`

Query parameters

Expand table

Parameter	Type	Required	Default	Description
`sessionCount`	`integer`	No	`100`	Number of sessions to fetch (range: 1–1000).
`continuationToken`	`string`	No	`null`	Token for pagination.
`startDate`	`DateTimeOffset`	No	`null`	Start date filter (inclusive, ISO format).
`endDate`	`DateTimeOffset`	No	`null`	End date filter (inclusive, ISO format).
`orderByDescending`	`boolean`	No	`false`	Order results by descending.

Request example

HTTP

```
GET /exports/evaluations?sessionCount=200&continuationToken=abc123
Authorization: Bearer <token>
```

Response schema

Success response - `(200 OK)`

JSON

```
{
  "workspaceId": "string",
  "workspaceName": "string",
  "tenantId": "string",
  "evaluations": [
    {
      // Evaluation object schema
    }
  ],
  "sessionsContinuationToken": "string?",
  "totalCount": "integer?",
  "sessionCount": "integer"
}
```

Error codes and messages

Expand table

Status Code	Error Code	Error Message
400	Bad Request	Invalid parameters or missing workspace/tenant information
404	Not Found	Admin export APIs not enabled
500	Internal Server Error	Server error during export

Data models

Base response properties

Both export APIs return responses with these common properties:

Expand table

Property	Type	Description
`workspaceId`	`string`	The workspace ID that is exported
`workspaceName`	`string`	The workspace name that is exported
`tenantId`	`string`	The tenant ID that is exported
`sessionsContinuationToken`	`string?`	Token for next page (null if no more data)
`totalCount`	`integer?`	Total count of items in current page
`sessionCount`	`integer`	Session count used for this request

Pagination

Both APIs support cursor-based pagination:

1. **Initial Request:** Make request without `continuationToken`.
2. **Subsequent Requests:** Use `sessionsContinuationToken` from the previous response.
3. **End of Data:** When `sessionsContinuationToken` is null, no more data is available.

Pagination example

First request

HTTP

```
GET /exports/prompts?sessionCount=100
```

Response includes continuationToken

JSON

```
{
  "sessionsContinuationToken": "[{\"compositeToken\":{\"token\":null,\"range\":
{\"min\":\"05C1D1D5378D58\",\"max\":\"05C1D3CFCBB964\"}},\"resumeValues\":[\"2025-08-
01T23:49:27.8981554+00:00\"],\"rid\":\"xQAMAIZUJBs7BEAAAADABQ==\",\"skipCount\":1}]",
  "prompts": [...],
  // ... other properties
}
```

Next request using token (URL encoded)

HTTP

```
GET /exports/prompts?
sessionCount=100&continuationToken=%5B%7B%22compositeToken%22%3A%7B%22token%22%3Anull%
2C%22range%22%3A%7B%22min%22%3A%2205C1D1D5378D58%22%2C%22max%22%3A%2205C1D3CFCBB964%22
%7D%7D%2C%22resumeValues%22%3A%5B%222025-08-
01T23%3A49%3A27.8981554%2B00%3A00%22%5D%2C%22rid%22%3A%22xQAMAIZUJBs7BEAAAADABQ%3D%3D%
22%2C%22skipCount%22%3A1%7D%5D
```

> **Note**
>
> - If your service principal has no subscriptions, `az login` might report that none were found; in this case, include `--allow-no-subscriptions`.
> - For tokens, you can use either `--resource https://api.securitycopilot.microsoft.com` (v1) or `--scope https://api.securitycopilot.microsoft.com/.default` (v2), depending on your auth flow. See, Get access token (CLI) and .default scope.

Related content

Use the following articles for guidance:

- Register an app (portal)
- Create SP (CLI tutorial)
- Add credentials (client secret/cert)
- Create client secret (explicit steps)
- (Alt) Create client secret in portal
- Security Copilot roles & assignments)
- Add Service Principal to a group (portal)
- Add Service Principal to a group (API)
- Sign in as SP (CLI)
- Get access token (CLI)

Last updated on 06/04/2026

Configure Copilot settings with the owner role

Give administrators the ability to manage the data and custom plugin use in Security Copilot.

Prerequisites

You need the Security Copilot owner role to manage owner settings.

For more information, see Understand authentication in Microsoft Security Copilot.

Manage owner settings

Manage the usage and data privacy for Copilot in the owner settings menu in Security Copilot.

- Manage capacity - Manage security compute unit association and creation.
- Help improve Copilot - Help improve Copilot by allowing Microsoft to capture data.
- Logging audit data in Microsoft Purview - Manage logging audit data by configuring whether to allow Microsoft Purview to access, process, copy, and store Customer Data.
- Manage who can upload files - Specify who can upload files. For more information, see File upload.

Manage capacity settings

Manage capacity association, overage units, and geo location evaluation options. Keep in mind, purchasing new security capacity units (SCUs), changing capacity, or associating with a different capacity all require Azure Owner or Contributor access to the capacity resource in the Azure portal.

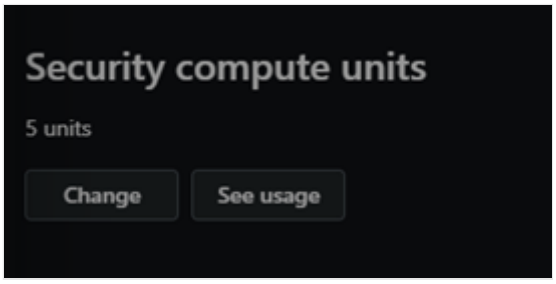

Figure shows owner setting for associating SCUs.

For more information on purchasing SCUs, see Onboarding to Security Copilot.

> **Tip**
>
> Use the Security Copilot Cost Estimator to estimate how many SCUs your organization may need based on the number of Security Copilot users.

Help improve Copilot

The following data sharing options are available:

Expand table

Setting	Description
Allow Microsoft to capture data from Security Copilot to validate product performance using human review	Such validations include but aren't limited to: - Ability of Security Copilot to successfully provide responses to user requests and understand capability gaps that need to be addressed based on user prompts. - Understand the types of tasks customers are using Security Copilot for. - Produce metrics surrounding the usability and quality of responses. - Validate Security Copilot capabilities involving other Microsoft products purchased and integrated by a customer. - Improve responses from plugins accessing other Microsoft products. For more information, see Set up location for prompt evaluation and opt-in (or out of) data sharing.
Allow Microsoft to capture and human review data from Security Copilot to build and validate Microsoft's security AI model	Such validations include but aren't limited to: - Captured data is used to develop security specific models built on top of Azure OpenAI foundational model, which would power more intelligent and personalized capabilities for Security Copilot and other Microsoft products that it integrates with. **NOTE:** Data isn't shared with OpenAI or used to train the Azure OpenAI foundational model.

For more information, see Data sharing and feedback options.

Manage logging audit data in Microsoft Purview

Manage logging audit data by configuring whether to allow Microsoft Purview to access, process, copy, and store Customer Data.

Security Copilot uses Microsoft Purview to process and store data such as administrator actions, user actions, and system responses.

In most scenarios, logs are available within 24 hours in Microsoft Purview after enabling the capability.

> **Important**
>
> Microsoft Purview will store your Customer Data in the region where your Microsoft 365 data is stored. For more information, see, Privacy and data security. The default retention period for audit logs is 180 days, but can be extended using audit log retention policies. For more information, see Manage audit log retention policies.

If you already use Microsoft Purview, no further set up is needed. Otherwise, you'll need to provision a free, limited experience. For more information, see Turn auditing on or off.

For more information, see Audit logging.

Manage who can upload files

As an owner, you can configure who can upload files into Security Copilot to allow Copilot to refer to contents of the file for responses. You can set it so that no one can upload files or contributors and owners can upload files.

Related content

Use the following Security Copilot articles for guidance:

- Manage plugins
- Assign roles

For more information on authentication in Security Copilot, see Understand authentication in Microsoft Security Copilot.

Last updated on 06/04/2026

Manage plugins in Microsoft Security Copilot

The plugin settings page helps owners control the functionality of plugins in Security Copilot.

Manage custom plugins

The following custom plugin options are configurable:

- Control whether contributors can add custom plugins for their sessions
- Control whether contributors can publish custom plugins for everyone in the organization

Owners must set the permissions for adding custom plugins before they can be accessed. By default, only owners can add and manage their own custom plugins. In addition, owners are able to specify who has the ability to add and manage custom plugins for everyone in their organization.

Plugins added by owners can be set to be available only to them or to anyone in the organization.

Plugins added by owners and set to be available to anyone in the organization can be accessed, turned on or off, and used by all Security Copilot users. However, only owners can make modifications to the plugin.

For instructions on creating your own plugins, see Custom plugins.

Options for managing plugins

To manage plugins as an owner:

1. On the left navigation, navigate to Owner. Select **Plugin settings**.

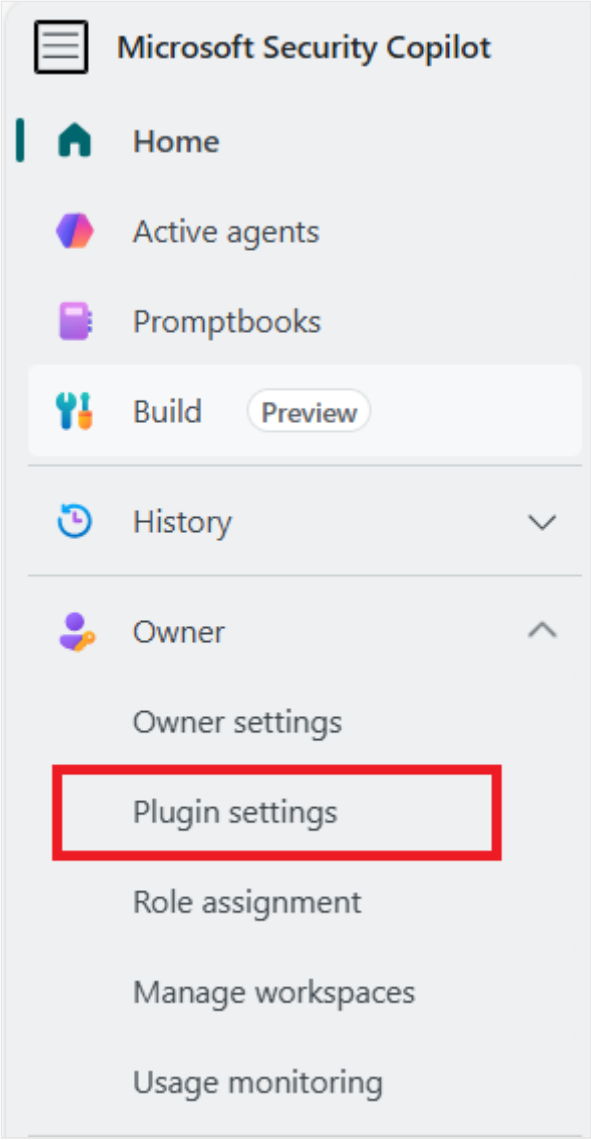

2. Choose who can add and manage their own custom plugins. Select **Owners only** or **Contributors and Owners**.

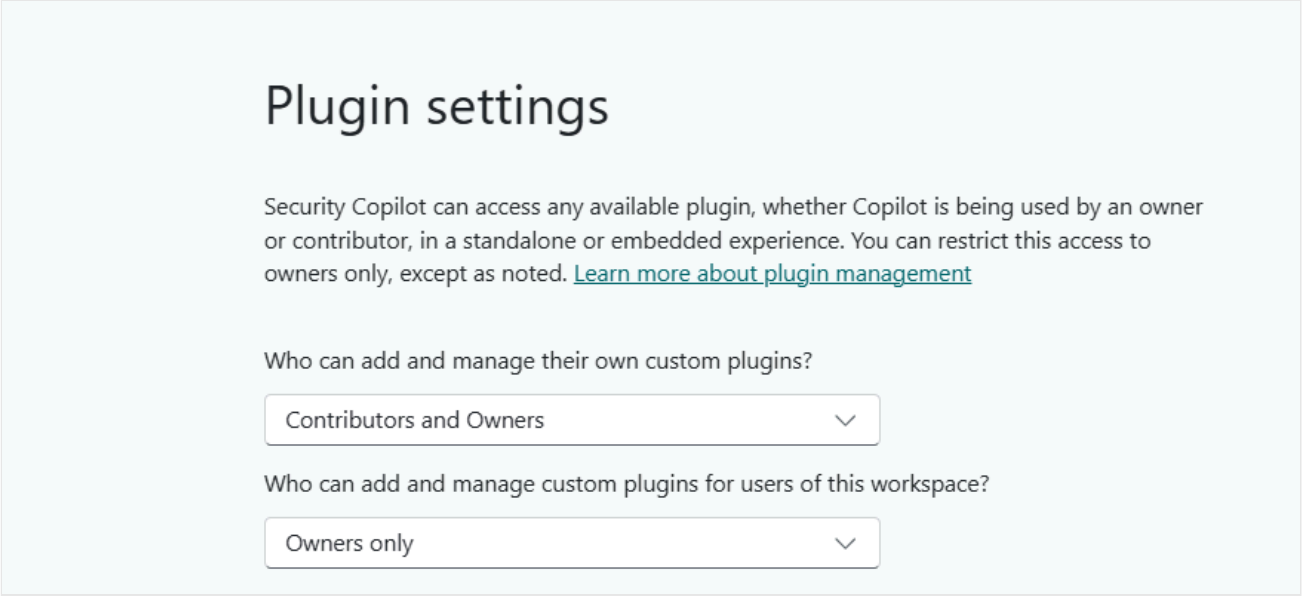

- **Owners only**:

When an owner sets this configuration value, only owners can add and manage custom plugins for themselves (user scope).

- **Owners and Contributors:**

  When an owner sets this configuration value, owners and contributors can add and manage custom plugins for themselves (user scope). This option enables the second permission, "Who can add and manage custom plugins for everyone in the organization?".

3. Choose who can add and manage custom plugins for users in this workspace.

   - **Owners only**:

     When an owner sets this configuration value, only owners can add and manage custom plugins for the organization (tenant scope).

   - **Owners and Contributors**:

     When an owner sets this configuration value, owners and contributors can add and manage custom plugins for the organization (tenant scope).

4. Choose to manage preinstalled plugin availability and restrict access.

   As an administrator, you have full control over which plugins are available across your organization. You can choose to make specific plugins accessible to everyone or restrict them to owners only. These settings apply to both standalone and embedded Security Copilot experiences.

   When you set up a Security Copilot agent the required plugins for that agent is automatically enabled. These plugins are activated solely for that agent and will not change the status of the plugins that you have made available or restricted to your organization. You can disable the agent at any time, giving you complete control over its operation.

Manage preinstalled plugins

By default, all Owners and Contributors have access to preinstalled Microsoft and non-Microsoft plugins. When an owner toggles plugin availability and restricts access, there are two options for plugin availability:

- All users
- Owners only

Once access is restricted, new preinstalled plugins are made available to Owners only until configured otherwise.

> ⚠ **Warning**
>
> Restricting access is an immediate change impacting all users of Security Copilot and embedded experiences. Exercise caution and notify users prior to impact.

Restricted plugins affect embedded experiences. Consider the example where an analyst has access to Security Copilot and Microsoft Defender XDR, but the plugins for Microsoft Defender XDR and Natural Language to KQL are restricted. The analyst experience for Copilot in Defender looks like this:

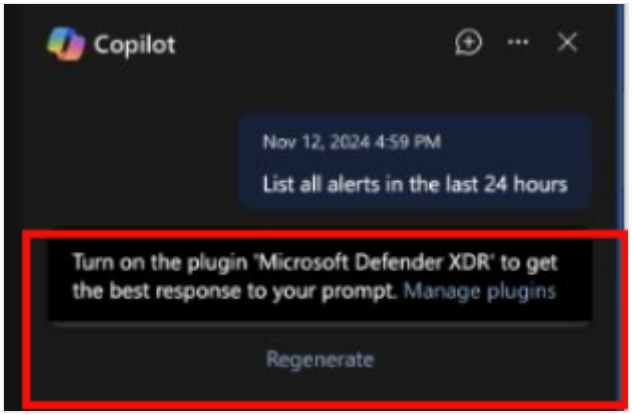

Preinstalled plugins, like **Microsoft Sentinel** and **Azure AI Search**, require more setup. Any plugin with the gear ⚙ or **Set up** button is configured per user. All users who have access to the plugin configure that plugin's setup for themselves.

Manage accessing data from Microsoft 365 services

Allow users to query information from Microsoft 365 services that your organization has licensed. For more information, see Accessing data from Microsoft 365 services.

Add custom plugins

Adding custom plugins extends the capabilities of Security Copilot.

> ⓘ **Important**
>
> You must take precaution and ensure that plugins you create or make available in your organization adhere to responsible AI practices. For more information, see Microsoft AI

principles.

You must also ensure that custom plugins you make available to your organization are fully vetted. It is good practice to vet custom plugins at a user level first, before making plugins available to the entire organization.

> **Note**
>
> This option is only available if an Owner has configured Security Copilot to allow Owners and Contributors to add and manage custom plugins for everyone in the organization.

The steps to add a custom plugin are as follows:

1. Select the **Security Copilot sources** icon.
2. In **Manage plugins**, scroll down to the **Custom** section.
3. Select, **Upload plugin**.

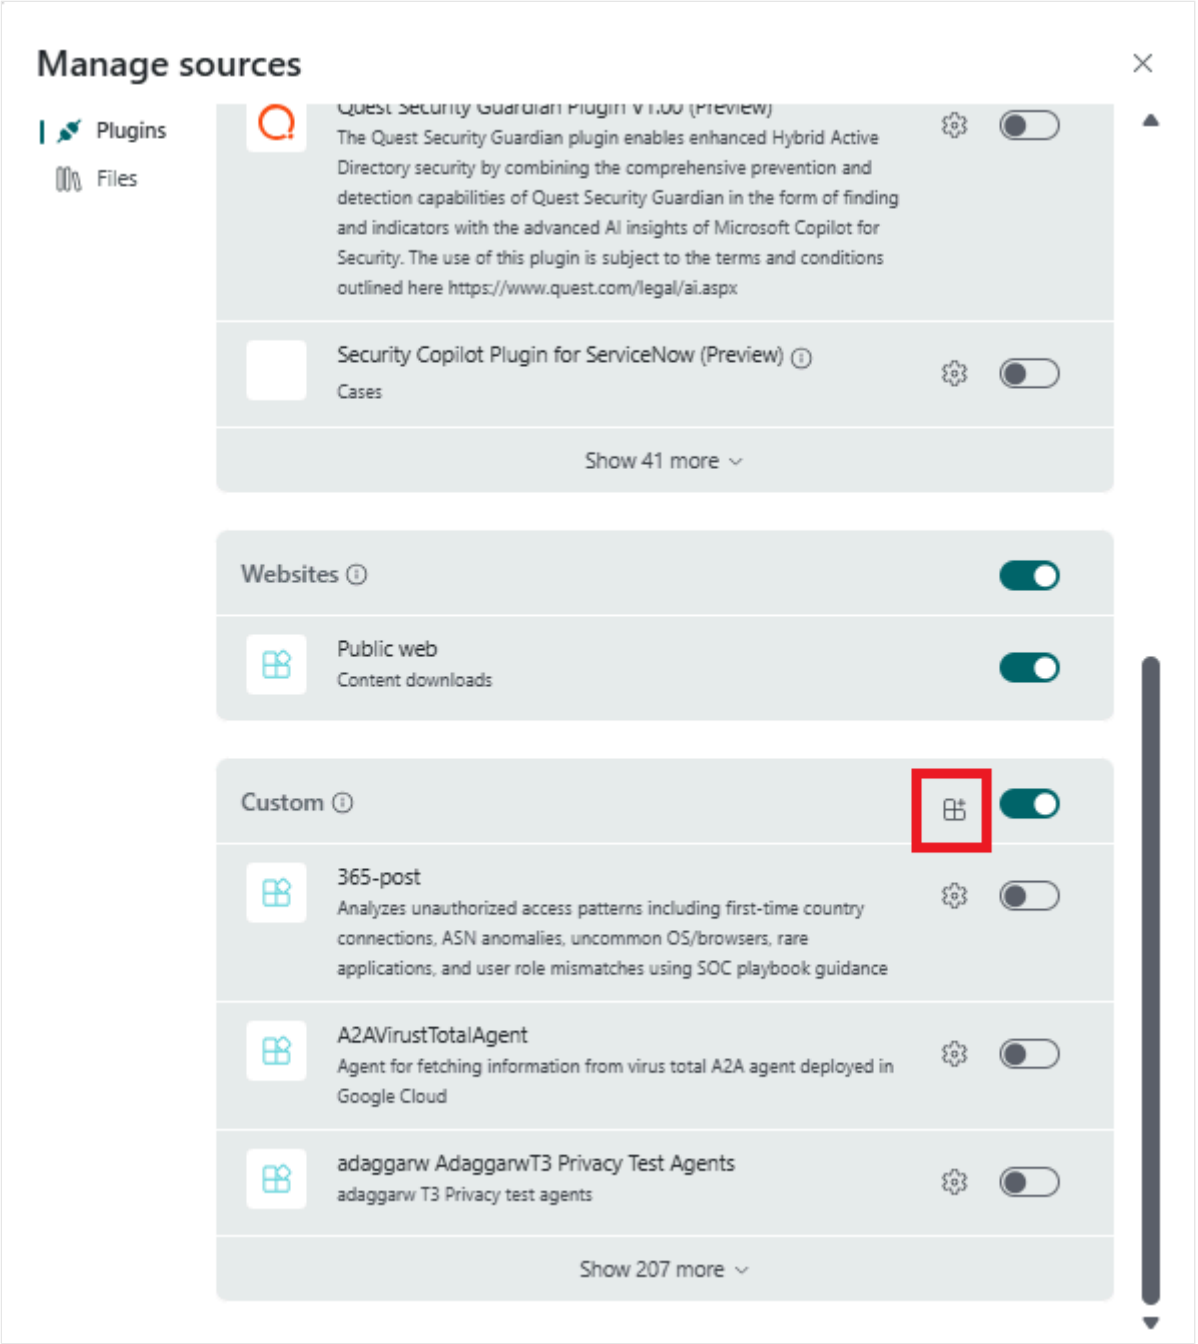

4. From the drop-down options, specify if you want to make the plugin available to just yourself or anyone in the organization. Then choose between adding a Security Copilot plugin or an OpenAI plugin.

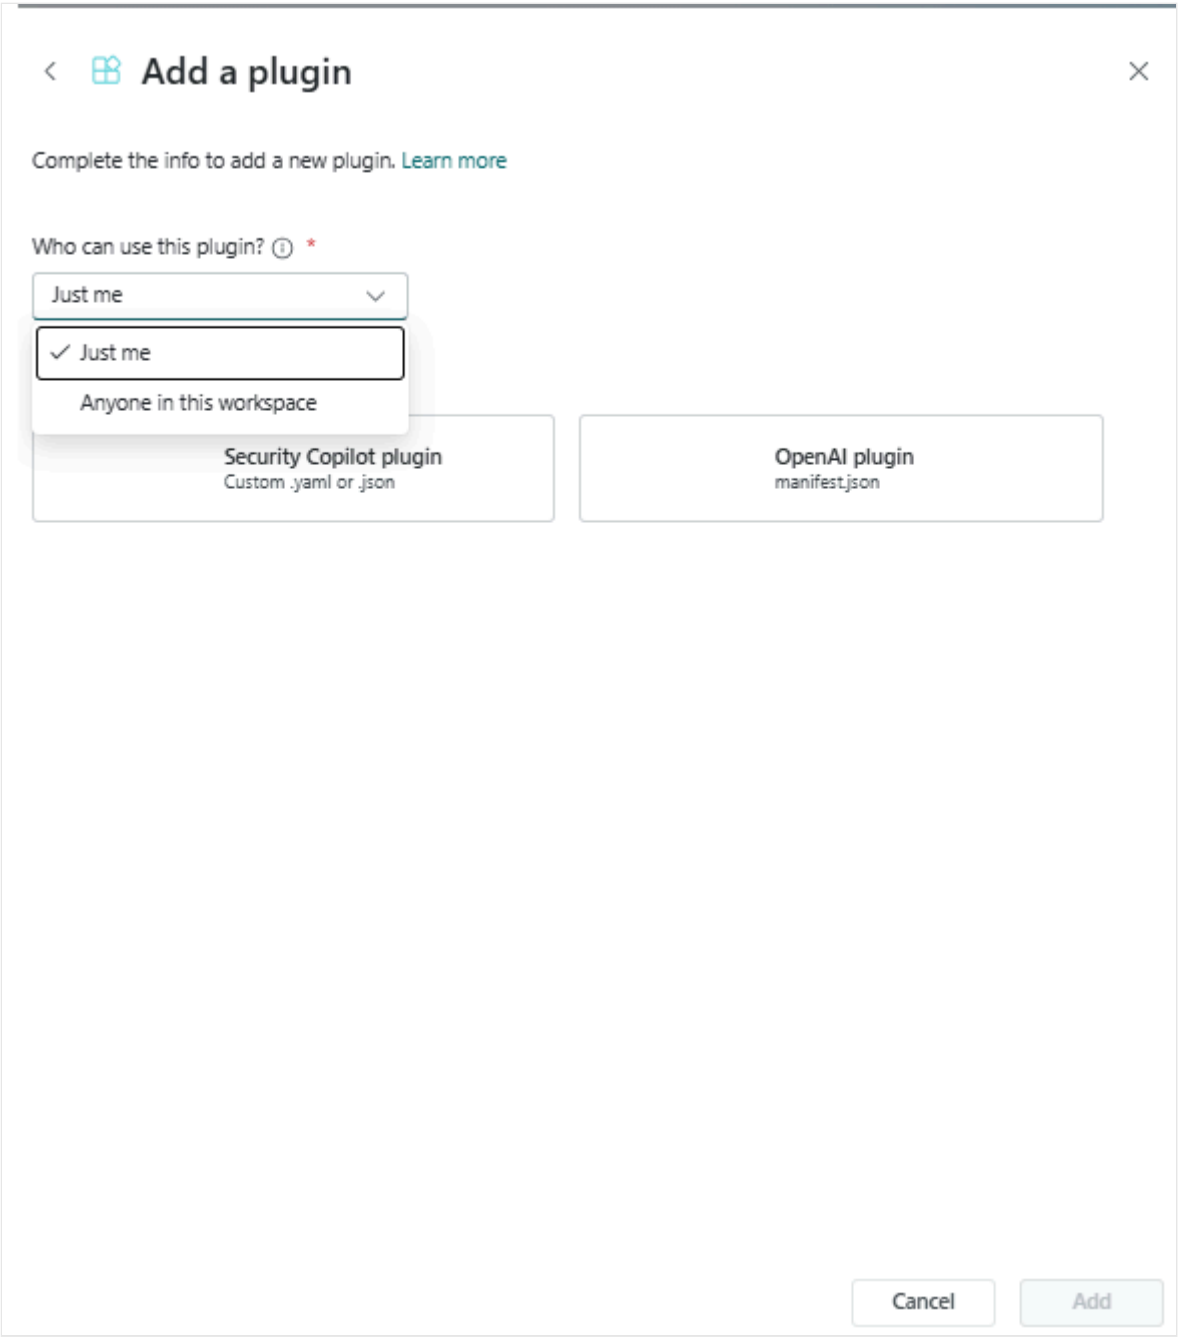

> **Note**
>
> If you choose to make a plugin available to just yourself, a private tag is added to the plugin. If you choose the option of making it available to anyone in the organization, then the plugin will be available to all Security Copilot users.

5. Depending on the plugin you'd like to add, follow the corresponding steps:

Security Copilot plugin

a. Select **Security Copilot plugin**.

b. You can choose to upload a file or upload as a link.

- To upload a file:
  i. Select the **Upload file** section.
  ii. Select the file to upload, then select **Add**.
- To upload a link:
  i. Turn on **Upload as a link** toggle.
  ii. Paste a link to a .yaml or .json file.
  iii. Specify the file type.
  iv. Select **Add**.

6. After the plugin is successfully added, some plugins might need to be set up. You can also choose to set up the plugin later.
7. Provide the appropriate values for the settings fields, then select **Setup**.

   After completing the setup steps, your plugin will be visible in the Custom section. You can turn the toggle on or off.

> ⓘ **Note**
>
> You need to complete the setup steps for the plugin to be available for use and appear in the **Custom** section.

Update a custom plugin

Depending on how owners manage custom plugins, you may be able to update your own custom plugins.

Custom plugins created by owners aren't editable by contributors.

If you need to make changes to the plugin that you added, you can edit the plugin file as follows:

1. Select the **Security Copilot plugin** icon.
2. In the **Custom** section, select the gear icon beside the plugin that you want to update.
3. Select **Edit file** and make the appropriate updates.
4. Select **Save** and confirm that you want to proceed.

Delete a custom plugin

Depending on how owners manage custom plugins, you may be able to delete custom plugins you added as follows:

1. Select the **Security Copilot plugin** icon.
2. In the **Custom** section, select the gear icon beside the plugin that you want to delete.
3. Select **Delete** and confirm that you want to proceed.

Last updated on 05/29/2026

Understand authentication in Microsoft Security Copilot

Copilot uses on-behalf-of authentication to access security-related data through active Microsoft plugins. Specific Security Copilot roles must be assigned in order for a group or individual to access the Security Copilot platform. After you're authenticated to the platform, your Microsoft Entra and Azure Role Based Access Control (RBAC) determines what plugins are available in prompts. Your Security Copilot role controls what other activities you have access to on the platform, such as configuring settings, assigning permissions, and performing tasks.

Security Copilot RBAC roles are not Microsoft Entra roles. Security Copilot roles are defined and managed within Copilot and only grant access to Security Copilot features.

Microsoft Entra RBAC grants access across the Microsoft portfolio of products including services that contain security data. These roles are managed through the Microsoft Entra admin center. For more information, see Assign Microsoft Entra roles to users.

Azure RBAC controls access to Azure resources like Security Capacity Units (SCU) in a resource group, or Microsoft Sentinel enabled workspaces. For more information, see Assign Azure roles.

Access Security Copilot platform

After Security Copilot is onboarded for your organization, configure Security Copilot RBAC to determine users' access to the Security Copilot platform. Then layer your security coverage with conditional access policies. For more information about securing access to the Security Copilot platform beyond RBAC, see Protect AI with Conditional Access policy.

Security Copilot roles

Security Copilot introduces two roles that function like access groups but aren't Microsoft Entra ID roles. Instead, they only control access to the capabilities of the Security Copilot platform and provide no access by themselves to security data.

- Security Copilot owner
- Security Copilot contributor

> **Note**

To prevent accidental removal of owners, Security Copilot enforces retention of two owners at all times. These two owners cannot be removed.

Microsoft Entra, Microsoft Intune, and Microsoft Purview roles

The following Microsoft Entra, Microsoft Intune, and Microsoft Purview roles automatically inherit **Copilot owner** access, ensuring Security Copilot always has at least two owners.

Microsoft Entra roles:

- Billing Administrator
- Entra Compliance Administrator
- Global Administrator
- Intune Administrator
- Security Administrator

For more information on continuity in Microsoft Entra, see Manage emergency access accounts.

Microsoft Intune role:

- The Microsoft Entra ID Intune Administrator

Microsoft Purview roles:

- Purview Compliance Administrator
- Purview Data Governance Administrator
- Purview Organization Management

As part of the Security Copilot inclusion in the Microsoft 365 E5 and E7 license, the Compliance pack is automatically added.

For capacity provisioning

- Azure Contributor or Owner role on subscription/resource group
- Security Administrator or higher in the tenant you're onboarding Security Copilot.

These roles are needed to attach the provision of SCU capacities to the Security Copilot workspace.

Recommended Security Copilot owner roles

You can update the set of Security Copilot at any time. The list of recommended roles includes Microsoft Entra ID roles and some service specific roles, like Microsoft Purview roles and Microsoft Defender roles.

> **Note**
>
> Only the selected owner roles will inherit the owner role on that workspace. For more information on the supported roles see, Microsoft Entra or Microsoft Purview role.

If a Microsoft plugin requires a different role to access its security data, you need to ensure that the appropriate service role is assigned as well. For example, the bundle works well for an analyst assigned the **Compliance Administrator** role. This role is one of the Microsoft Purview recommended roles and gives them access to Microsoft Purview plugin data. That same analyst needs additional role assignments to access security data like Microsoft Sentinel though. For more examples, see access the capabilities of Microsoft plugins.

Recommended Security Copilot contributor roles

Provision access to Security Copilot with the **Recommended Microsoft Security roles** group, which uses a balanced approach to security and administrative efficiency. Users who already have security permission through Microsoft Entra roles are given contributor access to the Security Copilot platform with this bundle. If the **Everyone** group is assigned, you must remove it before you add the recommended security roles.

Expand table

Role assignment	Advantage	Weakness
Recommended Microsoft Security roles (Default)	Quick, secure way to give users in your organization who already have access to security data the access they need to the platform. More roles, users and groups can still be added.	The group is all or nothing. If there's a role in the recommended group you don't want to have access, the entire group must be removed.
Everyone	This group was initially given contributor access by default to simplify provisioning.	If a user in the everyone group doesn't have access to any security data, the experience in Security Copilot is confusing.
Custom	Complete control over users and groups with access to the platform.	Requires more administrative complexity.

> **Note**
>
> The assignment of the recommended Microsoft security roles to Copilot is the default for new Security Copilot instances. Existing customers who have the everyone group assigned may continue to use the assignment. If the everyone group is removed though, it can't be assigned again. Security Copilot only supports assigning permissions to Role-assignable groups. For more information about these groups, see Creating role-assignable groups in Entra ID.

Access the capabilities of Microsoft plugins

Security Copilot doesn't go beyond the access you have, aligning with Microsoft's Security and privacy RAI principle. Each Microsoft plugin has its own role requirements that remain in effect to access the plugin's service and its data. Verify that you have the proper roles and licenses assigned to use the capabilities of the Microsoft plugins that are activated.

Consider these examples:

1. **Copilot contributor**

   As an analyst, you're assigned the **Copilot contributor** role for access to the Copilot platform and the ability to create sessions. This assignment by itself doesn't give you access to your organization's security data. Following the least privilege model, you don't have any sensitive Microsoft Entra roles like **Security Administrator**. In order to use Copilot to access security data like the Microsoft Sentinel plugin, you still need an appropriate Azure RBAC role like **Microsoft Sentinel Reader**. This role gives you access to incidents in your Microsoft Sentinel workspace from Copilot. For your Copilot session to access the devices, policies, and postures available through the Intune plugin, you need an Intune role like the **Endpoint Security Manager**. The same pattern is true for access to Microsoft Defender XDR data through your Copilot session or the embedded Security Copilot experiences.

   For more information on service specific RBAC, see the following articles:

   - Microsoft Sentinel roles and permissions
   - Microsoft Intune RBAC
   - Microsoft Defender XDR Unified RBAC

2. **Microsoft Entra security group**

Although the **Security Administrator** role inherits access to Copilot and certain plugin capabilities, this role includes PRIVILEGED permissions. Don't assign users this role purely for Copilot access. Instead, create a security group and add that group to the appropriate Copilot role (Owner or Contributor).

For more information, see Best practices for Microsoft Entra roles.

3. **Recommended Microsoft Security roles**

   An analyst with access to Microsoft's unified security operations platform has a custom Defender XDR role assigned that gives them access to multiple workloads in the Defender portal. If the custom role security operations include Security Copilot, then the **Recommende Microsoft Security roles** gives them access to the Security Copilot platform as well as Copilot in Microsoft Defender, simplifying security management.

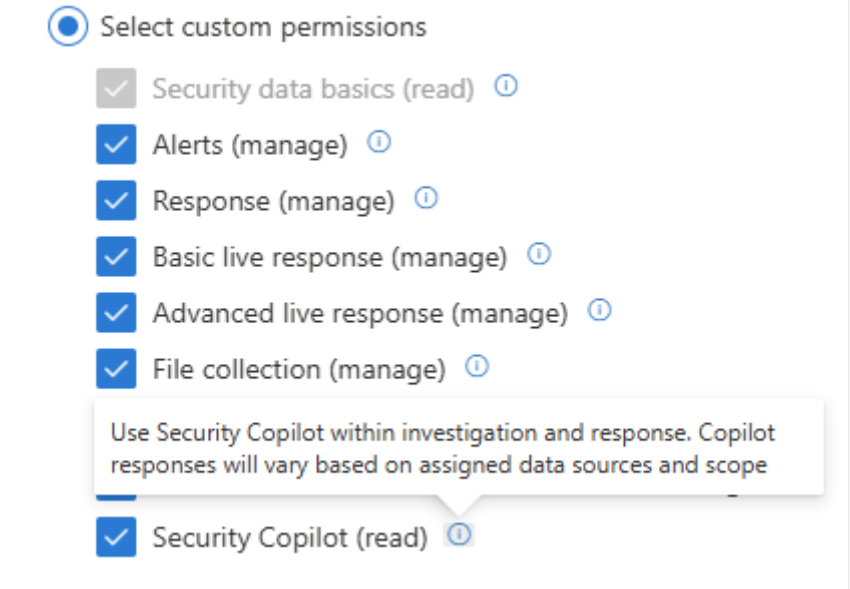

Access embedded experiences

In addition to the Copilot contributor role, verify the requirements for each Security Copilot embedded experience to understand what extra roles and licenses are required.

For more information, see Security Copilot experiences.

Assign roles

The following table illustrates the default access granted to starting roles.

> ⓘ **Note**

Some organizations may still have the **Everyone** group assigned to **Copilot contributor** access. Consider replacing this broad access with **Recommended Microsoft Security roles** group.

Expand table

Capability	Copilot owner	Copilot contributor
Create sessions	Yes	Yes
Manage personal custom plugins	Yes	Default No
Allow contributors to manage personal custom plugins	Yes	No
Allow contributors to publish custom plugins for the tenant	Yes	No
Change availability of preinstalled plugins for the tenant	Yes	No
Upload files	Yes	Default Yes
Manage upload file usage	Yes	No
Run promptbooks	Yes	Yes
Manage personal promptbooks	Yes	Yes
Share promptbooks with tenant	Yes	Yes
Update data sharing and feedback options	Yes	No
Capacity management	Yes*	No
View usage dashboard	Yes	No
Select language	Yes	Yes

Assign Security Copilot access

Assign Copilot roles within Security Copilot settings.

1. Select the home menu.
2. Select **Role assignment** > **Add members**.
3. Start typing the name of the person or group in the **Add members** dialog box.
4. Select the person or group.
5. Select the Security Copilot role to assign (Copilot owner or Copilot contributor).
6. Select **Add**.

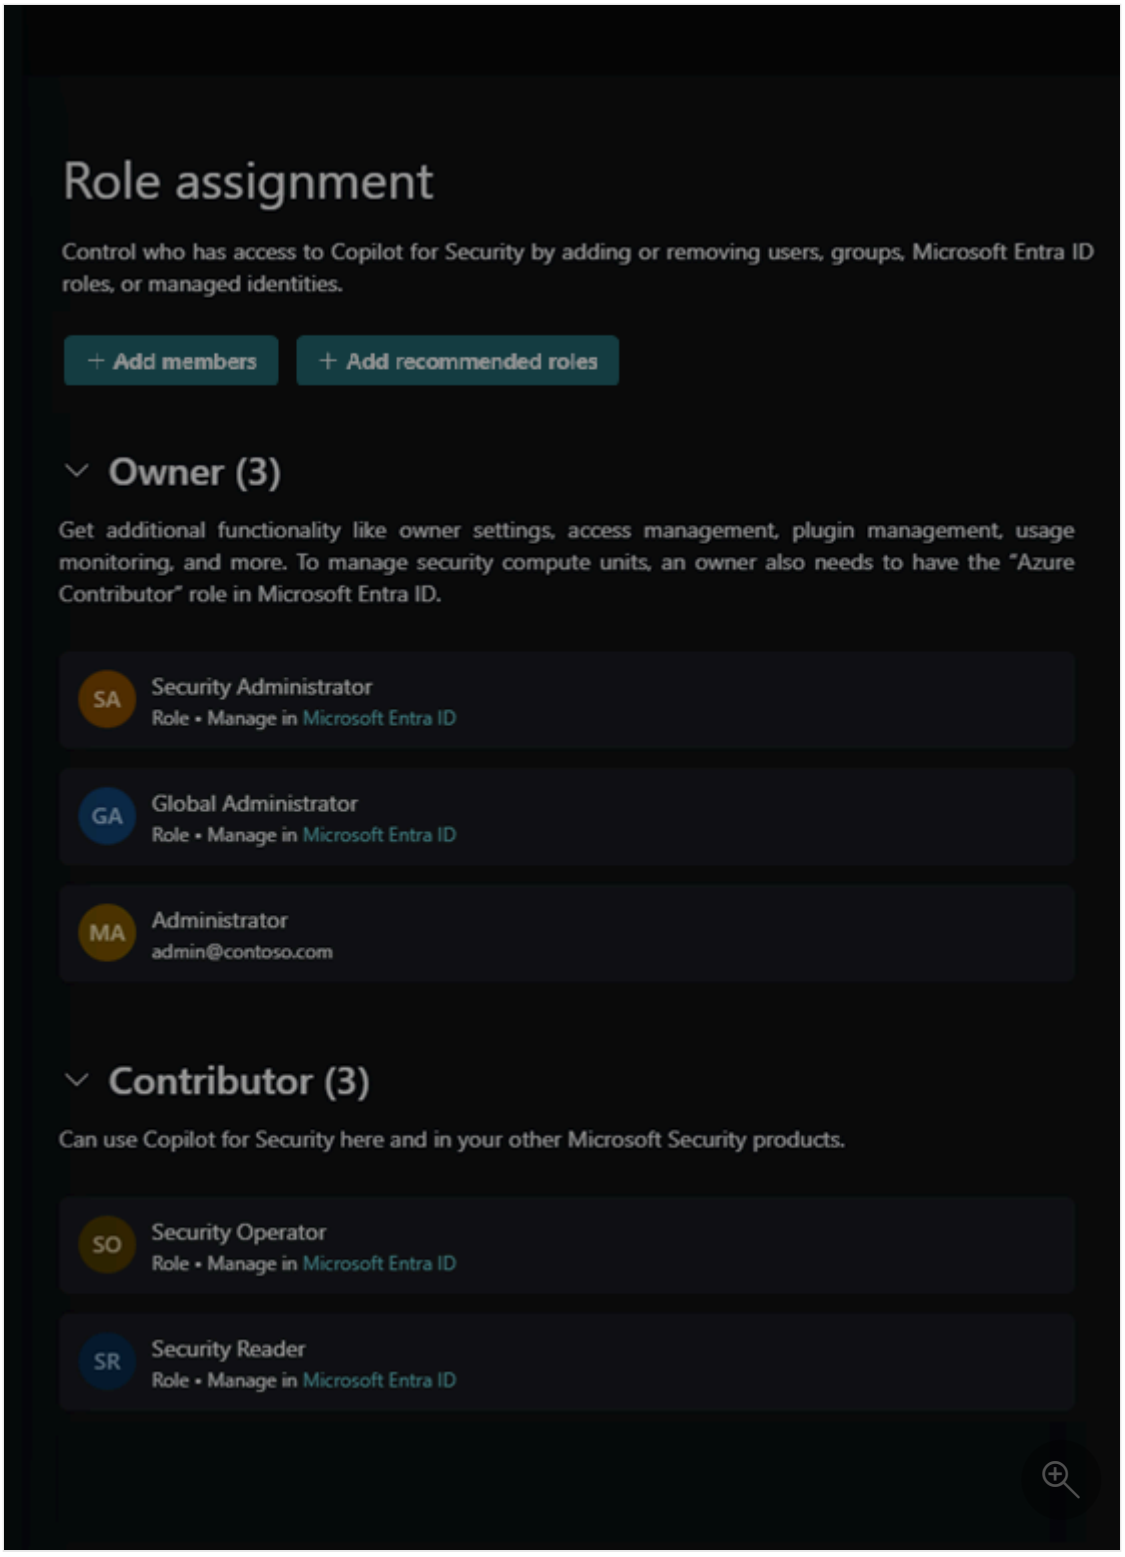

Tip

We recommend using security groups to assign Security Copilot roles instead of individual users. This reduces administrative complexity.

The **Everyone** group is removable from Contributor access. Consider adding the **recommended roles** after you remove the everyone group.

Microsoft Entra role membership is only manageable from the Microsoft Entra admin center. For more information, see Manage Microsoft Entra user roles.

Shared sessions

Copilot contributor role is the only requirement for sharing a session link or viewing it from that tenant.

When you share a session link, consider these access implications:

- Security Copilot needs to access a plugin's service and data to generate a response, but that same access isn't evaluated when viewing the shared session. For example, if you have access to devices and policies in Intune, and the Intune plugin is utilized to generate a response that you share, the recipient of the shared session link doesn't need Intune access to view the full results of the session.
- A shared session contains all the prompts and responses included in the session, whether it was shared after the first prompt or the last.
- Only the user that creates a session controls which Copilot users can access that session. If you receive a link for a shared session from the session creator, you have access. If you forward that link to someone else, it doesn't grant them access.
- Shared sessions are read-only.
- Sessions can only be shared with users in the same tenant that have access to Copilot.
- Some regions don't support session sharing via email.
  - `SouthAfricaNorth`
  - `UAENorth`

For more information on shared sessions, see Navigating Security Copilot.

Multitenant

If your organization has multiple tenants or you leverage a managed service provider partner (MSSP), Security Copilot can accommodate authentication across them to access security data where Security Copilot is provisioned. The tenant that is provisioned for Security Copilot doesn’t need to be the tenant your security analysts logs in from. We currently support the authentication models below through different access approaches.

- Azure B2B & Guest Accounts are able to leverage “Tenant Switching” which is natively built into the Security Copilot standalone portal. For more information, see Navigating Security Copilot tenant switching.

- Granular Delegated Admin Privileges (GDAP) through the Partner Center. For more information, see Grant MSSPs access to Microsoft Security Copilot.
- Azure Lighthouse through the Azure Portal. For more information, see Grant MSSPs access to Microsoft Security Copilot.

Tenant-switching is available through the Security Copilot standalone portal but only supports B2B collaboration accounts or guest account. For more information, see Navigating Security Copilot tenant switching.

Cross tenant sign-in example

Contoso recently merged with Fabrikam. Both tenants have security analysts, but only Contoso purchased and provisioned Security Copilot. Angus MacGregor, an analyst from Fabrikam wants to use their Fabrikam credential to use Security Copilot. Here are the steps to accomplish this access:

1. Ensure Angus MacGregor's Fabrikam account has an external member account in the Contoso tenant.
2. Assign the external member account the necessary roles to access Security Copilot and the desired Microsoft plugins.
3. Sign in to the Security Copilot portal with the Fabrikam account.
4. Switch tenants to Contoso.

For more information, see Grant MSSP access.

Manage promptbooks

Promptbook creation is available to all roles, including the ability to publish a custom promptbook for the tenant. Choose whether to publish a promptbook for yourself or the tenant at the time of creation.

For more information, see Build your own promptbook.

Preinstalled plugin authentication

Preinstalled plugins, like **Microsoft Sentinel** and **Azure AI Search**, require more setup. The plugin provider determines the type of authentication. Any plugin with the gear or **Set up** button is configured per user. Whether a preinstalled plugin is restricted, all users who have access to the plugin configure that plugin's setup for themselves.

> **Note**
>
> Website plugins use anonymous authentication to access content.

For more information, see Manage preinstalled plugins.

Last updated on 06/04/2026

Manage security compute unit usage in Security Copilot

The usage monitoring dashboard provides a comprehensive view into various data dimensions to help you keep track of Security Compute Units (SCUs) usage in Security Copilot. You have visibility into the number of provisioned and overage units used, the specific plugins employed during sessions, and the initiators of those sessions. The dashboard also allows you to apply filters and export usage data seamlessly. The dashboard includes up to 90 days of data, offering a robust window into recent activity.

> **Tip**
>
> To understand what Security Compute Units are and the available capacity models, see Understand Security Compute Units and capacity.

Monitor security compute units usage

Copilot owners can view the usage over time within the Microsoft Security Copilot usage monitoring dashboard.

To view the usage dashboard:

1. Sign in to Security Copilot (https://securitycopilot.microsoft.com) and go to the **Home menu** > **Owner settings**.
2. Select **Usage monitoring**.

The usage dashboard shows how many SCUs are consumed over a period of time by your Microsoft Security Copilot workloads.

Understand the data dimensions

Use the data dimensions to get a granular view of how SCUs are being used.

Capacities are captured on the tabs at the top of the dashboard. Select a capacity to view the usage information related to that specific capacity.

You can toggle between the various capacity resources that you own. You see information such as provisioned units used, overage units used, and the workspace that's using the capacity.

The following table provides information on what each column represents and helps you interpret data accurately:

Expand table

Column	Description	Examples/Types
Session ID	A unique identifier for each Copilot session.	a123b456-789-1011-1c2-d12345678e99
Date	Date when the user initiated the session, displayed in local time zone.	Sept 10, 08:30 AM
Units used	Number of SCUs used in the session.	0.6
Initiated by	The user that initiated the session.	Ben Heppard
Category	The format in which session was made. There are three category types: - Prompt - Promptbook - Agent	**Prompt**: A single prompt. **Promptbook**: A session related to a promptbook. For example, a user initiates "Script Analysis" promptbooks in which five prompts are called. Each prompt from promptbook call has a category as "promptbook". **Agent**: An interactive or a fully autonomous computational entity that works to achieve a specific goal. It's defined by its ability to perceive and react to its environment, make decisions, and produce outcomes to meet its objective.
Type	The method used for the session. There are two invocation types: - Manual Action - Automated Action.	**Manual Action**: User triggered calls. A user must manually initiate. For example, a user types and submits prompt in prompt bar. **Automated Action**: Calls that are automatically initiated/scheduled. For example, incident summary or scheduled logic app.
Copilot experience	The user interface where the session was initiated from. Users can use the standalone interface, an embedded interface, or through LogicApps.	- Security Copilot portal (Standalone experience) - Copilot in Microsoft products (Embedded experience); for example, inside Microsoft Defender XDR.

Column	Description	Examples/Types
		- Azure Logic Apps - Users can initiate Security Copilot workflows using Azure Logic Apps.
Plugin used	The plugin used in the session.	Microsoft Entra, Microsoft Defender XDR.
Status	Indicates if SCUs are used.	All available units used.

Apply filters

Use filters to narrow down the displayed information in the dashboard. Applying filters can help you analyze relevant data without reviewing the entire data set.

> **Note**
>
> Currently, filters are only applied to the table and not the bar chart.

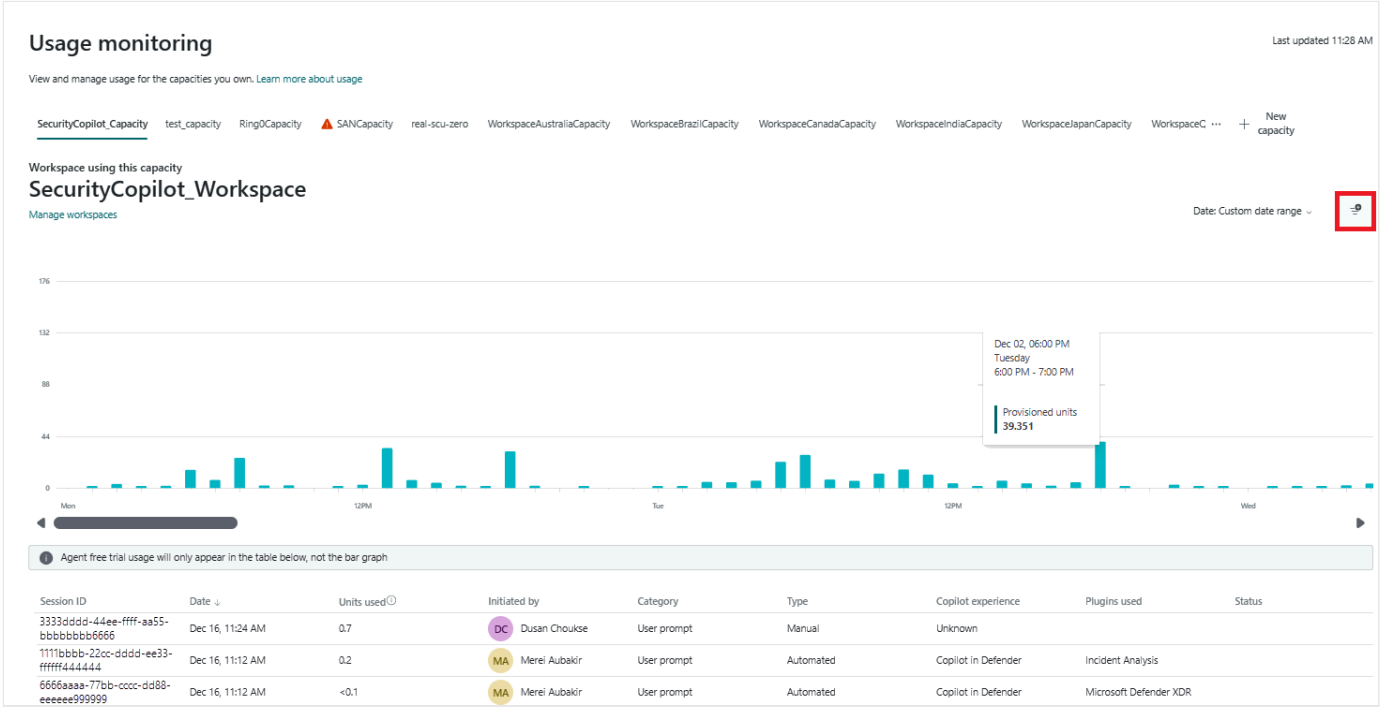

The following filters are available:

- Copilot experience
- Users
- Plugins used
- Type
- Category

> **Important**
>
> Selecting a date range impacts the data shown in the table. For instance, choosing **Last 24 hours** and then selecting a filter, such as a specific user, will display information related to that user from only the past 24 hours.

Selecting a bar in the chart filters the table to show information for that specific day.

Export usage dashboard data

Export usage dashboard data into Excel files for easier analysis, sharing, or offline access. Select the **Export** button on the corner of the dashboard.

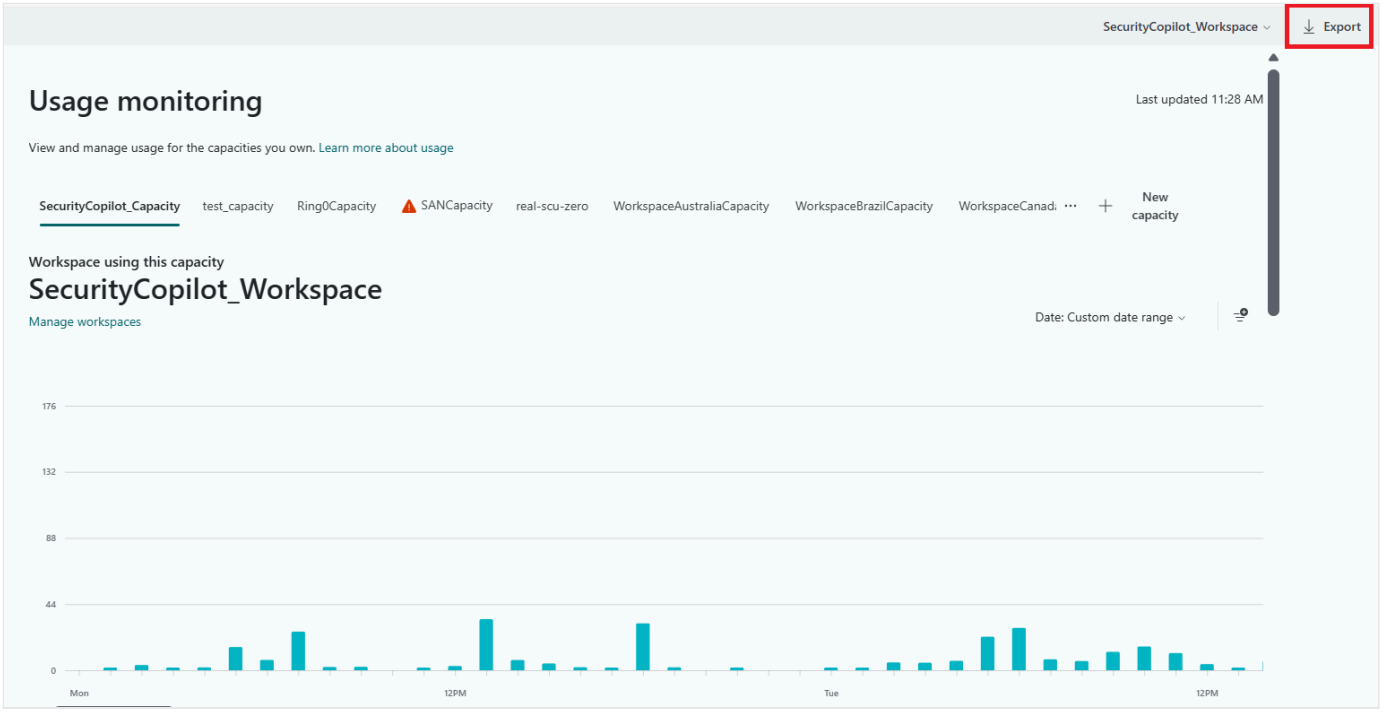

Nearing and crossing security compute unit limit

When the usage in your organization is nearing the limit, you're notified of it when submitting the prompts. To avoid disruption, you can contact the Azure capacity owner or contributor to increase the security compute units or limit the number of prompts.

Experience when nearing security compute unit limit

When you're in the middle of an investigation and the usage is nearing the provisioned capacity limit, a notification is displayed in response to a prompt.

The capacity limit notifications are also shown in Security Copilot embedded experiences and it informs the analyst to contact the Azure capacity owner or contributor to increase the capacity or limit the number of prompts to avoid disruptions.

The notification allows the analyst to react before they run out of capacity.

Experience when security compute unit limit is exceeded

When the provisioned and overage SCUs is crossed, the analyst will see an error message stating that due to high usage in the organization Copilot can't respond to requests. Analysts can't submit more prompts at this time. More capacity would become available in the next hour.

The analyst is asked to contact the Azure capacity owner or contributor to increase the provisioned SCUs.

The error message view only appears when your organization reaches 100% of capacity.

Update the provisioned and overage security compute units (SCUs)

You can update provisioned security compute units and overage units in the Azure portal and the Security Copilot portal. You also have the option to view billing in Azure while updating SCUs.

Azure portal required role: Azure capacity owner or contributor.

Security Copilot portal required role: Azure capacity owner or contributor who are also Security Copilot owners.

> **Note**
>
> Users can expect capacity adjustments to take effect within 30 minutes.

Process to update provisioned and overage security compute units (SCUs)

You can update the provisioned and overage SCUs as follows:

1. Sign in to Security Copilot (https://securitycopilot.microsoft.com) and go to **Settings**.
2. There are two ways you can update the provisioned and overage SCUs:

- Through the owner settings page: Under **Security compute units** or **Number of overage units**, select **Change**.
- Through the usage dashboard: Select **Change units** in the main dashboard.

In either option, update the number of security compute units (SCUs) for the provisioned and overage units. You also have the option to view billing in Azure by selecting **View billing in Azure**.

As the number of SCUs are increased or decreased, the estimated monthly cost changes accordingly, assuming the provisioned capacity remains the same 24 hours a day for the whole month.

- Through the owner settings page: Under **Security compute units** or **Number of overage units**, select **Change**.
- Through the usage dashboard: Select **Change units** in the main dashboard.

In either option, update the number of security compute units (SCUs) for the provisioned and overage units. You also have the option to view billing in Azure by selecting **View billing in Azure**.

As the number of SCUs is increased or decreased, the estimated monthly cost changes accordingly, assuming the provisioned capacity remains the same 24 hours a day for the whole month.

Offboard and delete capacity

To offboard from Security Copilot, you'll need to delete the provisioned capacity. You need to be at least a Security Administrator role to accomplish this task.

> ⓘ **Note**
>
> To export Security Copilot data, you will need to contact Security Copilot support. For more information, see Contact Security Copilot support.

Delete capacity from the Owner settings

> ⚠ **Warning**
>
> Deleting capacity and their internal data is a permanent action and cannot be undone.

You can delete capacity from the Owner settings page or the usage monitoring page.

1. Sign in to Security Copilot (https://securitycopilot.microsoft.com).
2. Select the **Home menu** icon.
3. Navigate to the **Owner settings** or **Usage monitoring** section.
4. In the units section, select **Change**.
5. Select the overflow menu (**...**).
6. Select **Delete the capacity**.
7. Confirm that you want to delete capacity. This action deletes the active capacity for the tenant.

Understand SCU charges by capability

SCUs are charged based on the capability release stage:

Expand table

Capability stage	Charged?
Generally available (GA)	Yes
Public preview	Yes
Private preview	No

Last updated on 06/04/2026

Microsoft Security Copilot Data and Compliance Frequently Asked Questions

This article provides a reference for how Microsoft Security Copilot handles, stores, and protects Customer Data. It highlights key data residency behaviors, security and privacy commitments, eligibility considerations, retention practices, and how Security Copilot complies with Microsoft policies and industry regulations.

FAQ

What is Security Copilot's Microsoft 365 E5 and E7 data handling policies?

All existing Security Copilot customers' data storage and handling will continue in the location where their current workspace resides. For more information, see Privacy and data security in Microsoft Security Copilot.

Can my organization elect to disable Security Copilot at this time?

Existing customers can remove Security Copilot capacity. For more information, see Delete capacity through Security Copilot.

For customers who do not currently have access to Security Copilot through Microsoft 365 E5 and E7, details will be shared once it is included in their Microsoft 365 E5 and E7 license.

Is Customer Data used to train Azure OpenAI Service foundation models?

No, Customer Data isn’t used to train Azure OpenAI Service foundation models, and this commitment is documented in our Product Terms. For more information on data sharing in the context of Security Copilot, see Privacy and data security.

What is the GDPR Guidance for EU Markets?

Microsoft complies with all laws and regulations applicable to its providing the Products and Service including security breach notification law and Data Protection Requirements (as defined

in the Microsoft DPA). However, Microsoft isn't responsible for compliance with any laws or regulations applicable to Customer or Customer's industry that aren't generally applicable to information technology service providers. Microsoft doesn't determine whether Customer's data includes information subject to any specific law or regulation. For more information, see Microsoft Products and Services Data Protection Addendum (DPA).

Are US Government Cloud (GCC) customers eligible?

Currently, Security Copilot isn't designed for use by customers using US government clouds, including but not limited to GCC, GCC High, DoD, and Microsoft Azure Government. For more information, see with your Microsoft representative.

Are US and Canada health care customers eligible?

US and Canada HLS customers are eligible to purchase Security Copilot. Microsoft Security Copilot is now listed and covered by Business Associate Agreement ("BAA"), which is important to healthcare providers who are subject to regulations under HIPAA. For additional information on compliance offerings currently covered for Microsoft Security Copilot can be found in the Service Trust Portal.

How do I export or delete data from Security Copilot?

You will need to contact support. For more information, see Contact support.

Where can I find more information on Data Protection and Privacy?

You can learn more at the Microsoft Trust Center.

The Azure OpenAI Service code of conduct includes "Responsible AI Mitigation Requirements". How do those requirements apply to Security Copilot customers?

These requirements don't apply to Security Copilot customers because Security Copilot implements these mitigations.

Why does Microsoft Copilot transfer data to a Microsoft tenant?

Microsoft Copilot is a SaaS (Software as a Service) offering that runs in the Azure production tenant. Users enter prompts and Security Copilot provides responses based on the insights sourced from other products such as Microsoft Defender XDR, Microsoft Sentinel, and Microsoft Intune. Security Copilot stores past prompts and responses for a user. The user can use the in-product experience to access prompts and responses. Data from a customer is logically isolated from the data of other customers. This data doesn't leave the Azure production tenant and is stored until customers ask to delete them or offboard from the product.

How is the transferred data secured in transit and at rest?

The data is encrypted both in transit and at rest as described in the Microsoft Products and Services Data Protection Addendum.

How is the transferred data protected from unauthorized access and what testing was done for this scenario?

By default, no human users have access to the database and the network access is restricted to the private network where the Microsoft Copilot application is deployed. If a human needs access to respond to an incident, then the on-call engineer needs elevated access and network access approved by authorized Microsoft employees.

Apart from regular feature testing, Microsoft also completed penetration testing. Microsoft Security Copilot complies with all the Microsoft Privacy, security and compliance requirements.

In "My Sessions" when an individual session is deleted, what happens to the session data?

Session data is stored for runtime purposes (to operate the service), and also in logs. In the runtime database, when a session is deleted via the in-product UX, all data associated with that session is marked as deleted and the time to live (TTL) is set to 30 days. After that TTL expires, queries can't access that data. A background process physically deletes the data after that time. In addition to the 'live' runtime database, there are periodic database backups. The backups will age out – these have short-lived retention periods (currently set to four days).

Logs, which contain session data aren't affected when a session is deleted via the in-product UX. These logs have a retention period of up to 90 days.

What Product Terms apply to Security Copilot? Is Security Copilot a "Microsoft Generative AI Service" within the meaning

of Microsoft's Product Terms?

The following Product Terms govern Security Copilot customers:

- Universal License Terms for Online Services terms in the Product Terms, which include the Microsoft Generative AI Services terms and the Customer Copyright Commitment.
- Privacy & Security Terms in the Microsoft Product Terms, which include the Data Protection Addendum.

Security Copilot is a Generative AI Service within the definition of the Product Terms. Additionally, Security Copilot is a "Covered Product" for purposes of the Customer Copyright Commitment. At this time, in the Product Terms there are no product-specific terms unique to Security Copilot. For more information, see Microsoft Azure Legal Information.

In addition to the Product Terms, customers' MBSA/EA and MCA agreements, for example, govern the parties' relationship. If a customer has specific questions about its agreements with Microsoft, engage the CE, the deal manager, or the local CELA supporting the deal.

What is the Microsoft Copilot Copyright Commitment?

The Microsoft Customer Copyright Commitment is a new commitment that extends Microsoft's existing intellectual property indemnity support to certain commercial Copilot services. The Customer Copyright Commitment applies to Security Copilot. If a third party sues a commercial customer for copyright infringement for using Microsoft's Copilots or the output they generate, Microsoft will defend the customer and pay the amount of any adverse judgments or settlements that result from the lawsuit, provided that the customer used the guardrails and content filters built into our products.

Can Security Copilot customers opt out of Azure OpenAI Service abuse monitoring? Does Security Copilot engage in any content filtering or abuse monitoring?

Azure OpenAI abuse monitoring is currently disabled service-wide for all customers.

Does Security Copilot make any location of data processing or data residency commitments?

For more information on Customer Data storage location and processing, see Privacy and data security.

Is Security Copilot a Microsoft EU Data Boundary service?

At the time of GA, all Microsoft Security Services are out of scope for EU data residency requirements and Security Copilot won't be listed as an EUDB service.

Where is EU Customer Data stored?

Security Copilot stores Customer Data and Personal Data such as user prompts and Microsoft Entra Object IDs in the tenant Geo. If a customer provisions their tenant in the EU and isn't opted in to data sharing, all Customer Data and pseudonymized personal data are stored at rest within the EU. Processing of Customer Data and Personal Data prompts can occur in the designated Security GPU Geo. For more information on Security GPU geography selection, see Get Started with Security Copilot. If a customer is opted in to data sharing, prompts can be stored outside of the EU Data Boundary. For more information on data sharing, see Privacy and data security..

Are customer prompts (such as input content from the customer) considered Customer Data within the terms of the DPA and the Product Terms?

Yes, customer prompts are considered Customer Data. Under the Product Terms, customer prompts are considered Inputs. Inputs are defined as "all Customer Data that Customer provides, designates, selects, or inputs for use by a generative artificial intelligence technology to generate or customize an output".

Is "Output Content" considered Customer Data within the terms of the DPA and the Product Terms?

Yes, Output Content is Customer Data under the Product Terms.

Is there a transparency note or transparency documentation for Security Copilot?

Yes, the Responsible AI transparency document can be found here: Responsible AI FAQ.

What are the Compliance Offerings for Microsoft Security Copilot?

Microsoft Security Copilot is dedicated to upholding the highest standards of security, privacy, and operational excellence, as demonstrated by its extensive array of industry certifications.

These include ISO 27001 for information security management, ISO 27018 for the protection of personal data in the cloud, ISO 27017 for cloud-specific security controls, and ISO 27701 for privacy information management.

Additionally, Security Copilot holds certifications for ISO 20000-1 in IT service management, ISO 9001 in quality management, and ISO 22301 in business continuity management. It also complies with SOC2 requirements for security, availability, and confidentiality, underscoring our commitment to delivering secure and reliable services. For healthcare-related services, Security Copilot is certified under the HiTrust CSF framework, further enhancing its security and compliance stance, and is covered by HIPAA Business Associate Agreements (BAA), ensuring adherence to healthcare regulations and the protection of sensitive health information.

For more information on compliance offerings currently covered for Microsoft Security Copilot see, the Service Trust Portal.

Last updated on 05/28/2026

Privacy and data security in Microsoft Security Copilot

> **Important**
>
> The privacy and data security settings described on this page apply to customers who are not using a Microsoft 365 E5 license. If you are a Microsoft 365 E5 customer, note that your default settings differ from those described here, but this page still includes important information about how those settings work. For more information, see Auto provisioned for Microsoft 365 E5 customers.

When you use Microsoft Security Copilot, Customer Data and system-generated logs are stored and processed as part of the service.

Data sharing is turned on by default. Global Administrators and Security Administrators with the Capacity Contributor role (for example, capacity write access) for the capacity linked to a Workspace, can configure Customer Data sharing settings. Administrators can do this during the first run experience and at any time afterward. For more information on roles, see Security Copilot roles.

> **Important**
>
> Microsoft recommends that you use roles with the fewest permissions. Using lower permissioned accounts helps improve security for your organization. Global Administrator is a highly privileged role that should be limited to emergency scenarios when you can't use an existing role.

This article compares Security Copilot's Customer Data to system-generated logs, describes data sharing options, and summarizes how data is protected.

Customer Data and system-generated logs

As defined in the Microsoft Product Terms, *Customer Data* means all data, including all text, sound, video, or image files, and software, that are provided to Microsoft by, or on behalf of, the Customer through use of the Online Service. Customer Data doesn't include Professional

Services Data or information used to configure resources in the Online Services such as technical settings and resource names.

Microsoft online services create *system-generated logs* as part of the regular operation of the services. System-generated logs continuously record system activity over time to allow Microsoft to monitor whether systems are operating as expected. "Logging" (the storage and processing of logs) is essential to identify, detect, respond to, and prevent operational problems, policy violations, and fraudulent activity. Logging is also essential to optimize system, network, and application performance, as well as to help with security investigations and resilience activities and to comply with laws and regulations.

The following table compares Security Copilot's Customer Data to system-generated logs.

Expand table

Customer Data	System-generated logs
- Prompts that users submit to Security Copilot. - Information retrieved to generate responses. - Responses. - Content of pinned items. - File uploads.	- Account information (tenant ID, account ID, licensing, and others). - Usage data. - Performance information. - Internal system behavior information.

Customer Data sharing preferences

Copilot owners can change data sharing settings for Customer Data during the first run experience, and at any time thereafter.

Enabling or disabling these Customer Data sharing preferences described in the following table won't affect Microsoft's rights or responsibilities under the Microsoft Products and Services Data Protection Addendum.

The following data sharing options are available:

Expand table

Setting	Description
Allow Microsoft to capture data from Security Copilot to validate	Such validations include but aren't limited to: - Ability of Security Copilot to successfully provide responses to

Setting	Description
product performance using human review	user requests and understand capability gaps that need to be addressed based on user prompts. - Understand the types of tasks customers are using Security Copilot for. - Produce metrics surrounding the usability and quality of responses and agents. - Validate Security Copilot capabilities involving other Microsoft products purchased (including those products agents) and integrated by a customer. - Improve responses from plugins and agents accessing other Microsoft products. - Improve responses from Microsoft Security Copilot agents within Security Copilot, including agents accessing other Microsoft products. For more information, see Set up location for prompt evaluation and opt-in (or out of) data sharing.
Allow Microsoft to capture and human review data from Security Copilot to build and validate Microsoft's security AI model	Such validations include but aren't limited to: - Captured data* is used to develop security specific models built on top of Azure OpenAI foundational model or Microsoft built models, which would power more intelligent and personalized capabilities for Security Copilot and other Microsoft products that it integrates with. *Captured data includes data from Microsoft plugins and Microsoft Security Copilot agents utilized by you while using Security Copilot. **NOTE:** Data isn't shared with OpenAI or used to train the Azure OpenAI foundational model.

Accessing data from Microsoft 365 services

Security Copilot seamless integrates with multiple Microsoft 365 and Microsoft security services that your organization has licensed. You have the option of allowing users to use those products to query information directly from those services in both the standalone and embedded experiences.

> **Note**
>
> Currently, Security Copilot only accesses Microsoft 365 services data processed by Microsoft Purview, and Customer Data generated by Microsoft Purview (for example, DLP alerts).
>
> In Microsoft Purview, services such as data loss prevention (DLP), Insider Risk Management (IRM), or communication compliance are configured by the administrator to run on Microsoft 365 services data (or other data types).
>
> The data types that Security Copilot can access is dictated by what an administrator has configured for Microsoft Purview.

The following table summarizes the Microsoft 365 services data accessed by Security Copilot.

Expand table

Microsoft Purview Product or Services	What is accessed by Security Copilot
Data Loss Prevention	DLP alert data associated with a DLP match
Microsoft Purview Information Protection	Activity logs associated with labeling activity
eDiscovery	Data captured within a review set of an eDiscovery search
Insider Risk Management	IRM alert data associated with an IRM policy alert
Communication Compliance	Data captured within a policy match of a Communication Compliance Policy
Data Security Posture Management	Activity and alerts data associated with MIP, DLP, IRM

Microsoft 365 services data accessed by Security Copilot, including Customer Data generated by Microsoft Purview, is processed and stored according to the data processing activities described herein. This means, Microsoft 365 data accessed by Security Copilot will be processed and stored in location(s) described herein, regardless of the location in which the data was processed or stored pursuant to EU Data Boundary Services and the data residency commitments under the "Customer Data at Rest for Core Online Services" section of the Product Terms before the data was accessed by Security Copilot. This also means, Microsoft

365 data accessed by Security Copilot will be processed pursuant to the security practices and policies applicable to Security Copilot, regardless of the security practices and policies applicable to the data under "Security Practices and Policies for Core Online Services" section of the Product Terms before the Microsoft 365 data was accessed by Security Copilot.

To learn more about information captured, recorded, and retained by Microsoft Purview, see Learn about auditing solutions in Microsoft Purview. For information about activities that are audited in Microsoft 365, Audit log activities.

Expand table

Setting	Description
Allow Security Copilot to access data from your Microsoft 365 services	When turned on: - Security Copilot can retrieve your data from a Microsoft 365 service on your behalf if you're a customer of both Security Copilot and the Microsoft 365 service, and you allow Security Copilot access to your Microsoft 365 services. See the note in the prior section for more information about the Microsoft 365 services data accessed by Security Copilot. - All data, including Microsoft 365 data returned to answer your queries in Security Copilot is retained pursuant to Security Copilot's data retention policy. For more information, see Customer Data storage location.

Configuring Microsoft 365 services data access

Use the following steps to turn on or off Security Copilot's access to Microsoft 365 services.

1. In Security Copilot, go to **Settings** > **Owner settings**.
2. Update your data sharing selection.

> ⚠ **Warning**
>
> Turning off Microsoft 365 data access does not mean that any data retrieved from those Microsoft 365 services are deleted at that instant. Microsoft 365 data accessed up to that time will be deleted pursuant to Security Copilot's data retention policy. For more information, see **Data retention and deletion**.

Customer Data storage location

Data is stored in the location you select during Workspace creation. This happens during the onboarding steps as you set up your first Workspace, or whenever a new Workspace is created. Whenever you create a Workspace, you can choose the new Workspace's data storage location to be the same or different from that of your other existing workspaces. Once you have created a Workspace, you cannot change its data storage location. For Government Community Cloud (GCC) customers adopting the Security Copilot Commercial SKU, their relevant Customer Data associated with, and feedback shared from GCC Product SKUs will be stored in the US Geo of the Azure Public Cloud.

When data sharing is opted in, Customer Data such as prompts and responses are shared with Microsoft to enhance product performance, improve accuracy, and address response latency. In this case, Customer Data (except uploaded files) may be stored outside of the tenant's chosen Workspace geo. While uploaded files are stored in the tenant's chosen Workspace, if content from uploaded files is part of information retrieved to generate responses during sessions, that retrieved content can be stored outside of the Workspace. For more information, see Data, privacy, and security for Azure OpenAI Service.

> **Note**
>
> Communications initiated by skills through Microsoft Security Copilot agents that surface results through emails or Microsoft Teams may traverse global Microsoft network infrastructure during transit, and therefore may cross geographic boundaries, even when data at rest remains within the selected Workspace region.

File upload storage and processing

While uploaded files are stored in the tenant's chosen Workspace, if content from uploaded files is part of information retrieved to generate responses during sessions, that retrieved content can be stored outside of the Workspace. Uploaded files are only available to the user account that uploaded them, and not available to other users within or outside the tenant.

When data sharing is opted in, Microsoft may only capture and human-review content from uploaded files when that content is part of information retrieved to generate responses.

Data retention and deletion

Security Copilot stores Customer Data necessary for in-product functionality (such as your session data (for example your prompts and responses)) so long as you have an active

subscription to Security Copilot.

Customer Data can be deleted in the following scenarios:

- **When you delete all provisioned capacity**
  Customer Data is deleted within 180 days of when you delete all provisioned capacity.

- **When you request for your Customer Data to be deleted**
  You can also request that Security Copilot delete your Customer Data through the portal (https://securitycopilot.microsoft.com) or by requesting deletion through customer support . This Customer Data will be deleted within 30 days of that request.

When you opt-in to sharing your Customer Data with Security Copilot this Customer Data is only retained by Security Copilot for 90 days before being deleted by that team evaluating that Customer Data.

If you opt out of data sharing, Security Copilot deletes all Customer Data shared within 30 days. Customer Data is retained by you in your tenant so long as you have an active subscription to Security Copilot and have not requested it be deleted.

Data Subject Requests (DSRs) under the GDPR

For DSRs, please contact the support team: customer support. The Security Copilot team will then guide you through the next steps and process your request in accordance with GDPR guidelines. If you have any further questions or need additional assistance, feel free to contact support.

Location for prompt evaluation

With any Microsoft Copilot product, prompts refer to the text-based, natural language input you provide in the prompt bar that instructs Security Copilot to generate a response. Prompts are the primary input Copilot needs to generate answers that help you in your security-related tasks. Prompts are evaluated using GPU resources in Azure datacenters protected with Azure security and privacy controls.

You can choose to select where the prompts are evaluated from any of the following locations:

- Australia (ANZ)
- Europe (EU)
- United Kingdom (UK)

- United States (US)

You can opt in to having prompts evaluated anywhere in the world to mitigate potential disruptions in case your primary location experiences high activity.

Microsoft recommends having prompts evaluated anywhere with available GPU capacity, which enables the Copilot system to determine the optimal location based on load, latency, and responsiveness.

> **Note**
>
> Data is stored in the location you select during Workspace creation. For more information, see Customer Data storage location.

Set up location for prompt evaluation and opt in (or out of) data sharing

During initial setup, Copilot owners are prompted to set data sharing and prompt evaluation options. For more information, see Get started with Security Copilot. Copilot owners can change these settings during the first run experience, or at any time thereafter.

Authorized role
You need to be a Copilot owner to change the data sharing options. For more information on roles, see Understand authentication.

Set up data sharing

During initial setup, a Copilot owner is provided with the following data sharing options:

Expand table

Setting	Description
Allow Microsoft to capture data from Security Copilot to validate product performance using human review	Such validations include but aren't limited to: - Ability of Security Copilot to successfully provide responses to user requests and understand capability gaps that need to be addressed based on user prompts. - Understand the types of tasks customers are using Security Copilot for.

Setting	Description
	- Produce metrics surrounding the usability and quality of responses. - Validate Security Copilot capabilities involving other Microsoft products that a customer has purchased and integrated. - Improve responses from plugins accessing other Microsoft products.
Allow Microsoft to capture and human review data from Security Copilot to build and validate Microsoft's security AI model	Such validations include but aren't limited to: - Captured data is used to develop security specific models built on top of Azure OpenAI foundational model, which would power more intelligent and personalized capabilities for Security Copilot and other Microsoft products that it integrates with. **NOTE:** Data isn't shared with OpenAI or used to train the Azure OpenAI foundational model.

- When you opt in to data sharing, your Customer Data is shared with Microsoft from that point forward.
- When you opt out of data sharing, no further Customer Data is shared. Customer Data that was shared previously is retained for not more than 180 days.

Updating data sharing

1. In Security Copilot, go to **Settings** > **Owner settings**.

2. Update your data sharing selection.

How Microsoft protects your data

Microsoft uses comprehensive controls to protect your data. All Security Copilot data is handled according to Microsoft's commitments to privacy, security, compliance, and responsible AI practices. Access to the systems that house your data is governed by Microsoft's certified processes.

Security Copilot runs queries as the user, so it never has elevated privileges beyond what the user has.

If you opt in to share Customer Data, your data is:

- Not shared with OpenAI

- Not used for sales
- Not shared with third parties
- Not used to train Azure OpenAI foundational models

Security Copilot meets all Azure production data compliance standards.

All data stored in Azure is automatically encrypted at rest and uses AES-256 encryption. For more information, see Data encryption and Encryption at rest.

Microsoft security products data handling

Microsoft Security products purchased by you may share data, including Customer Data, as described in the product documentation. Customer Data shared with Security Copilot is governed by the Product Terms, Data Protection Addendum, and documentation applicable to Security Copilot. For Microsoft 365 services, an administrator needs to enable Security Copilot in the sharing preference option detailed in Accessing data from Microsoft 365 services and users will need to enable a plugin for those Microsoft 365 Services. For other Microsoft services, such plugins are enabled by default for users. Users can turn off plugins at any time. For more information, see Manage plugins.

Feedback from Security Copilot users

Microsoft collects feedback on the response produced by Microsoft Security Copilot from users of the product. A Copilot owner can turn off feedback collection for their tenant by contacting Microsoft Support through a support ticket. For more information, see Contact support.

See also

Data, privacy, and security for Azure OpenAI Service

Microsoft responsible AI principles .

Last updated on 04/08/2026

Microsoft Security Copilot egress IPs

To ensure communication between Security Copilot and certain security solutions, ensure to allow Security Copilot's egress IPs to contact the solution.

The egress IP addresses enable Security Copilot to securely send requests and receive responses from the plugin.

Allow the following IPs based on the instance type you are using.

Expand table

Region	Egress IP address
Australia East	4.237.91.160/29
Brazil South	4.160.64.120/29
Canada Central	4.173.72.88/29
Canada East	4.248.114.8/29
Central India	4.188.45.56/29
East US	48.211.114.168/29
East US 2	72.153.16.200/29
France Central	4.251.29.8/29
Germany West Central	4.182.156.104/29
Japan East	48.210.117.64/29
Korea Central	4.230.181.72/29
North Europe	72.145.24.32/29
Qatar Central	4.171.44.192/29
South Africa North	4.222.245.0/29
South Central US	23.98.246.88/29
Southeast Asia	57.155.157.72/29
Sweden Central	135.225.147.168/29
Switzerland North	74.242.234.136/29
UAE North	74.243.192.152/29

Region	Egress IP address
UK South	4.250.47.152/29
West Central US	172.208.176.8/29
West Europe	72.145.132.208/29
West US	172.178.185.8/29
West US 2	72.154.34.248/29
West US 3	57.154.125.160/29

Last updated on 12/19/2025

Supported languages in Microsoft Security Copilot

Security Copilot supports prompts and providing responses in various languages. This topic provides guidance on the supported languages and how to configure a preferred language.

Supported languages

The following languages are supported in Security Copilot:

Expand table

Languages	Availability
English (US, GB, AU, CA, IN)	Prompting and product interface
Spanish (Spain, Mexico)	Prompting and product interface
Japanese	Prompting and product interface
French (France, Canada)	Prompting and product interface
German	Prompting and product interface
Portuguese (Brazil)	Prompting and product interface
Italian	Prompting and product interface
Chinese Simplified	Prompting and product interface
Korean	Prompting and product interface
Dutch	Product interface
Swedish	Product interface
Polish	Product interface
Portuguese (Portugal)	Product interface
Norwegian	Product interface
Turkish	Product interface
Danish	Product interface
Finnish	Product interface
Chinese Traditional	Product interface

Languages	Availability
Arabic	Product interface
Thai	Product interface
Hebrew	Product interface
Czech	Product interface
Hungarian	Product interface
Ukrainian	Product interface
Russian	Product interface

Configure language preference

1. Sign in to Microsoft Security Copilot.

2. In the left navigation pane, navigate to your tenant settings.

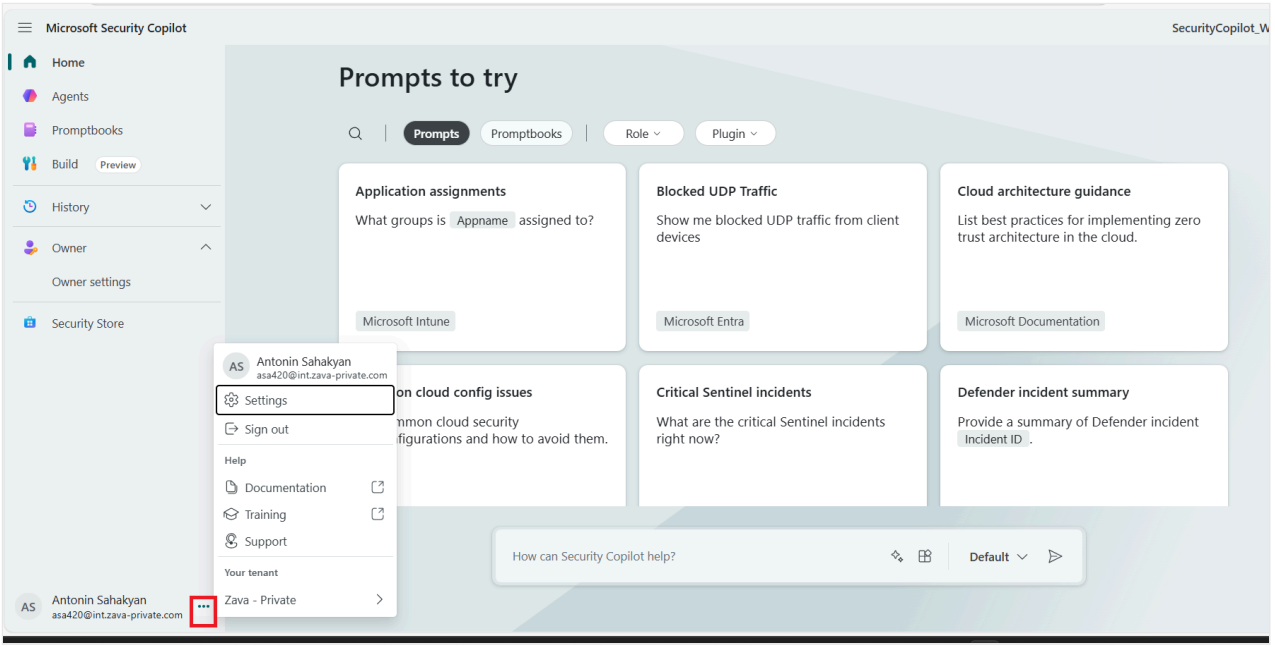

3. Select **Settings** > **Preferences**, select your preferred language.

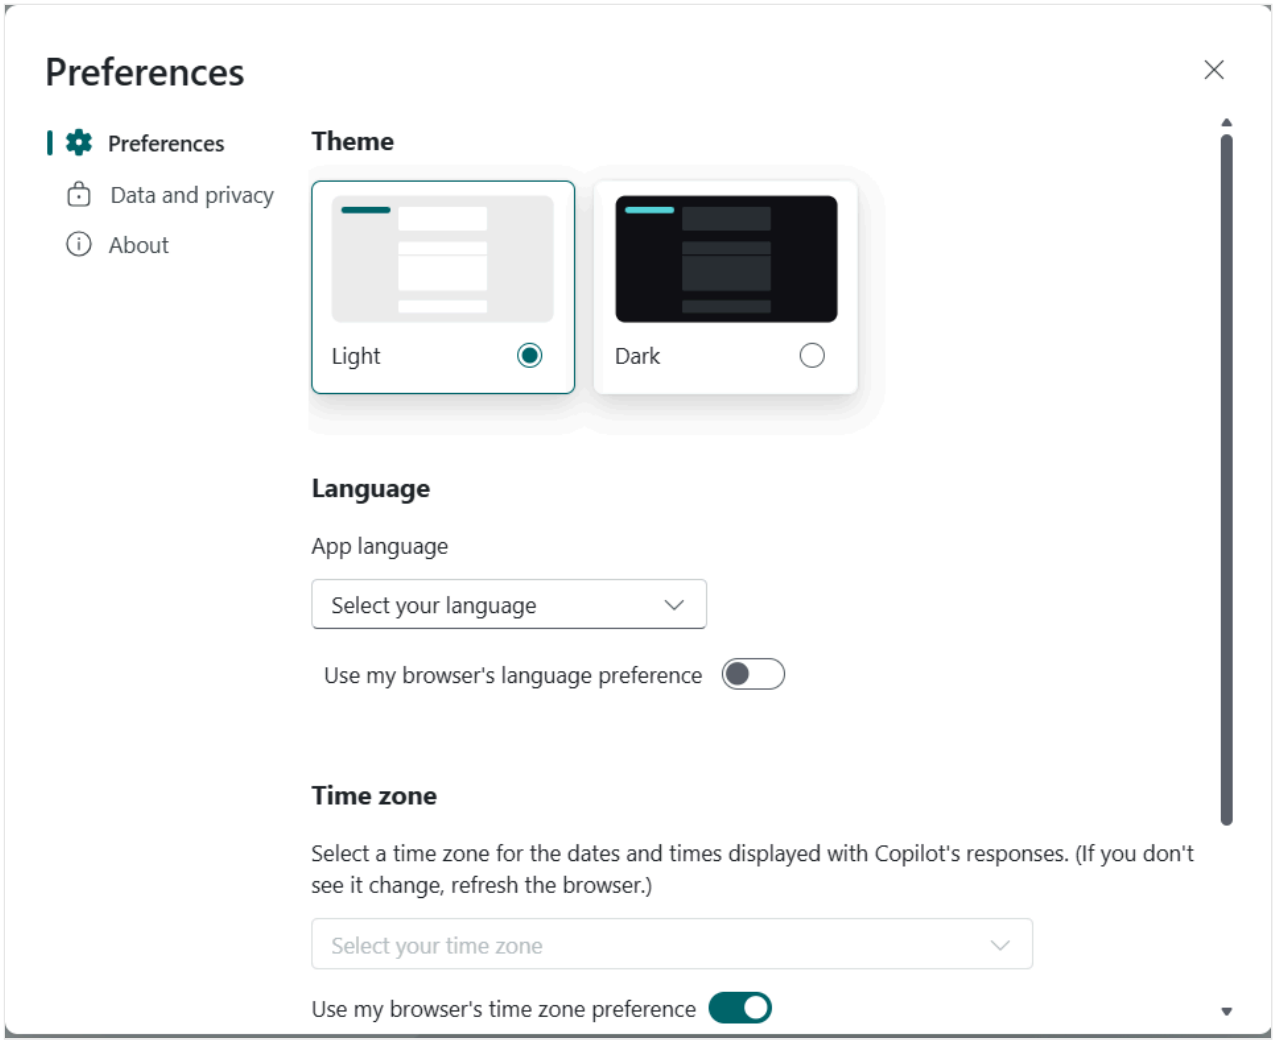

Create prompts using supported languages

Scenario 1: Prompt using a supported language

If you set a language in the preferences setting and create a prompt using the selected language, the resulting response aligns with your chosen language setting.

Scenario 2: Prompt in a different language other than the one set in the preferences settings

If you create a prompt using a supported language other than the one set in the preferences setting, the resulting response aligns with the language you used when you created the prompt - even if a different language was configured in the preferences settings.

For example, if Spanish was set in the preferences settings, but you created a prompt in English, the resulting response will be in English.

Scenario 3: Prompt in a different language other than the one set in the preferences settings and request for a response in another supported language

If you create a prompt that explicitly requests a response in a different supported language, the resulting response will adhere to the language you specified in the prompt. This holds true

even if you set a different language in the preferences settings or used a different language to create a prompt. The system prioritizes the language explicitly requested within the prompt.

For example, if Spanish was set in the preferences settings, but you created a prompt in English requesting the response be made in Dutch, the resulting response will be in Dutch.

Scenario 4: Prompt in a non-supported language

If you create a prompt using a language that isn't supported, an error will be displayed.

Last updated on 01/12/2026

Microsoft Security Copilot Frequently Asked Questions

General information

What is Microsoft Security Copilot?

Microsoft Security Copilot is a generative AI-powered assistant for security and IT. It provides tailored insights and recommendations using global threat intelligence, industry best practices, and organizations' data from Microsoft and partner tools. Teams can also automate workflows with autonomous Security Copilot agents, accelerating responses, prioritizing risks, and reducing manual workloads - all while remaining firmly in control.

What are the use cases and capabilities that Security Copilot unlocks for customers?

Security Copilot focuses on making the following use cases easy to accomplish.

- Investigating and remediating security threats
- Building KQL queries and analyzing suspicious scripts
- Understanding risks and managing organization's security posture
- Troubleshooting IT issues faster
- Defining and managing security policies
- Configuring secure lifecycle workflows
- Developing reports for stakeholders
- Automating tasks with autonomous agents

Visit the Security Copilot adoption hub to delve deeper into how Security Copilot benefits roles like CISOs, threat intelligence analysts, IT admins, data security admins, and more.

Does Microsoft Security Copilot work with other Microsoft products?

Yes. Security Copilot is integrated with and embedded in other Microsoft Security products. These products include, but aren't limited to:

- Azure Firewall

- Microsoft Defender Attack Surface Management
- Microsoft Defender for Cloud
- Microsoft Defender Threat Intelligence
- Microsoft Defender XDR
- Microsoft Intune
- Microsoft Purview
- Microsoft Sentinel

Security Copilot can access data from these products and provide genAI assistance and agentic automation to increase the effectiveness and efficiency of security professionals using those solutions.

Who are the intended users of Security Copilot?

SOC analysts, compliance analysts, IT admins, data security admins, identity admins and CISOs are some of the intended users of Security Copilot. Visit the Security Copilot adoption hub to learn about the key scenarios.

What languages are supported?

Security Copilot supports multiple languages. The model is available in eight languages* and the user experience is available in 25 languages.**

*Model: English, German, Japanese, Spanish, French, Italian, Portuguese, and Chinese

**UX: Above languages plus Korean, Dutch, Swedish, Polish, Norwegian, Turkish, Danish, Finnish, and more in the user experience.

For more information, see Supported languages.

What is the difference between Security Copilot and generic LLM?

A generic LLM lacks key benefits such as:

- Real-time signal processing with structured log data, to detect and analyze threats as they happen rather than after the fact.
- Investigative reasoning to trace the source and impact of incidents. AI needs to pivot across many data points, just as human analysts would, and adjust conclusions as new information emerges.

- Precision with evidence to ensure that the AI is accurate and backed by reliable data.
- Persistent data collection to maintain a continuous, comprehensive view.

Without these capabilities and a deep understanding of your specific security environment, a generic model simply can't provide the depth of analysis required.

Microsoft is uniquely positioned to tackle the challenges of using AI in security – our solution is built on the most advanced models available, and builds on:

- A hyperscale infrastructure, providing the scalability and infrastructure needed to process massive volumes of data in real-time.
- A security-specific orchestrator that streamlines, contextualizes, and coordinates responses across tools and teams.
- Seamless integration through plugins, connecting your existing systems and extending what our solution can do.
- Evergreen threat intelligence, informed by 84 trillion daily signals, for unmatched visibility across your whole threat landscape.
- Built-in cyber skills training and promptbooks to empower teams with the knowledge and resources they need.

By building on all these components, we provide an AI security solution that understands your security environment, delivers insights in natural language, and handles sensitive data securely.

For Customer Data handling FAQs, see Data and Compliance.

Purchasing and billing information

Are there any prerequisites to purchase?

An Azure subscription and Microsoft Entra ID (formerly known as Azure Active Directory) are prerequisites for using Security Copilot; there are no other product prerequisites. For more information, see Get started with Security Copilot.

Is deployed Microsoft Entra ID (formerly known as Azure Active Directory) a requirement for Security Copilot?

Yes. Security Copilot is a SaaS application and requires Microsoft Entra ID to authenticate the users who have access.

How is Security Copilot priced?

Security Copilot is priced based on Security Compute Units (SCUs).

- Provisioned SCUs support regular workloads and are billed monthly.
- Overage SCUs offer flexible, on-demand capacity and are billed only when used.

Use the in-product dashboard to monitor SCU usage and adjust capacity as needed.

For more information, see Microsoft Security Copilot - Pricing

Does Security Copilot support tenant or subscription transfers?

No, at this time Security Copilot doesn't support moving Security Copilot resources across Microsoft Entra tenants or subscription transfers.

Can SCUs be shared between a tenant's Security Copilot Workspaces?

SCUs, whether Provisioned or Overage, can't be shared between Workspaces. For example, if an organization has Workspace A (1 provisioned, three overage) and Workspace B (2 provisioned, five overage), Workspace A can't use Workspace B's SCUs to prevent throttling if it exhausts its own total of 4 SCUs. Similarly, Workspace B can't use Workspace A's SCU capacity if it exhausts its own SCUs.

How can I estimate SCU provisioning and budget for Security Copilot?

Estimating SCU needs and budgeting depends on how your organization uses generative AI across Microsoft Security products. Use the SCU capacity calculator to get a starting point based on the number of users and workloads across Microsoft Defender, Microsoft Intune, Microsoft Purview, Microsoft Entra, and the Security Copilot standalone experience. Because every prompt and workflow varies in complexity, SCU consumption isn't fixed. Start small, experiment, and refine based on your actual usage. Track usage in real time with the in-product dashboard, which helps you monitor SCU consumption and adjust capacity as needed.

How do I interpret the results on the SCU capacity calculator?

The results show the maximum expected number of SCUs per hour, based on the number of monthly users per experience and the inclusion of automation through Logic Apps and Promptbooks. In addition to the maximum SCUs per hour, the calculator also displays all possible

combinations of provisioned and overage SCUs, along with the corresponding monthly cost range. The highest monthly cost is calculated based on continuous 24/7 capacity usage.

What if my needs aren't covered by the SCU capacity calculator (for example, I have more than 50 Intune users).

Feel free to Contact us or reach out to your Microsoft account manager for more information.

What are the minimum provisioning requirements for SCUs in Workspaces?

To use Microsoft Security Copilot, each tenant must provision a minimum of one SCU. The one provisioned SCU enables access to at least one Workspace.

- Provisioned SCUs are required to activate capacity and are billed hourly.
- Overage SCUs can be set from 0 to 999 and are billed only when used.

Technical and product information

What partner tools are integrated with Security Copilot?

Microsoft Security Copilot supports many plugins, including Microsoft and non-Microsoft plugins. For more information, see Plugins.

> **Note**
>
> Products that integrate with Security Copilot need to be purchased separately.

Does Security Copilot make recommendations for IoT/OT scenarios?

No, Security Copilot doesn't currently support IoT/OT.

Does Security Copilot offer dashboarding?

Security Copilot offers an in-product usage dashboard where customers can dive deep into their SCU consumption.

Can Security Copilot execute workflows - from triaging to using pinned messages, to governing how the customer should label the incident and whether an incident should be closed?

Security Copilot provides Promptbooks that are ready-to-use workflows that can serve as templates to automate repetitive steps—for instance, regarding incident response or investigations. Additionally, there are connectors in Security Copilot that are a wrapper around the API that allows developers and users to call out to the Microsoft Security Copilot platform to perform specialized tasks. For example, the Logic Apps connector allows you to call into Copilot from an Azure Logic Apps workflow. Similarly, the Copilot Studio connector enables you to access Security Copilot to perform actions like 'submit a Security Copilot prompt' or 'fetch a Security Copilot prompt status'.

What role-based access control or delegation features does Security Copilot have? How are user permissions kept in Security Copilot aligned to user permission configurations in other solutions?

Copilot uses on-behalf-of authentication to access security-related data through active Microsoft plugins. Specific Security Copilot roles must be assigned in order for a group or individual to access the Security Copilot platform. For more information, see Understand authentication.

How is Security Copilot dealing with a "token limit"?

Large language models (LLMs) including GPT have limits on how much information they can process at once. This limit is known as a "token limit", and roughly correlates to 1.2 words per token. Security Copilot uses the latest GPT models from Azure OpenAI to ensure we can process as much information as possible in a single session. In some cases, large prompts, long sessions, or verbose plugin output may overflow the token space. When this scenario happens, Security Copilot attempts to apply mitigations to ensure an output is always available, even if the content in that output isn't optimal. Those mitigations aren't always effective, and it might be necessary to stop processing the request and direct the user to try a different prompt or plugin.

What are Security Copilot agents?

Microsoft Security Copilot agents enhance security and IT operations with autonomous and adaptive automation. Agents seamlessly integrates with Microsoft Security solutions and third-party partner ecosystem to handle high-volume security tasks. Purpose-built for security, these

agents learn from feedback, adapt to organizational workflows with your team fully in-control, and operate securely within Microsoft's Zero Trust framework—accelerating responses, prioritizing risks, and driving efficiency. By reducing manual workloads, they enhance operational effectiveness and strengthen your organization's overall security posture.

Visit the adoption hub to learn more about Security Copilot agents.

For information on how to build custom agents, see Developer content.

What compute resources do Security Copilot agents use?

Agents utilize SCUs to operate just like other features in Security Copilot. They integrate seamlessly with Microsoft Security solutions and the broader supported partner ecosystem.

Where can I find Security Copilot agents?

You can easily discover Microsoft Security Copilot agents from both the standalone and embedded experiences.

For the standalone experience, you can select Go to agents from the banner. You can also navigate to the agent library from the home menu. In the embedded experiences, you'll see agents within the portal and explore their capabilities.

Visit the adoption hub to learn more.

What are connectors in Security Copilot?

The connectors in Security Copilot are a wrapper around the API that allows the developers and users to call out to the Microsoft Security Copilot platform to perform specialized tasks. Currently, Logic Apps and Copilot Studio connectors are supported. For more information, see Connectors.

How does the Copilot Studio connector work?

Copilot Studio connector enables you to access Security Copilot while creating automation workflows. Using the Security Copilot connector, you can:

- Submit a Security Copilot prompt - Submit a natural language prompt to create a new Security Copilot investigation. After completion, the evaluation result will then be returned to your workflow.
- Fetch a Security Copilot prompt status - Submit a natural language prompt to pull the status of a Security Copilot evaluation. After completion, the evaluation result will then be

returned to your workflow.

How does the Logic Apps connector work?

The Microsoft Security Copilot Logic Apps connector allows you to call into Copilot from an Azure Logic Apps workflow. The connector exposes two connector actions:

- Submit a Security Copilot prompt - Submit a natural language prompt to create a new Security Copilot investigation. After completion, the evaluation result will then be returned to your workflow.
- Submit a Security Copilot promptbook - Given a promptbook, invoke a new Security Copilot promptbook evaluation and return the output to your Azure Logic Apps workflow.

Partner information

What are the use cases for Partners?

Partners can provide signals or build complementary solutions around Security Copilot scenarios.

If a customer works with a managed security service provider (MSSP), can the MSSP use and manage Security Copilot on the customer's behalf?

Yes, MSSPs that provide SOC services for customers are able to access the customer's Security Copilot environment if the customer elects to provide access. Available options include:

- Azure Lighthouse
- B2B Collaboration / Guest Accounts
- Granular Delegated Admin Privileges (GDAP)

There currently isn't a CSP or reseller multitenant model for MSSPs. Each customer is responsible for purchasing their own SCUs and setting up their MSSPs with the necessary access if they're accessing within the customer's tenant. If Azure Lighthouse delegated access is provided to the customer's Microsoft Sentinel Workspaces then the partner is able to use their partner capacity plan SCUs to perform prompting against the customer's Microsoft Sentinel Workspace data.

If an MSSP prompts across a customer's tenant, using Azure Lighthouse, will the MSSP use their own SCUs or the customer's SCUs?

Azure Lighthouse allows partners to gain Security Copilot permissions for customer's Microsoft Sentinel Workspaces and other supported Azure resources. The capacity plan (SCUs) used is the partner tenant's capacity plan.

Can MSSPs, use a single instance of Security Copilot to manage multiple tenants?

Azure Lighthouse is supported for invoking Sentinel-based skills from the partner tenant against a single customer's Microsoft Sentinel Workspace at a time that the partner has been provided delegated access to through Azure Lighthouse. The partner tenant will leverage their SCUs to perform Security Copilot invoked Microsoft Sentinel skills against the customer's tenant without the need for the customer tenant to be provisioned for Security Copilot or having SCUs of their own.

Are there third-party integrations available today?

Microsoft Security Copilot supports many plugins, including Microsoft and non-Microsoft plugins. For more information, see Non-Microsoft plugins overview. Additionally, there are agents from our partners that are available in Security Copilot. For more information, see Partner agents.

> **Note**
>
> Products that integrate with Security Copilot need to be purchased separately.

Is there a marketplace for the plugins or services?

There isn't a plugin marketplace. ISVs can publish their solutions to GitHub. All partners are required to publish their solutions or managed services to the Microsoft Commercial marketplace. For more information on publishing to marketplace, see:

Publish solution to the Microsoft Commercial Marketplace:

- Create a commercial marketplace account in Partner Center - Partner Center
- How to manage a listing in Partner Center

MSSP Specific: Must have a security designation in Microsoft AI Cloud Partner Program.

- Consulting Service
- Managed Services

SaaS Specific:

- Plan a SaaS offer for the Microsoft commercial marketplace - Marketplace publisher
- Create a SaaS offer in the commercial marketplace - Marketplace publisher
- ISV Success Program Overview
- Build and publish with ISV Success - Partner Center

What if MSSPs aren't using Microsoft Defender XDR or Microsoft Sentinel?

Microsoft Security Copilot doesn't have any specific Microsoft security product requirement for provisioning or use since the solution is built on aggregating data sources both from Microsoft and third party services. With that said, there's a significant value in having Microsoft Defender XDR and Microsoft Sentinel enabled as supported plugins for enriching investigations. Security Copilot only uses skills and accesses data from enabled plugins.

Does an MSSPs SOC Solution need to be hosted on Azure?

It's recommended that the solution is hosted on Azure but not required.

Is there a product roadmap that can be shared with Partners?

Not at this time.

Last updated on 04/27/2026

Availability and recovery of Security Copilot

Summarize this article for me

This article describes the availability and recovery support in Microsoft Security Copilot experiences.

Availability zone support

Azure availability zones are at least three physically separate groups of datacenters within each Azure region. Datacenters within each zone are equipped with independent power, cooling, and networking infrastructure. When a local zone fails, availability zones are designed so that if one zone is affected, regional services, capacity, and high availability are supported through the remaining two zones.

Failures can range from software and hardware failures to events such as earthquakes, floods, and fires. Tolerance for failures is achieved with redundancy and logical isolation of Azure services. For more detailed information on availability zones in Azure, see Regions and availability zones.

Security Copilot makes commercially reasonable efforts to use zone-redundant availability resources; underlying resources automatically replicate across zones, without any need for you to set up or configure.

> **Note**
>
> Availability zones protect service uptime; they don't determine where prompts are processed.

Availability zones

Availability zones are supported in various Geos as listed in the table:

Expand table

Geo	Copilot Sessions
ANZ	✅
Asia	✅

Geo	Copilot Sessions
Brazil	✅
Canada	✅
Europe	✅
India	Not supported due to capacity constraints
Japan	✅
Qatar	✅
South Africa	✅
Switzerland	✅
UAE	Not supported due to capacity constraints
UK	✅
US	✅

> **Note**
>
> Processing location is unrelated to high availability location. It is determined dynamically per-user traffic.

Zonal outage recovery

During a zone-wide outage, no action is required during zone recovery. Copilots for Security capabilities in supported Geos self-heal and rebalance automatically to take advantage of the healthy zone. For more information, see What are Azure availability zones?.

Disaster recovery and business continuity

Disaster recovery (DR) is about recovering from high-impact events, such as natural disasters or failed deployments that result in downtime and data loss. Regardless of the cause, the best remedy for a disaster is a well-defined and tested DR plan and an application design that actively supports DR. Before you begin to think about creating your disaster recovery plan, see Recommendations for designing a disaster recovery strategy.

When it comes to DR, Microsoft uses the shared responsibility model. In a shared responsibility model, Microsoft ensures that the baseline infrastructure and platform services are available. At

the same time, many Azure services don't automatically replicate data or fall back from a failed region to cross-replicate to another enabled region. There may be scenarios where a specific Geo does not support replication or services have system errors.

For Security Copilot, best effort is made to back up and restore your data if there is an outage. However, it's important to note that you should have a plan ready to deal with outages, even if it involves manual intervention. This means that you should have a disaster recovery plan in place that suits your workload.

Most services that run on Azure platform as a service (PaaS) offerings also provide features and guidance to support DR and you can use service-specific features to support fast recovery to help develop your DR plan.

Recovery time objective and Recovery point objective

The following table lists the supported geographies for Security Copilot. There is 72-hour Recovery Time Objective (RTO) and 15-minute Recovery Point Objective (RPO) commitment for regional outages within a geography. If a particular region experiences an outage, the RTO and RPO commitments are applied for that Geo.

Expand table

Geo	Copilot Sessions
ANZ	✅
Asia	Not replicated due to capacity constraints
Brazil	Not replicated due to capacity constraints
Canada	✅
Europe	✅
India	Not replicated due to capacity constraints
Japan	✅
Qatar	✅
South Africa	Not replicated due to capacity constraints
UAE	✅
UK	✅
US	✅

Last updated on 02/25/2026

Contact support for Microsoft Security Copilot

If you encounter issues or need to seek assistance, Security Copilot provides a modern and advanced support experience.

Depending on your role, the widget allows you to:

- Find solutions to common problems
- Submit a support case to the Microsoft support team

Prerequisites

It's important to know the specific roles that have permission to open support cases.

At a minimum, you must have a Service Support Administrator **OR** Helpdesk Administrator role.

For more information on which roles have permission, see, Security Administrator permissions. Roles that include the action `microsoft.office365.supportTickets/allEntities/allTasks` can submit a case.

For general information on administrator roles, see About admin roles.

Use the self help widget

Anyone with access to Security Copilot can access the self help widget.

1. In Security Copilot, navigate to the help icon, then select **Help**.

2. Select the **Self help** tab.

3. Type your question in the prompt bar. Articles related to your search will be surfaced.

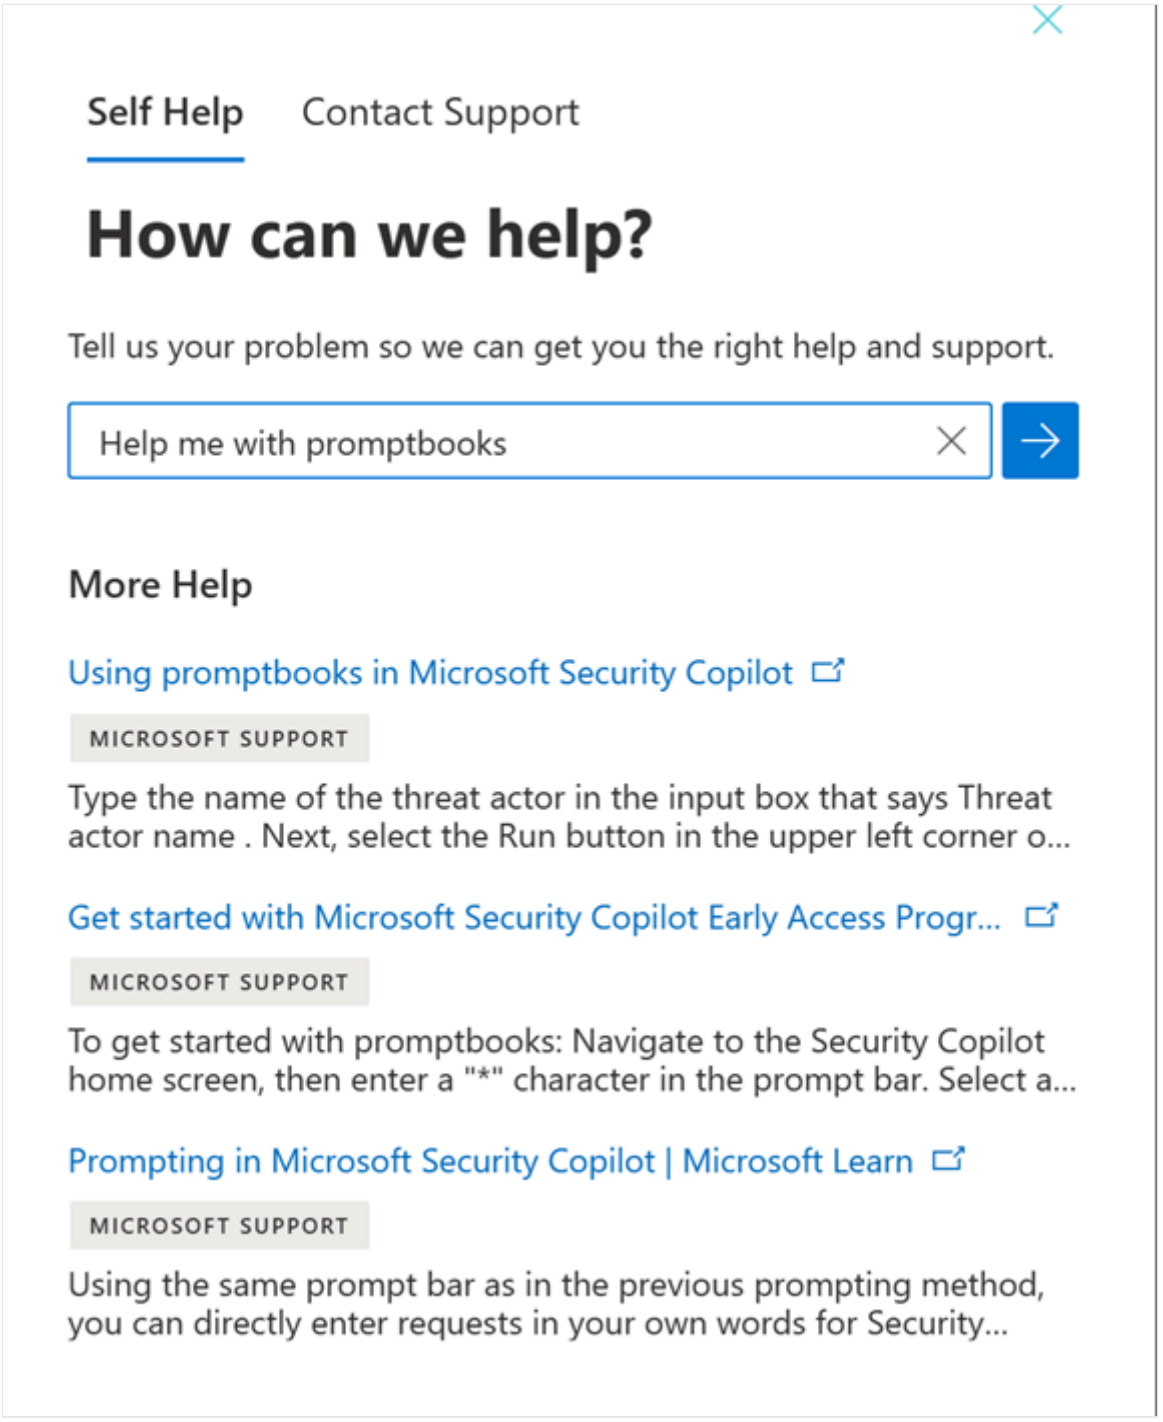

Create a support ticket

If you require more support, you can open support tickets by contacting the Security Copilot support.

Required role

You'll need an administrator role to create support tickets. For more information on roles, see Understand authentication.

1. In Security Copilot, navigate to the help icon, then select **Help**.

2. Select the **Contact support** tab.
3. Fill in the details related to your case.

> ⓘ **Note**
>
> Leave the **Product** field blank. Do NOT select any of the options from the drop-down list.

4. (Optional) Include up to five attachments that are relevant to the issue to provide additional context for the support case.
5. Select **Contact me**.

View previously submitted tickets

You can view the status of your previously submitted tickets in the Support History tab.

1. In Security Copilot, navigate to the help icon, then select **Help**.

2. Select the **Support history** tab.
3. View, edit, and track the status of the support tickets you submitted.

Last updated on 06/04/2026

Understand why Azure Private Link isn't needed for Security Copilot

Security Copilot is a multitenant, cloud-native SaaS solution. Unlike single-tenant services that benefit from Azure Private Link's network isolation, Copilot is designed to operate securely over public endpoints. It can't be deployed into a customer's private network or virtual network, and Private Link isn't supported or required.

Secure by design architecture

- Security Copilot achieves the goals of Private Link secure, isolated data flows through its architecture:
- **Encrypted Communications**: All traffic is secured with TLS 1.2+ encryption, ensuring confidentiality and integrity across public endpoints.
- **Identity-Based Access Control**: Azure AD authentication and role-based permissions enforce strict tenant-level isolation. Every request is authenticated and authorized.
- **Microsoft Backbone Routing**: Service-to-service communication (for example, with Microsoft Sentinel or Defender) occurs over Microsoft's internal cloud network, not the public internet. This ensures data stays within Microsoft's trusted infrastructure.
- **No Inbound Exposure**: Customers don't need to expose internal systems. For plugin integrations, outbound calls from Copilot can be restricted to known IP ranges, maintaining Zero Trust principles.

Compliance and data protection

Security Copilot complies with Microsoft's enterprise-grade security and privacy standards:

- **Encryption at Rest and in Transit**: All customer data is protected using AES-256 encryption at rest and TLS 1.2+ in transit, ensuring confidentiality and integrity throughout the data lifecycle.
- **Tenant and Region Isolation**: Data is stored within the customer's home region and logically isolated by tenant. This guarantees that no data is commingled across organizations, supporting regulatory and operational boundaries.
- **No Data Shared with OpenAI**: Although Security Copilot leverages AI models, customer data isn't sent to OpenAI or external providers. All model processing occurs entirely within Microsoft's compliance boundary, using infrastructure managed and secured by Microsoft.
- **Compliance Frameworks**: Security Copilot aligns with Microsoft's broader compliance commitments, including SOC 2, ISO 27001, and GDPR. Customers can apply their own

governance policies-such as retention, eDiscovery, and access controls through integrations with Microsoft Purview.

This architecture ensures that even without Azure Private Link, Security Copilot meets the security and compliance expectations of regulated and security-conscious organizations.

Understand why Private Link isn't needed

Implementing Private Link in a multitenant SaaS would require managing thousands of private endpoints-an operationally unscalable model. Moreover, most users access Copilot from varied locations, making network-level isolation less relevant than application-layer security. Copilot's design already meets the intent of Private Link through secure defaults and cloud-native protections.

Summary

Security Copilot achieves the goals of Private Link secure, isolated data flows through its architecture. For technical customers, this means strong security without the need for network tunneling or endpoint management.

Last updated on 01/12/2026

Microsoft Purview roles that inherit Security Copilot contributor access

The following Microsoft Purview roles with inherited Security Copilot contributor access:

Expand table

Microsoft Purview Role
Communication Compliance
Communication Compliance Analysts
Communication Compliance Investigators
Compliance Administrator
Compliance Data Administrator
Data Security Management
Information Protection
Information Protection Admins
Information Protection Analysts
Information Protection Investigators
Insider Risk Management
Insider Risk Management Admins
Insider Risk Management Analysts
Insider Risk Management Approvers
Insider Risk Management Auditors
Insider Risk Management Investigators
Organization Management
Data Catalog Reader
Data Governance Administrator
Data Health Owner
Data Health Reader

Microsoft Purview Role
Data Product Owner
Data Profile Reader
Data Profile Steward
Data Quality Metadata Reader
Data Quality Reader
Data Quality Stewards
Data Steward
Governance Domain Creator
Governance Domain Owner
Governance Domain Reader

Last updated on 01/05/2026

Application card: Microsoft Security Copilot

What is an Application or Platform card?

Microsoft's Application and Platform cards are intended to help you understand how our AI technology works, the choices application owners can make that influence application performance and behavior, and the importance of considering the whole application, including the technology, the people, and the environment. Application cards are created for AI applications and platform cards are created for AI platform services. These resources can support the development or deployment of your own applications and can be shared with users or stakeholders impacted by them.

As part of its commitment to responsible AI, Microsoft adheres to six core principles: fairness, reliability and safety, privacy and security, inclusiveness, transparency, and accountability. These principles are embedded in the Responsible AI Standard, which guides teams in designing, building, and testing AI applications. Application and Platform cards play a key role in operationalizing these principles by offering transparency around capabilities, intended uses, and limitations. For further insight, readers are encouraged to explore Microsoft's Responsible AI Transparency Report and the Microsoft Enterprise AI Services Code of Conduct, which outlines how to engage with AI responsibly.

Overview

Microsoft Security Copilot is a generative AI-powered security solution that helps increase the efficiency and capabilities of defenders to improve security outcomes at machine speed and scale. It provides a natural language, assistive copilot experience that helps security professionals and IT administrators handle a wide range of end-to-end scenarios, including incident response, threat hunting, intelligence gathering, and posture management.

Security Copilot is designed with integration in mind. It offers an immersive standalone experience at https://securitycopilot.microsoft.com and integrates with Microsoft security products including Microsoft Defender XDR, Microsoft Sentinel, Microsoft Intune, Microsoft Entra, and Microsoft Purview, as well as supported third-party services.

The intended users include SOC analysts, IT administrators, data security and identity administrators, compliance analysts, and security leaders such as Chief Information Security Officers (CISOs).

Security Copilot has earned the ISO 42001 certification that confirms that an independent third party reviewed Microsoft's application of the necessary framework and capabilities to effectively manage risks and opportunities associated with the continuous development, deployment, and operation of Microsoft AI systems.

For more information, see What is Microsoft Security Copilot?, Microsoft Security Copilot experiences, and Application card for Security Copilot agents.

Key terms

The following table provides a glossary of key terms related to Microsoft Security Copilot.

Expand table

Term	Definition
Agent	A Security Copilot agent processes signals from the customer environment through integrated data sources and plugins, analyzes over data, and generates recommendations. Agents may also perform scoped actions within configured permissions when prompted, which require an appropriate user or administrator approval. Agents can range from simple prompt-and-response experiences to more automated, semi-autonomous workflows with human oversight. For example, the Phishing Triage Agent and the Vulnerability Remediation Agent.
Agent identity	The credential an agent uses to authenticate with Microsoft services and access the data it needs to perform its tasks. During setup, an administrator chooses either to create a dedicated identity (using Microsoft Entra Agent ID) or to let the agent inherit the credentials of the configuring user. The choice of identity governs what data the agent can access.
Data source	Structured or unstructured data accessed through plugins or integrations to ground responses.
Embedded experience	Accessing Security Copilot capabilities from within another Microsoft security product, such as Microsoft Defender XDR or Microsoft Sentinel. The Security Copilot sidecar panel surfaces AI assistance directly in the context of that product.
Grounding	The process of providing contextual input sources to the large language model related to a user's prompt. By enabling Security Copilot to access organizational data through plugins and Microsoft security products, Security Copilot can deliver more accurate and contextually relevant responses.
Large language model (LLM)	AI models trained on large amounts of text data to predict words in sequences. LLMs are capable of performing various tasks such as text generation, summarization, translation, classification, and more.
Plugin	A collection of related tools that extends Security Copilot's capabilities by giving it access to resources from Microsoft and non-Microsoft services and public websites through APIs.

Term	Definition
	Plugins add more context to the responses and outputs that Security Copilot generates.
Post-processing	The set of actions Security Copilot performs to refine and prepare the LLM response before returning it to the user. This post-processing includes additional grounding calls through plugins, responsible AI checks, security, compliance, and privacy checks.
Prompt	The natural language text a user sends to Security Copilot to execute a specific task or obtain information. For example, *Summarize this incident and suggest remediation steps.*
Promptbook	A series of prompts that run in sequence, building on previous responses, to accomplish specific security-related tasks. Promptbooks can be used from the library or built and shared by users.
Red team testing	Techniques used by experts to assess the limitations and vulnerabilities of a system and to test the effectiveness of planned mitigations. Red team testing is used to identify potential risks and is distinct from systematic measurement of risks.
Responsible AI	Microsoft's policy, research, and engineering practices that are grounded in its AI principles and operationalized through the Responsible AI standard. For more information, see the Fluent RAI guidance.
Security Compute Unit (SCU)	SCUs are the units of compute capacity used to run Security Copilot workloads and deliver consistent performance across its experiences. Security Copilot capacity is measured in SCUs and can be consumed through provisioned or overage capacity models. For more information, see Understand SCUs.
Security Copilot response	The AI-generated output returned to a user prompt, including summaries, recommendations, or actions that may include code and visuals.
Security Operations Center (SOC)	A dedicated security team or facility focused on continuous monitoring, analysis, and response to cybersecurity incidents within an organization. SOC analysts are among the primary intended users of Security Copilot.
Security Store	A marketplace-like experience where users can discover and enable partner-built agents and extensions for Security Copilot.
Session	A bounded interaction context that includes prompts, responses, and associated data used to generate outputs. Session context is isolated and isn't shared across sessions. Access to session data is governed by user permissions and workspace roles, and prompts and responses are treated as Customer Data that isn't used to train foundation models.
Standalone experience	The immersive Security Copilot portal experience accessed directly at https://securitycopilot.microsoft.com.
Tenant	The organizational boundary in Microsoft Entra ID that isolates identity, access, and data for Security Copilot. All workspaces, users, and interactions operate within this boundary and are governed by tenant-level permissions and security controls.

Key features or capabilities

The key features and capabilities in the following table describe what Microsoft Security Copilot is designed to do and how it performs across supported tasks.

Expand table

Feature or capability	Description
Incident investigation and response	Security Copilot helps security professionals triage and investigate incidents by generating summaries of complex security alerts, correlating signals across Microsoft Defender XDR, Microsoft Sentinel, and other integrated products, and providing step-by-step remediation guidance.
Threat intelligence	Security Copilot can search across Microsoft Defender threat intelligence articles and profiles, threat analytics reports, and vulnerability disclosure publications to surface relevant intelligence aligned to a prompt.
Security Alert Triage	Security Copilot includes the Security Alert Triage Agent to help security teams triage alerts at scale by identifying real attacks and false positives, enabling analysts to focus on investigating real threats.
Script analysis and KQL query generation	Security Copilot can analyze suspicious scripts or malware and translate natural language into KQL queries, enabling team members at all skill levels to perform advanced hunting and technical analysis tasks.
Security posture management	Security Copilot helps users understand prioritized risks across their environment and identify opportunities to improve posture through integration with Microsoft Defender XDR, Microsoft Entra, and Microsoft Intune.
Security policy creation and management	Users can define new policies, cross-reference them with existing policies for conflicts, and summarize policies in plain language to manage complex organizational context.
Promptbooks	Promptbooks are sequences of prompts that run in order to accomplish specific security tasks. Users can run promptbooks from a shared library or create and share their own.
Agents	Security Copilot supports agents that can automate and assist with security and IT operations tasks within the permissions granted by administrators. Agent actions are governed by configured identities, access controls, and triggers, and are designed to operate with human oversight as part of security workflows. Microsoft-built agents span the security product portfolio, covering SOC operations, threat hunting, threat intelligence, identity management, endpoint management, and data security. Administrators configure each agent's identity, permissions, and trigger. For details on specific agents and their use cases, see Intended uses.
File upload	Users can upload files directly into Security Copilot sessions, allowing Copilot to analyze, summarize, or cross-reference uploaded content when generating responses.

Feature or capability	Description
Multi-language support	Security Copilot supports prompting and responses in multiple languages. For more information, see Supported languages.

Intended uses

Microsoft Security Copilot is designed for use by security professionals and IT administrators who need AI-assisted support across a broad range of security and IT operations tasks. The assistive copilot experience helps users work more efficiently by surfacing relevant context, generating actionable guidance, and reducing the time spent on manual analysis. Some examples of intended use cases include:

- **Incident investigation and response**: Generates summaries of complex alerts, correlates signals across integrated security products such as Microsoft Defender XDR and Microsoft Sentinel, and provides step-by-step remediation guidance. Analysts can ask follow-up questions to progressively refine their investigation.
- **Threat intelligence gathering**: Surfaces relevant intelligence about threat actors, malware, and vulnerabilities based on natural language prompts, consolidating information into tailored summaries.
- **Script analysis and reverse engineering**: Explains suspicious scripts or command-line activity in plain language, identifies malicious behaviors, and highlights indicators of compromise.
- **KQL query generation**: Translates natural language prompts into ready-to-run KQL queries for advanced hunting and log analysis.
- **Security posture management**: Identifies prioritized risks and recommends actions to improve organizational security posture.
- **Security policy creation and management**: Assists in drafting, reviewing, and summarizing policies, and identifying gaps or conflicts.
- **Stakeholder reporting**: Generates reports tailored to different audiences, including technical summaries and executive briefings.
- **Promptbooks for repeatable security workflows**: Enables users to run or create multi-step workflows to standardize recurring security tasks.

Models and training data

Microsoft Security Copilot uses Azure OpenAI large language models (LLMs) from Foundry Models sold by Azure to power natural language experiences. These models aren't trained on Security Copilot Customer Data. Model capabilities vary in reasoning, speed, limitations, and supported scenarios.

Security Copilot also incorporates security-specific knowledge and context through plugins and grounding, which provide the LLM with relevant organizational data, threat intelligence, and authoritative content at inference time rather than through model training.

Performance

Security Copilot is designed to operate in enterprise security environments where large volumes of real-time security signals are generated across Microsoft Security products and other data sources configured by the organization.

Unlike general-purpose LLMs, Security Copilot provides:

- Real-time signal processing across structured security data
- Investigative reasoning that correlates multiple data sources
- Evidence-backed outputs grounded in customer data
- Continuous data collection for ongoing visibility

Users interact with Security Copilot using natural language prompts. The system processes these inputs through active plugins, data sources, and LLMs to generate responses.

Outputs are primarily textual including summary explanations, recommended actions, step-by-step guidance, code snippets (such as KQL queries), and formatted reports tailored to stakeholders. As the system generates a response, it displays intermediate steps in a process log, providing opportunities to double-check its processes and sources. Users can cancel, edit, rerun, or delete a prompt at any time, and responses can be pinned, shared, and exported to facilitate collaboration.

Limitations

Understanding the limitations of Microsoft Security Copilot is important to ensure it is used effectively and responsibly. While Security Copilot enhances security workflows, it isn’t designed for every scenario. Refer to the Microsoft Enterprise AI Services Code of Conduct as well as the following considerations when choosing a use case:

- **Accuracy and completeness**: Security Copilot can produce responses that are inaccurate, incomplete, or outdated. Output quality depends on available data sources, enabled integrations, and user-provided context. Users should apply human judgment and validate critical outputs.
- **Bias, stereotyping, and ungrounded content**: Despite safeguards, outputs may contain bias, stereotypes, or unsupported conclusions due to the probabilistic nature of large language models. Users should critically evaluate responses, especially in sensitive or high-impact scenarios.
- **Domain-specific scope**: Security Copilot is optimized for security-related tasks such as incident investigation and threat analysis. Prompts outside this domain may result in less accurate or less relevant responses.
- **Usage limits and latency**: Use of Security Copilot may be subject to capacity constraints and performance considerations. Generating responses, including executing integrations and applying safety checks, may introduce latency. Organizations should monitor their usage (SCUs) to maintain consistent performance.
- **Preview status**: Some Security Copilot capabilities may be in preview. Preview features should be treated as prerelease functionality, and outputs should be reviewed before taking action.
- **Prompt and context constraints**: Security Copilot operates within token and context limits inherent to large language models. Long prompts or extended interactions may exceed these limits, resulting in truncated or less optimal responses. Reframing or simplifying prompts may improve outcomes.
- **Dependence on data sources and configuration**: Responses are grounded in available data, including connected Microsoft services, third-party integrations, and user input. If relevant data sources aren’t available, enabled, or current, results may lack completeness or accuracy.
- **Script and code generation**: Security Copilot may generate code or include code in responses. Responses may appear valid but might not be semantically or syntactically correct, or might not accurately reflect the intent of the requester. Generated code shouldn't be deployed into production environments without appropriate validation, testing, and review procedures. Users must also verify that any parameters used by the generated code align with the original request. For example, if an agent operates on alerts within a specific time range, confirm that the time range in the generated code matches the time range specified in the natural language prompt.
- **Data access and permission boundaries**: Security Copilot operates within existing organizational permissions and data access controls. Responses are limited to the data a

user is authorized to access. Misconfigured permissions in underlying systems may affect the relevance or completeness of results.

- **Government cloud environments**: Security Copilot isn't currently designed for use in certain government cloud environments.

Evaluations

Performance and safety evaluations assess whether AI applications are operating reliably and securely by examining factors like groundedness, relevance, and coherence while identifying the risks of generating harmful content. The following evaluations were conducted with safety components already in place, which are also described in Safety Components and Mitigations.

Evaluation data for quality and safety

Our evaluation data is custom-built to assess AI application performance across key areas of safety and quality, simulating real-world scenarios, and risks. We begin by identifying relevant evaluation aspects of concern based on multi-disciplinary research and expert input. These concerns are translated into targeted evaluation objectives and guide formulation of evaluation metrics.

For safety, we create adversarial prompts to elicit undesirable or edge-case responses, which are then scored using AI-assisted annotators trained to assess alignment with Microsoft's safety standards. For quality, we craft rubric-based prompts relevant to scenarios including evaluating retrieval-augmented generation (RAG) applications and agents.

Datasets are curated from diverse sources including synthetic and public datasets to simulate real-world user scenarios. Using the curated datasets, both evaluations undergo iterative refinement and human alignment to improve metric efficacy and reliability. This methodology forms the foundation of repeatable, rigorous assessments that reflect how customers use evaluations to build better and safer AI.

Custom evaluations

Custom evaluations were conducted to validate model performance across grounding, adversarial robustness, and harmful content scenarios using regression testing, curated prompt datasets, and production-aligned examples. The evaluation compared outputs between GPT models, using internal tools to assess **groundedness** and **Azure OpenAI content filtering** to validate protections against **jailbreak**, **prompt injection**, and **intellectual property violations**.

Results show consistent or improved performance, including strong protection rates across adversarial scenarios and improved grounding accuracy.

Harmful content handling remains consistent across models and operates in annotation mode to support security-focused use cases, with additional large-scale testing confirming high protection rates across categories. Regression tests are conducted to verify that the content, which is not harmful isn't getting classified as harmful.

User feedback is critical to improving Security Copilot. Users can provide feedback using: **Needs improvement**, **Inappropriate** or the **Looks right** options. This feedback goes directly to Microsoft and is used to improve the platform's performance through ongoing iterative refinement.

Safety components and mitigations

As we identified potential risks and misuse through processes like red team testing and measured them, we developed mitigations to reduce the potential for harm. In the following list, we describe some of those mitigations. We will continue to evaluate the Microsoft Security Copilot experience to improve product performance and mitigations.

- **Harmful content filtering and guardrails**: Security Copilot integrates Microsoft-developed guardrails (content filters) and abuse detection models as part of the Azure OpenAI Service foundation. These neural classification models detect and filter harmful content across categories including hate, sexual, violence, and self-harm at multiple severity levels. Optional classification models also detect jailbreak risks, known text or code material, and indirect prompt injection attacks. These layered controls help prevent the AI from producing responses that violate Microsoft's safety standards.

- **Safety system design**: Microsoft developed a safety system for Security Copilot that is designed to mitigate failures and prevent misuse, including harmful content annotation, operational monitoring, and other safeguards. The Azure OpenAI Service Responsible AI Mitigation Requirements do not apply to Security Copilot customers directly because Security Copilot implements these mitigations on the customer's behalf.

- **User feedback loop**: The platform provides built-in feedback mechanisms: **Needs improvement**, **Inappropriate**, and **Looks right** options that allow users to report problematic and useful outputs directly to Microsoft. This feedback drives continuous improvement cycles, enabling rapid fixes and model adjustments when errors or misalignments are discovered.

- **Data encryption and access protection**: Customer data handled by Security Copilot is encrypted both in transit and at rest as described in the Microsoft Products and Services Data Protection Addendum. By default, no human users have access to the database and network access is restricted to the private network where the Security Copilot application is deployed; if human access is needed (for incident response), elevated access and network access must be approved by authorized Microsoft employees. See Compliance.

- **Phased deployment approach**: Security Copilot releases features through an invite-only early access program, allowing Microsoft to gather feedback and refine features before broader availability.

Our approach to mapping, measuring, and managing risks will continue to evolve as we learn more, and we're already making improvements based on feedback received from customers.

Best practices for deploying and adopting Microsoft Security Copilot

Responsible AI is a shared commitment between Microsoft and its customers. While Microsoft builds AI systems with safety, fairness, and transparency at the core, customers play a critical role in deploying and using these technologies responsibly within their own contexts.

Security Copilot agents are designed to augment human expertise, not replace. Customers remain responsible for reviewing outputs, validating decisions, and ensuring compliance with applicable laws, regulations, and organizational policies.

Deployers and end users should:

- **Exercise caution and evaluate outcomes when using Security Copilot for consequential decisions or in sensitive domains**: Consequential decisions are those that may have a legal or significant impact on a person's access to employment, legal services, healthcare, or that could result in physical, psychological, or financial harm. Sensitive domains such as financial services, healthcare, and legal require particular care due to the potential for disproportionate impact on different groups of people. When using AI for decisions in these areas, customers should ensure that impacted stakeholders can understand how decisions are made, appeal decisions, and update any relevant input data.

- **Evaluate legal and regulatory considerations**: Customers need to evaluate potential specific legal and regulatory obligations when using any AI services and solutions, which may not be appropriate for use in every industry or scenario. Additionally, AI services or

solutions are not designed for and may not be used in ways prohibited in applicable terms of service and relevant codes of conduct.

- **Enable and maintain relevant plugins**: The quality and accuracy of Security Copilot responses depend significantly on the plugins that are enabled. Administrators should ensure that appropriate Microsoft and third-party plugins are configured and maintained so that users receive grounded, contextually relevant responses.

End users should:

- **Write effective prompts**: Writing clear, specific prompts is key to getting better outcomes with Security Copilot. Include relevant context such as incident IDs, asset names, or time ranges. Iterate and regenerate prompts as needed, and always review and verify AI-generated responses. For more information, see Prompting tips for Security Copilot.

- **Exercise human oversight when appropriate**: Human oversight is an important safeguard when interacting with AI systems. While we continuously improve Security Copilot, AI systems may make mistakes. The output generated may be inaccurate, incomplete, biased, or not fully aligned with your intended goals due to ambiguity in inputs or limitations of the underlying models. Users should review the responses generated by Security Copilot and verify that they match their expectations and requirements before taking action.

- **Be aware of the risk of overreliance**: Overreliance on AI occurs when users accept incorrect or incomplete AI outputs, mainly because mistakes in AI outputs may be hard to detect. For security professionals, overreliance could result in missed threats, incorrect incident conclusions, or policy changes based on flawed recommendations. Security Copilot includes AI disclosure and cites source materials to help mitigate this risk, but users should still make sure to verify the accuracy of responses. Users can review the agent node map that provides a high-level view of the steps performed during an agent's workflow.

- **Exercise caution when deploying or designing agentic AI in sensitive domains**: Users must implement appropriate human oversight, when configuring and deploying agentic AI systems in domains where agent actions are irreversible or highly consequential. Additional precautions should be taken when creating autonomous agentic AI as described in the Microsoft Enterprise AI Services Code of Conduct.

Deployers should:

- **Configure RBAC and agent permissions carefully**: Administrators are responsible for configuring role-based access controls for both users and agents. Permissions should follow

the principle of least privilege. Agents should only be granted access to the data and actions necessary for their designated task.

- **Monitor usage and review activity**: Administrators (owners) can use the Security Copilot usage monitoring dashboard to review session-level data such as usage over time, session initiators, and plugins used during sessions. This visibility helps organizations understand how Security Copilot is used across prompts, promptbooks, and agents. For more information, see Manage usage.

- **Manage data sharing settings**: Owners can configure Customer Data sharing preferences at any time and must review and update these settings in accordance with their organization's privacy and compliance requirements. For more information, see Privacy and data security in Microsoft Security Copilot.

- **Educate users on capabilities and limitations**: Effective and responsible use of Security Copilot requires users to understand what the system can and cannot do. Deployers should provide training and guidance to help users interact with Security Copilot effectively, including the importance of verifying AI-generated outputs before taking action.

Learn more about Security Copilot

For additional guidance on the responsible use of Microsoft Security Copilot, see the following documentation:

- What is Microsoft Security Copilot?
- Privacy and data security in Microsoft Security Copilot

Learn more about responsible AI

- Microsoft AI principles
- Microsoft responsible AI resources
- Microsoft Azure Learning courses on responsible AI
- AI Literacy

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 06/22/2026

Application card: Microsoft Security Copilot agents

What is an Application or Platform card?

Microsoft's Application and Platform cards are intended to help you understand how our AI technology works, the choices application owners can make that influence application performance and behavior, and the importance of considering the whole application, including the technology, the people, and the environment. Application cards are created for AI applications and platform cards are created for AI platform services. These resources can support the development or deployment of your own applications and can be shared with users or stakeholders impacted by them.

As part of its commitment to responsible AI, Microsoft adheres to six core principles: fairness, reliability and safety, privacy and security, inclusiveness, transparency, and accountability. These principles are embedded in the Responsible AI Standard, which guides teams in designing, building, and testing AI applications. Application and Platform cards play a key role in operationalizing these principles by offering transparency around capabilities, intended uses, and limitations. For further insight, readers are encouraged to explore Microsoft's Responsible AI Transparency Report and the Microsoft Enterprise AI Services Code of Conduct, which outlines how to engage with AI responsibly.

Overview

Microsoft Security Copilot is a generative AI-powered security solution that helps increase the efficiency and capabilities of defenders to improve security outcomes at machine speed and scale. Security Copilot provides a natural language, assistive copilot experience that helps security professionals and IT administrators handle a wide range of end-to-end scenarios, including incident response, threat hunting, intelligence gathering, and posture management.

Security Copilot is designed with integration in mind and offers an immersive standalone experience at https://securitycopilot.microsoft.com. The platform seamlessly integrates with products in the Microsoft security portfolio, such as Microsoft Defender XDR, Microsoft Sentinel, Microsoft Intune, Microsoft Entra, and Microsoft Purview and supported third-party services.

A Security Copilot agent generates outputs and can perform actions based on configured logic, permissions, and triggers defined by the customer. Customers can use autonomous Security Copilot agents to automate security workflows, accelerate response times, prioritize risks, and reduce manual workloads, all while retaining full control. Agent autonomy refers to an agent's ability to take actions independently within defined boundaries.

An administrator (Security Copilot owner) discovers and deploys agents through the Security Copilot portal and Security Store, with some agents also surfaced in embedded Microsoft security product experiences. The administrator sets the agent's identity and configures role-based access control (RBAC) for the agent.

The intended users include SOC analysts, IT administrators, data security and identity administrators, compliance analysts, and security leaders such as Chief Information Security Officers (CISOs).

Security Copilot holds ISO 42001 certification, which confirms that an independent third party reviewed Microsoft's application of the necessary framework and capabilities to effectively manage risks and opportunities associated with the continuous development, deployment, and operation of Microsoft AI systems.

For more information, see What is Microsoft Security Copilot?, Microsoft Security Copilot experiences, and Application card for Security Copilot.

Key terms

The following table provides a glossary of key terms related to Microsoft Security Copilot agents.

Expand table

Term	Definition
Agent	A Security Copilot agent processes signals from the customer environment through integrated data sources and plugins, analyzes over data, and generates recommendations. Agents may also perform scoped actions within configured permissions when prompted, which require an appropriate user or administrator approval. Agents can range from simple prompt-and-response experiences to more automated, semi-autonomous workflows with human oversight.
Agent action	An operation executed by an agent, such as retrieving data, generating outputs, or modifying configurations, based on permissions.
Agent manifest	A configuration file that defines an agent's capabilities, tools, and behavior. It controls how the agent operates within the Security Copilot and is represented in a YAML file format.

Term	Definition
Embedded experience	Accessing Security Copilot capabilities from within another Microsoft security product, such as Microsoft Defender XDR or Microsoft Sentinel. The Security Copilot sidecar panel surfaces AI assistance directly in the context of that product.
Grounding	The process of providing contextual input sources to the large language model related to a user's prompt. By enabling Security Copilot to access organizational data through plugins and Microsoft security products, Security Copilot can deliver more accurate and contextually relevant responses.
Large language model (LLM)	AI models trained on large amounts of text data to predict words in sequences. LLMs are capable of performing various tasks such as text generation, summarization, translation, classification, and more.
Partner-built agent	An agent published by Microsoft partners and made available through the Security Store to address specific security use cases.
Plugin	A collection of related tools that extends Security Copilot's capabilities by giving it access to resources from Microsoft and non-Microsoft services and public websites through APIs. Plugins add more context to the responses and outputs that Security Copilot generates.
Post-processing	The set of actions Security Copilot performs to refine and prepare the LLM response before returning it to the user. This post-processing includes additional grounding calls through plugins, responsible AI checks, security, compliance, and privacy checks.
Prompt	The natural language text a user sends to Security Copilot to execute a specific task or obtain information. For example, *Summarize this incident and suggest remediation steps.*
Promptbook	A series of prompts that run in sequence, building on previous responses, to accomplish specific security-related tasks. Promptbooks can be used from the library or built and shared by users.
Red team testing	Techniques used by experts to assess the limitations and vulnerabilities of a system and to test the effectiveness of planned mitigations. Red team testing is used to identify potential risks and is distinct from systematic measurement of risks.
Responsible AI	Microsoft's policy, research, and engineering practices that are grounded in its AI principles and operationalized through the Responsible AI standard. For more information, see the Fluent RAI guidance.
Security Compute Unit (SCU)	SCUs are the units of compute capacity used to run Security Copilot workloads and deliver consistent performance across its experiences. Security Copilot capacity is measured in SCUs and can be consumed through provisioned or overage capacity models. For more information, see Understand SCUs.
Security Operations Center (SOC)	A dedicated security team or facility focused on continuous monitoring, analysis, and response to cybersecurity incidents within an organization. SOC analysts are among the primary intended users of Security Copilot.
Standalone	The immersive Security Copilot portal experience accessed directly at

Term	Definition
experience	https://securitycopilot.microsoft.com .
Tenant	The organizational boundary in Microsoft Entra ID that isolates identity, access, and data for Security Copilot. The tenant governs all workspaces, users, and interactions through tenant-level permissions and security controls.
Trigger	An event or condition that tells an agent to start executing its workflow. Triggers can be time-based (for example, a weekly schedule for the Threat Intelligence Briefing Agent) or manual (run on demand by a user or administrator). The trigger is configured during agent setup.

Key features or capabilities

The key features and capabilities in the following table describe what Microsoft Security Copilot is designed to do and how it performs across supported tasks.

Expand table

Feature or capability	Description
Incident investigation and response	Security Copilot helps security professionals triage and investigate incidents by generating summaries of complex security alerts, correlating signals across Microsoft Defender XDR, Microsoft Sentinel, and other integrated products, and providing step-by-step remediation guidance.
Threat intelligence	Security Copilot can search across Microsoft Defender threat intelligence articles and profiles, threat analytics reports, and vulnerability disclosure publications to surface relevant intelligence aligned to a prompt.
Security Alert Triage	Security Copilot includes the Security Alert Triage Agent to help security teams triage alerts at scale by identifying real attacks and false positives, enabling analysts to focus on investigating real threats with transparent, step-by-step reasoning to support every decision.
Script analysis and KQL query generation	Security Copilot can analyze suspicious scripts or malware and translate natural language into KQL queries, enabling team members at all skill levels to perform advanced hunting and technical analysis tasks.
Security posture management	Security Copilot helps users understand prioritized risks across their environment and identify opportunities to improve posture through integration with Microsoft Defender XDR, Microsoft Entra, and Microsoft Intune.
Security policy creation and management	Users can define new policies, cross-reference them with existing policies for conflicts, and summarize policies in plain language to manage complex organizational context.

Feature or capability	Description
Stakeholder reporting	Security Copilot can generate reports that summarize context, open issues, and protective measures tailored to the intended audience, such as executives or security teams.
Promptbooks	Promptbooks are sequences of prompts that run in order to accomplish specific security tasks. Users can run promptbooks from a shared library or create and share their own.
Agents	Security Copilot supports agents that automate and assist with security and IT operations tasks within the permissions granted by administrators. An agent takes actions through configured identities, access controls, and triggers, and operate with human oversight as part of security workflows. Microsoft-built agents span the security product portfolio, covering SOC operations, threat hunting, threat intelligence, identity management, endpoint management, and data security. Administrators configure each agent's identity, permissions, and trigger during setup. Users can review an agent's triggers, data access, identity, and action permissions (such as read or write) to understand how the agent operates within its defined scope. For details on specific agents and their use cases, see Intended uses.
Multi-language support	Security Copilot supports prompting and responses in multiple languages. For more information, see Supported languages.

Security Copilot is an AI-powered security solution that operates both assistively and as an autonomous agentic system. To understand agent autonomy, consider:

- **Activation triggers** – what conditions or user actions cause the agent to run
- **Access permissions** – what data, systems, or resources the agent can use
- **Action rights** – what actions the agent is authorized to take on its own

The following sections describe the core agentic capabilities that underpin how agents reason, plan, remember, adapt, and extend their reach.

Reasoning

Security Copilot agents use the underlying large language model to analyze available context, evaluate signals, and determine the most appropriate course of action. For example, the Phishing Triage Agent assesses email content, sender reputation, and behavioral signals to produce a classification verdict with a natural language rationale. Agents surface their reasoning transparently so that analysts can review, validate, or override conclusions before acting on them.

Planning

Agents operate against defined triggers that instruct the agentic system to initiate a structured sequence of actions toward a goal. Agents can be configured to:

- Run automatically on a schedule (for example, the Threat Intelligence Briefing Agent runs every seven days).
- Run manually on demand when needed.

This design gives agents a goal-directed execution model where the system analyzes when and how to act to complete its task.

Memory

Security Copilot agents can retain information over time, referred to as memory. Memory allows an agent to incorporate past inputs into future behavior, depending on how the agent is designed and configured.

Memory can include feedback provided by users. Agents can use this feedback to adjust their responses or actions in subsequent interactions.

Adaptability

Security Copilot agents are designed to adapt based on user feedback and operational context, while continuing to operate within the scope defined by their configured identity, permissions, and triggers.

- **Feedback loop**: Security Copilot Owners and Contributors can provide feedback on an agent's responses. This feedback is stored in the agent's memory and can influence future outputs, depending on the agent design and configuration.
- **Contextual grounding**: During prompt processing, Security Copilot enriches prompts using grounding. This process incorporates relevant organizational data, enabled plugins, and threat intelligence so that responses reflect the current context.
- **Configurable identity and permissions**: Agents can be updated after setup to modify identity, triggers, and parameters. This enables agents to align with evolving workflows and requirements.

Extensibility

- **Plugins**: Agents use plugins to reach external services via APIs, including reputation lookups, threat intelligence, and endpoint data. Both Microsoft-built and partner-built plugins are supported. For more information, see Plugins overview.

- **Connectors**: Logic Apps and Copilot Studio connectors wrap the Security Copilot API, enabling developers and users to call into the platform from external automation workflows. For more information, see Connectors overview.
- **Custom agents**: Developers can build custom agents tailored to specific use cases and add them to the Security Copilot ecosystem via the developer platform. For more information, see Develop custom agents.
- **Security Store**: Prebuilt Microsoft and partner agents can be discovered and deployed from the built-in agent library and Security Store.
- **Embedded experiences**: Agents operate not only in the Security Copilot standalone portal but are also embedded across the broader Microsoft security ecosystem, including Microsoft Defender XDR, Microsoft Sentinel, Microsoft Intune, Microsoft Entra, and Microsoft Purview.

Intended uses

Microsoft Security Copilot agents are designed for security professionals and IT administrators to support security workflows, such as intelligence gathering and correlation. Agents can operate within administrator-defined permissions, and help streamline tasks that might otherwise require significant manual effort, while keeping humans in control. Some examples of intended agent use cases include:

- **Threat intelligence briefing**: The **Threat Intelligence Briefing Agent** generates timely, detailed threat intelligence reports by correlating Microsoft Defender Threat Intelligence data, Defender External Attack Surface Management (EASM) signals, and real-time customer context. Security analysts can use this agent to replace hours or days of manual intelligence gathering and correlation with a report generated in minutes.

- **Security data analysis (Microsoft Defender XDR)**: The **Security Analyst Agent** helps security analysts quickly identify, assess, and prioritize risks across large volumes of security data. The agent performs both basic analysis tasks such as pattern analysis, trend analysis, and visualization, and advanced tasks such as anomaly detection, clustering, risk scoring, and predictive modeling. It integrates data from Microsoft Defender XDR, Microsoft Sentinel Log Analytics, Microsoft Sentinel Data Lake, and uploaded CSV files, generating prioritized insights with a full evidence trail in an interactive, code-free chat experience.

- **Security operations and incident response (Microsoft Defender XDR)**:
  - The **Phishing Triage Agent** evaluates user-reported phishing emails as they're submitted, classifies verdicts with transparent reasoning, and incorporates analyst feedback over

time. The **Security Alert Triage Agent (Preview)** extends the same triage approach beyond phishing to a subset of identity alerts and cloud container alerts.
    - The **Dynamic Threat Detection Agent** runs continuously in the background to uncover hidden threats and gaps across Defender and Microsoft Sentinel environments by correlating alerts, events, anomalies, and threat intelligence.

- **Threat hunting (Microsoft Defender XDR)**: The **Threat Hunting Assistant** enables end-to-end threat hunting using natural language. It generates KQL queries, interprets results, surfaces insights, and guides analysts through full hunting sessions to find threats faster and with greater confidence.

- **Identity and access management (Microsoft Entra)**:
    - The **Conditional Access Optimization Agent** analyzes conditional access policies and recommends improvements based on Microsoft best practices and Zero Trust principles.
    - The **Identity Risk Management Agent** helps identity administrators investigate potential risks and take action to protect critical assets.

- **Endpoint management (Microsoft Intune)**:
    - The **Vulnerability Remediation Agent** uses Defender data to monitor vulnerabilities and prioritize remediation with AI-driven risk assessments.
    - The **Policy Configuration Agent** lets administrators import documents or write plain-language instructions to find matching settings in the Intune settings catalog and create policies.
    - The **Change Review Agent** evaluates the effect of approval requests in Intune and makes recommendations for actions administrators can take.

- **Data security investigation (Microsoft Purview)**:
    - The **Triage Agent in Insider Risk Management** evaluates alerts based on user, file, and activity risk, automatically sorting them into prioritized categories to help security teams focus on the highest-risk cases.
    - The **Purview Triage Agent in Data Loss Prevention** evaluates DLP alerts based on sensitivity risk, exfiltration risk, and policy risk, helping data security administrators act on the most critical incidents.

- **Custom agentic workflows**: Developers can build and deploy custom agents using natural language processing, uploading the agent manifest, or using the MCP tools tailored to their organization's specific security use cases, extending agent capabilities through plugins and connectors to Microsoft and non-Microsoft services. For more information, see Custom agents.

You can also build an interactive chat agent when agents and users need to collaborate for a guided experience to solve something. For more information, see Interactive Agent.

Models and training data

Microsoft Security Copilot uses Azure OpenAI large language models (LLMs) from Foundry Models sold by Azure to power natural language experiences. These models aren't trained on Security Copilot Customer Data. Model capabilities vary in reasoning, speed, limitations, and supported scenarios.

Security Copilot also incorporates security-specific knowledge and context through plugins and grounding, which provide the LLM with relevant organizational data, threat intelligence, and authoritative content at inference time rather than through model training.

Performance

Security Copilot is designed to operate in enterprise security environments where large volumes of real-time security signals are generated across Microsoft Security products and other data sources configured by the organization. It integrates signals from these sources, combining real-time processing of structured log data with investigative reasoning. This is designed to help detect, analyze, and trace the source and impact of security incidents across multiple data sources.

Security Copilot runs on Microsoft's hyperscale infrastructure and a security-specific orchestration layer, which are intended to support scalable, resilient performance across repeated threat detection and investigation scenarios.

Agents extend Security Copilot's ability to act on security signals by executing defined security tasks through administrator-configured permissions. Agents reason over available signals, plan and execute structured actions, and produce outputs such as triage decisions, intelligence reports, and remediation guidance grounded in real-time organizational and threat intelligence data. Agent execution is transparent through step-by-step node maps, allowing analysts to review generated steps (not an in-depth summary of the specific actions taken at each node) and validate outcomes across runs.

Performance reliability is further supported through preconfigured Microsoft-built agents, which are tested against representative workflows, and custom agents, where a manifest YAML file can help promote more consistent behavior by specifying capabilities, tools, triggers, and operational boundaries. When manifest instructions are clearly defined, the agent may be better able to

interpret its task, select appropriate actions, and operate within its intended scope. For more information, see Custom agents.

For interactive agents, performance is shaped primarily through the iterative exchange between the agent and the user. The agent produces an initial response, and the user refines context through follow-up prompts based on what the agent surfaces. This back-and-forth interaction allows the user to incrementally guide the agent toward more accurate and targeted outcomes than a single autonomous run would produce. For more information, see Interactive agents.

Limitations

Understanding Security Copilot's limitations are important to ensure it's used within safe and effective boundaries. While customers are encouraged to use Security Copilot in their security workflows, it's important to note that Security Copilot wasn't designed for every possible scenario. Refer to the Microsoft Enterprise AI Services Code of Conduct as well as the following considerations when choosing a use case:

- **Public preview status**: Some agents are in public preview and may be substantially modified before general availability. Microsoft makes no warranties, express or implied, with respect to these capabilities. As with any AI output, customers must review the decision making of the agent prior to acting upon its outputs.
- **Accuracy and completeness**: Like any AI-powered technology, Security Copilot doesn't get everything right. Responses may be inaccurate, incomplete, or outdated, particularly if the relevant plugins aren't enabled or if the most current data isn't available through user input or organizational context. Users should always exercise human judgment and verify critical outputs rather than relying solely on AI-generated response.
- **Domain-specific scope**: Security Copilot is designed to respond to prompts related to the security domain, such as incident investigation and threat intelligence. Prompts outside the scope of security may result in responses that lack accuracy and comprehensiveness.
- **Script and code generation**: Security Copilot may generate code or include code in responses. Responses may appear valid but might not be semantically or syntactically correct, or might not accurately reflect the intent of the requester. Generated code shouldn't be deployed into production environments without appropriate validation, testing, and review procedures. Users must also verify that any parameters used by the generated code align with the original request. For example, if an agent operates on alerts within a specific time range, confirm that the time range in the generated code matches the time range specified in the natural language prompt.

- **Prompt-length constraints**: The system may not be able to process long prompts, such as those containing hundreds of thousands of characters. The underlying LLM has a token limit, and overly verbose queries or extended sessions can overflow the token space. When this scenario happens, Security Copilot attempts to apply mitigations to ensure an output is always available, even if the content isn't optimal; however, those mitigations are not always effective, and it might be necessary to try a different prompt or plugin.
- **Usage limits and latency**: Use of the platform may be subject to usage limits or capacity throttling. Generating responses, including making API calls through plugins and checking responses before displaying them, can take time and require high GPU capacity. Organizations should monitor their SCU consumption and adjust provisioned capacity as needed to avoid unexpected service interruptions.
- **Bias, stereotyping, and ungrounded content**: Despite the implementation of responsible AI controls on both user prompts and LLM outputs, AI services are fallible and probabilistic. This makes it challenging to comprehensively block all inappropriate content, which may lead to potential biases, stereotypes, or ungrounded content in AI-generated output.
- **Government and Microsoft Sovereign Public/Private Cloud**: Security Copilot isn't supported in these environments at this time.
- **Task-specific agent boundaries**: The agents are suitable only for the specific task they're designed to perform. For more information, see the intended use cases section. They aren't suitable for any other task and shouldn't be repurposed beyond their defined scope.
- **Node map detail**: The agent node map provides a high-level view of the steps performed during an agent workflow. Each node represents a step in the process and displays the title of the skill used, along with basic metadata such as completion status, duration, and timestamp. The node map is designed to show the sequence of actions, but it doesn't provide detailed information about the specific operations or decisions made within each step. Because the node map presents only summarized information, it may not fully capture the context or complexity of each action.
- **Agent feedback and memory transparency**: When users submit feedback to an agent for storage in memory, the agent doesn't provide a summary of its interpretation of the feedback. To increase accuracy of future outputs, users should provide feedback that's clear, concise, and specific. Feedback submitted to an agent may be stored and used to influence future outputs. However, visibility into this stored feedback can vary depending on the agent experience and user role. For more information, see Provide feedback.

Evaluations

Performance and safety evaluations assess whether AI applications are operating reliably and securely by examining factors like groundedness, relevance, and coherence while identifying the risks of generating harmful content. The following evaluations were conducted with safety components already in place, which are also described in Safety Components and Mitigations.

Evaluation data for quality and safety

Our evaluation data is custom-built to assess AI application performance across key areas of safety and quality, simulating real-world scenarios, and risks. We begin by identifying relevant evaluation aspects of concern based on multi-disciplinary research and expert input. These concerns are translated into targeted evaluation objectives and guide formulation of evaluation metrics.

For safety, we create adversarial prompts to elicit undesirable or edge-case responses, which are then scored using AI-assisted annotators trained to assess alignment with Microsoft's safety standards. For quality, we craft rubric-based prompts relevant to scenarios including evaluating retrieval-augmented generation (RAG) applications and agents.

Datasets are curated from diverse sources including synthetic and public datasets to simulate real-world user scenarios. Using the curated datasets, both evaluations undergo iterative refinement and human alignment to improve metric efficacy and reliability. This methodology forms the foundation of repeatable, rigorous assessments that reflect how customers use evaluations to build better and safer AI.

Custom evaluations

Custom evaluations were conducted to validate model performance across grounding, adversarial robustness, and harmful content scenarios using regression testing, curated prompt datasets, and production-aligned examples. The evaluation compared outputs between GPT models, using internal tools to assess **groundedness** and **Azure OpenAI content filtering** to validate protections against **jailbreak**, **prompt injection**, and **intellectual property violations**. Results show consistent or improved performance, including strong protection rates across adversarial scenarios and improved grounding accuracy.

Harmful content handling remains consistent across models and operates in annotation mode to support security-focused use cases, with additional large-scale testing confirming high protection rates across categories. Regression tests are conducted to verify that the content, which isn't harmful isn't getting classified as harmful.

Security Copilot agents were evaluated by its product and research team with use cases and design inputs from customers. The security of the agent system was also assessed through a dedicated red teaming exercise. Microsoft also completed penetration testing on the Security Copilot service to validate protection against unauthorized access.

Now that Security Copilot is released, user feedback is critical in helping Microsoft improve the system. Users have the option of providing feedback on the agent response from Security Copilot. This feedback goes directly to Microsoft and is used to improve the platform's performance through ongoing iterative refinement. For more information, see Provide feedback.

Safety components and mitigations

As we identified potential risks and misuse through processes like red team testing and measured them, we developed mitigations to reduce the potential for harm. We'll continue to evaluate the Microsoft Security Copilot experience to improve product performance and mitigations. The following list describes some of those mitigations:

- **Harmful content filtering and guardrails**: Security Copilot integrates Microsoft-developed guardrails (content filters) and abuse detection models as part of the Azure OpenAI Service foundation. These neural classification models detect and filter harmful content across categories including hate, sexual, violence, and self-harm at multiple severity levels. Optional classification models also detect jailbreak risks, known text or code material, and indirect prompt injection attacks. These layered controls help prevent the AI from producing responses that violate Microsoft's safety standards.
- **Designed to minimize irreversible actions**: Microsoft-designed agent scenarios are intended to minimize irreversible actions and keep users in control of critical decisions. For custom agents, non-Microsoft designed agents the agent developers can modify agent behavior to ensure irreversible actions are minimized.
- **Safety system design**: Microsoft developed a safety system for Security Copilot that is designed to mitigate failures and prevent misuse, including harmful content annotation, operational monitoring, and other safeguards. The Azure OpenAI Service Responsible AI Mitigation Requirements don't apply to Security Copilot customers directly because Security Copilot implements these mitigations on the customer's behalf.
- **User feedback loop**: After an agent returns a response, users can provide feedback. Depending on the configuration, users may also submit additional written feedback to provide context on their experience. Submitted feedback is collected and used by Microsoft to improve product quality, identify issues, and prioritize enhancements to Security Copilot experiences.

- **Agent identity governance and Role-based access control (RBAC)**: Each Security Copilot agent runs under either a managed identity or a user account, enabling the administrator to govern the data it has access to. Each agent has RBAC controls, and the agents may be further restricted as to what data they process. By confining each agent's permissions, the system mitigates risks of unauthorized data exposure and ensures all automated actions are auditable and traceable.
- **Data encryption and access protection**: Customer data handled by Security Copilot is encrypted both in transit and at rest as described in the Microsoft Products and Services Data Protection Addendum. By default, no human users have access to the database and network access is restricted to the private network where the Security Copilot application is deployed; if human access is needed (for incident response), elevated access and network access must be approved by authorized Microsoft employees. See Compliance.
- **Phased deployment approach**: Security Copilot releases features through an invite-only early access program, allowing Microsoft to gather feedback and refine features before broader availability.

Our approach to mapping, measuring, and managing risks continue to evolve as we learn more, and make improvements based on feedback received from customers.

Best practices for deploying and adopting Microsoft Security Copilot

Responsible AI is a shared commitment between Microsoft and its customers. While Microsoft builds AI systems with safety, fairness, and transparency at the core, customers play a critical role in deploying and using these technologies responsibly within their own contexts.

Security Copilot agents are designed to augment human expertise, not replace. Customers remain responsible for reviewing outputs, validating decisions, and ensuring compliance with applicable laws, regulations, and organizational policies.

Deployers and end users should:

- **Exercise caution and evaluate outcomes when using Security Copilot for consequential decisions or in sensitive domains**: Consequential decisions are those that may have a legal or significant impact on a person's access to employment, legal services, healthcare, or that could result in physical, psychological, or financial harm. Sensitive domains such as financial services, healthcare, and legal require particular care due to the potential for disproportionate impact on different groups of people. When using AI for decisions in these

areas, customers should ensure that impacted stakeholders can understand how decisions are made, appeal decisions, and update any relevant input data.

- **Evaluate legal and regulatory considerations**: Customers need to evaluate potential specific legal and regulatory obligations when using any AI services and solutions, which may not be appropriate for use in every industry or scenario. Additionally, AI services or solutions are not designed for and may not be used in ways prohibited in applicable terms of service and relevant codes of conduct.

- **Enable and maintain relevant plugins**: The quality and accuracy of Security Copilot responses depend significantly on the plugins that are enabled. Administrators should ensure that appropriate Microsoft and third-party plugins are configured and maintained so that users receive grounded, contextually relevant responses.

End users should:

- **Write effective prompts**: Writing clear, specific prompts is key to getting better outcomes with Security Copilot. Include relevant context such as incident IDs, asset names, or time ranges. Iterate and regenerate prompts as needed, and always review and verify AI-generated responses. For more information, see Prompting tips for Security Copilot.

- **Exercise human oversight when appropriate**: Human oversight is an important safeguard when interacting with AI systems. While we continuously improve Security Copilot, AI systems may make mistakes. The output generated may be inaccurate, incomplete, biased, or not fully aligned with your intended goals due to ambiguity in inputs or limitations of the underlying models. Users should review the responses generated by Security Copilot and verify that they match their expectations and requirements before taking action.

- **Be aware of the risk of overreliance**: Overreliance on AI occurs when users accept incorrect or incomplete AI outputs, mainly because mistakes in AI outputs may be hard to detect. For security professionals, overreliance could result in missed threats, incorrect incident conclusions, or policy changes based on flawed recommendations. Security Copilot includes AI disclosure and cites source materials to help mitigate this risk, but users should still make sure to verify the accuracy of responses. Users can review the agent node map that provides a high-level view of the steps performed during an agent's workflow.

- **Exercise caution when deploying or designing agentic AI in sensitive domains**: Users must implement appropriate human oversight, when configuring and deploying agentic AI systems in domains where agent actions are irreversible or highly consequential. Additional

precautions should be taken when creating autonomous agentic AI as described in the Microsoft Enterprise AI Services Code of Conduct.

Deployers should:

- **Configure RBAC and agent permissions carefully**: Administrators are responsible for configuring role-based access controls for both users and agents. Permissions should follow the principle of least privilege. Agents should only be granted access to the data and actions necessary for their designated task.
- **Monitor usage and review activity**: Administrators (owners) can use the Security Copilot usage monitoring dashboard to review session-level data such as usage over time, session initiators, and plugins used during sessions. This visibility helps organizations understand how Security Copilot is used across prompts, promptbooks, and agents. For more information, see Manage usage.
- **Manage data sharing settings**: Owners can configure Customer Data sharing preferences at any time and must review and update these settings in accordance with their organization's privacy and compliance requirements. For more information, see Privacy and data security in Microsoft Security Copilot.
- **Educate users on capabilities and limitations**: Effective and responsible use of Security Copilot requires users to understand what the system can and cannot do. Deployers should provide training and guidance to help users interact with Security Copilot effectively, including the importance of verifying AI-generated outputs before taking action.

Learn more about Security Copilot agents

For more information on the responsible use of Microsoft Security Copilot, see the following documentation:

- What is Microsoft Security Copilot?
- Build custom agents
- Privacy and data security in Microsoft Security Copilot

Learn more about responsible AI

- Microsoft AI principles
- Microsoft responsible AI resources
- Microsoft Azure Learning courses on responsible AI

- AI Literacy

ⓘ**Note:** The author created this article with assistance from AI. Learn more

---

Last updated on 08/09/2026

Microsoft Security Copilot - Developer content

Learn how to create, publish, and test agents in Security Copilot using natural language, agent builder forms, YAMLs, and MCPs. Designed for developers to streamline security workflows.

About Security Copilot

OVERVIEW

What is Microsoft Security Copilot?

What are the Security Copilot experiences?

Custom agent overview

Plan

GET STARTED

Plan and design your agent

Understand use cases

Agent components

Cost considerations

Agent manifest

CONCEPT

Manifest YAML

Best practices

Changelog

Build agents and promptbooks

HOW-TO GUIDE

Build an agent using natural language

Build an agent using the agent builder form

Build an agent by uploading a YAML

Build an agent using a manifest

Build an agent using MCP

Build an interactive agent

Extend your agent with tools

Build promptbooks

Build plugins

HOW-TO GUIDE

API plugin

GPT plugin

KQL plugin

MCP plugin

Logic apps plugin

Samples collection

HOW-TO GUIDE

Add an API tool to the agent

Add a GPT tool to the agent

Build an agent using multiple tools

Test and publish your agent

HOW-TO GUIDE

Test your agent

Publish your agent

Reference

REFERENCE

Agent error messages

Plugin error codes

Application card for Security Copilot

Application card for Security Copilot agents